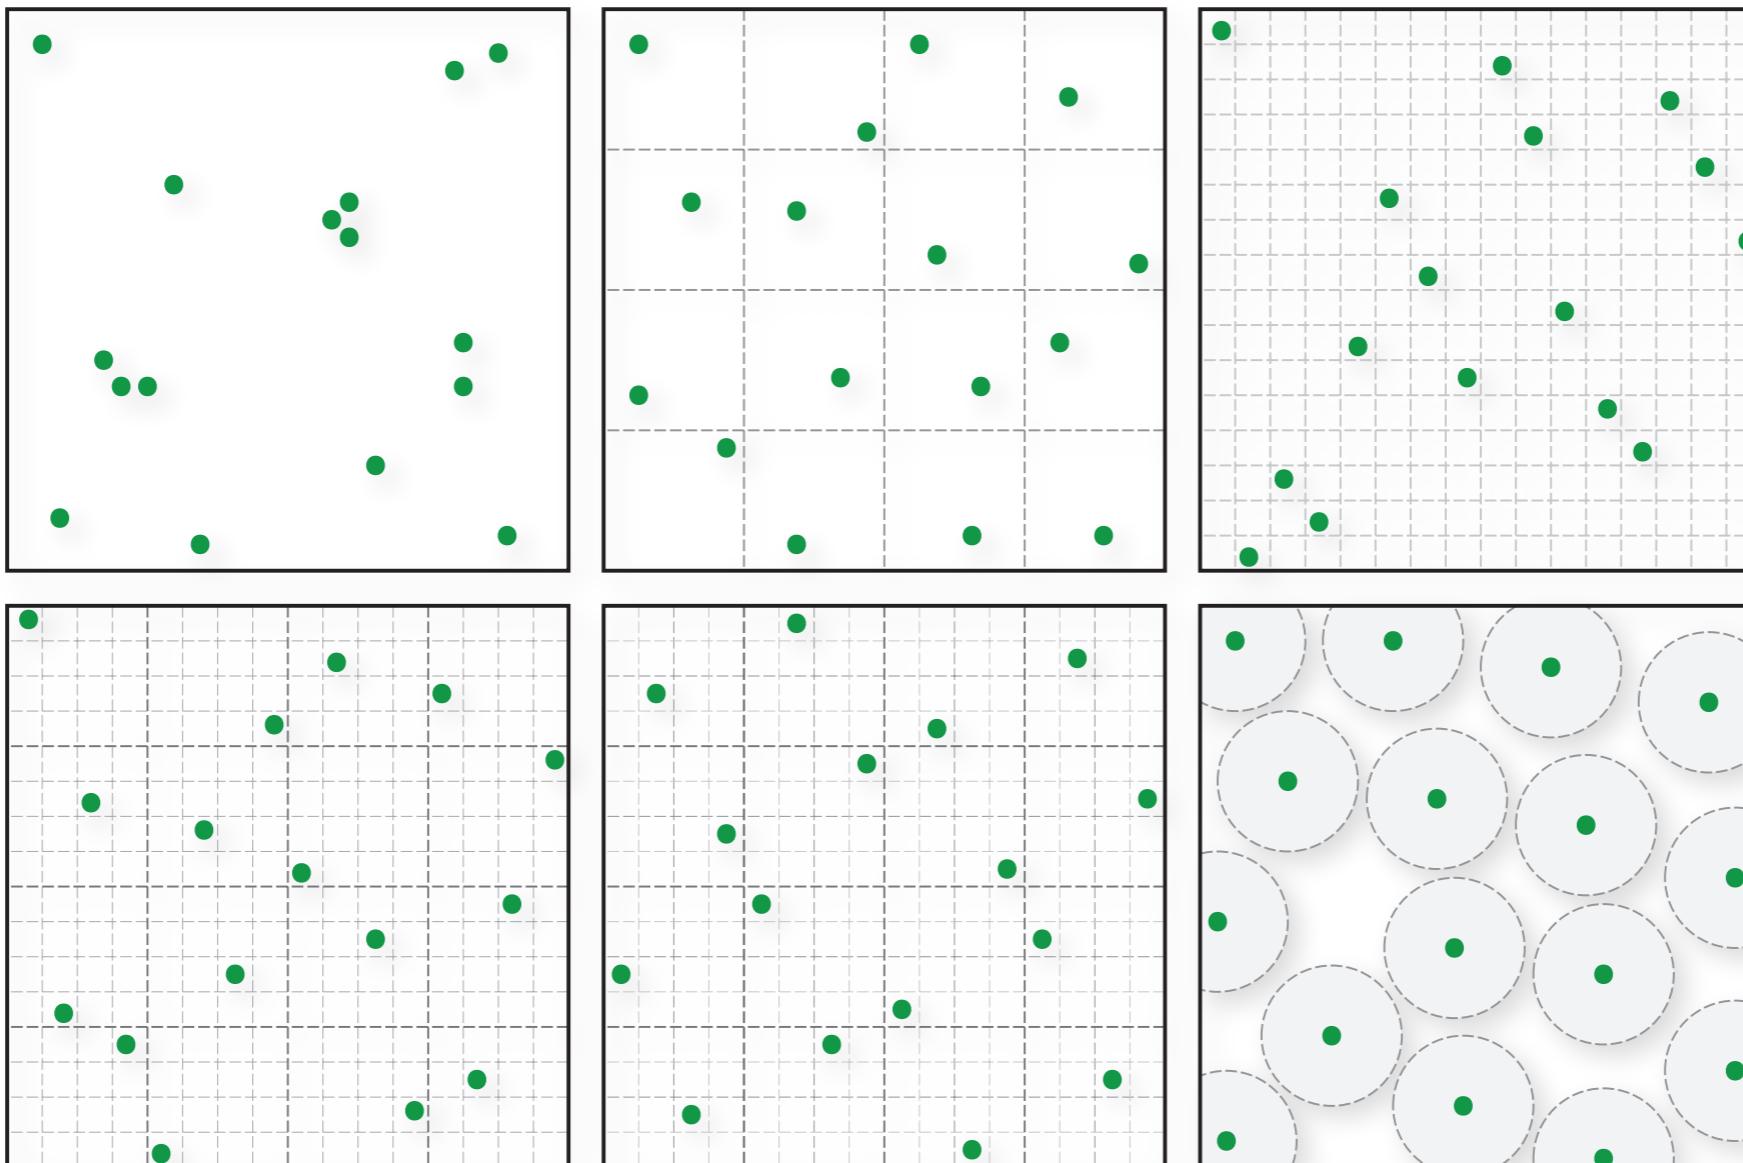


RENDERING ALGORITHMS

Monte Carlo II: Improved Sampling & Variance Reduction



Prof. Wojciech Jarosz
wojciech.k.jarosz@dartmouth.edu

Last Lecture

- Realistic lighting & realistic cameras

Environment Lighting



HENRIK WANN DENSEN - 2002

Environment Lighting



$$L_r(\mathbf{x}, \vec{\omega}_r) = \int_{\Omega} f_r(\vec{\omega}_i, \vec{\omega}_r) L_{\text{env}}(\vec{\omega}_i) V(\mathbf{x}, \vec{\omega}_i) \cos \theta_i d\vec{\omega}_i$$

Importance Sampling L_{env}



$$p(\vec{\omega}_i) \propto L_{\text{env}}(\vec{\omega}_i)$$

Marginal/Conditional CDF

- Assume the lat/long parameterization
- Draw samples from joint $p(\theta, \phi) \propto L_{\text{env}}(\theta, \phi) \sin \theta$
 - Step 1: create scalar version $L'(\theta, \phi)$ of $L_{\text{env}}(\theta, \phi) \sin \theta$

- Step 2: compute marginal PDF

$$p(\theta) = \int_0^{2\pi} L'(\theta, \phi) d\phi$$

- Step 3: compute conditional PDF

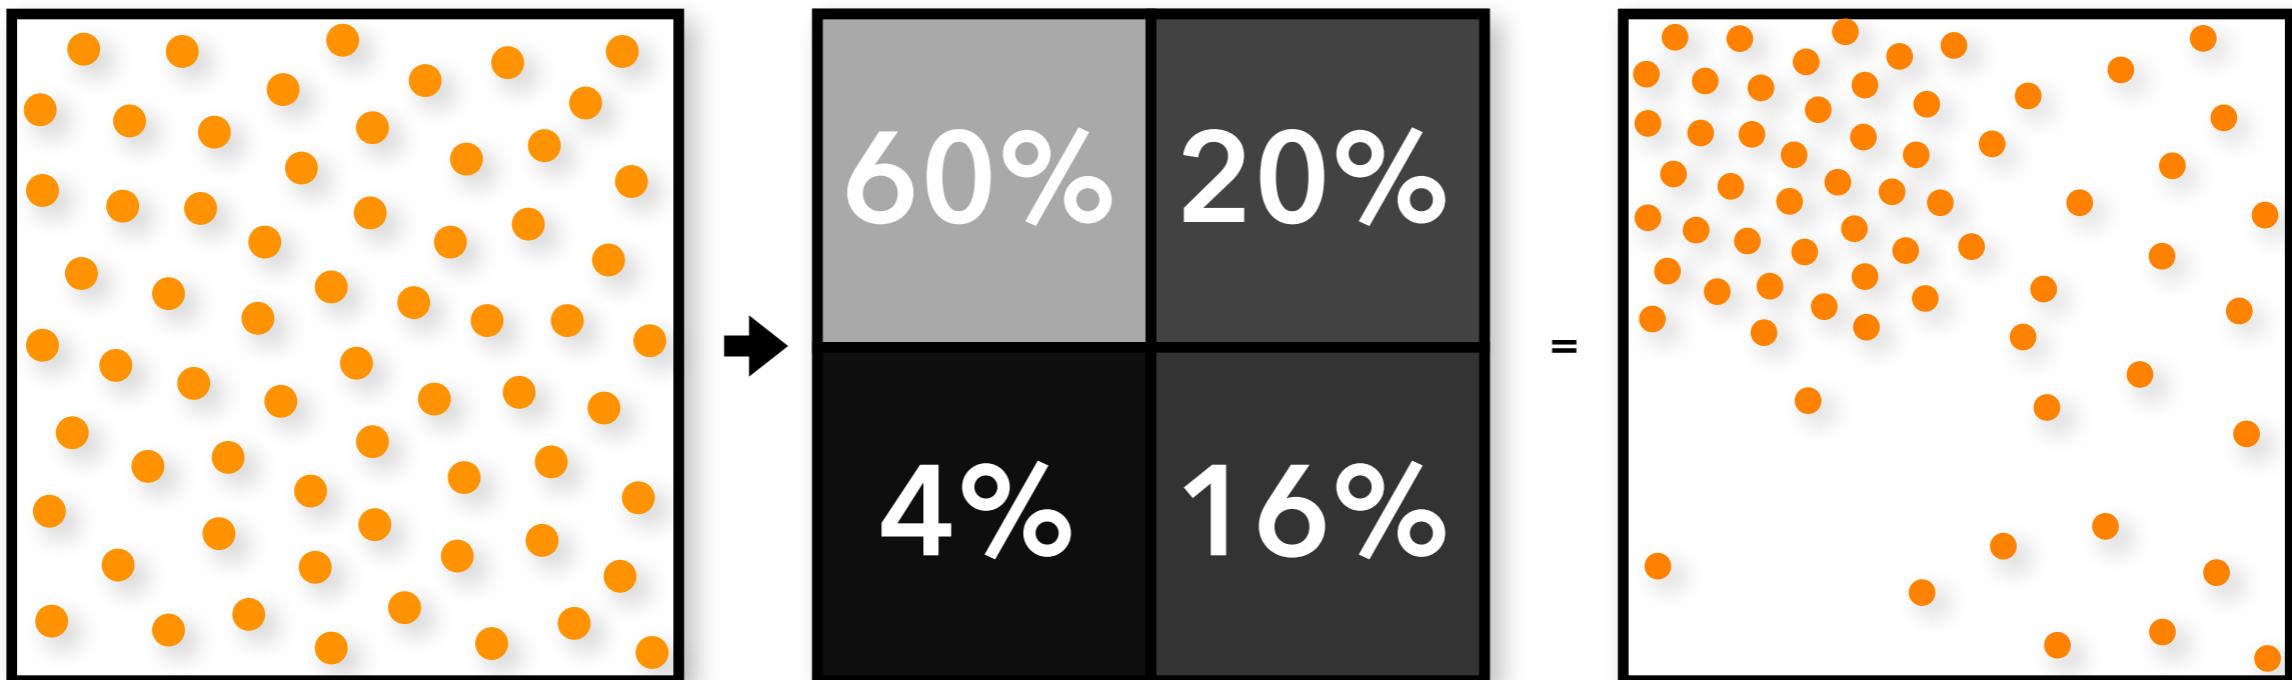
$$p(\phi|\theta) = \frac{p(\theta, \phi)}{p(\theta)}$$

- Step 4: draw samples $\theta_i \sim p(\theta)$ and $\phi_i = p(\phi|\theta)$

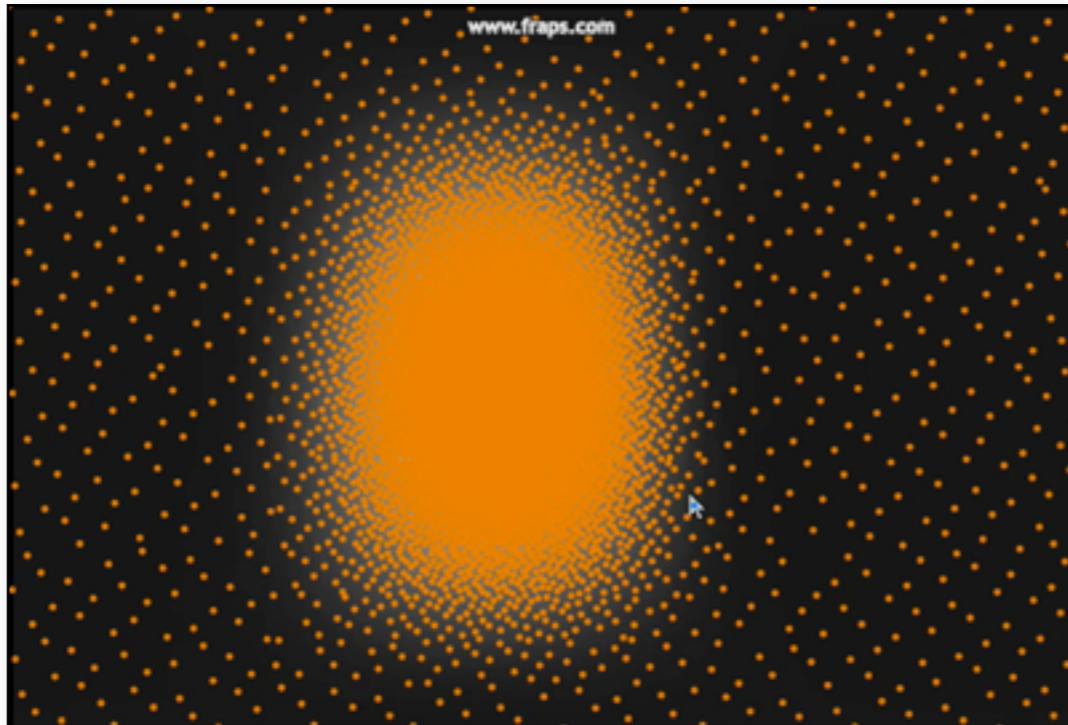
Hierarchical Sample Warping

Clarberg, Jarosz, Akenine-Möller, Jensen. "Wavelet Importance Sampling," 2005.

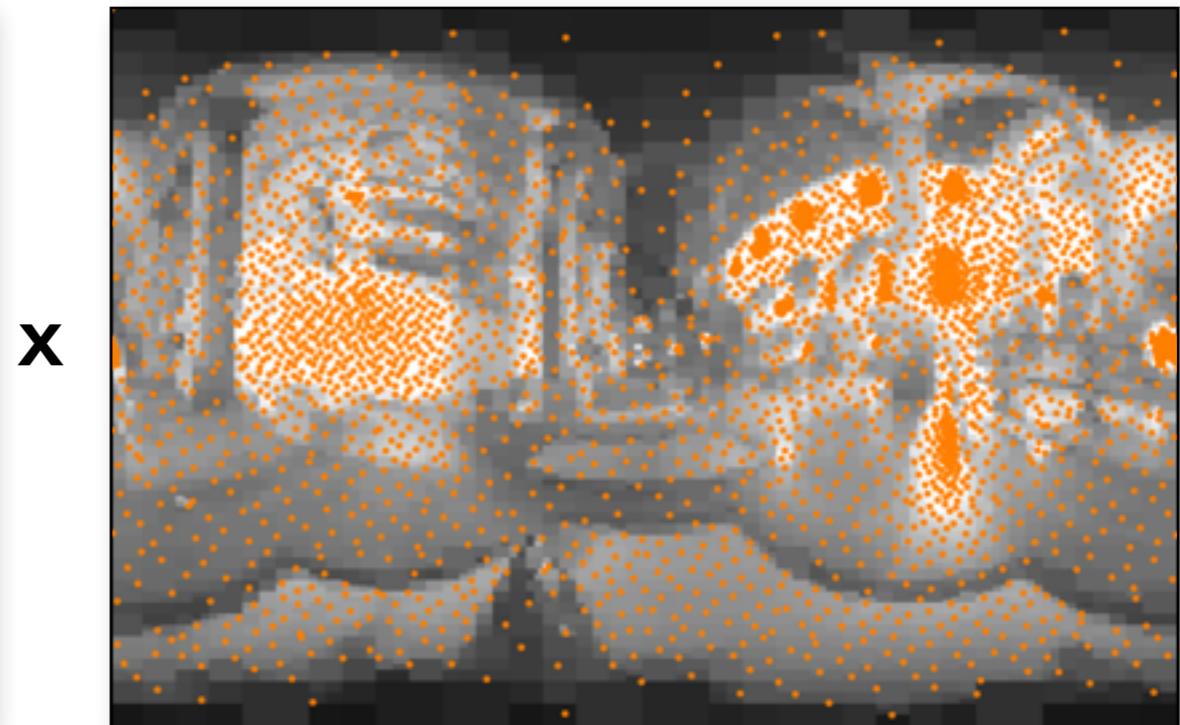
- Given:
 - input point set, and
 - hierarchical representation of density function (mip-map)
- Recursively warp point set to match the importance function



Product Sampling



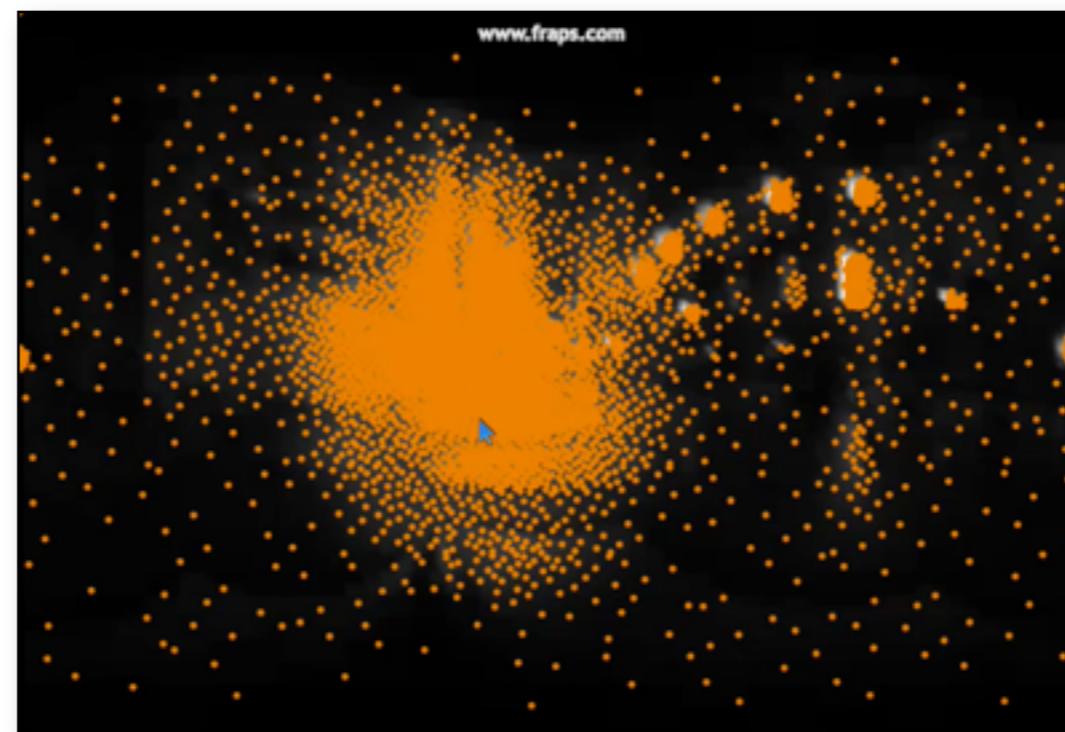
BRDF



x

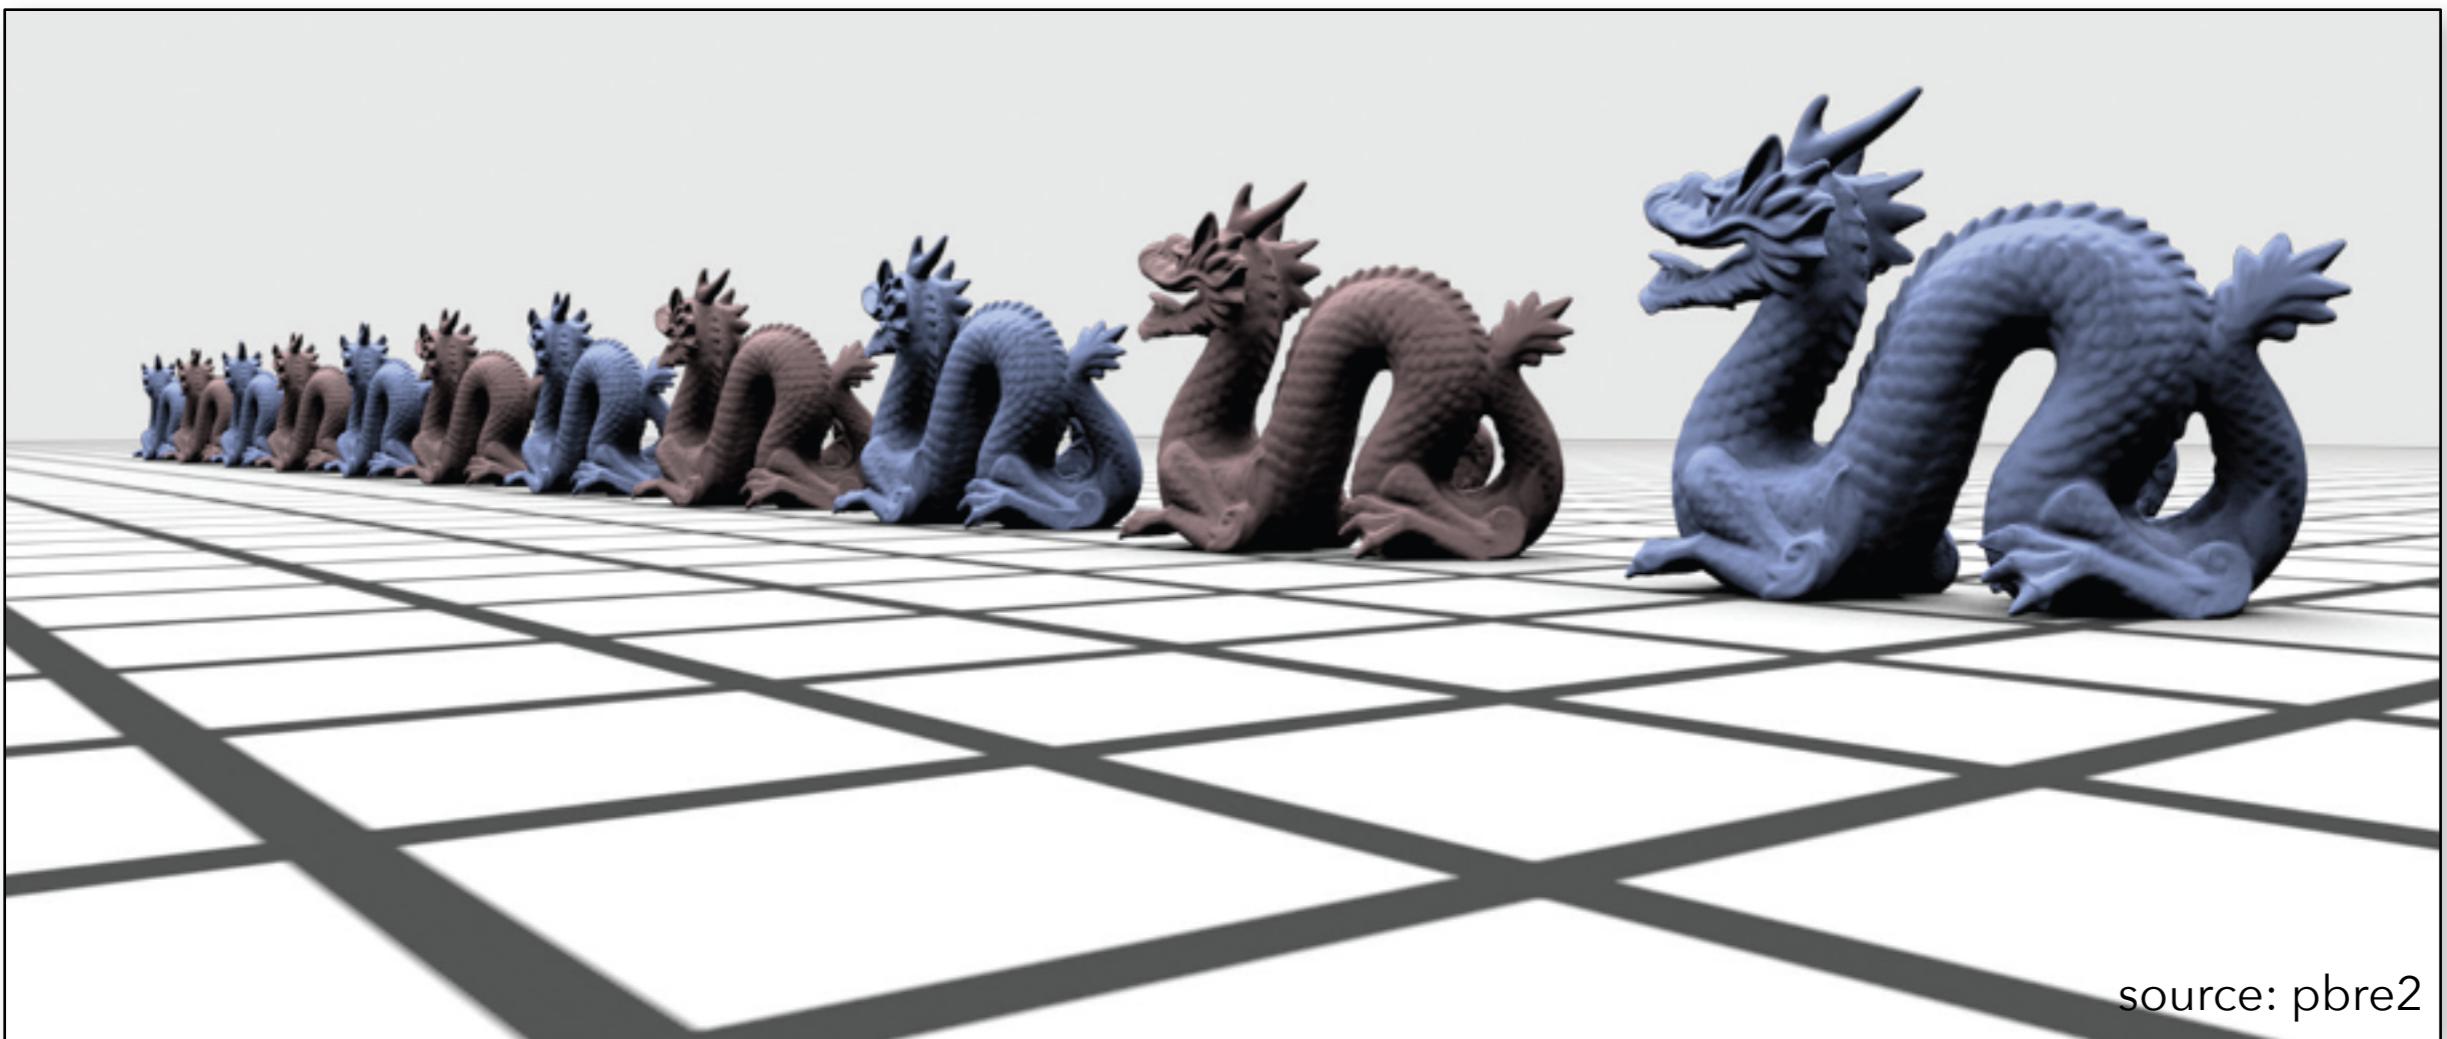
II

Lighting



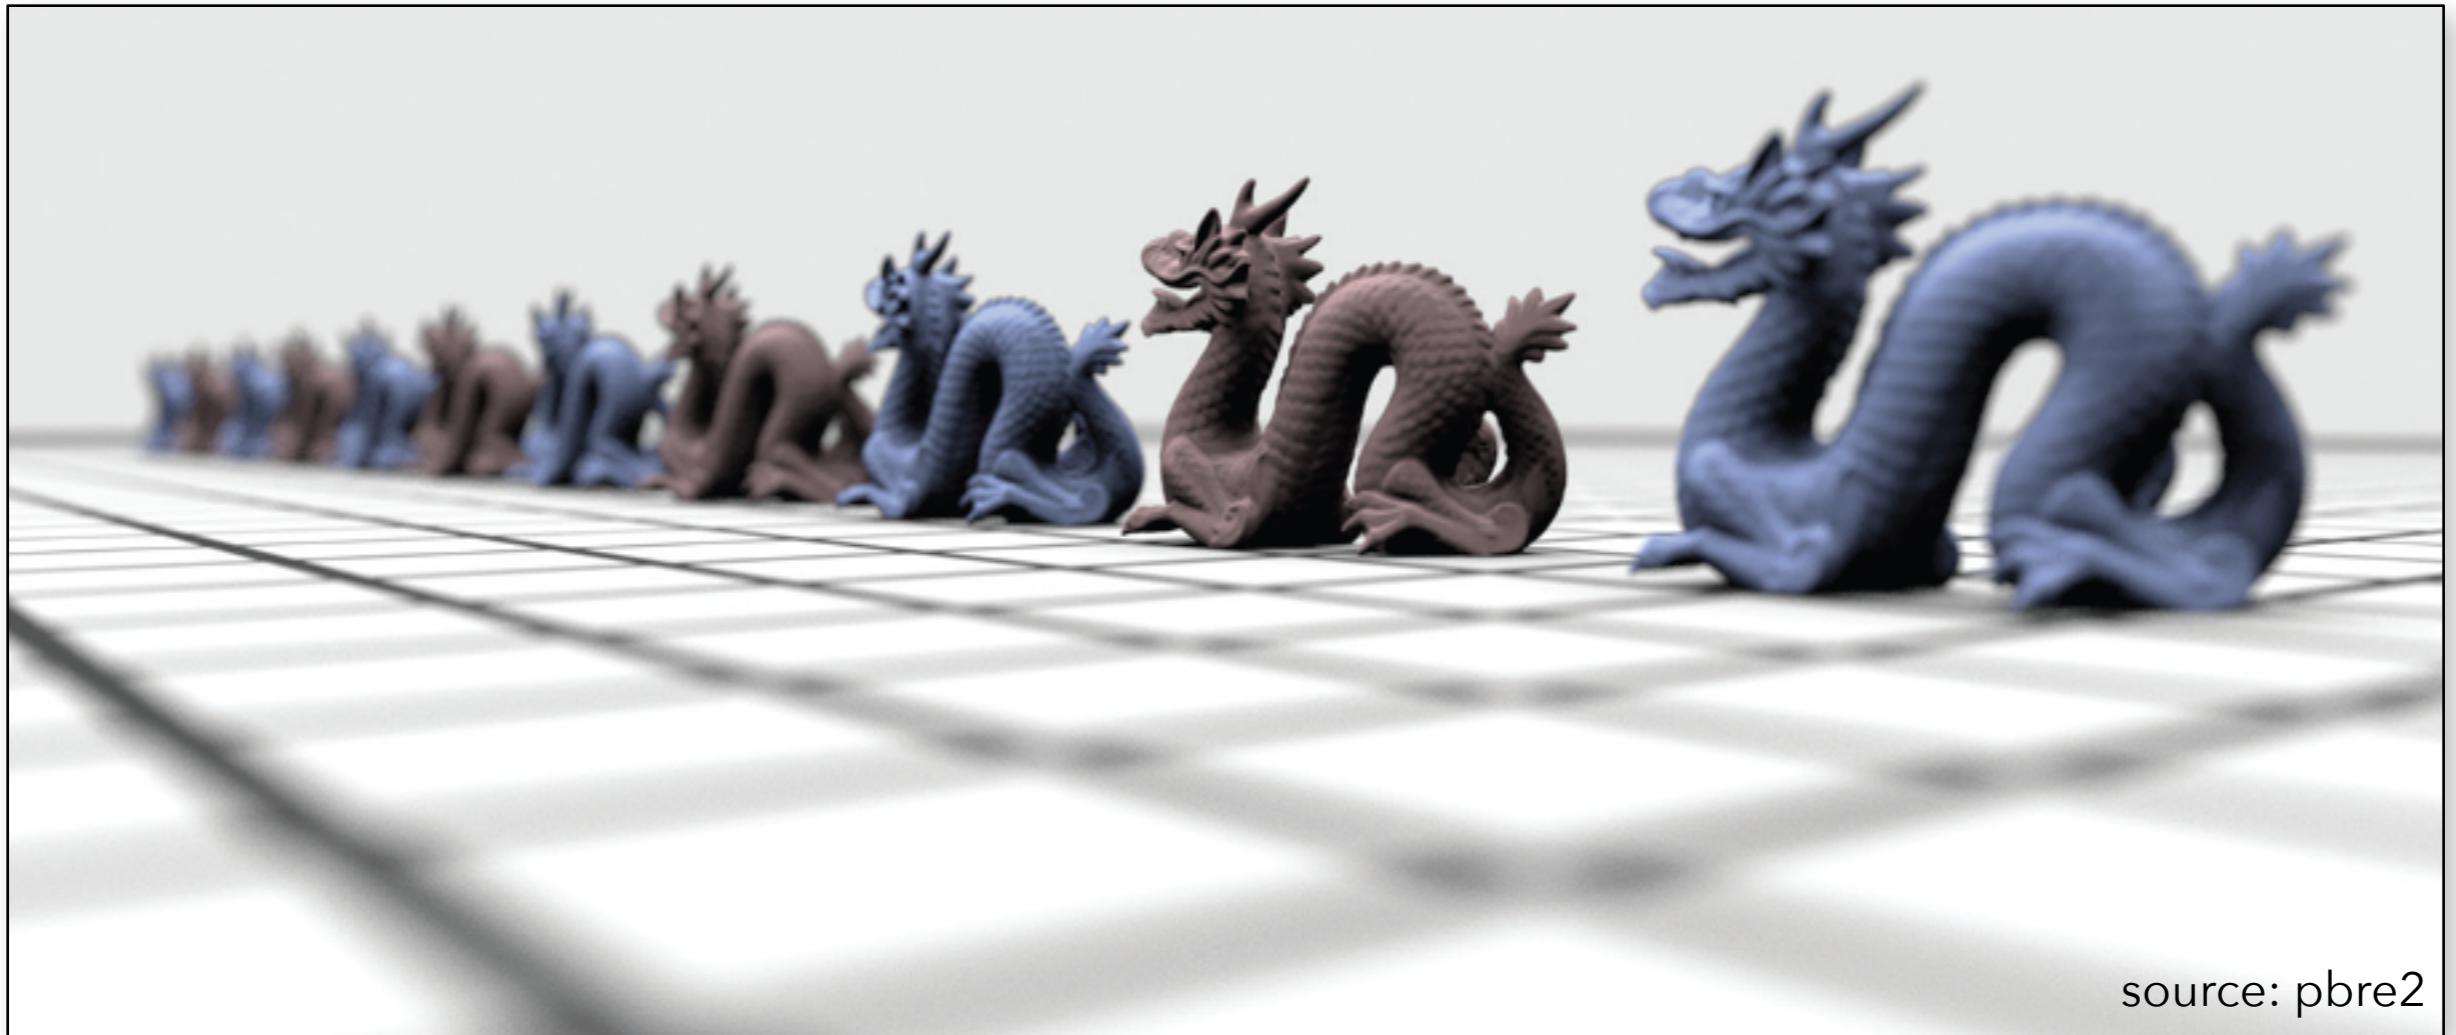
BRDF x Lighting

Pinhole Camera



source: pbre2

Thin-Lens Camera (larger aperture)



source: pbre2

Monte Carlo Integration Recap

- Goal: evaluate integral $\int_a^b f(x)dx$
- Random variable $X_i \sim p(x)$
- Monte Carlo Estimator $F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(X_i)}{p(X_i)}$
- Expectation $E[F_N] = \int_a^b f(x)dx$

Today's Menu

- Variance reduction techniques
 - Control variates
- Better random variables using
 - Stratified Sampling
 - Quasi Monte Carlo

Monte Carlo

- The standard MC estimator:

$$F = \int_{\mu(x)} f(x) \, d\mu(x)$$

$$\langle F^N \rangle = \frac{1}{N} \sum_{i=0}^{N-1} \frac{f(X_i)}{\text{pdf}(X_i)}$$

Monte Carlo

- The standard MC estimator:

$$F = \int_{\mu(x)} f(x) \, d\mu(x)$$

$$V[\langle F^N \rangle] = V \left[\frac{1}{N} \sum_{i=0}^{N-1} \frac{f(X_i)}{\text{pdf}(X_i)} \right]$$

Monte Carlo Error

- Variance:

$$\begin{aligned} V[\langle F^N \rangle] &= V\left[\frac{1}{N} \sum_{i=0}^{N-1} \frac{f(X_i)}{\text{pdf}(X_i)}\right] \quad \leftarrow \text{assume uncorrelated samples} \\ &= \frac{1}{N^2} \sum_{i=0}^{N-1} V\left[\frac{f(X_i)}{\text{pdf}(X_i)}\right] \\ &= \frac{1}{N^2} \sum_{i=0}^{N-1} V[Y_i] \\ &= \frac{1}{N} V[Y] \end{aligned}$$

Monte Carlo Error

- Variance:

$$\begin{aligned} V[\langle F^N \rangle] &= V\left[\frac{1}{N} \sum_{i=0}^{N-1} \frac{f(X_i)}{\text{pdf}(X_i)}\right] \\ &= \frac{1}{N^2} \sum_{i=0}^{N-1} V\left[\frac{f(X_i)}{\text{pdf}(X_i)}\right] \\ &= \frac{1}{N^2} \sum_{i=0}^{N-1} V[Y_i] \\ &= \frac{1}{N} V[Y] \end{aligned}$$

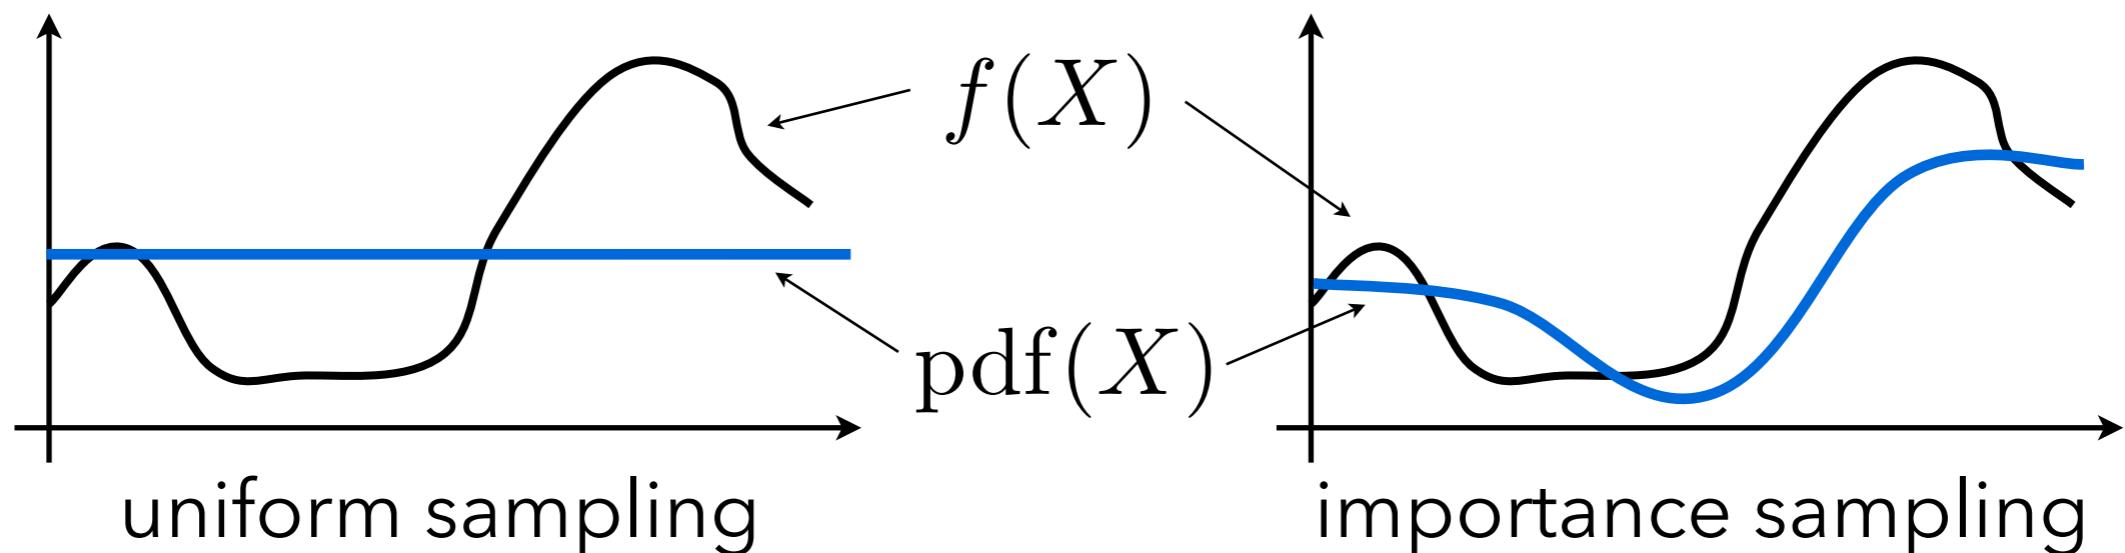
- Std. deviation:

$$\sigma[\langle F^N \rangle] = \boxed{\frac{1}{\sqrt{N}} \sigma[Y]}$$

Strategies for Reducing Variance

$$\sigma [\langle F^N \rangle] = \frac{1}{\sqrt{N}} \sigma [Y]$$

- Reduce the variance of Y
 - Importance sampling



Strategies for Reducing Variance

$$\sigma [\langle F^N \rangle] = \frac{1}{\sqrt{N}} \sigma [Y]$$

- Reduce the variance of Y
 - Importance sampling
 - Control Variates

Control Variates

- Say you have a function g , which you can efficiently (e.g. analytically) integrate to G

$$F = \int_a^b f(x) \, dx,$$

- Control variates estimator:

$$\langle F_c^N \rangle = \left(\frac{1}{N} \sum_{i=0}^{N-1} \frac{f(X_i) - g(X_i)}{\text{pdf}(X_i)} \right) + G$$

Control Variates

- Control variates estimator:

$$\langle F_c^N \rangle = \left(\frac{1}{N} \sum_{i=0}^{N-1} \frac{f(X_i) - g(X_i)}{\text{pdf}(X_i)} \right) + G$$

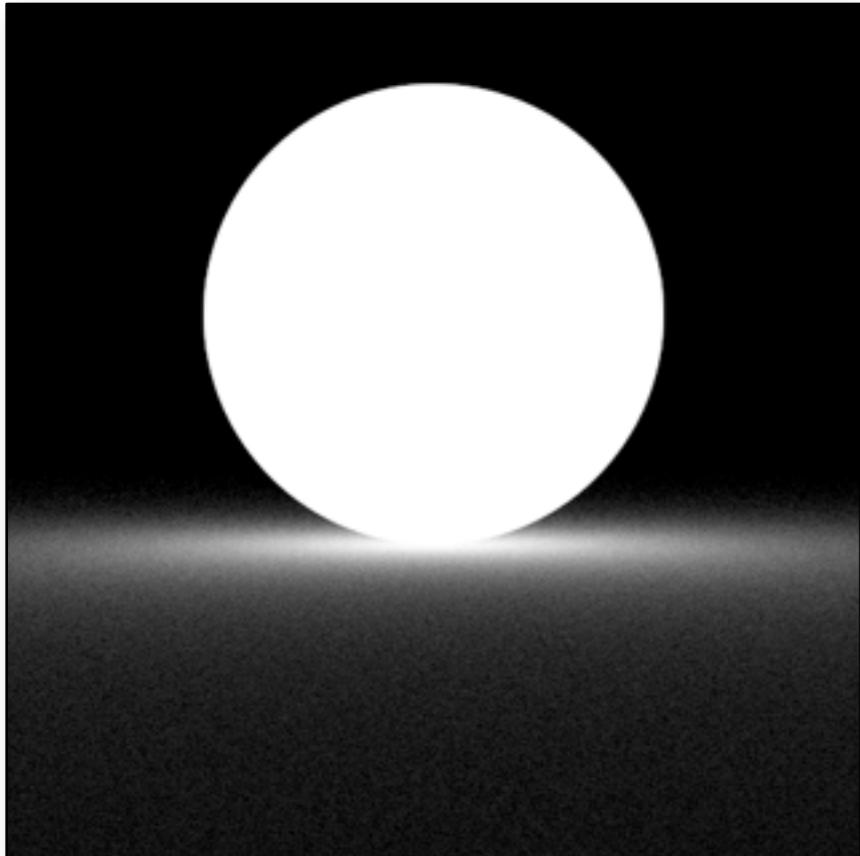
- Recall that:

$$\sigma [\langle F^N \rangle] = \frac{1}{\sqrt{N}} \sigma [Y]$$

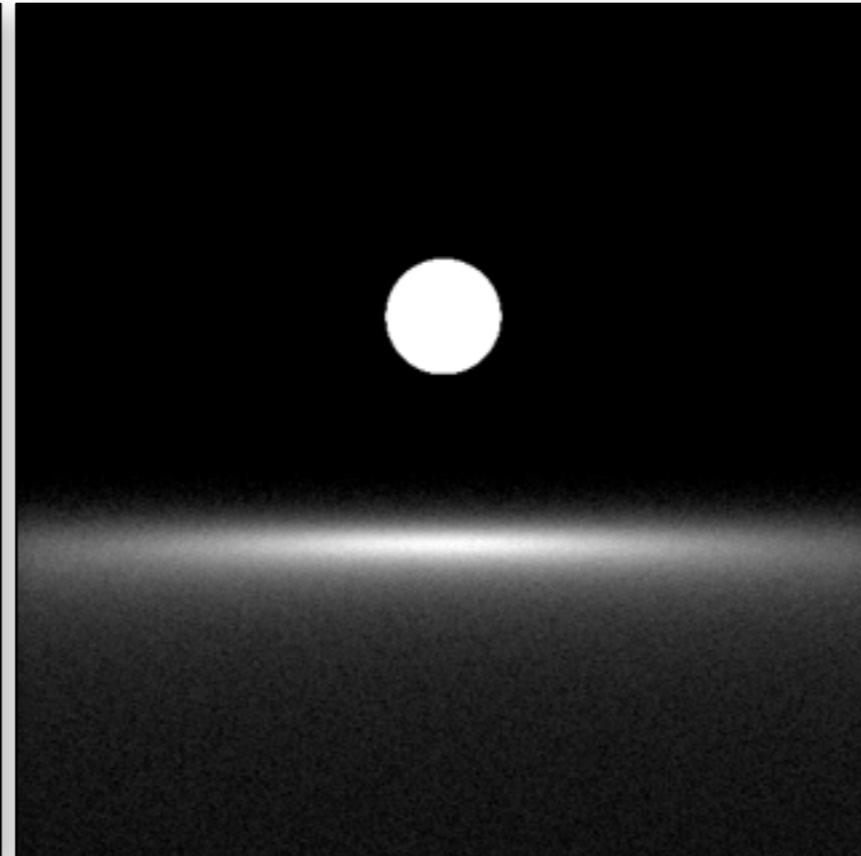
- and in this case:

$$Y \equiv \frac{f(X) - g(X)}{\text{pdf}(X)}$$

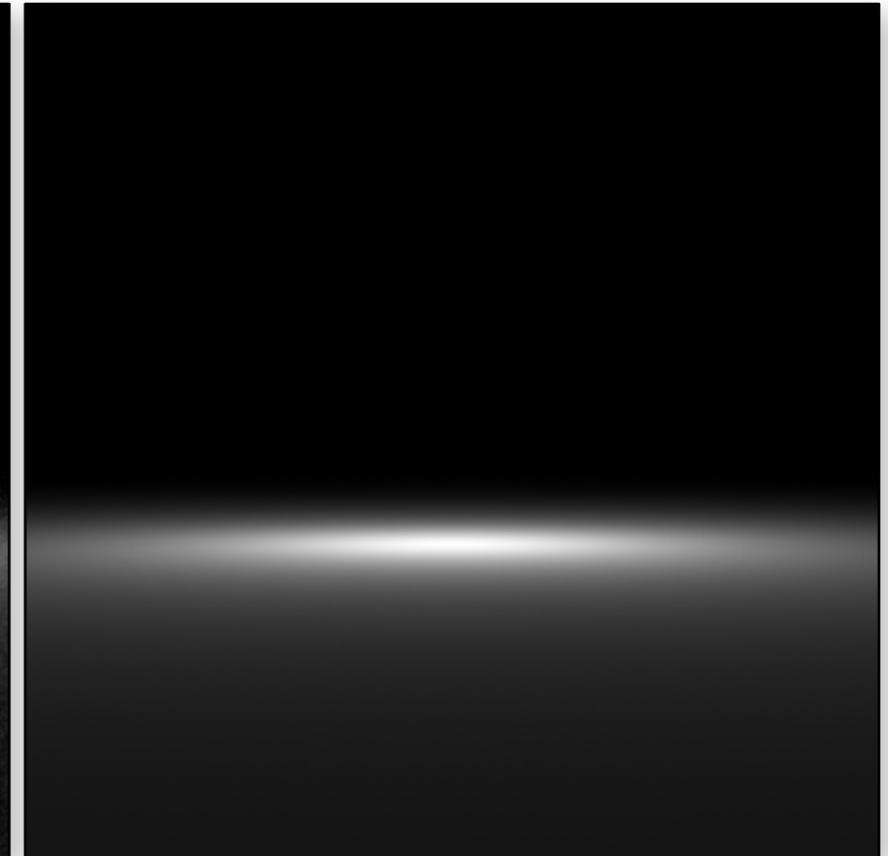
Sphere Light



A sphere light



A smaller sphere light



A point light

shading of diffuse surface independent of sphere light radius

idea: use analytic point light as control variate for sphere lights

Control Variates (Sphere light)

- Exploiting analytic sphere lights:

$$L_r(\omega_o) = \int_{\Omega} f_r(\omega_i, \omega_o) L_i(\omega_i, \omega_o) V(\omega_i) \cos \theta_i \, d\omega_i$$
$$\langle L_r(\omega_o) \rangle = \left(\frac{1}{N} \sum_{j=0}^{N-1} \frac{f(\omega_{i,j}) - g(\omega_{i,j})}{\text{pdf}(\omega_{i,j})} \right) + G$$

- where:

$$f(\omega_i) = f_s(\omega_i, \omega_o) L_i(\omega_i, \omega_o) V(\omega_i) \cos \theta_i$$

$$g(\omega_i) = L_i(\omega_i, \omega_o) \cos \theta_i$$

$$G = \frac{\Phi_e \cos \theta}{4\pi r^2}$$

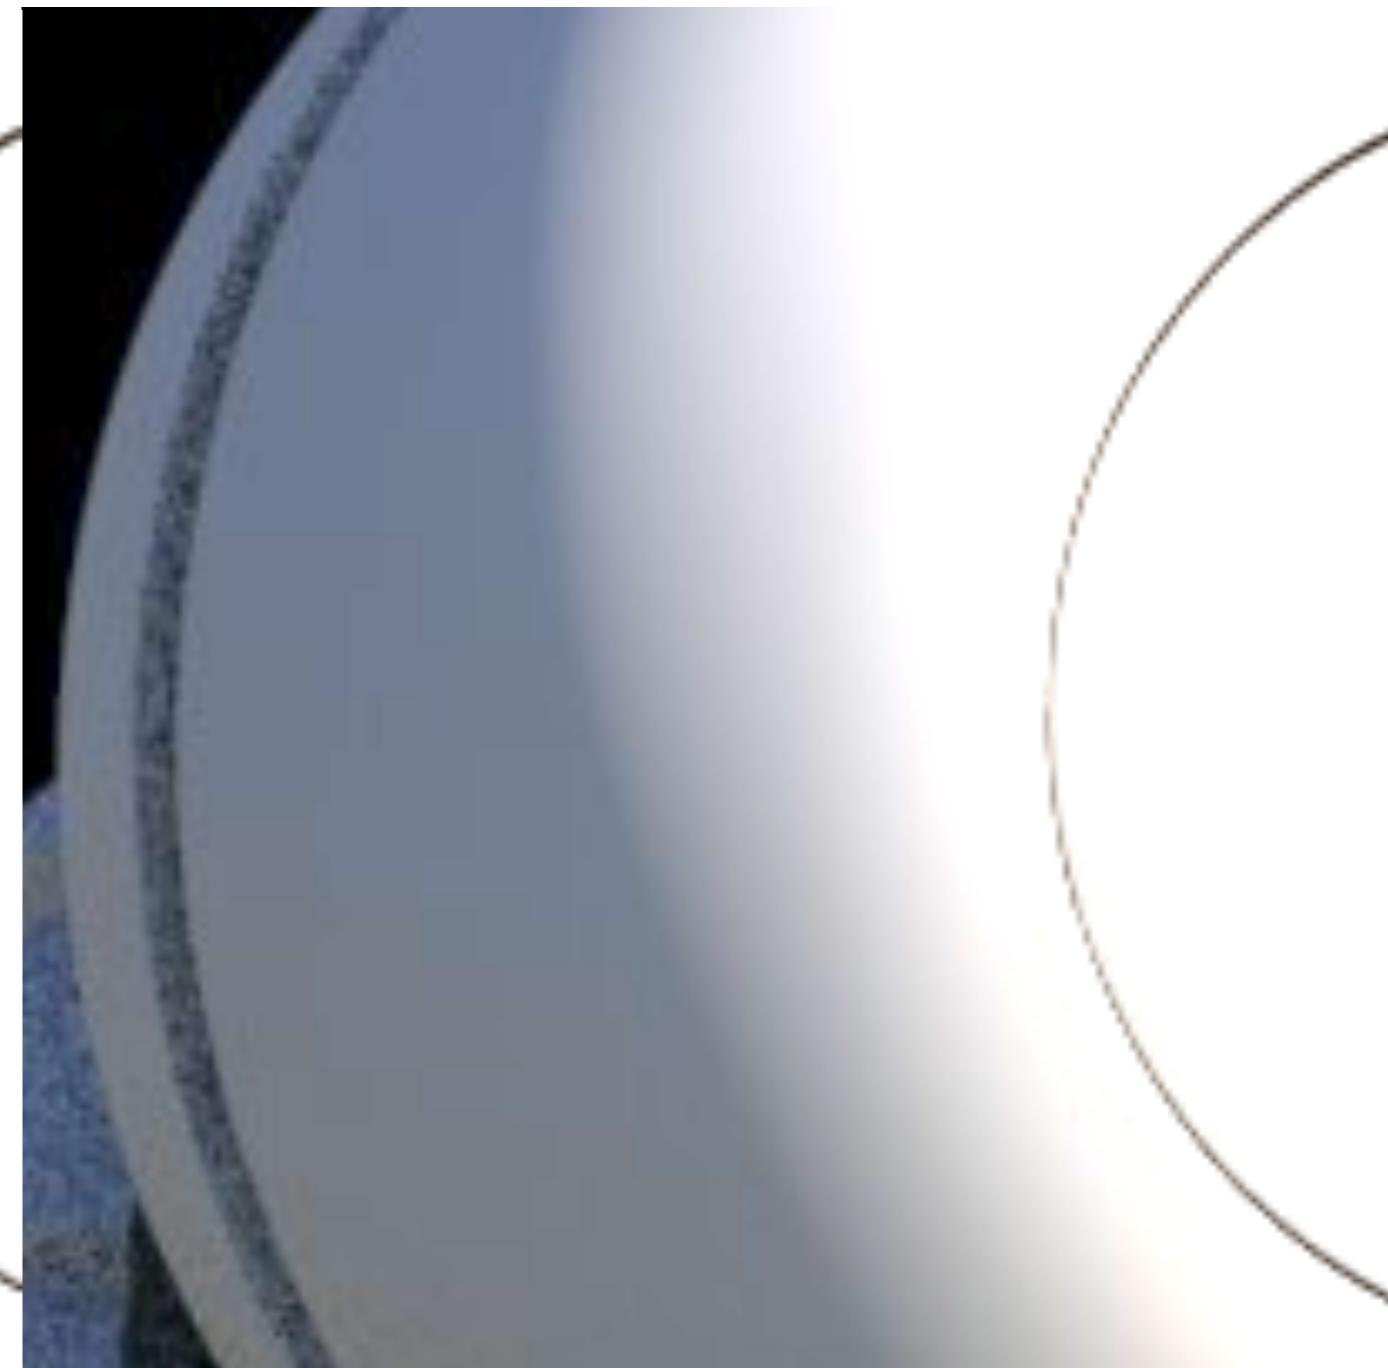
For sphere lights clipping the horizon see: Hery and Villemain. "Physically Based Lighting at Pixar." SIGGRAPH 2013.

Control Variates (Sphere light)

Control variates off



Control variates on



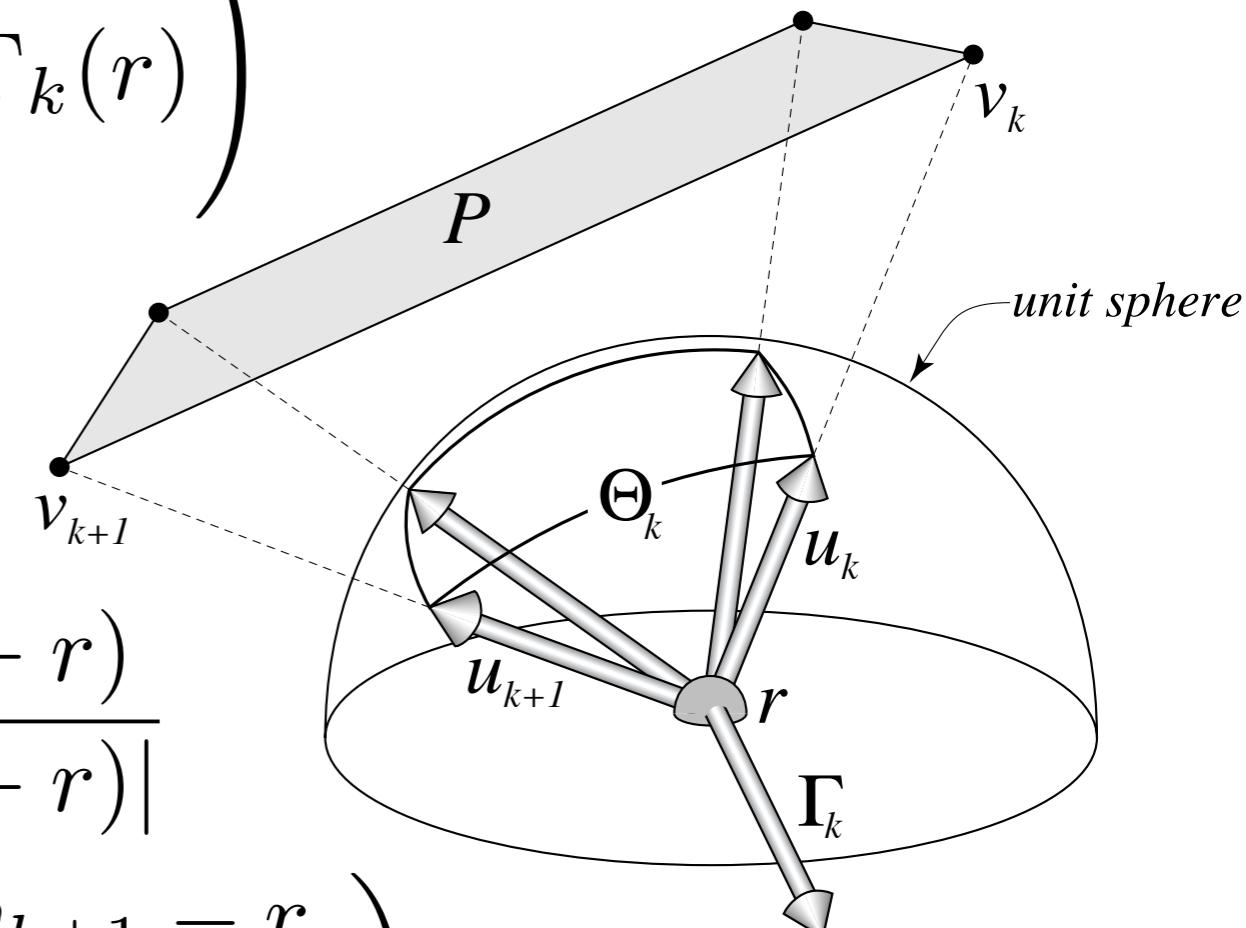
from Hery and Villemin. "Physically Based Lighting at Pixar." SIGGRAPH 2013.

Analytic polygon light

- Can obtain the irradiance due to a diffuse, polygonal source in closed form using Stoke's theorem

$$E(r) = -\mathbf{n} \cdot \left(\frac{M}{2\pi} \sum_{k=1}^n \Theta_k(r) \Gamma_k(r) \right)$$

$$\Gamma_k(r) = \frac{(v_k - r) \times (v_{k+1} - r)}{|(v_k - r) \times (v_{k+1} - r)|}$$
$$\Theta_k(r) = \cos^{-1} \left(\frac{v_k - r}{|v_k - r|} \cdot \frac{v_{k+1} - r}{|v_{k+1} - r|} \right)$$



J. Arvo. 1994. "The Irradiance Jacobian for Partially Occluded Polyhedral Sources"

Control Variates (Polygon light)

- Exploit analytic polygonal irradiance

$$L_r(\omega_o) = \int_{\Omega} f_r(\omega_i, \omega_o) L_i(\omega_i, \omega_o) V(\omega_i) \cos \theta_i \, d\omega_i$$
$$\langle L_r(\omega_o) \rangle = \left(\frac{1}{N} \sum_{j=0}^{N-1} \frac{f(\omega_{i,j}) - g(\omega_{i,j})}{\text{pdf}(\omega_{i,j})} \right) + G$$

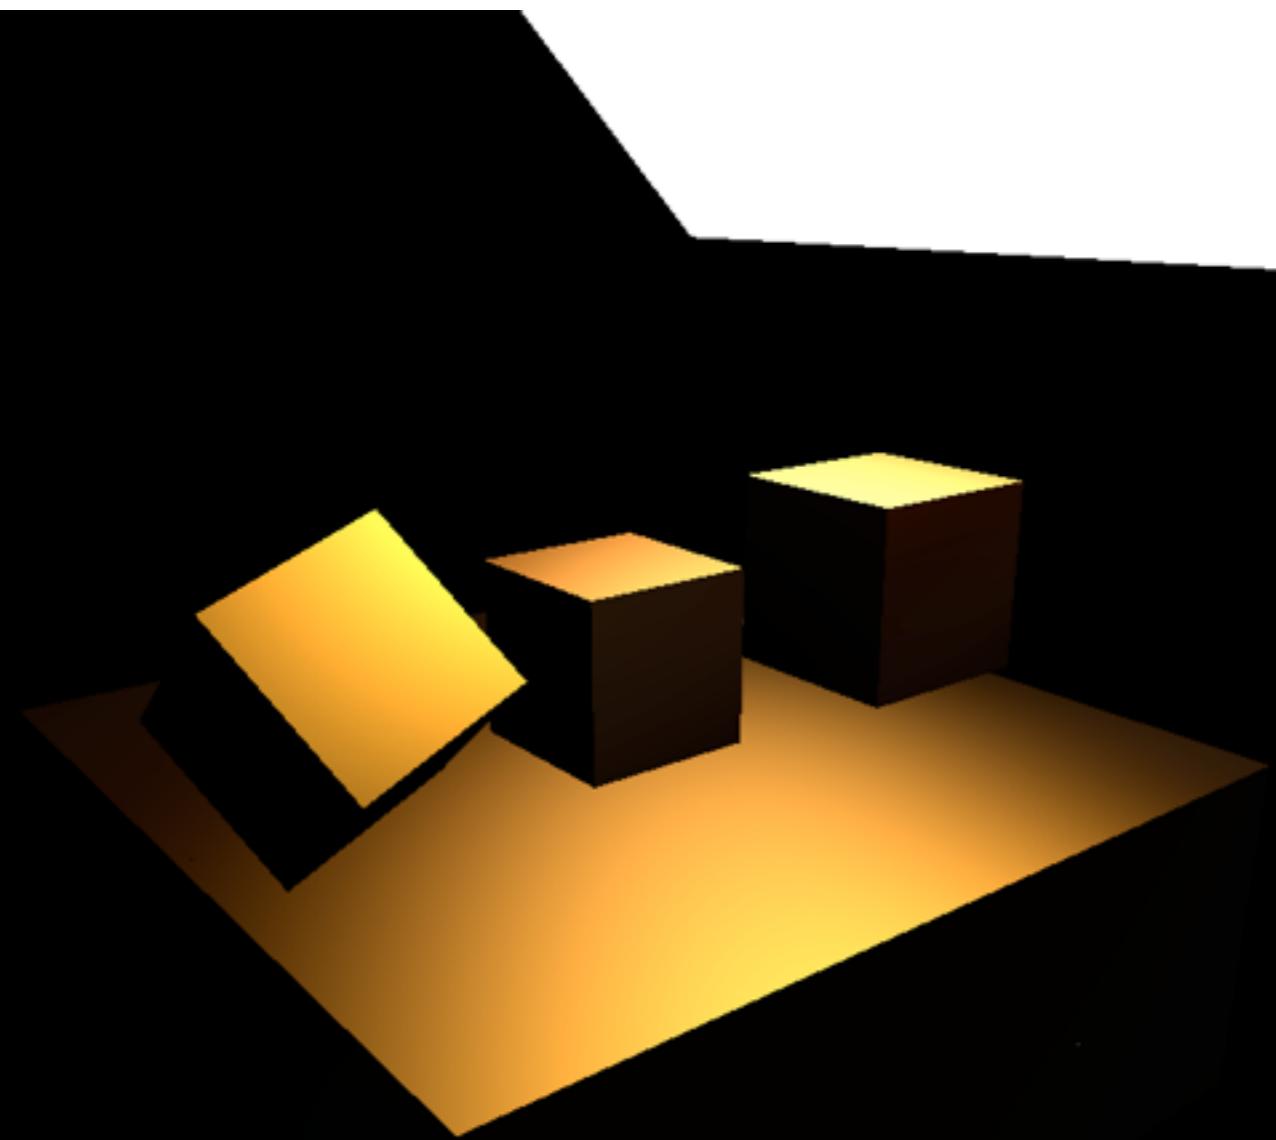
- where:

$$f(\omega_i) = f_s(\omega_i, \omega_o) L_i(\omega_i, \omega_o) V(\omega_i) \cos \theta_i$$

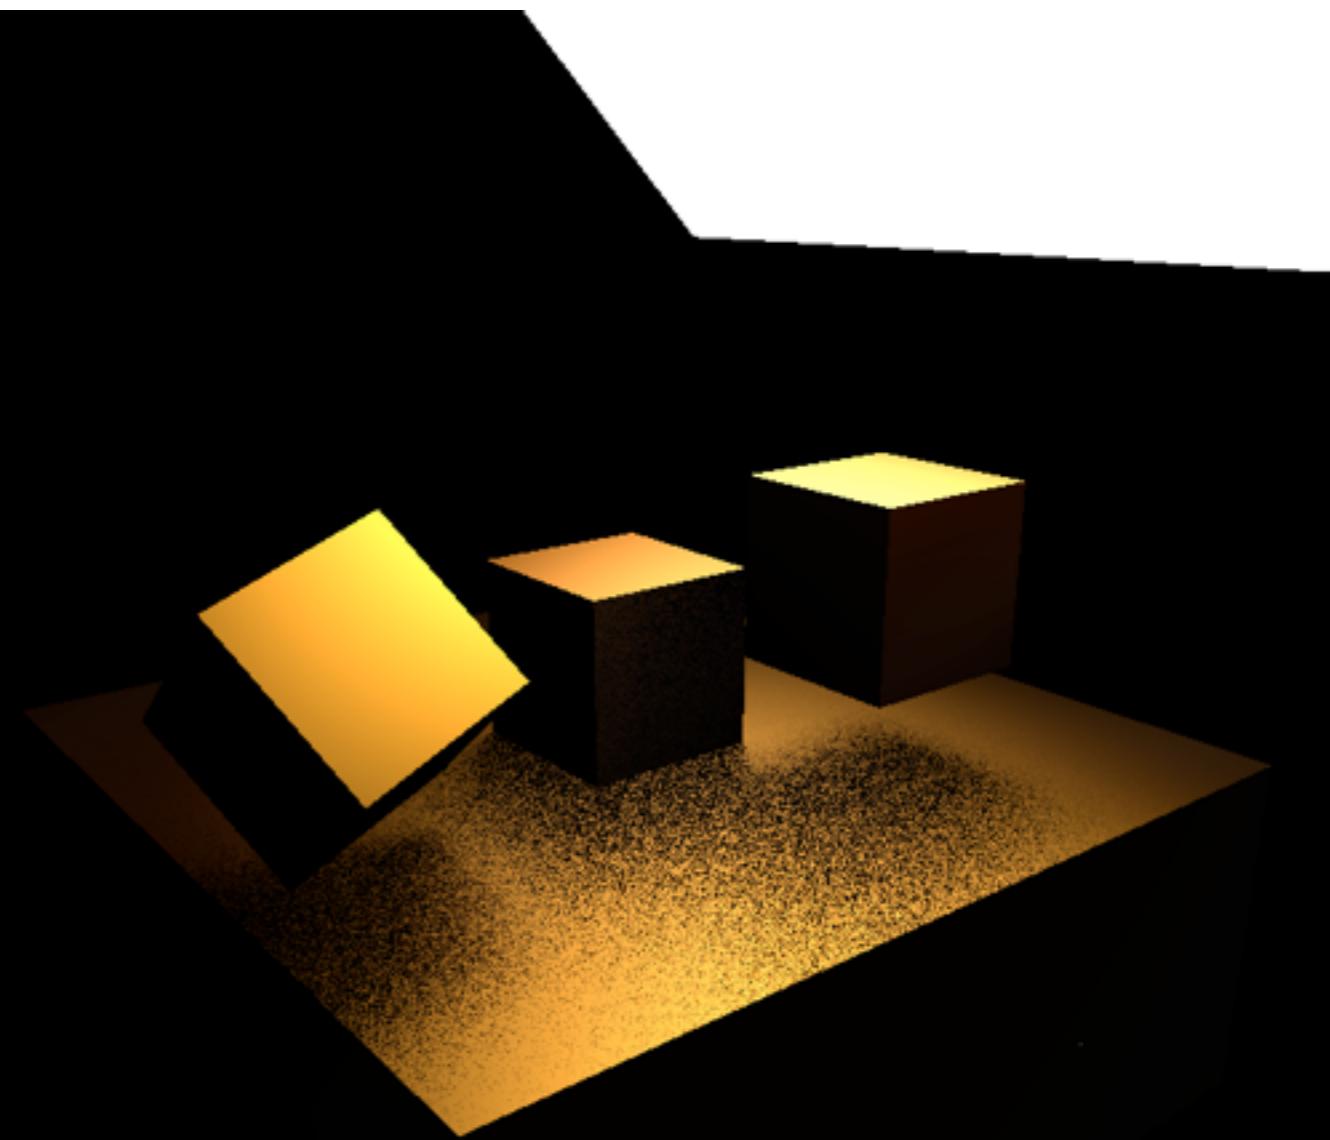
$$g(\omega_i) = L_i(\omega_i, \omega_o) \cos \theta_i$$

$$G = -\mathbf{n} \cdot \left(\frac{M}{2\pi} \sum_{k=1}^n \Theta_k \Gamma_k \right)$$

Analytic area lights as control variates



analytic result



control variate +
MC shadows

Control Variates (Env map)

$$L_r(\omega_o) = \int_{\Omega} f_r(\omega_i, \omega_o) L_{\text{env}}(\omega_i) V(\omega_i)(\omega_i \cdot \vec{n}) \, d\omega_i$$

- Precompute diffuse convolution for an env map

$$L_{\text{env}}(\omega_i)$$

$$G(\vec{n}) = \int_{\Omega} L_{\text{env}}(\omega_i)(\omega_i \cdot \vec{n}) \, d\omega_i$$



Control Variates (Env map)

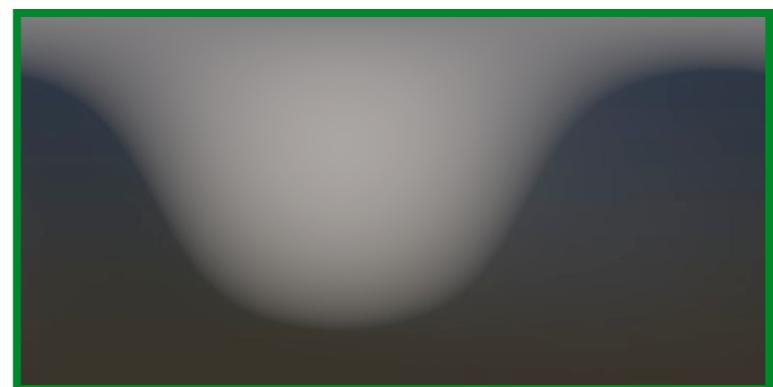
$$L_r(\omega_o) = \int_{\Omega} f_r(\omega_i, \omega_o) L_{\text{env}}(\omega_i) V(\omega_i)(\omega_i \cdot \vec{n}) \, d\omega_i$$

$$\langle L_r(\omega_o) \rangle = \left(\frac{1}{N} \sum_{j=0}^{N-1} \frac{f(\omega_{i,j}) - \frac{\rho}{\pi} g(\omega_{i,j})}{\text{pdf}(\omega_{i,j})} \right) + \frac{\rho}{\pi} G(\vec{n})$$

$$f(\omega_i) = f_r(\omega_i, \omega_o) \boxed{L_{\text{env}}(\omega_i)} V(\omega_i)(\omega_i \cdot \vec{n})$$

$$g(\omega_i) = \boxed{L_{\text{env}}(\omega_i)} (\omega_i \cdot \vec{n})$$

$$G(\vec{n}) = \int_{\Omega} L_{\text{env}}(\omega_i)(\omega_i \cdot \vec{n}) \, d\omega_i$$



Variance reduction summary

- If I have a function g which is “similar” to f , should I use it for importance sampling or control variates?
 - If $f-g$ is nearly constant, use g as a control variate
 - If f/g is nearly constant, use g as a PDF
- Importance sampling (IS)
 - must be able to sample from g (typically involves integrating to CDF and *inverting*)
- Control variates (CV)
 - only requires being able to evaluate g and efficiently compute G (does not actually need to be analytic!)

Variance reduction summary

- If I have a function g which is “similar” to f , should I use it for importance sampling or control variates?
- Can I employ both IS and CV?
 - Yes, but recall...

$$\langle F_c^N \rangle = \left(\frac{1}{N} \sum_{i=0}^{N-1} \frac{f(X_i) - g(X_i)}{\text{pdf}(X_i)} \right) + G$$

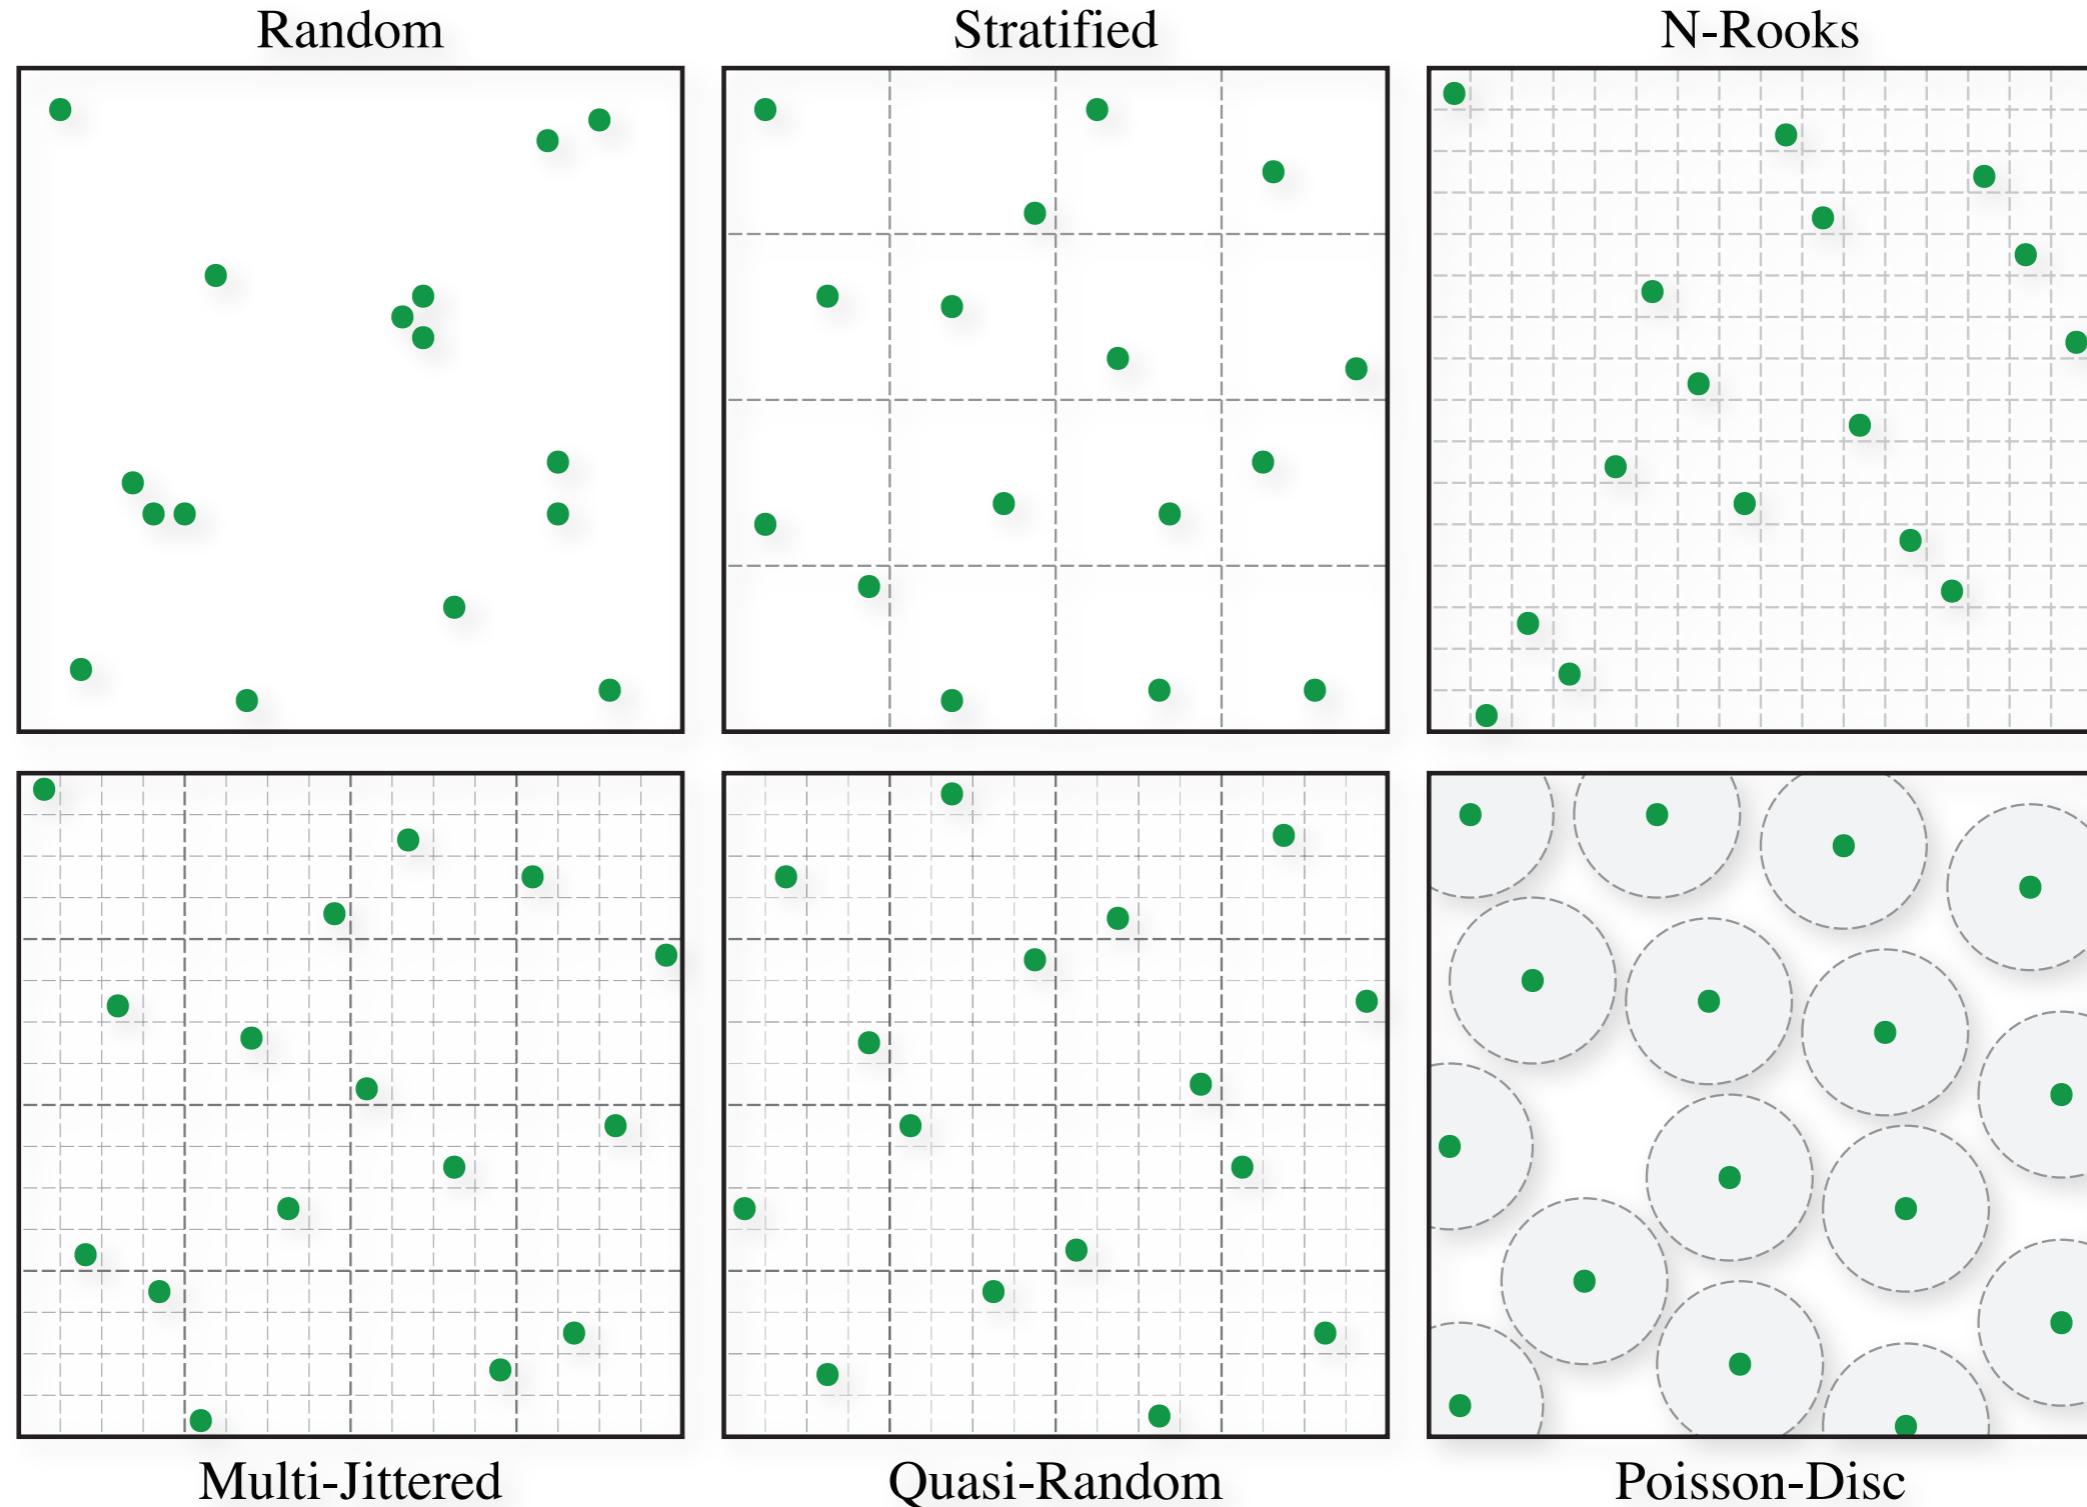
$$Y = \frac{f(X) - g(X)}{\text{pdf}(X)} \quad \sigma [\langle F^N \rangle] = \frac{1}{\sqrt{N}} \sigma [Y]$$

Strategies for Reducing Variance

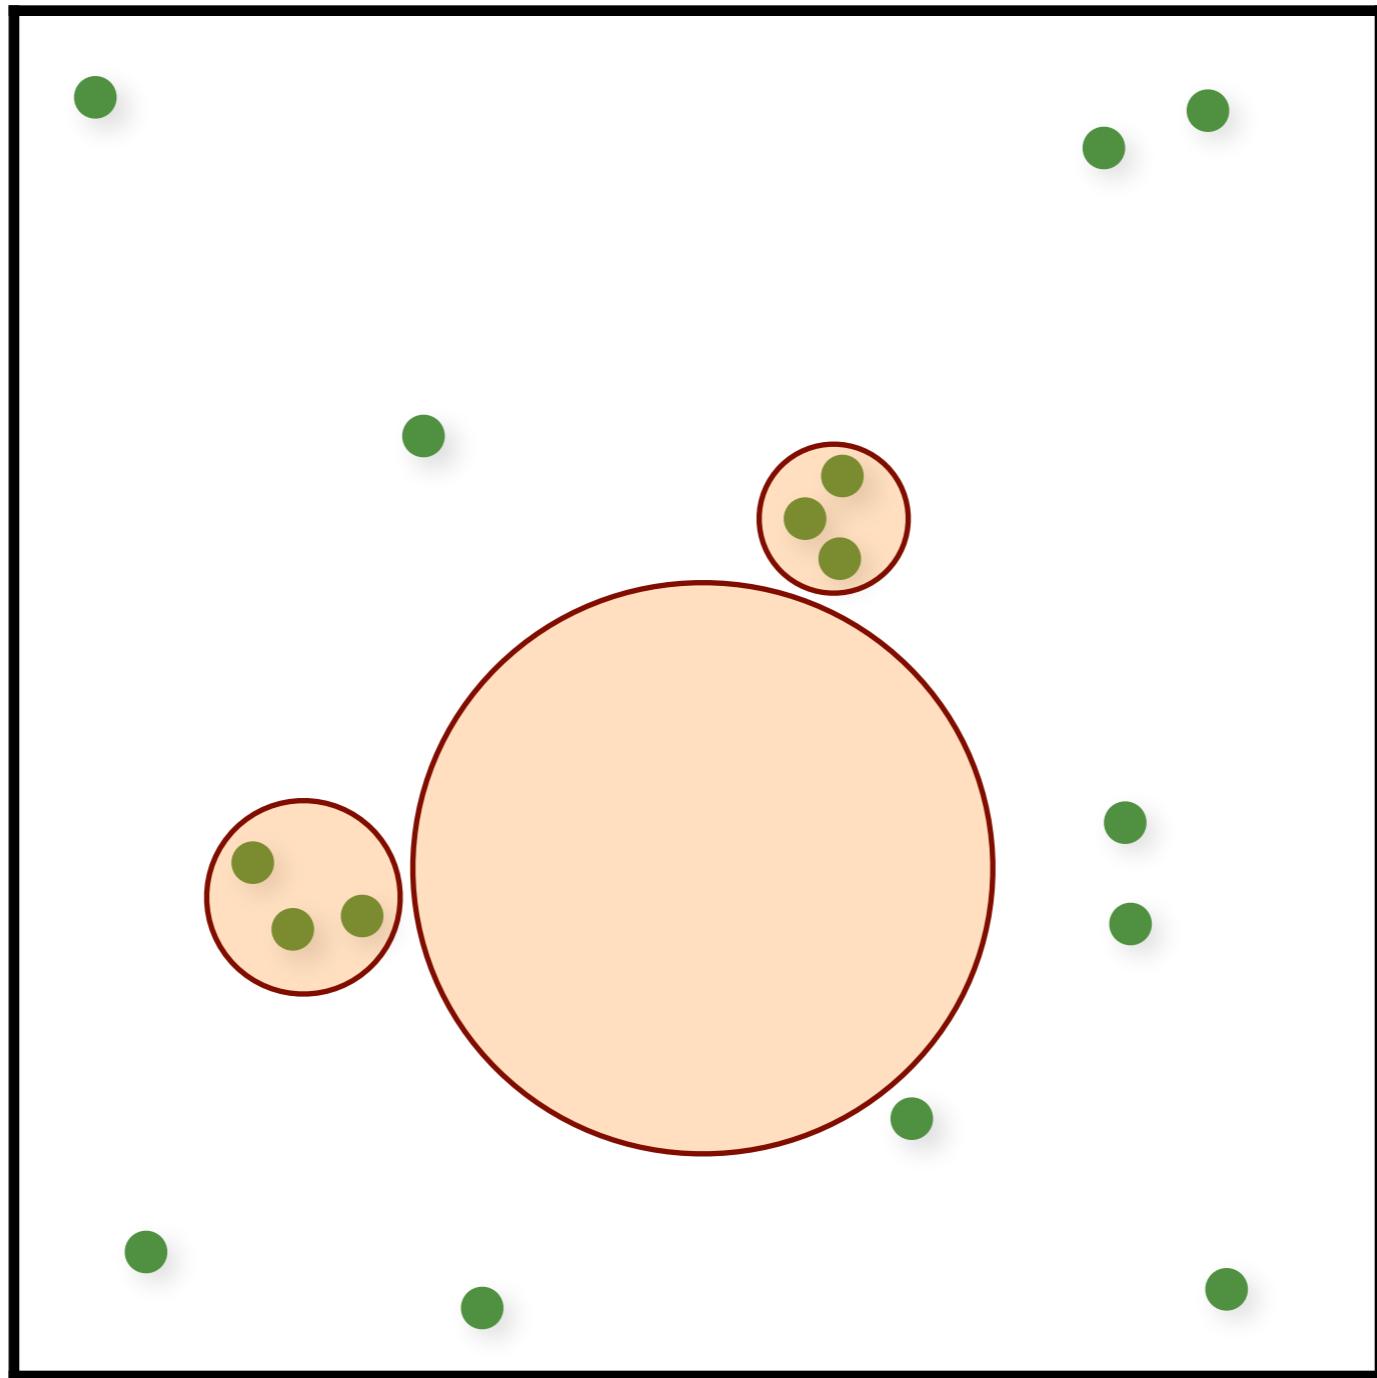
$$\sigma [\langle F^N \rangle] = \frac{1}{\sqrt{N}} \sigma [Y] \leftarrow \text{remember, this assumed uncorrelated samples}$$

- Reduce the variance of Y
 - Importance sampling
 - Control variates
- Relax assumption of uncorrelated samples

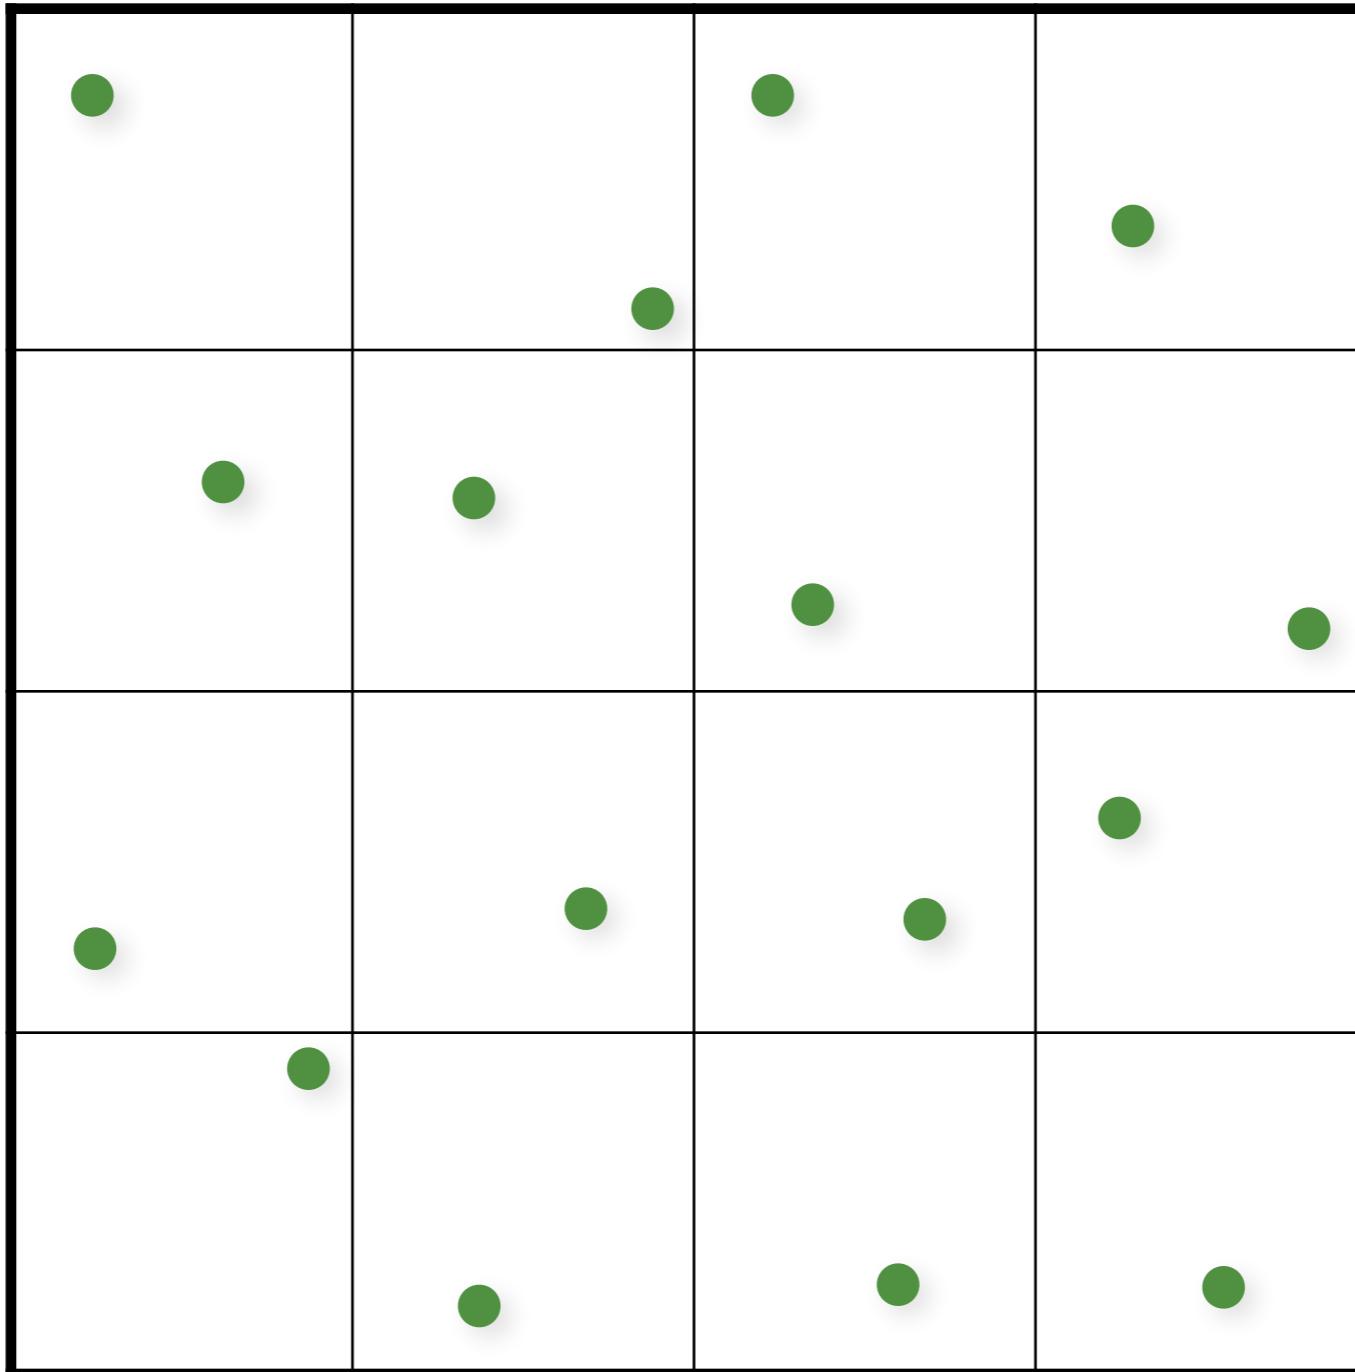
Careful Sample Placement



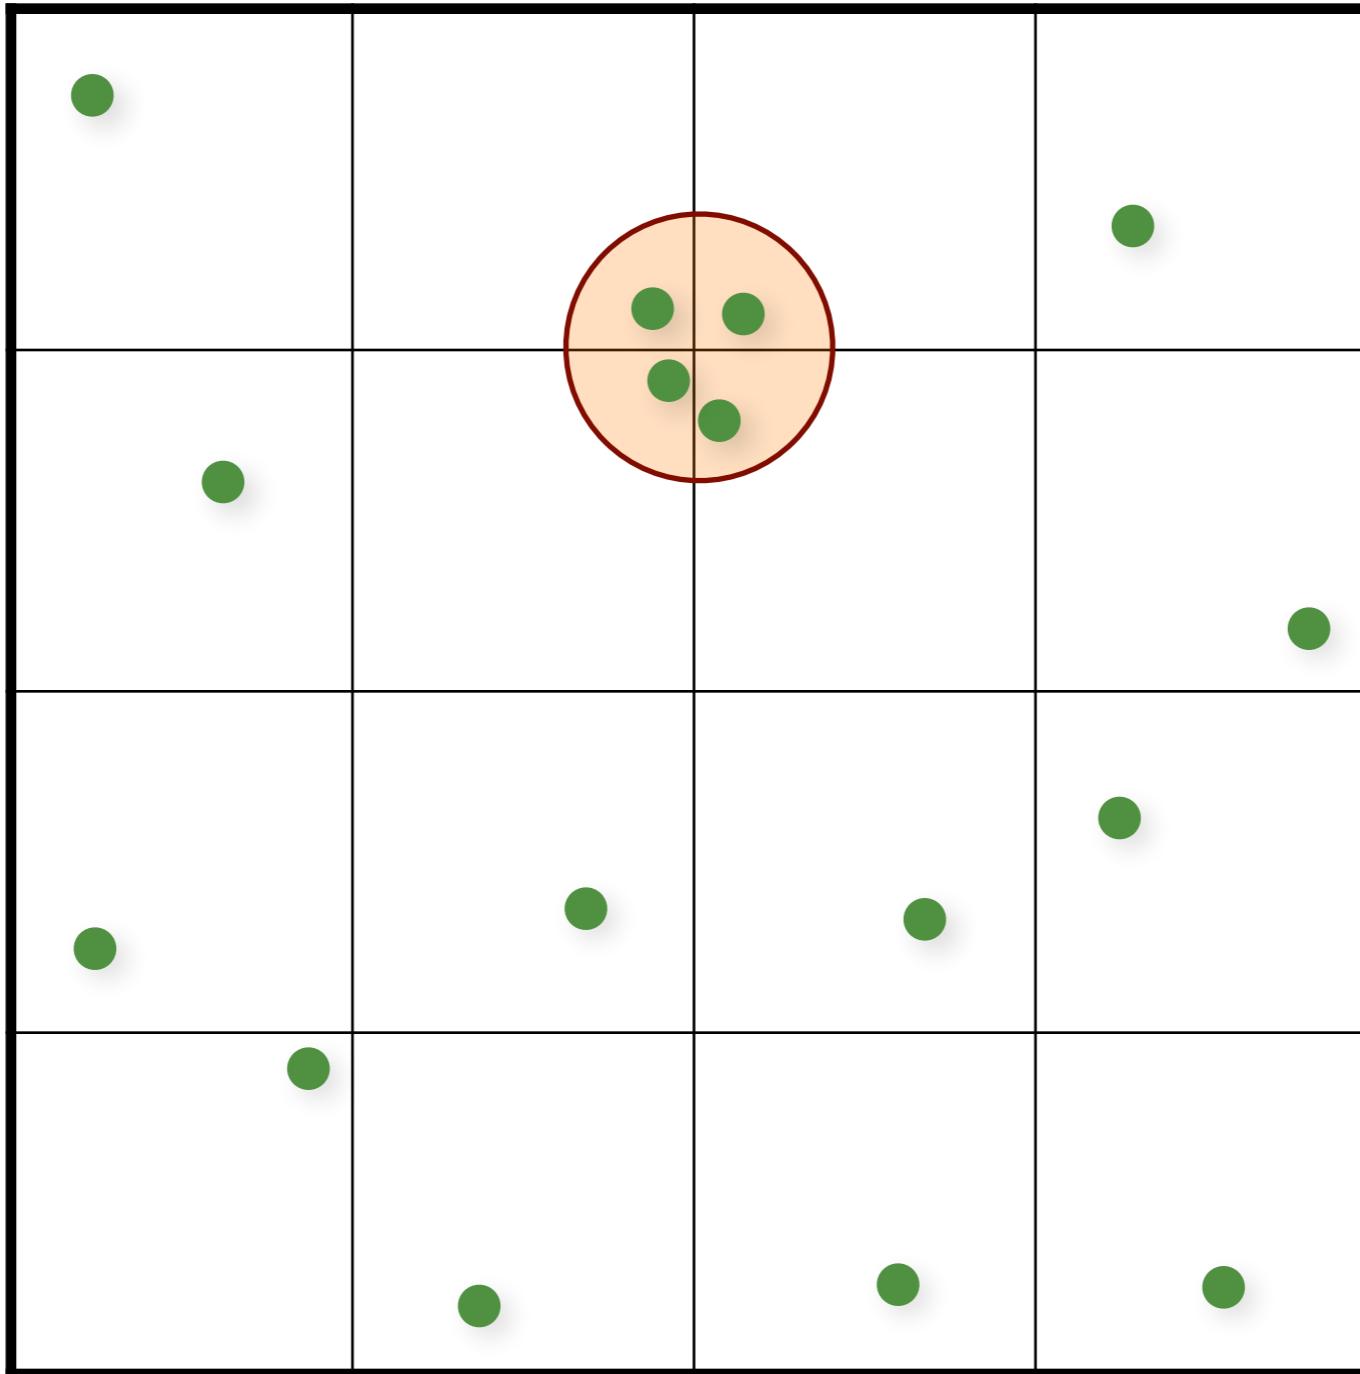
Random Sampling



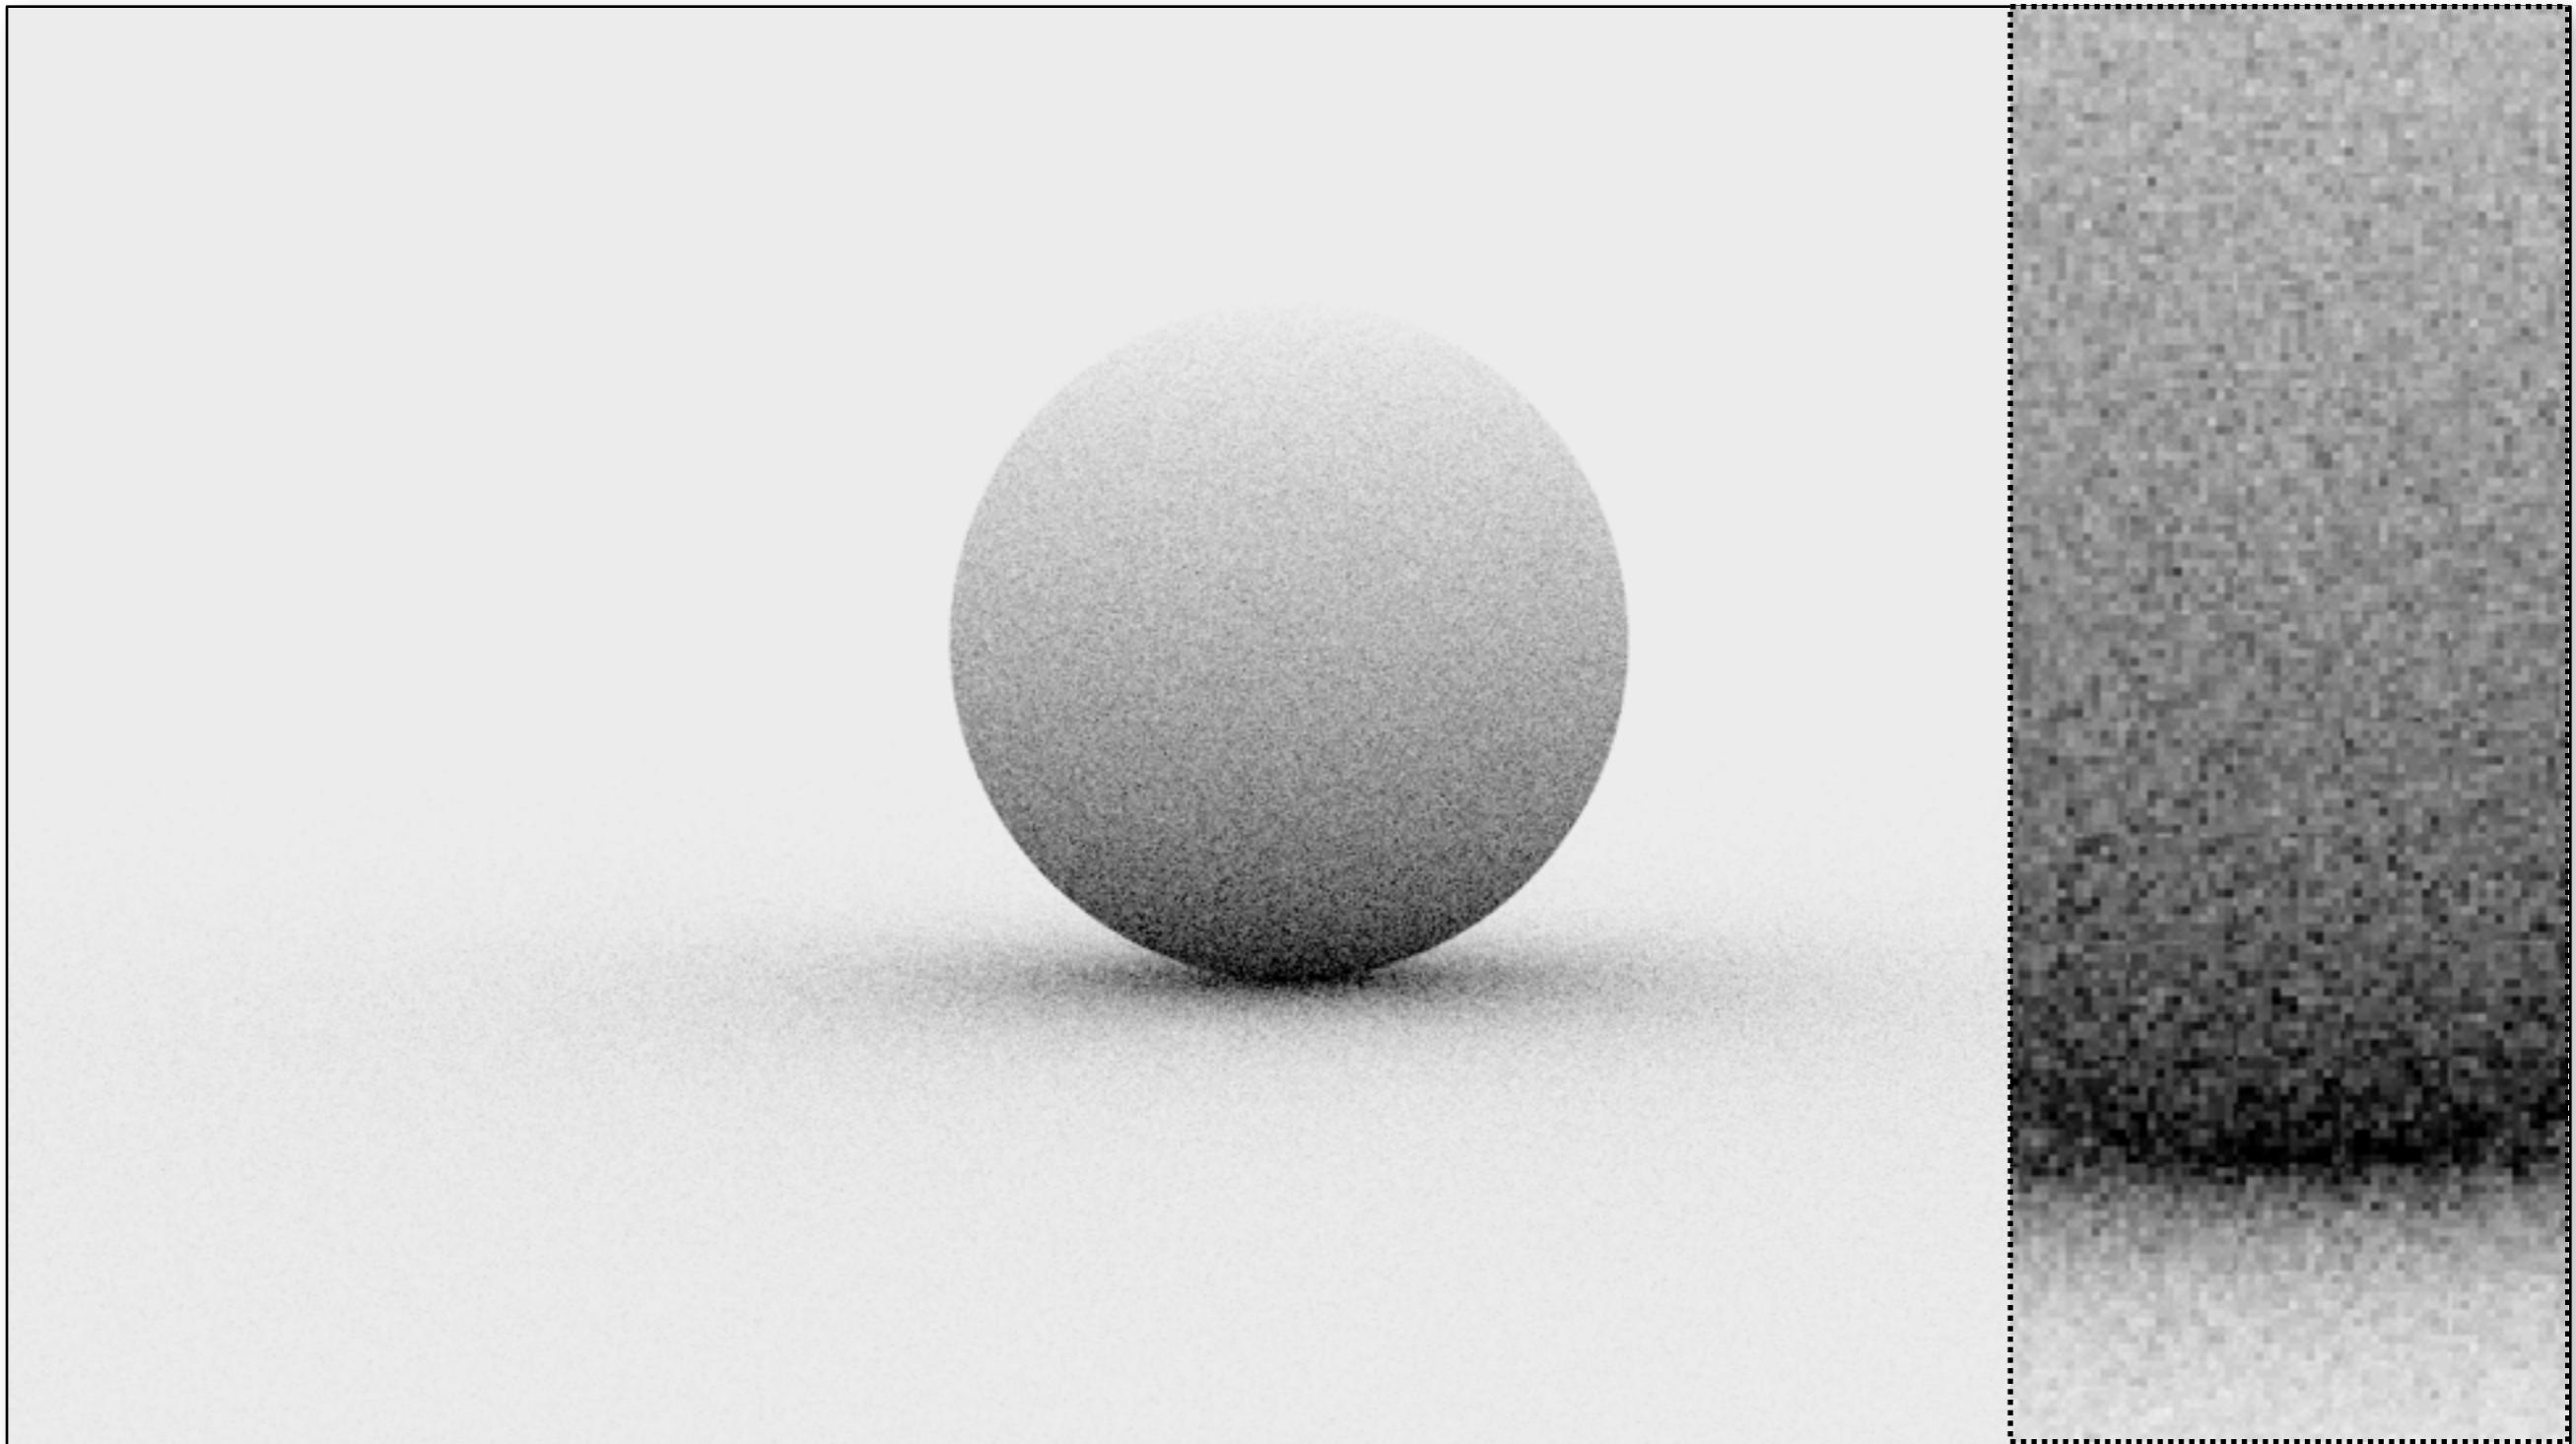
Stratified (Jittered) Sampling



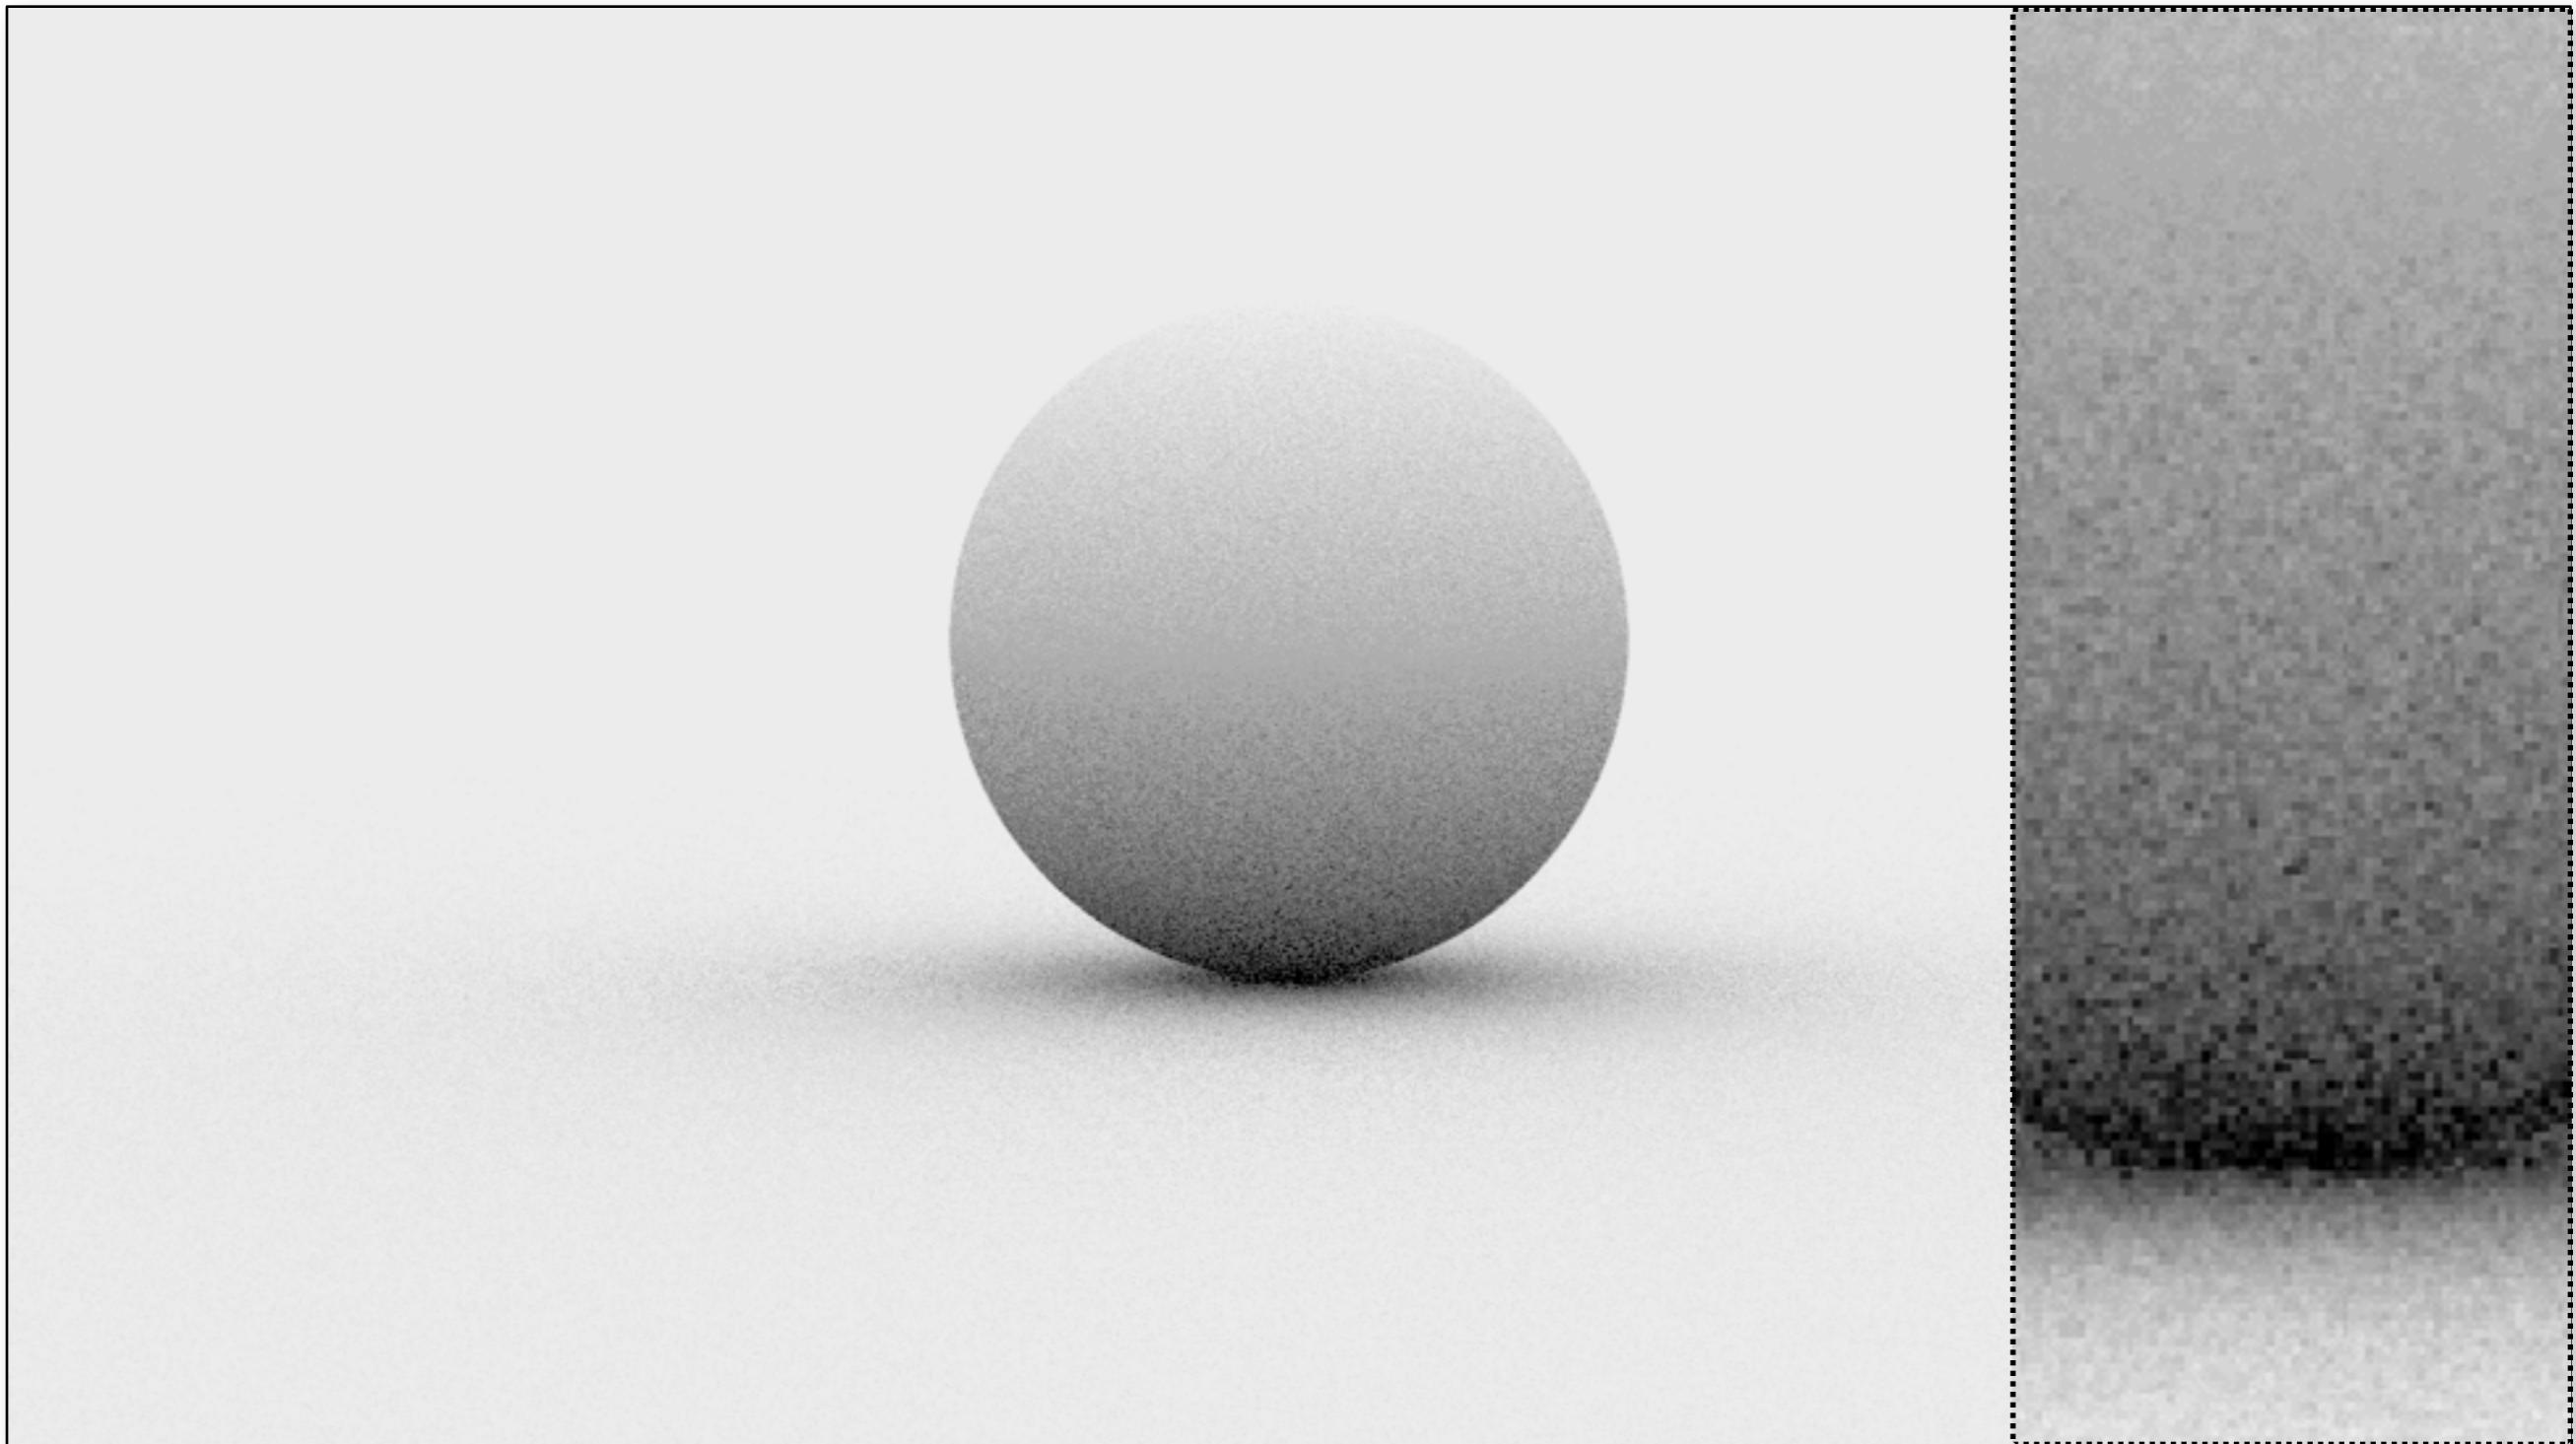
Stratified (Jittered) Sampling



Monte Carlo (16 random samples)

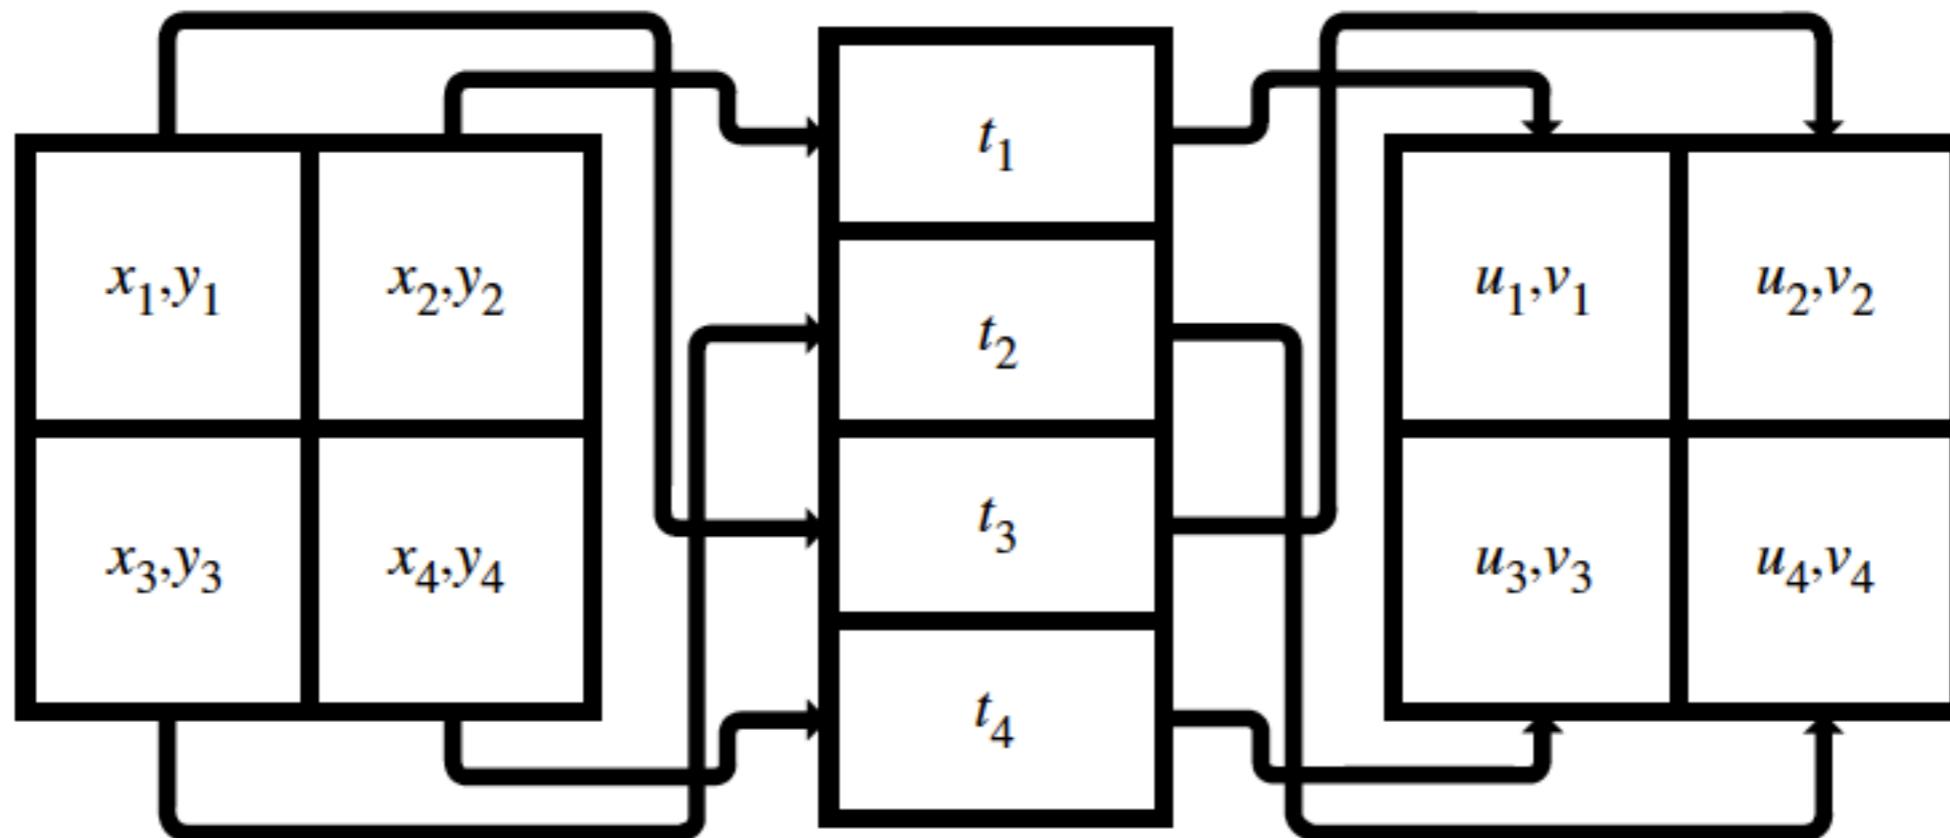


Monte Carlo (16 stratified samples)

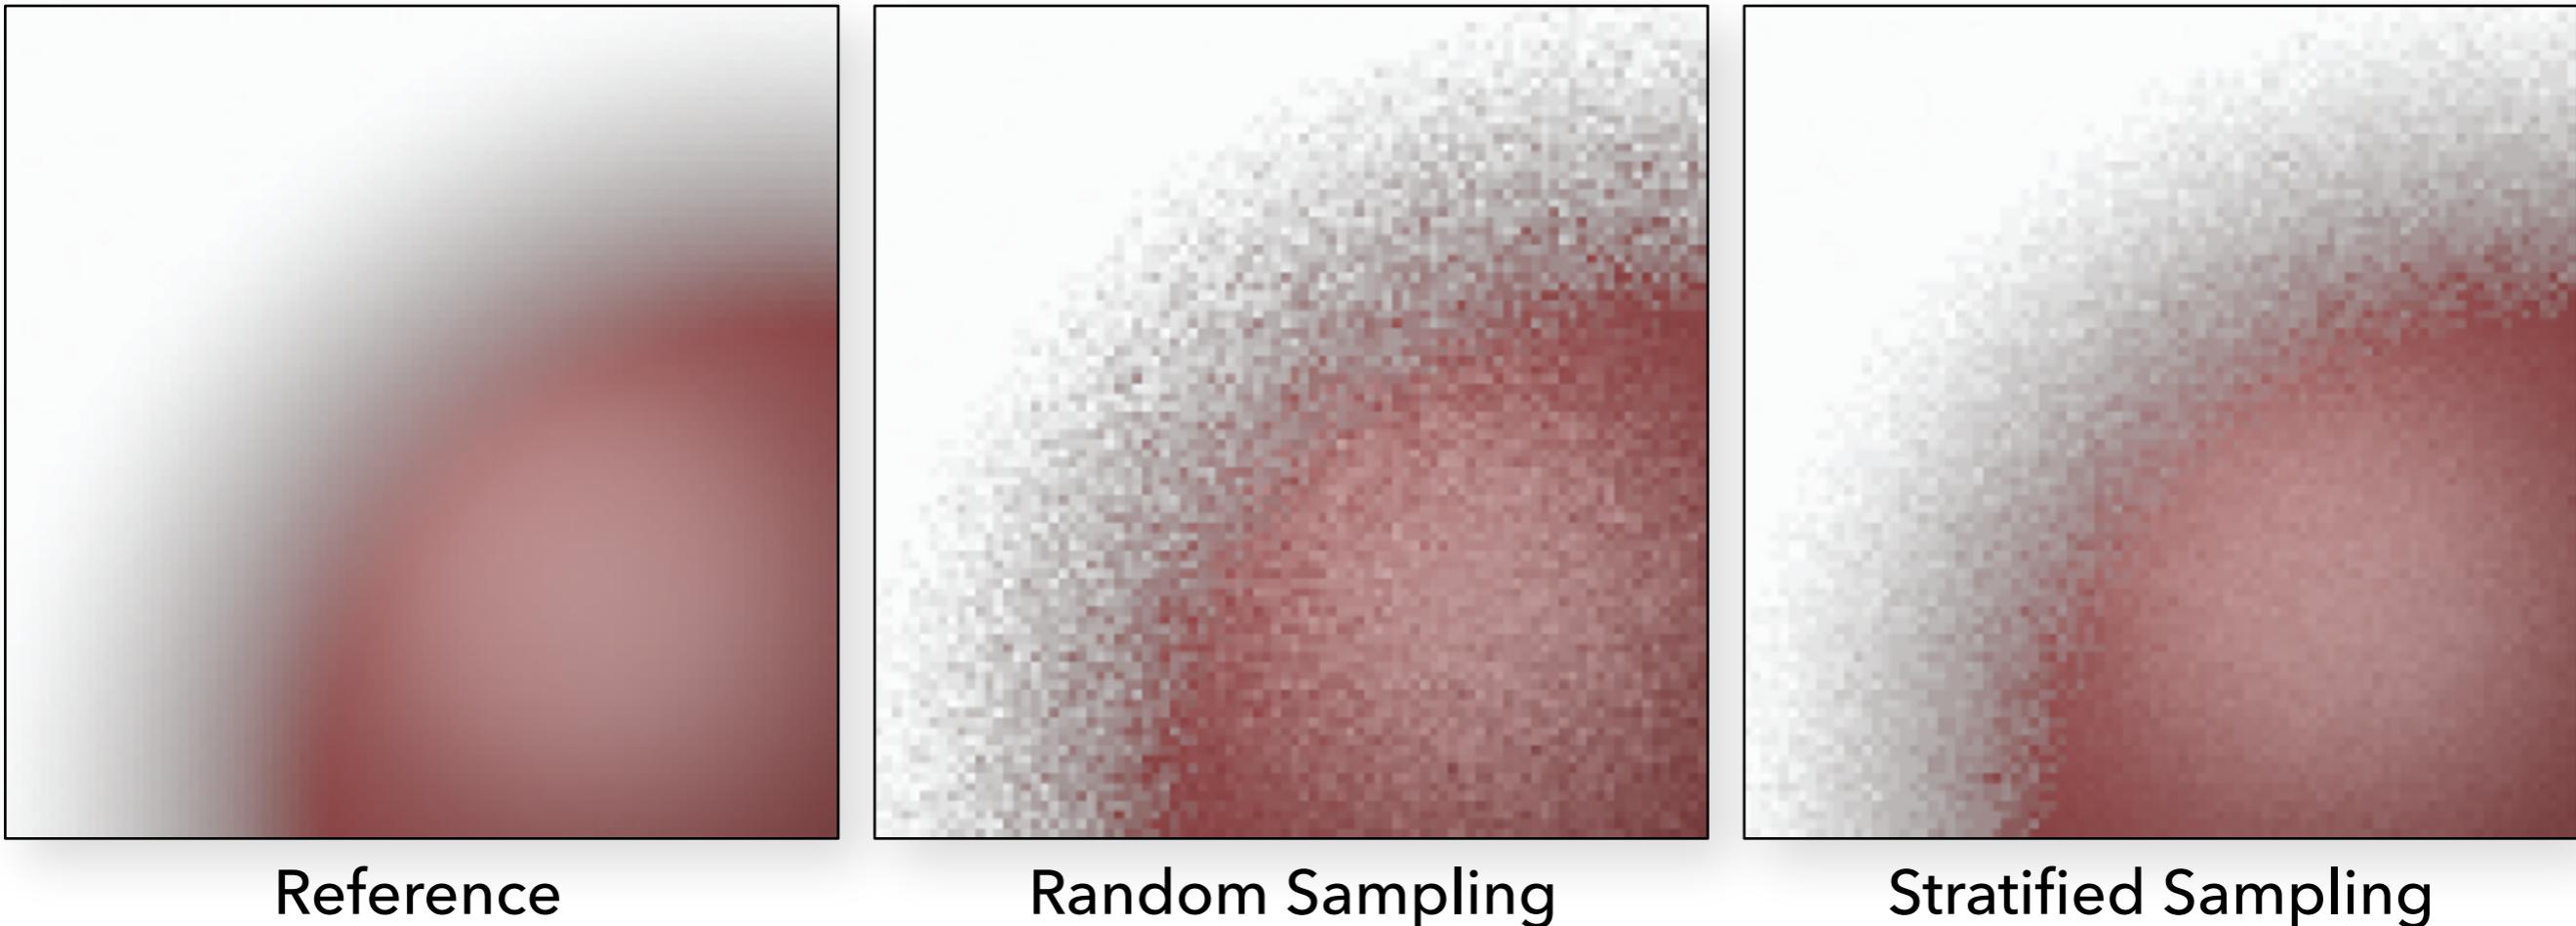


Stratifying in Higher Dimensions

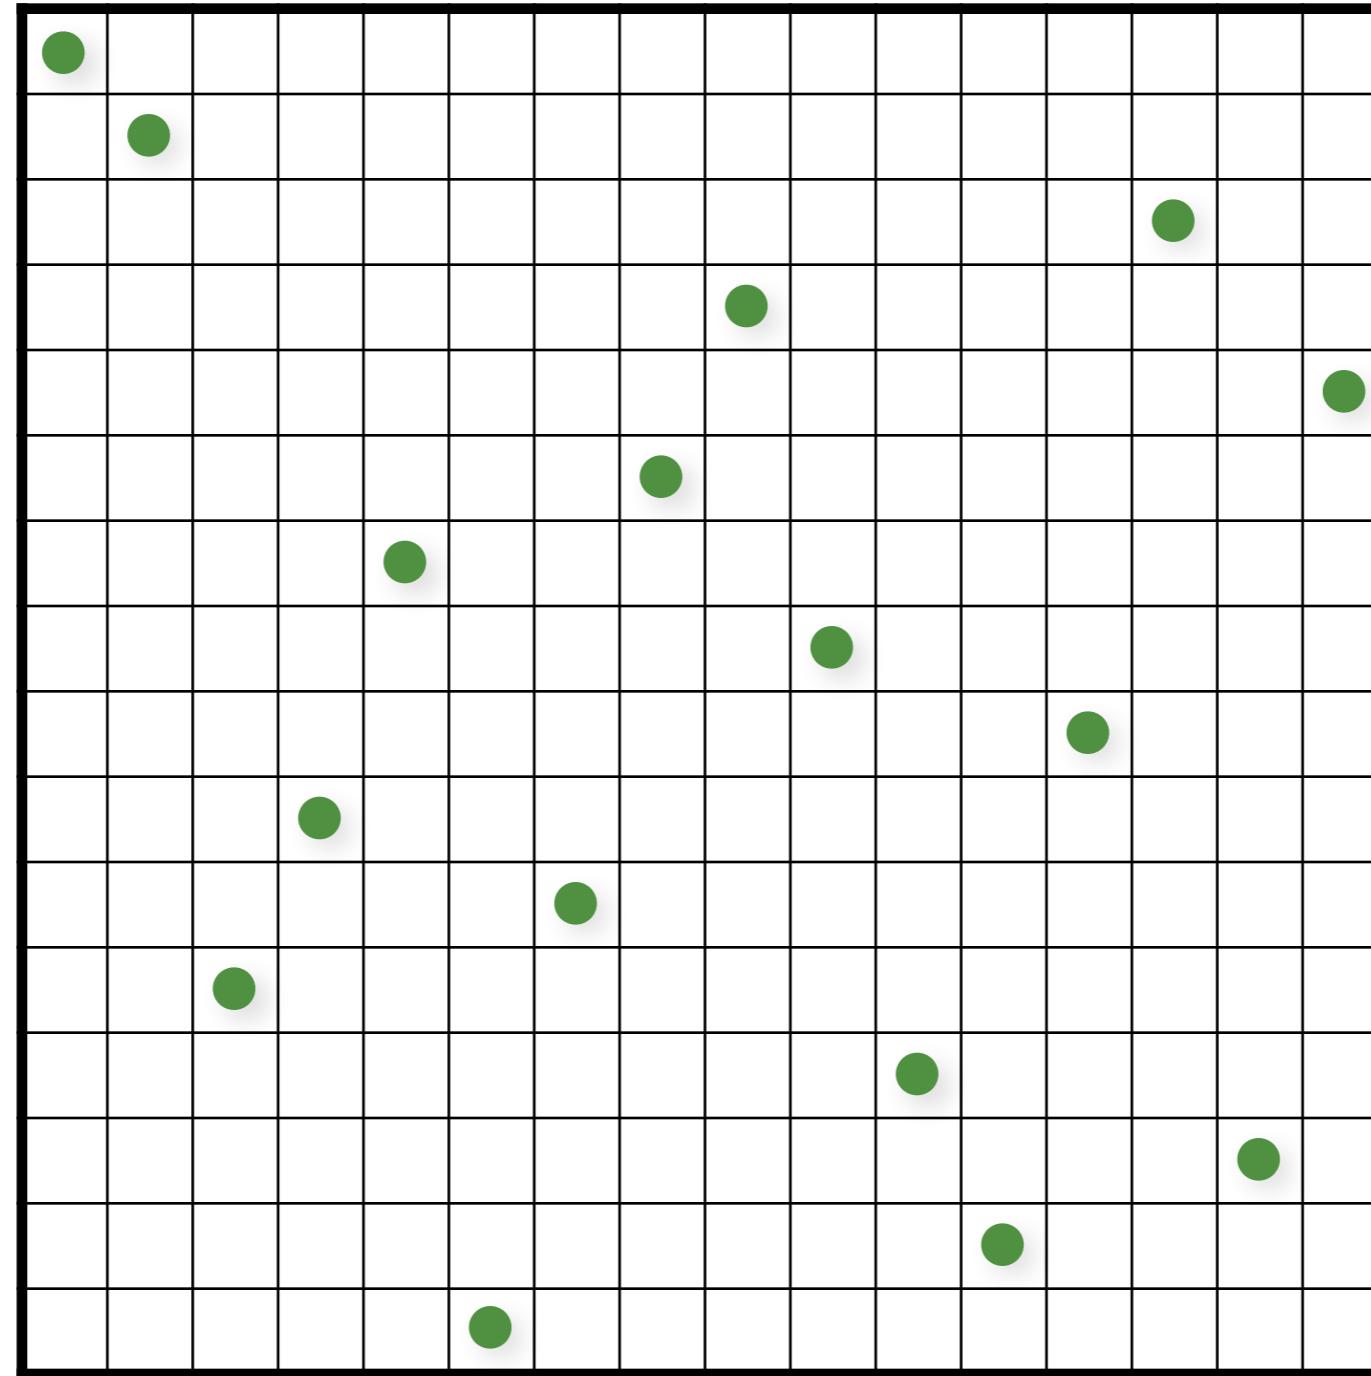
- Stratification requires N^d samples.
- Inconvenient for large d
- Compute stratified samples in lower-dimensions (e.g. (x,y) , t , (u,v)), and randomly combine



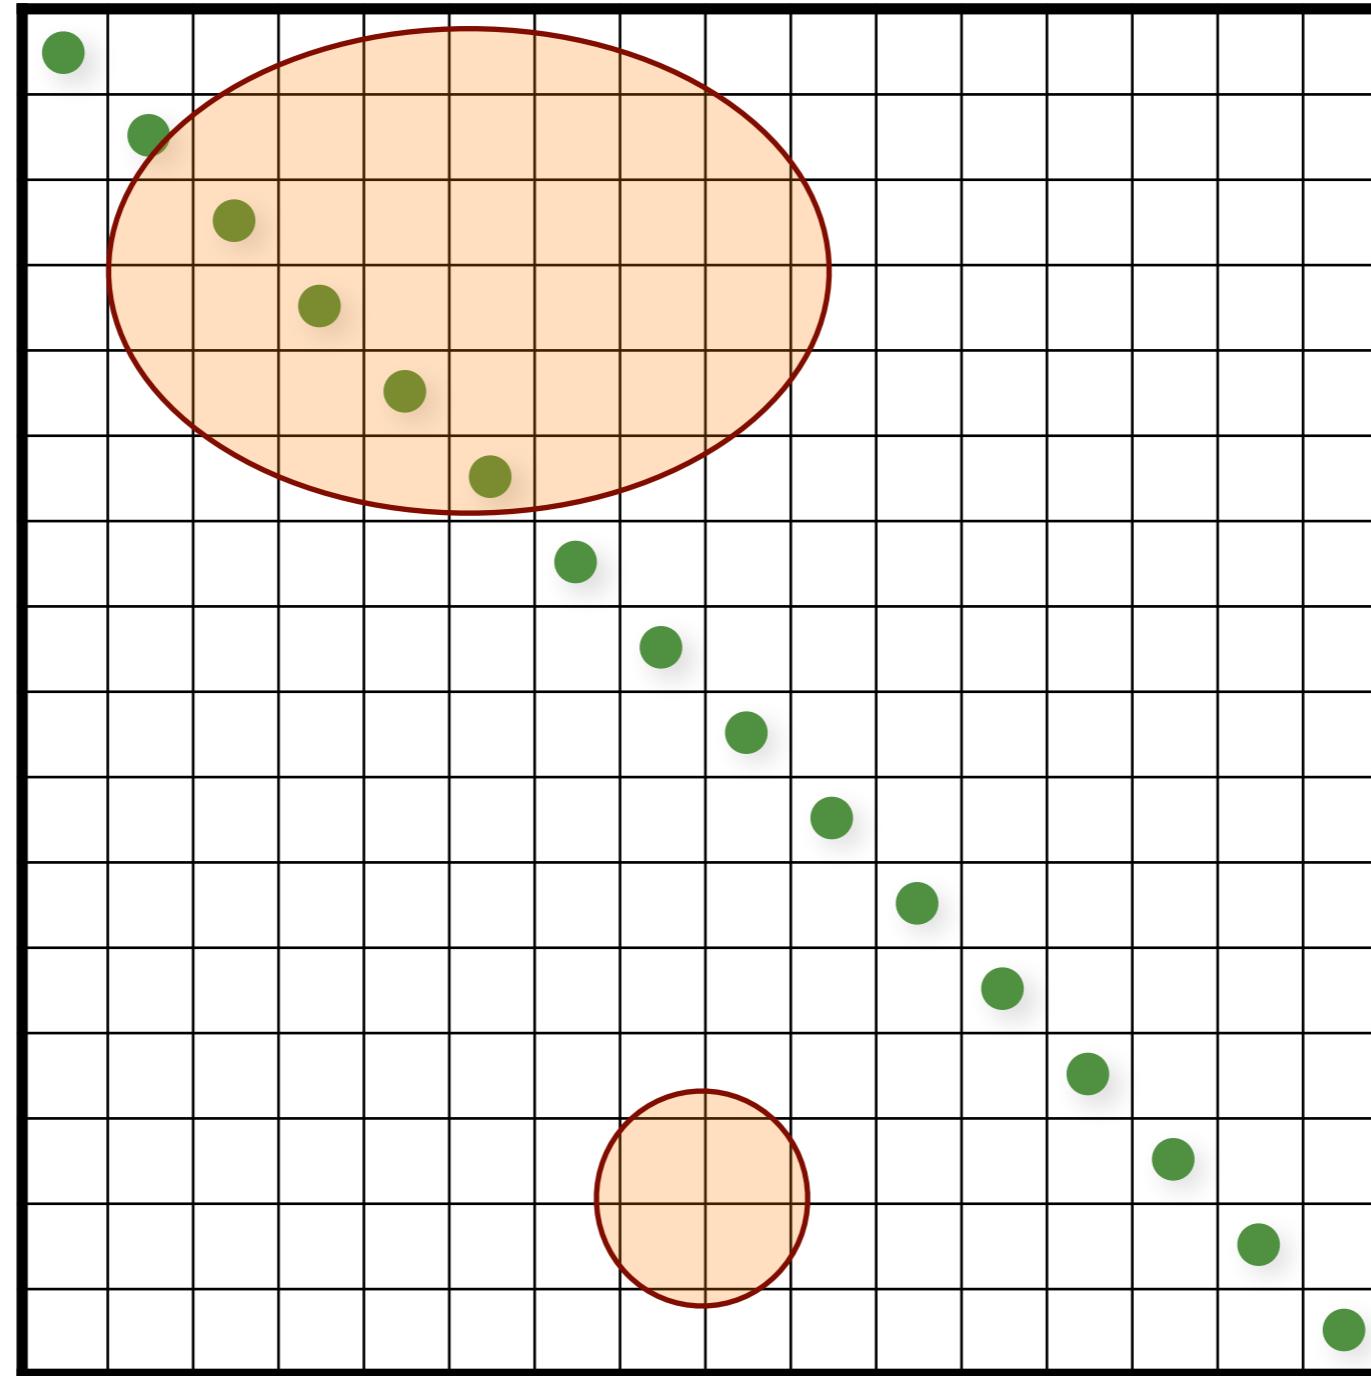
Depth of Field (4D)



Latin Hypercube (N-Rooks) Sampling

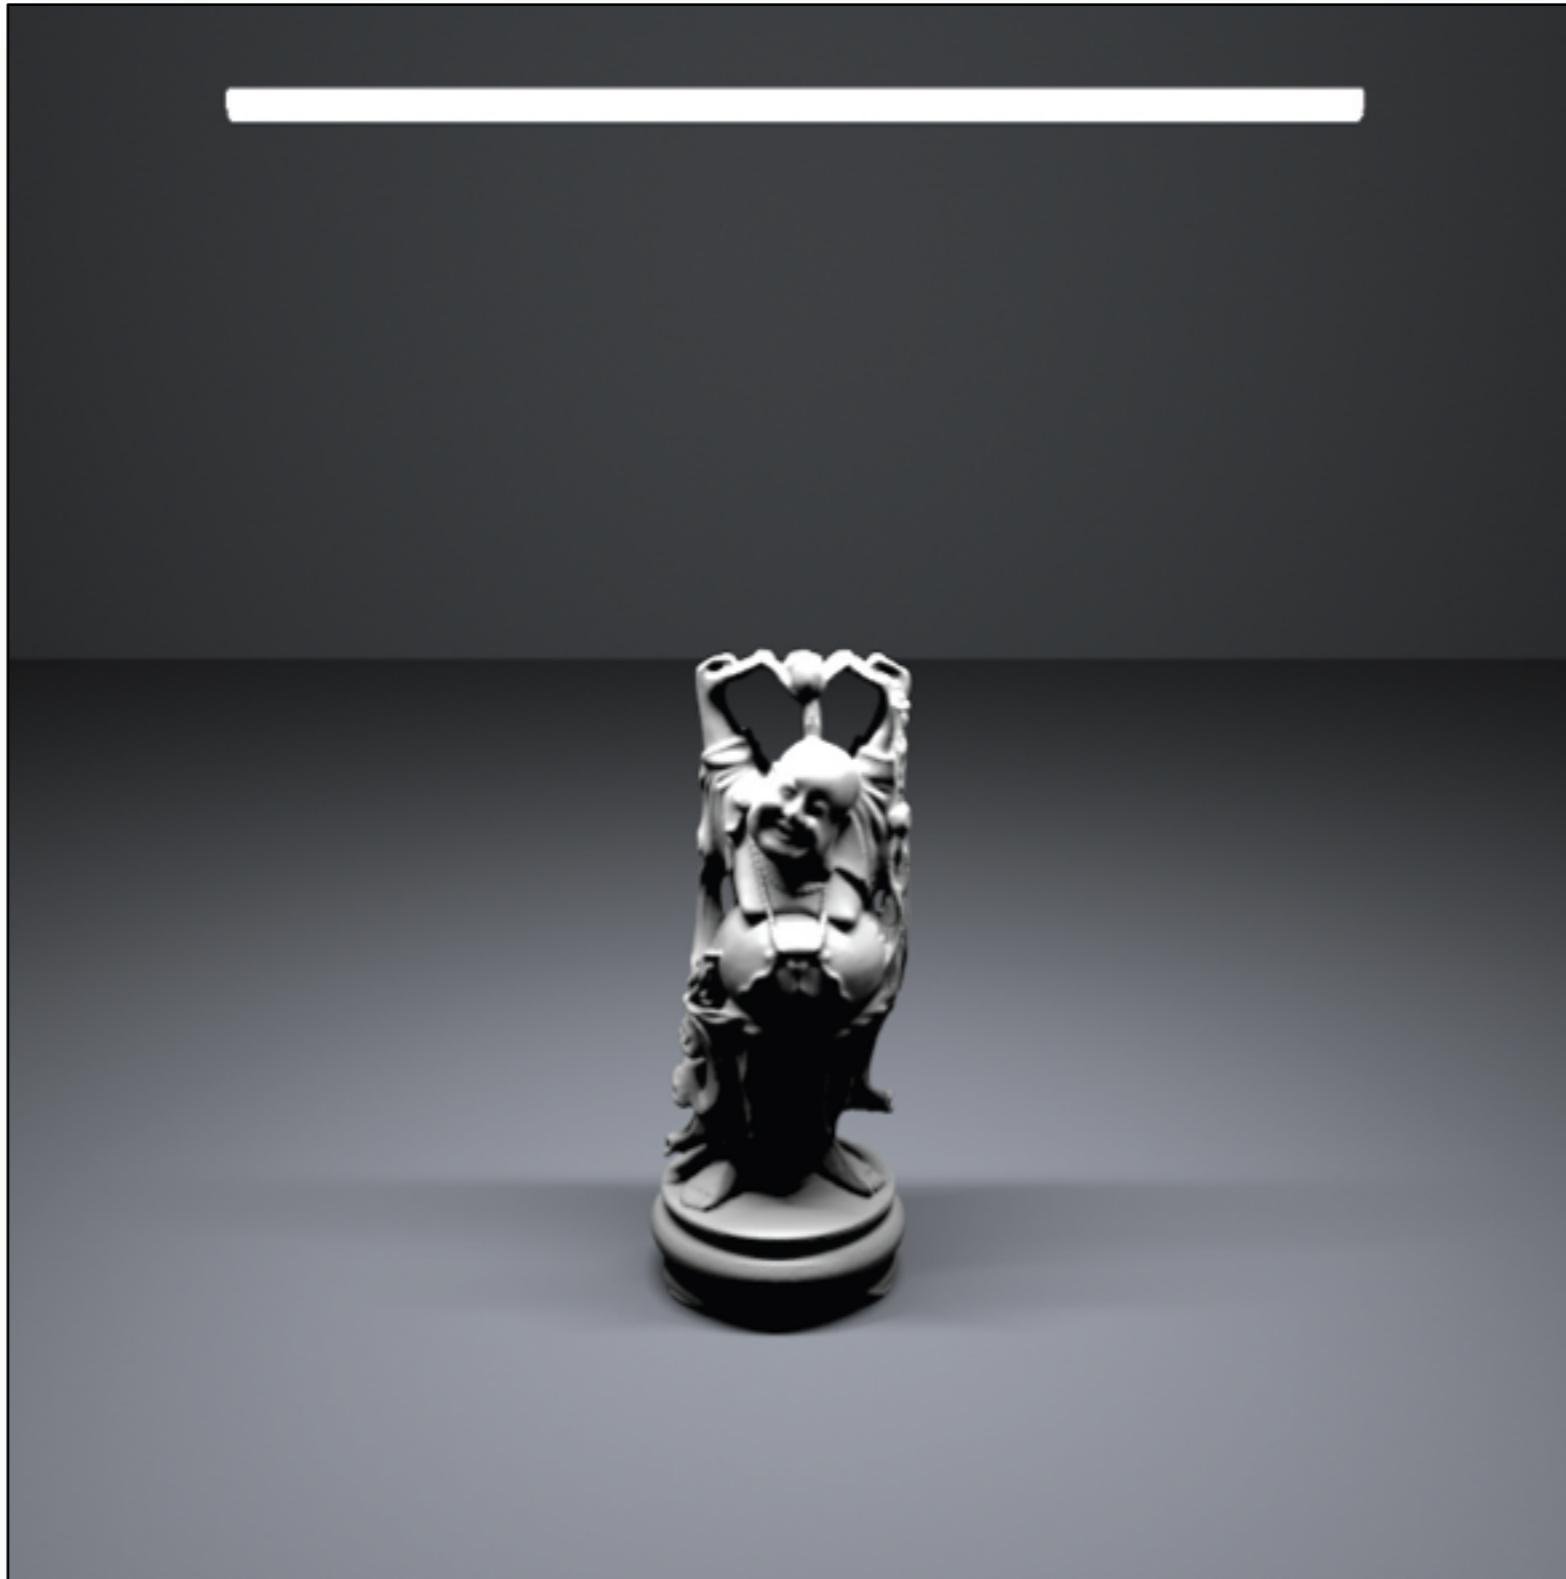


Latin Hypercube (N-Rooks) Sampling



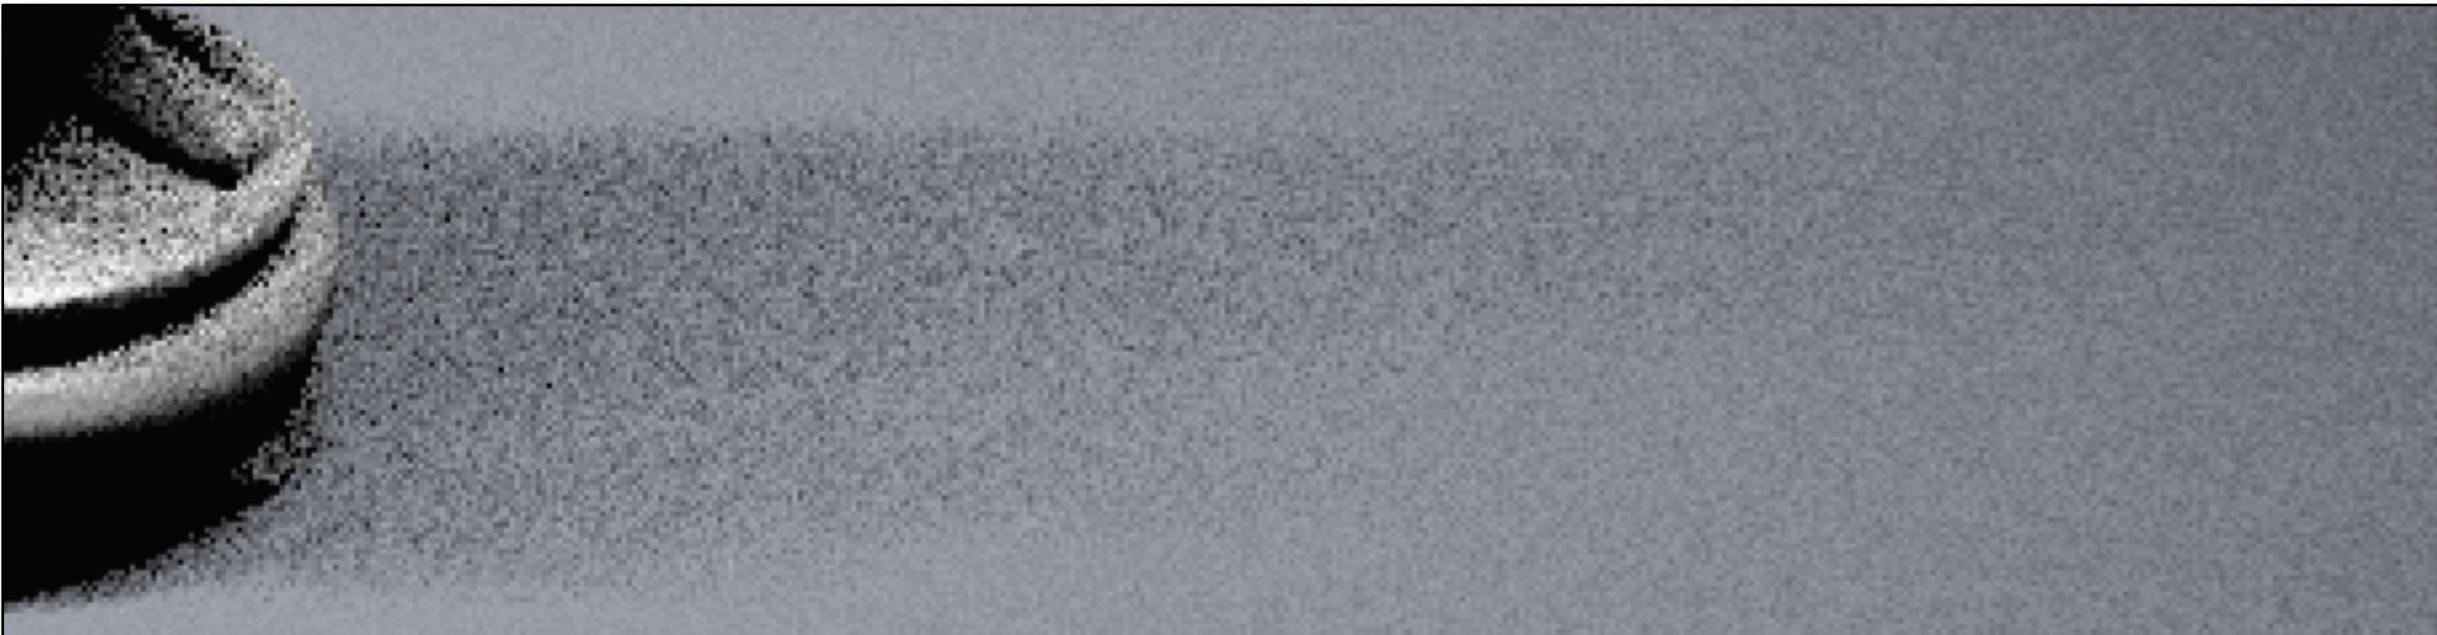
Shuffled rows

Soft Shadows



Soft Shadows: Sample Allocation

1 image sample x 16 light samples

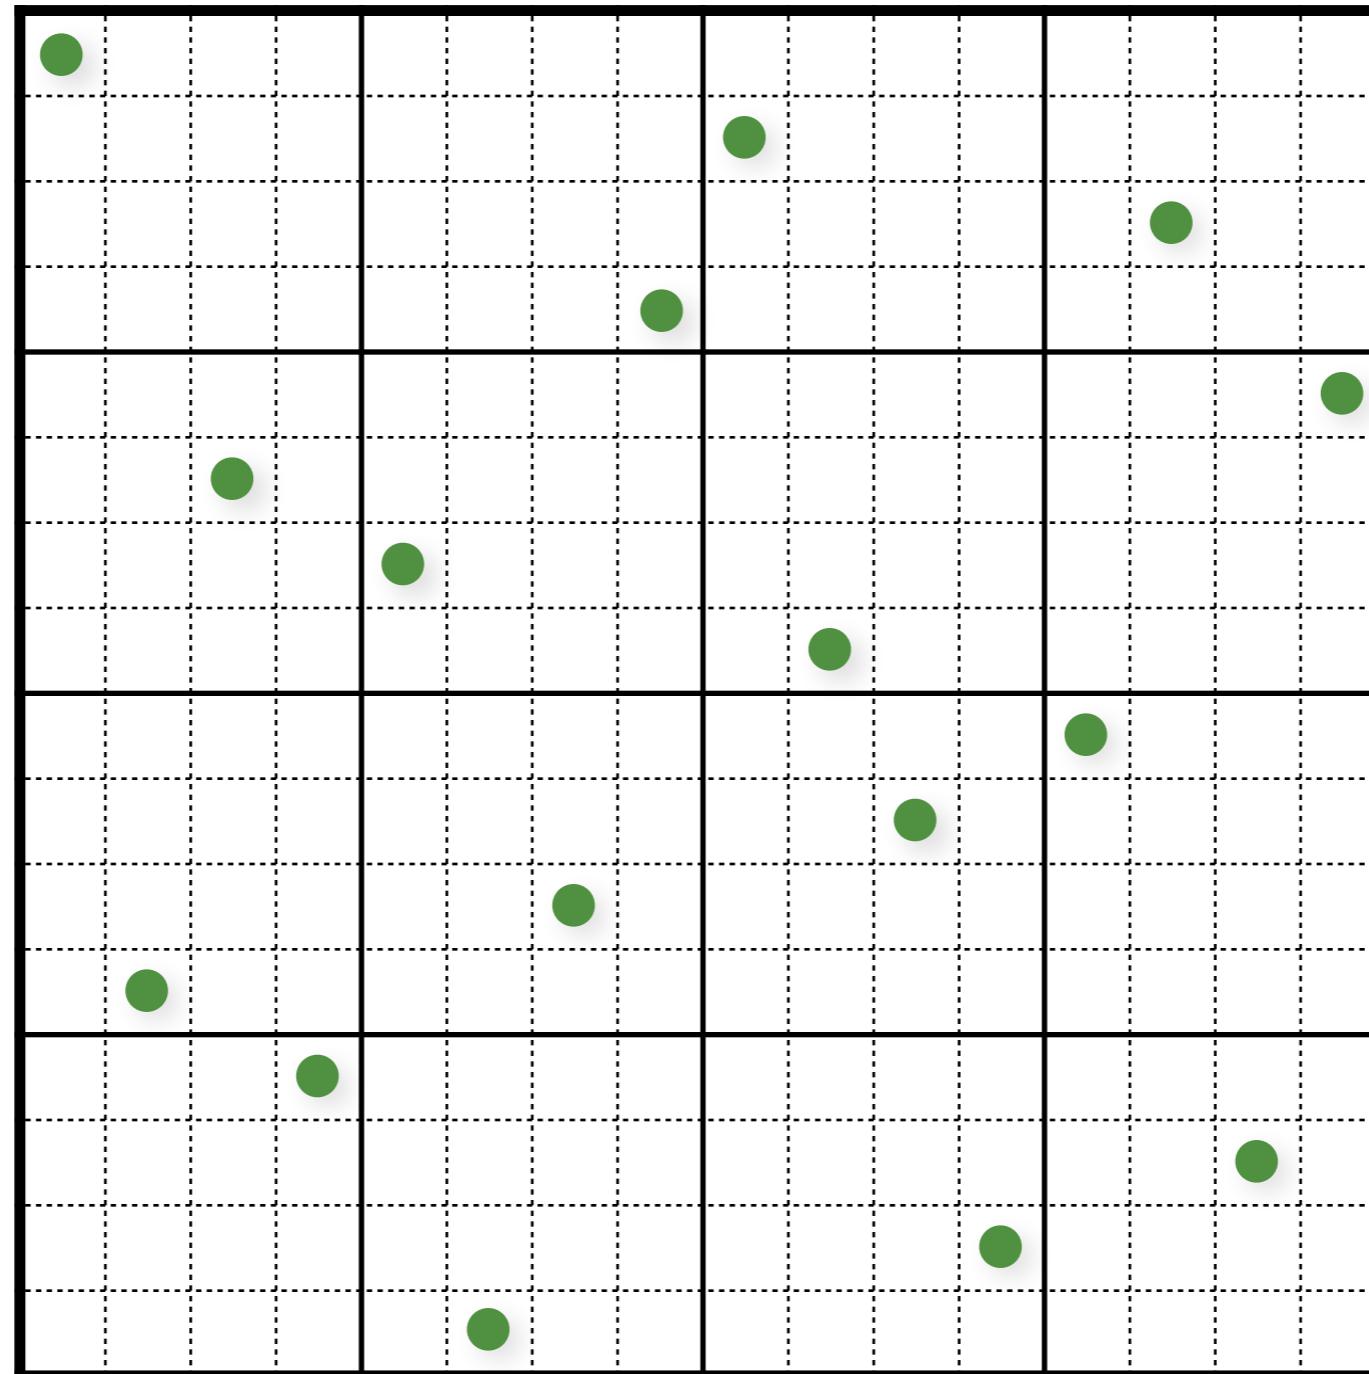


16 image samples x 1 light sample

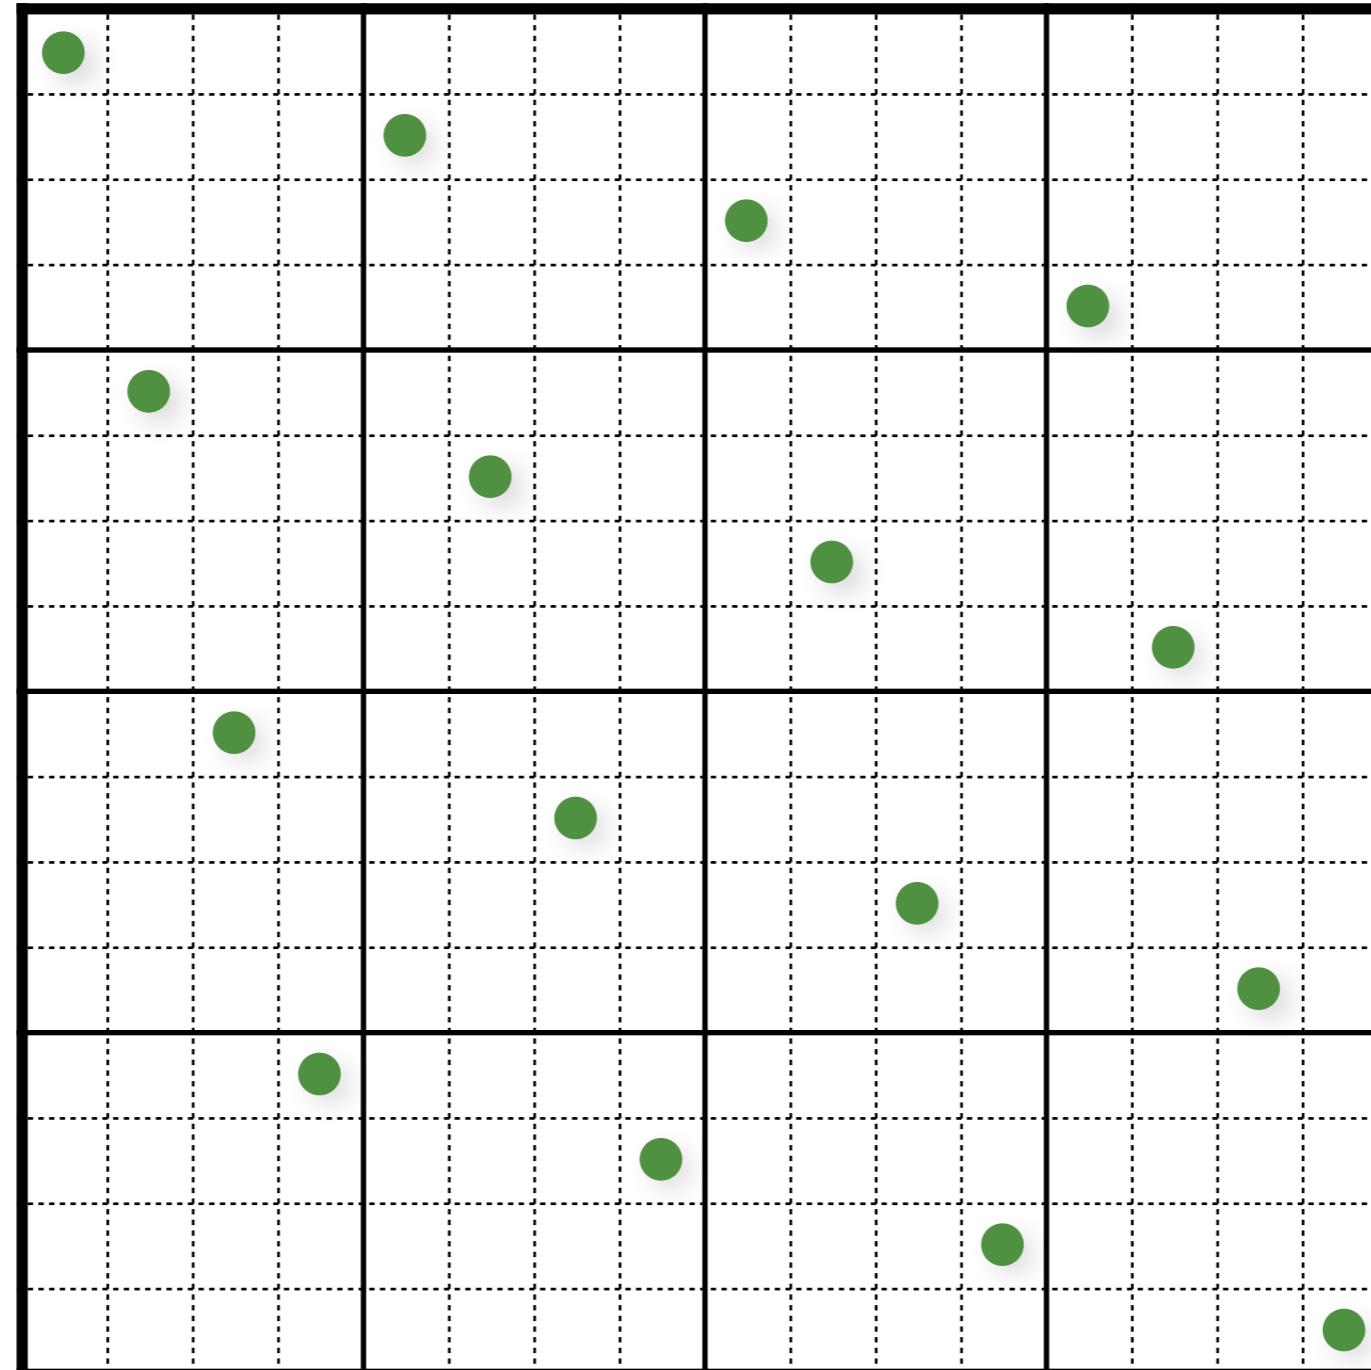
Multi-Jittered Sampling

- Kenneth Chiu, Peter Shirley, and Changyaw Wang. “Multi-jittered sampling.” In *Graphics Gems IV*, pp. 370-374. Academic Press, May 1994.
 - combine N-Rooks and Jittered stratification constraints

Multi-Jittered Sampling

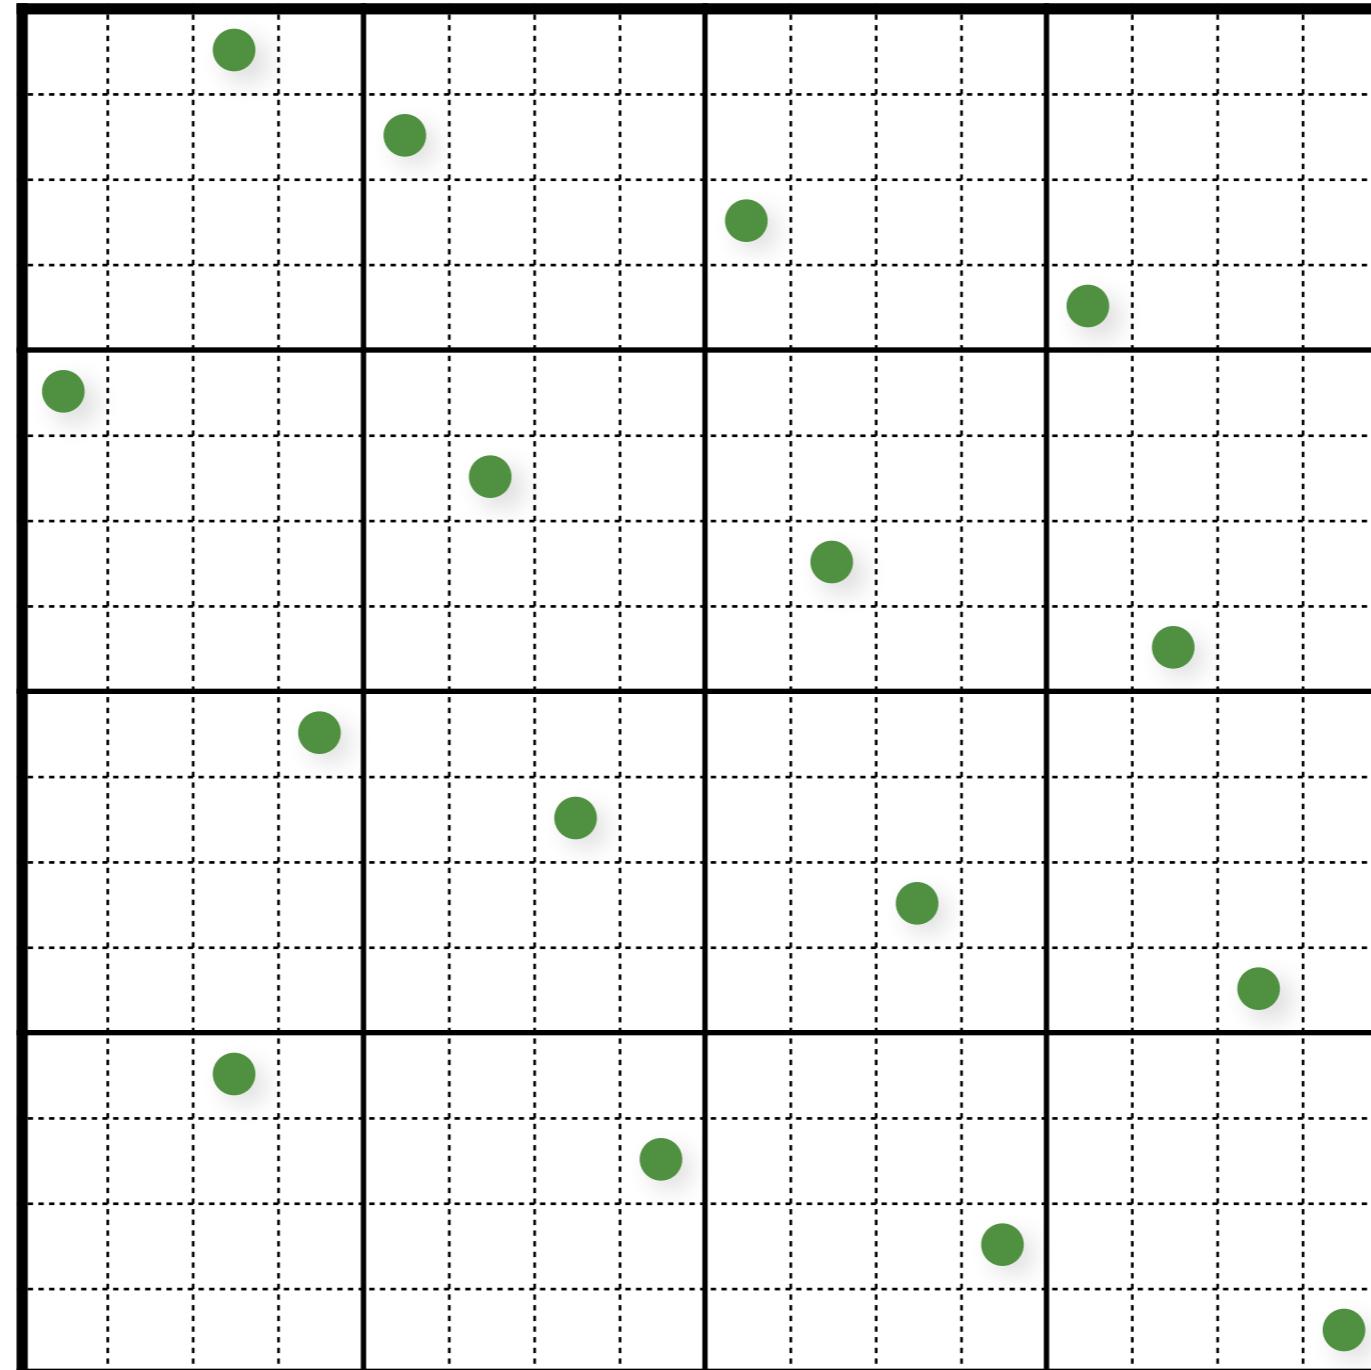


Multi-Jittered Sampling



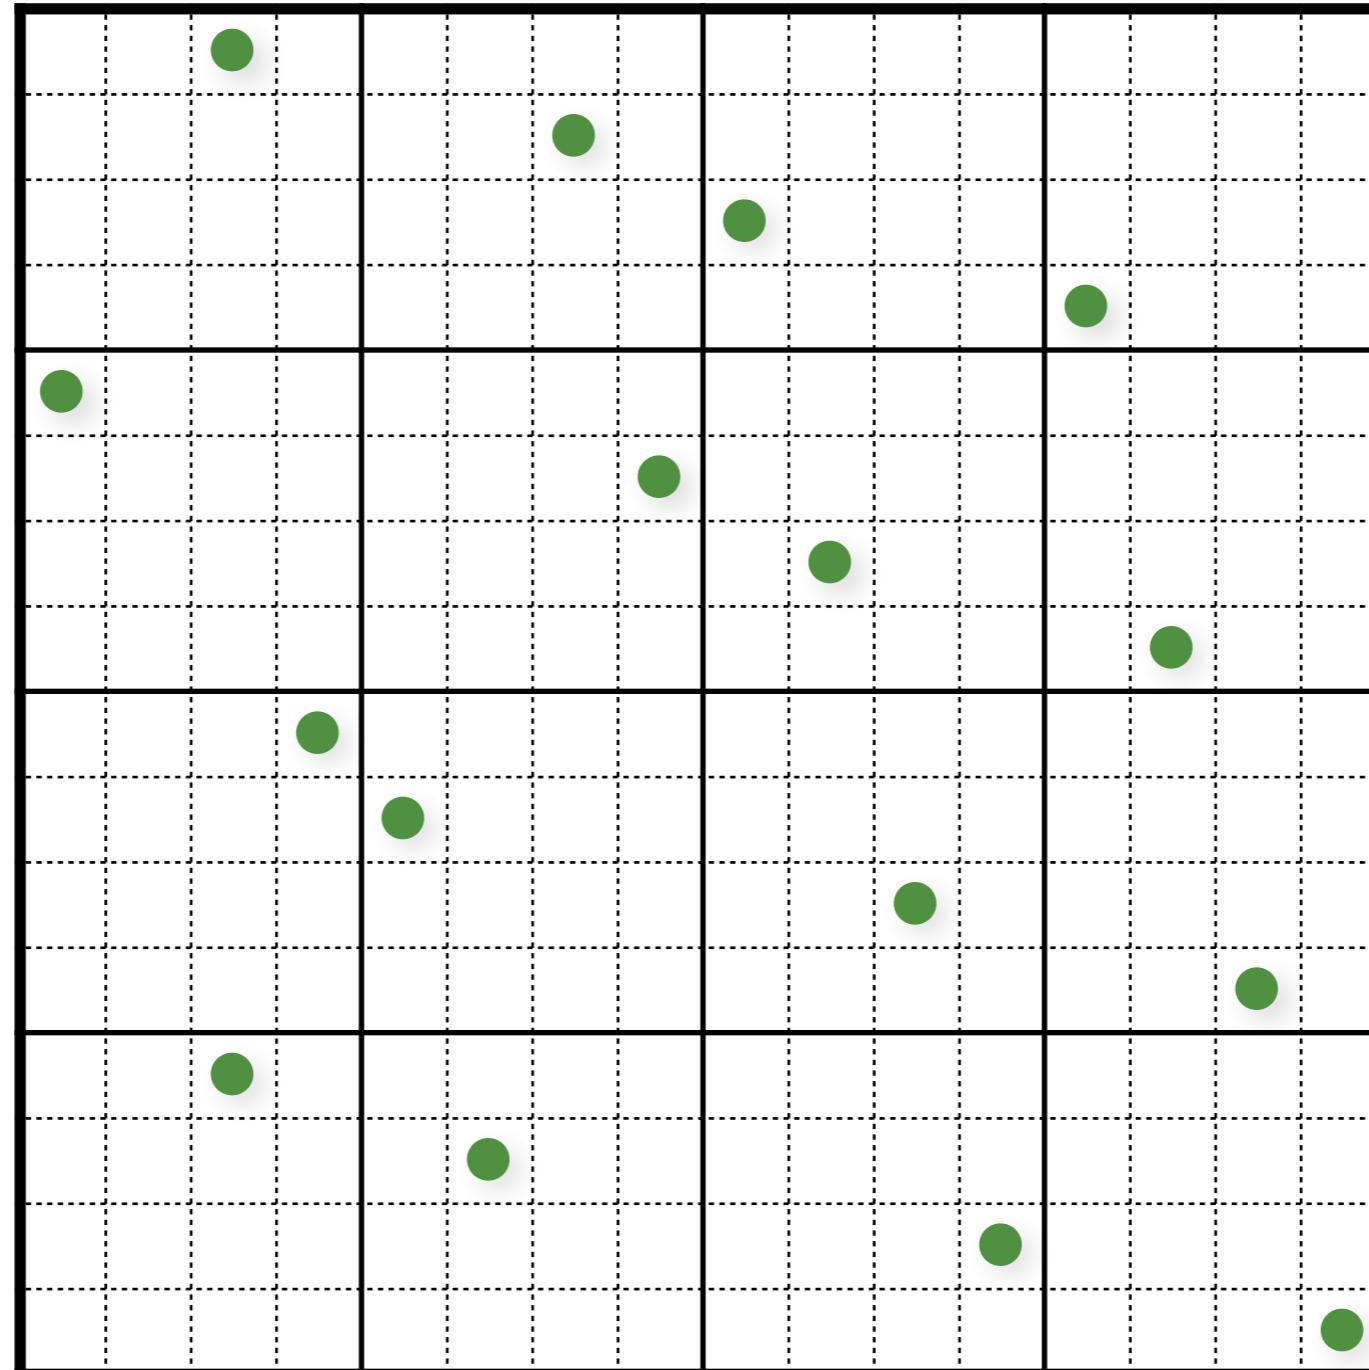
Shuffled coordinates

Multi-Jittered Sampling



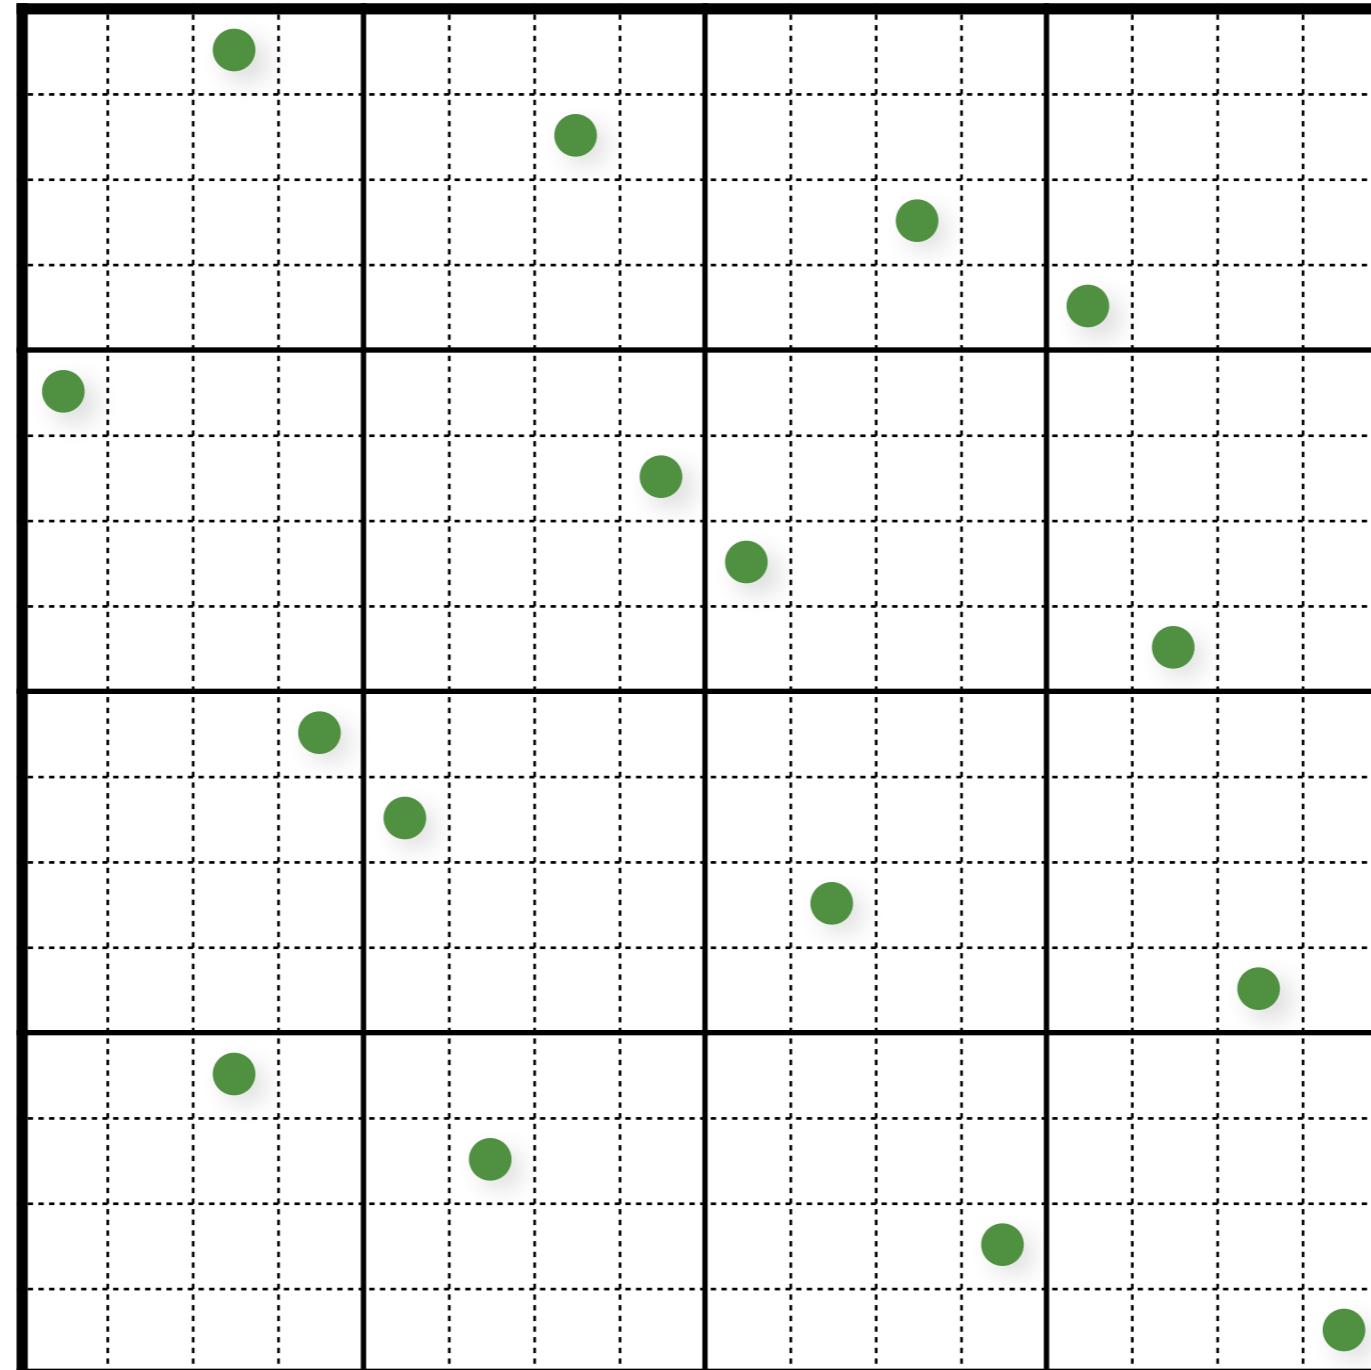
Shuffle x-coords

Multi-Jittered Sampling



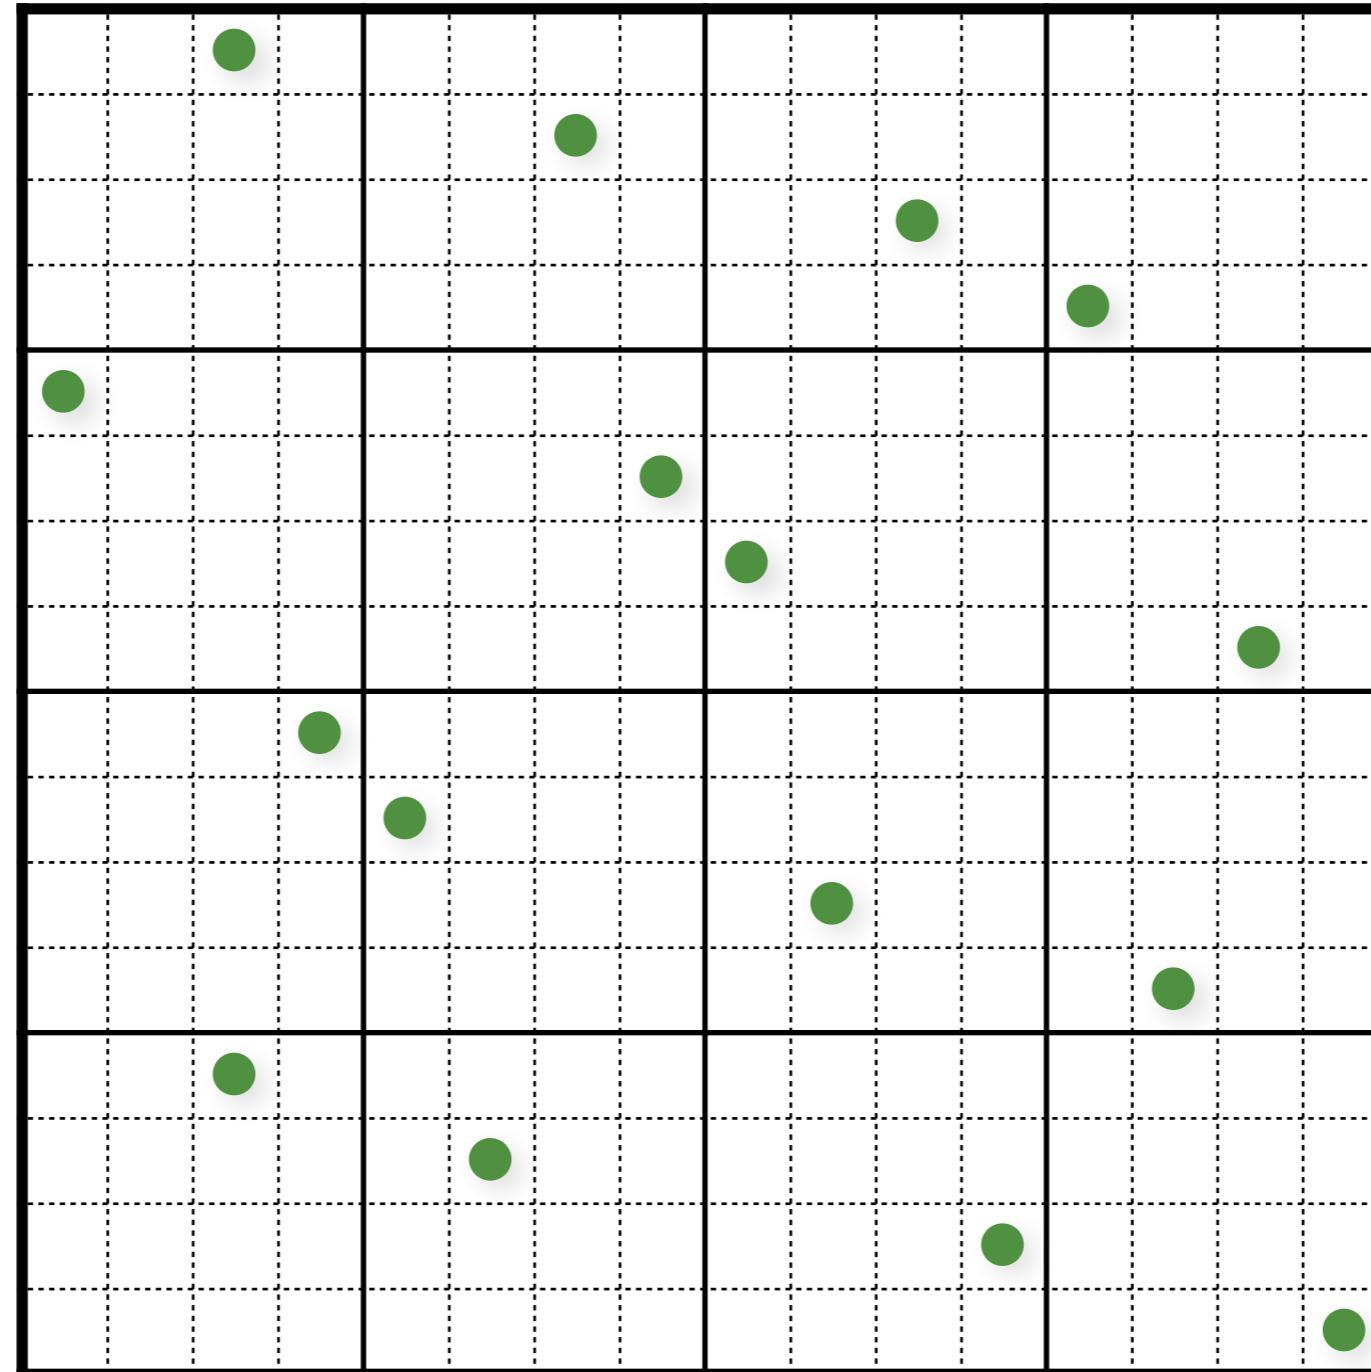
Shuffle x-coords

Multi-Jittered Sampling



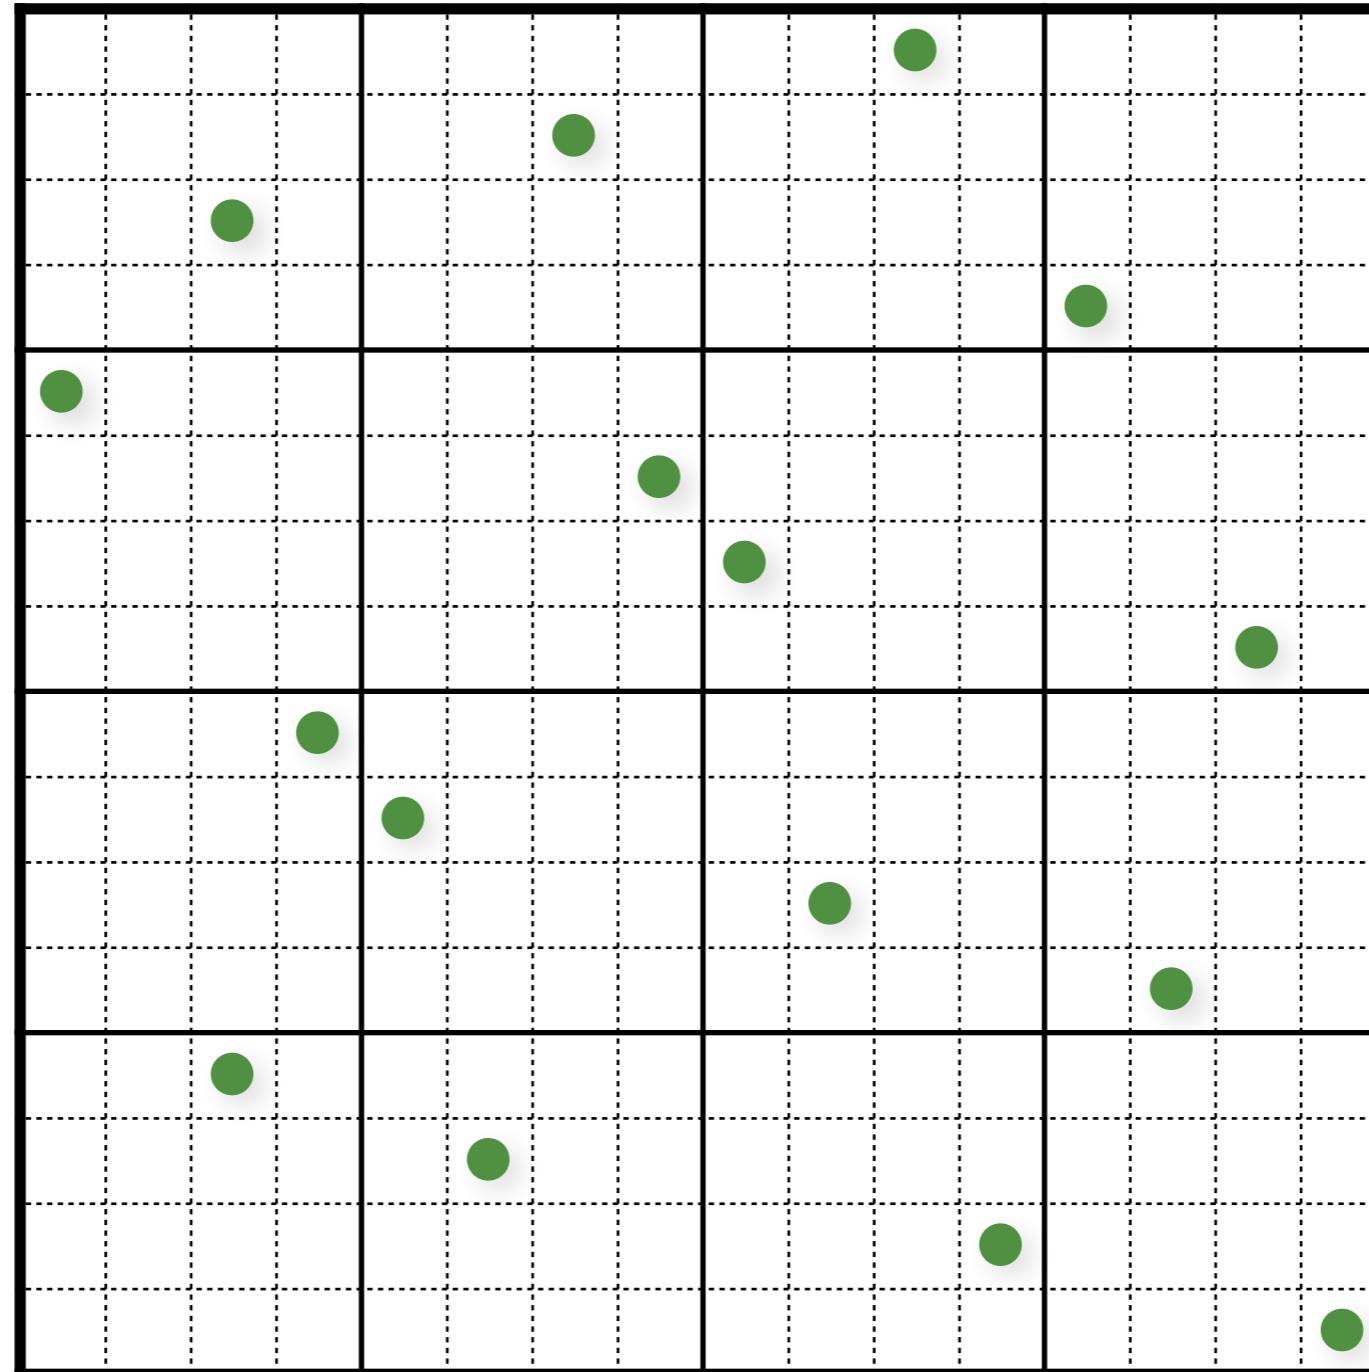
Shuffle x-coords

Multi-Jittered Sampling



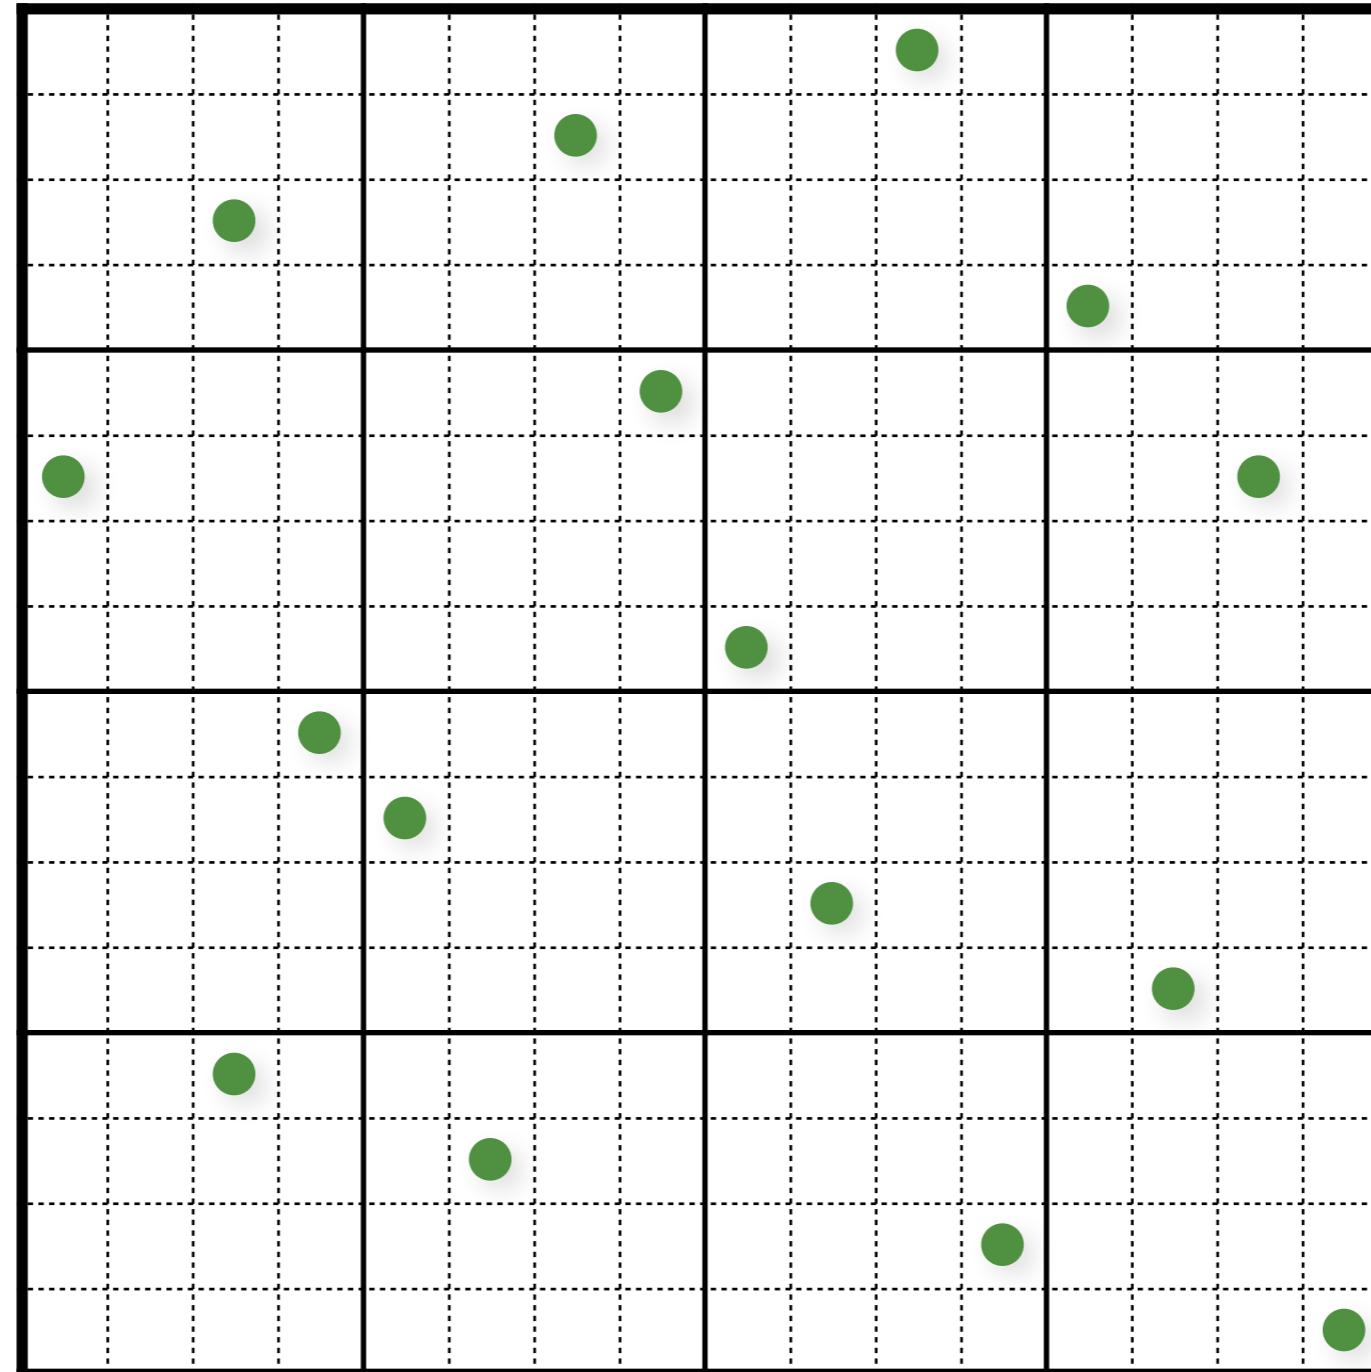
Shuffle x-coords

Multi-Jittered Sampling



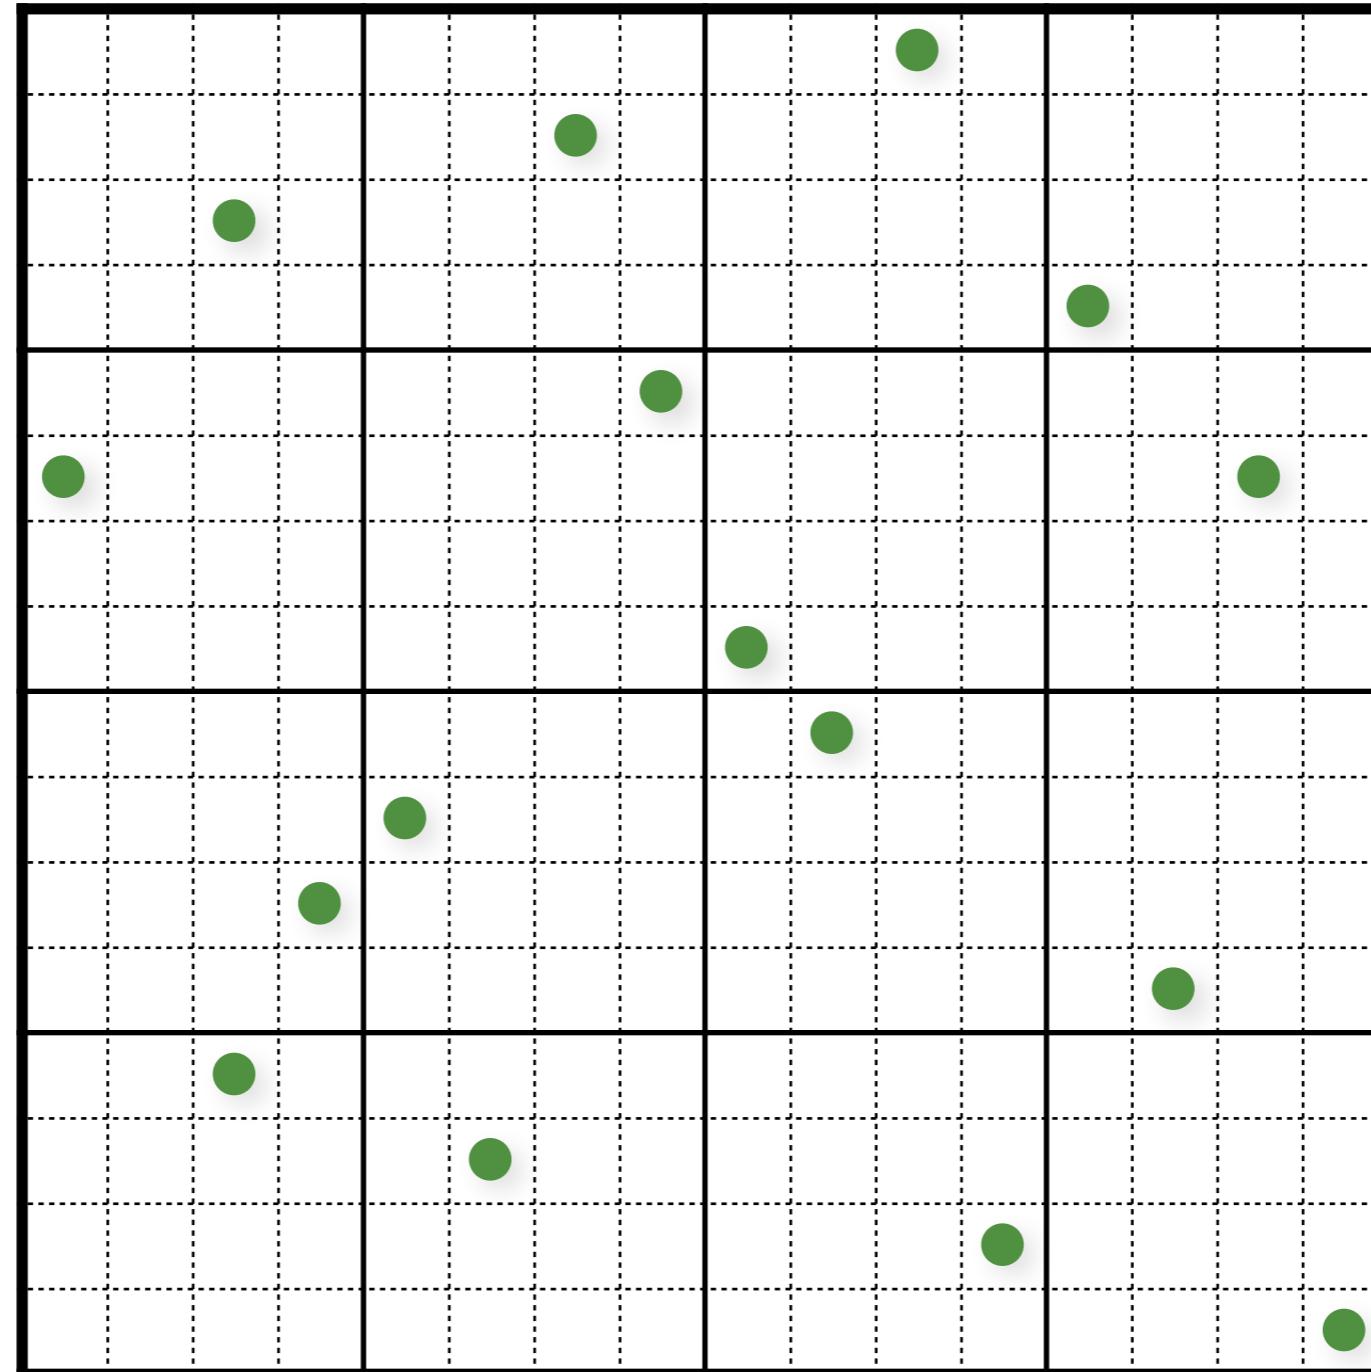
Shuffle y-coords

Multi-Jittered Sampling



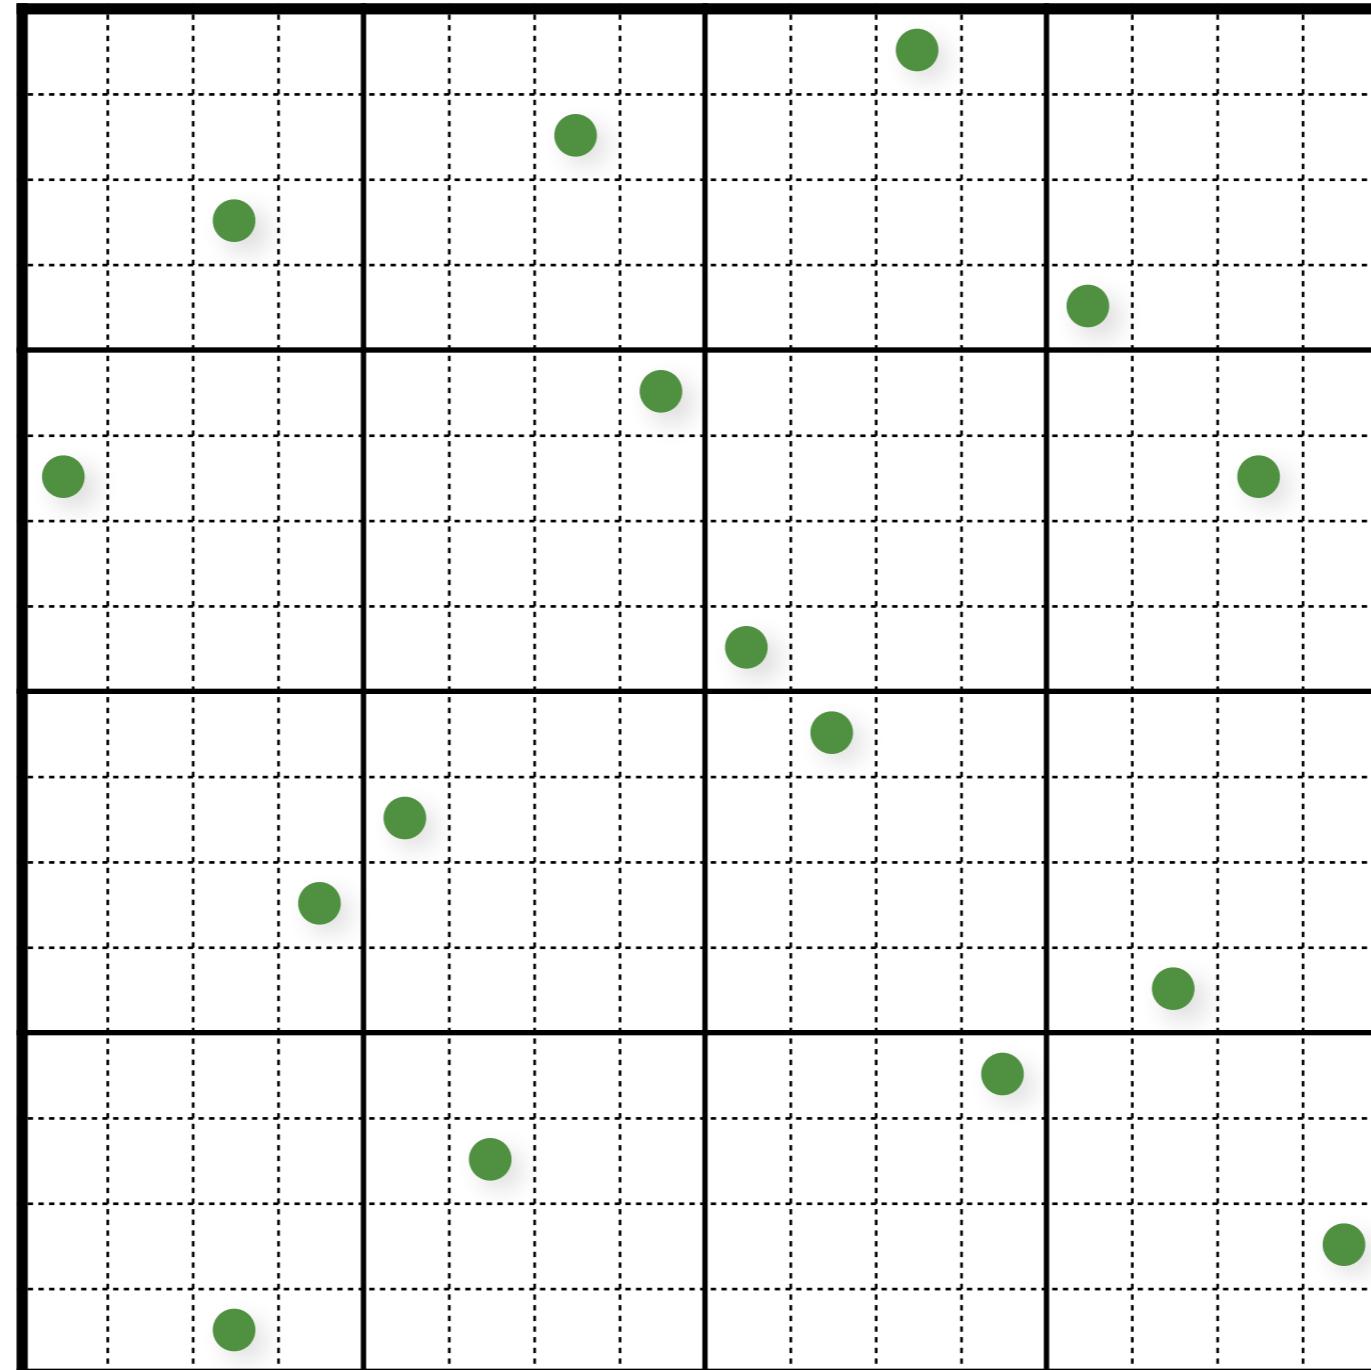
Shuffle y-coords

Multi-Jittered Sampling



Shuffle y-coords

Multi-Jittered Sampling



Shuffle y-coords

Visual Break

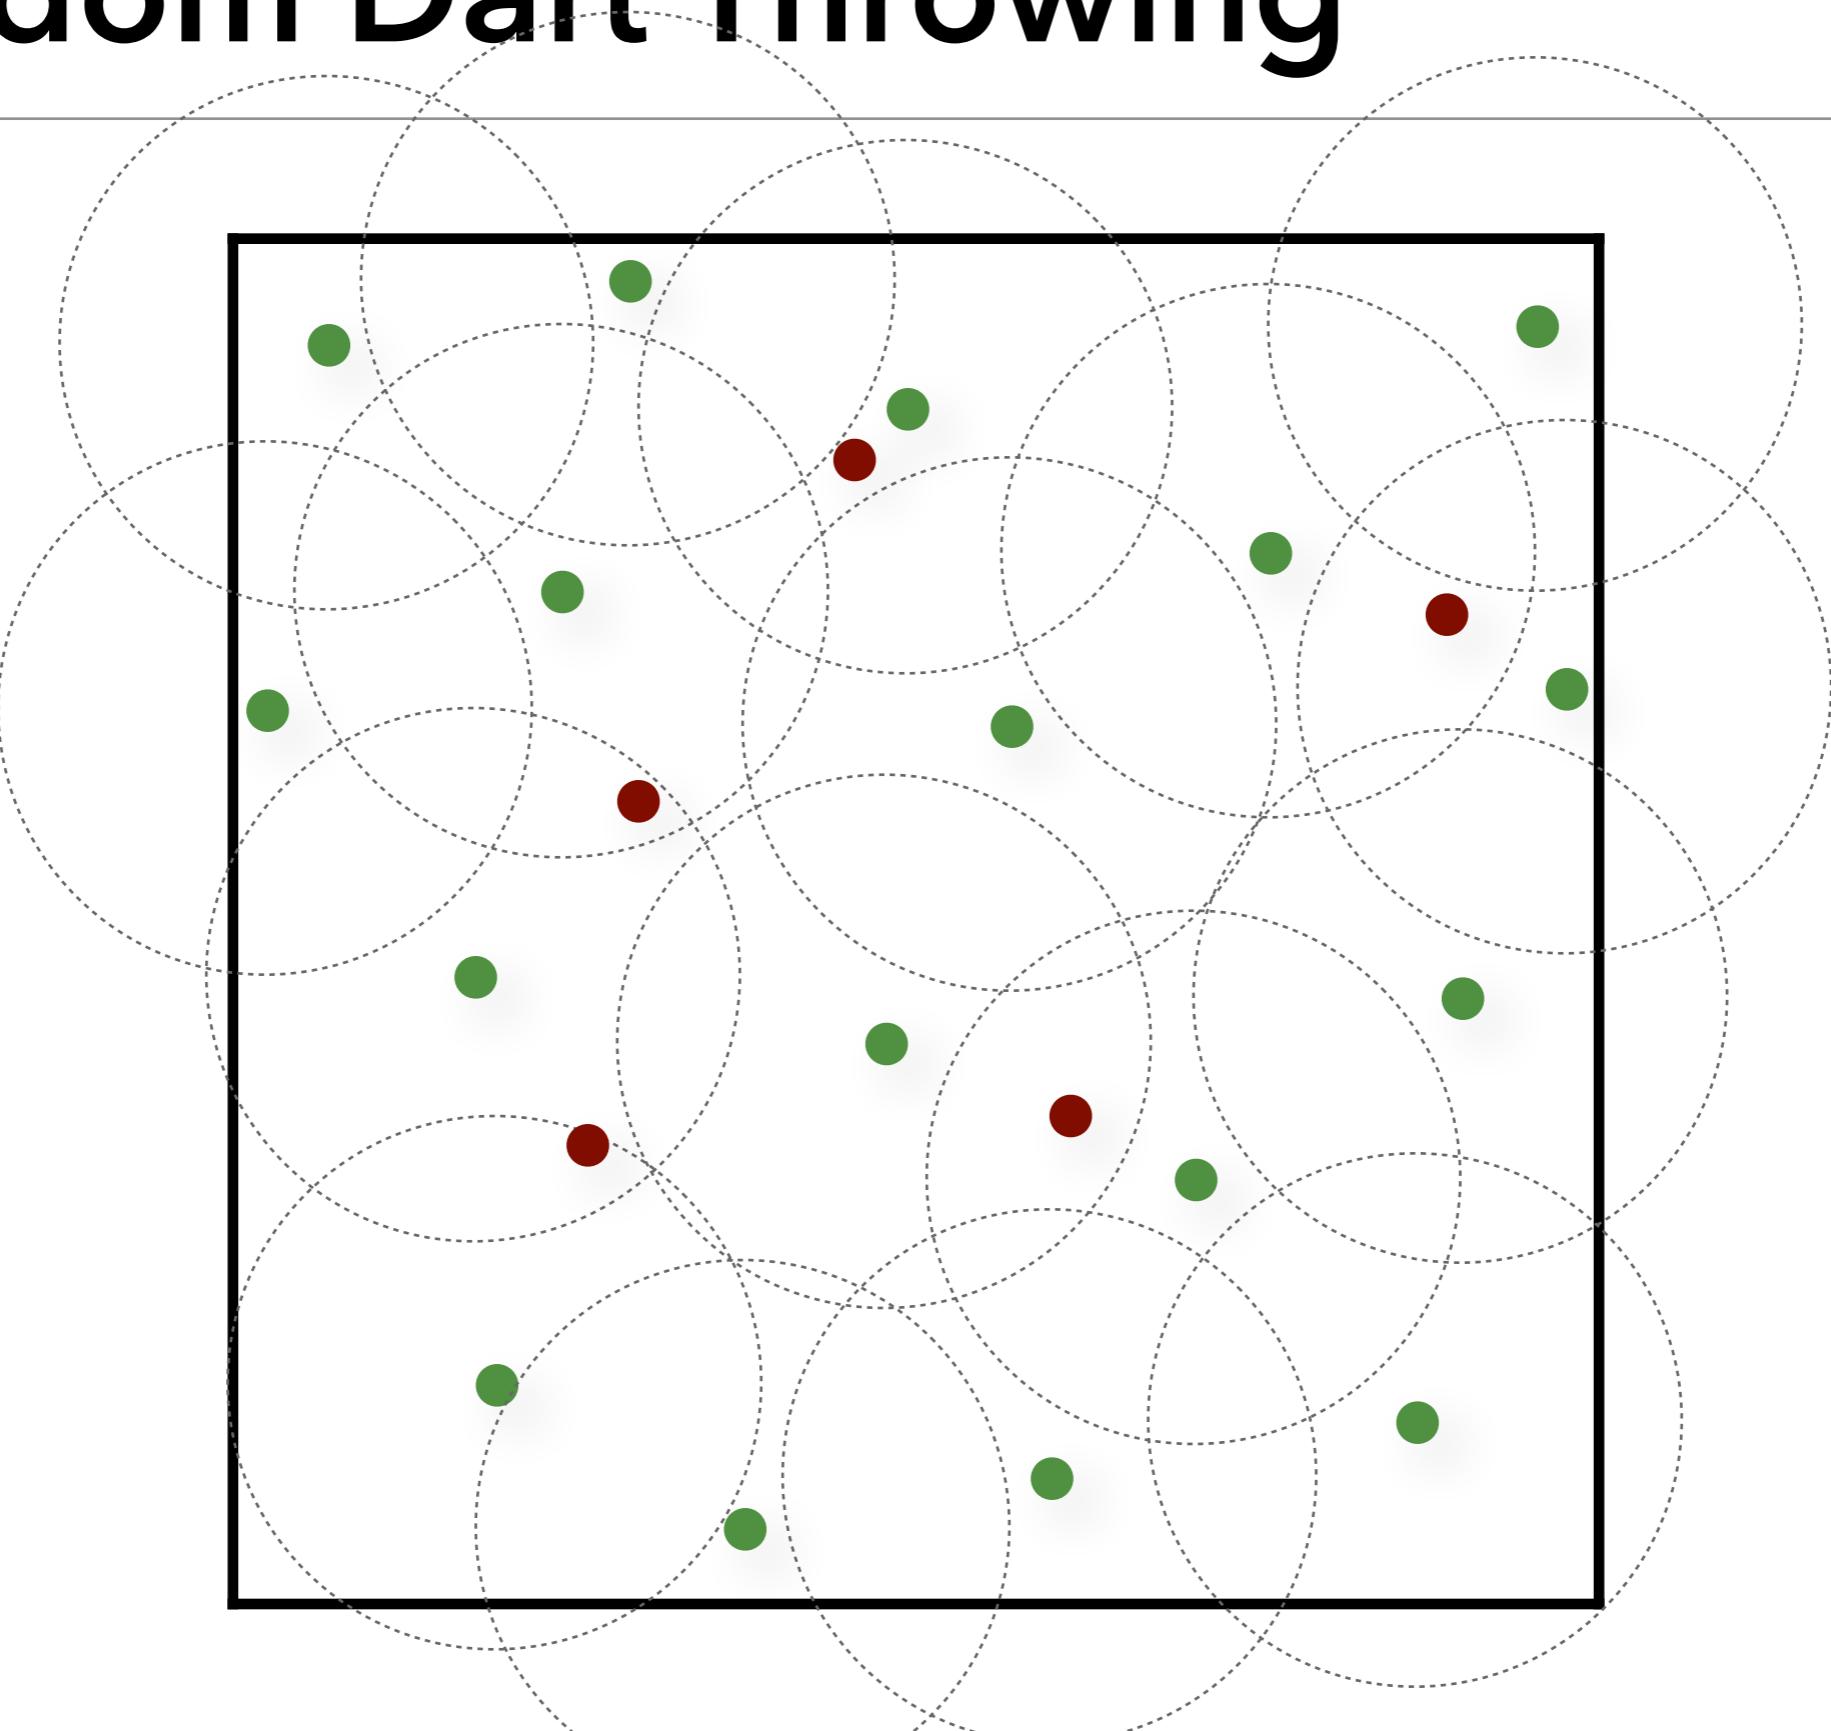


© ReinOwader

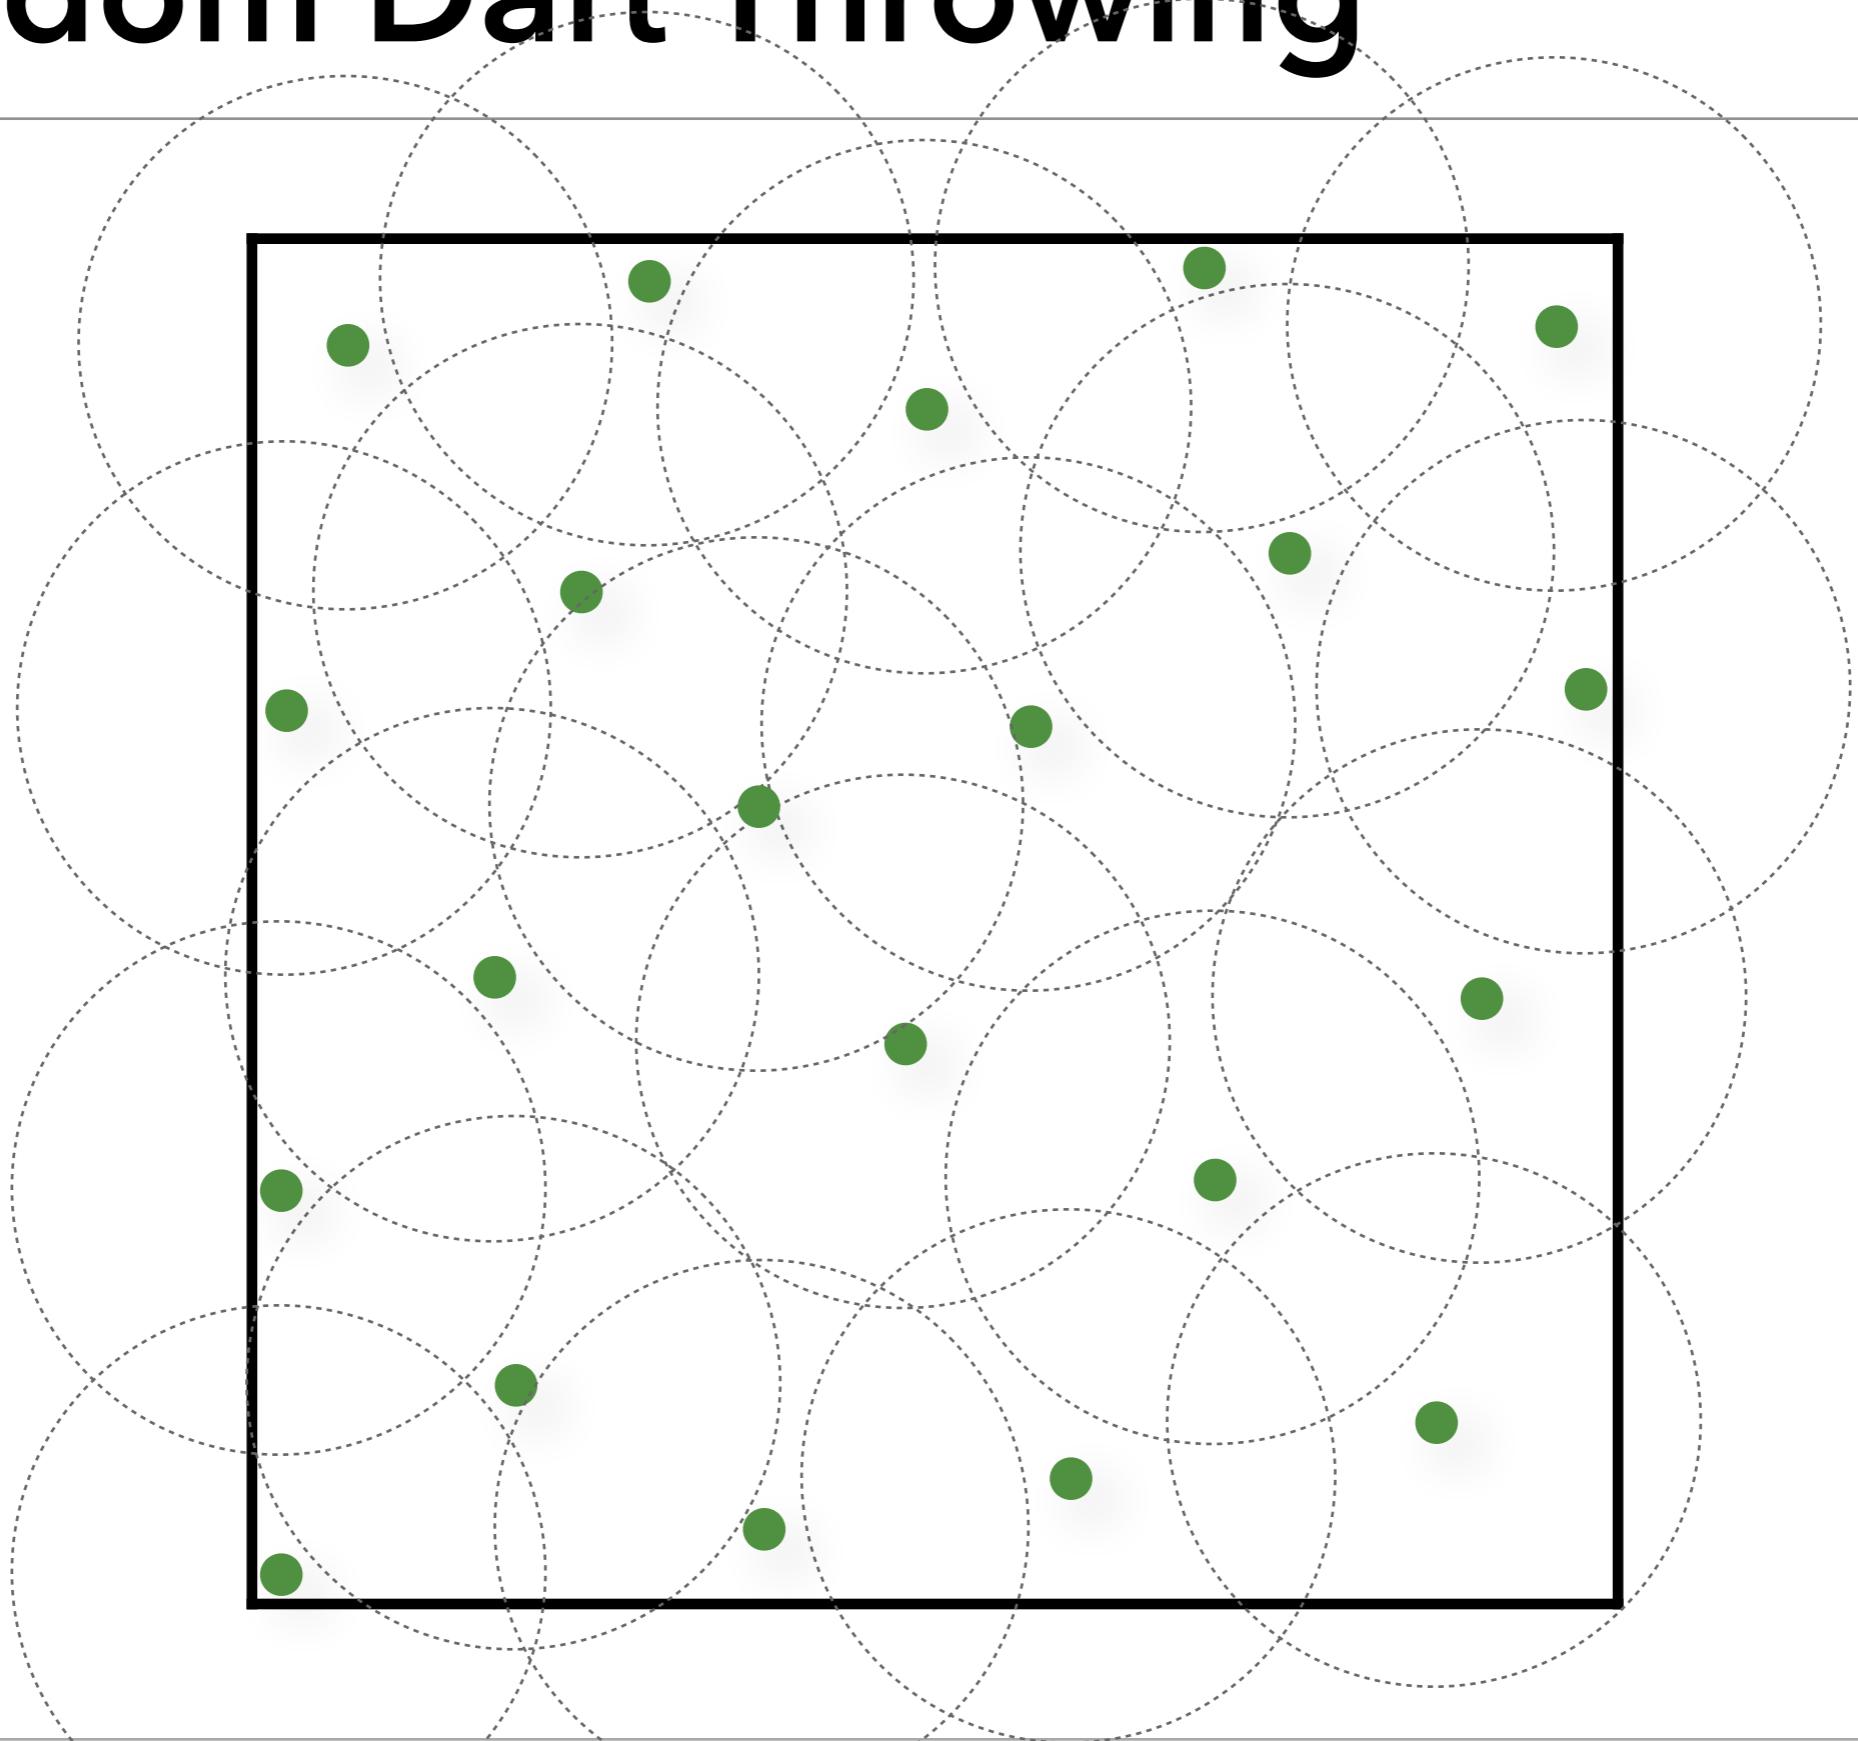
Poisson-Disk/Blue-Noise Sampling

- Enforce a minimum distance between points
- Poisson-Disk Sampling:
 - Robert L. Cook. "Stochastic sampling in computer graphics." *ACM Transactions on Graphics*. 1986.
 - Ares Lagae and Philip Dutré. "A comparison of methods for generating Poisson disk distributions." *Computer Graphics Forum*, 2008.

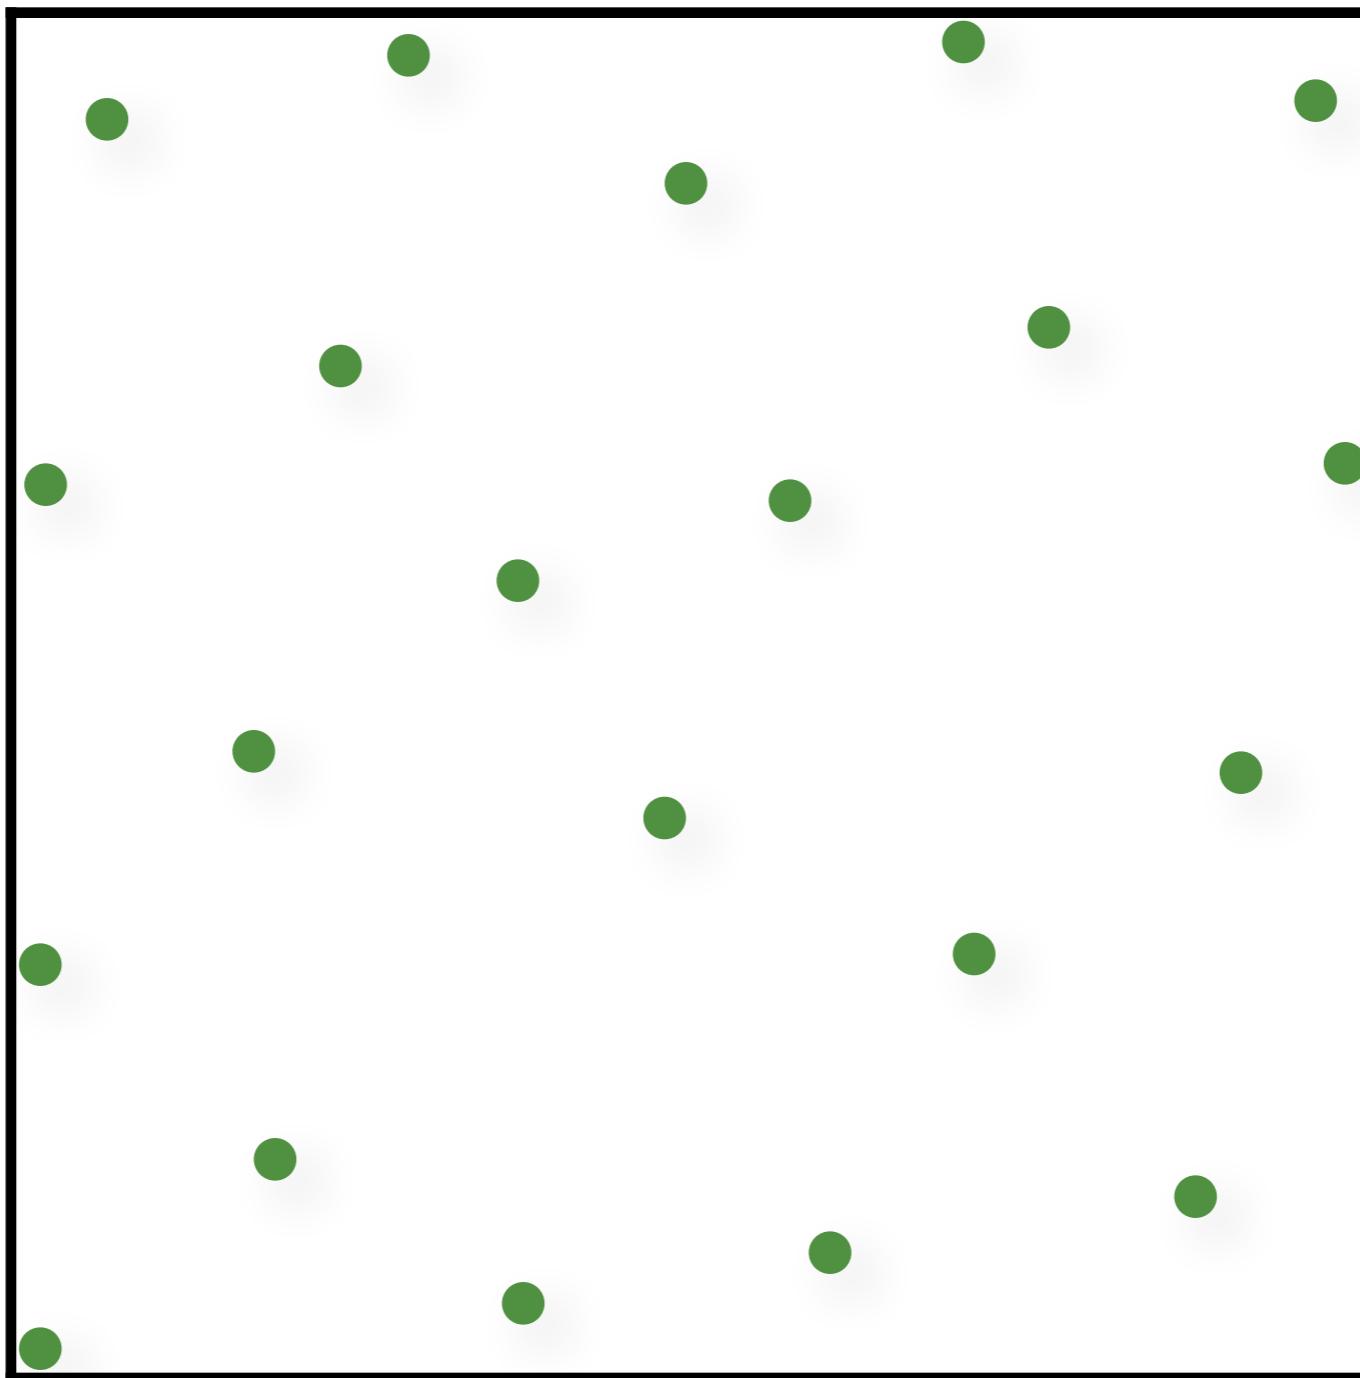
Random Dart Throwing



Random Dart Throwing



Random Dart Throwing



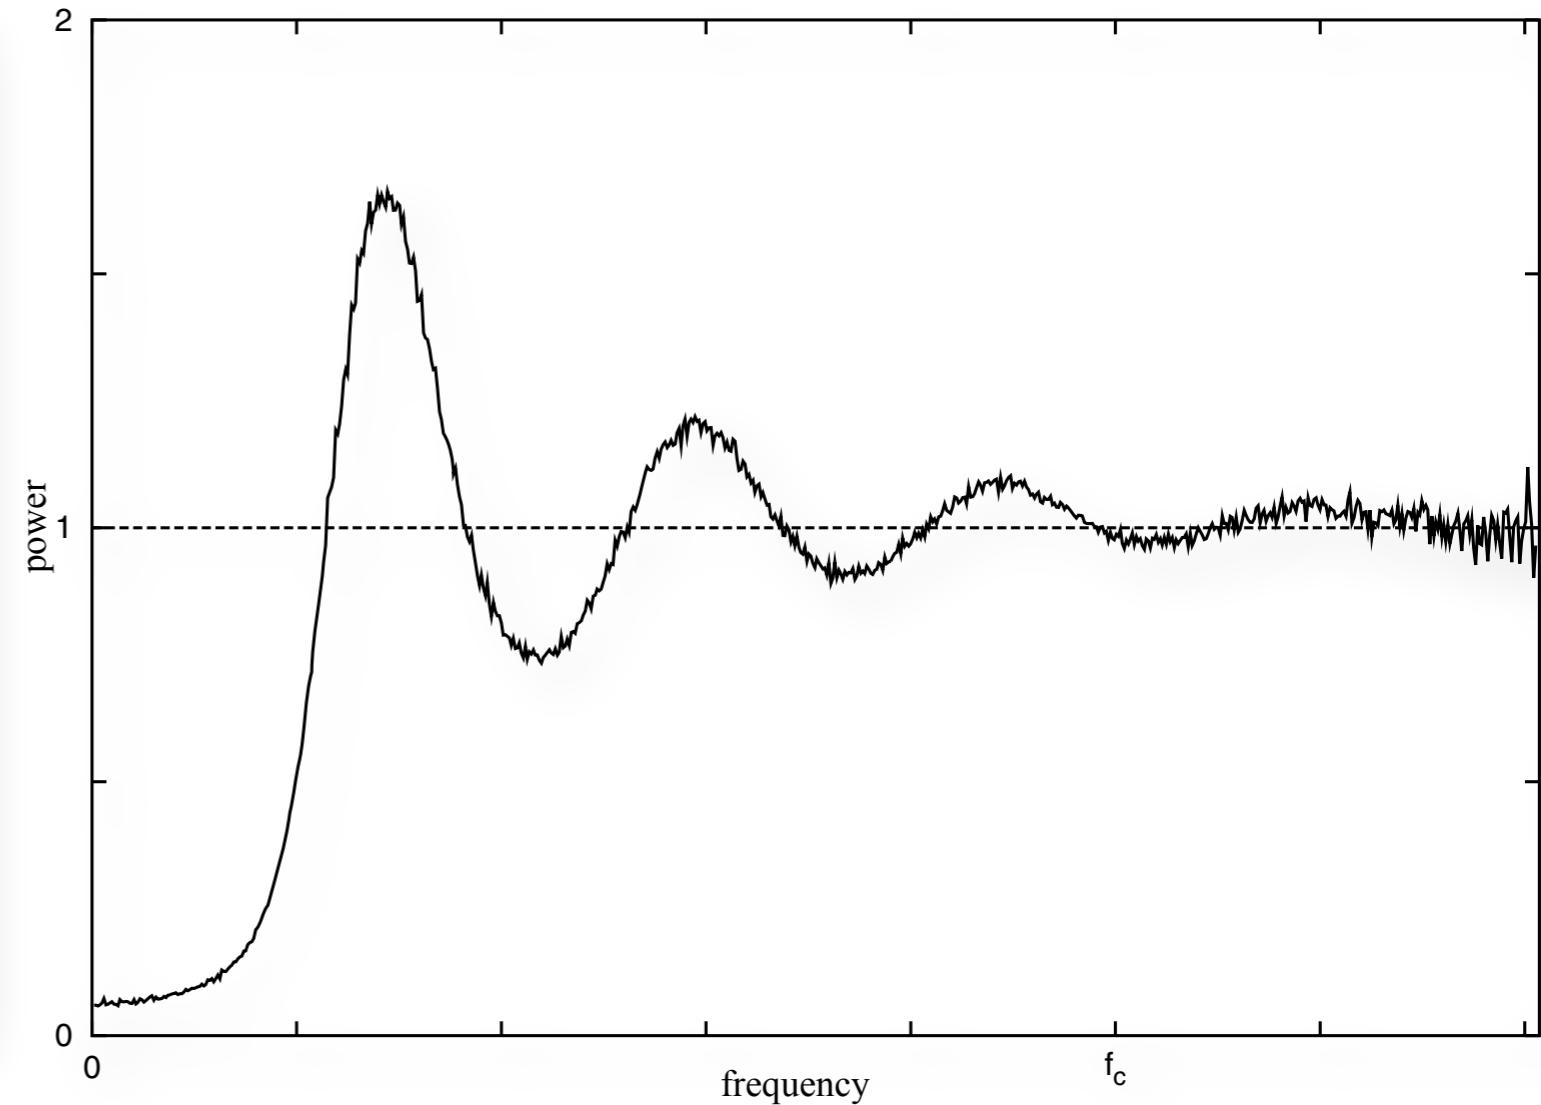
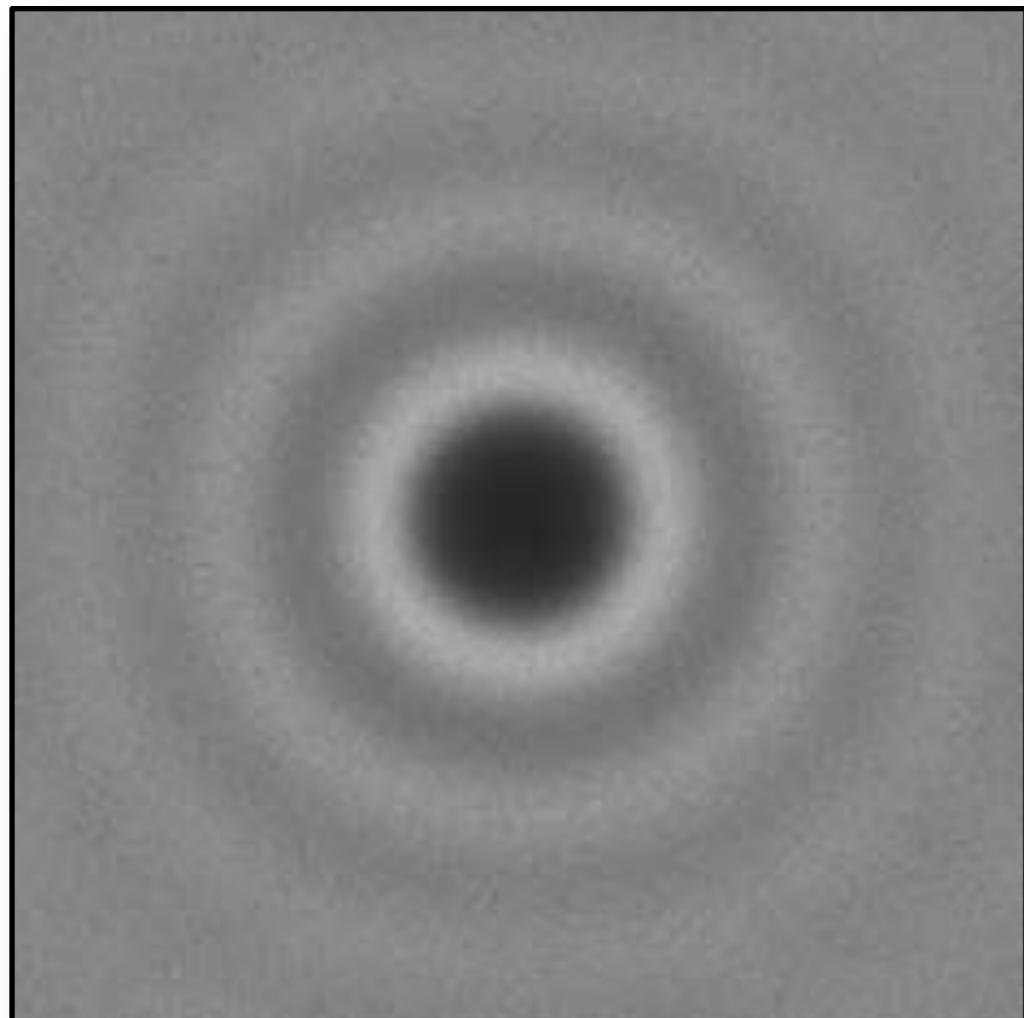
Stratified Sampling



Best Candidate Sampling



Power Spectrum



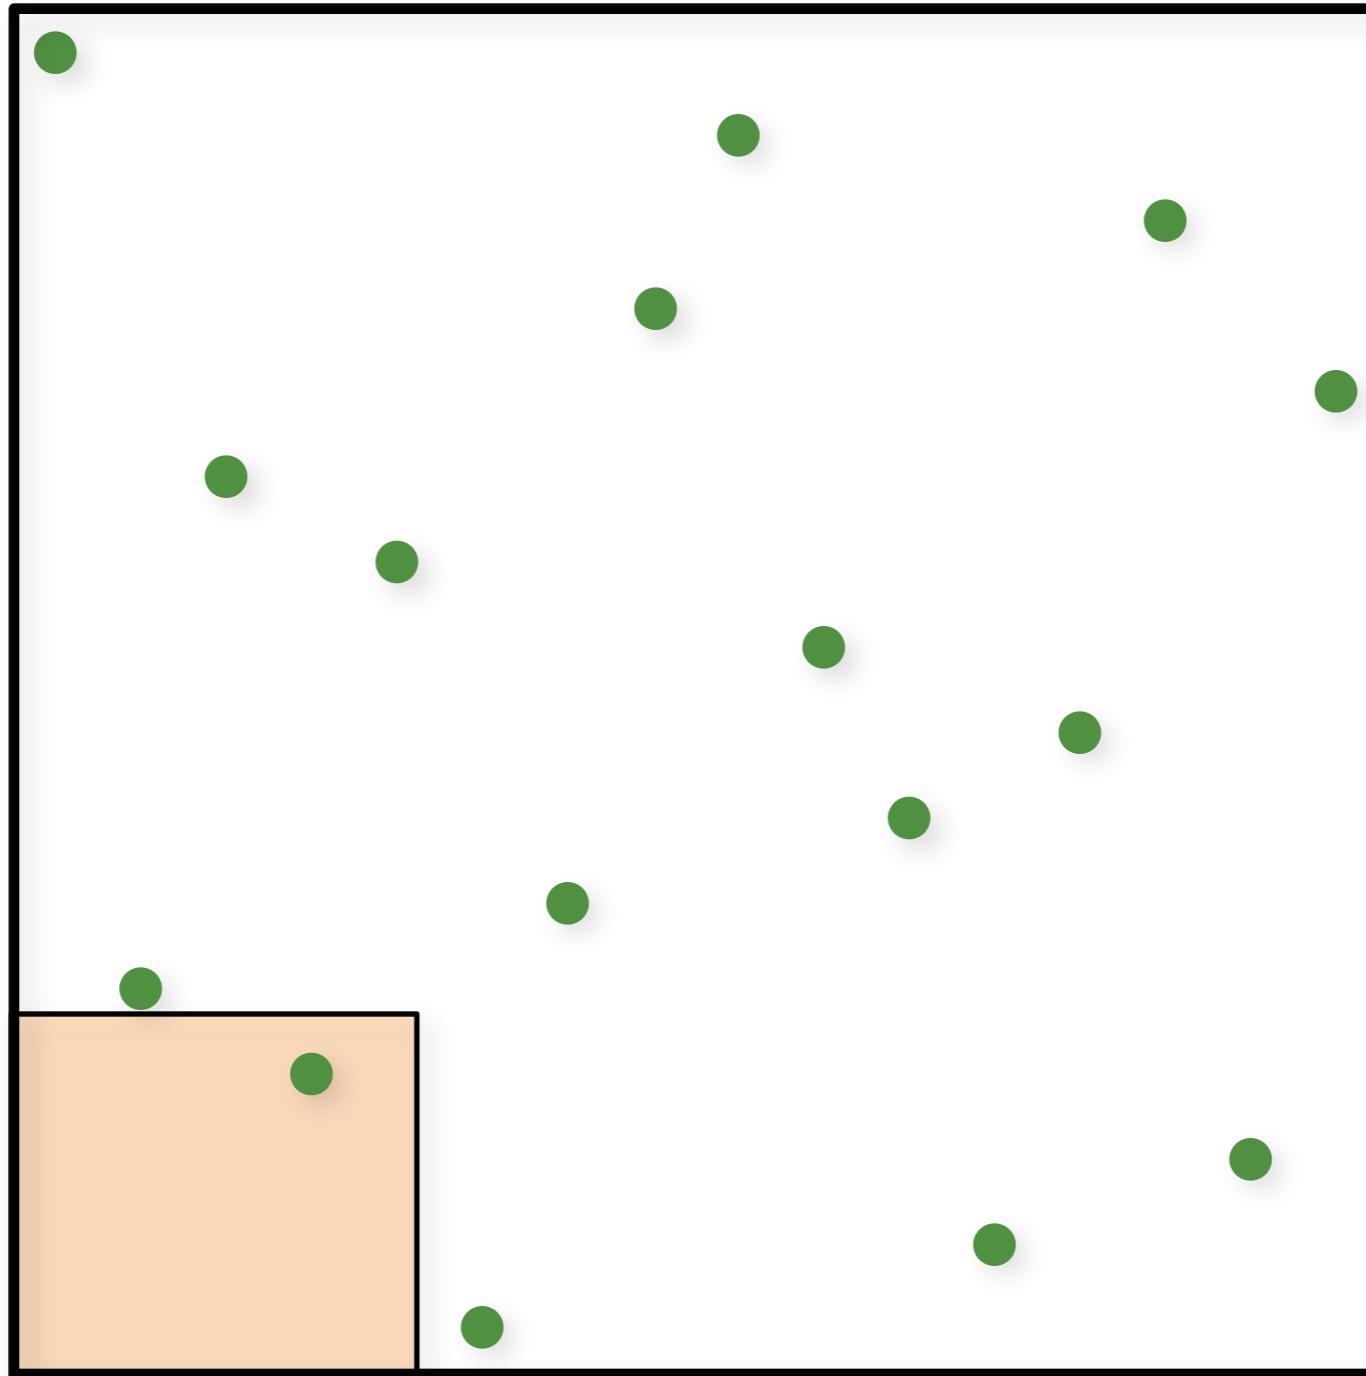
Visual Break



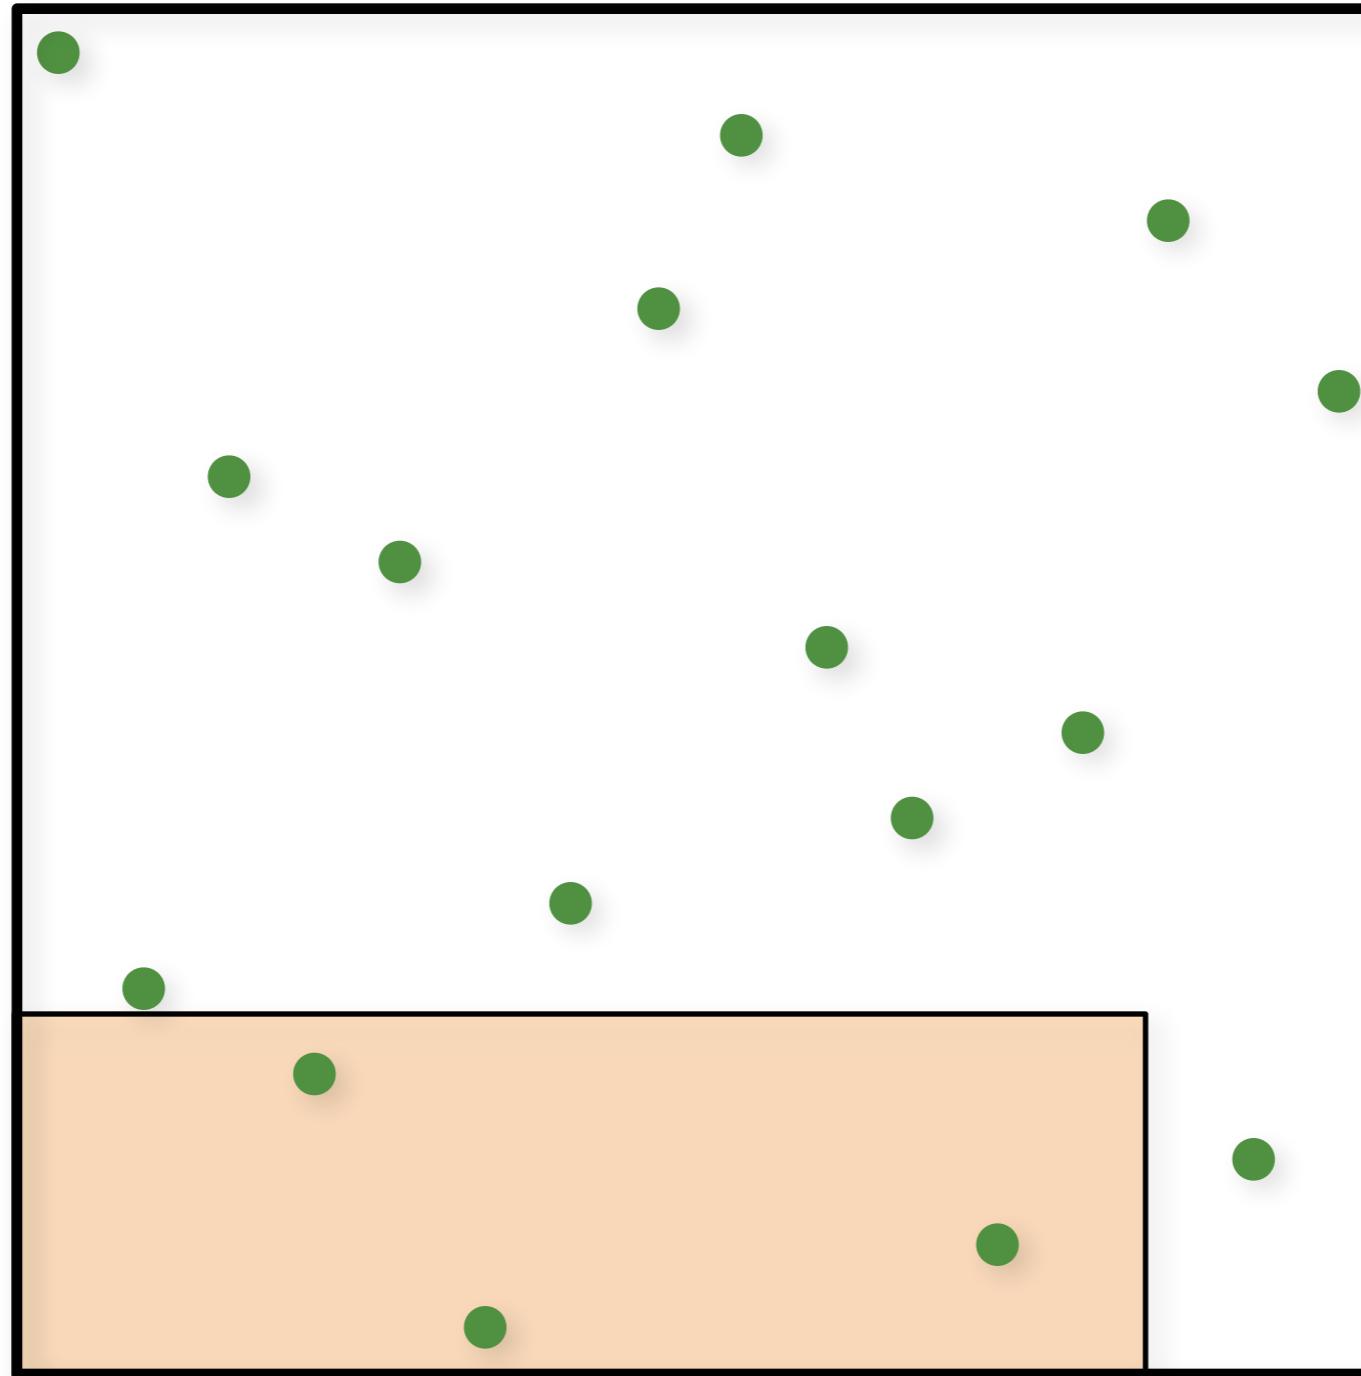
Discrepancy

- Previous stratified approaches try to minimize “clumping”
- “Discrepancy” is another possible formal definition of clumping: $D^*(x_1, \dots, x_n)$
 - for every possible subregion compute the maximum absolute difference between:
 - fraction of points in the subregion
 - volume of containing subregion

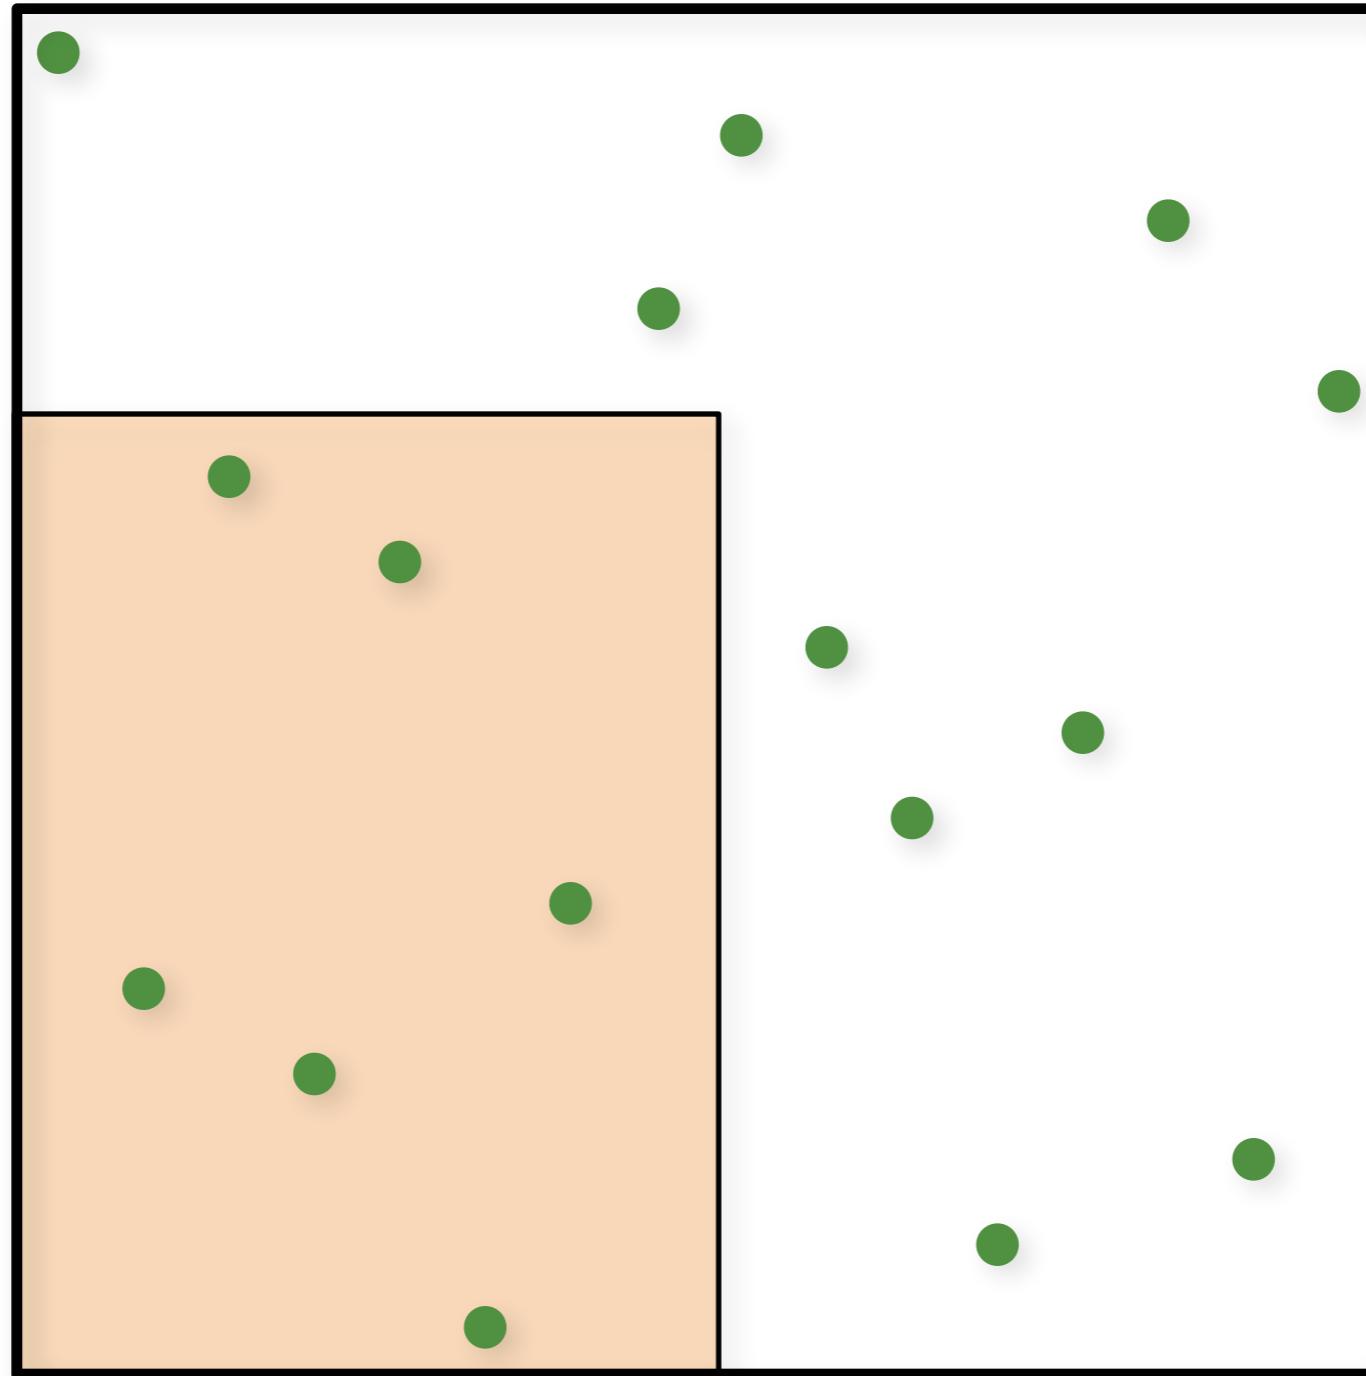
Discrepancy



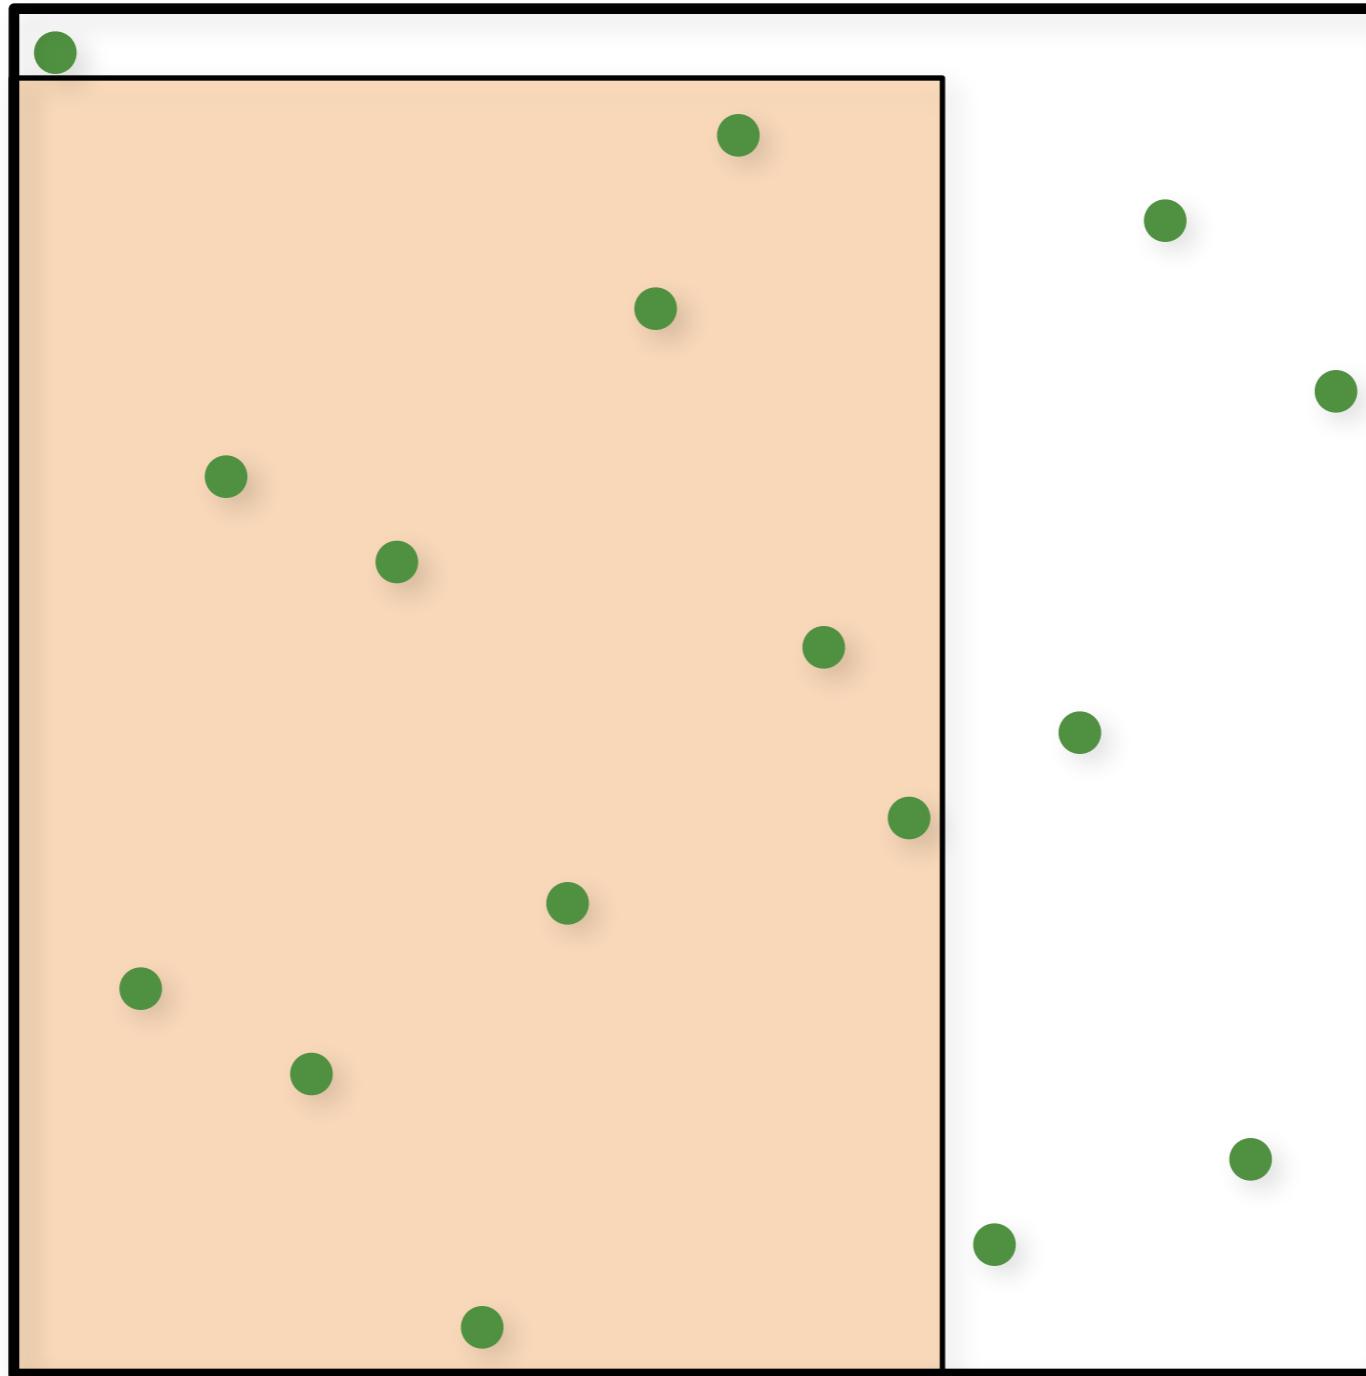
Discrepancy



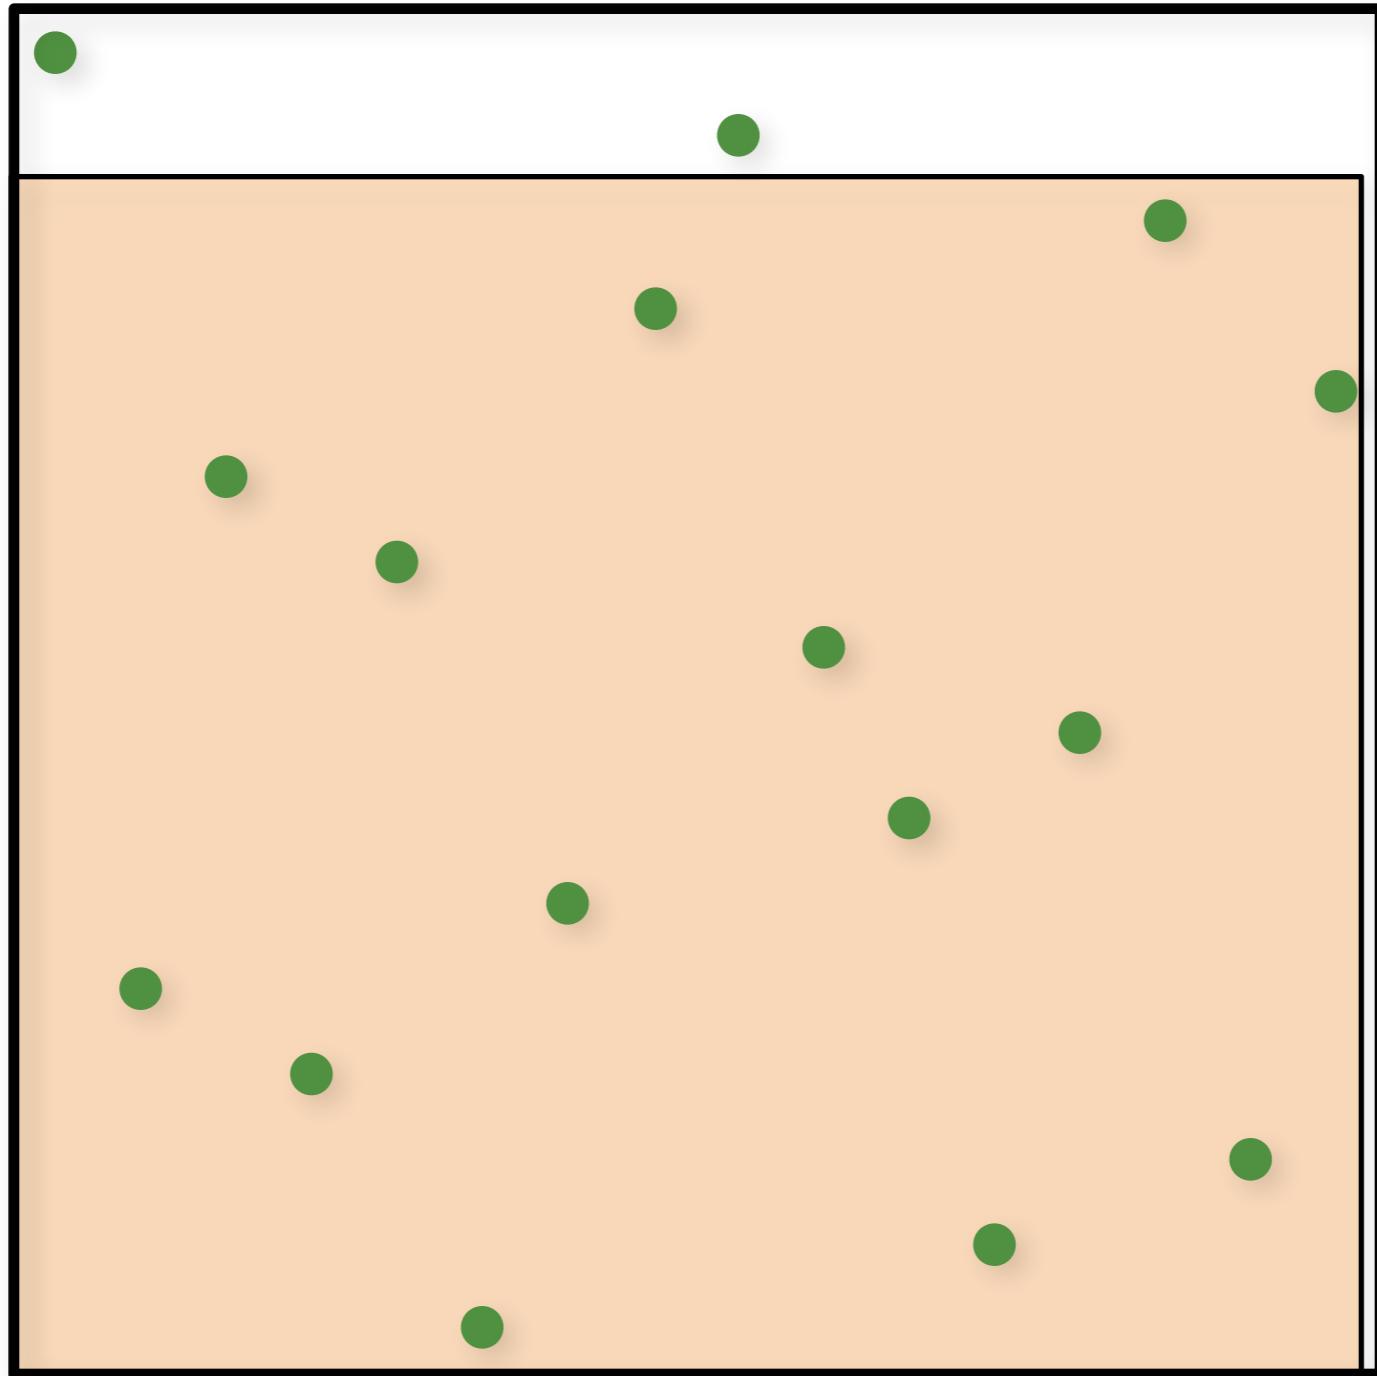
Discrepancy



Discrepancy



Discrepancy



Low-Discrepancy Sampling

- Low-Discrepancy sequences are specially crafted to have small discrepancy values.
- Entire field of study called Quasi-Monte Carlo (QMC)

Koksma-Hlawka inequality

$$\left| \frac{1}{n} \sum_{i=1}^n f(x_i) - \int f(u) \, du \right| \leq V(f) D^*(x_1, \dots, x_n)$$

The Radical Inverse

- A positive integer value n can be expressed in a base b with a sequence of digits $d_m \dots d_2 d_1$ uniquely determined by:

$$n = \sum_{i=1}^{\infty} d_i b^{i-1}$$

- The radical inverse function Φ_b in base b converts a nonnegative integer n to a floating-point value in $[0, 1)$ by reflecting these digits about the decimal point:

$$\Phi_b(n) = 0.d_1 d_2 \dots d_m$$

- Subsequent points “fall into biggest holes”

The Radical Inverse

```
float radicalInverse(int n, int base, float inv)
{
    float v = 0.0f;
    for (float p = inv; n != 0; p *= inv, n /= base)
        v += (n % base) * p;
    return v;
}

float radicalInverse(int n, int base)
{
    return radicalInverse(n, base, 1.0f / base);
}
```

More efficient version available for base 2

The Van der Corput Sequence

- Radical Inverse in base two

n	Base 2	Φ_b
1	1	.1 = 1/2
2	10	.01 = 1/4
3	11 100	.11 = 3/4 .001 = 1/8
4	101	.101 = 5/8
5	110	.011 = 3/8
6	111	.111 = 7/8
...		



The Radical Inverse (Base 2)

```
float vanDerCorputRIU(uint n)
{
    n = (n << 16) | (n >> 16);
    n = ((n & 0x00ff00ff) << 8) | ((n & 0xff00ff00) >> 8);
    n = ((n & 0x0f0f0f0f) << 4) | ((n & 0xf0f0f0f0) >> 4);
    n = ((n & 0x33333333) << 2) | ((n & 0xcccccccc) >> 2);
    n = ((n & 0x55555555) << 1) | ((n & 0xaaaaaaaa) >> 1);
    return n / float (0x100000000LL);
}
```

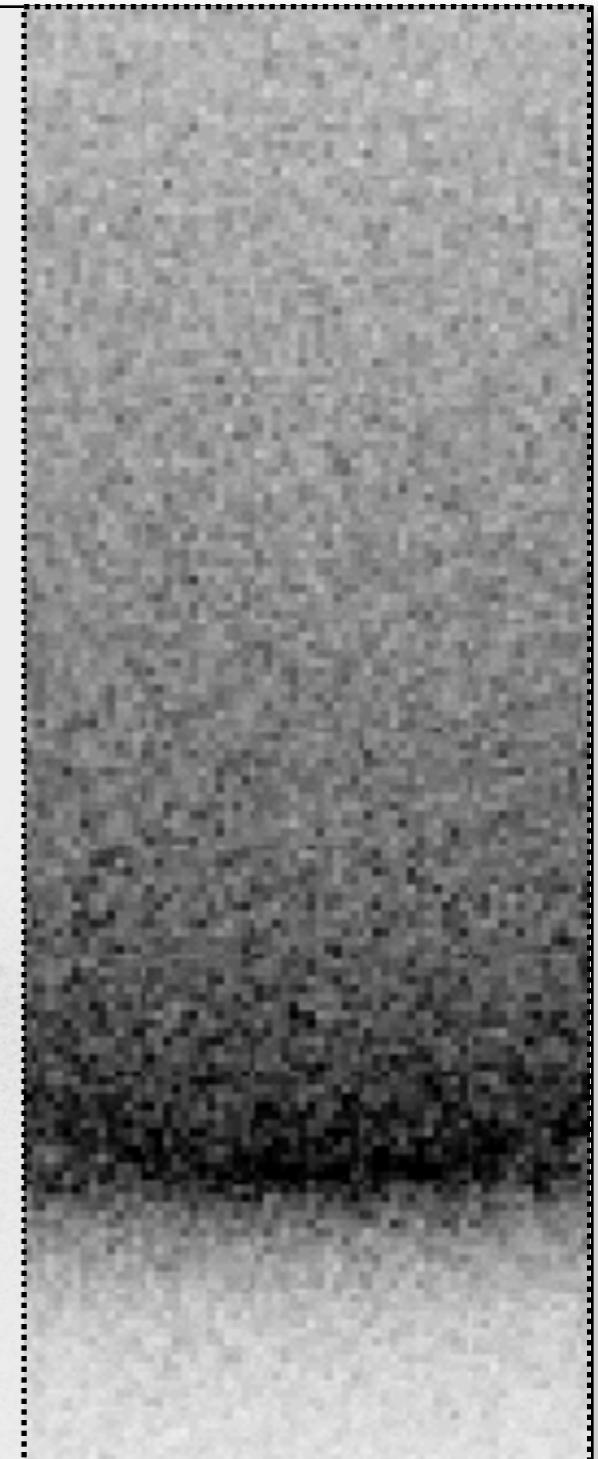
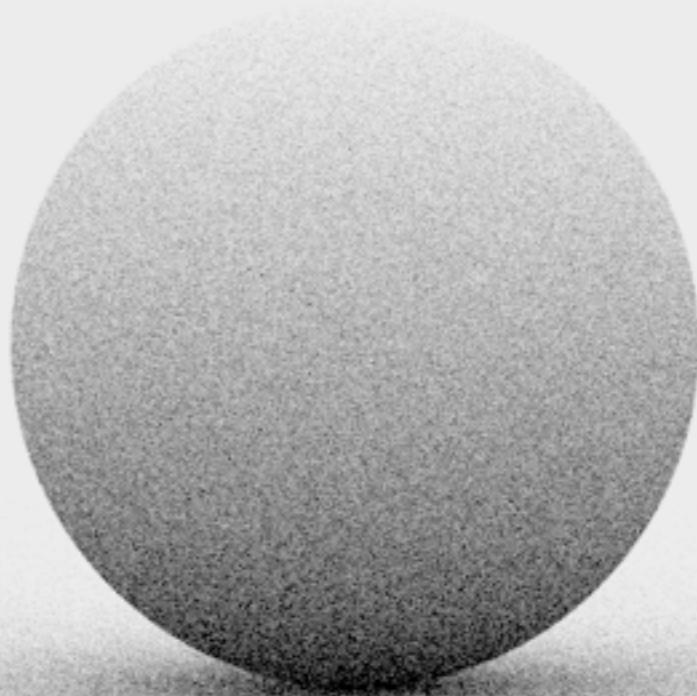
The Halton Sequence

- An n-dimensional Halton sequence uses the radical inverse with a different base b for each dimension of the pattern.
- The bases must all be relatively prime.

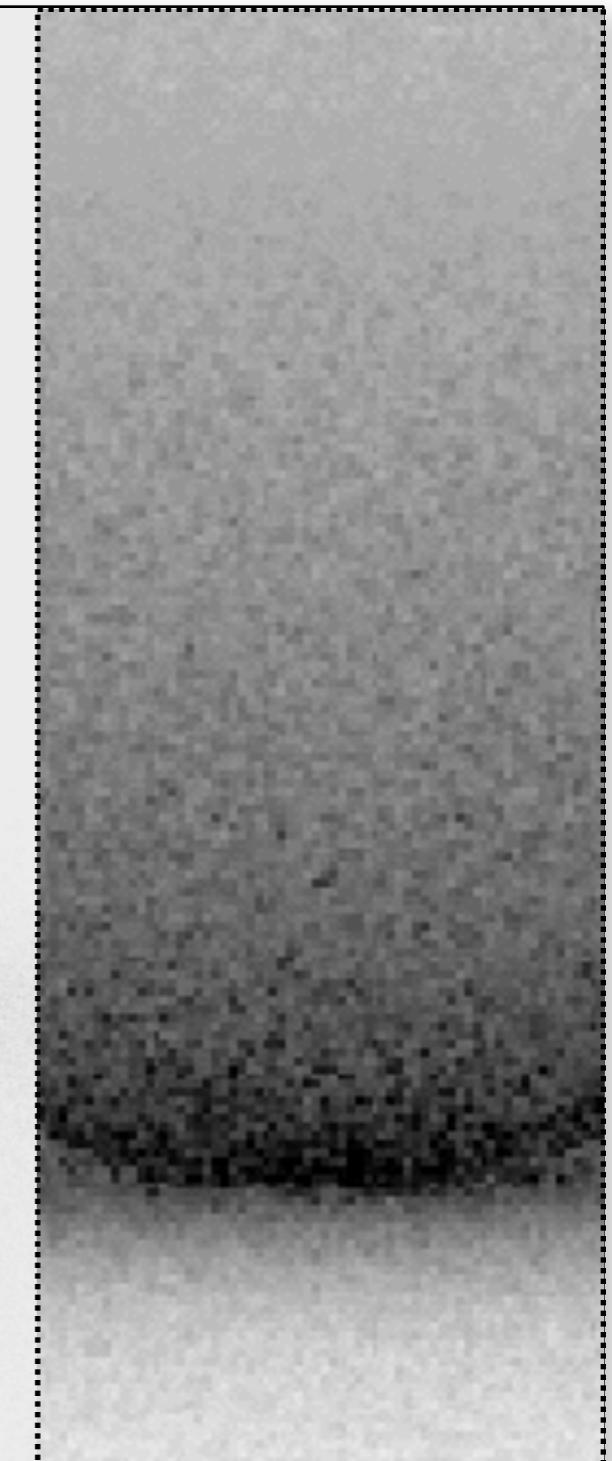
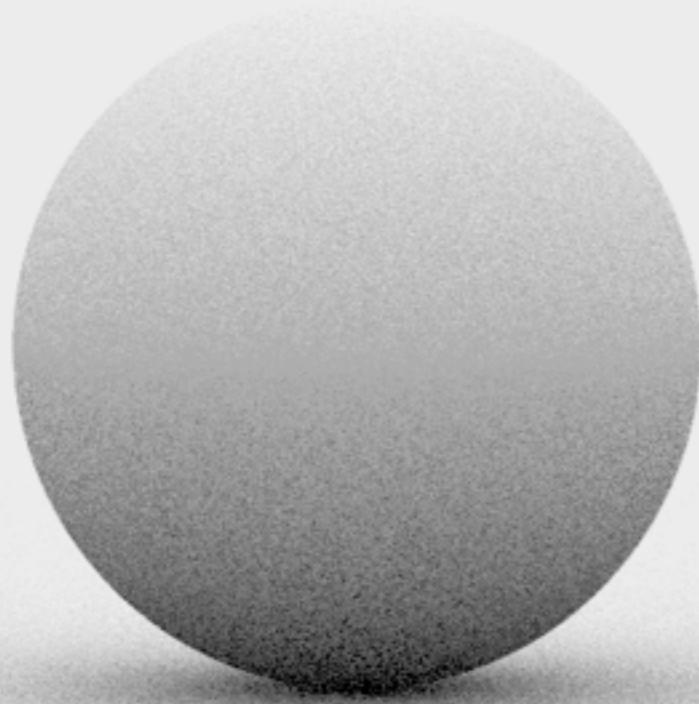
$$x_i = (\Phi_2(i), \Phi_3(i), \Phi_5(i), \dots, \Phi_{p_n}(i))$$

- Incremental/progressive generation of samples

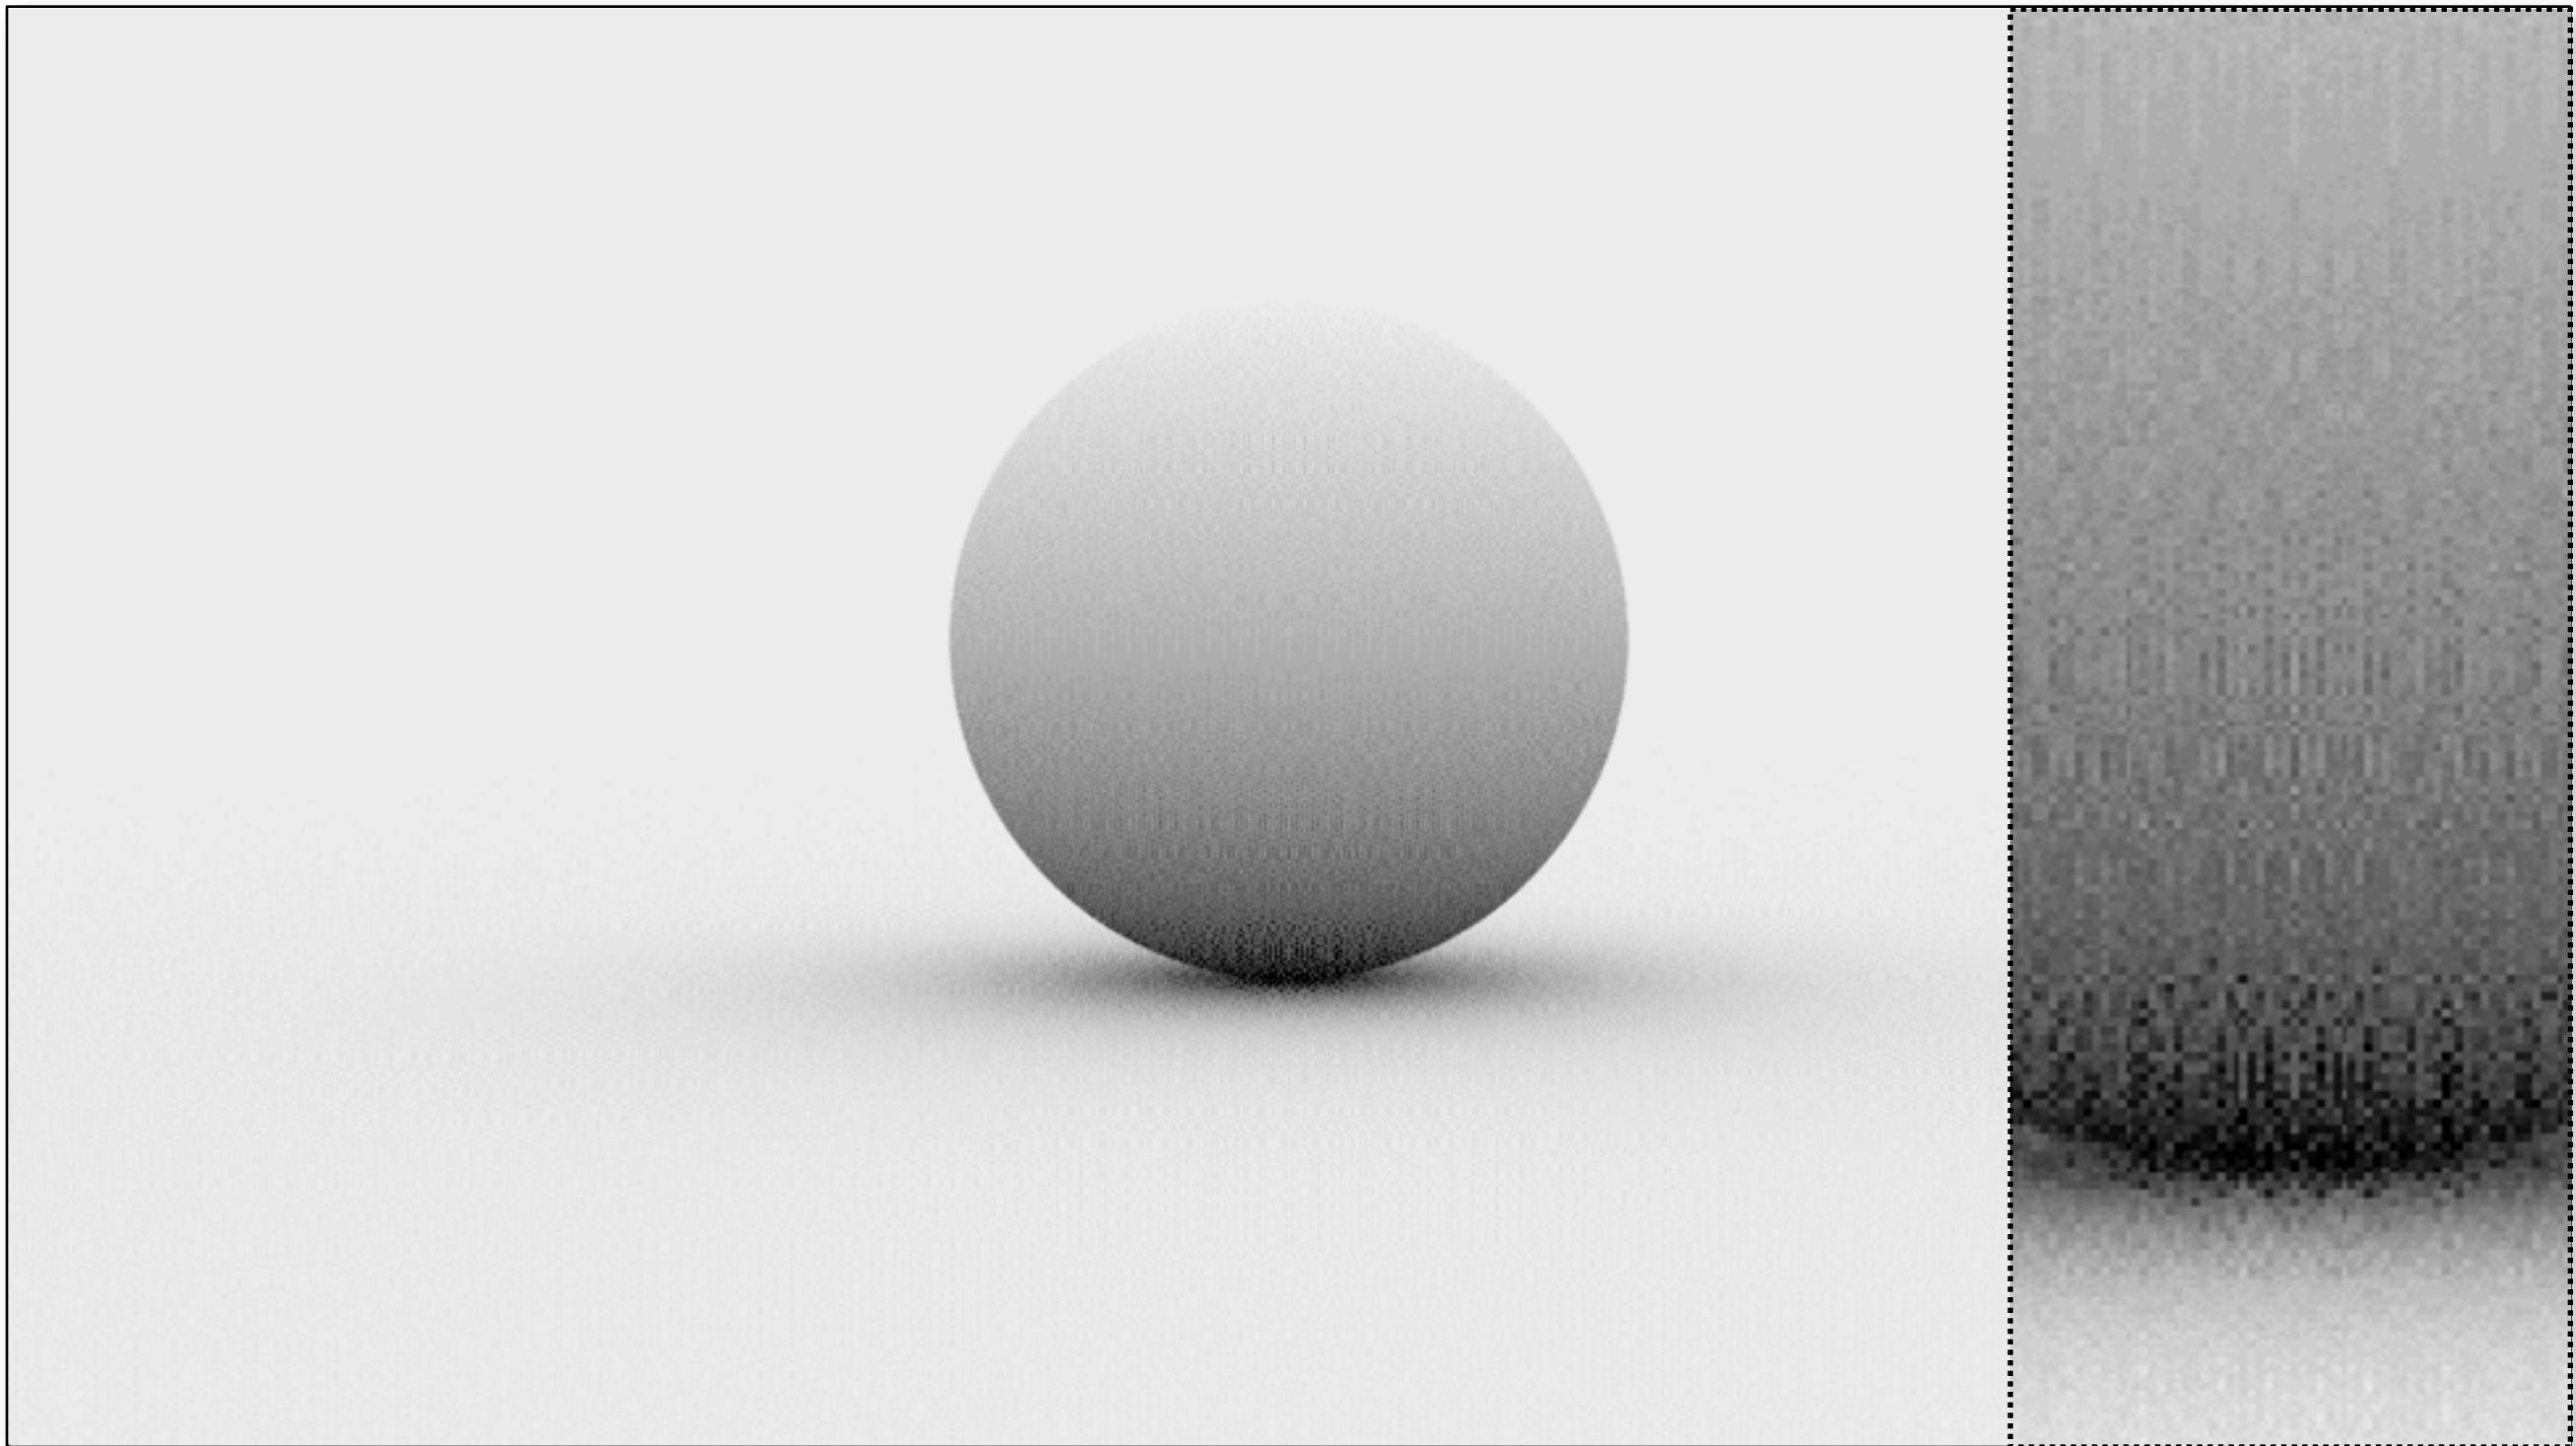
Monte Carlo (16 random samples)



Monte Carlo (16 stratified samples)



Quasi-Monte Carlo (16 Halton samples)

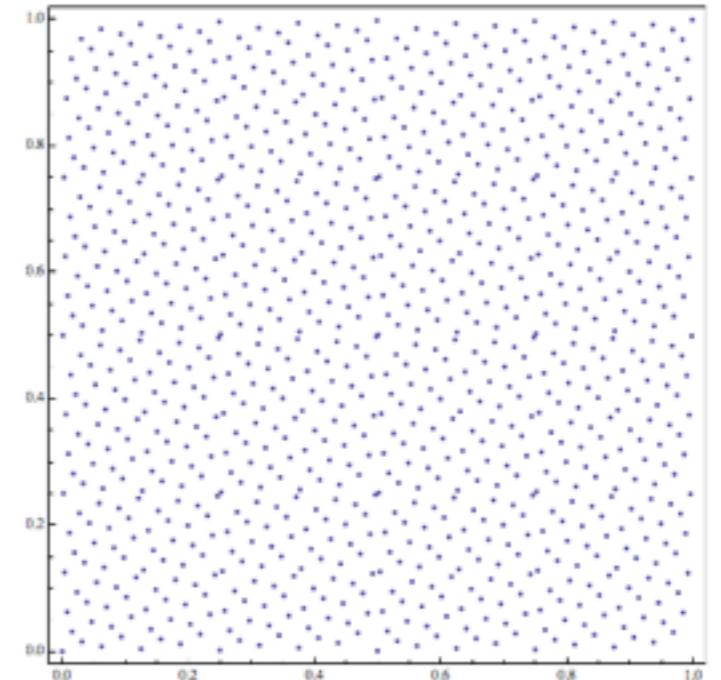


The Hammersley Sequence

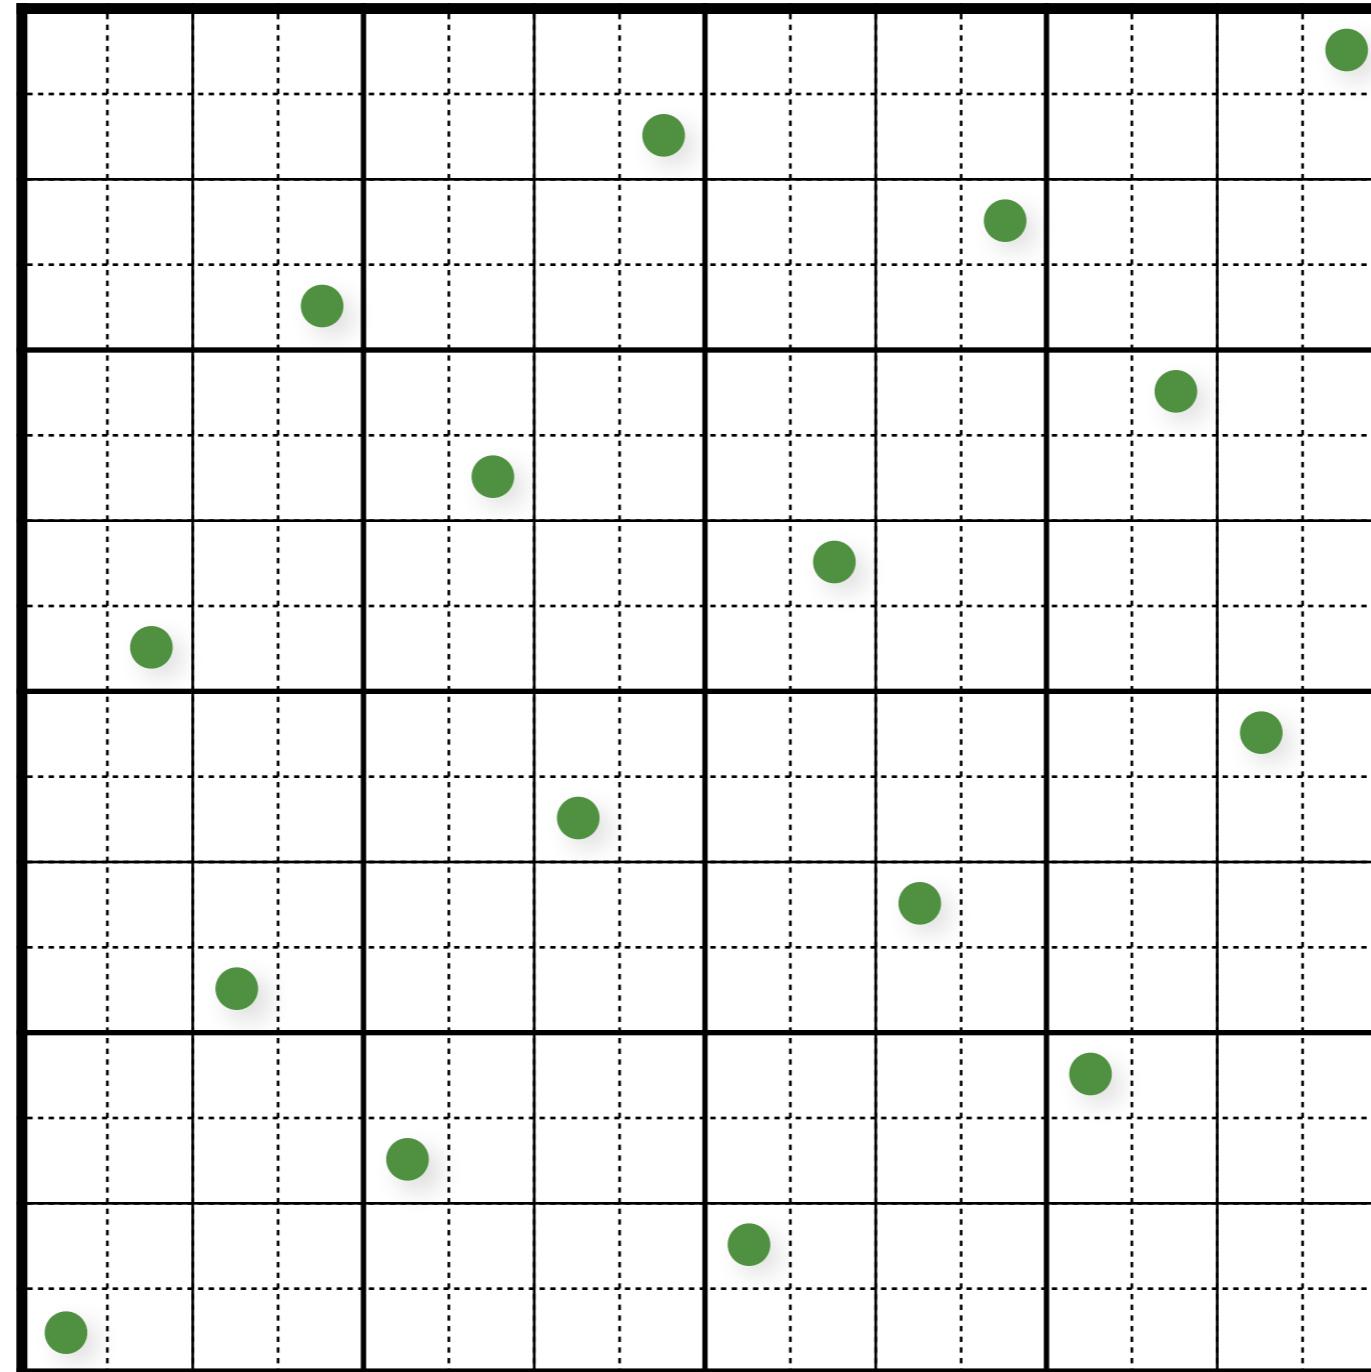
- Same as Halton, but uses i/N for first dimension:

$$x_i = \left(\frac{i}{N}, \Phi_2(i), \Phi_3(i), \dots, \Phi_{p_n}(i) \right)$$

- Provides slightly lower discrepancy
- Not incremental, need to know total number of samples, N , in advance

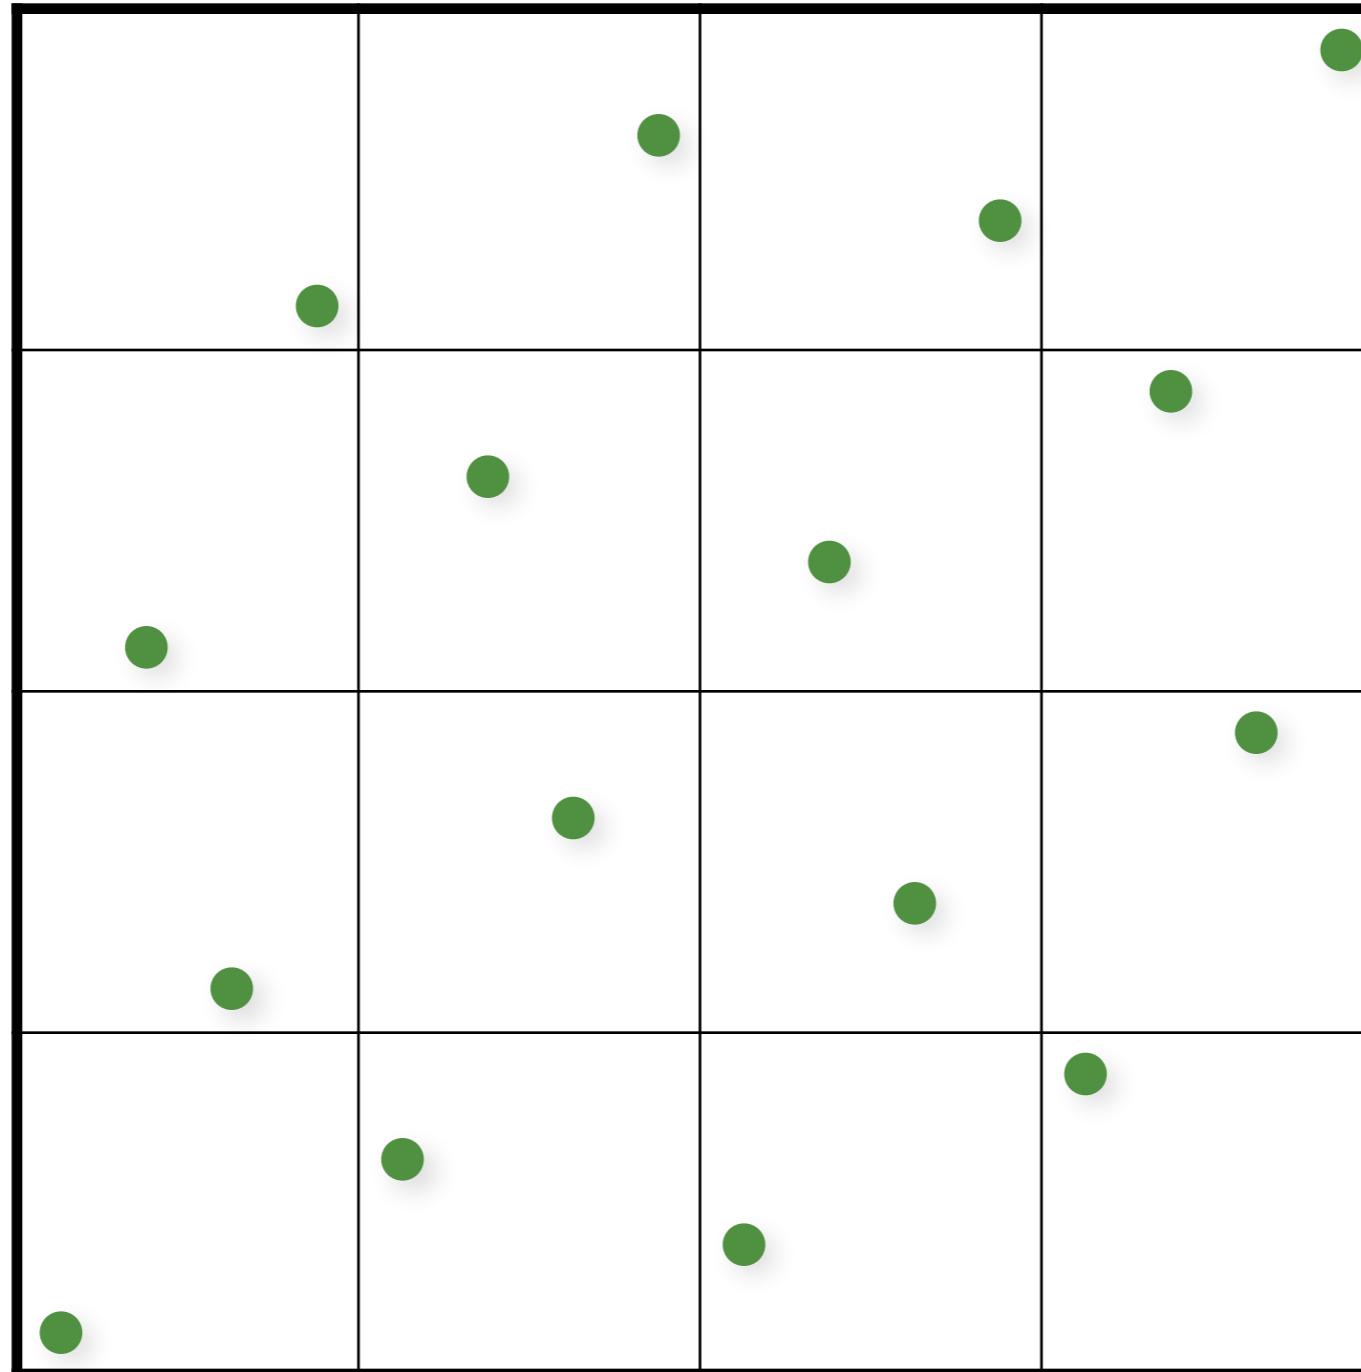


The Hammersley Sequence



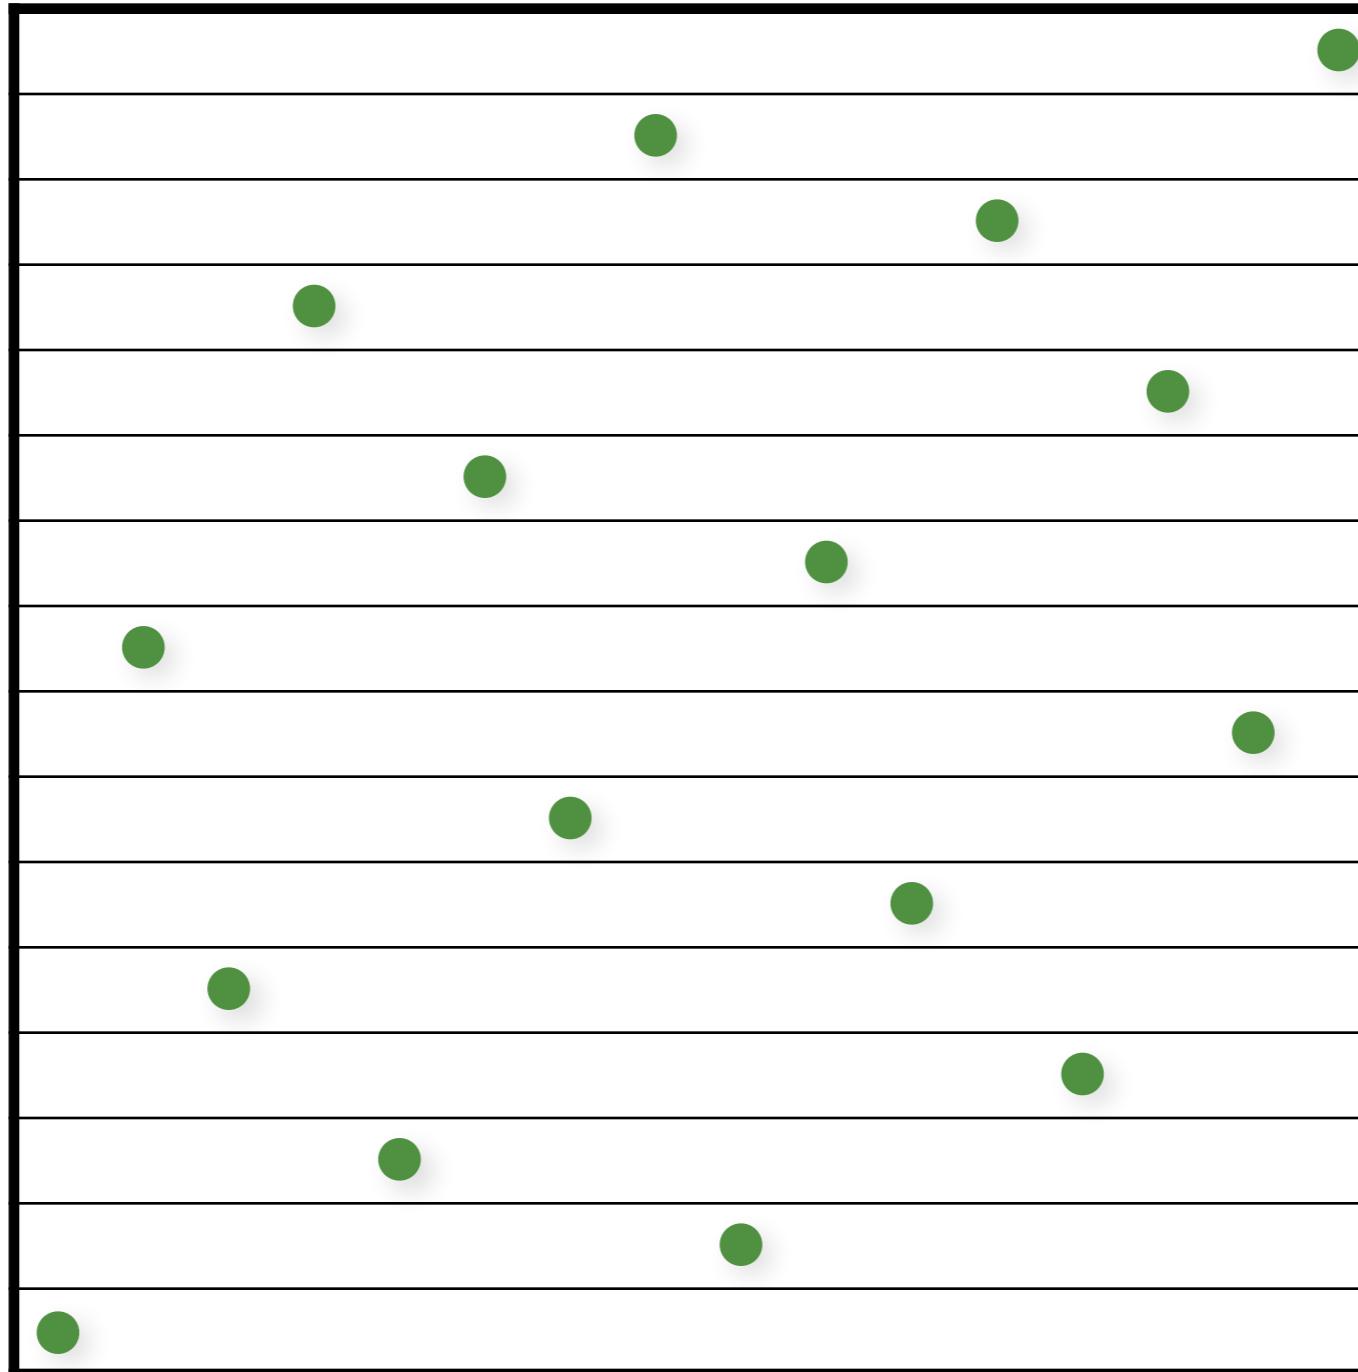
1 sample in each “elementary interval”

The Hammersley Sequence



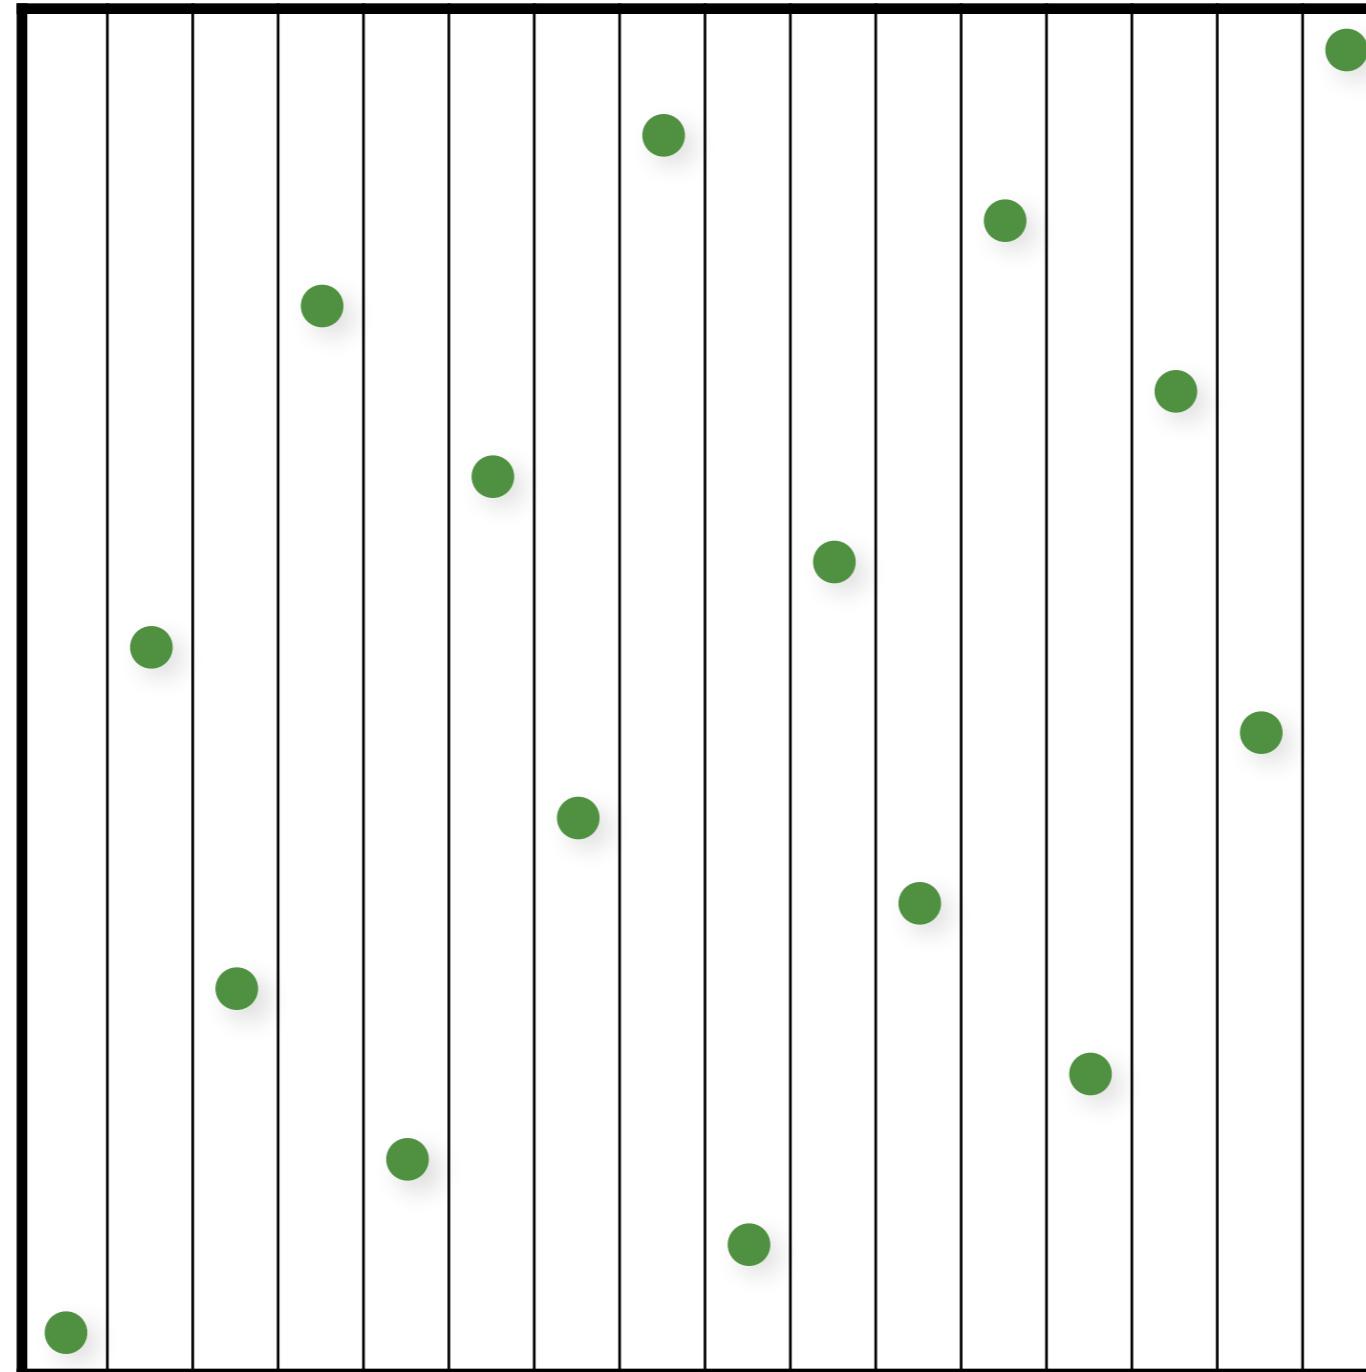
1 sample in each “elementary interval”

The Hammersley Sequence



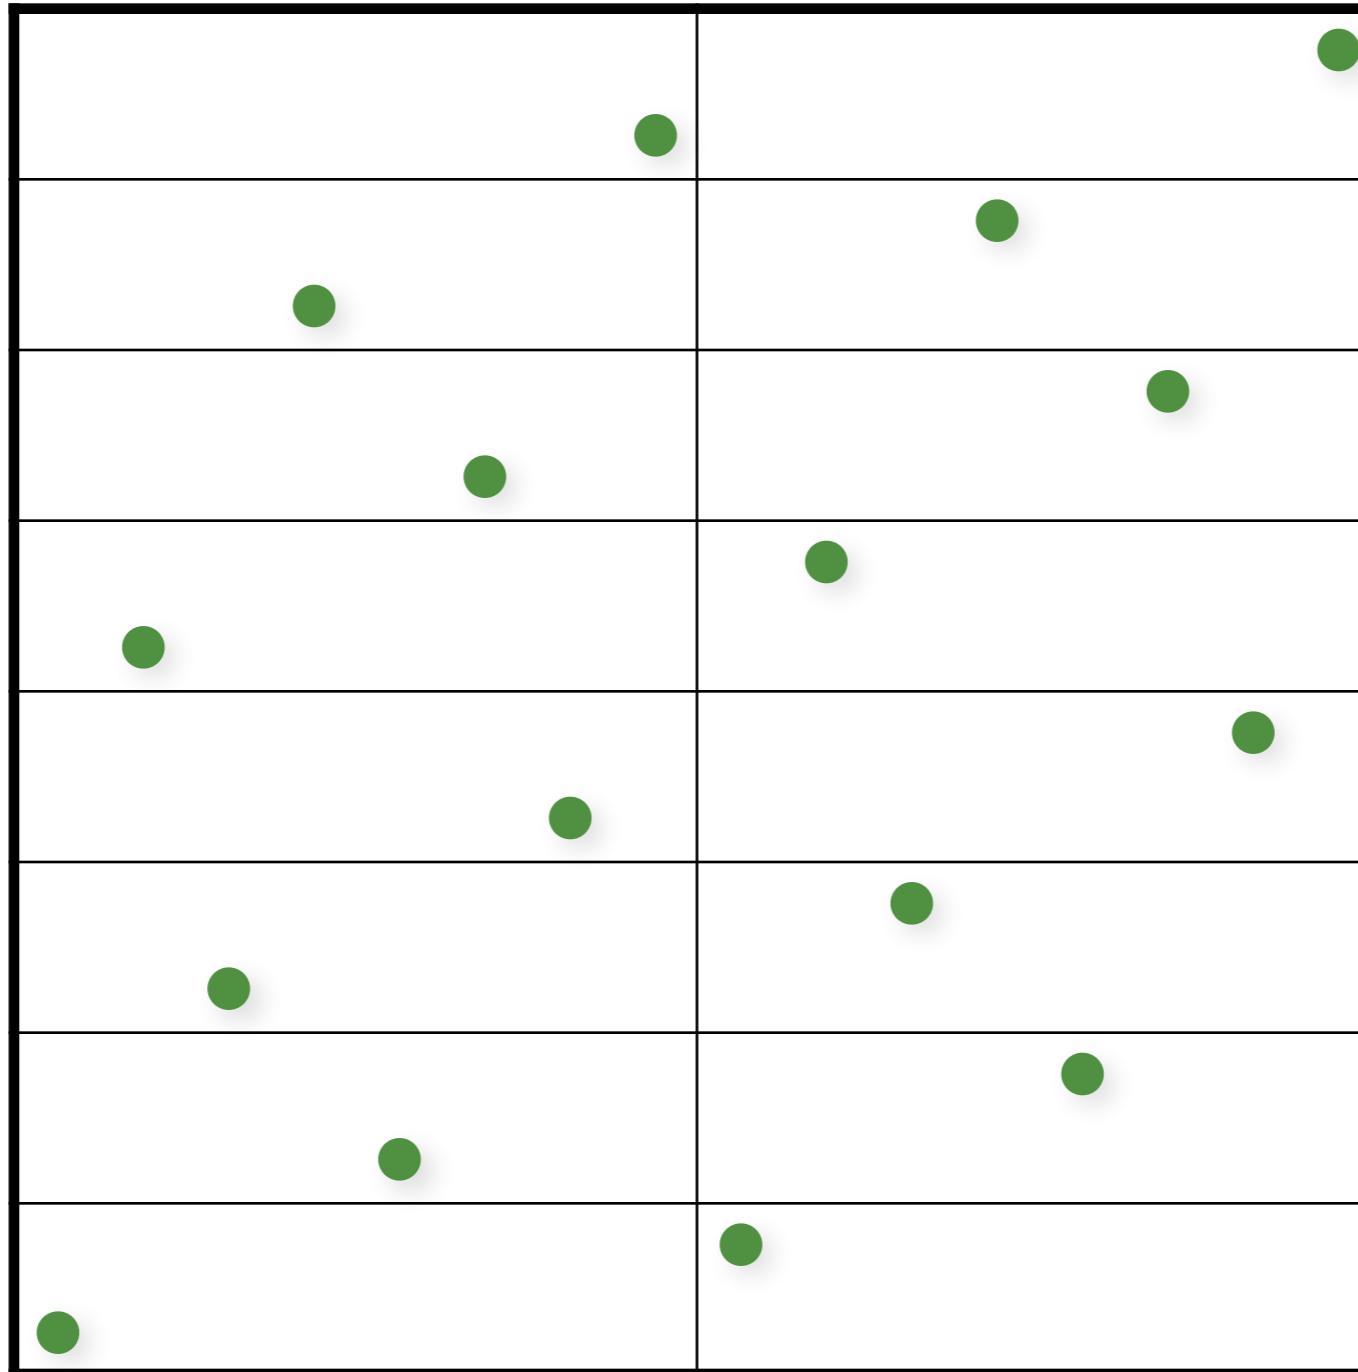
1 sample in each “elementary interval”

The Hammersley Sequence



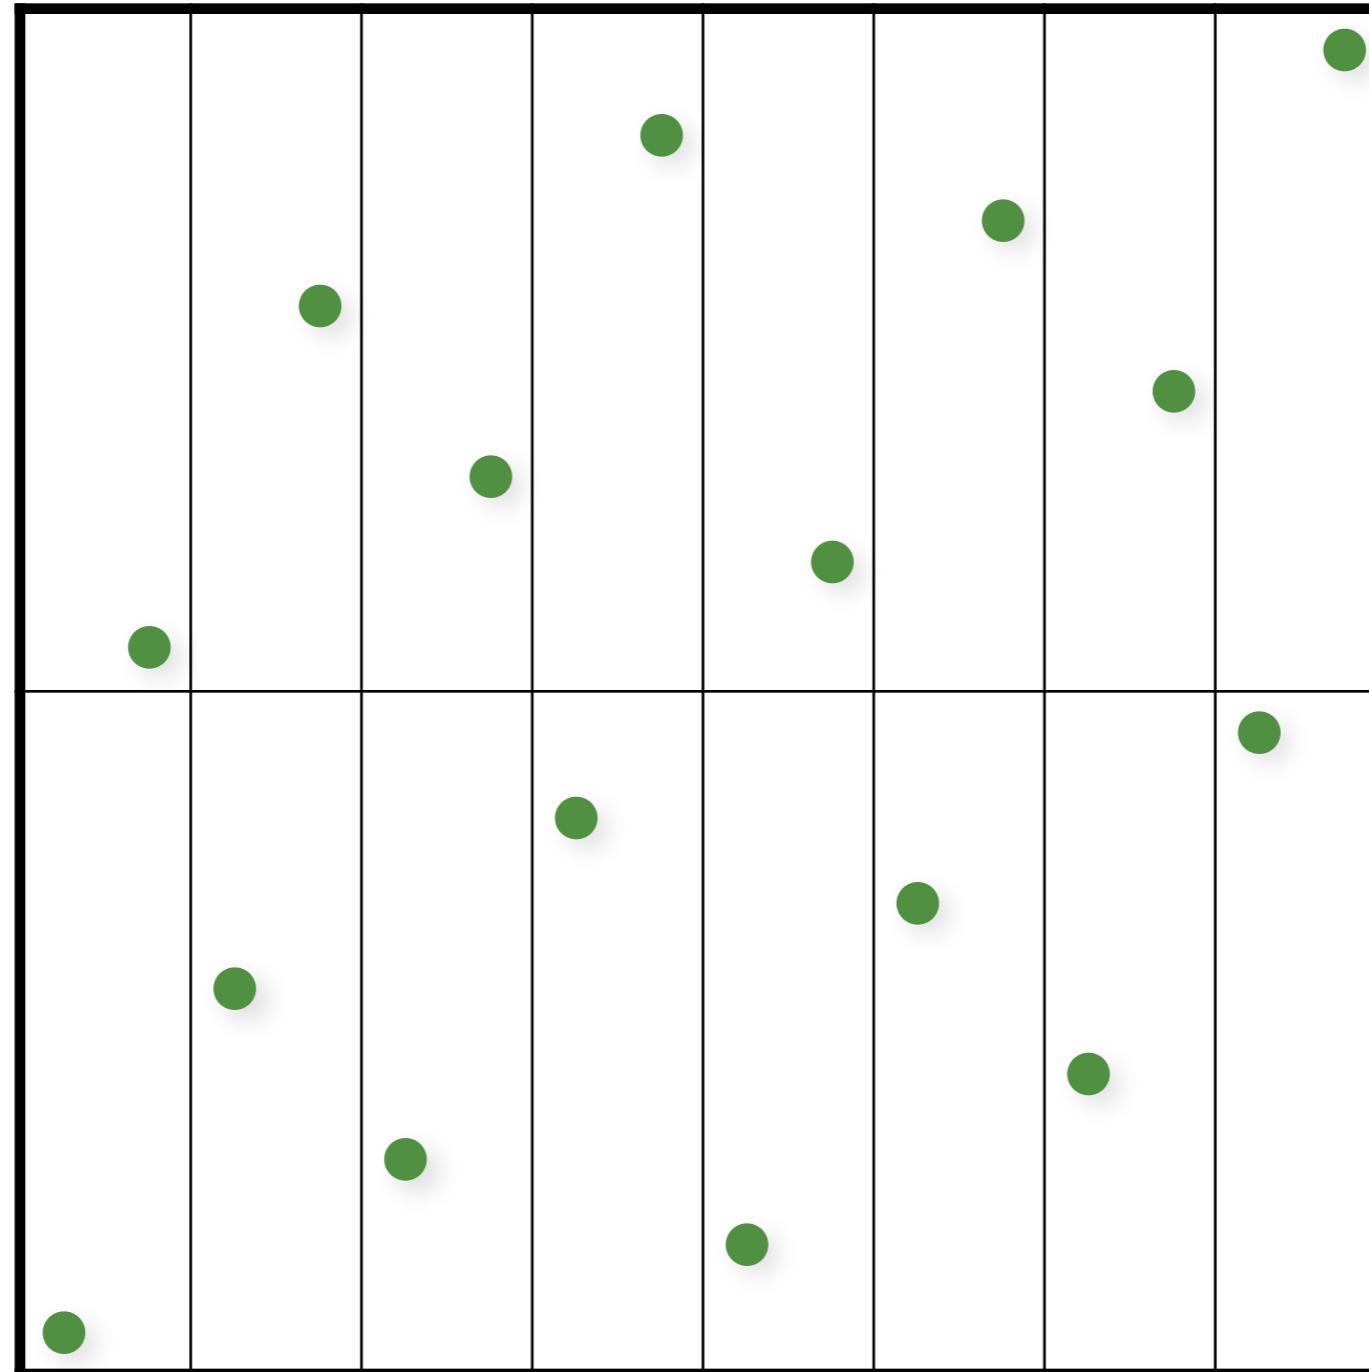
1 sample in each “elementary interval”

The Hammersley Sequence



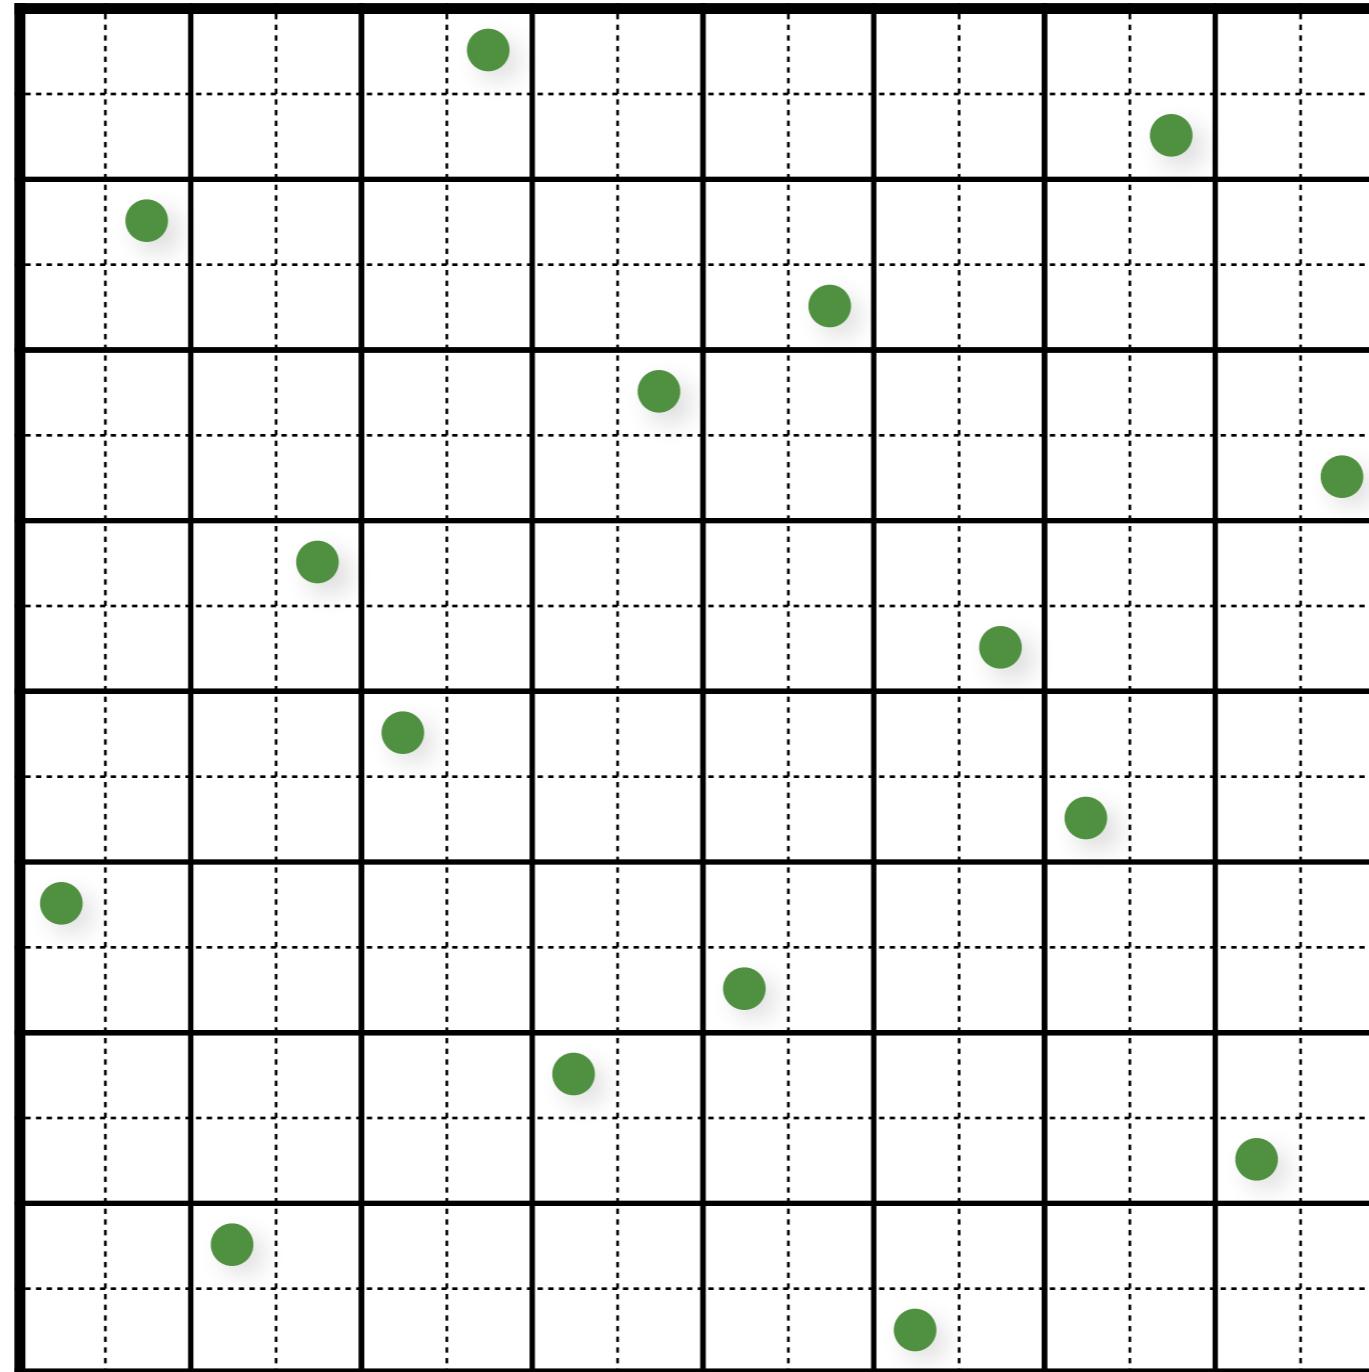
1 sample in each “elementary interval”

The Hammersley Sequence



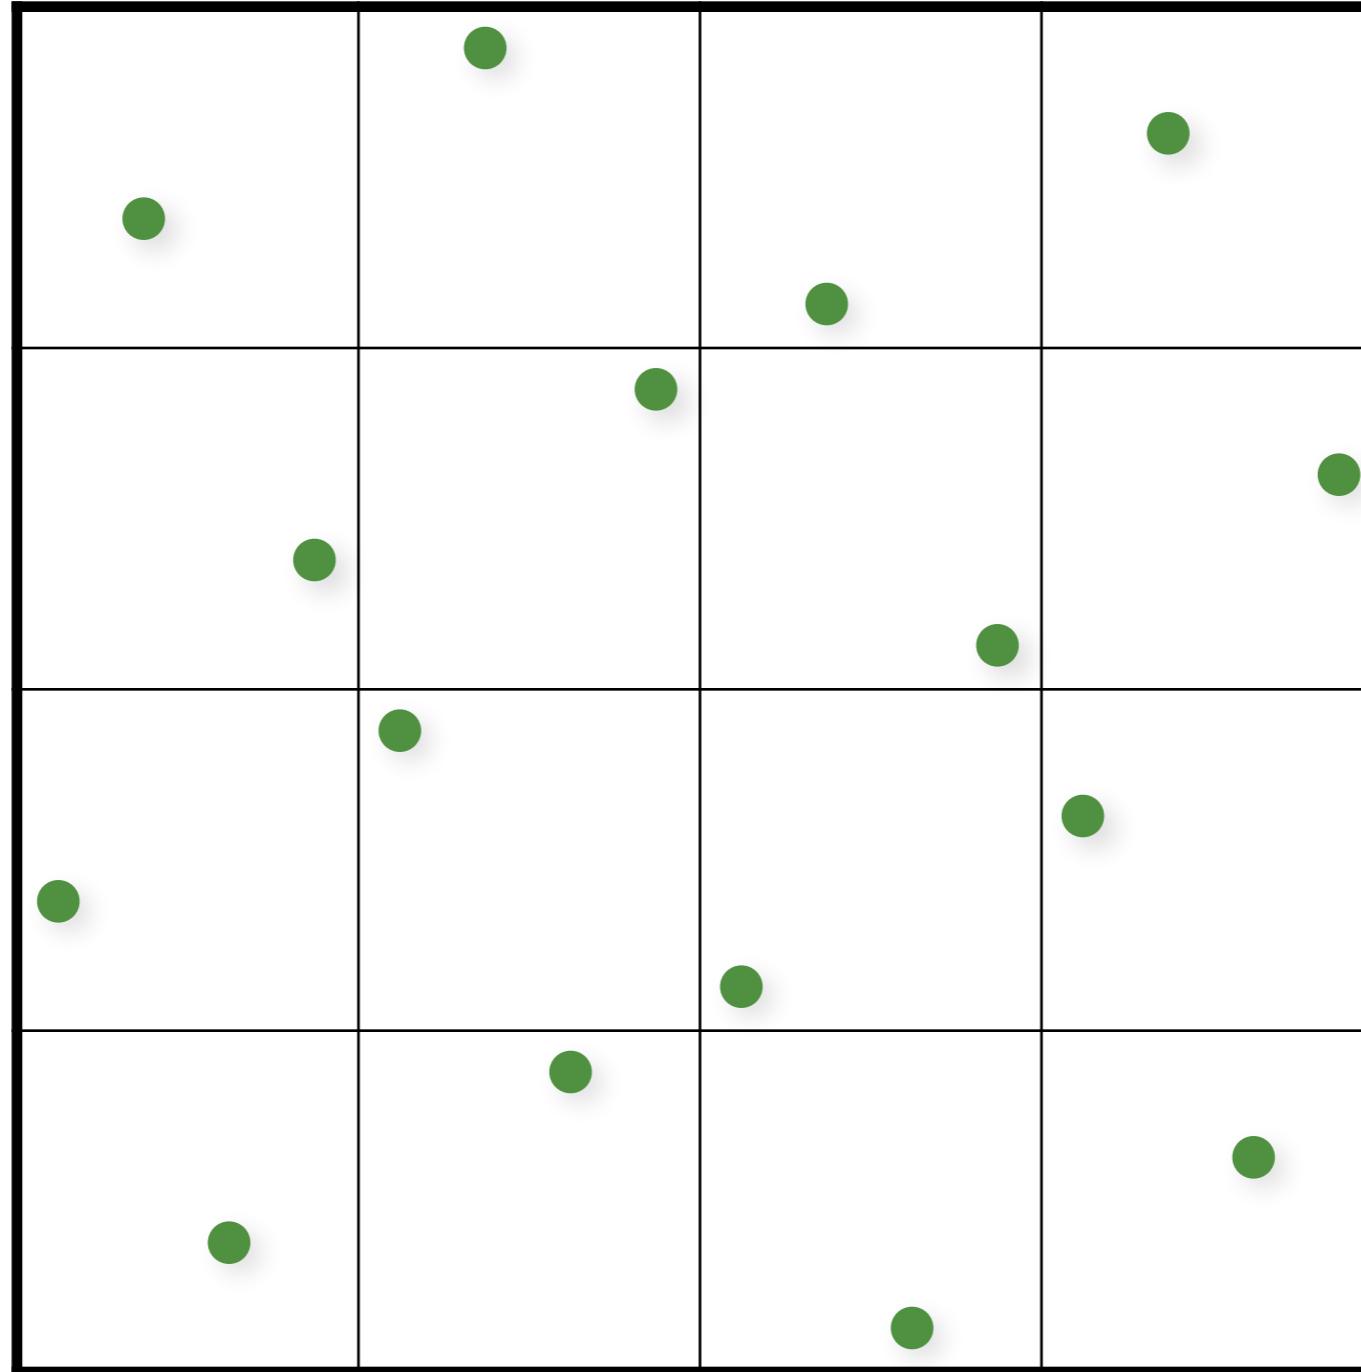
1 sample in each “elementary interval”

(0,2)-Sequences



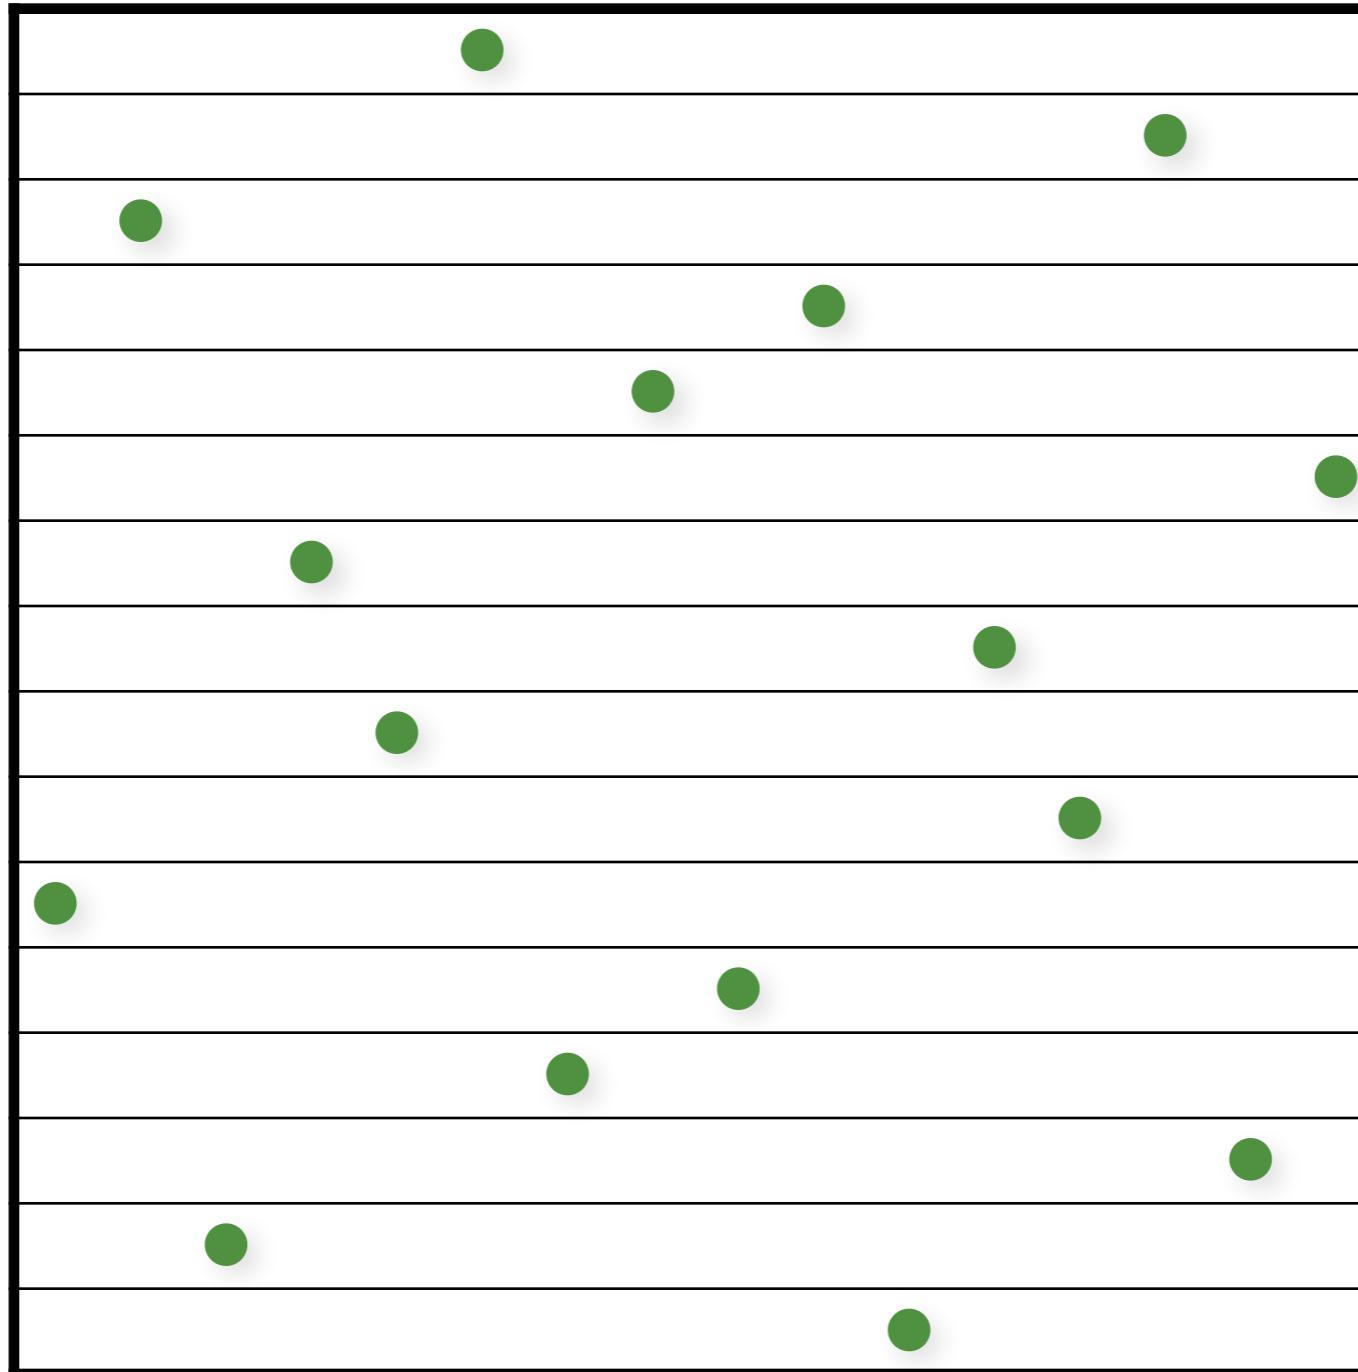
1 sample in each “elementary interval”

(0,2)-Sequences



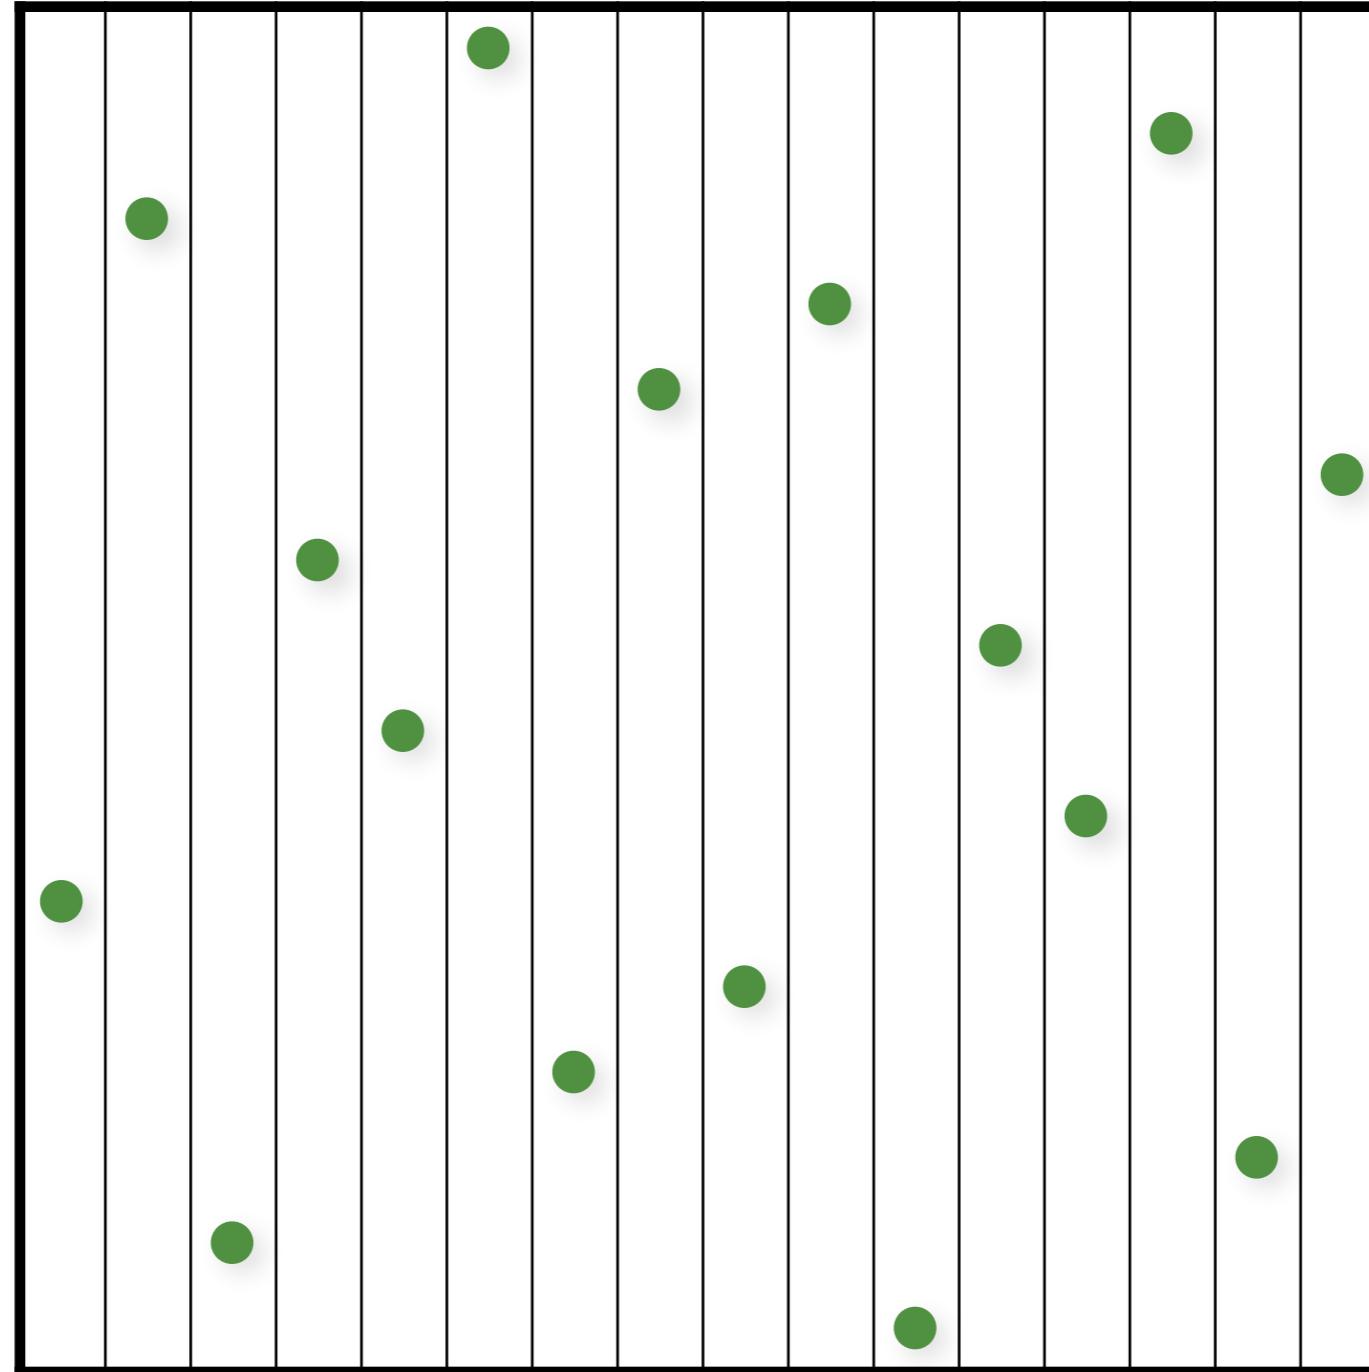
1 sample in each “elementary interval”

(0,2)-Sequences



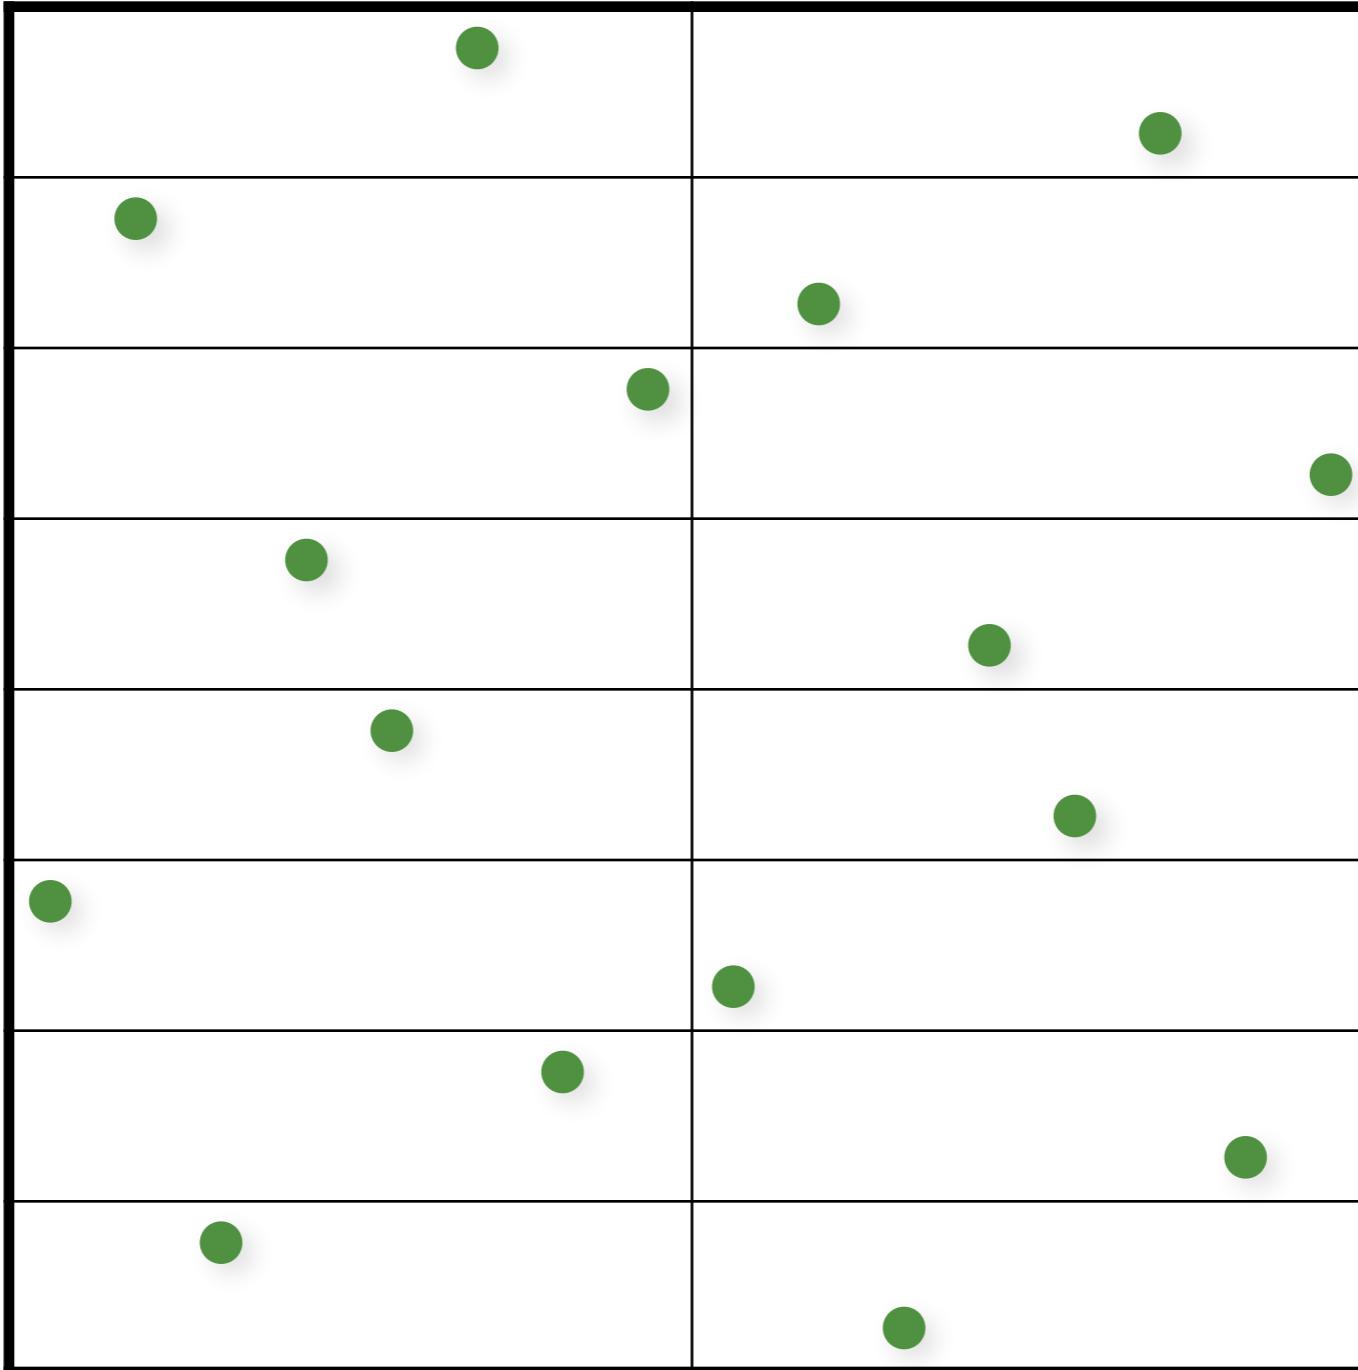
1 sample in each “elementary interval”

(0,2)-Sequences



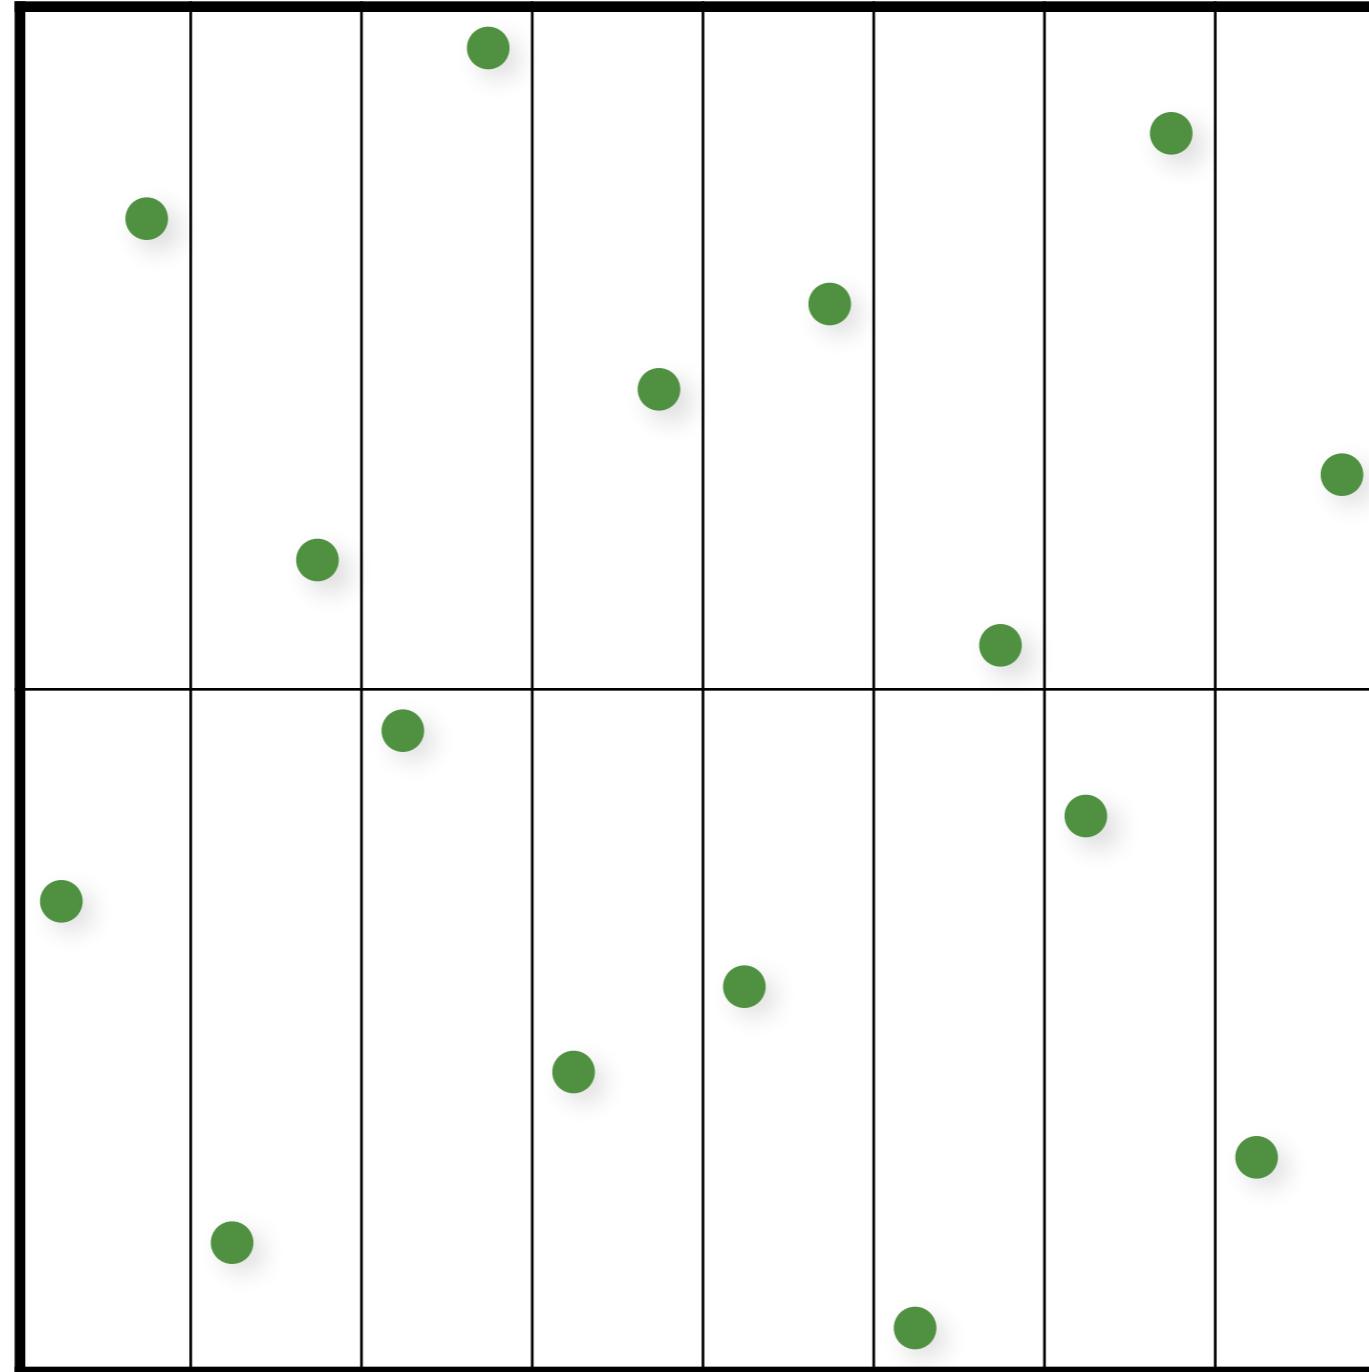
1 sample in each “elementary interval”

(0,2)-Sequences



1 sample in each “elementary interval”

(0,2)-Sequences



1 sample in each “elementary interval”

Additional resources

"Enumerating Quasi-Monte Carlo Point Sequences
in Elementary Intervals."

L. Grünschloß, M. Raab and A. Keller. In *Monte Carlo and Quasi-Monte Carlo Methods 2010*

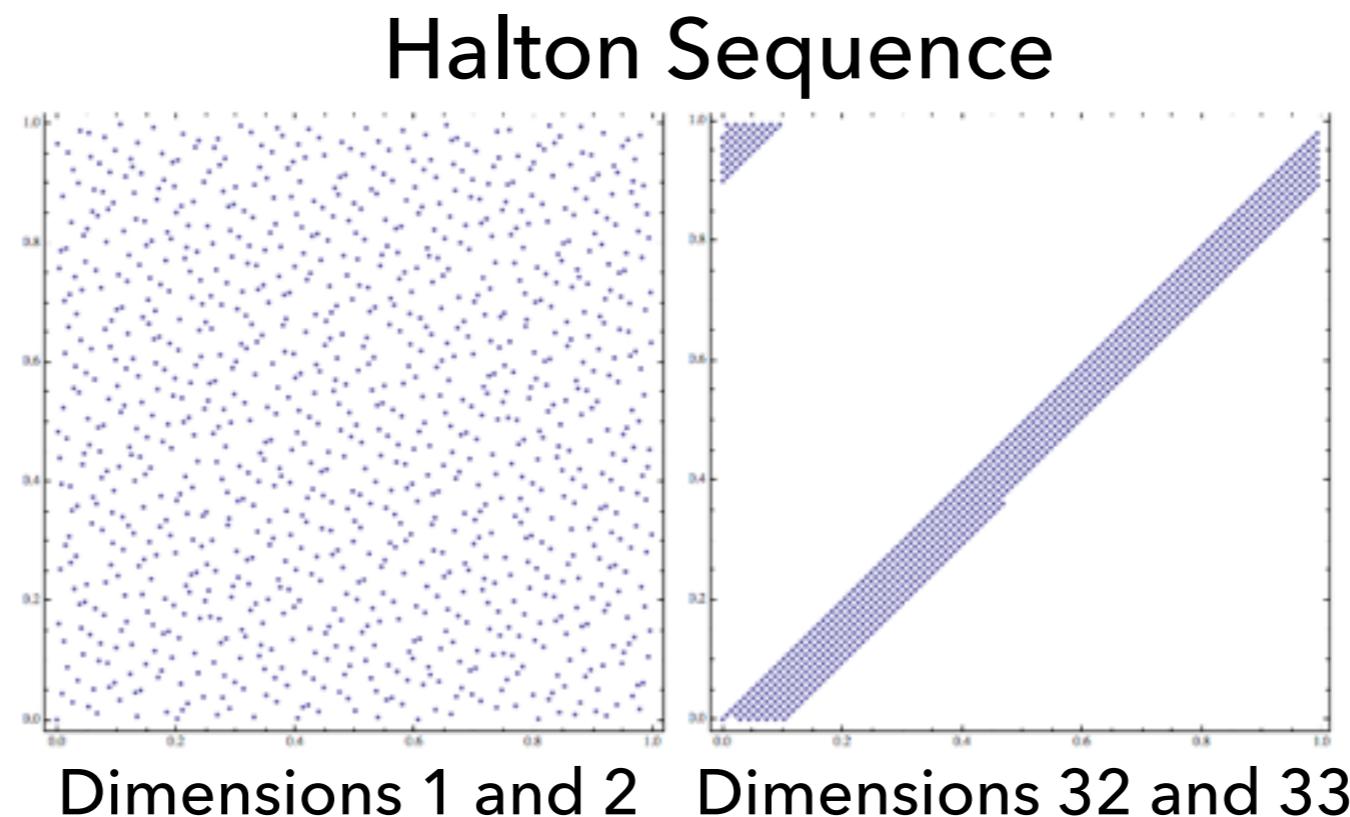
<http://gruenschloss.org/>

Many more...

- Sobol
- Faure
- Larcher-Pillichshammer
- Folded Radical Inverse
- (t,s) -sequences & (t,m,s) -nets
- Scrambling/randomization
- much more...

Challenges

- LD sequence identical for multiple runs
 - cannot average independent images!
 - no “random” seed
- Quality decreases in higher dimensions

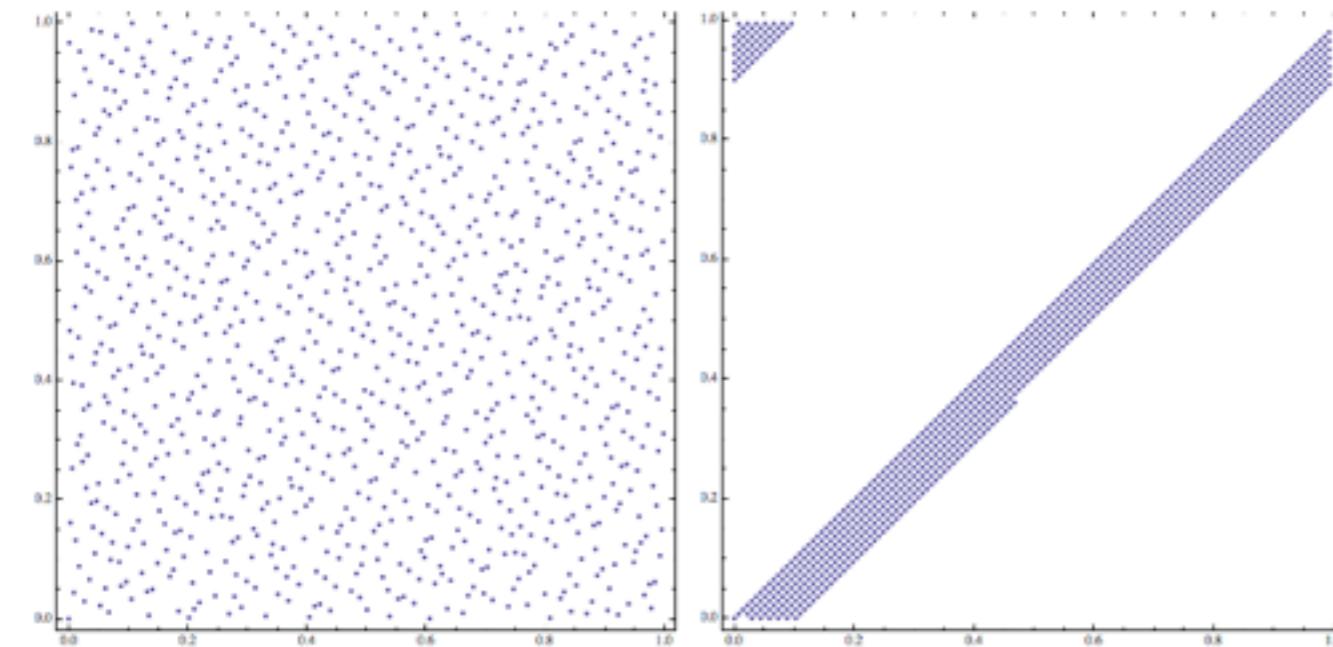


Randomized/Scrambled Sequences

- Random permutations: compute a permutation table for the order of the digits and use it when computing the radical inverse

$$\Phi_b(n) = 0.\pi(d_1)\pi(d_2)\dots\pi(d_m)$$

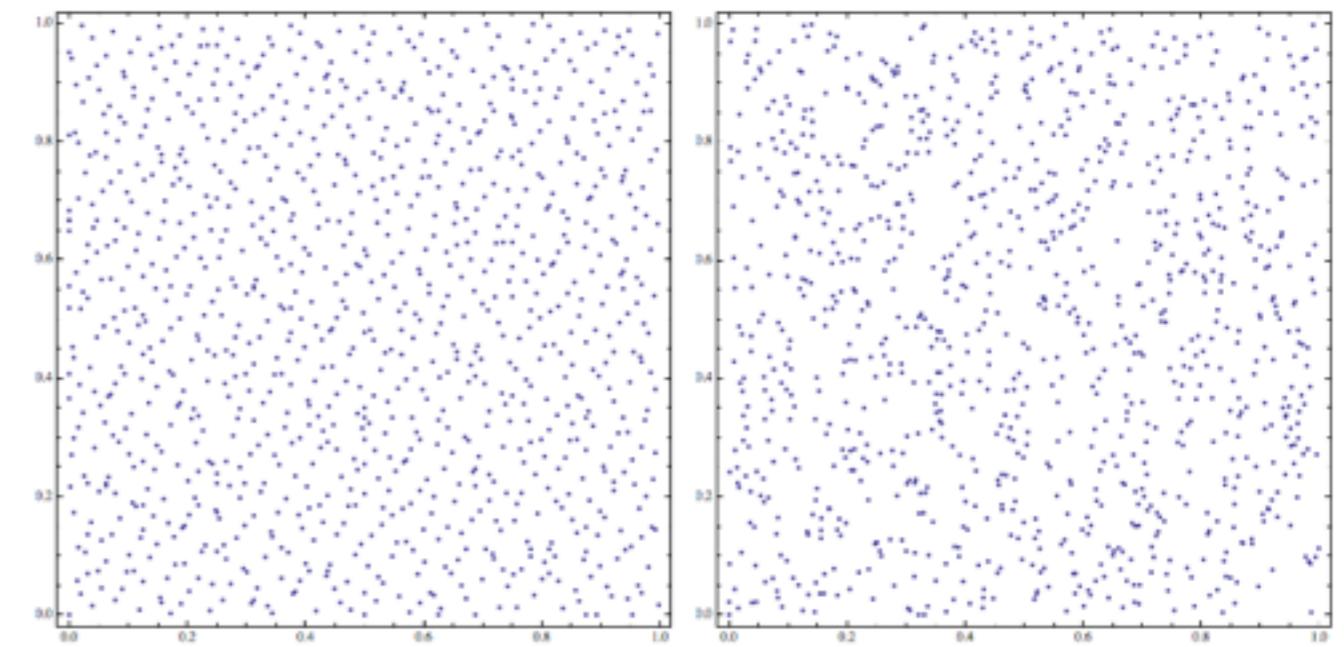
Without scrambling



Dimensions 1 and 2

Dimensions 32 and 33

With scrambling



Dimensions 1 and 2

Dimensions 32 and 33

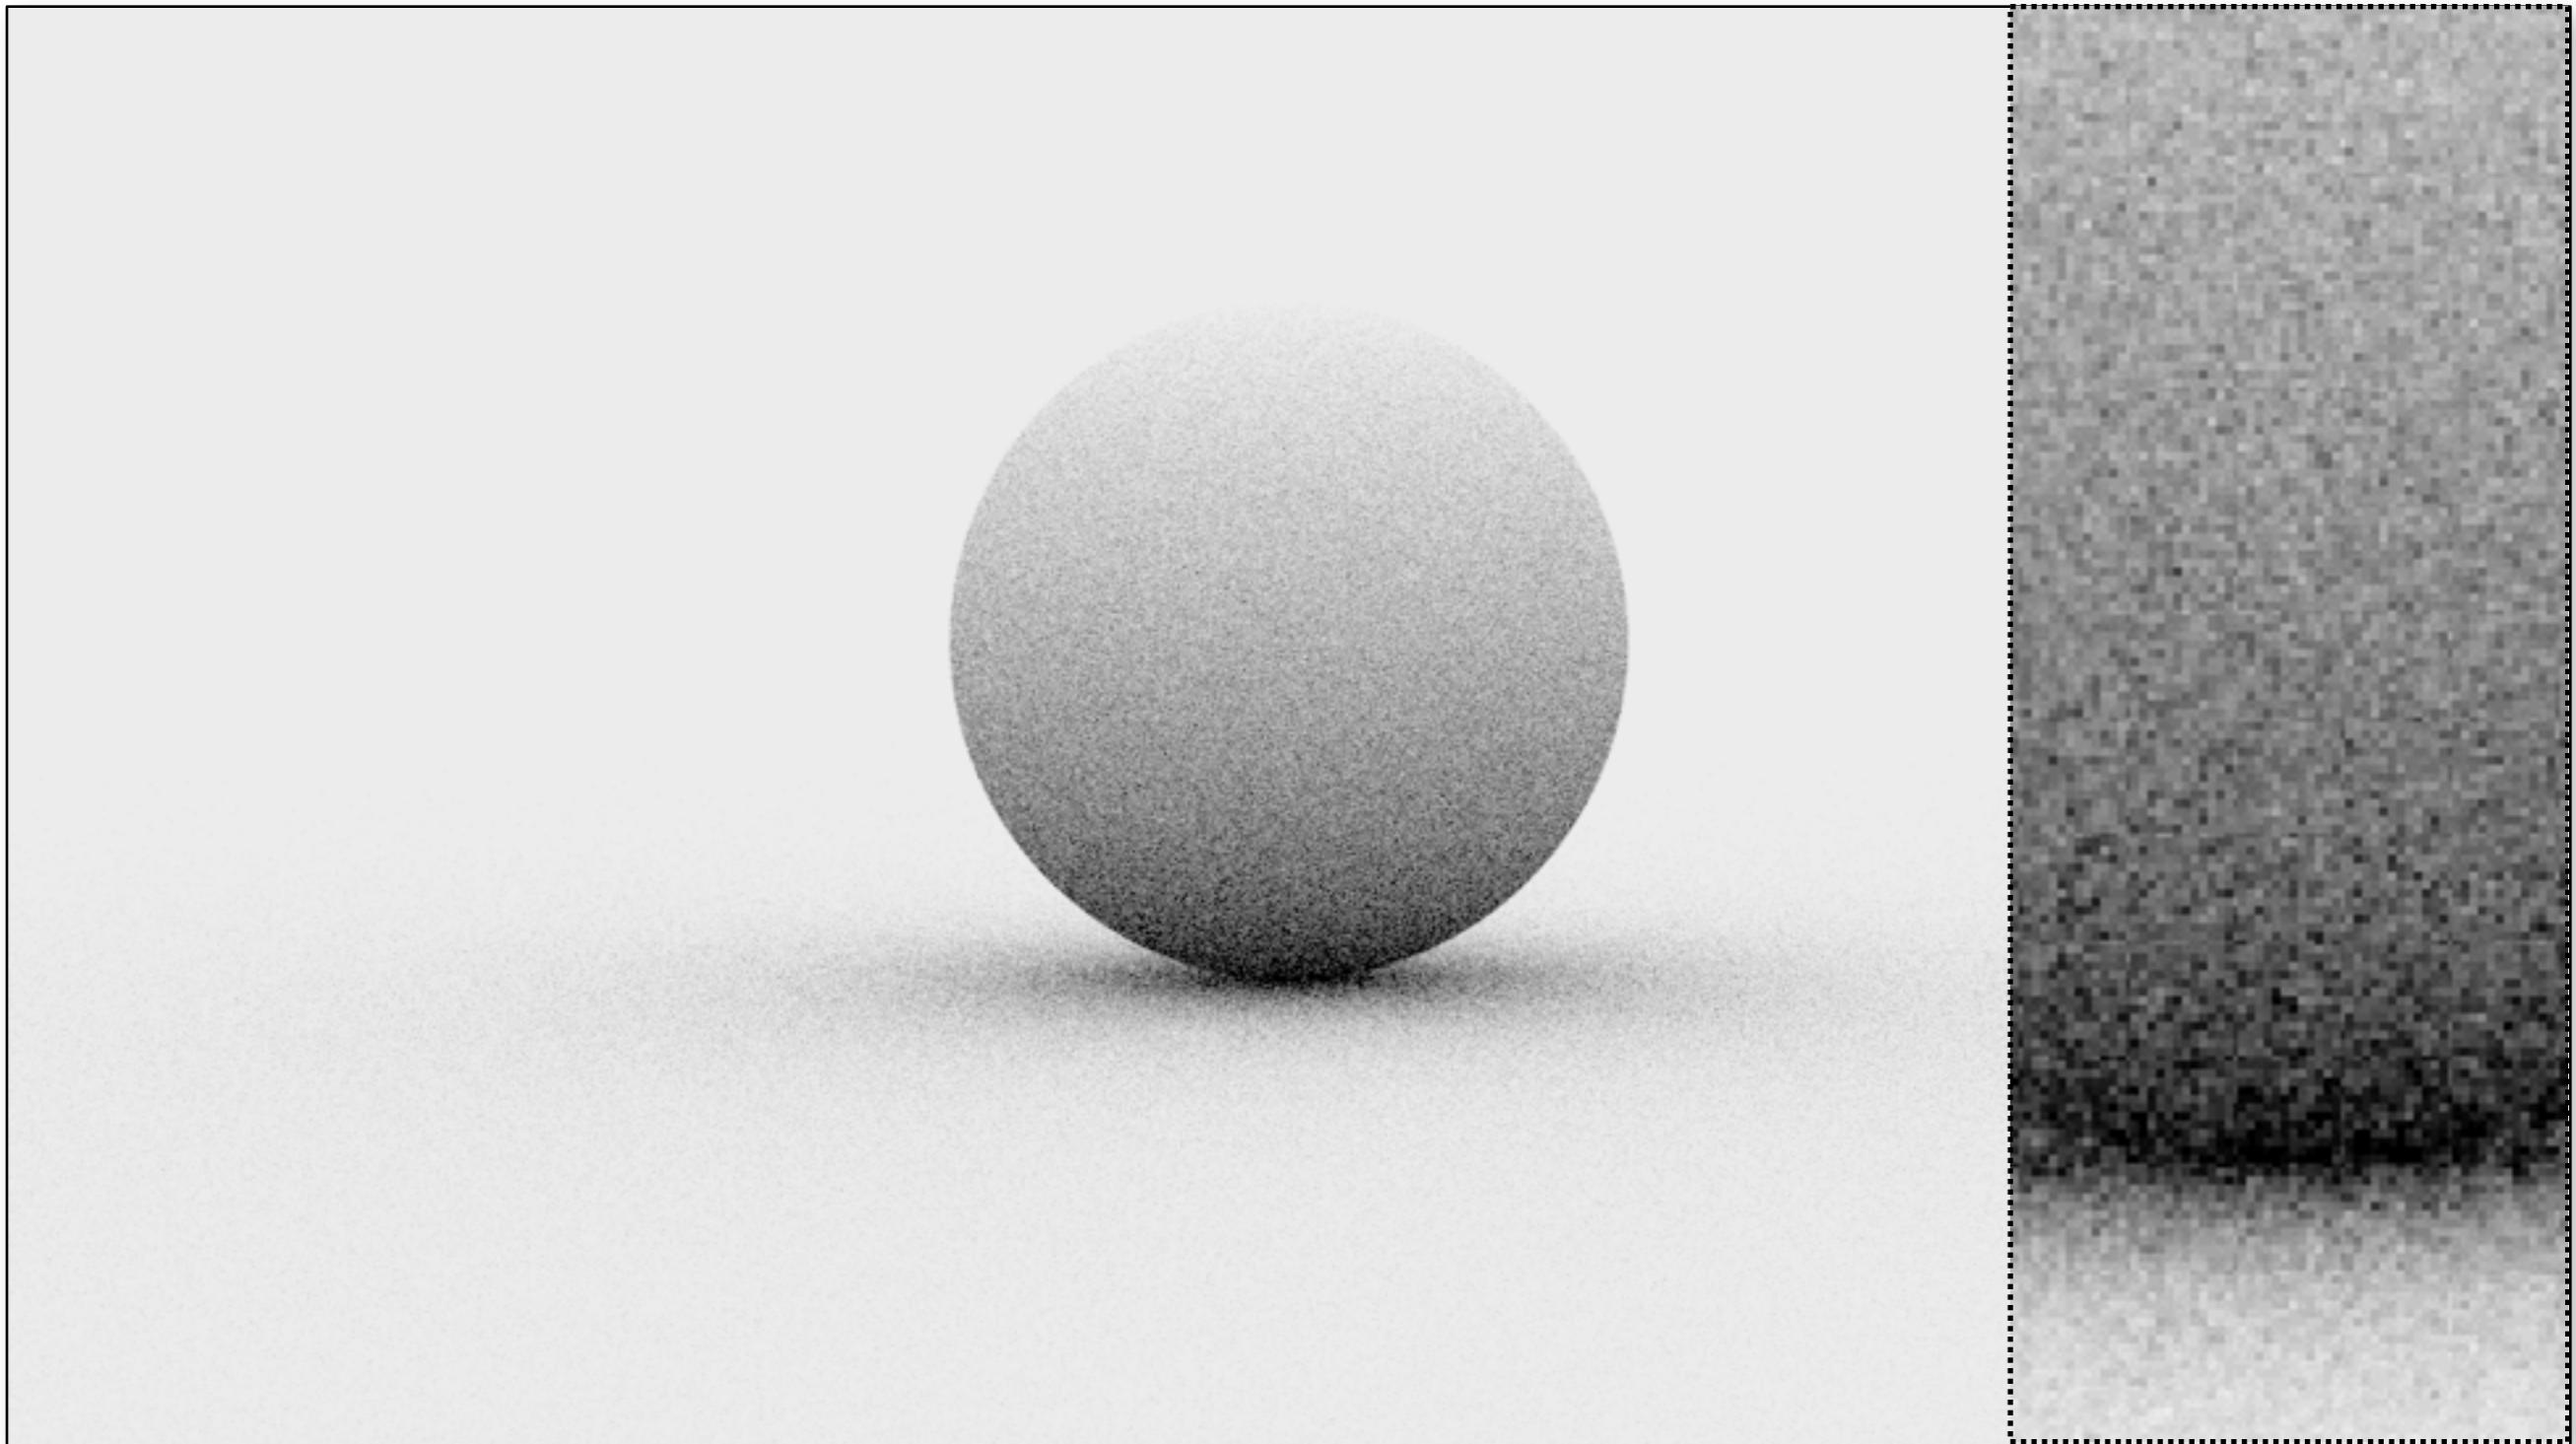
Randomized/Scrambled Sequences

- Random permutations: compute a permutation table for the order of the digits and use it when computing the radical inverse
 - Can be done very efficiently for base 2 with XOR operation
- See PBR2 Ch7 for details

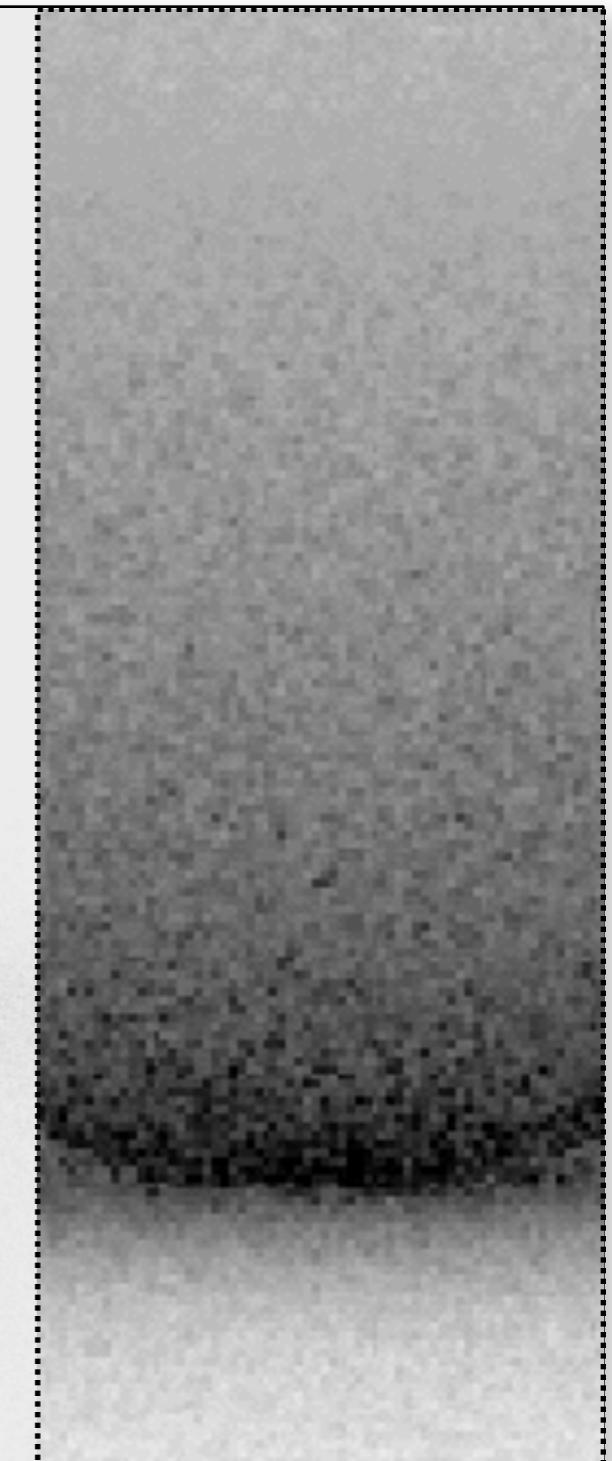
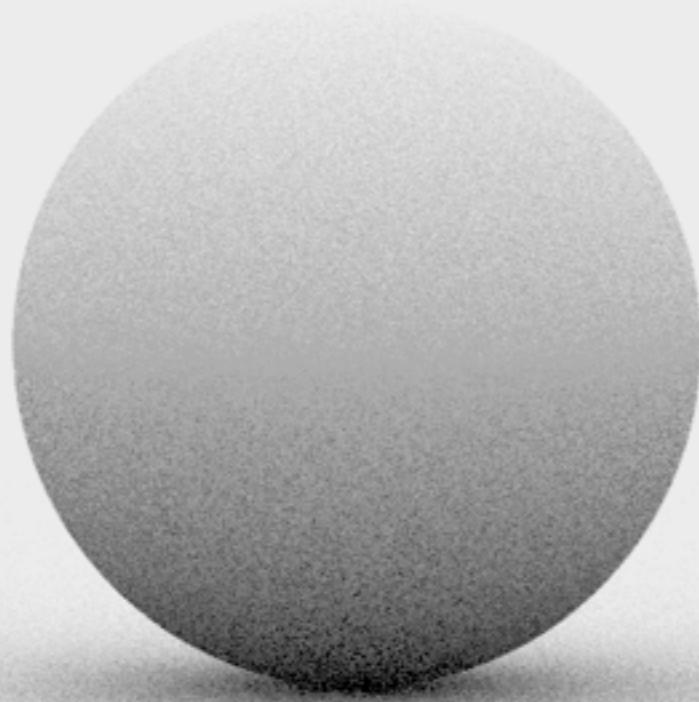
Scrambled Radical Inverse (Base 2)

```
float vanDerCorputRIU(uint n, uint scramble = 0)
{
    n = (n << 16) | (n >> 16);
    n = ((n & 0x00ff00ff) << 8) | ((n & 0xff00ff00) >> 8);
    n = ((n & 0x0f0f0f0f) << 4) | ((n & 0xf0f0f0f0) >> 4);
    n = ((n & 0x33333333) << 2) | ((n & 0xcccccccc) >> 2);
    n = ((n & 0x55555555) << 1) | ((n & 0xaaaaaaaa) >> 1);
    n ^= scramble;
    return n / float (0x100000000LL);
}
```

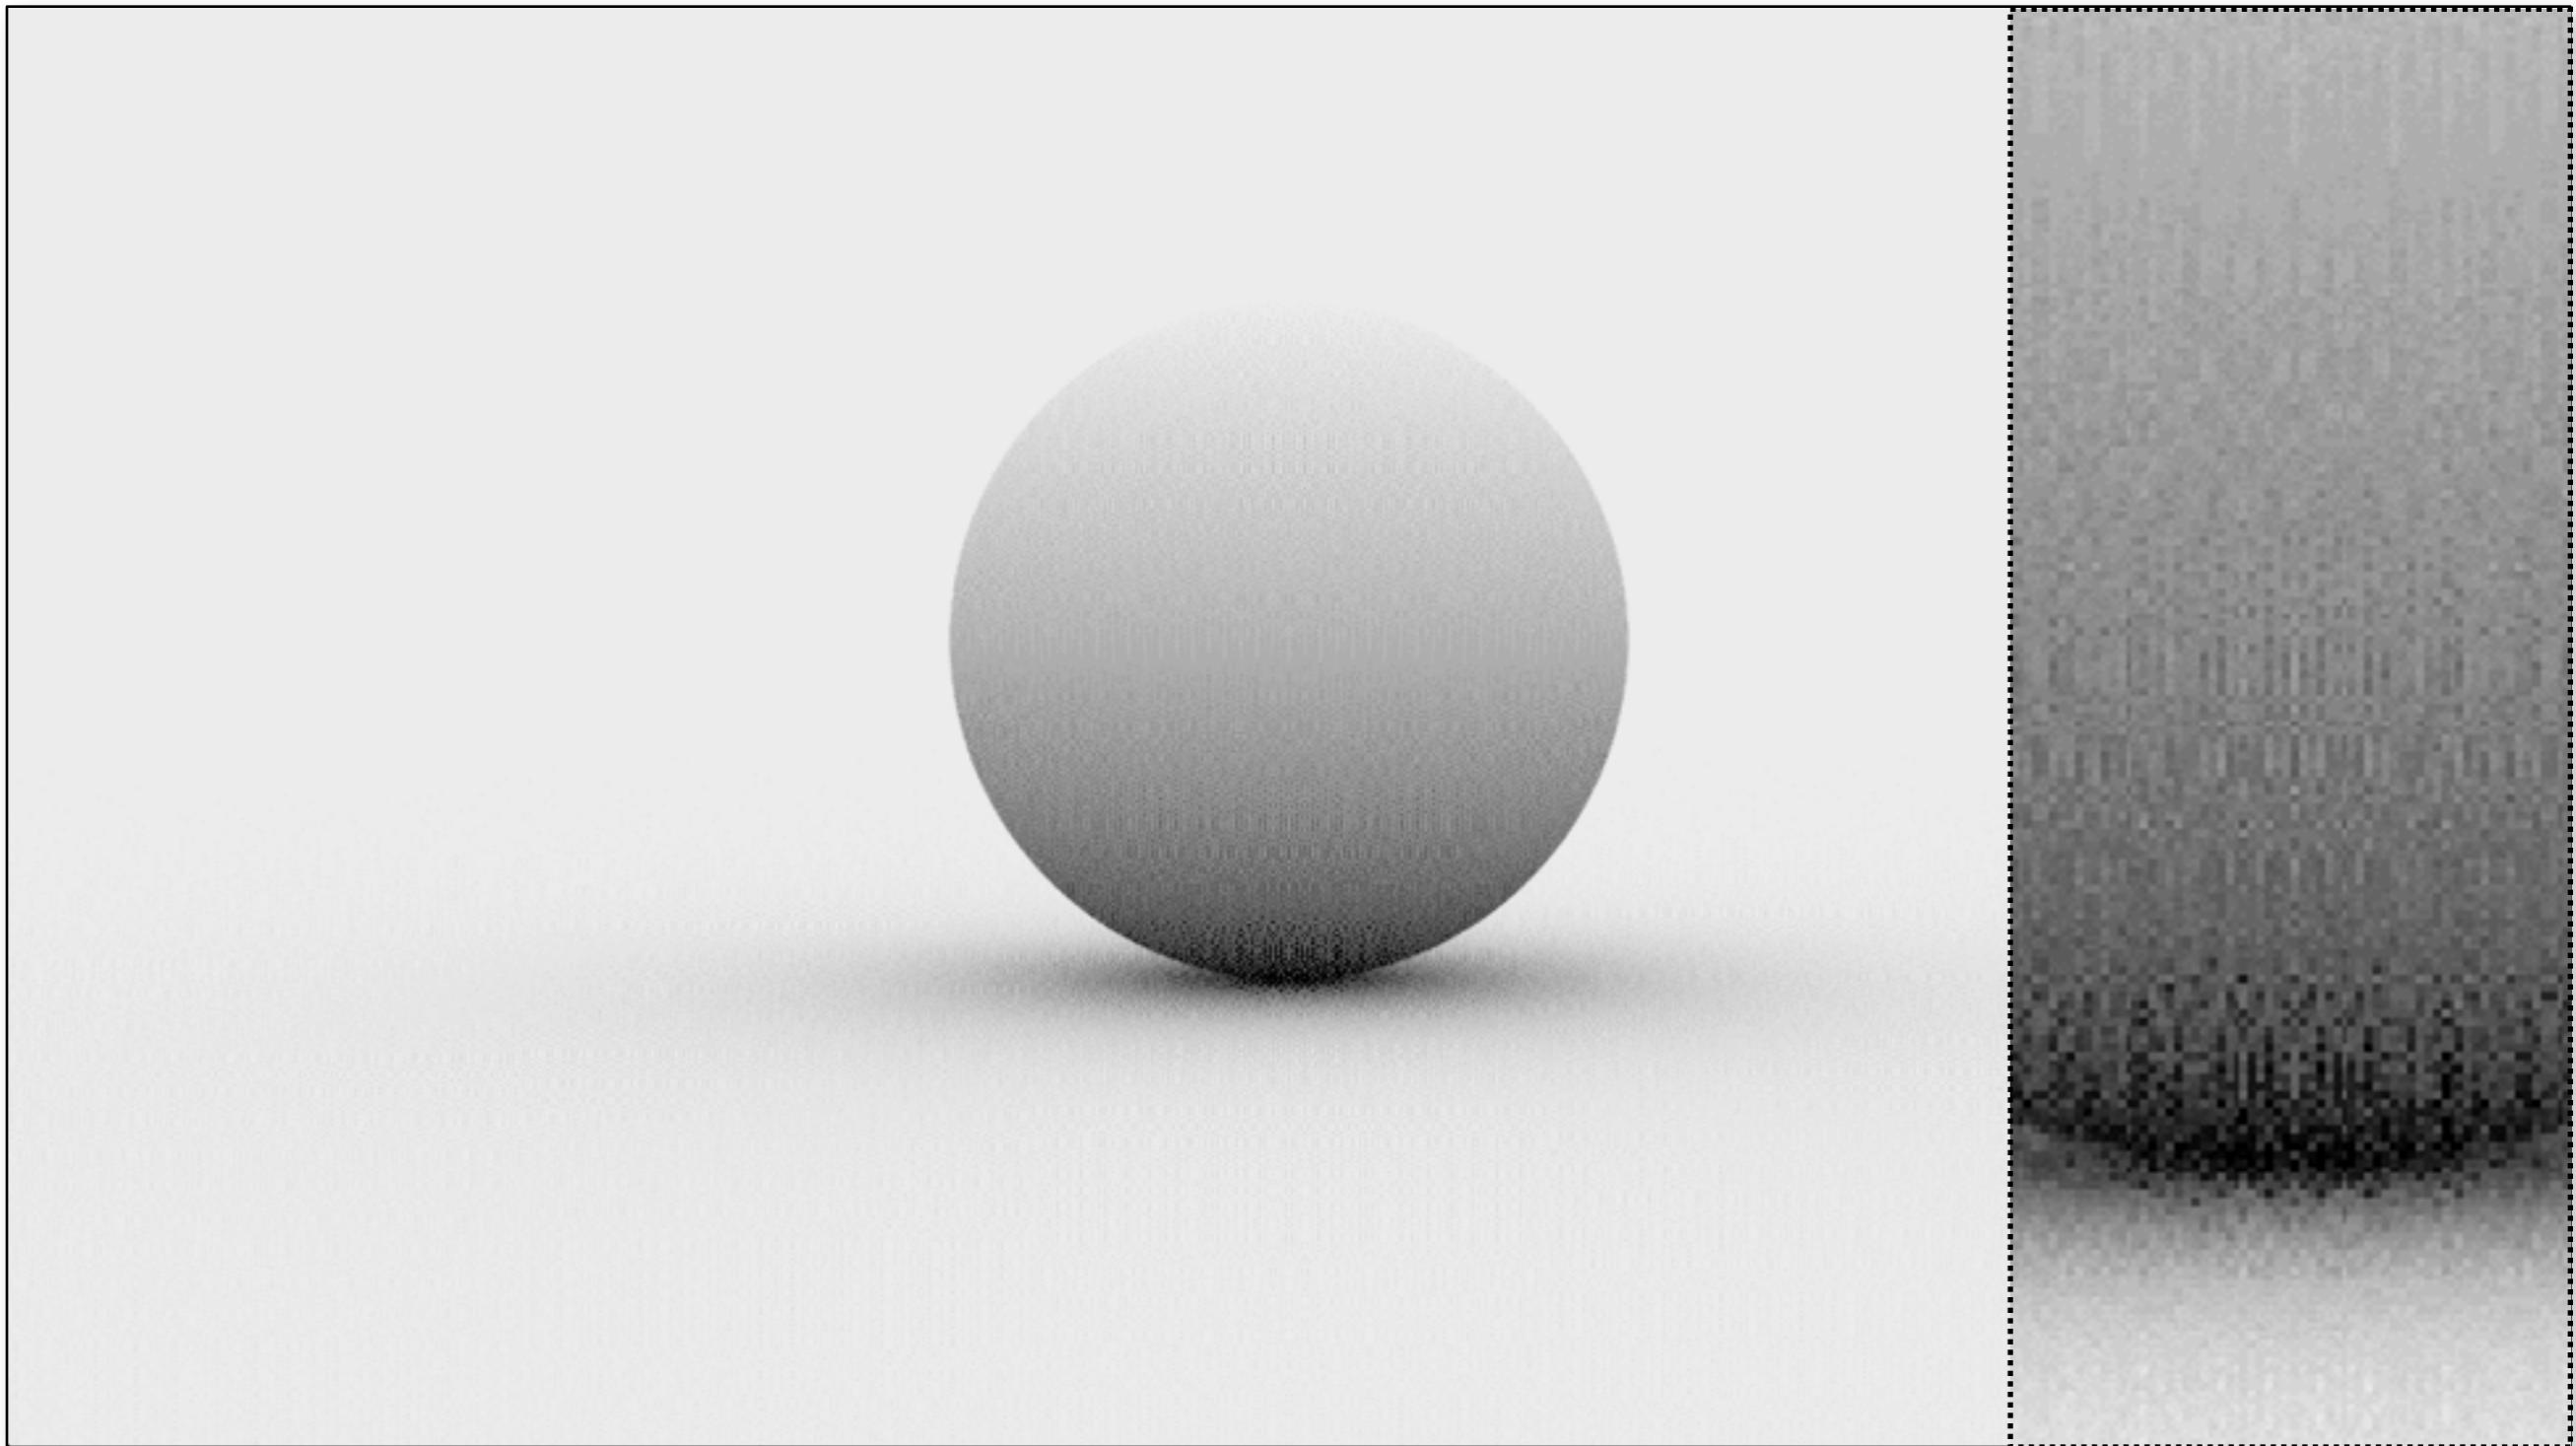
Monte Carlo (16 random samples)



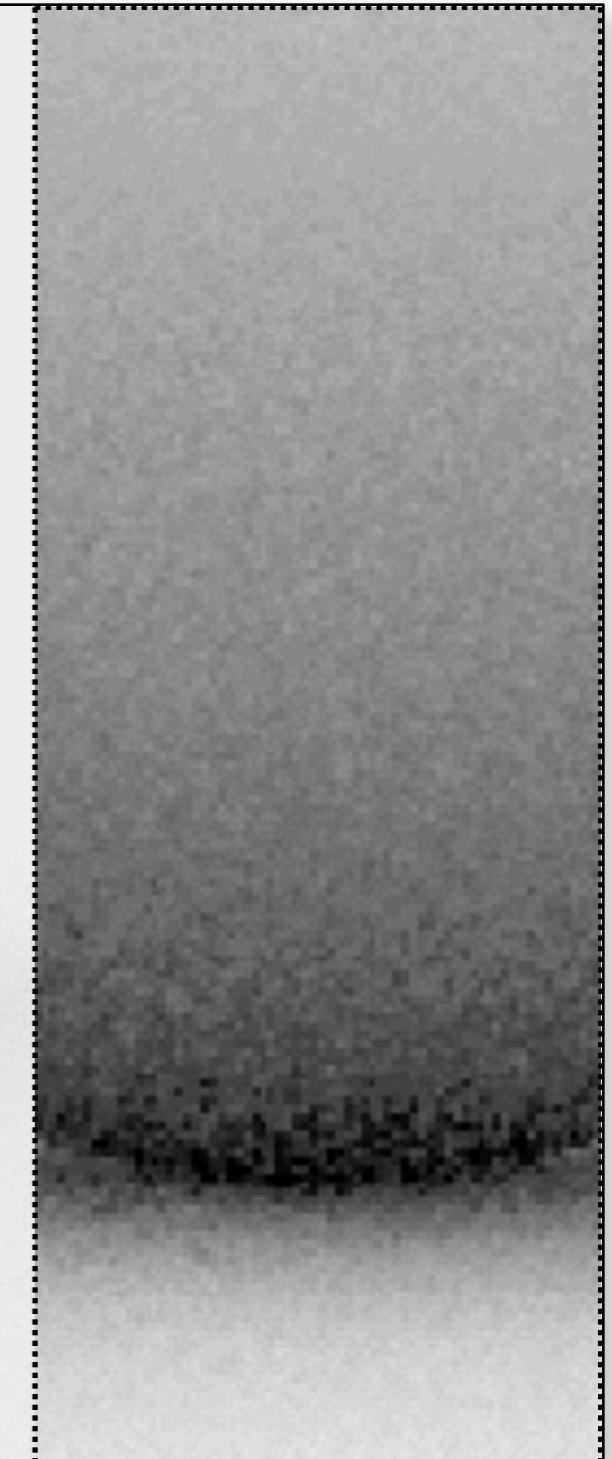
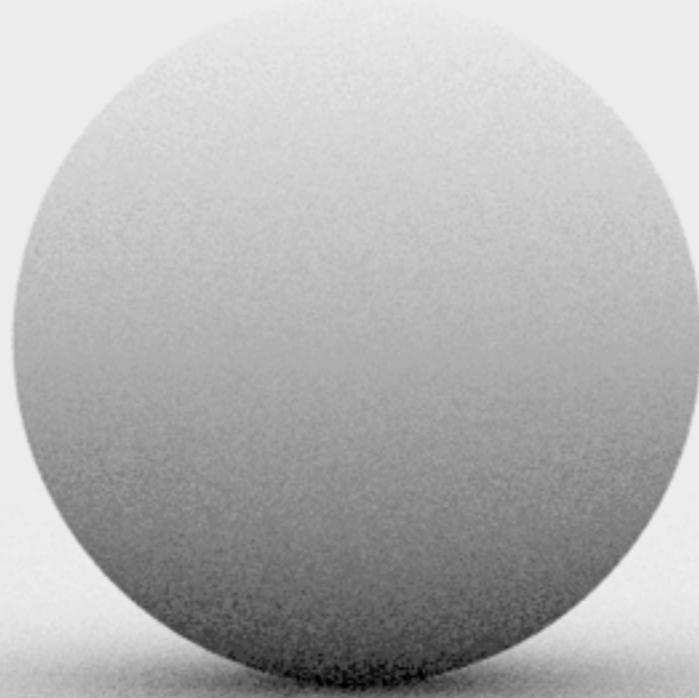
Monte Carlo (16 stratified samples)



Quasi-Monte Carlo (16 Halton samples)



Scrambled Quasi-Monte Carlo



scrambled Larcher-Pillichshammer sequence

Implementation tips

- Using QMC can often lead to unintuitive, difficult-to-debug problems.
 - Always code up MC algorithms first, using random numbers, to ensure correctness
 - Only after confirming correctness, slowly incorporate QMC into the mix

Questions?
