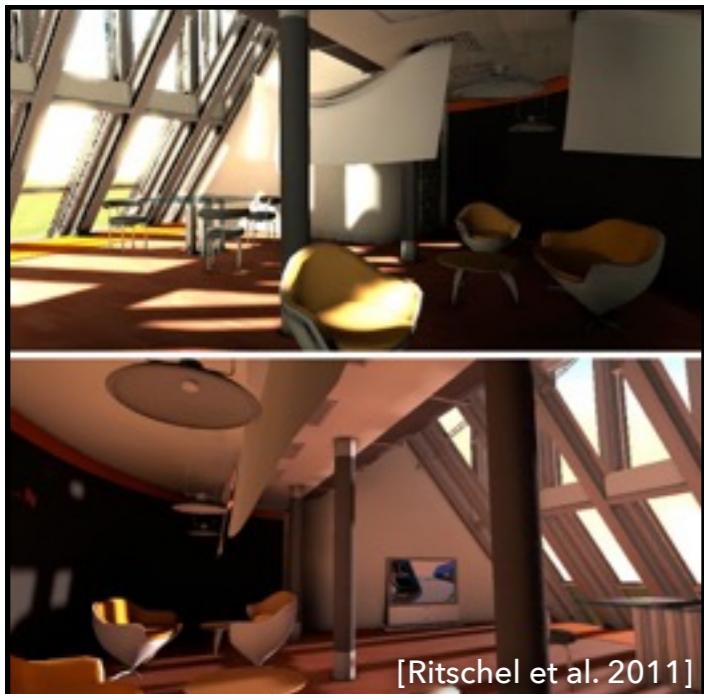


CS 87/187, Spring 2016

RENDERING ALGORITHMS

Many-Light Rendering



Prof. Wojciech Jarosz

wojciech.k.jarosz@dartmouth.edu

(most slides courtesy of Jan Novák)

Administrivia

- Assignment 3 due tomorrow
- Assignment 4 will be posted tonight or tomorrow
 - due in 2 weeks, May 11
- Midterm (tentative) in class on Thursday May 12

Characteristics of Estimators

- *Unbiased* estimator

- expected value equals the true value being estimated

$$E[F] = \int f(x) dx$$

- variance (noise) is the only error
 - averaging infinitely many estimates (each with finite number of samples) also yields the correct answer

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{i=1}^N \langle F^k \rangle = \int f(x) dx$$

Characteristics of Estimators

- *Bias* of an estimator

- difference between the expected value of the estimator and the true value being estimated

$$\beta = E[F] - \int f(x) dx$$

- expected average difference, expected error
- averaging infinitely many estimates yields the correct answer plus the bias

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{i=1}^N \langle F^k \rangle = \int f(x) dx + \beta$$

Characteristics of Estimators

- *Consistent* estimator

- bias disappears in the limit

$$\lim_{N \rightarrow \infty} E[F] = \int f(x) dx$$

- *Consistent* estimators and *unbiased* estimators

- are asymptotically equivalent

- both need an infinite number of samples to reduce the error to zero

Characteristics of Estimators

- *Mean Squared Error* (MSE) of an estimator
 - combines variance and squared bias
$$\text{MSE}[F] = \text{Var}[F] + \text{Bias}[F]^2$$
- *Root Mean Squared Error* (RMSE)
 - has the same units as the quantity being estimated
 - for unbiased estimators equal to std. deviation

$$\text{RMSE}[F] = \sqrt{\text{MSE}[F]}$$

Rendering Techniques

- Examples of unbiased methods
 - Path tracing
 - Light tracing
 - Bidirectional path tracing
 - Many-light rendering (in theory)
- Examples of biased/consistent methods
 - Many-light rendering (practical implementations)
 - Photon mapping
 - Progressive photon mapping
 - Irradiance caching

Last time

Equations

- Rendering equation:

$$L_o(\mathbf{x}, \vec{\omega}_o) = L_e(\mathbf{x}, \vec{\omega}_o) + \int_{H^2} f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_o) L_i(\mathbf{x}, \vec{\omega}_i) \cos \theta_i d\vec{\omega}_i$$

- Measurement equation:

$$I_j = \int_{A_{\text{film}}} \int_{H^2} W_e(\mathbf{x}, \vec{\omega}) L_i(\mathbf{x}, \vec{\omega}) \cos \theta d\vec{\omega} d\mathbf{x}$$

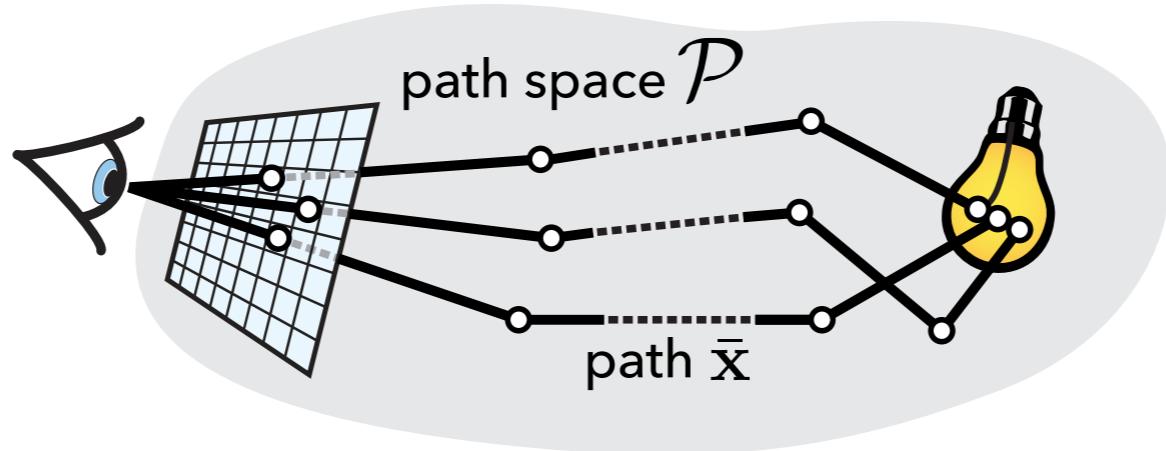
Duality of Radiance & Importance

$$\begin{aligned} I_j &= \int_{A_{\text{film}}} \int_{H^2} W_e(\mathbf{x}, \vec{\omega}) L_i(\mathbf{x}, \vec{\omega}) \cos \theta d\vec{\omega} d\mathbf{x} \\ &= \int_{A_{\text{light}}} \int_{H^2} W_i(\mathbf{z}, \vec{\omega}) L_e(\mathbf{z}, \vec{\omega}) \cos \theta d\vec{\omega} d\mathbf{z} \end{aligned}$$

The diagram illustrates the duality between radiance and importance terms. It shows two equivalent forms of the rendering equation. The top form uses \mathbf{x} and $\vec{\omega}$, while the bottom form uses \mathbf{z} and $\vec{\omega}$. Curved arrows point from the labels to the corresponding terms in each equation. The top row shows 'emitted importance' pointing to $W_e(\mathbf{x}, \vec{\omega})$ and 'incident radiance' pointing to $L_i(\mathbf{x}, \vec{\omega})$. The bottom row shows 'emitted radiance' pointing to $L_e(\mathbf{z}, \vec{\omega})$ and 'incident importance' pointing to $W_i(\mathbf{z}, \vec{\omega})$.

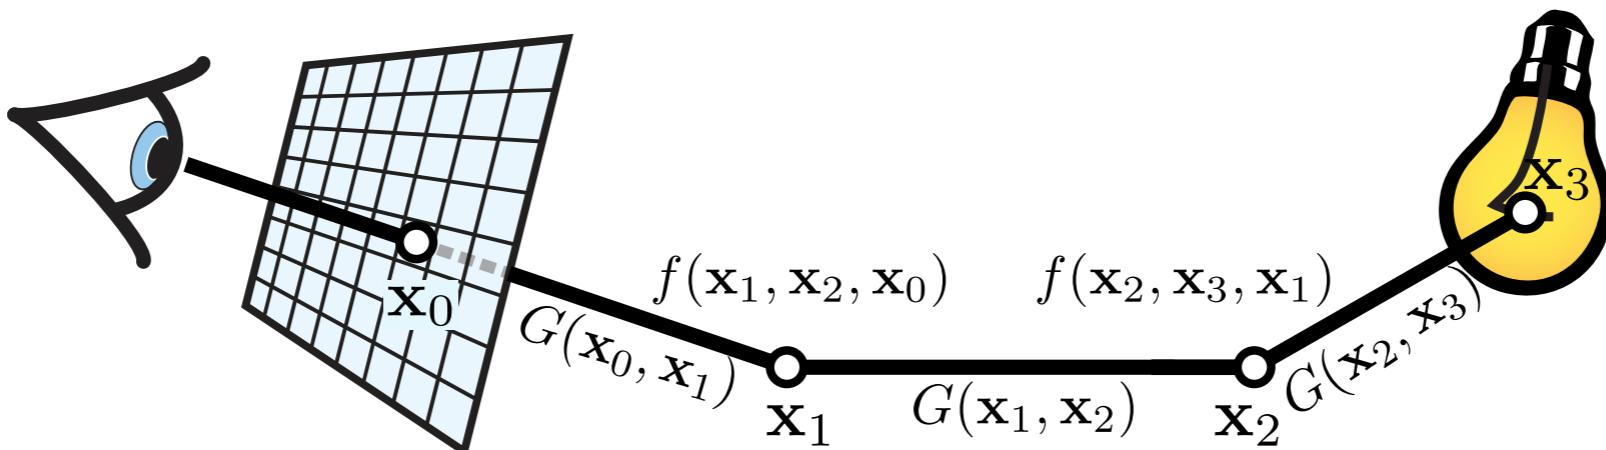
Path Integral Form of Measurement Eq.

$$I_j = \int_{\mathcal{P}} W_e(\mathbf{x}_0, \mathbf{x}_1) L_e(\mathbf{x}_k, \mathbf{x}_{k-1}) T(\bar{\mathbf{x}}) d\bar{\mathbf{x}}$$

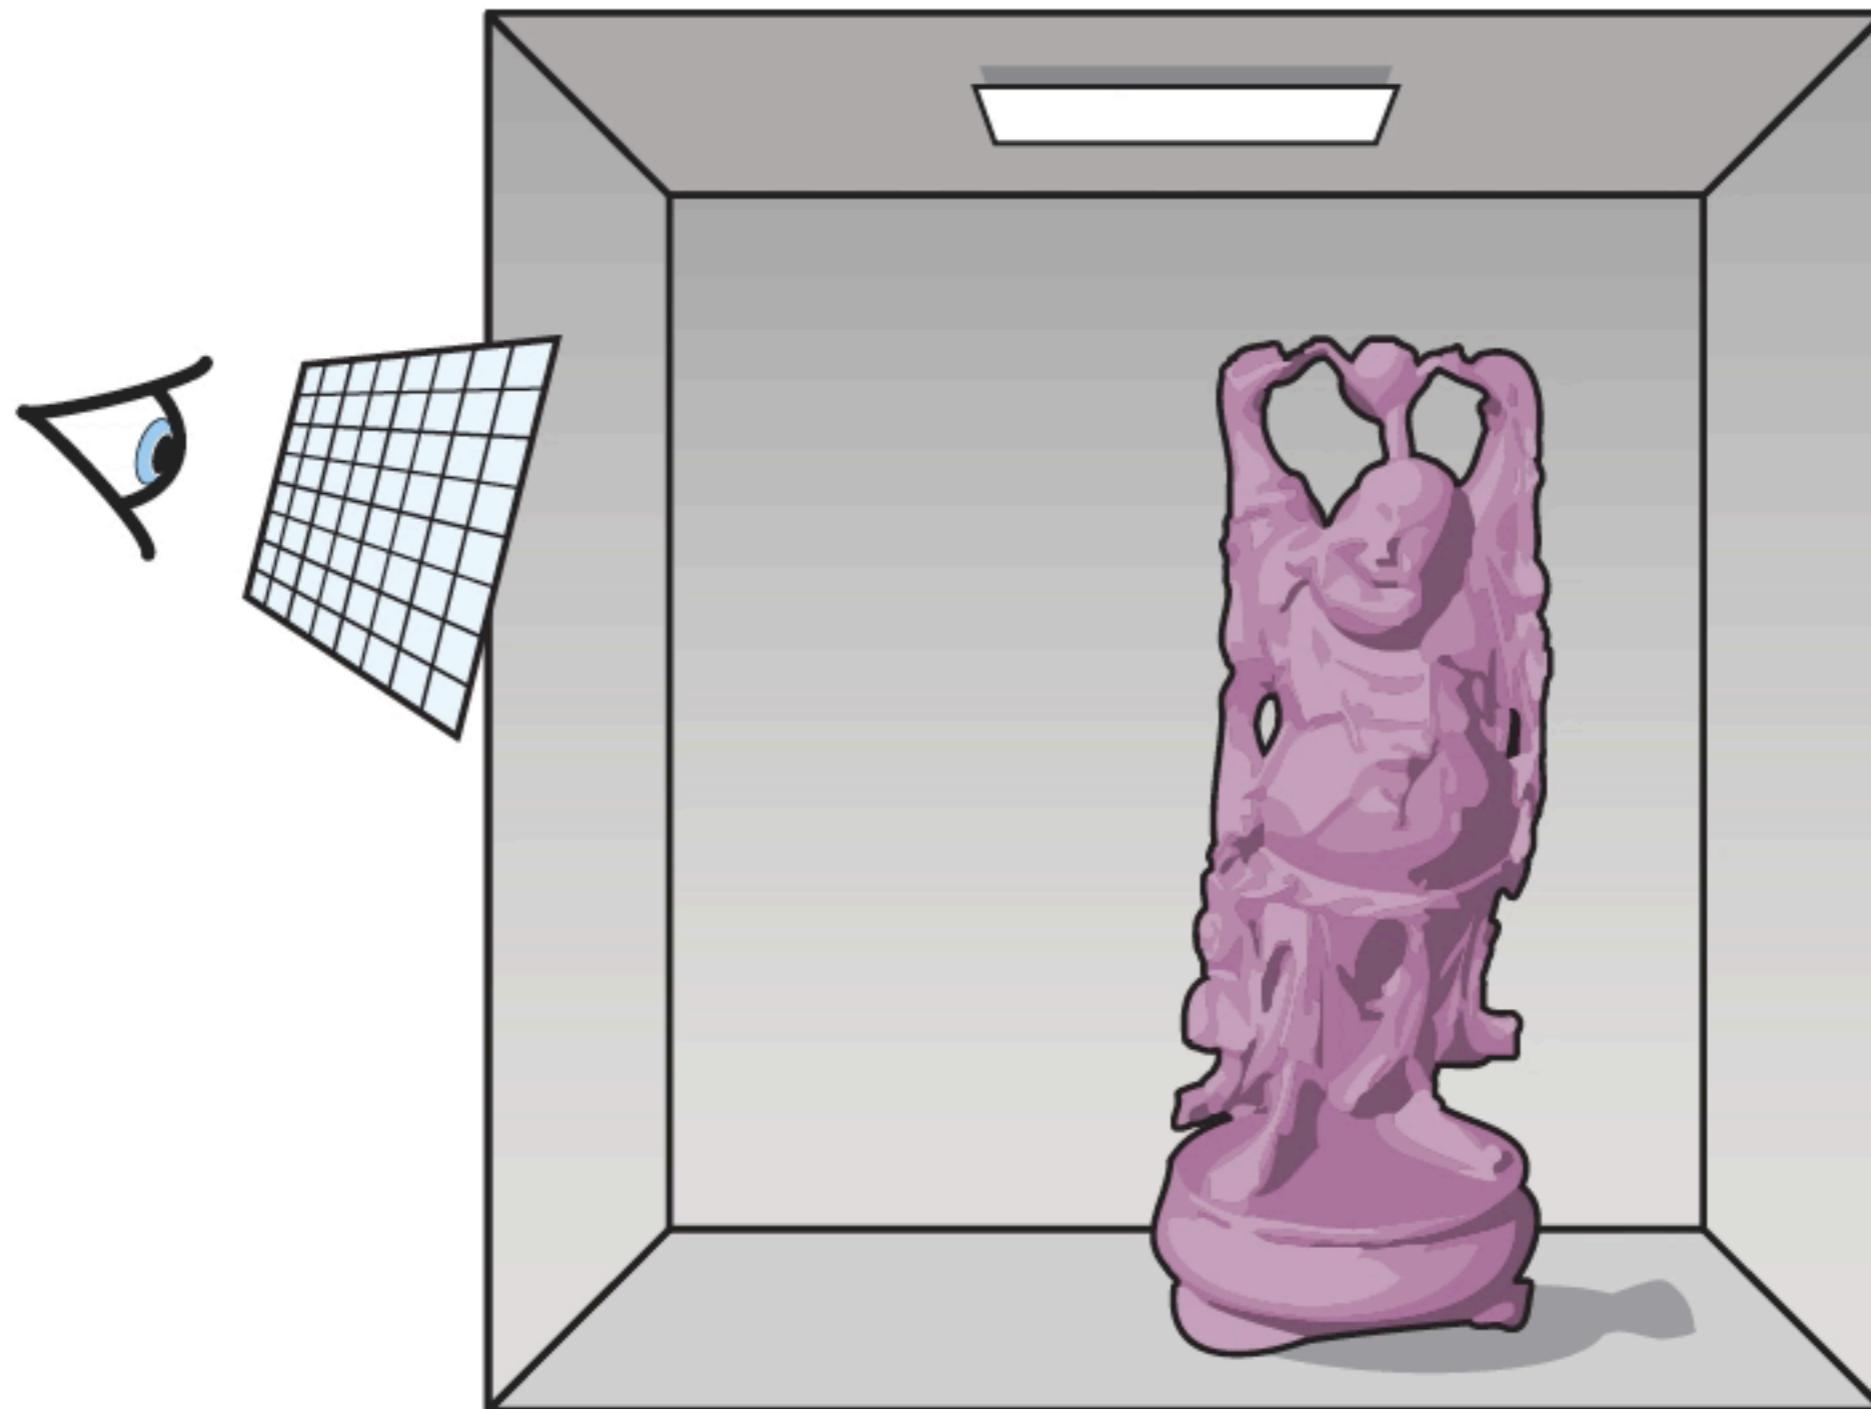


path throughput

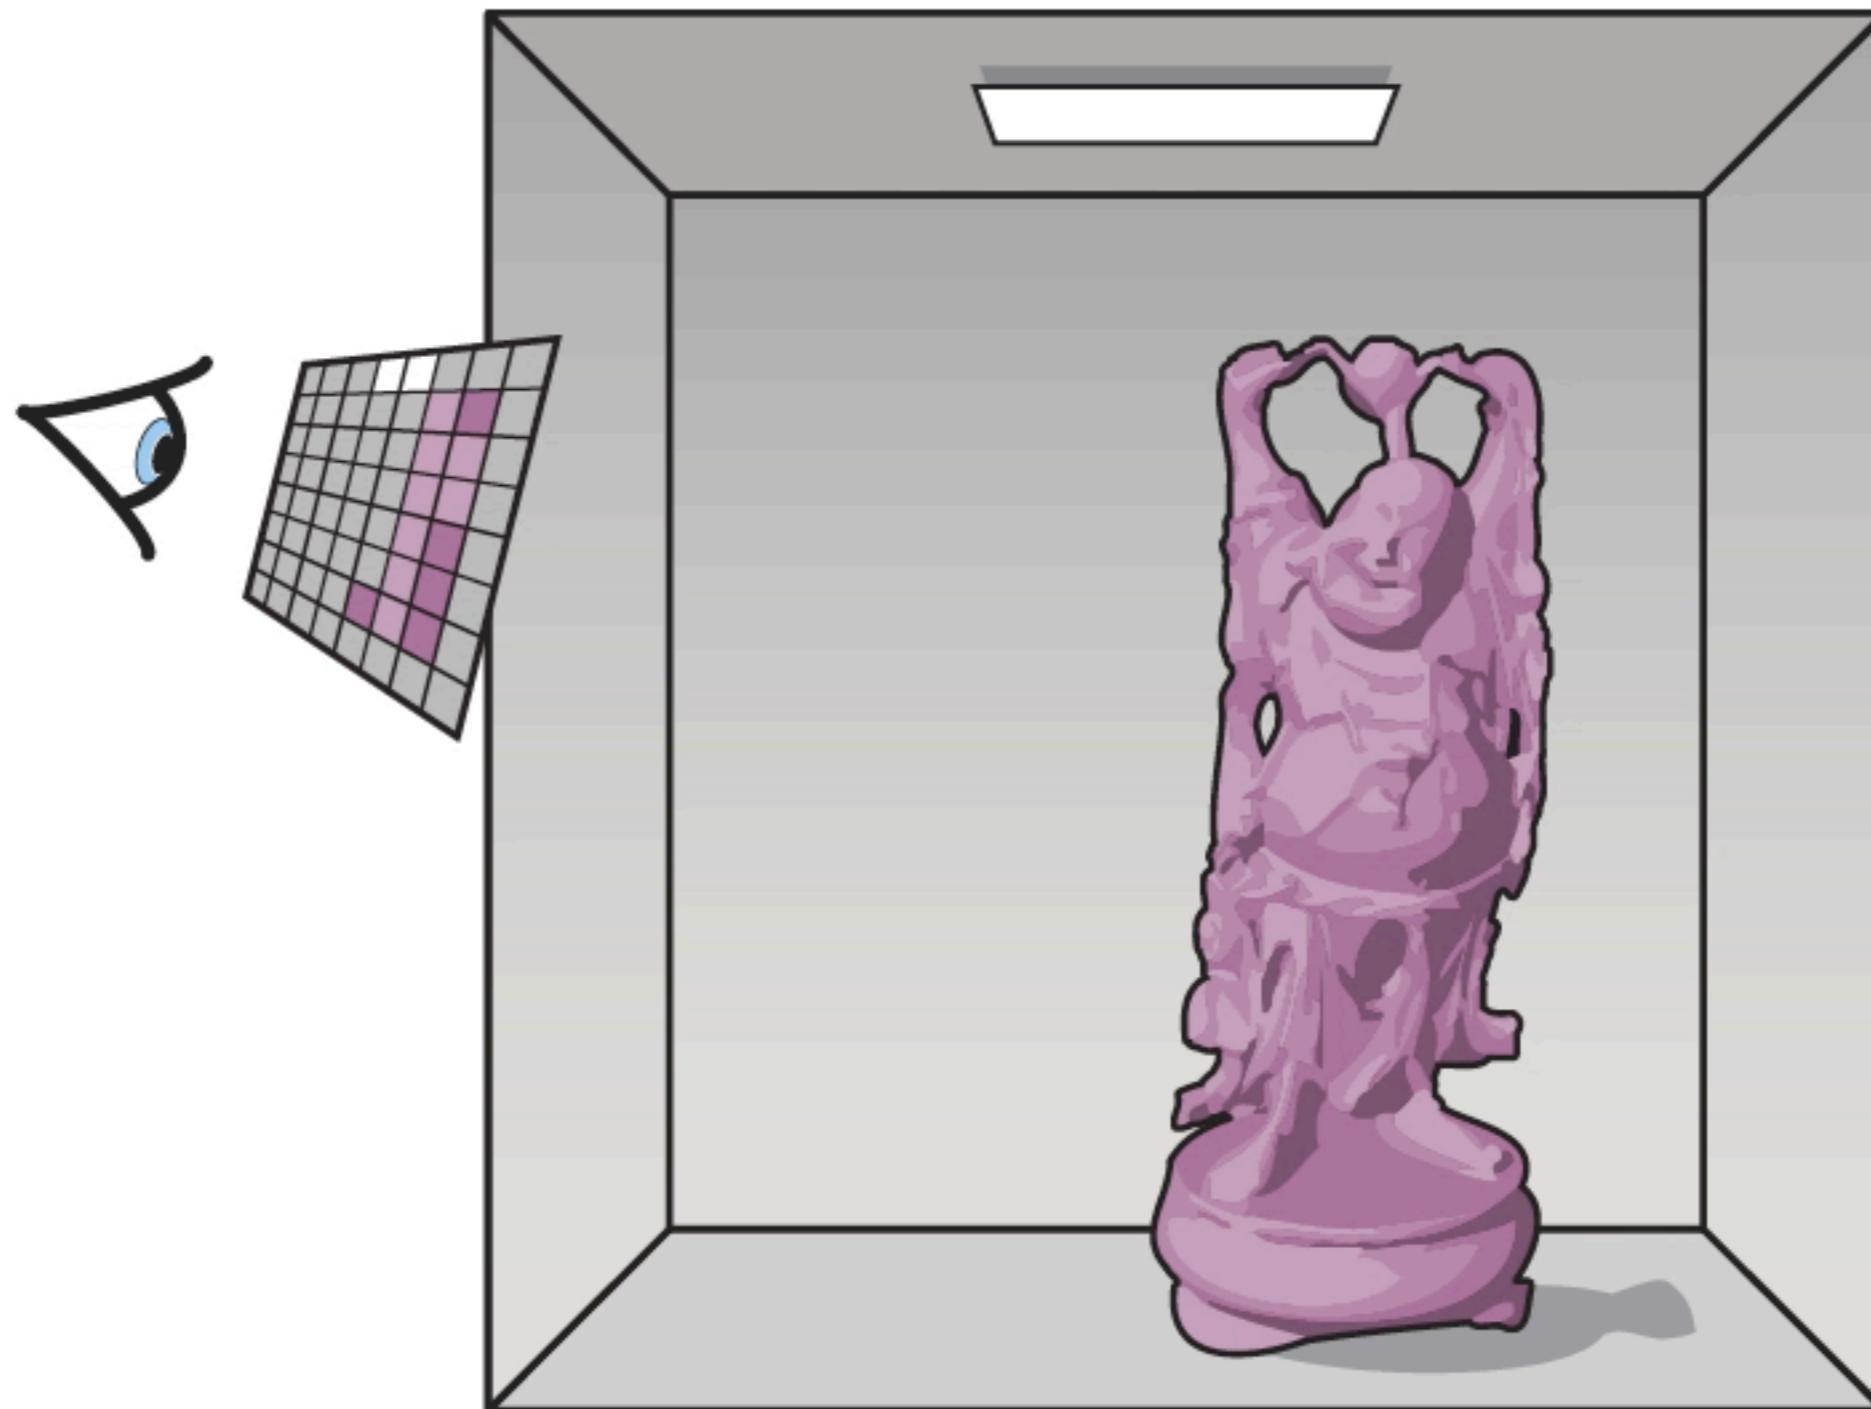
$$T(\bar{\mathbf{x}}) = G(\mathbf{x}_0, \mathbf{x}_1) \prod_{j=1}^{k-1} f(\mathbf{x}_j, \mathbf{x}_{j+1}, \mathbf{x}_{j-1}) G(\mathbf{x}_j, \mathbf{x}_{j+1})$$



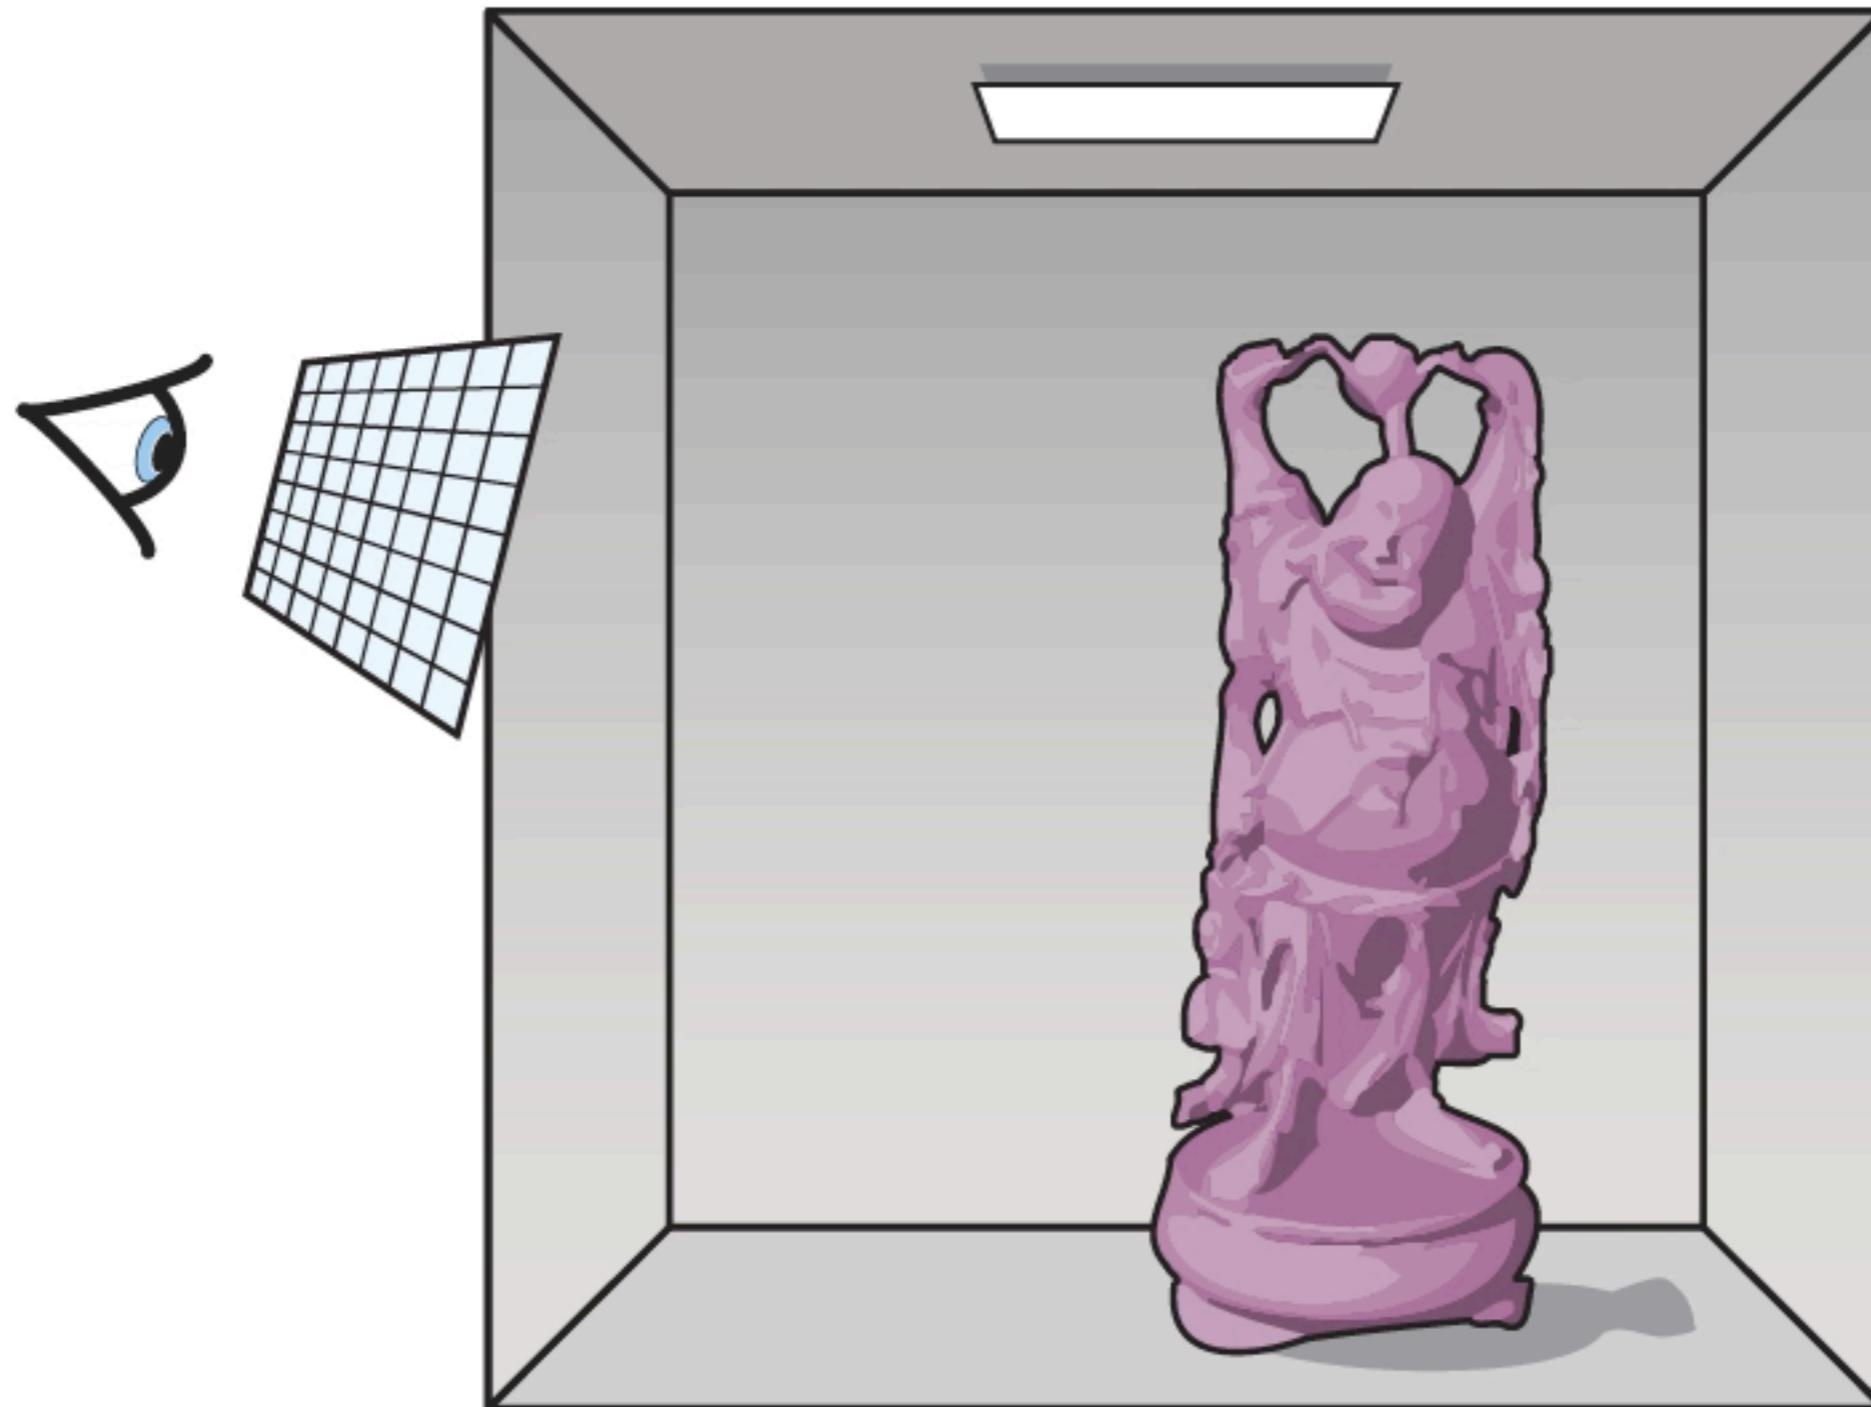
Light Tracing



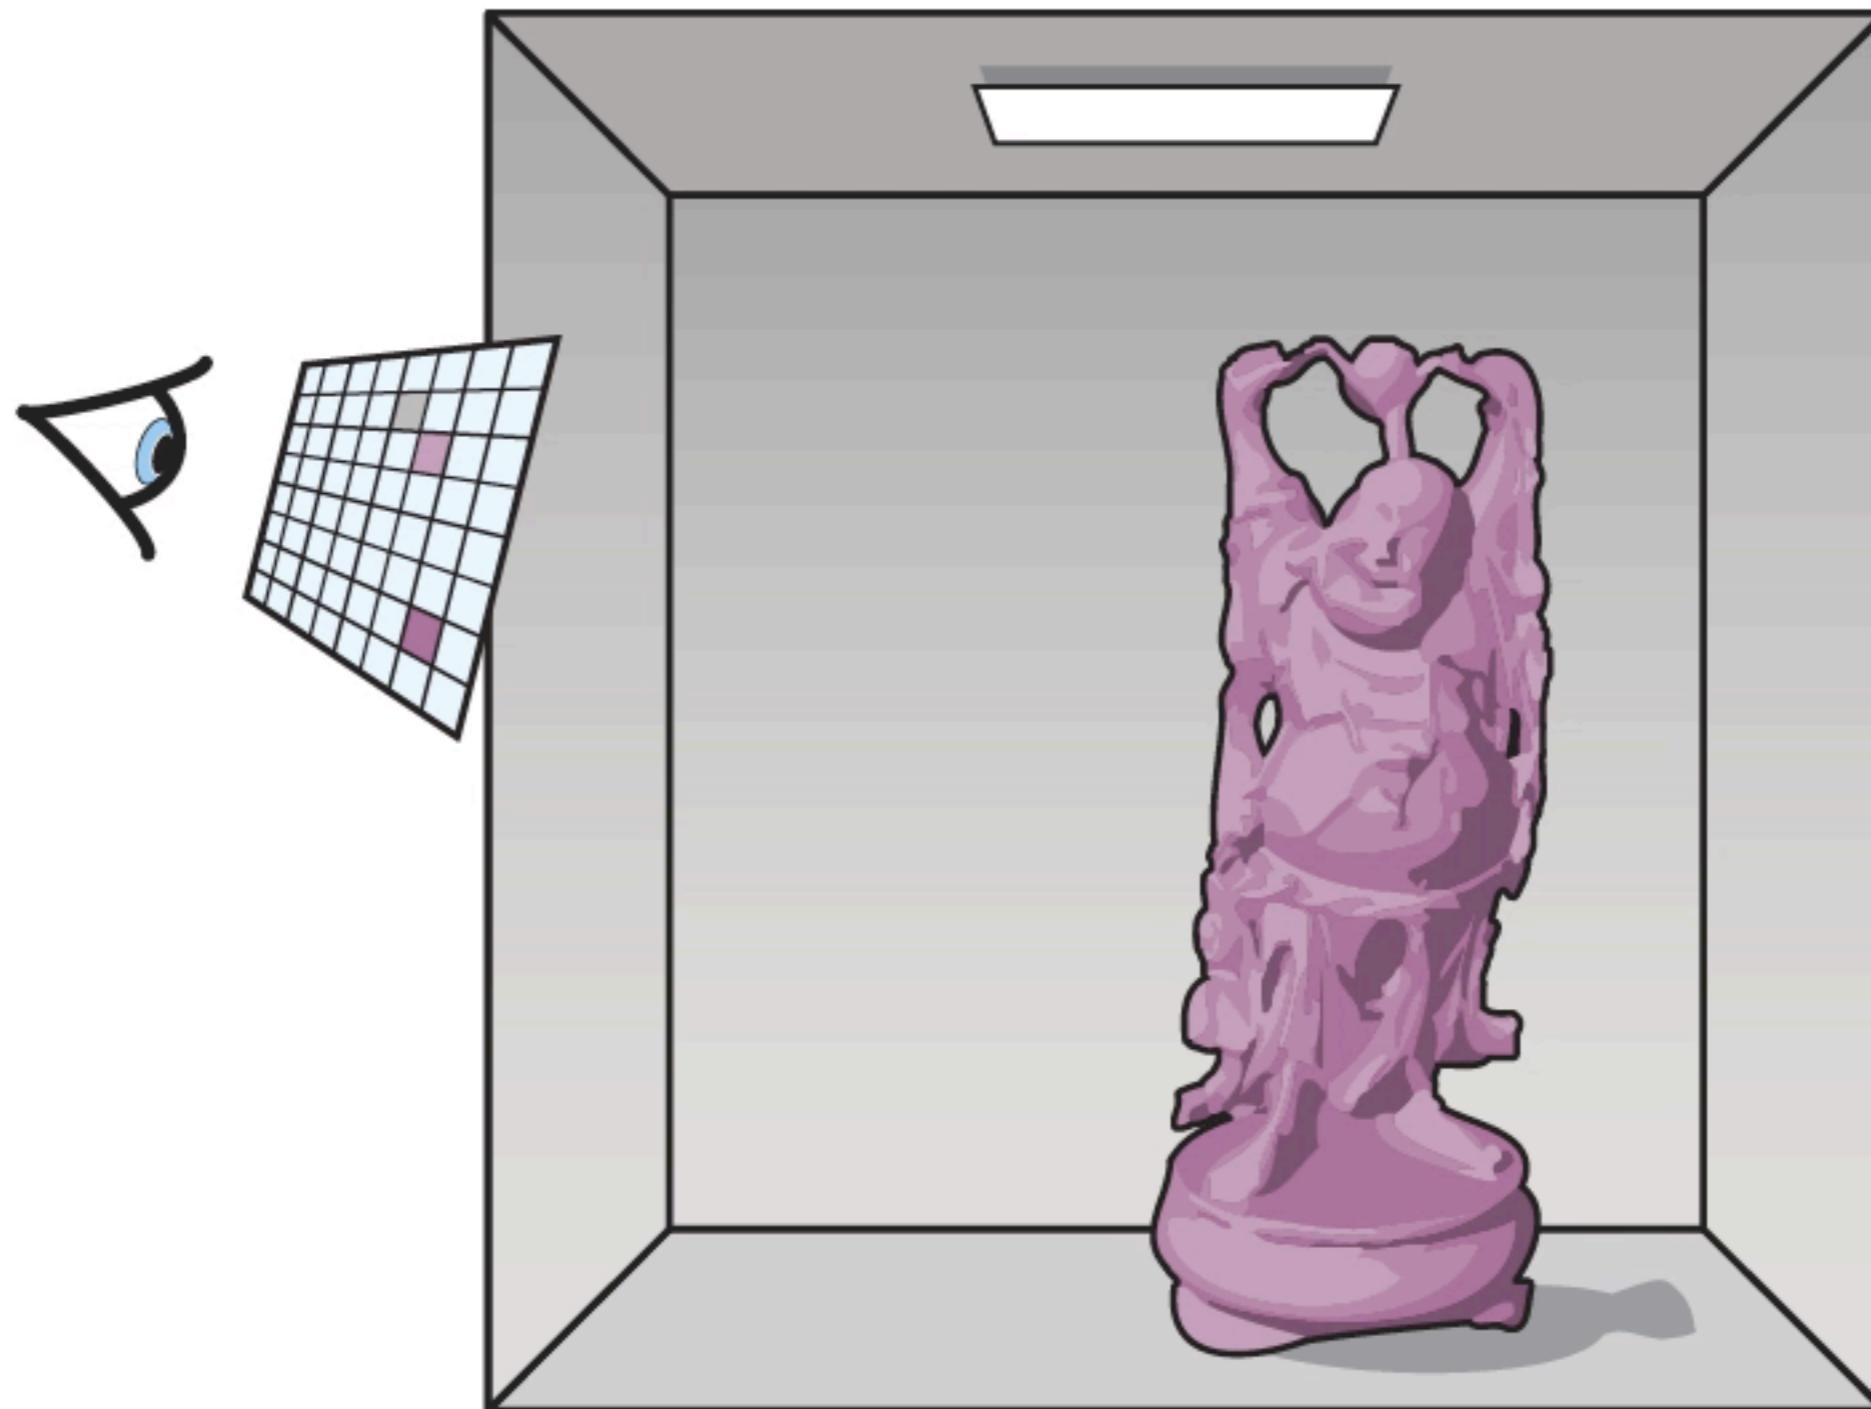
Light Tracing



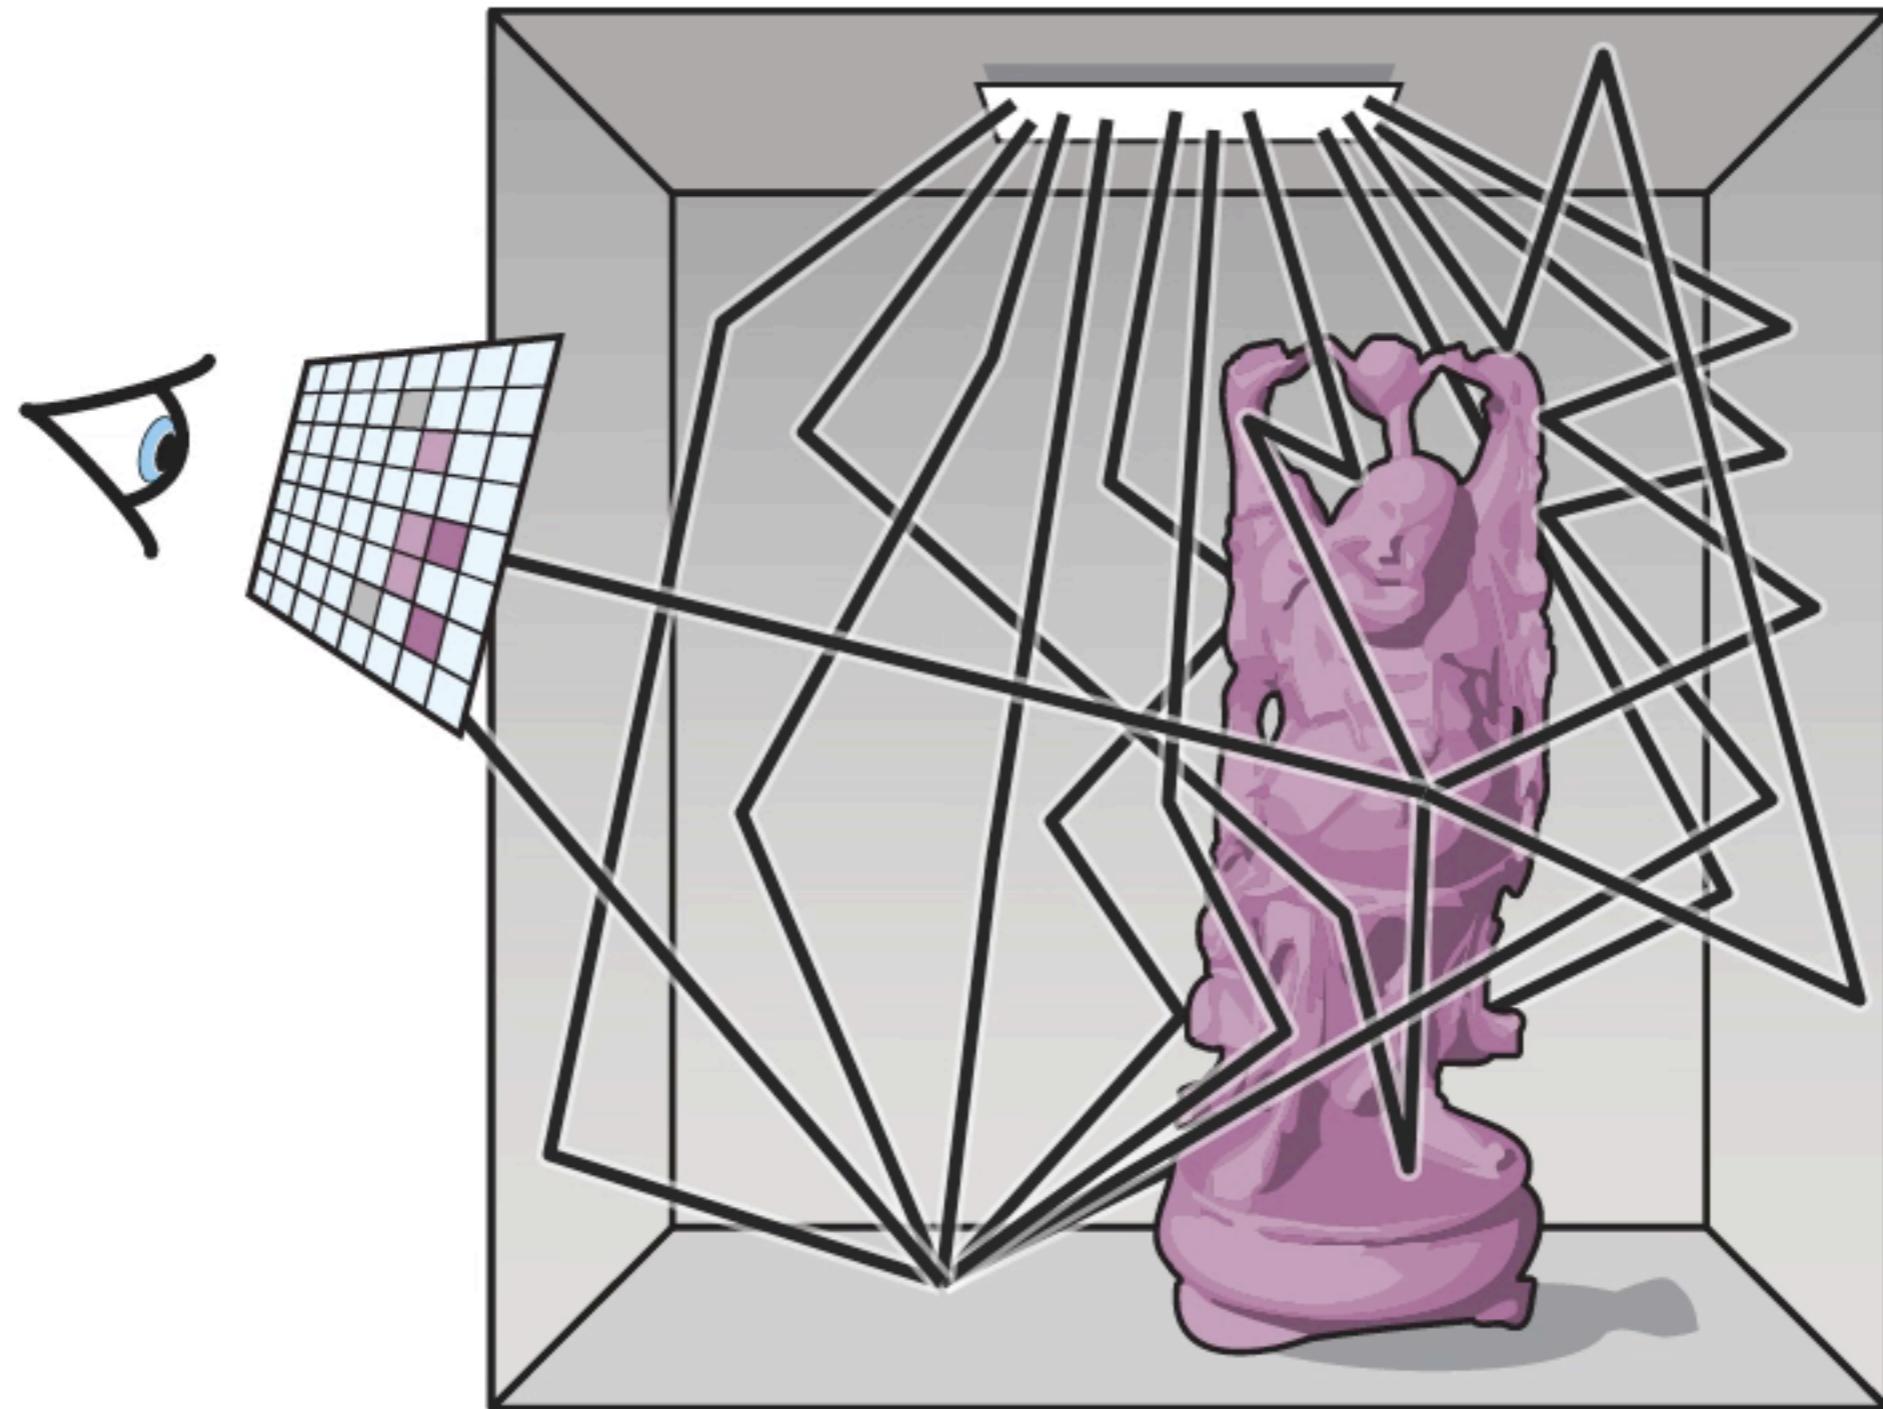
Path Tracing



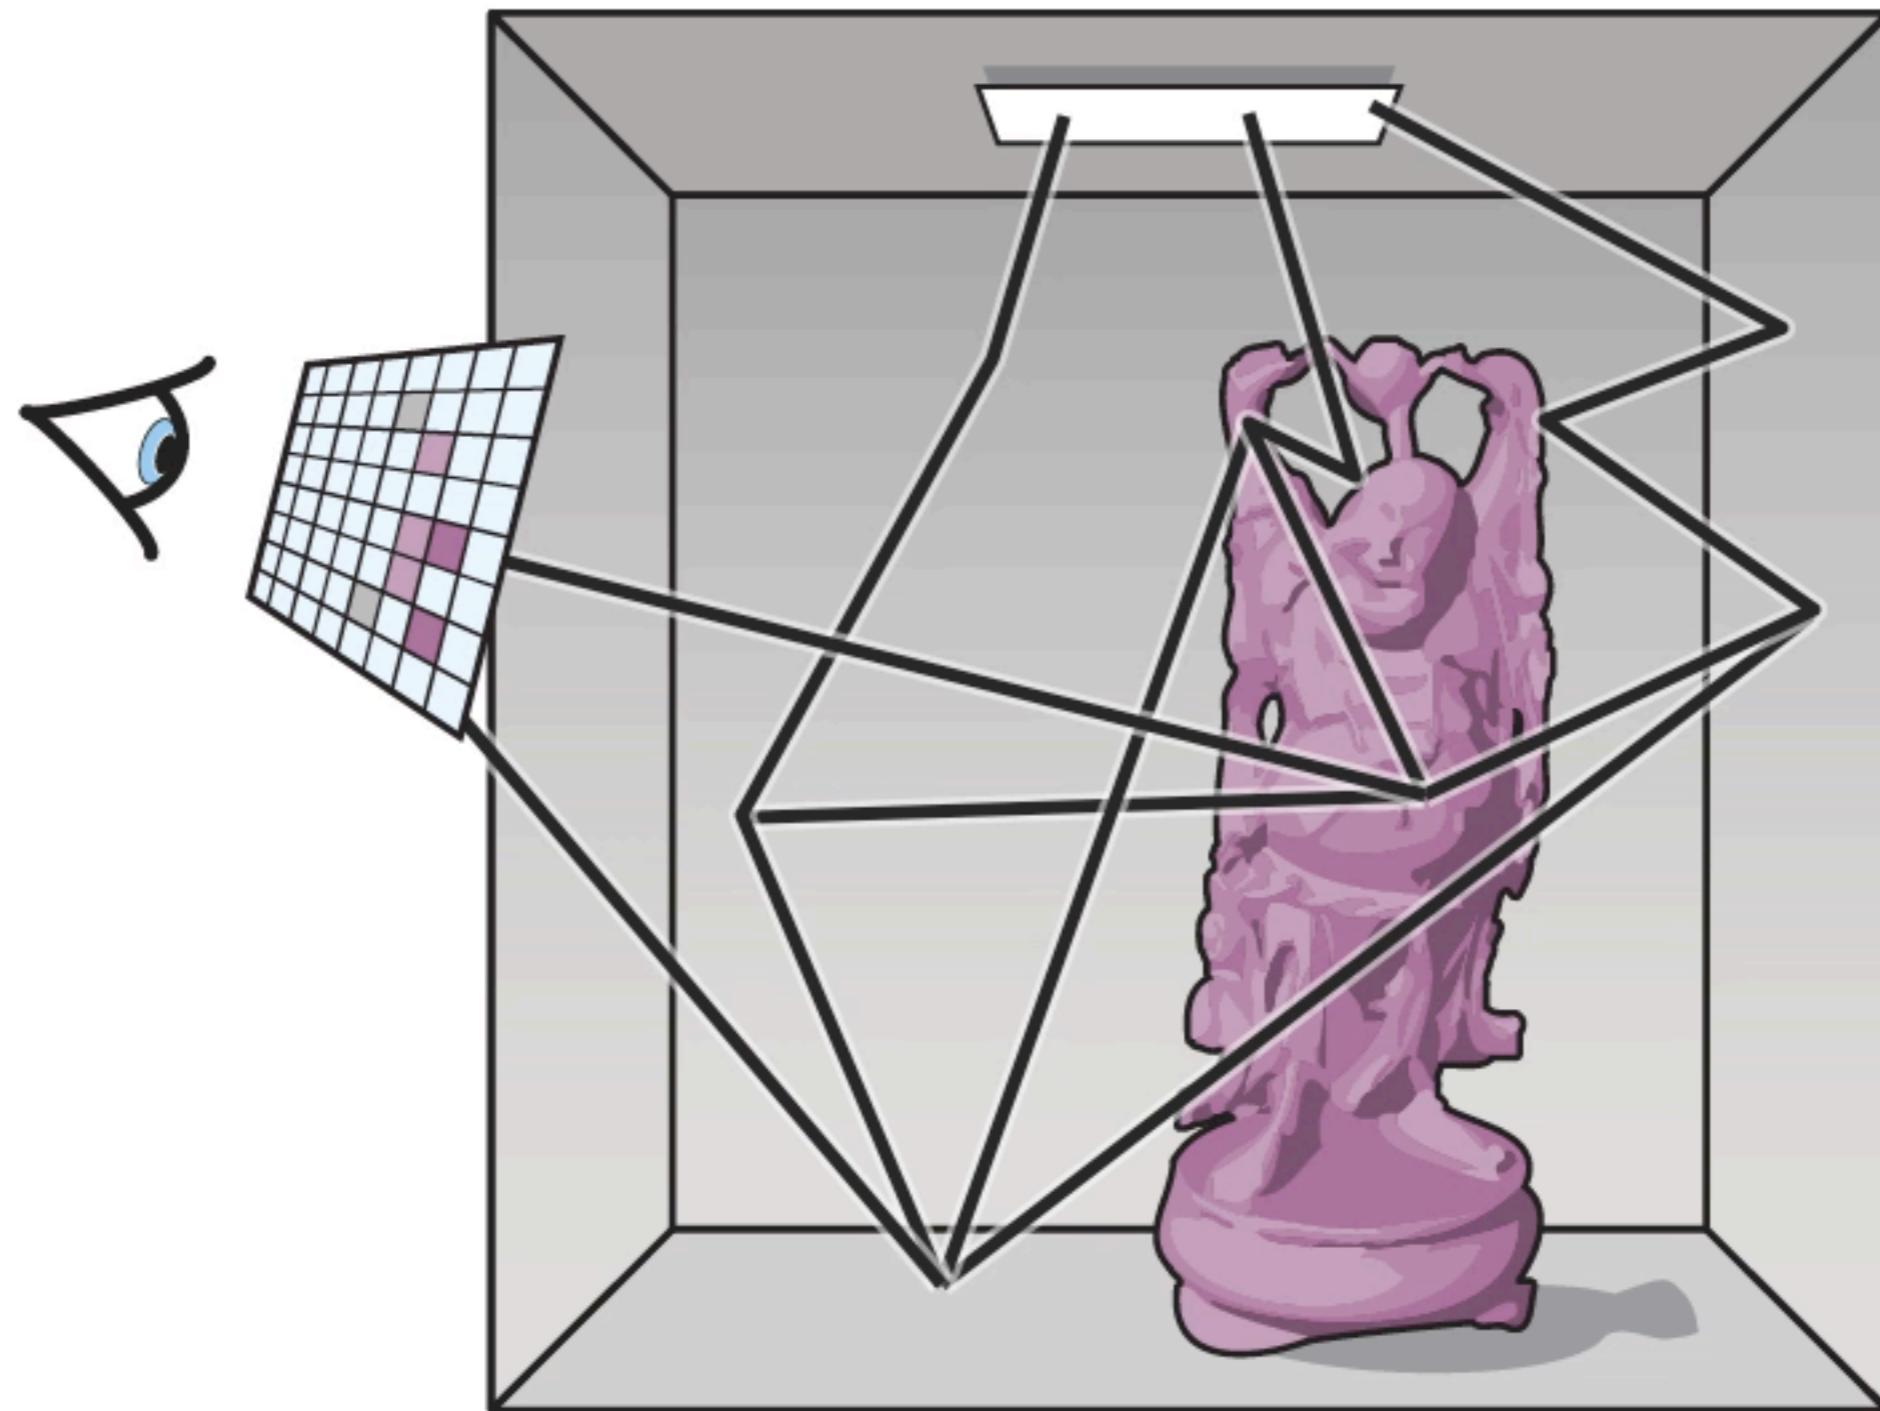
Bidirectional Path Tracing



Many-Light Rendering



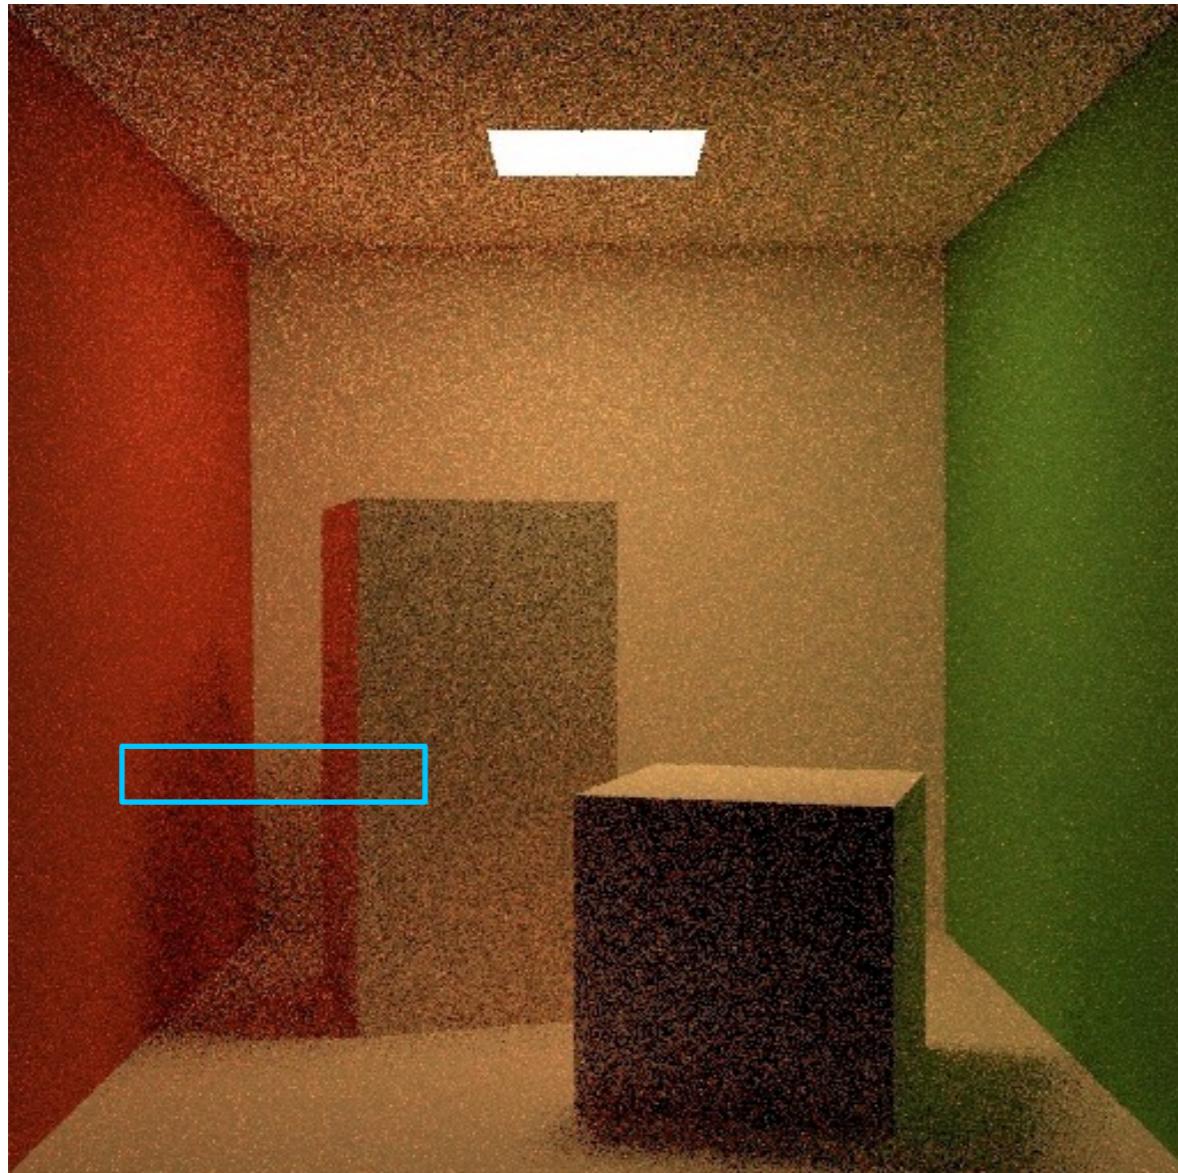
Many-Light Rendering



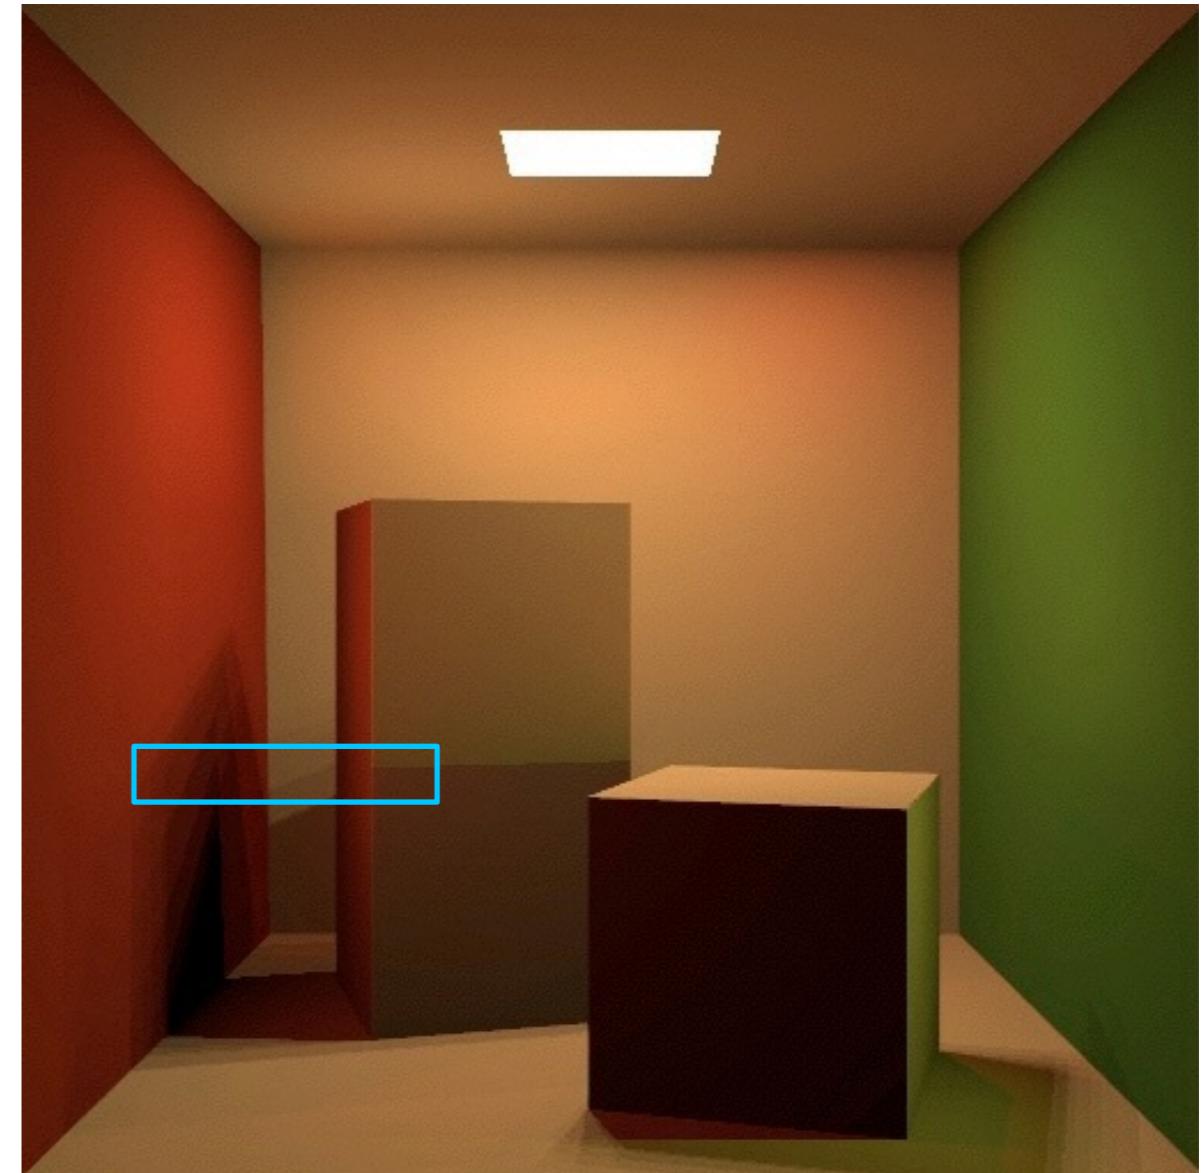
Many-Light (VPL) Rendering

- Main idea: convert the *global* illumination problem into a *direct* illumination problem

Path Tracing vs. Many-Lights



4 paths/pixel
(random error shows up as noise)



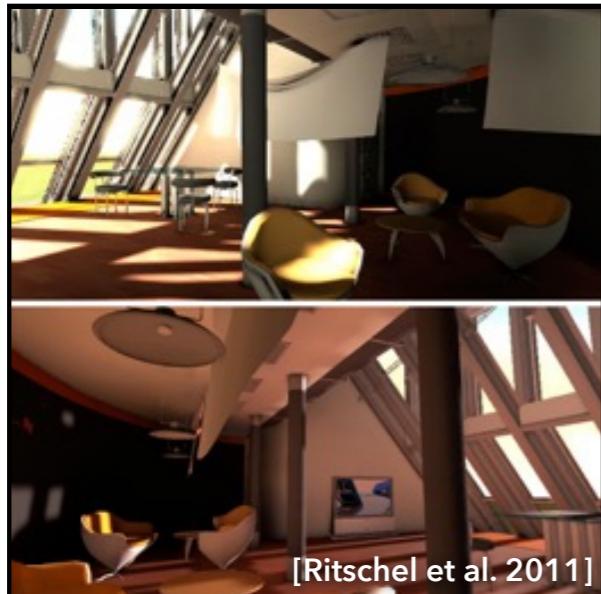
16 VPLs
(correlated error shows up as structured artifacts)

Today's Menu

- Many-Light Rendering
- Generation of Virtual Point Lights (VPLs)
- Lighting with VPLs
- Scalability

Many-Light Rendering

- Original idea presented as *Instant Radiosity* [Keller 1997]
- Now more than 30 extensions



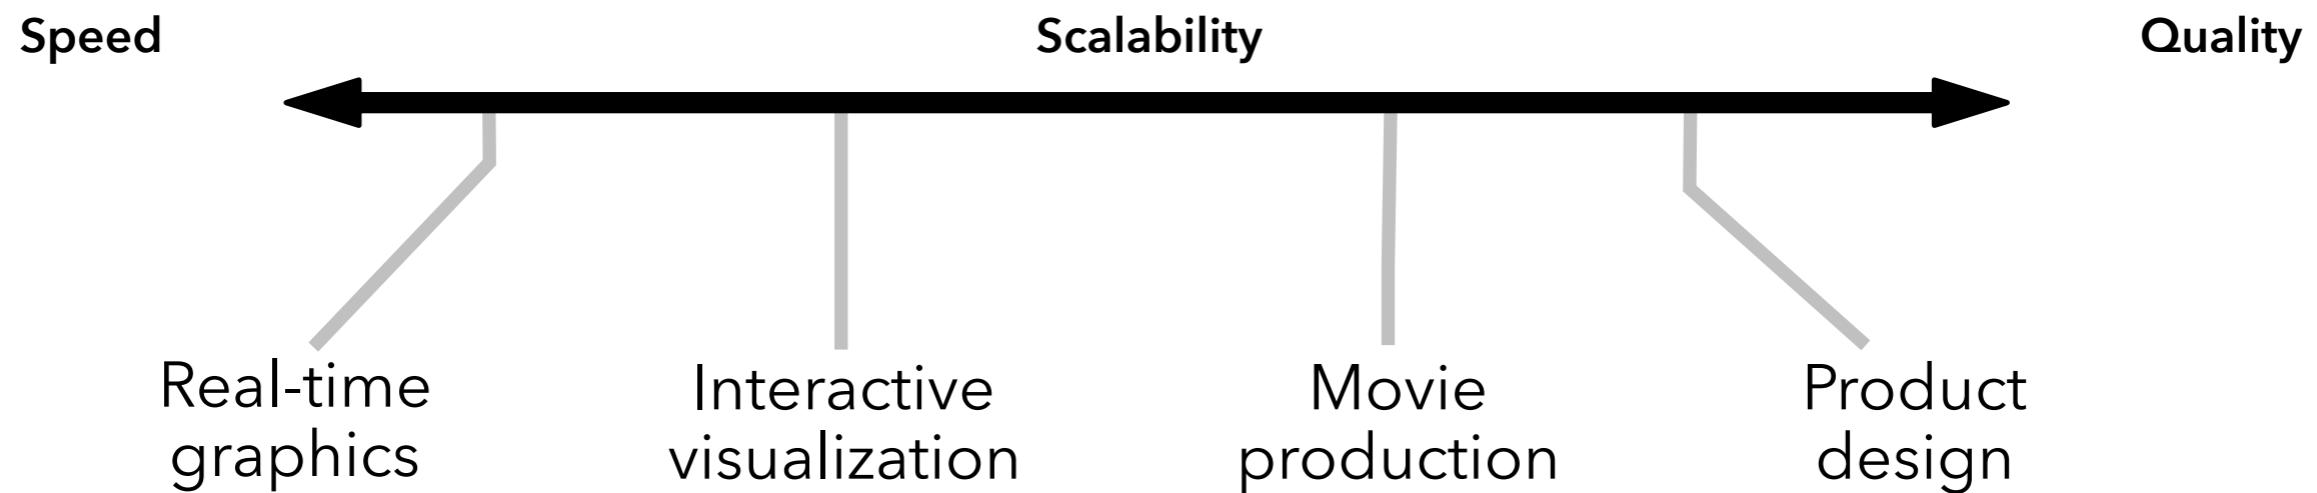
[Ritschel et al. 2011]



[Engelhardt et al. 2011]

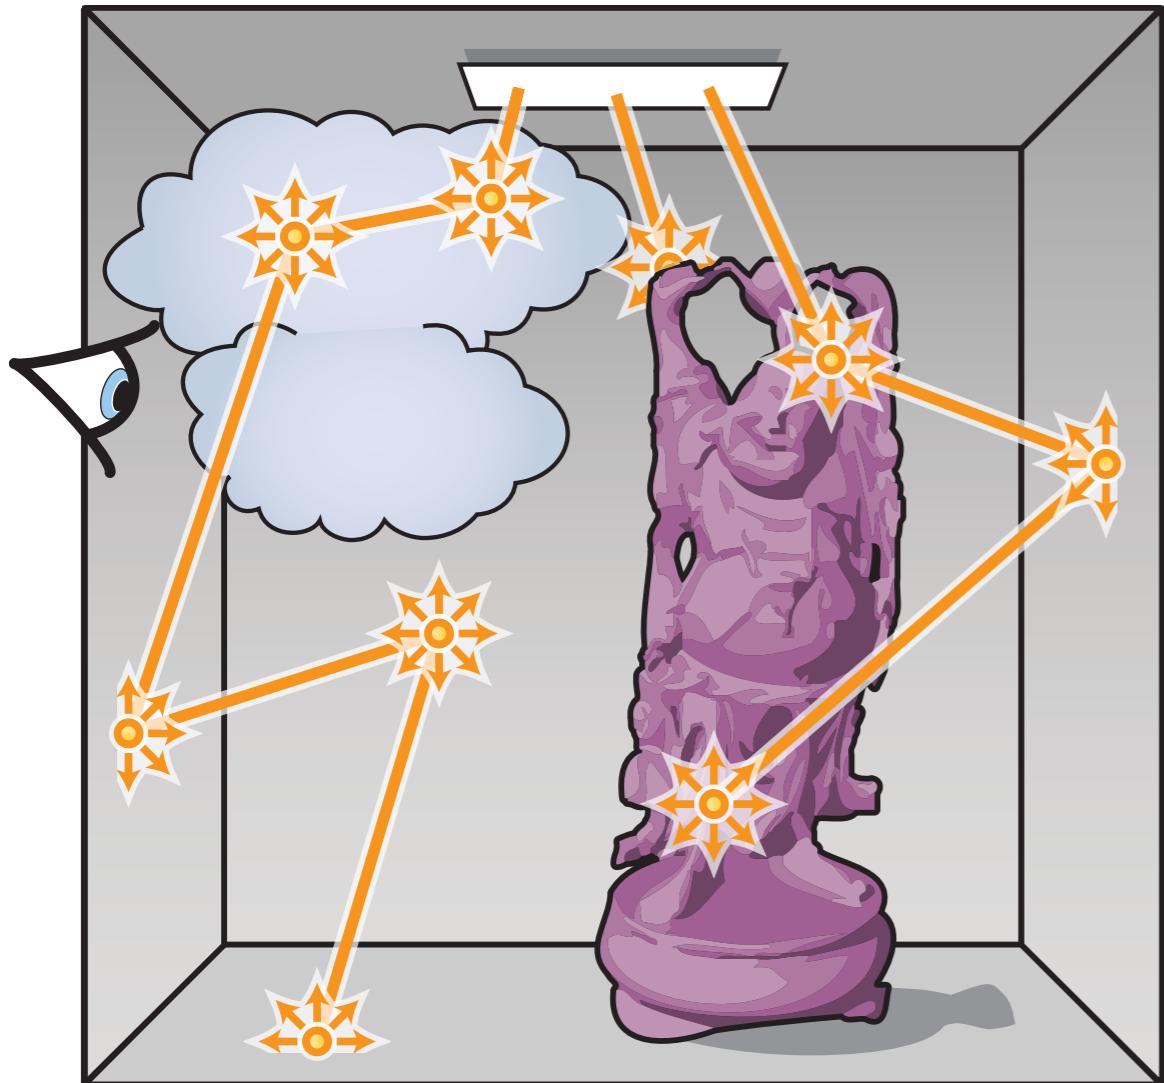


[Walter et al. 2012]

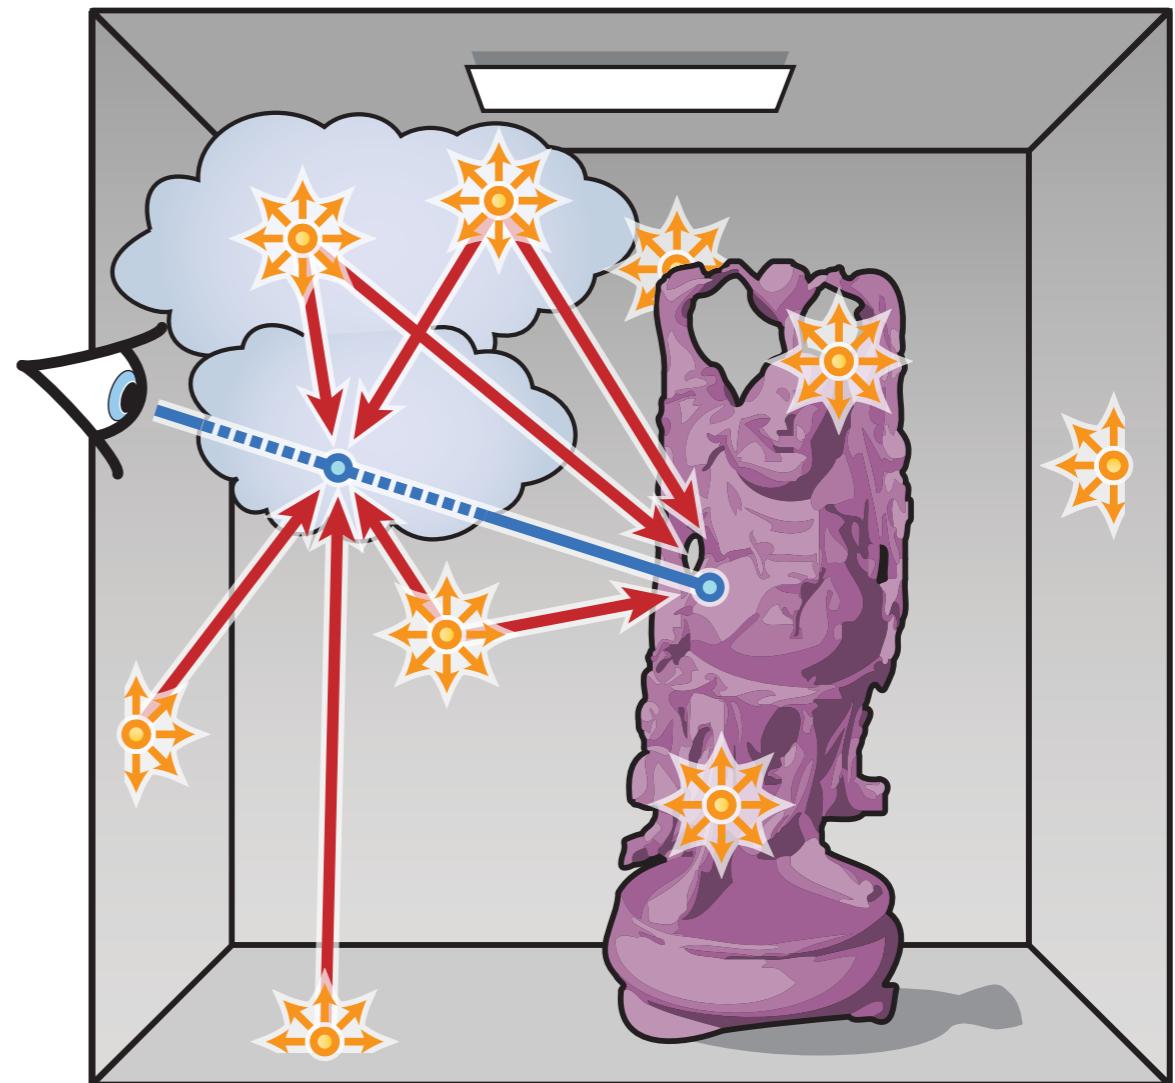


Many-Light Rendering

Generation of VPLs



Lighting with VPLs

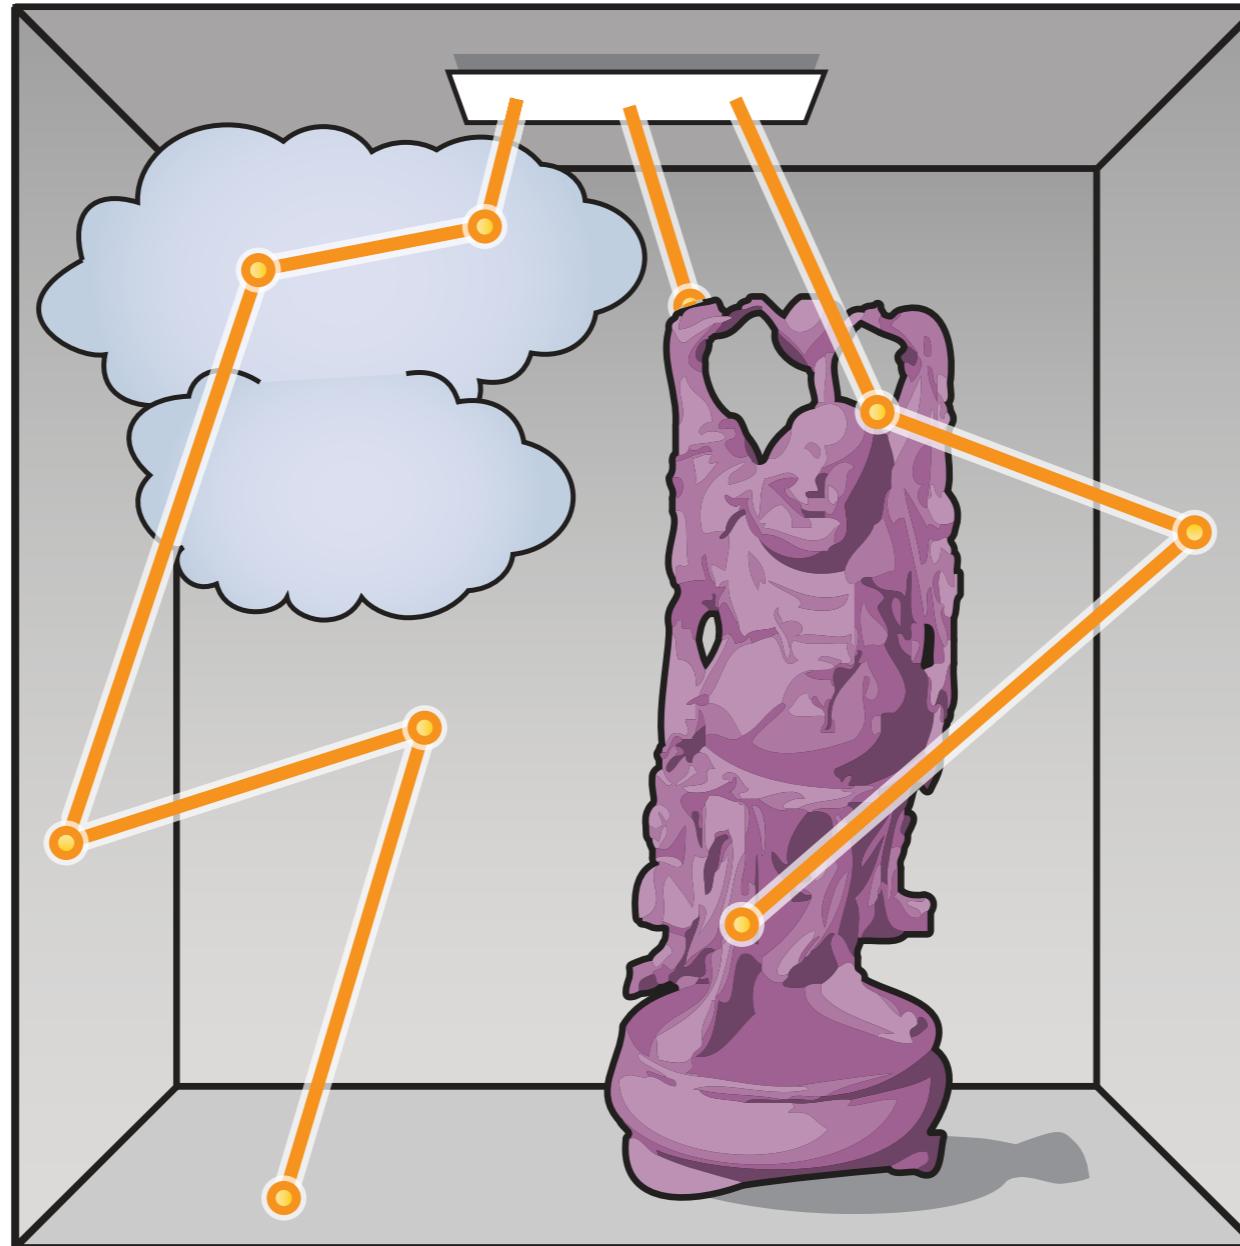


VPL = Virtual Point Light

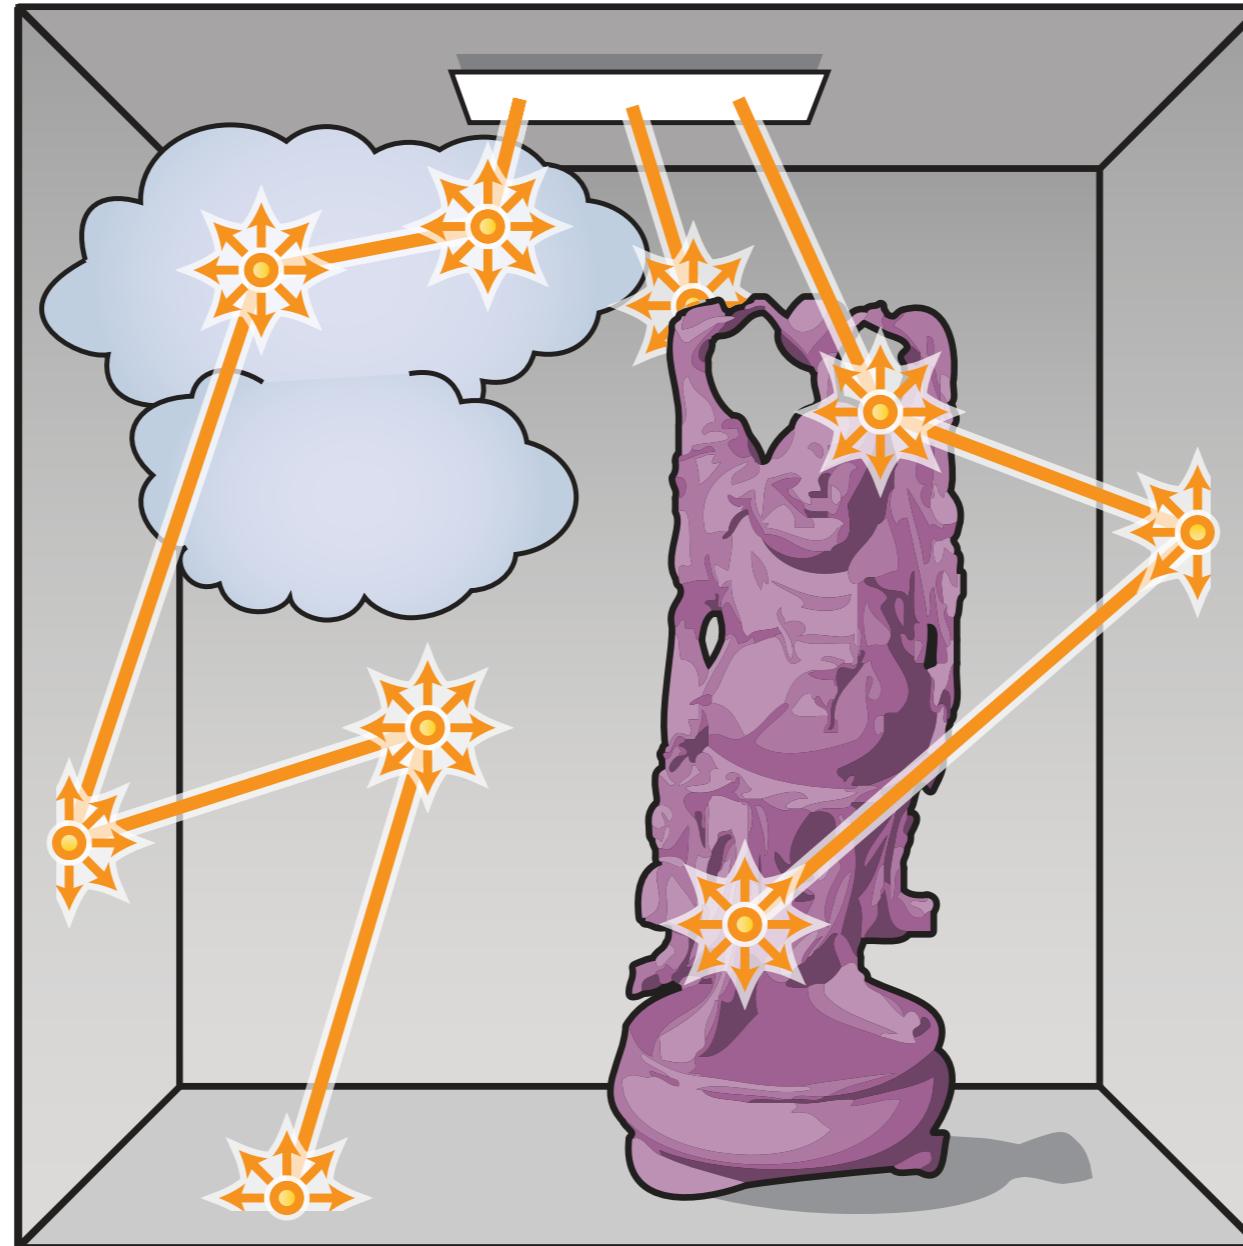
Basic VPL Generation



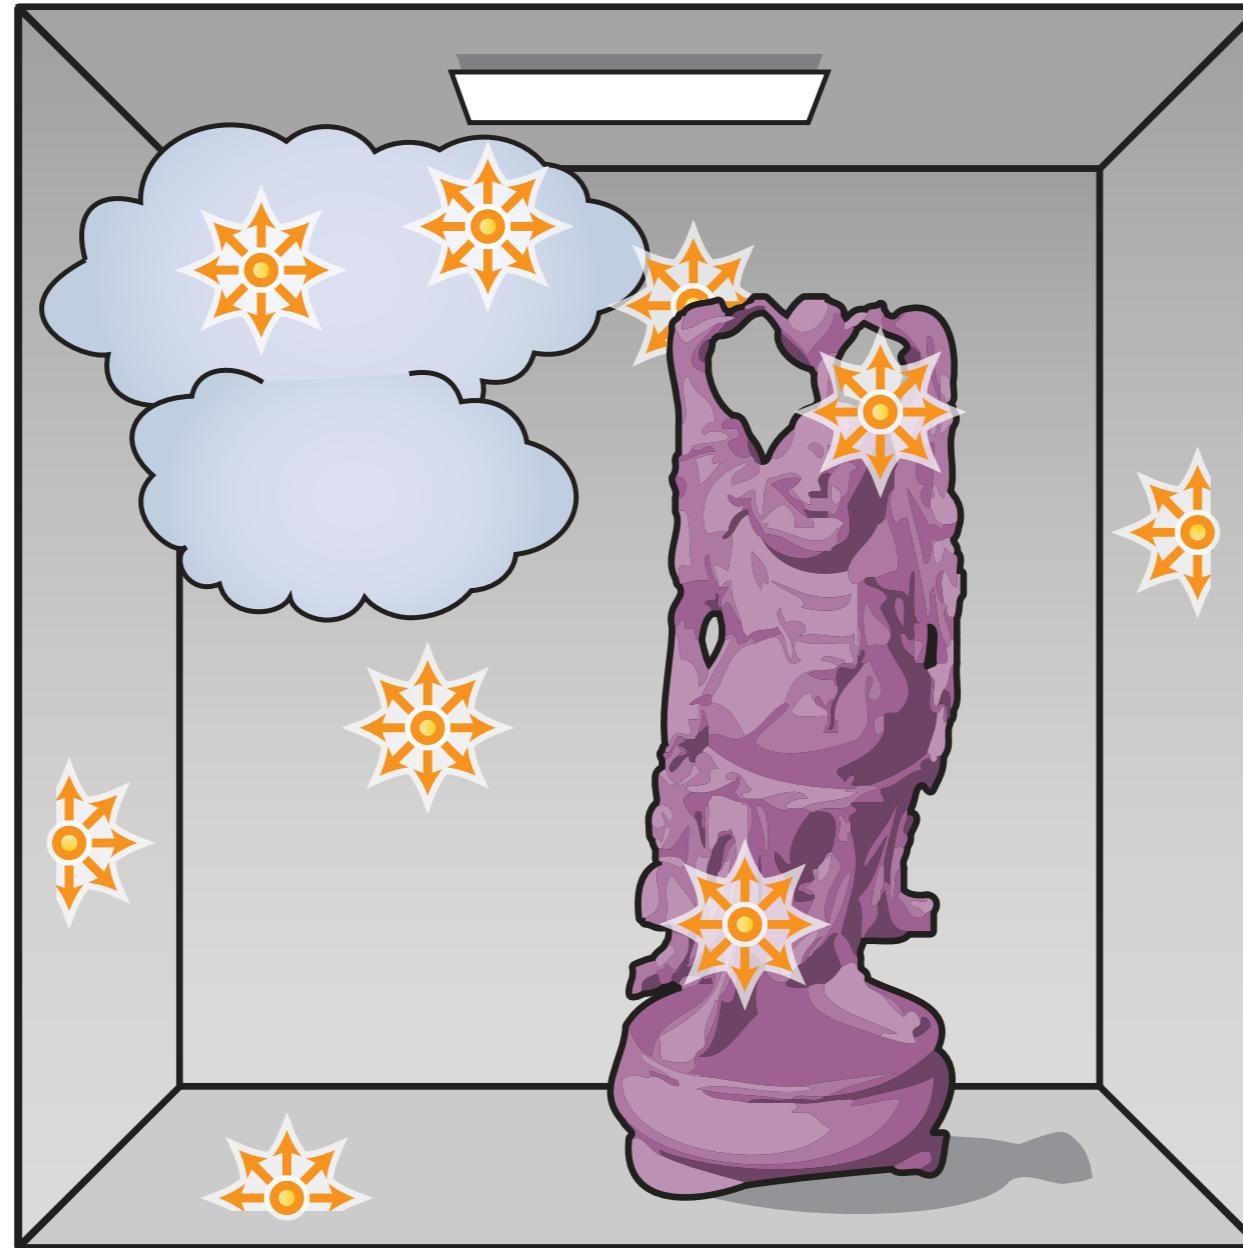
Basic VPL Generation



Basic VPL Generation

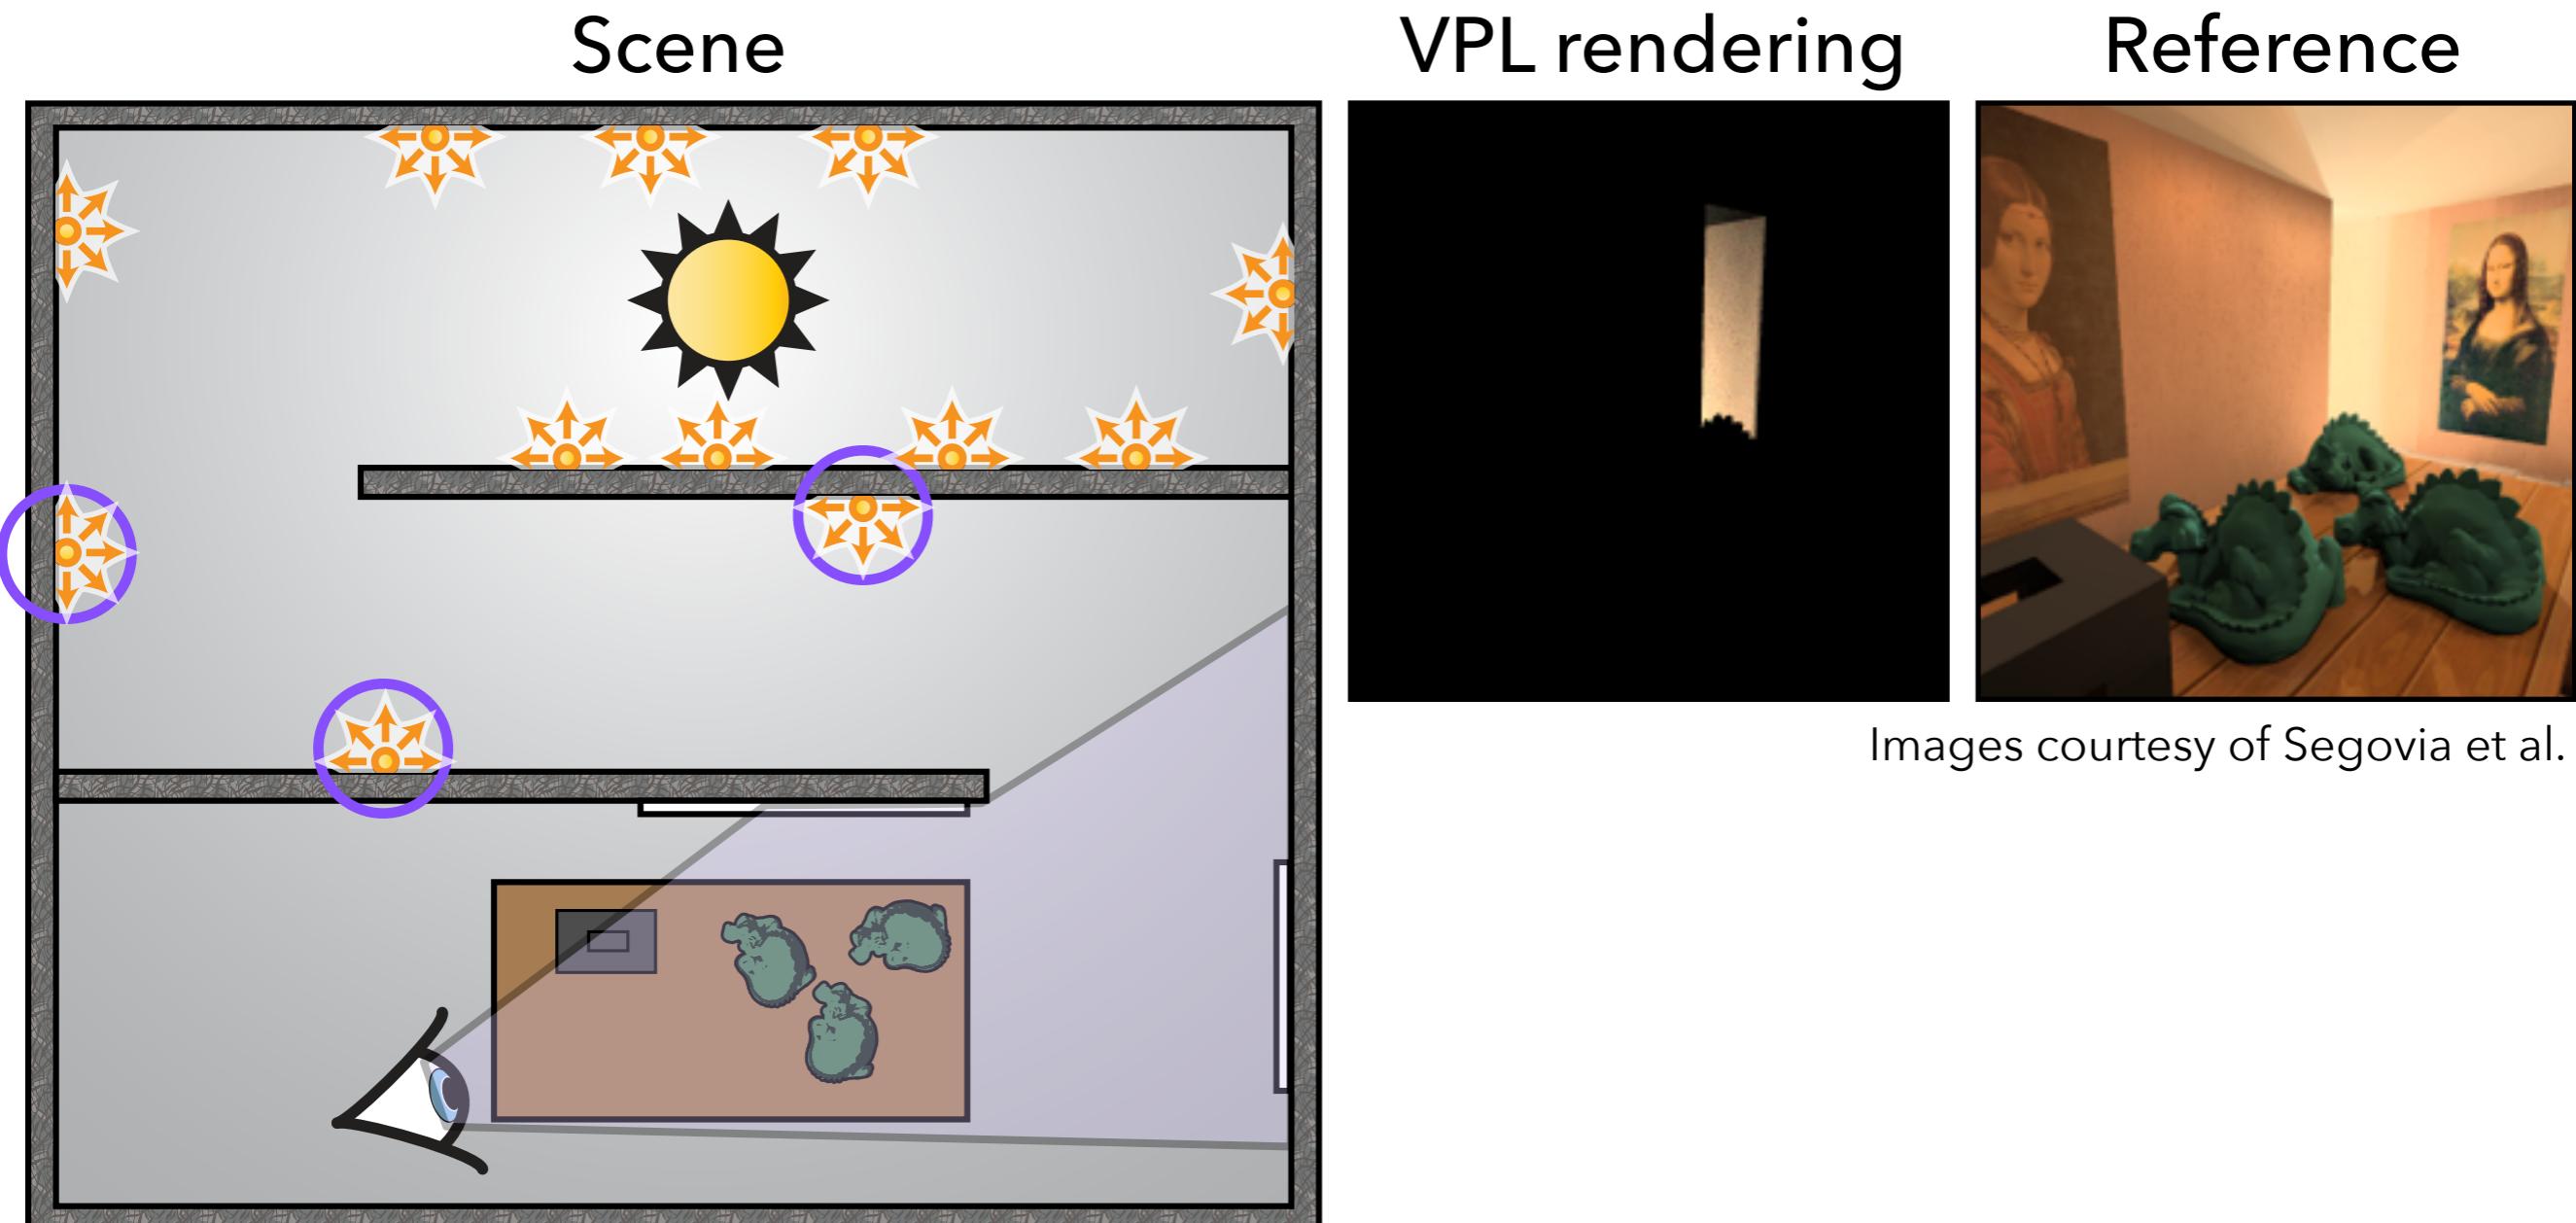


Basic VPL Generation



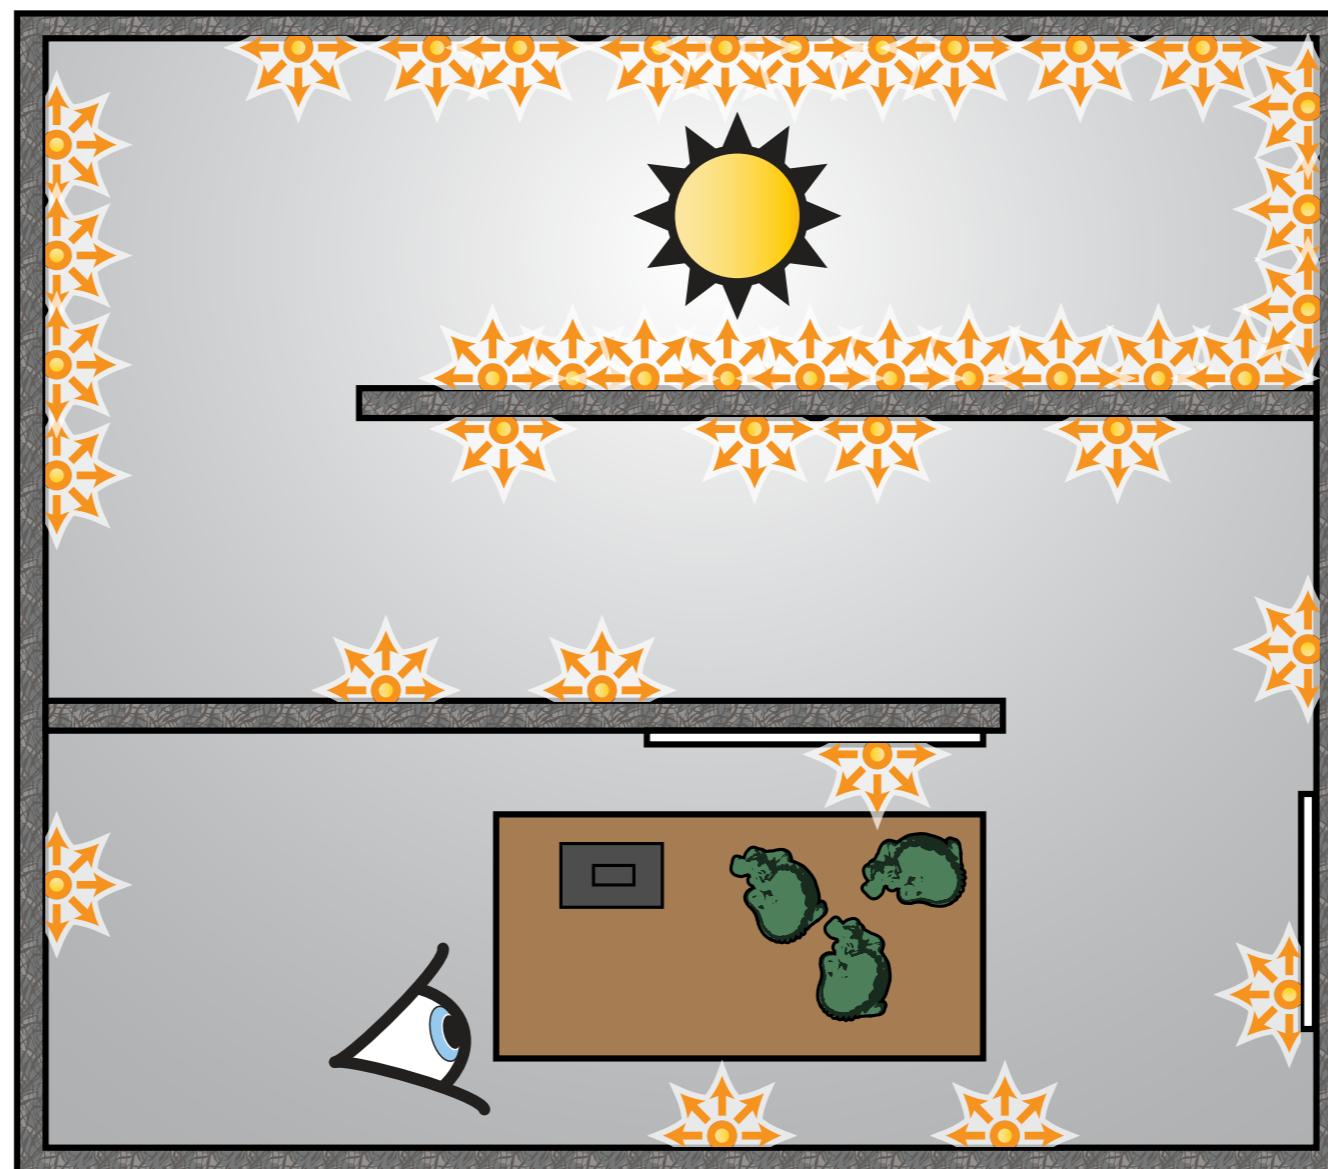
Basic Generation of VPLs

- In complex scenes, many VPLs do not contribute



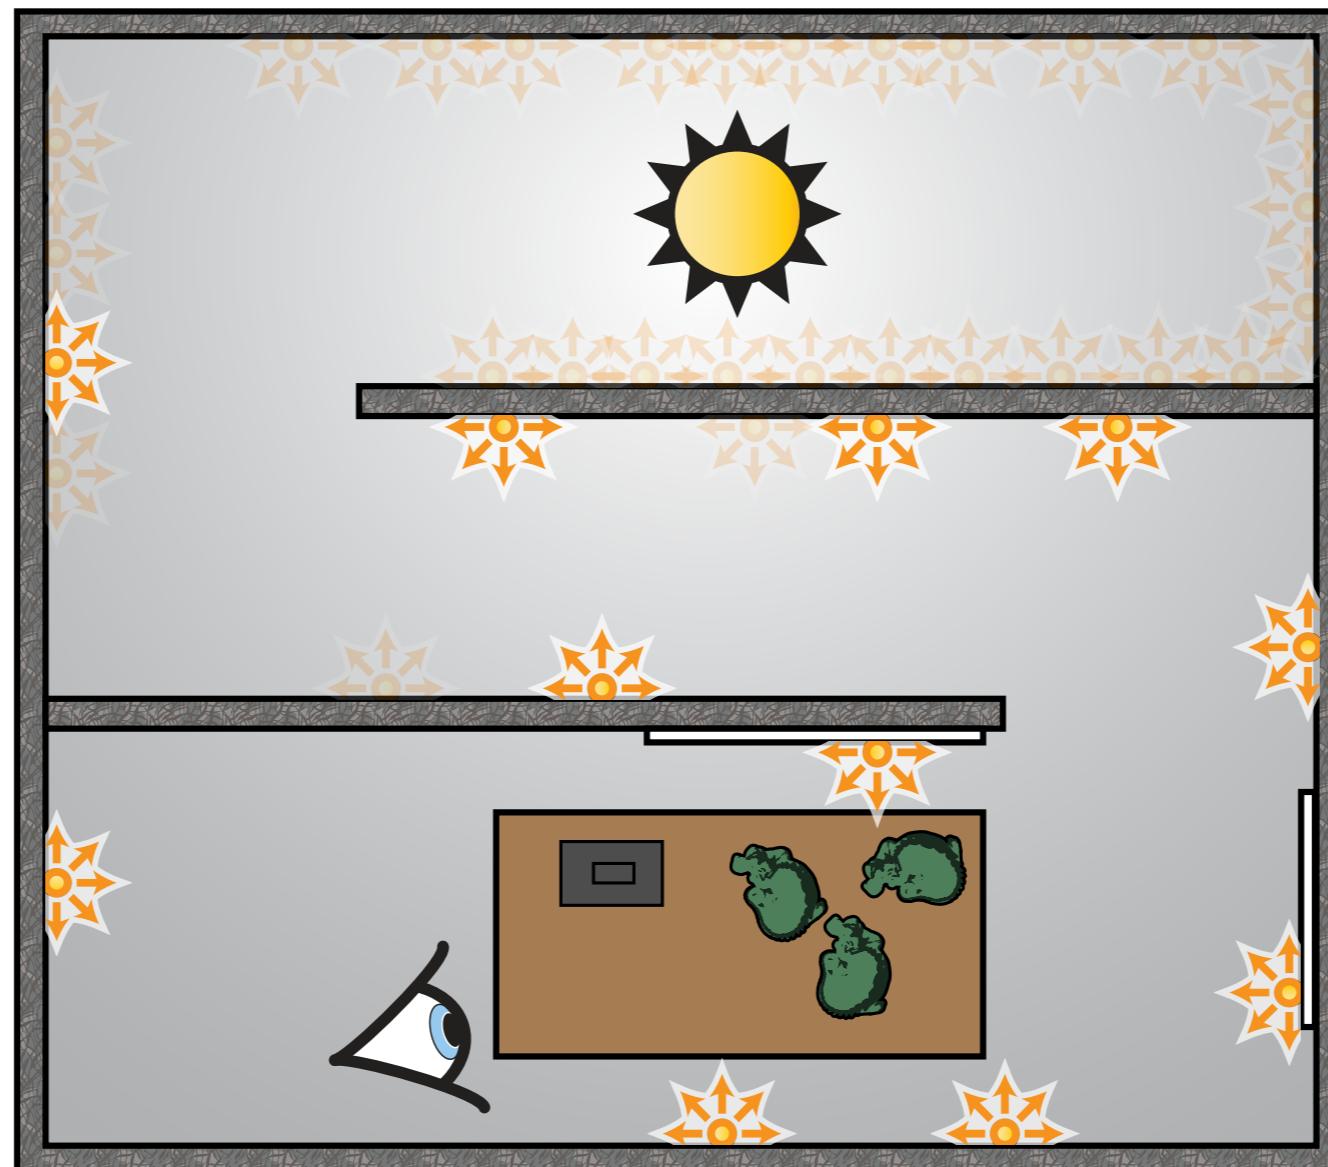
Improved Generation of VPLs

- Rejection of unimportant VPLs [Georgiev and Slusallek 2010]
 - generate many, reject those with low expected contribution



Improved Generation of VPLs

- Rejection of unimportant VPLs [Georgiev and Slusallek 2010]
 - generate many, reject those with low expected contribution



Improved Generation of VPLs

- Rejection of unimportant VPLs [Georgiev and Slusallek 2010]
 - generate many, reject those with low expected contribution
- Step 1: estimate average contribution of a VPL
 - a few pilot VPLs illuminate a few surface points seen by the camera
- Step 2: generate VPLs
 - for each (candidate) VPL:
 - estimate the contribution (to a few shading points)
 - use Russian roulette to either accept or reject

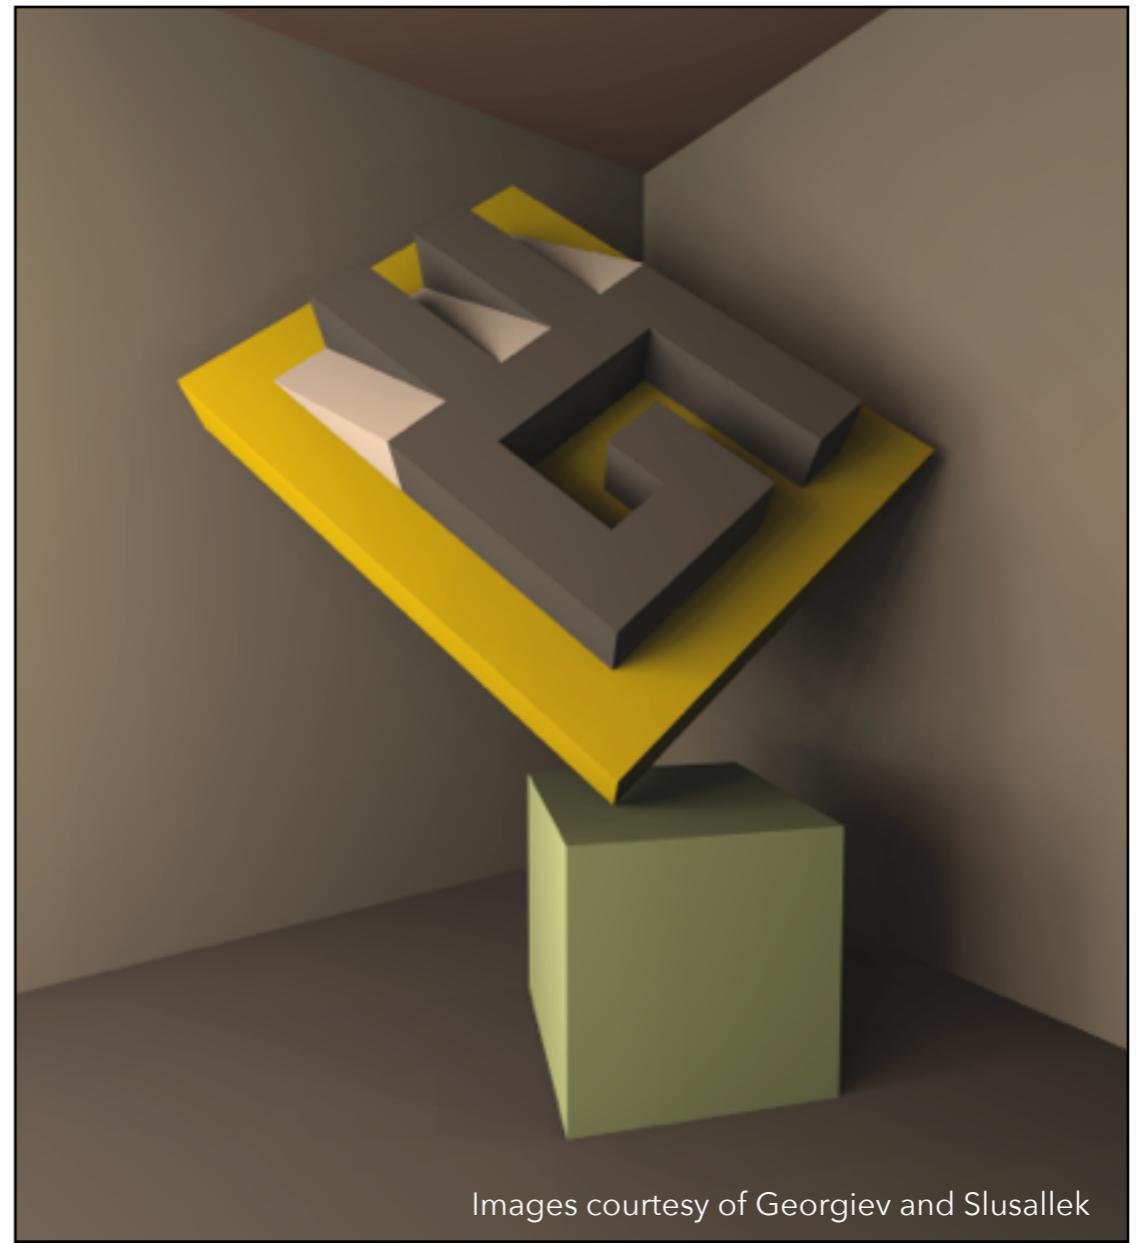
Improved Generation of VPLs

- Rejection of unimportant VPLs [Georgiev and Slusallek 2010]

Without rejection



With rejection (7% acceptance)



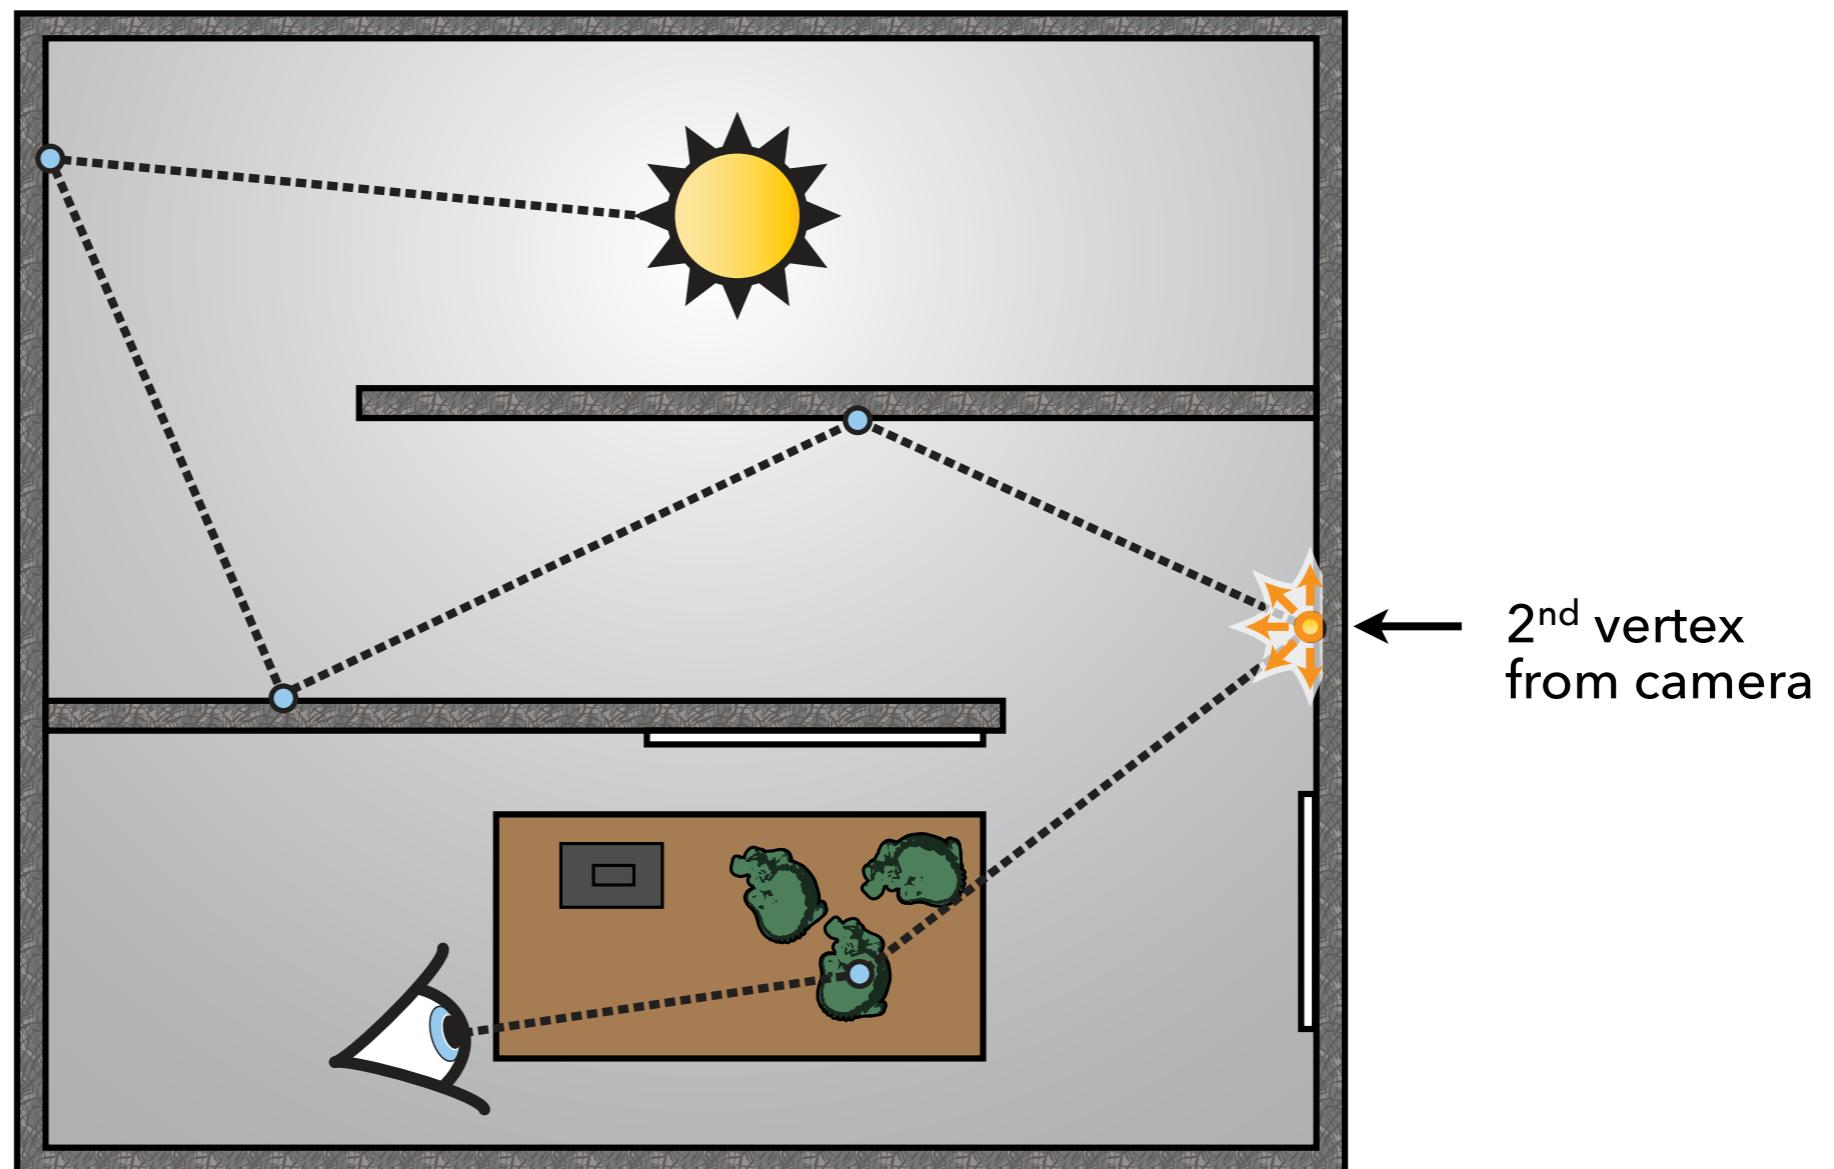
Images courtesy of Georgiev and Slusallek

Improved Generation of VPLs

- Rejection of unimportant VPLs [Georgiev and Slusallek 2010]
 - easy to implement
 - all VPLs contribute roughly equally
(thanks to the Russian roulette)
 - can become costly!

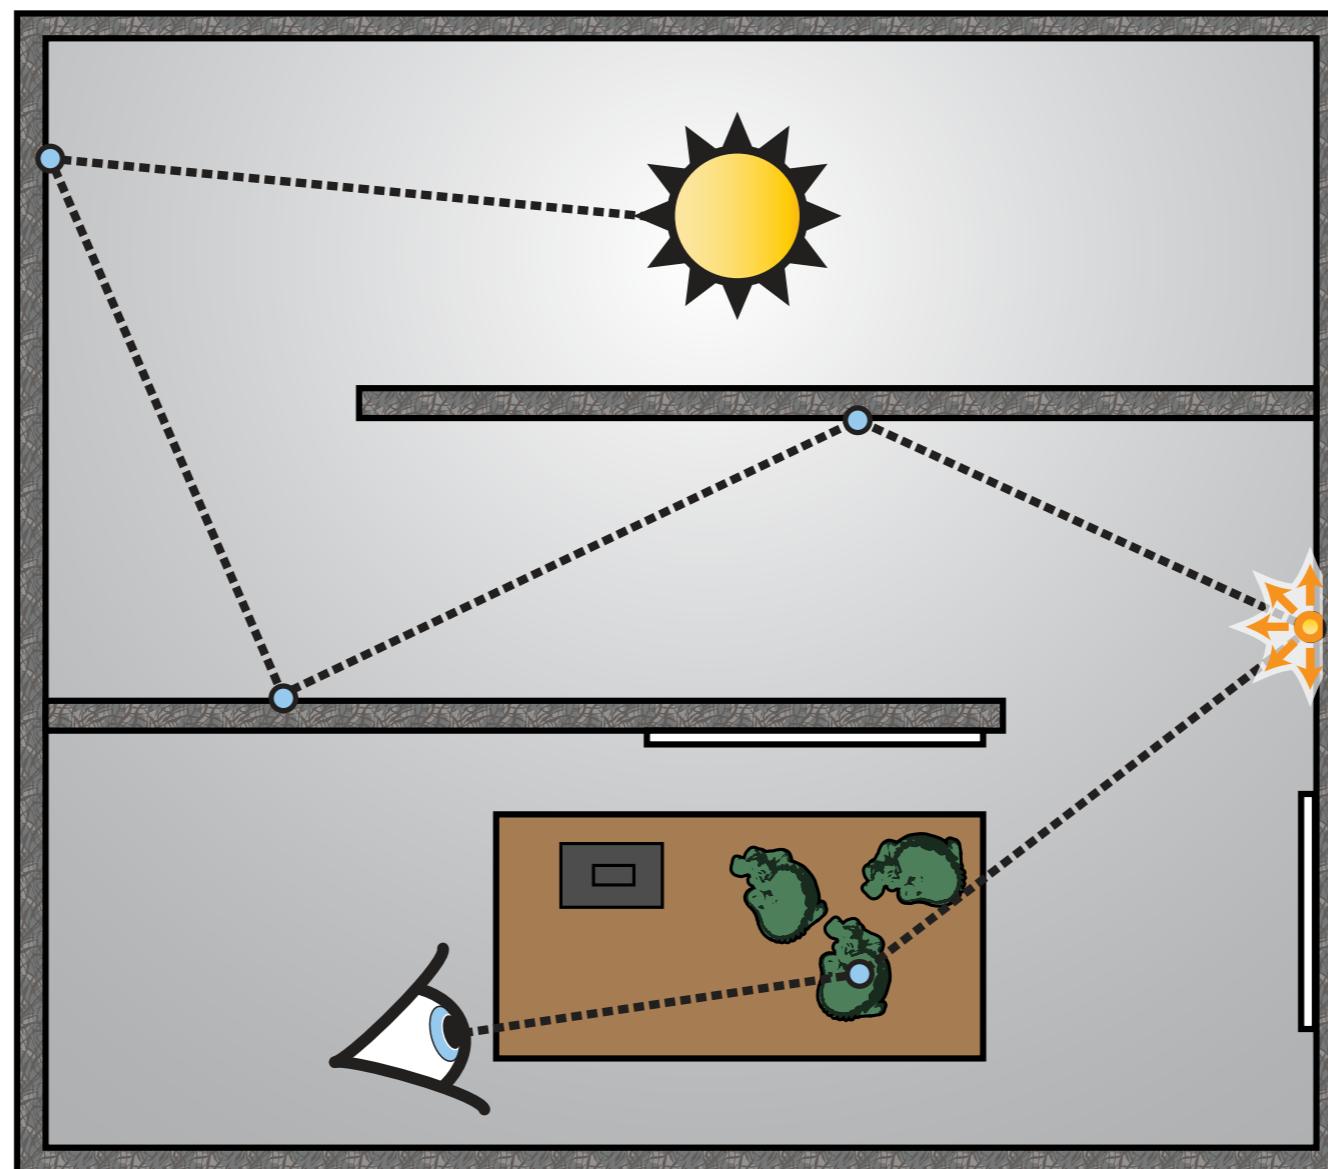
Improved Generation of VPLs

- Bidirectional generation of VPLs [Segovia et al. 2006]
 - construct a path (e.g. using bidirectional path tracing)
 - create a VPL at the 2nd bounce from the camera



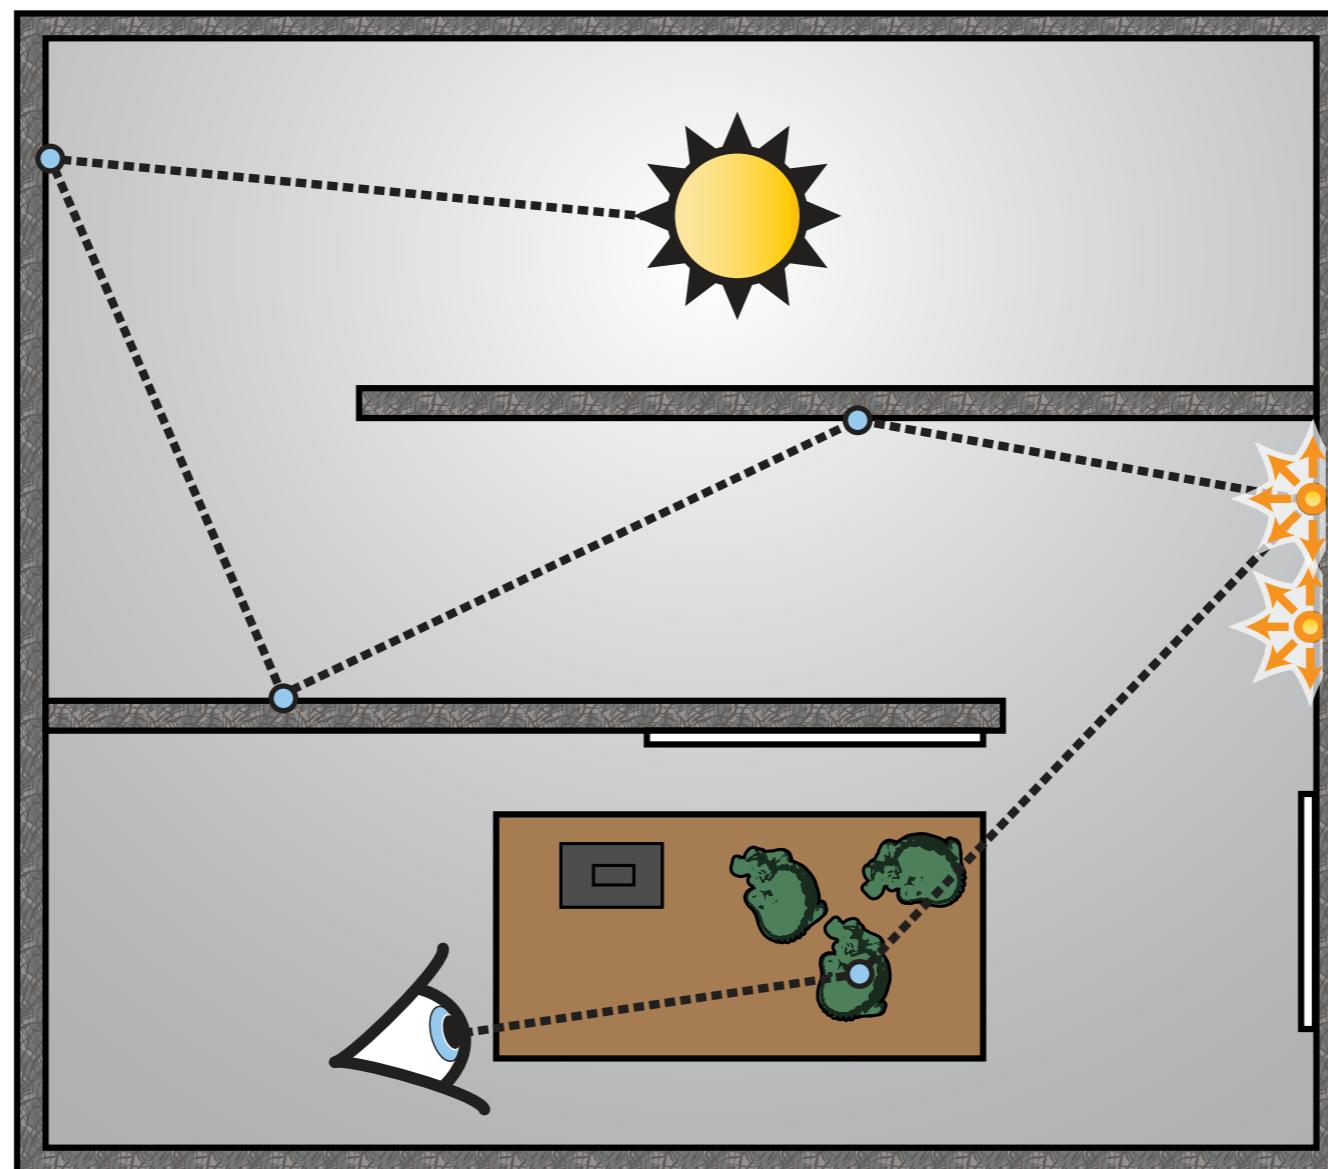
Improved Generation of VPLs

- Metropolis sampling of VPLs [Segovia et al. 2007]
 - generate additional VPLs by mutating paths



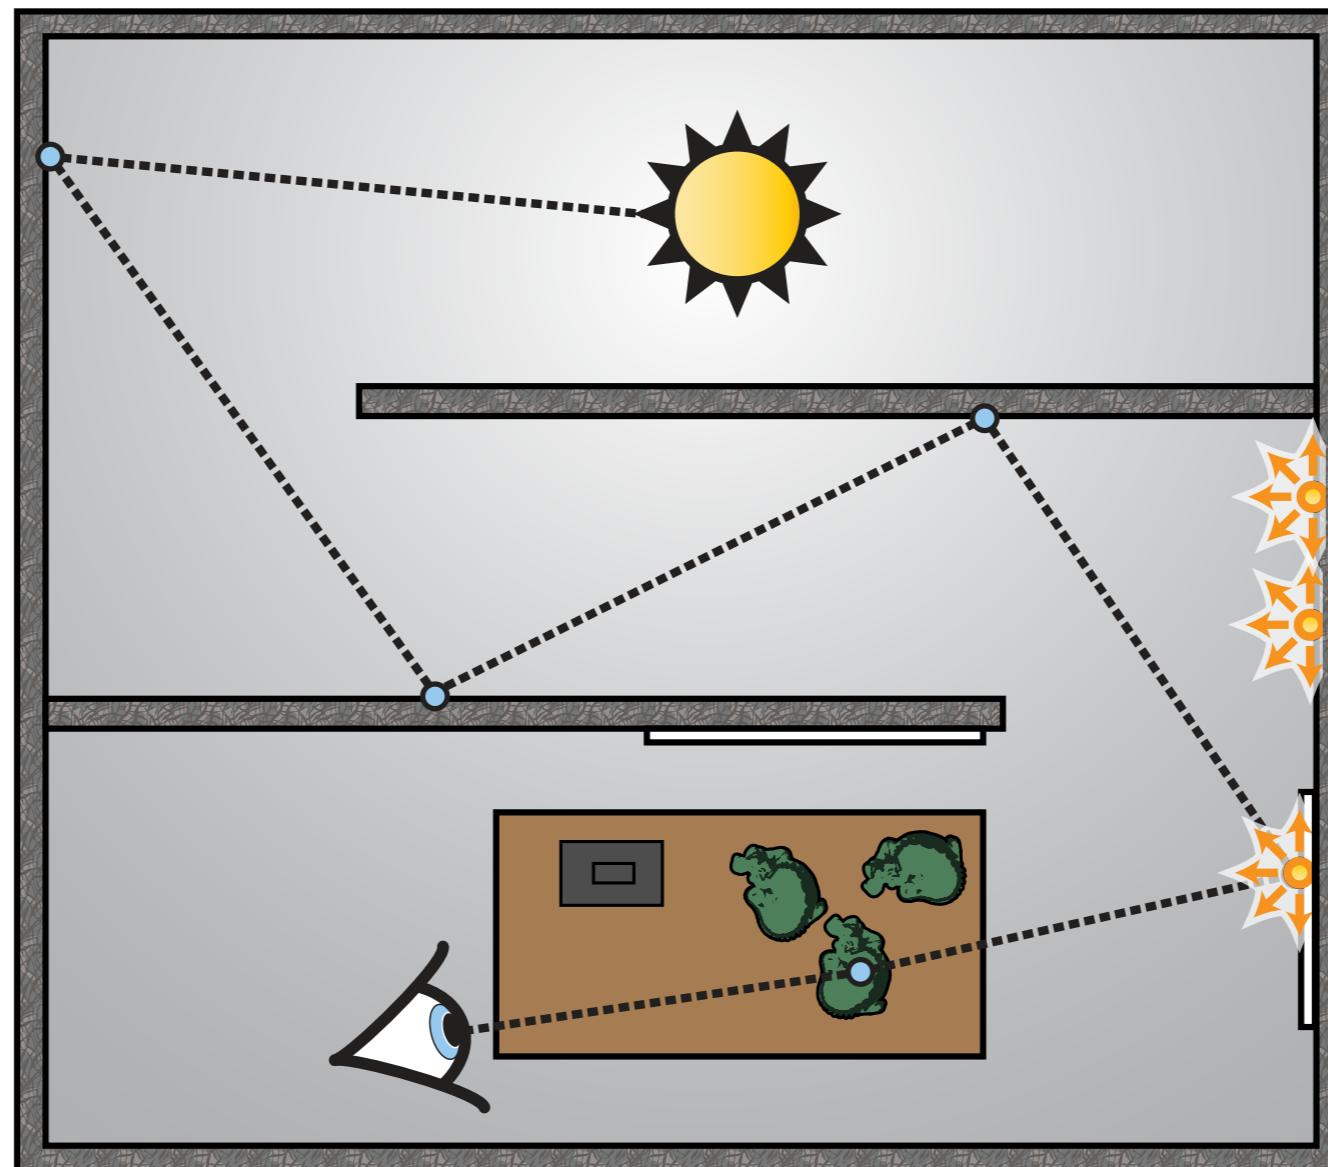
Improved Generation of VPLs

- Metropolis sampling of VPLs [Segovia et al. 2007]
 - generate additional VPLs by mutating paths



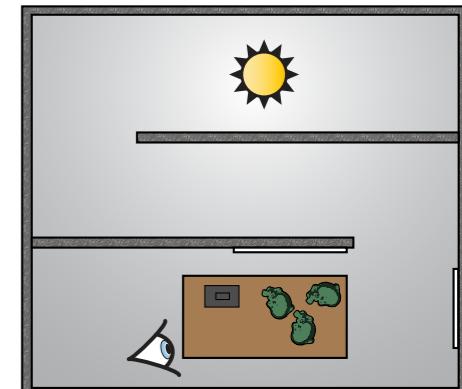
Improved Generation of VPLs

- Metropolis sampling of VPLs [Segovia et al. 2007]
 - generate additional VPLs by mutating paths

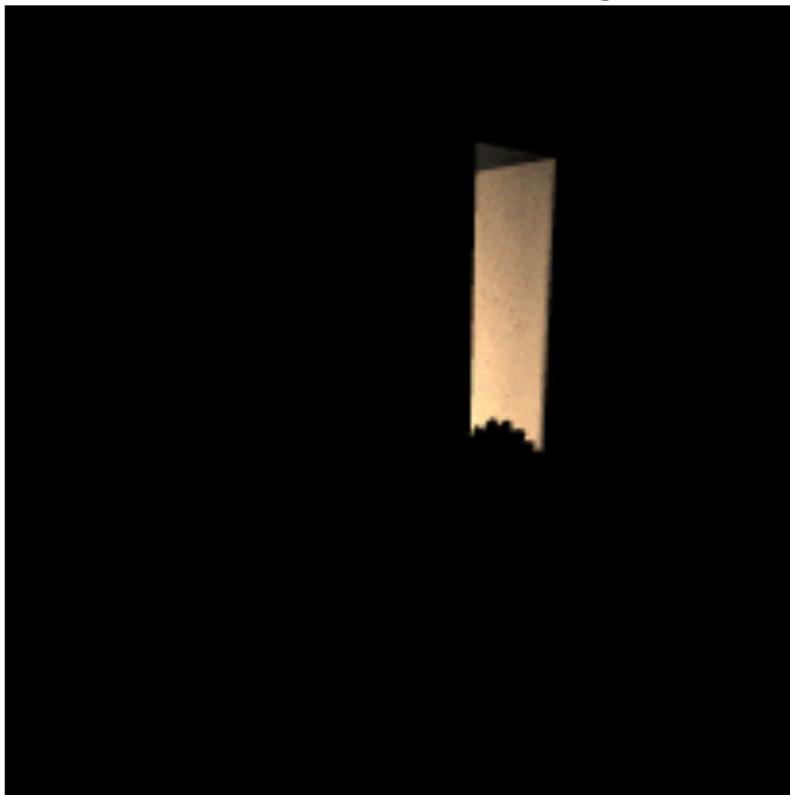


Improved Generation of VPLs

- Metropolis sampling of VPLs [Segovia et al. 2007]
 - generate additional VPLs by mutating paths



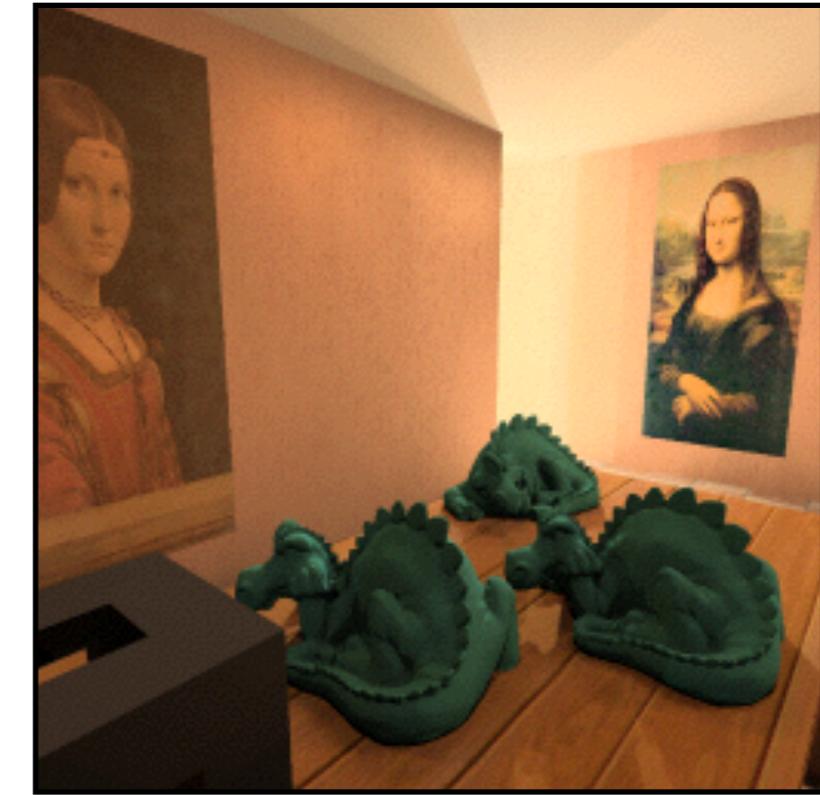
Instant Radiosity (IR)



Bidirectional IR



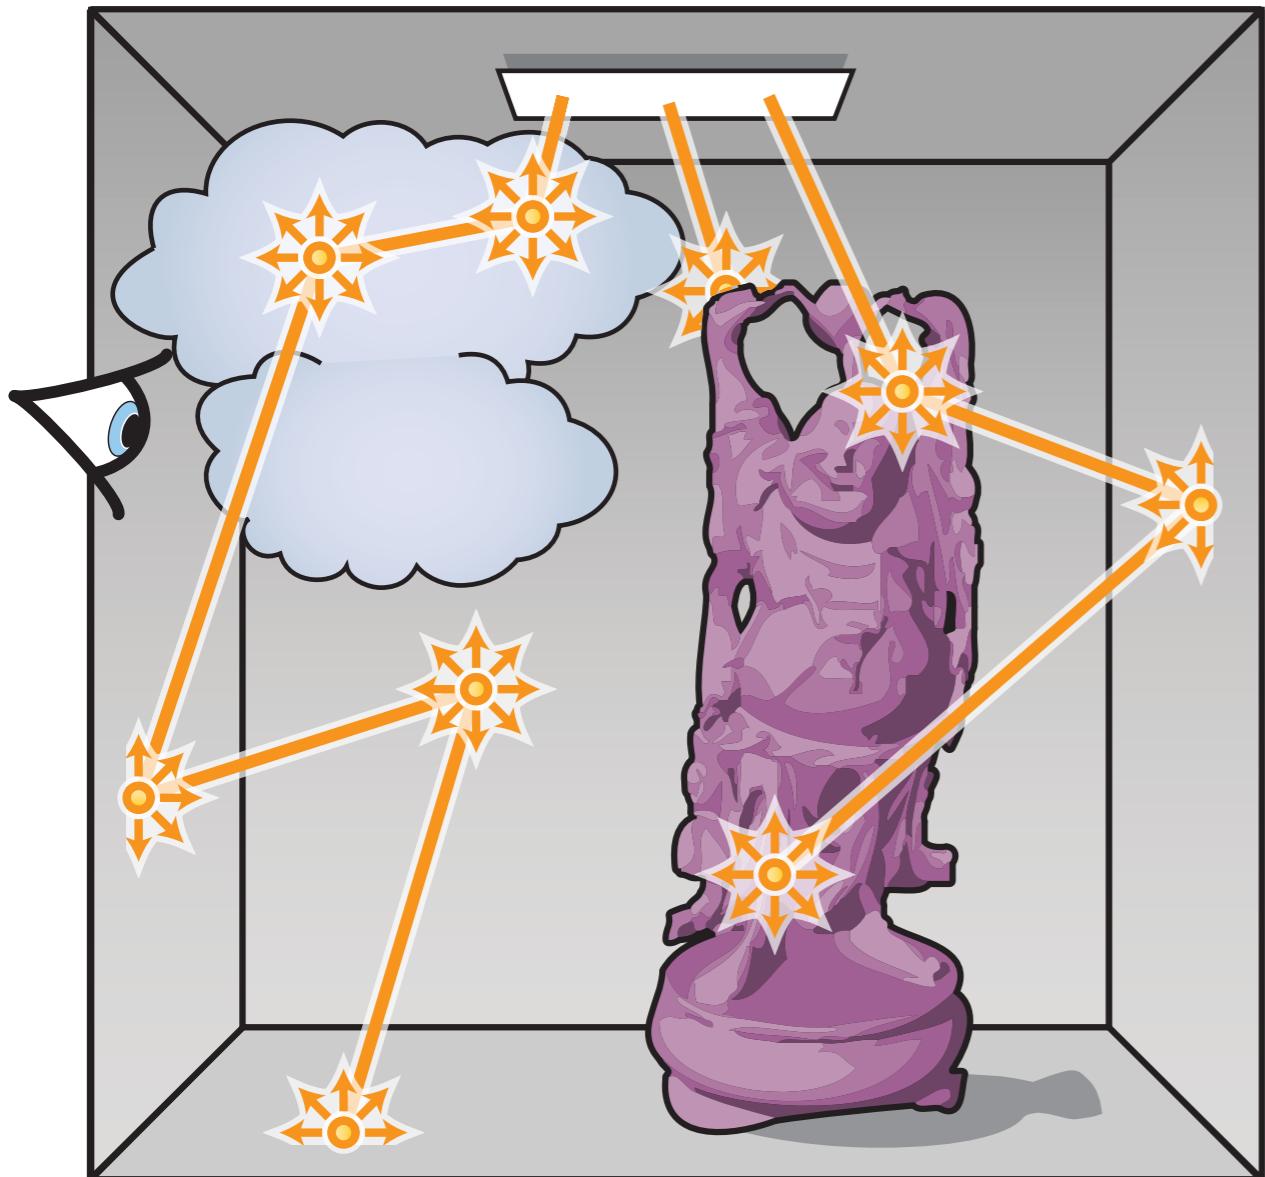
Metropolis IR



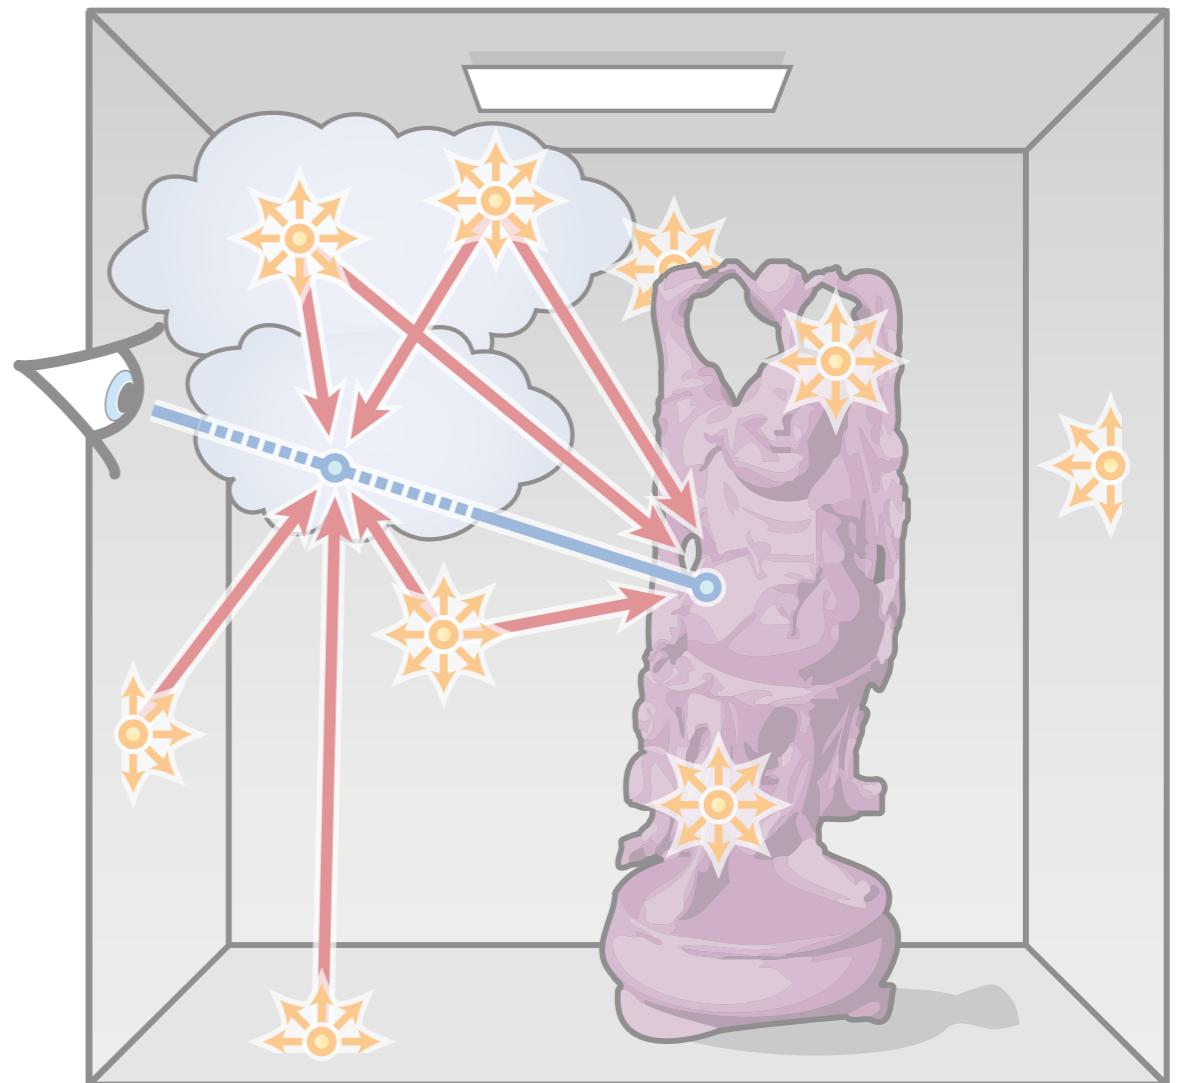
All approaches would ultimately yield the same image, but Metropolis gives good results already with a few VPLs.

Many-Light Rendering

Generation of VPLs

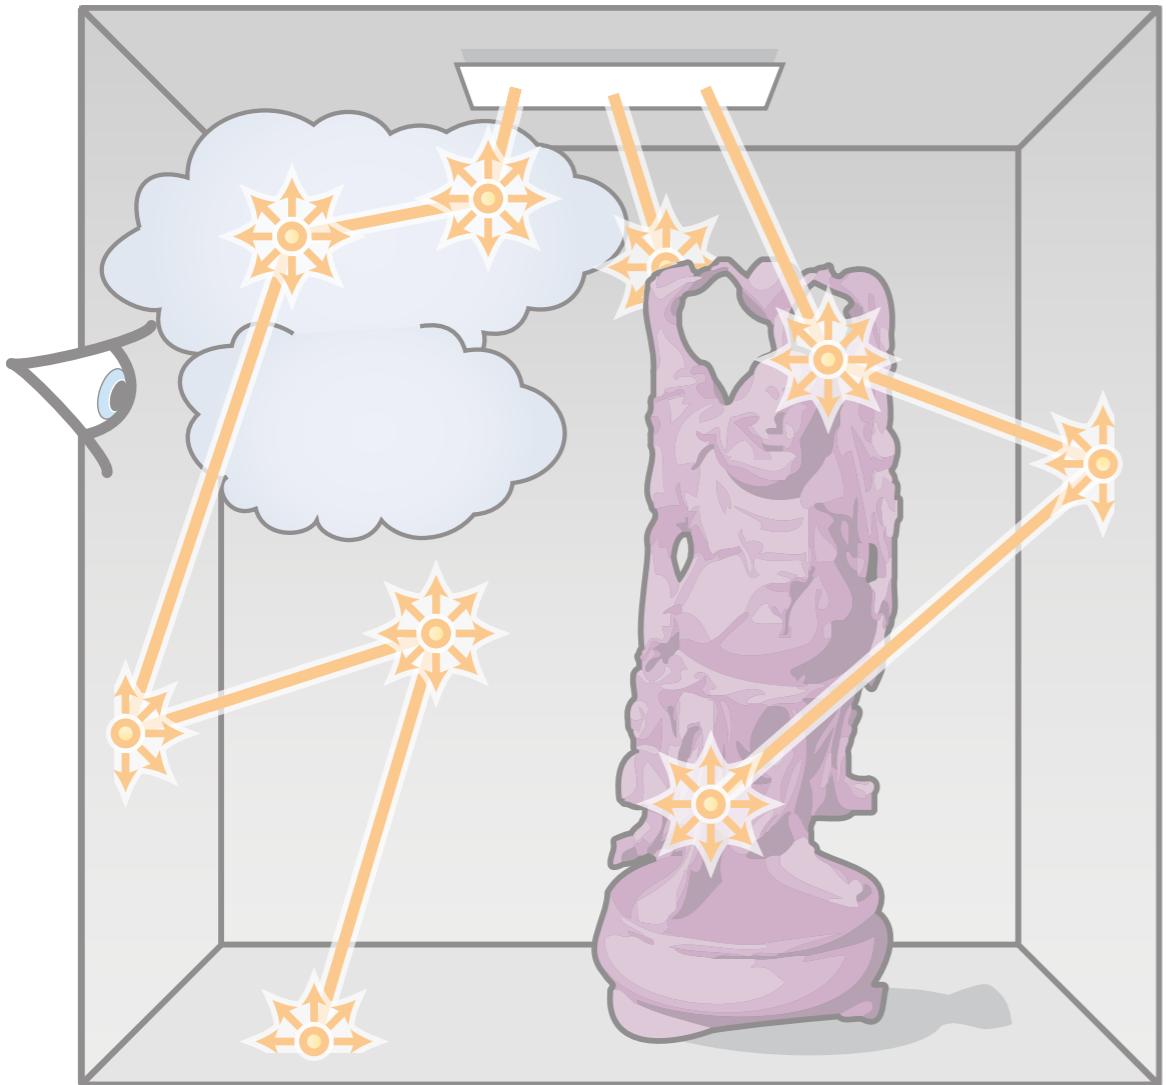


Lighting with VPLs

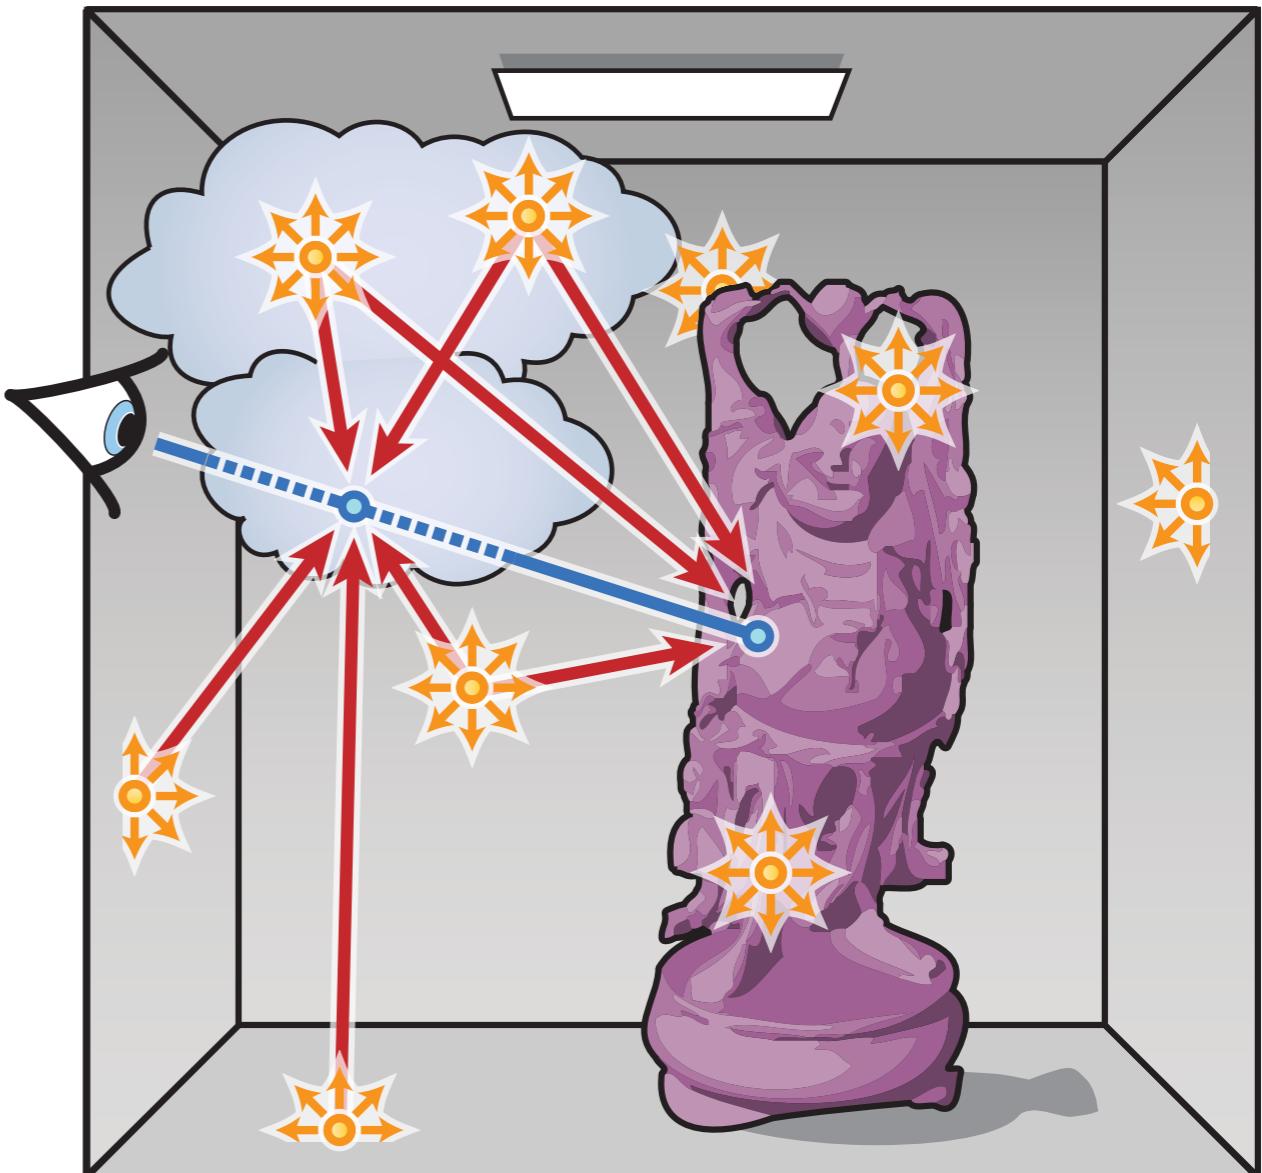


Many-Light Rendering

Generation of VPLs



Lighting with VPLs



Lighting with VPLs

- Rendering equation:

$$L(\mathbf{x}_1, \mathbf{x}_0) = L_e(\mathbf{x}_1, \mathbf{x}_0) + L_r(\mathbf{x}_1, \mathbf{x}_0)$$
$$L_r(\mathbf{x}_1, \mathbf{x}_0) = \int_A f_r(\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_0) G(\mathbf{x}_1, \mathbf{x}_2) L(\mathbf{x}_2, \mathbf{x}_1) dA(\mathbf{x}_2)$$

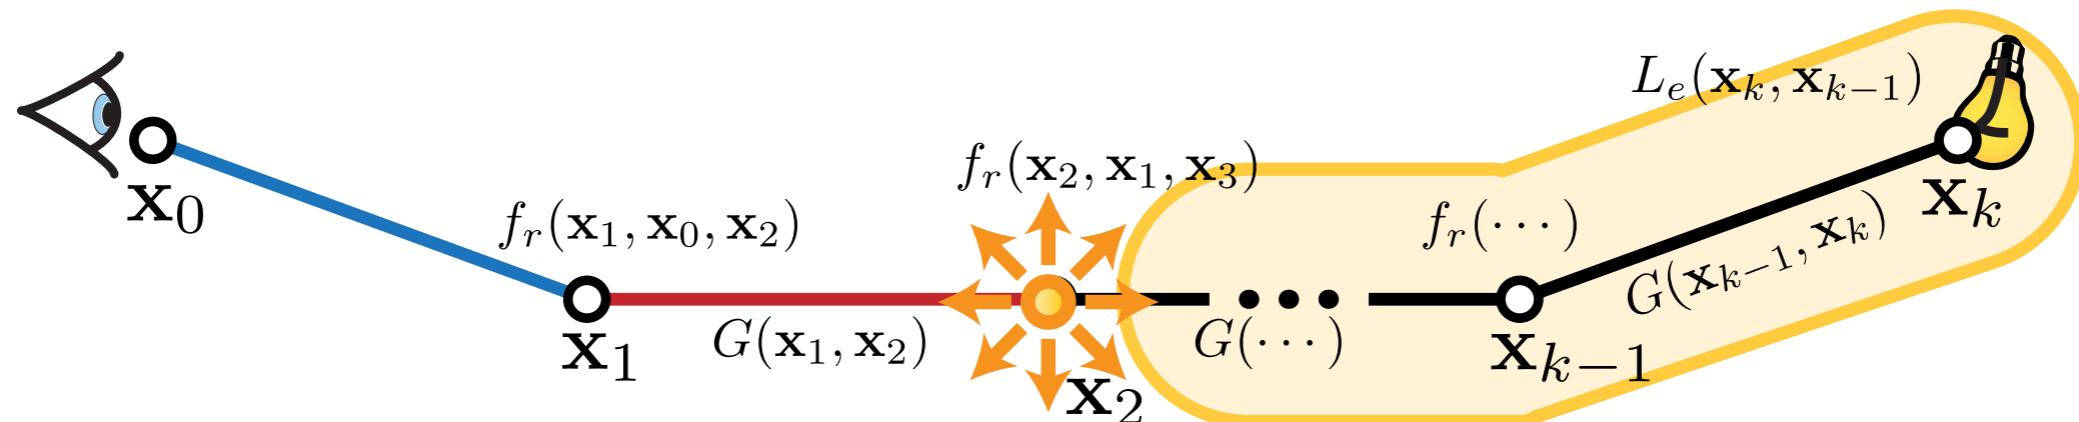
$\approx \sum_{i=1}^N f_r(\mathbf{x}_1, \mathbf{x}_{2,i}, \mathbf{x}_0) G(\mathbf{x}_1, \mathbf{x}_{2,i}) f_r(\mathbf{x}_{2,i}, \mathbf{x}_{3,i}, \mathbf{x}_1) I_i$

Approximation using N VPLs

recursion is hidden in the generation of VPLs

BRDF at i-th VPL

"Intensity" of i-th VPL



Lighting with VPLs

- Rendering equation:

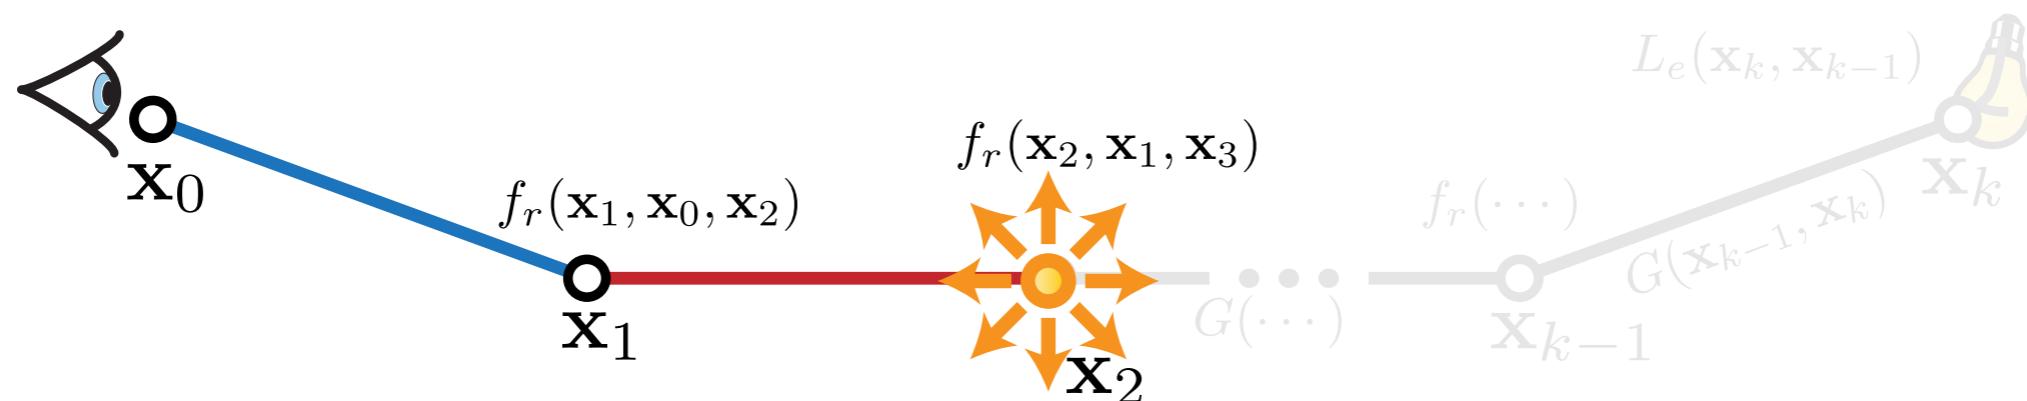
$$L(\mathbf{x}_1, \mathbf{x}_0) = L_e(\mathbf{x}_1, \mathbf{x}_0) + L_r(\mathbf{x}_1, \mathbf{x}_0)$$

$$L_r(\mathbf{x}_1, \mathbf{x}_0) = \int_A f_r(\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_0) G(\mathbf{x}_1, \mathbf{x}_2) L(\mathbf{x}_2, \mathbf{x}_1) dA(\mathbf{x}_2)$$

$$\approx \sum_{i=1}^N f_r(\mathbf{x}_1, \mathbf{x}_{2,i}, \mathbf{x}_0) G(\mathbf{x}_1, \mathbf{x}_{2,i}) f_r(\mathbf{x}_{2,i}, \mathbf{x}_{3,i}, \mathbf{x}_1) I_i$$

Intensity of a VPL: $I = \frac{L_e(\mathbf{x}_k, \mathbf{x}_{k-1}) T(\mathbf{x}_2, \dots, \mathbf{x}_k)}{p(\mathbf{x}_2, \dots, \mathbf{x}_k)}$

Throughput of
the light path



Lighting with VPLs

- Rendering equation:

$$L(\mathbf{x}_1, \mathbf{x}_0) = L_e(\mathbf{x}_1, \mathbf{x}_0) + L_r(\mathbf{x}_1, \mathbf{x}_0)$$

$$L_r(\mathbf{x}_1, \mathbf{x}_0) = \int_A f_r(\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_0) G(\mathbf{x}_1, \mathbf{x}_2) L(\mathbf{x}_2, \mathbf{x}_1) dA(\mathbf{x}_2)$$

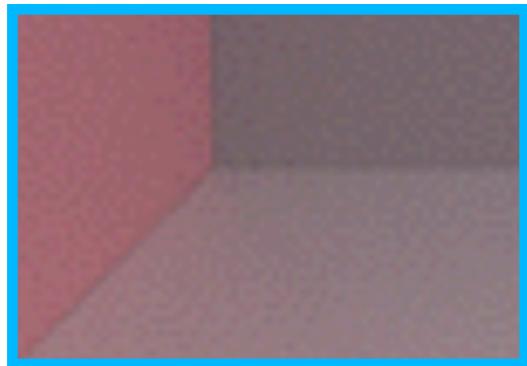
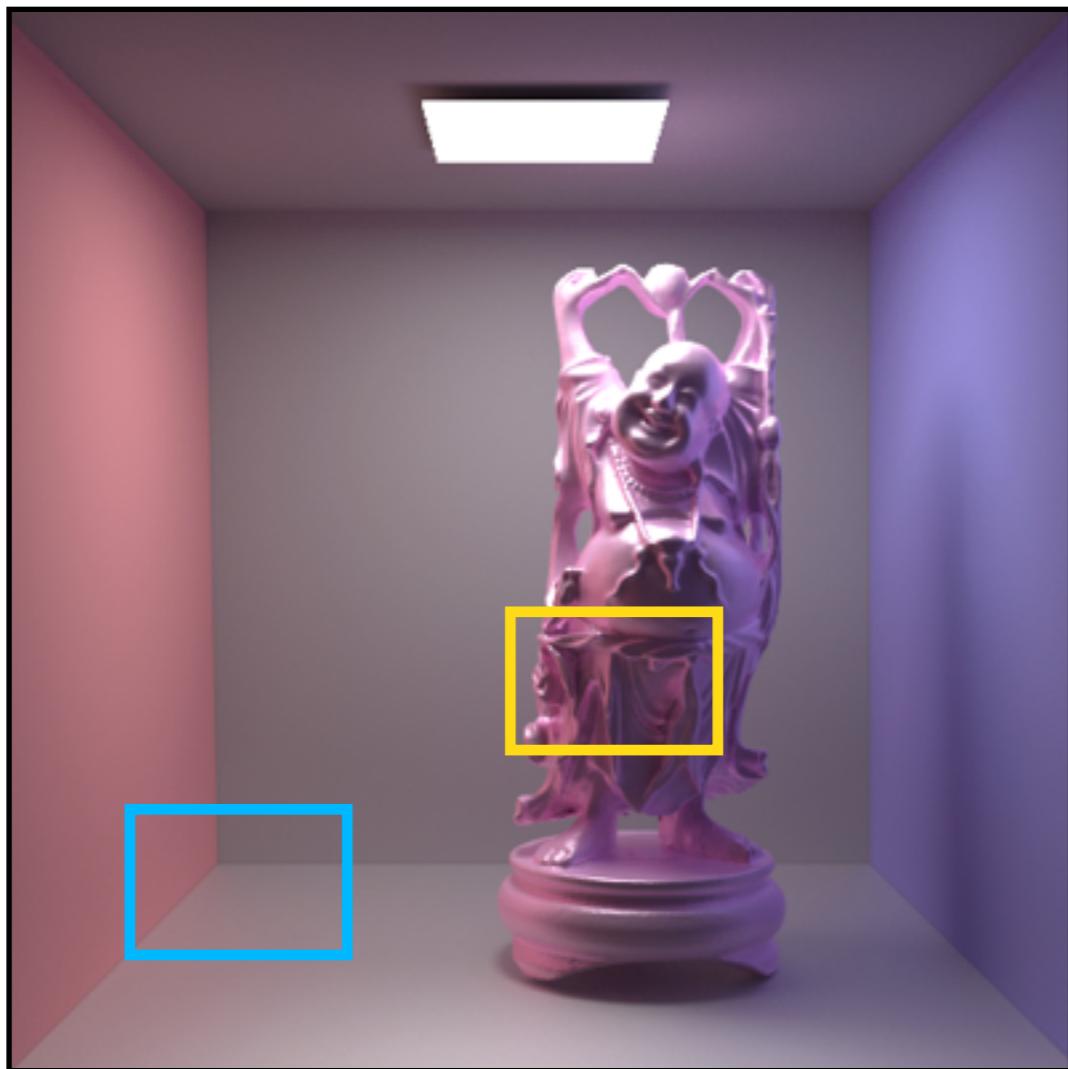
$$\approx \sum_{i=1}^N f_r(\mathbf{x}_1, \mathbf{x}_{2,i}, \mathbf{x}_0) G(\mathbf{x}_1, \mathbf{x}_{2,i}) f_r(\mathbf{x}_{2,i}, \mathbf{x}_{3,i}, \mathbf{x}_1) I_i$$

- For each VPL we store:

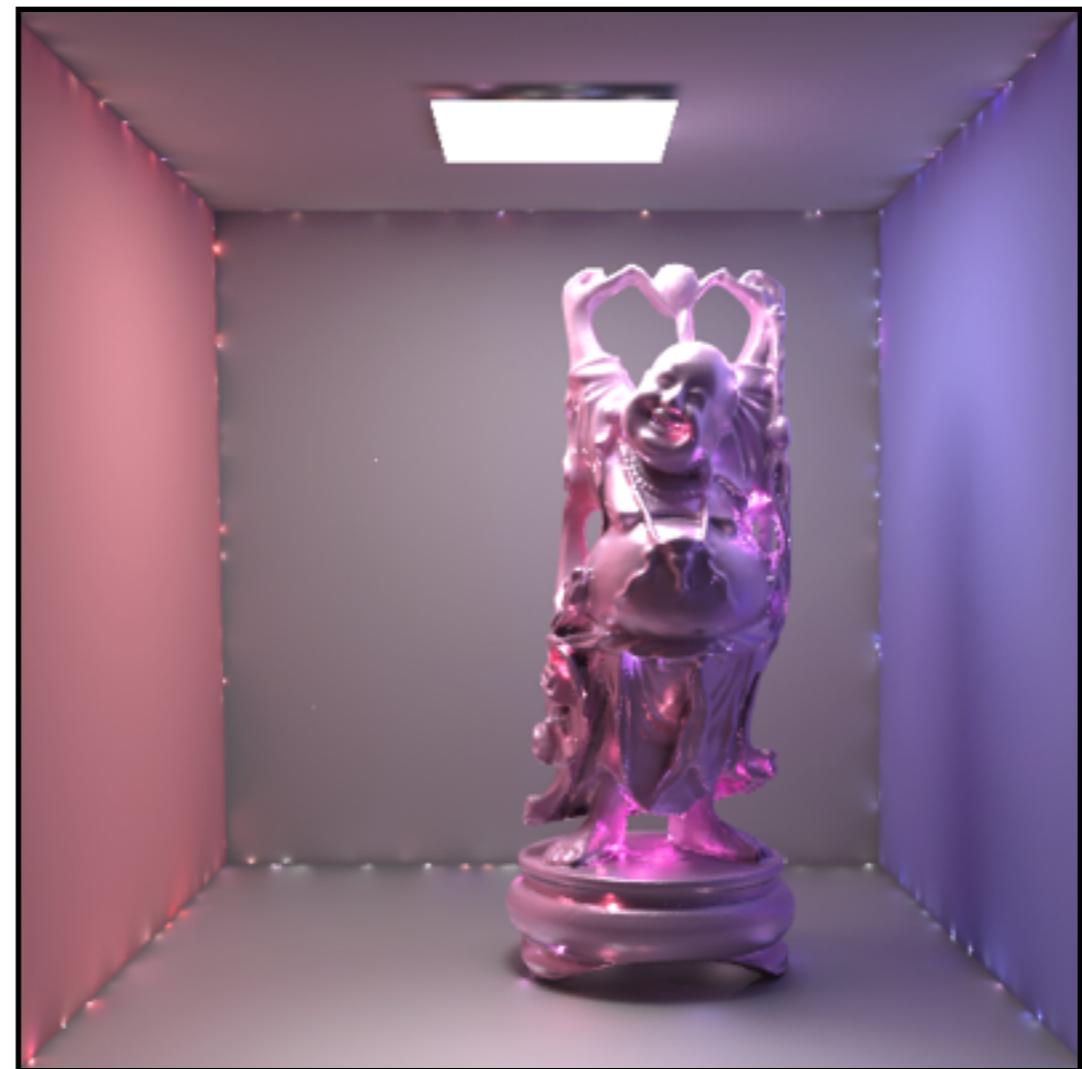
- position
- intensity
- direction towards the previous vertex on the light path
- BRDF & local shading frame (normal, tangent)

Lighting with VPLs

Reference



Approximation with VPLs

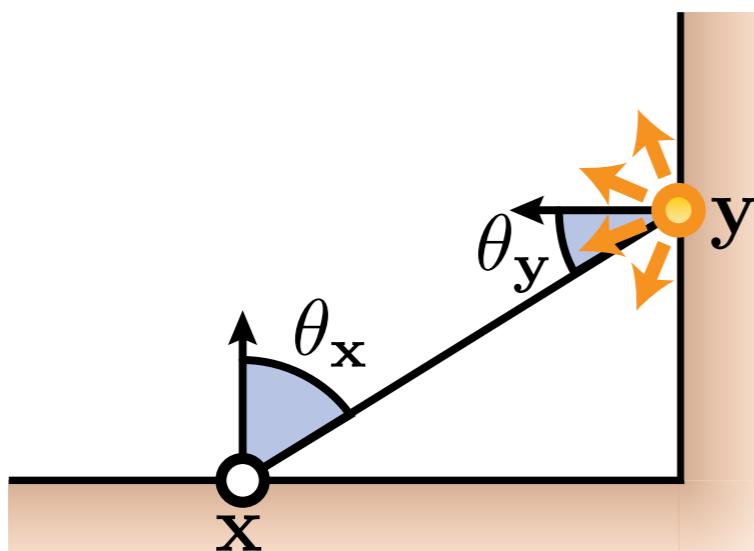


Lighting with VPLs

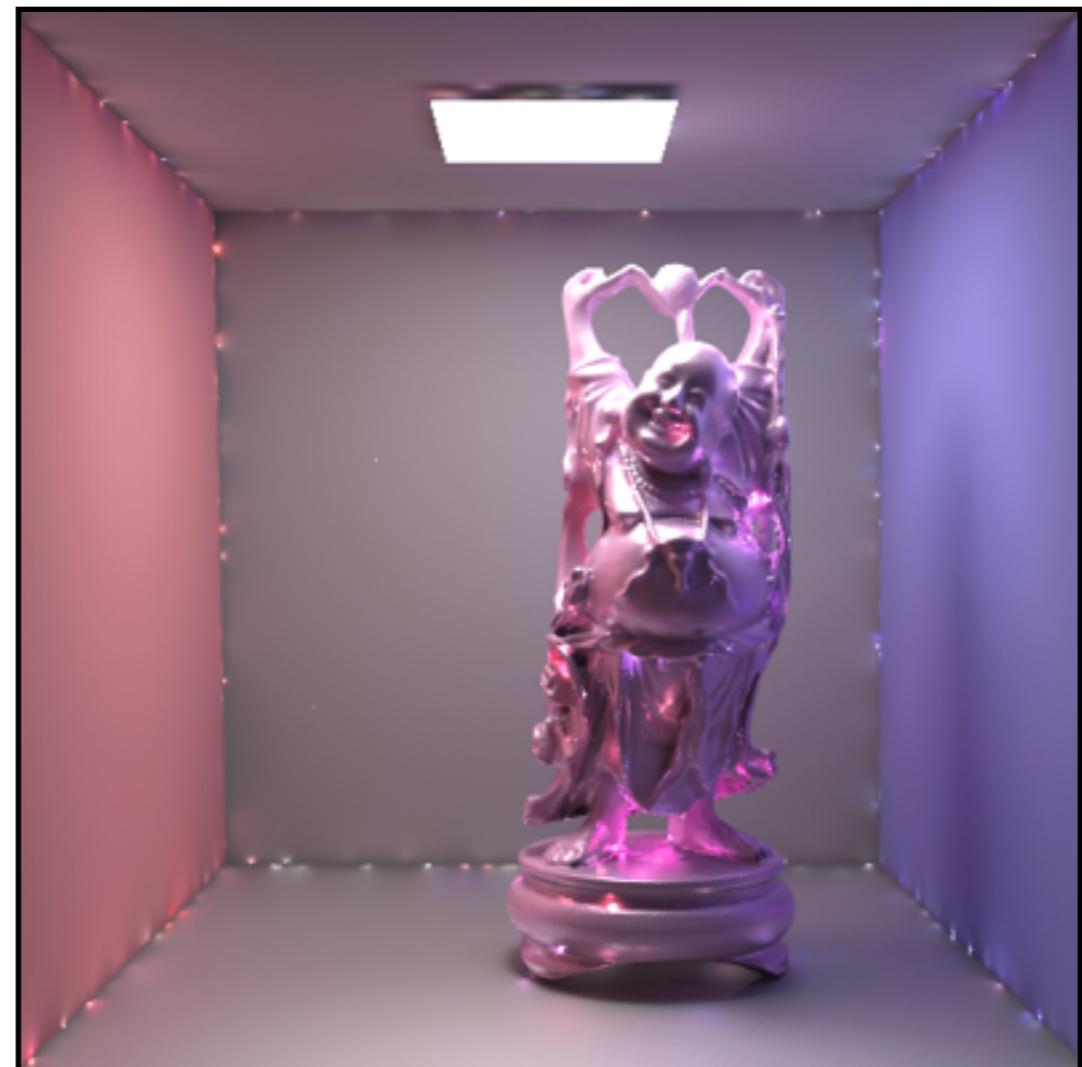
Splotches!!!

Reason: singularity in G-term

$$G(\mathbf{x}, \mathbf{y}) = V(\mathbf{x}, \mathbf{y}) \frac{\cos \theta_{\mathbf{x}} \cos \theta_{\mathbf{y}}}{\|\mathbf{x} - \mathbf{y}\|^2}$$



Approximation with VPLs

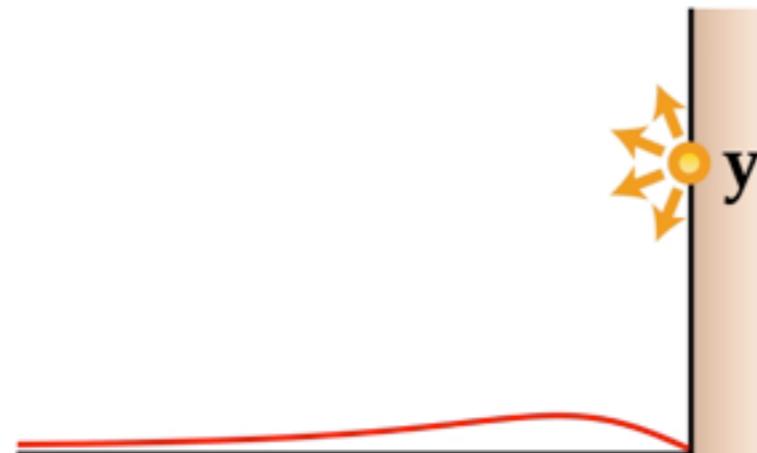


Lighting with VPLs

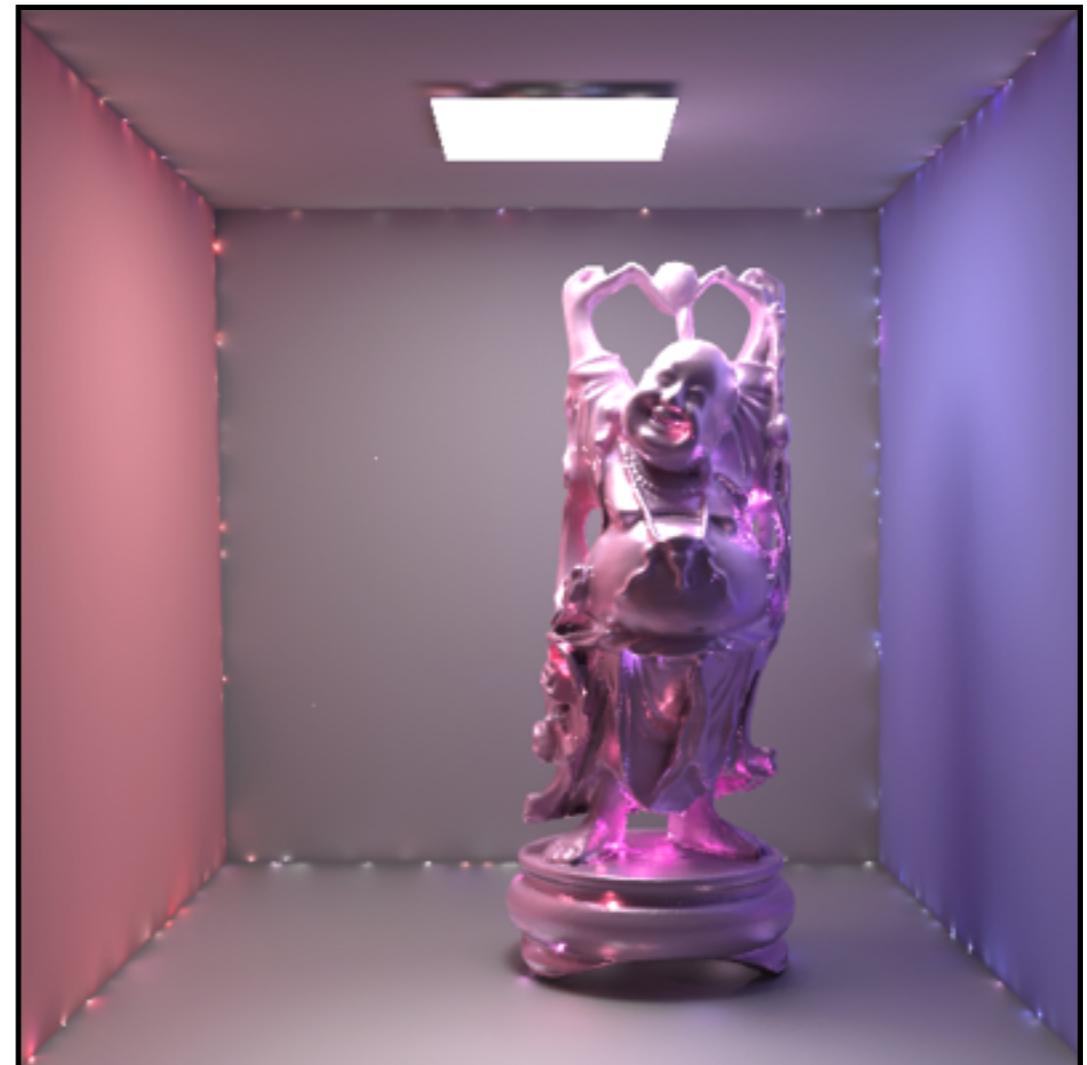
Splotches!!!

Reason: singularity in G-term

$$G(\mathbf{x}, \mathbf{y}) = V(\mathbf{x}, \mathbf{y}) \frac{\cos \theta_{\mathbf{x}} \cos \theta_{\mathbf{y}}}{\|\mathbf{x} - \mathbf{y}\|^2}$$



Approximation with VPLs



All points are lit by the same set of VPLs

Correlation shows up as splotches



Bounding the G-term

- Remove the singularity by bounding (clamping) the geometry term to some maximum

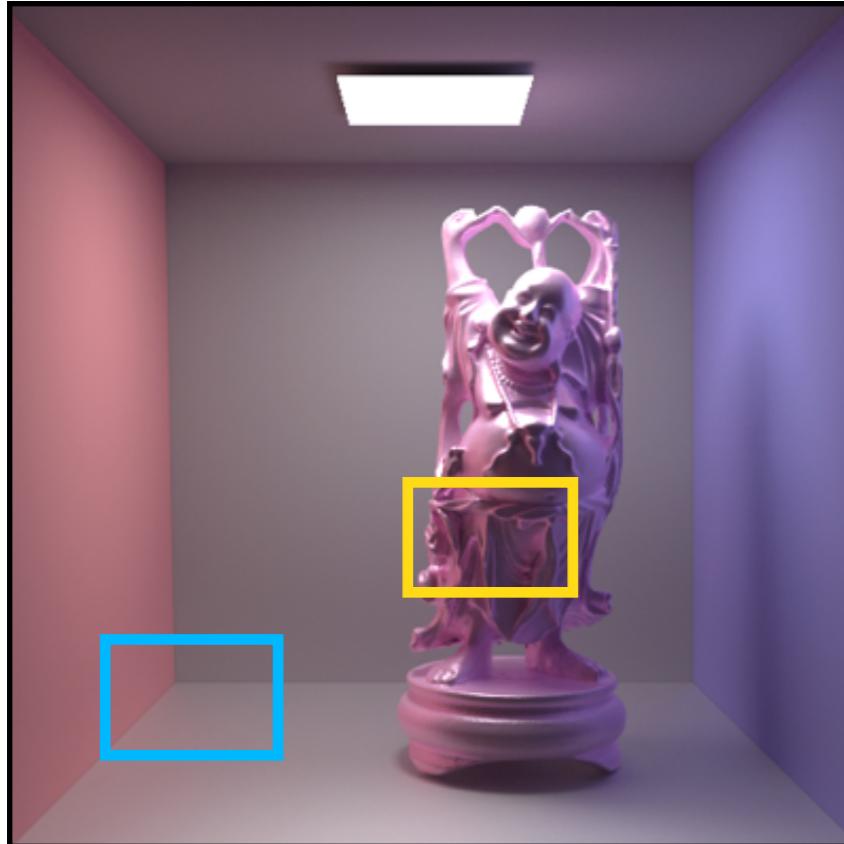
$$G_b(\mathbf{x}, \mathbf{y}) = \min(G(\mathbf{x}, \mathbf{y}), b)$$

User-defined
maximum value

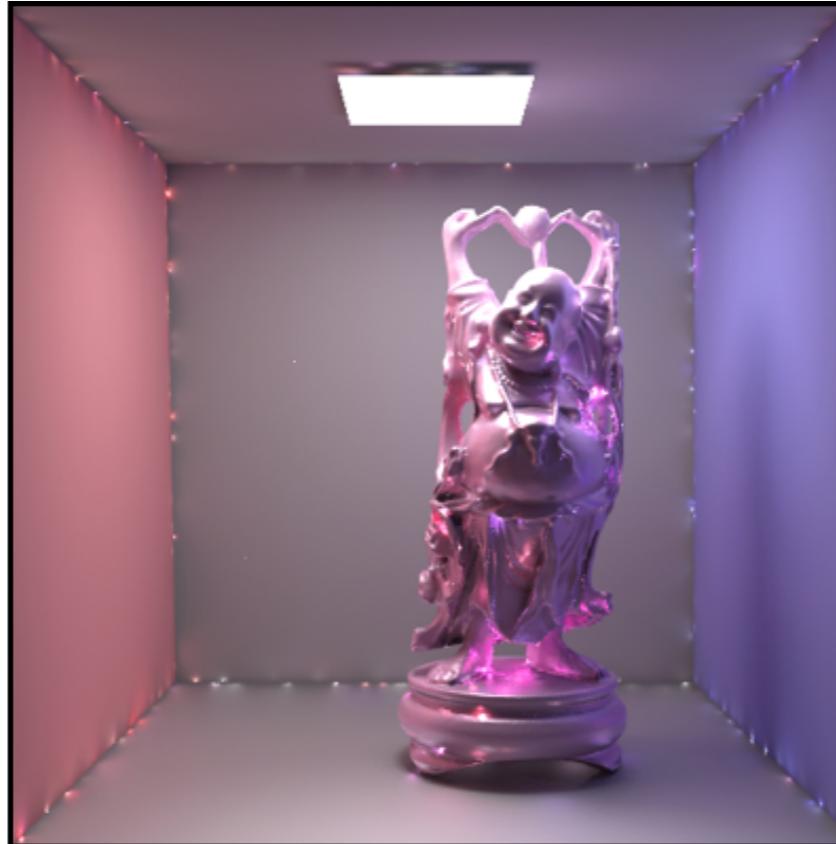
- Extremely simple and cheap 
- Removes short-distance light transport,  cavities appear darker

Bounding the G-term

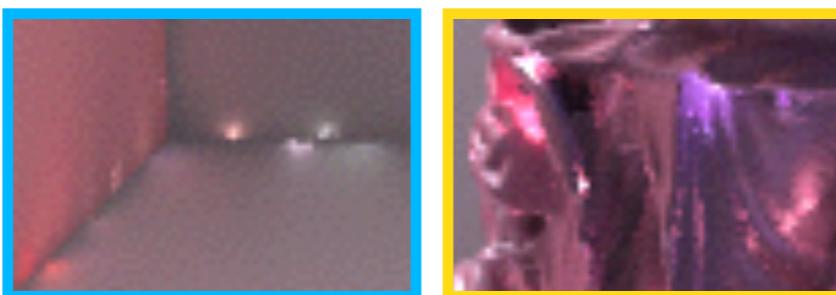
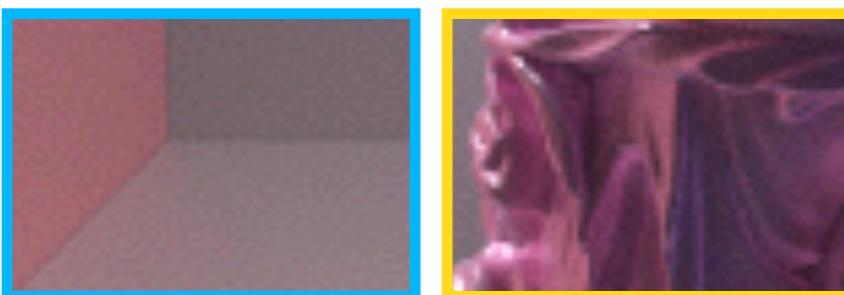
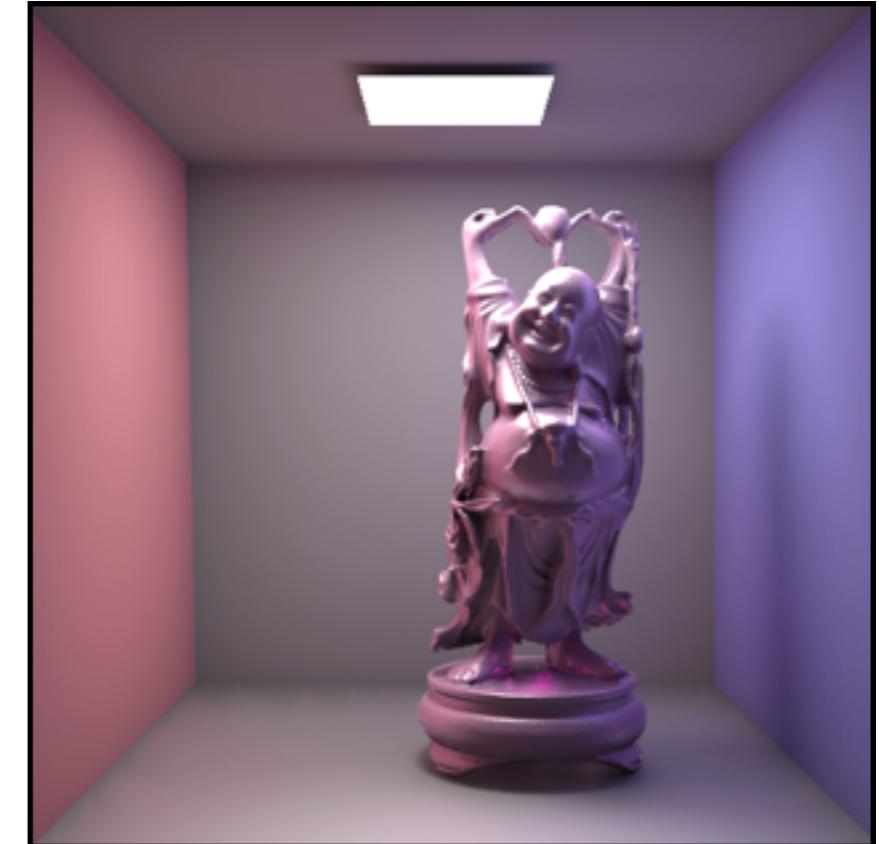
Reference



VPLs



VPLs with bounded G



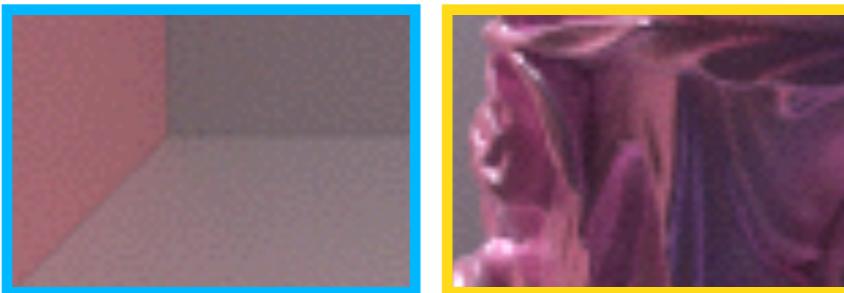
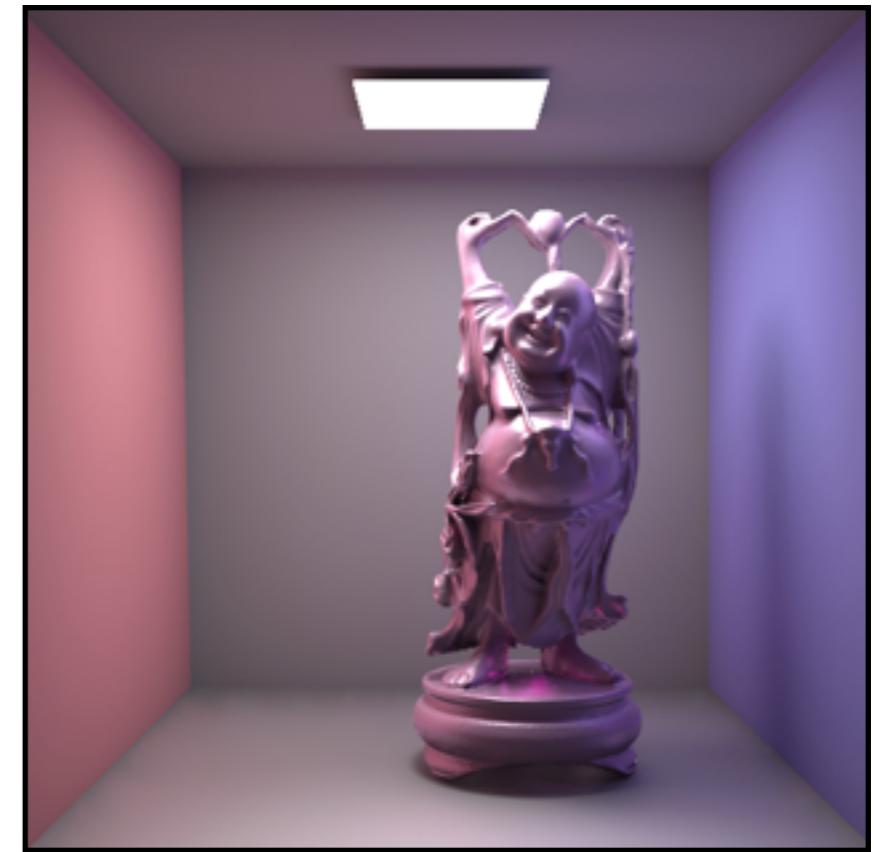
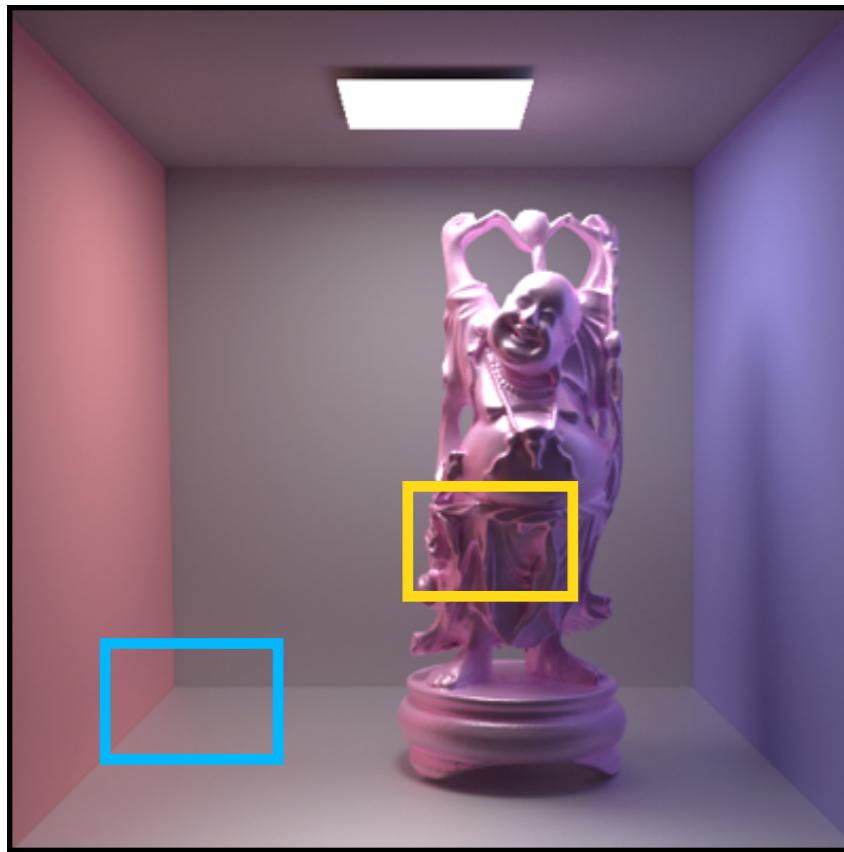
using $G(x, y)$

Bounding the G-term

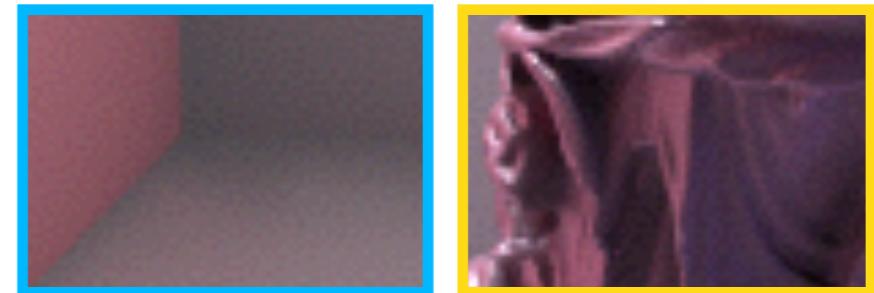
Reference

Difference

VPLs with bounded G



$$G(\mathbf{x}, \mathbf{y}) - G_b(\mathbf{x}, \mathbf{y})$$



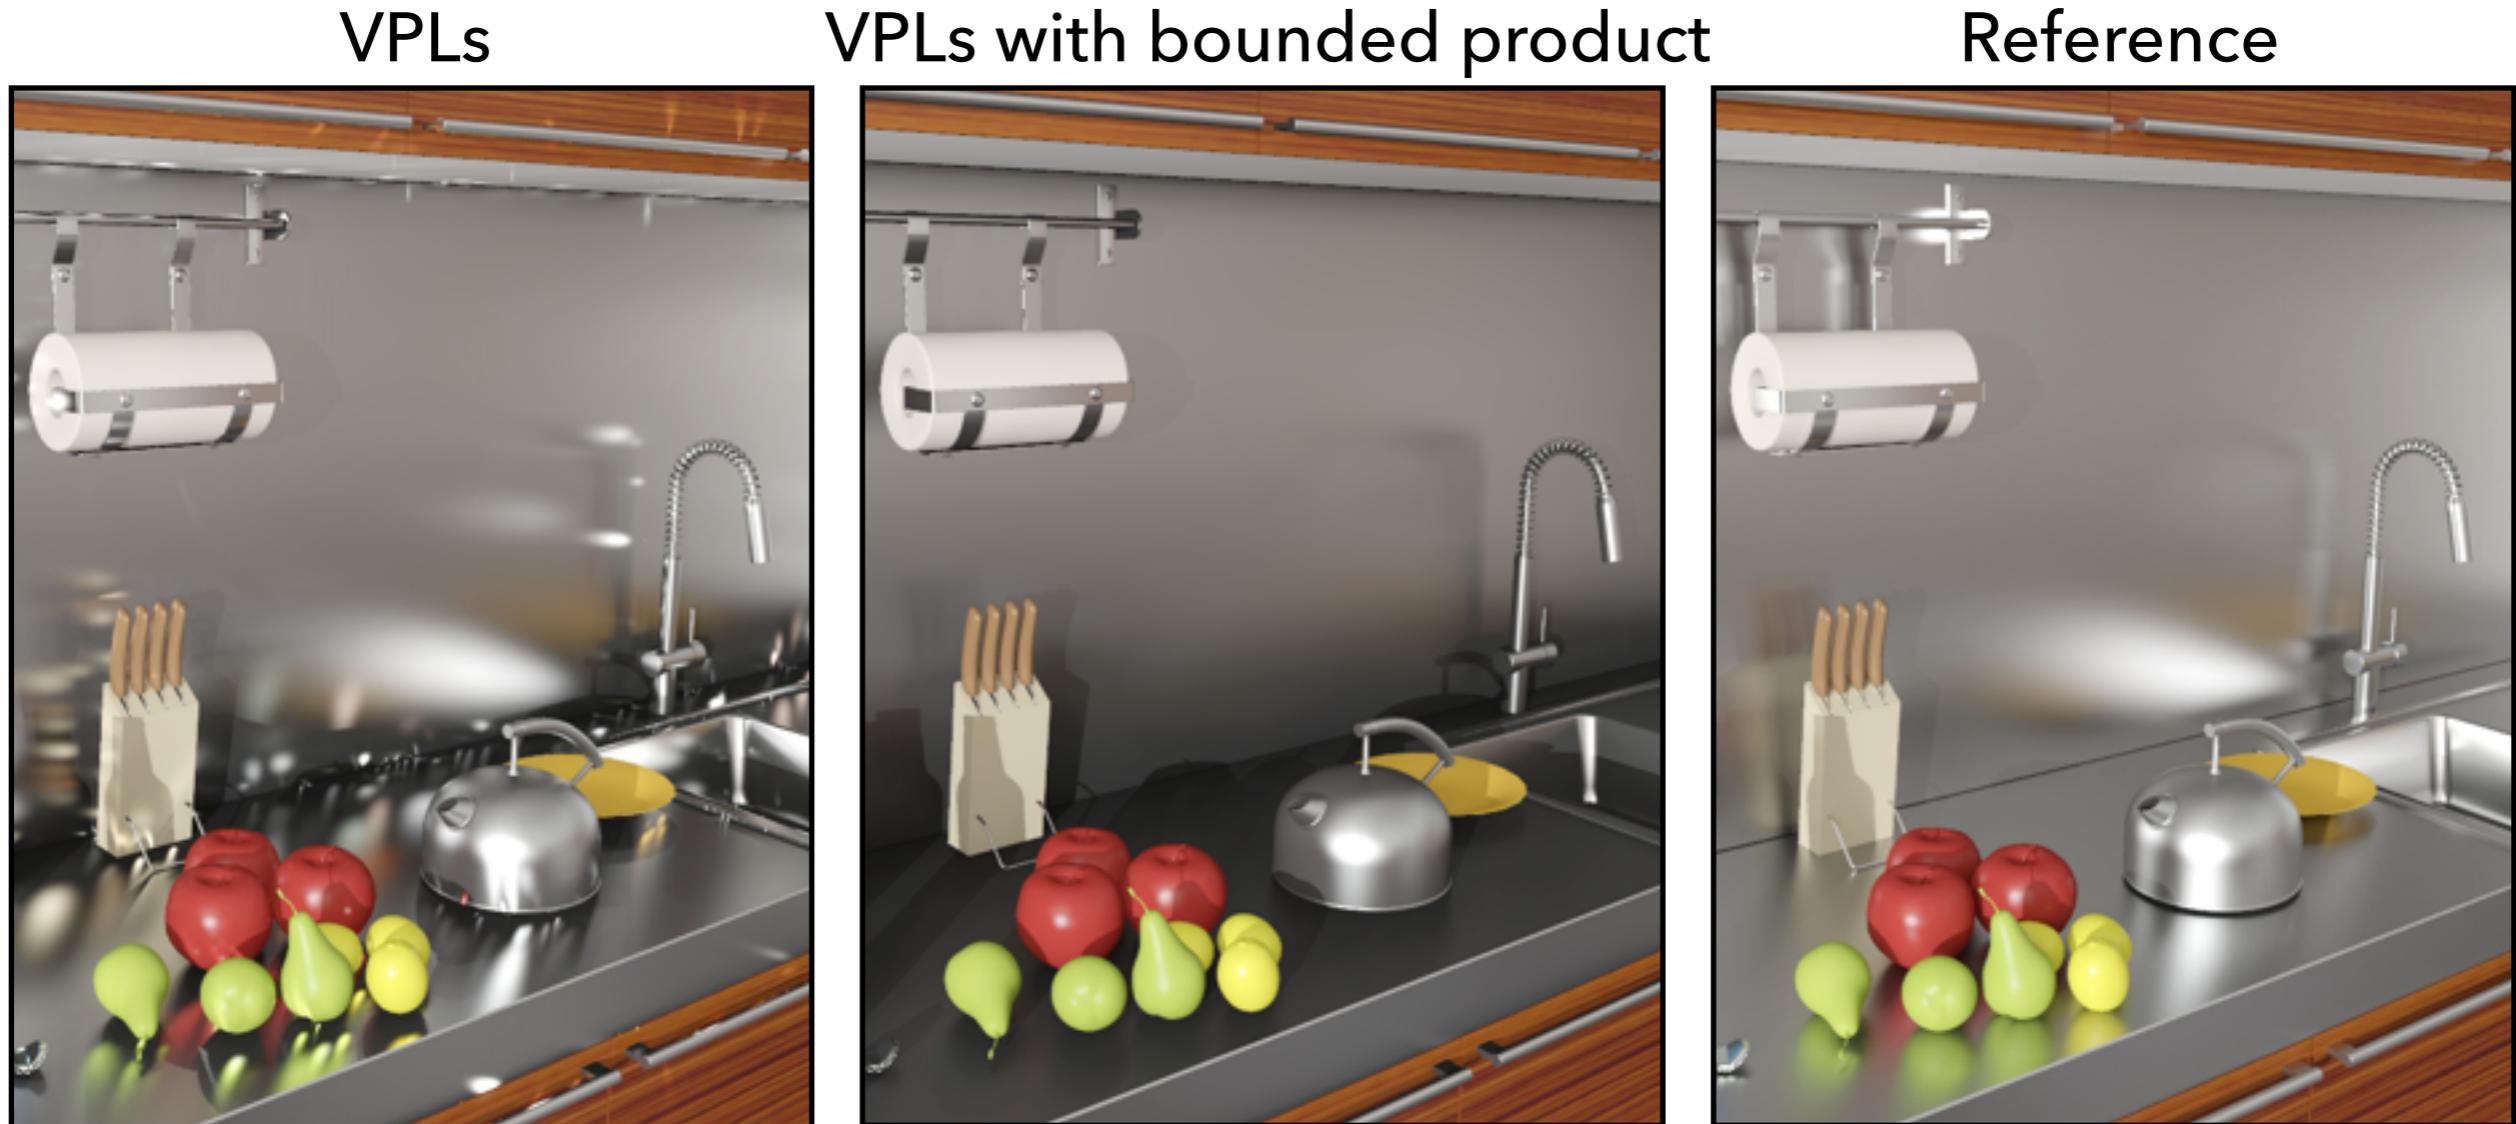
We need to compensate for the energy loss (bias)!

Bias Compensation

- Use bounded VPLs to obtain a noise free, *biased* solution
- Recover the lost energy using different technique
 - [Kollig and Keller 2004] (use path tracing for bias compensation: slow),
[Raab et al. 2008] (path-traced bias compensation for volumes),
[Davidovič et al. 2010] (create local lights, approximate),
[Novák et al. 2011] (bias compensation in screen-space, approximate)
[Engelhardt et al. 2012] (approximate bias compensation for volumes)
- Always some trade-off... speed, quality, elegance

Glossy Surfaces

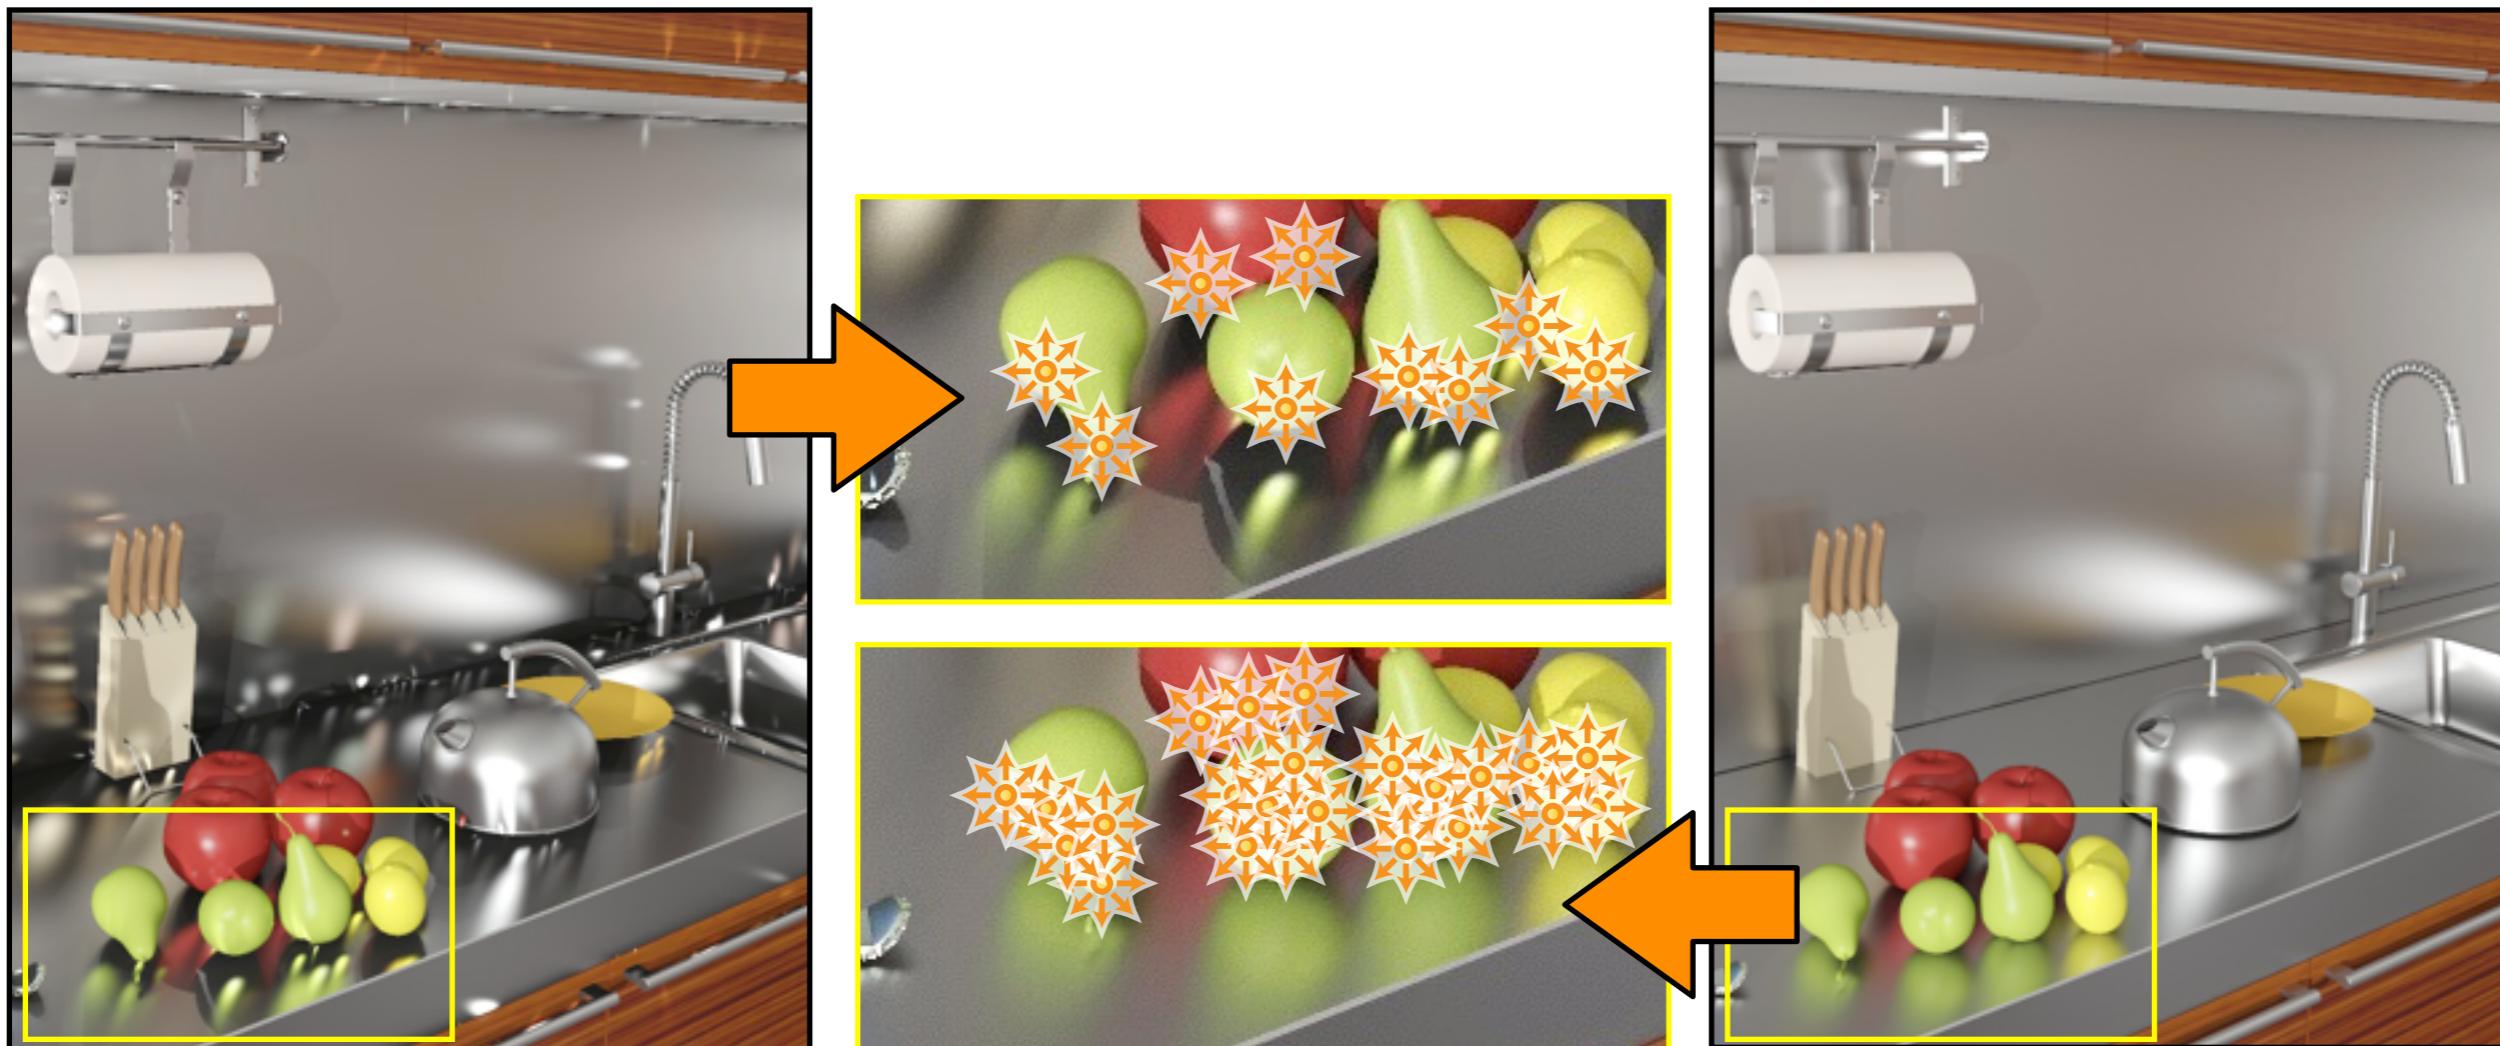
- Glossy surfaces further emphasize splotches
 - Naive solution: bound the product $f_r(\mathbf{x}_1 \dots) G(\mathbf{x}_1, \mathbf{x}_2) f_r(\mathbf{x}_2 \dots)$



Local VPLs

- Solution 1: create additional (local) VPLs [Davidović et al. 2010]
 - Global (bounded) VPLs distributed traditionally
 - Local VPLs distributed from camera, add the missing energy only

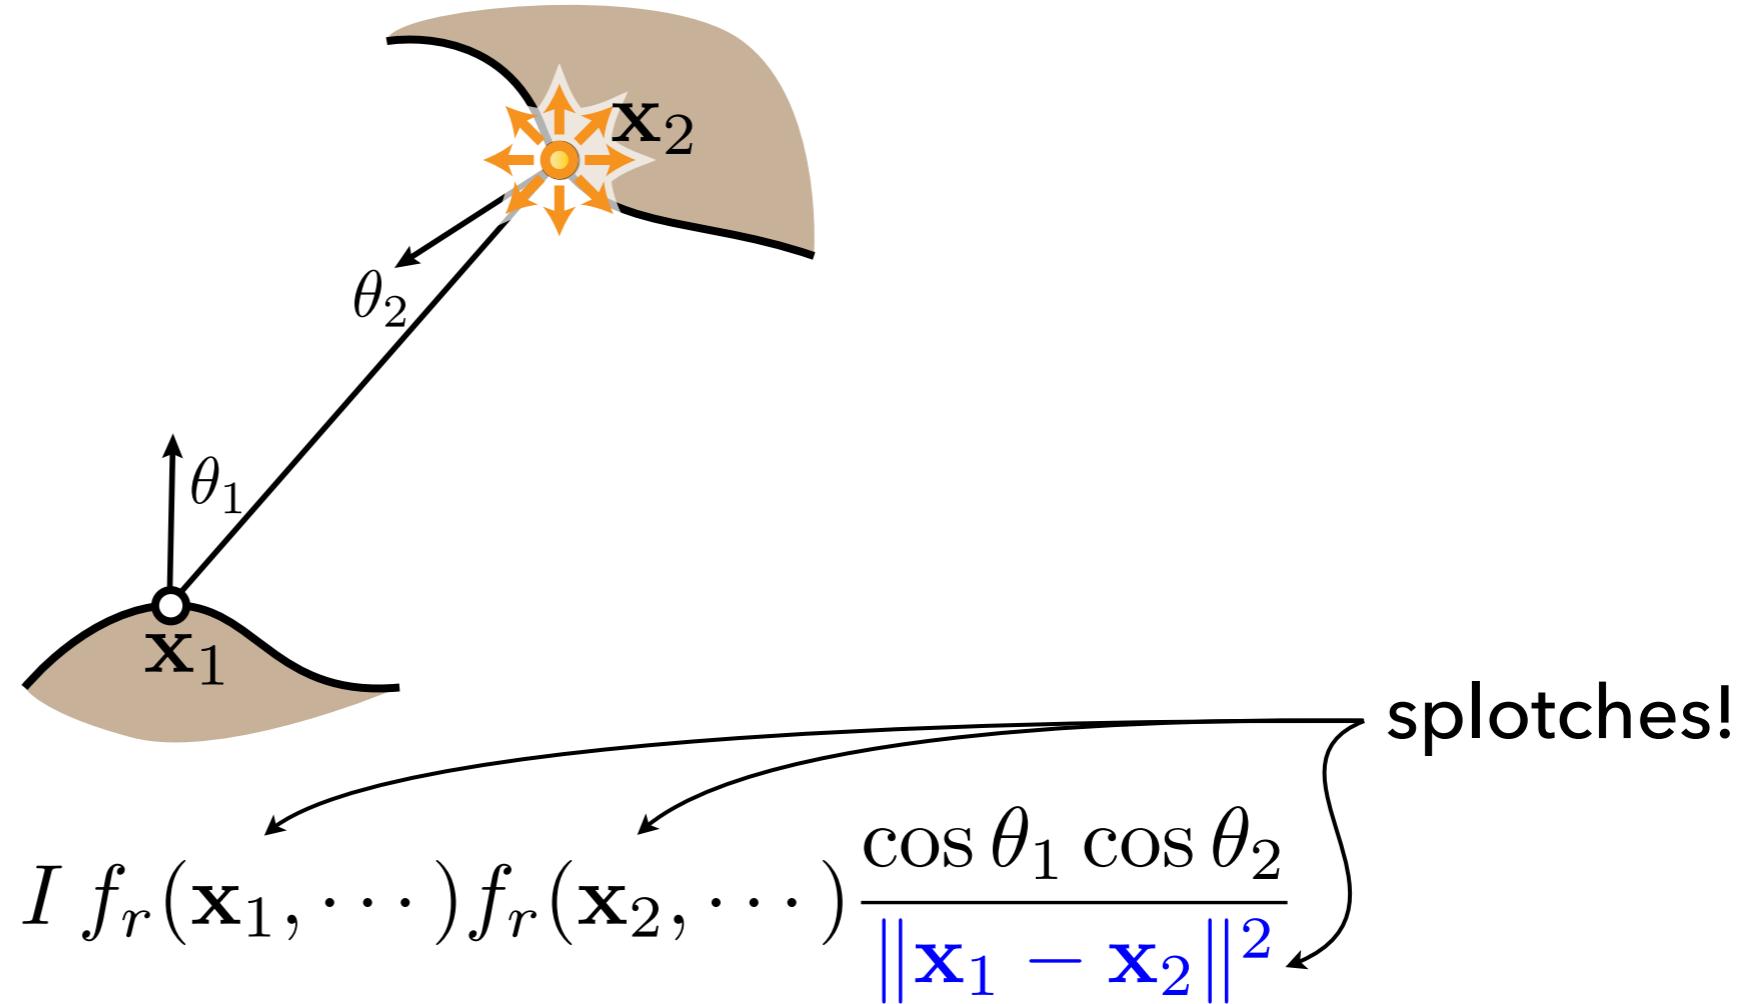
Reference



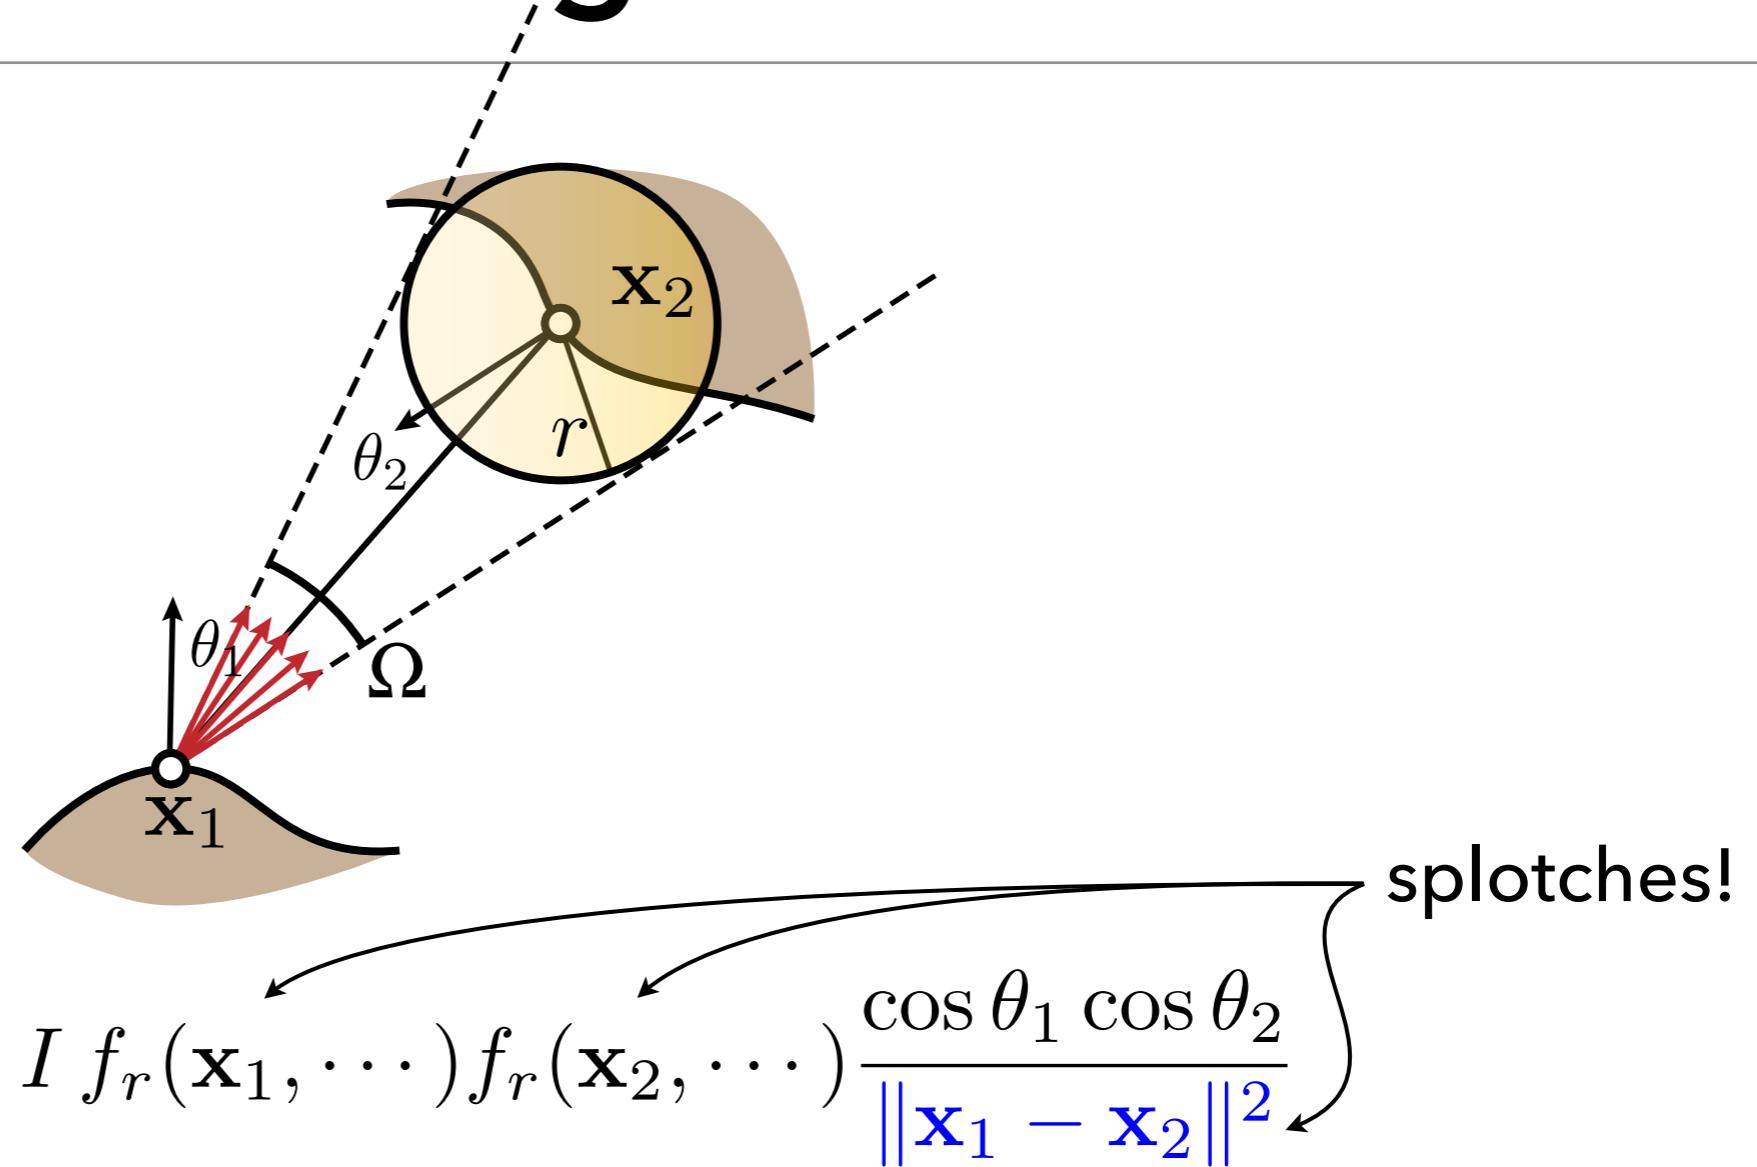
Finite-size Lights

- Solution 2: spread the energy spatially instead of concentrating it at single points
- Virtual *Spherical* Lights [Hašan et al. 2009]
 - energy distributed over surfaces inside a sphere
- Virtual *Ray/Beam* Lights [Novák et al. 2012ab]
 - energy is distributed along infinitesimal rays or beams with finite thickness [only for volumes]

Virtual Spherical Lights



Virtual Spherical Lights

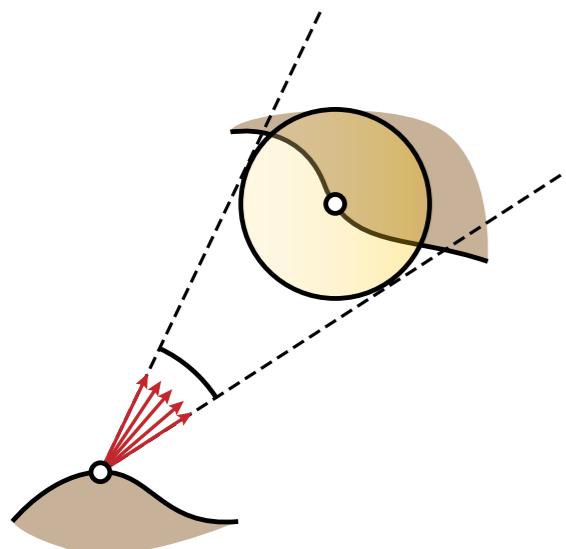


approx. disk-to-point: $\frac{I}{\pi r^2} \int_{\Omega} f_r(\mathbf{x}_1, \dots) f_r(\mathbf{x}_2, \dots) \cos \theta_1 \cos \theta_2 d\omega$

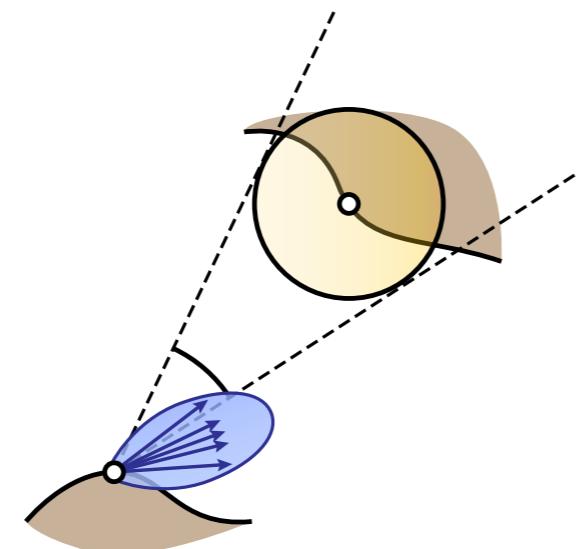
Surfaces inside the sphere are approximated by a disk facing \mathbf{x}_1

Virtual Spherical Lights

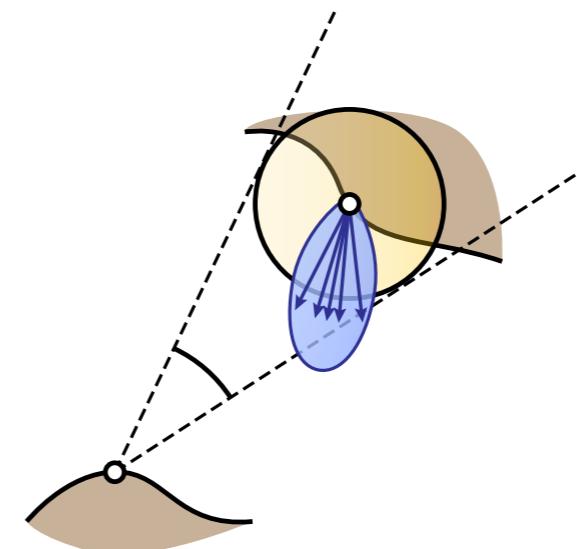
- Importance sampling the integral:



cone sampling

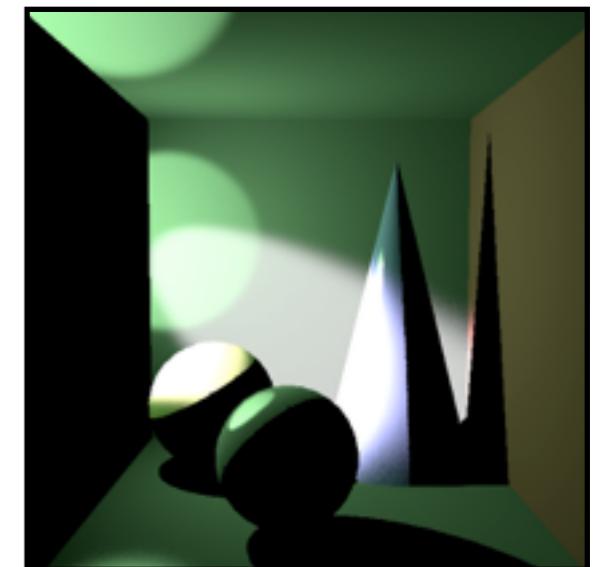
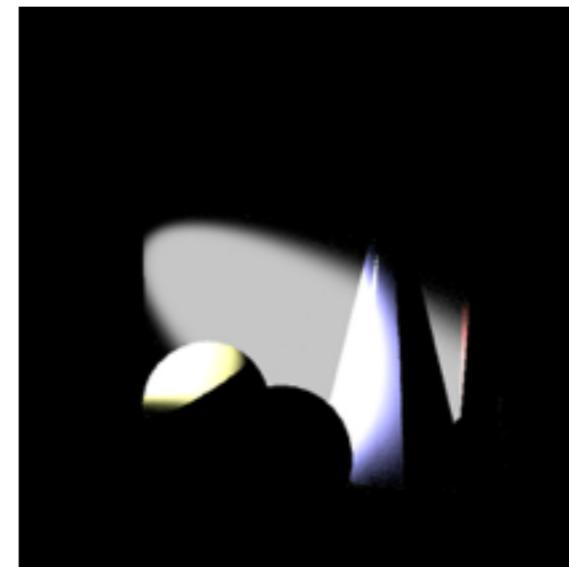
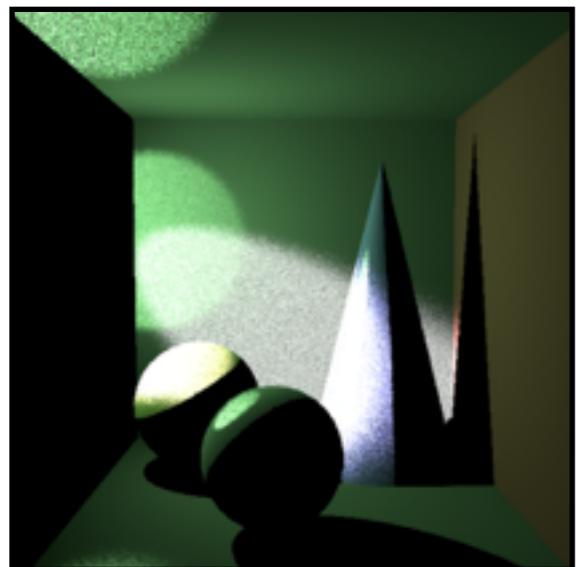


BRDF1 sampling



BRDF2 sampling

Multiple importance sampling

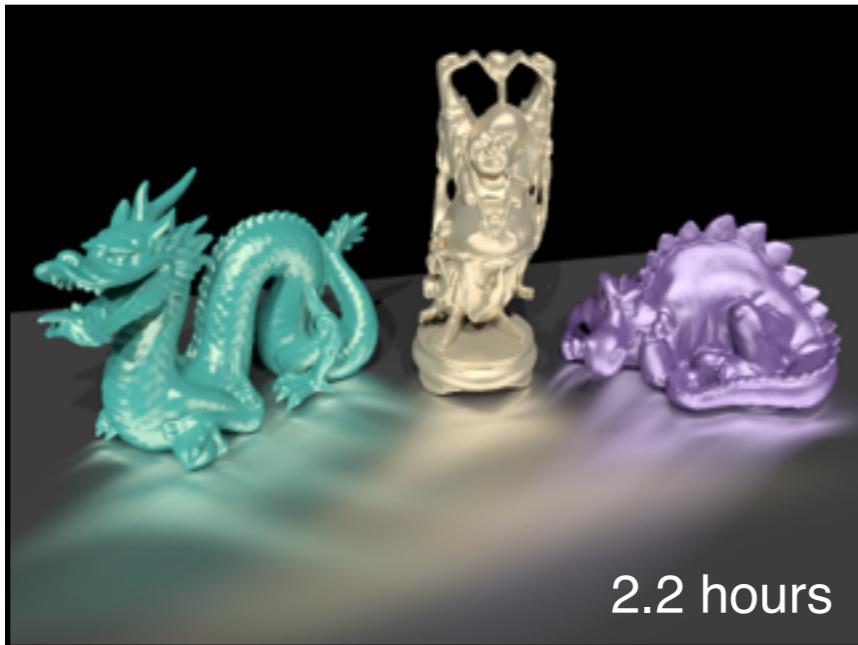


Virtual Spherical Lights

- Idea: spread the energy spatially instead of concentrating it at single points
- Advantages:
 - energy is blurred, not clamped
- Disadvantages:
 - blurring of energy introduces (another kind of) bias
 - requires an integral over the solid angle

Virtual Spherical Lights

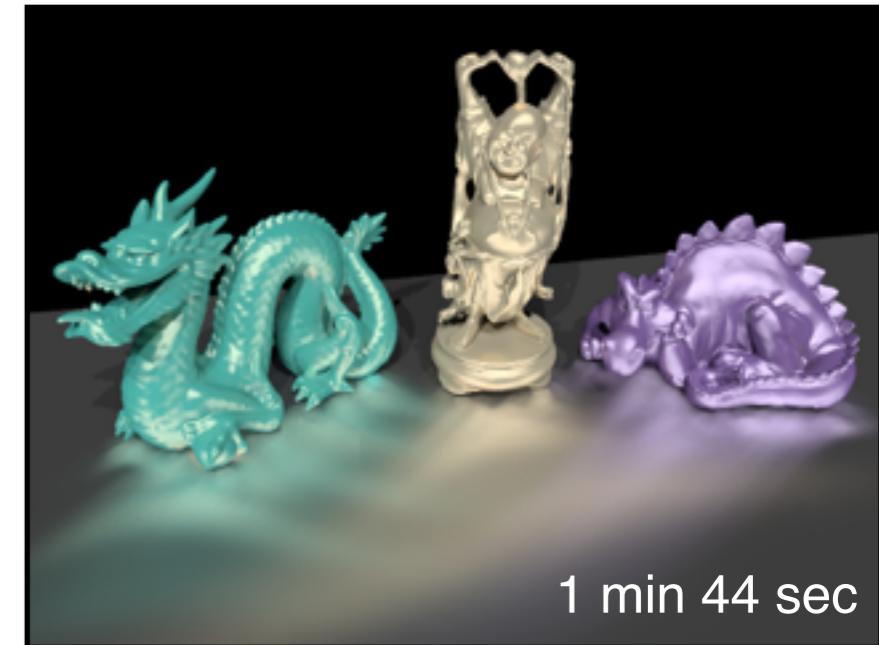
Reference



Bounded



VSLs

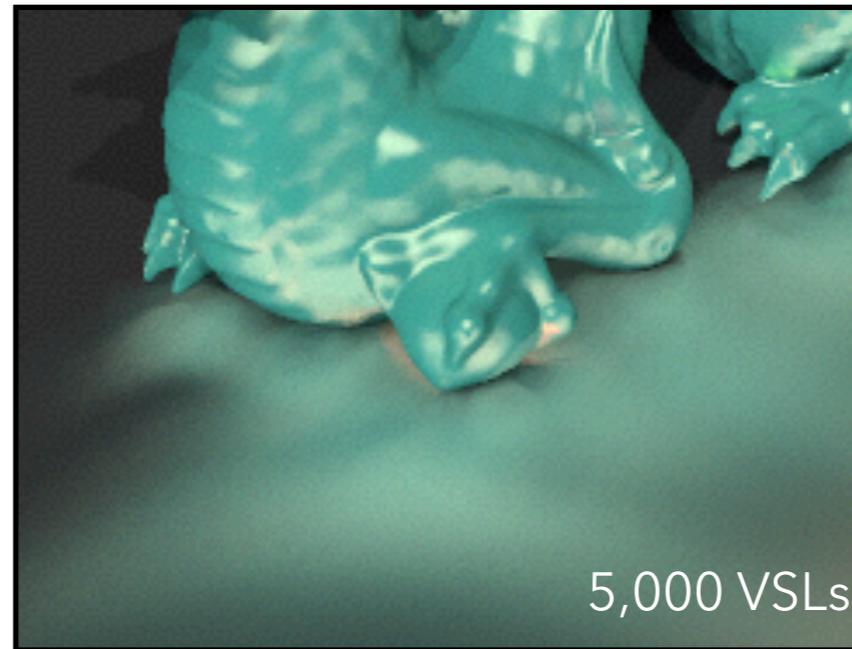


Virtual Spherical Lights

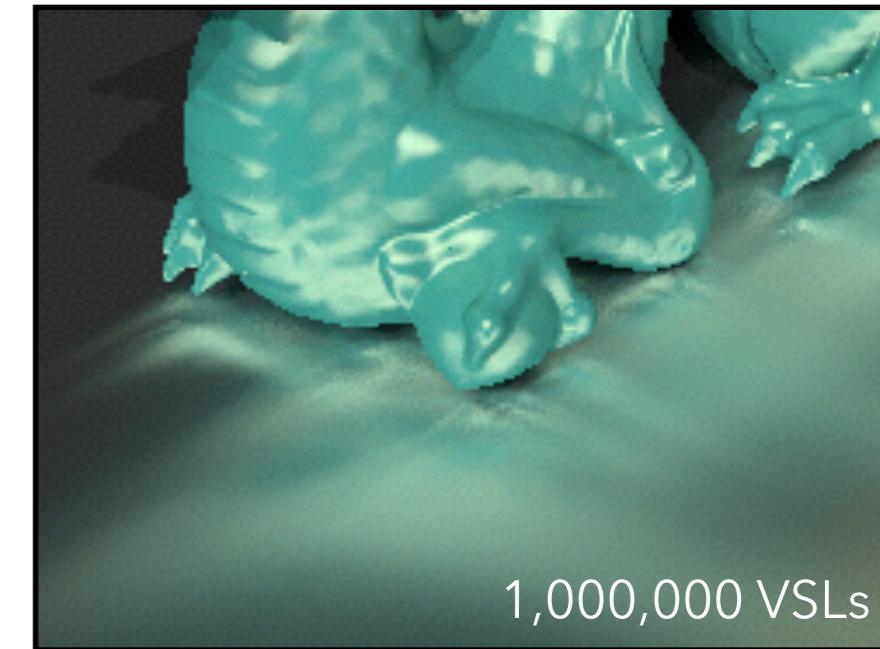
Reference



VSLs



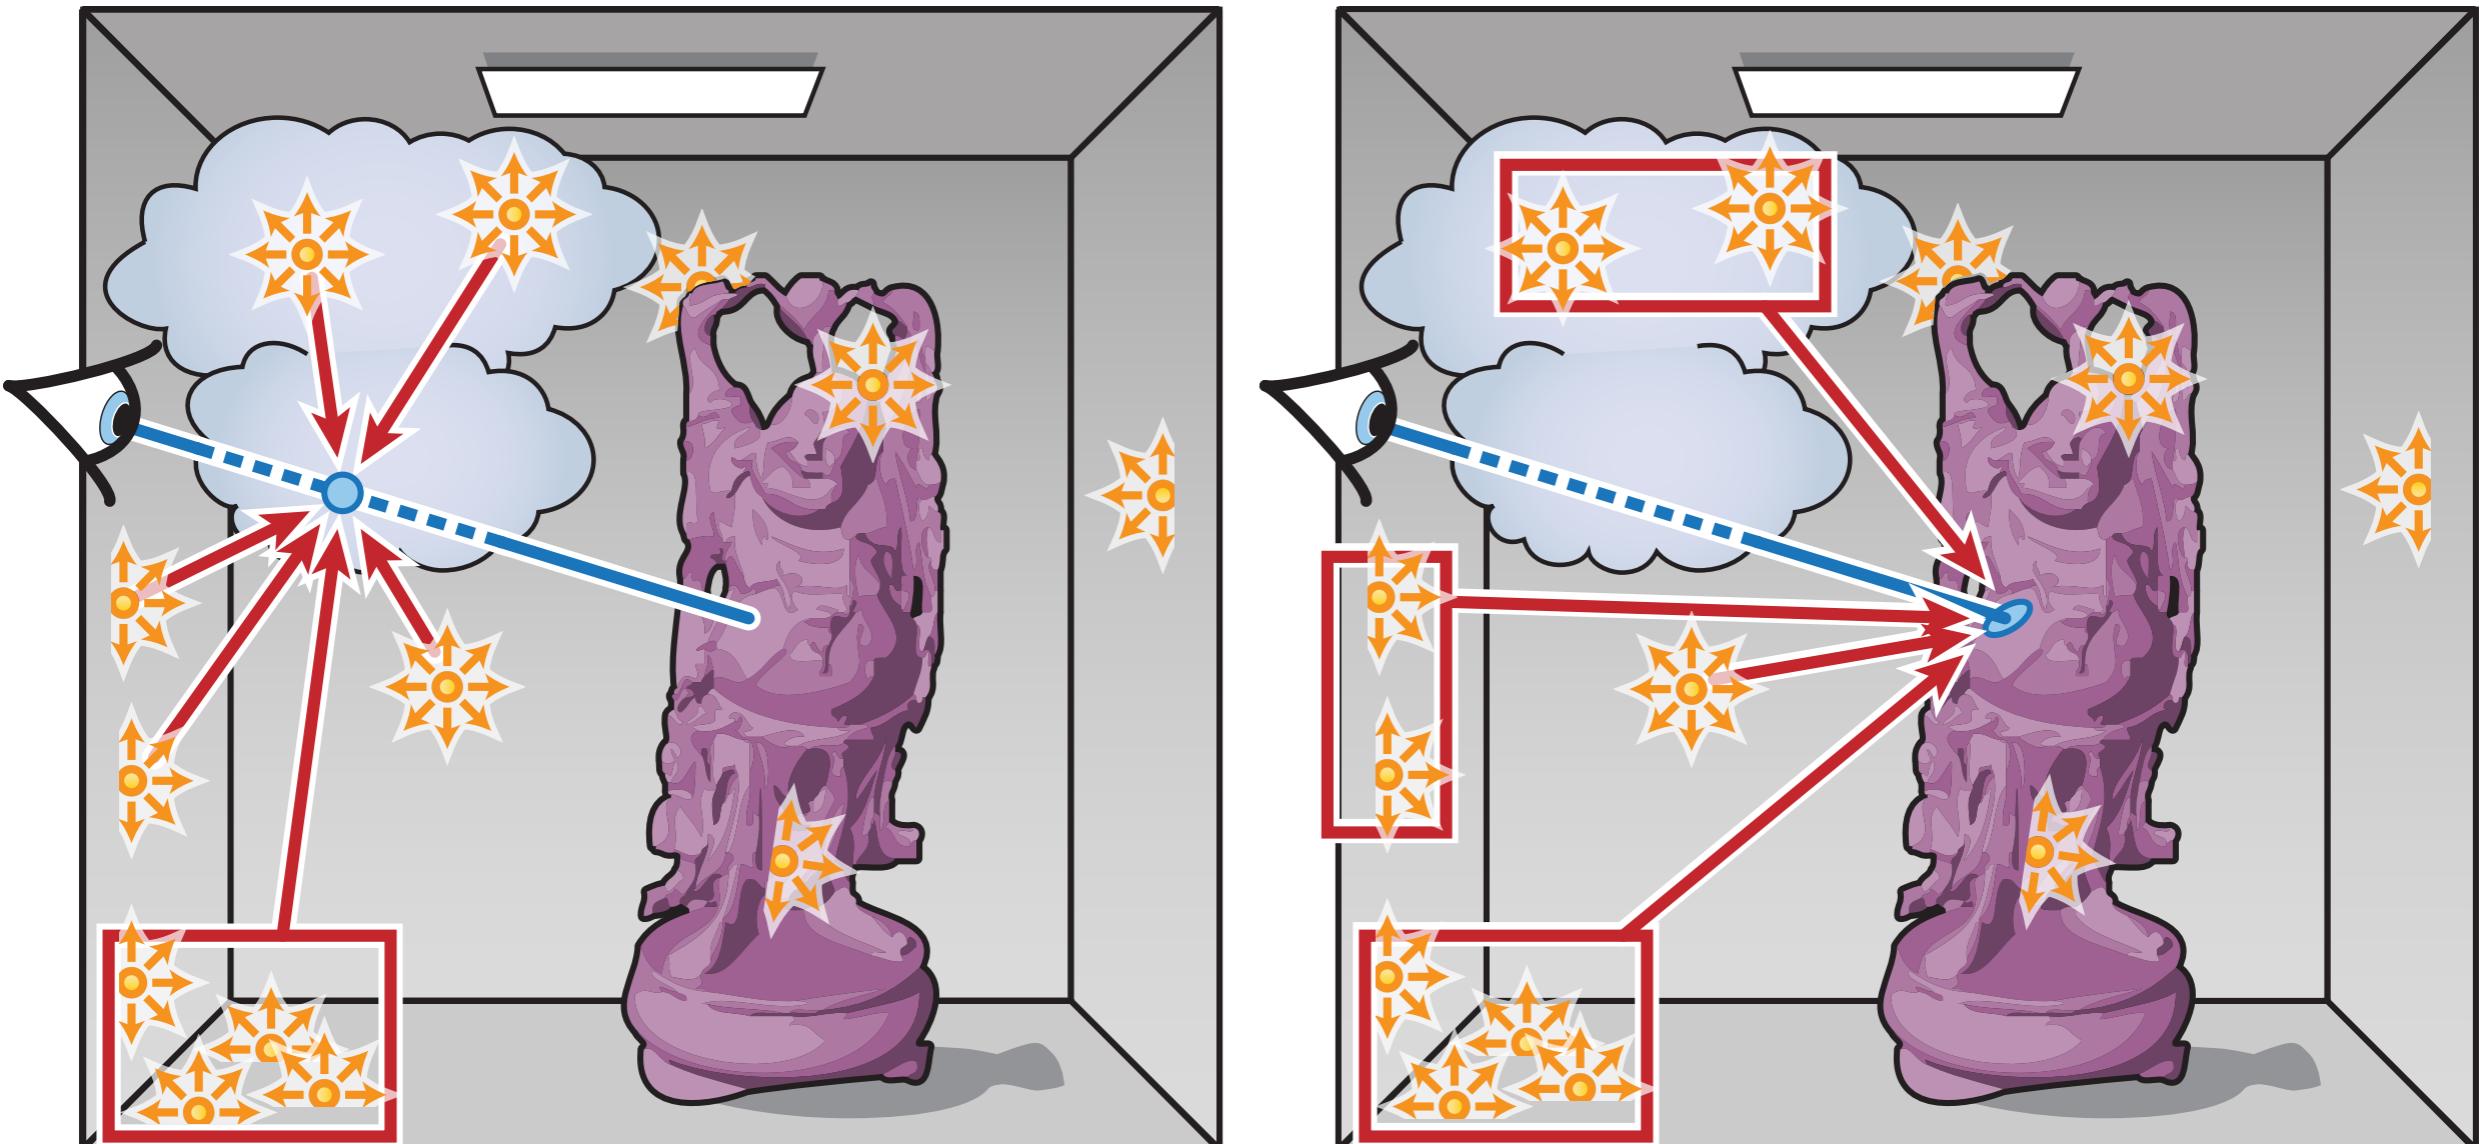
VSLs converged



Linear cost = too expensive!

Scalability

- Goal: sub-linear evaluation cost

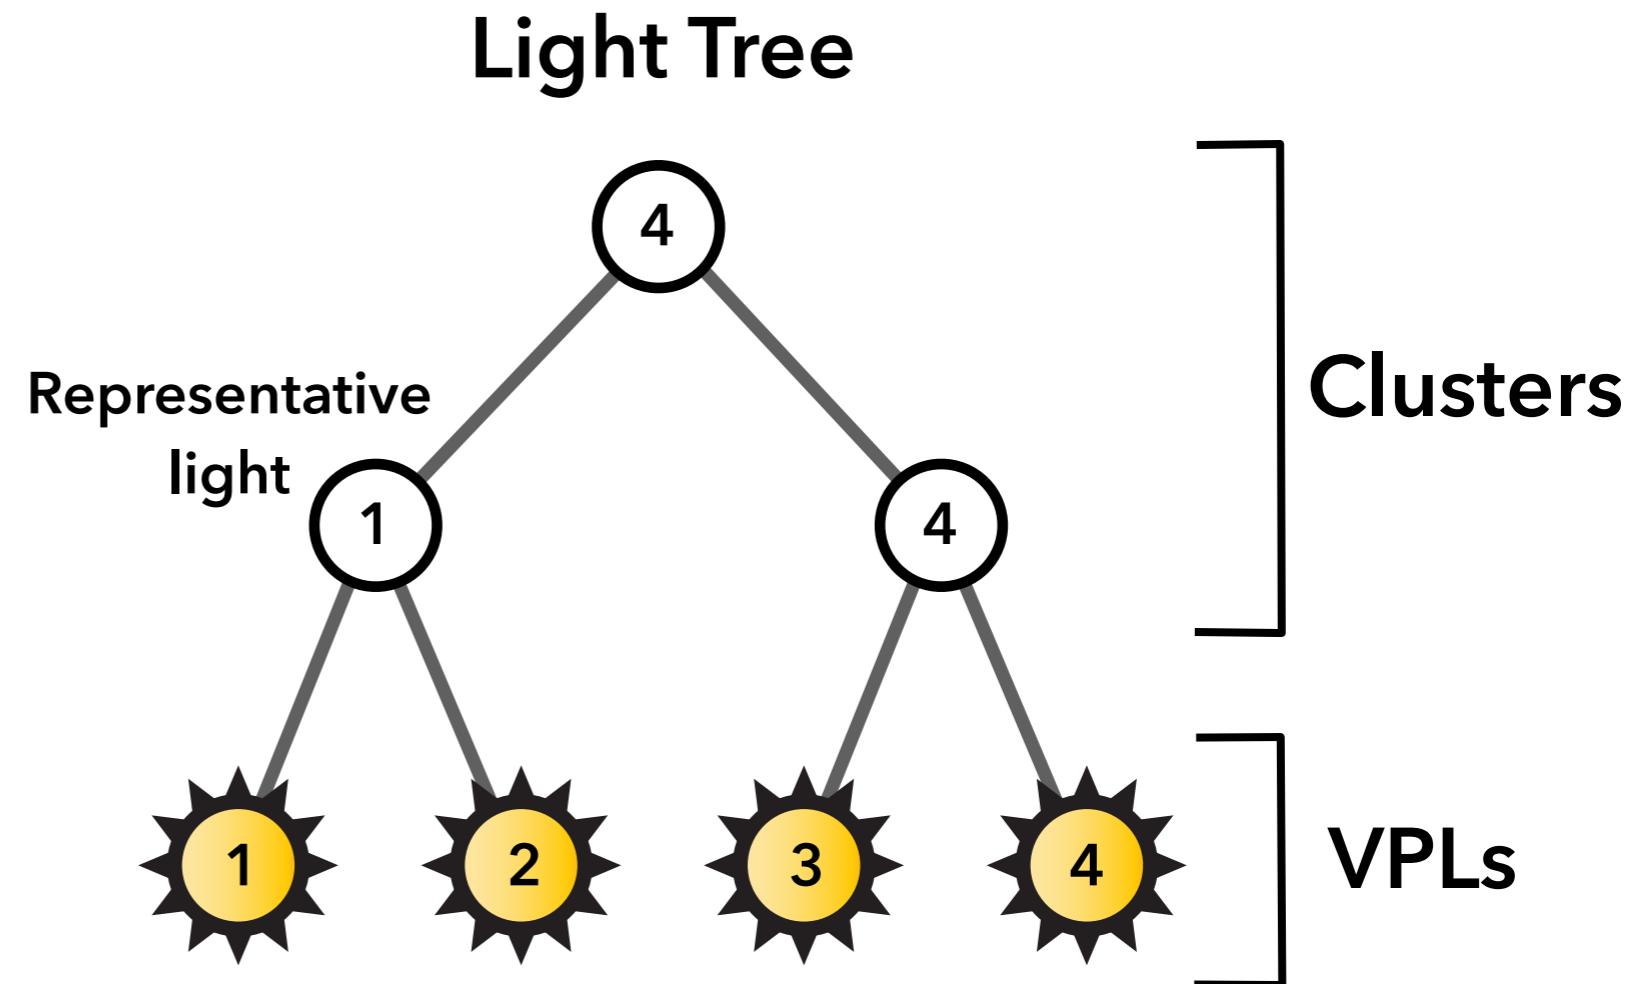
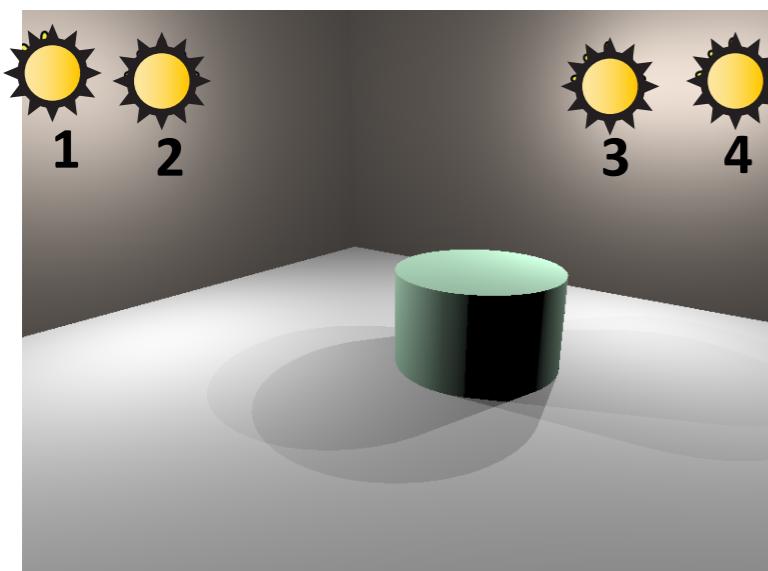


Scalability

- Goal: sub-linear evaluation cost
- Clustering of virtual lights:
 - Lightcuts [Walter et al. 2006]
 - Matrix Row-Column Sampling [Hašan et al. 2007]
 - Lightslice [Ou and Pellacini 2011]
- Importance sampling VPLs [Georgiev et al. 2012]

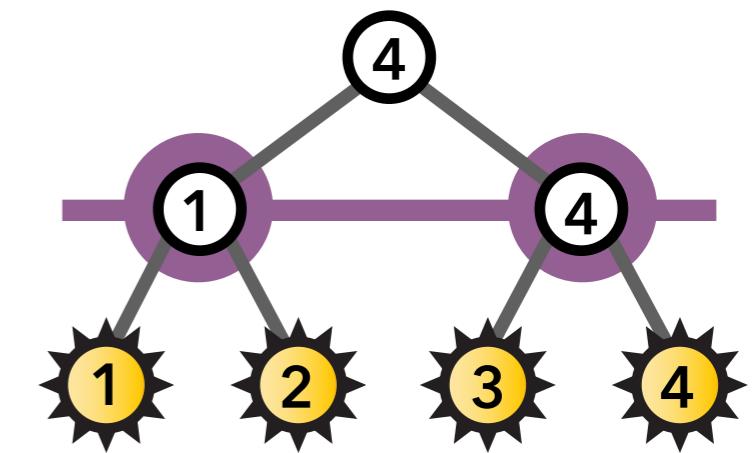
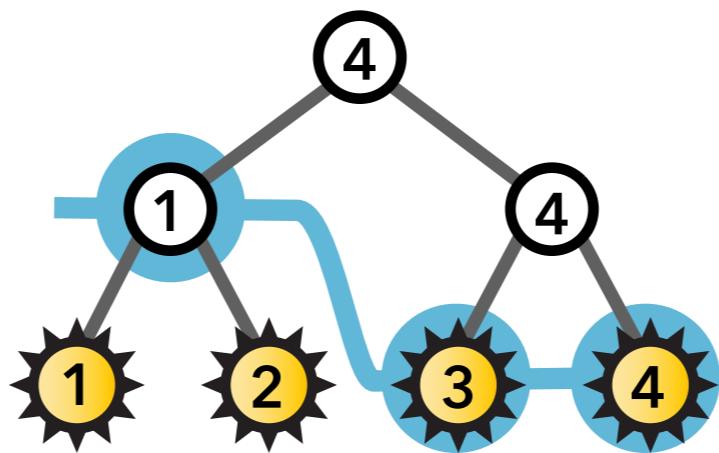
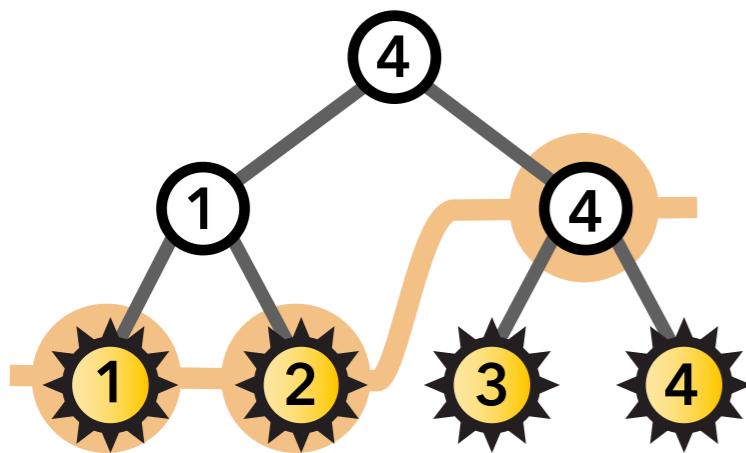
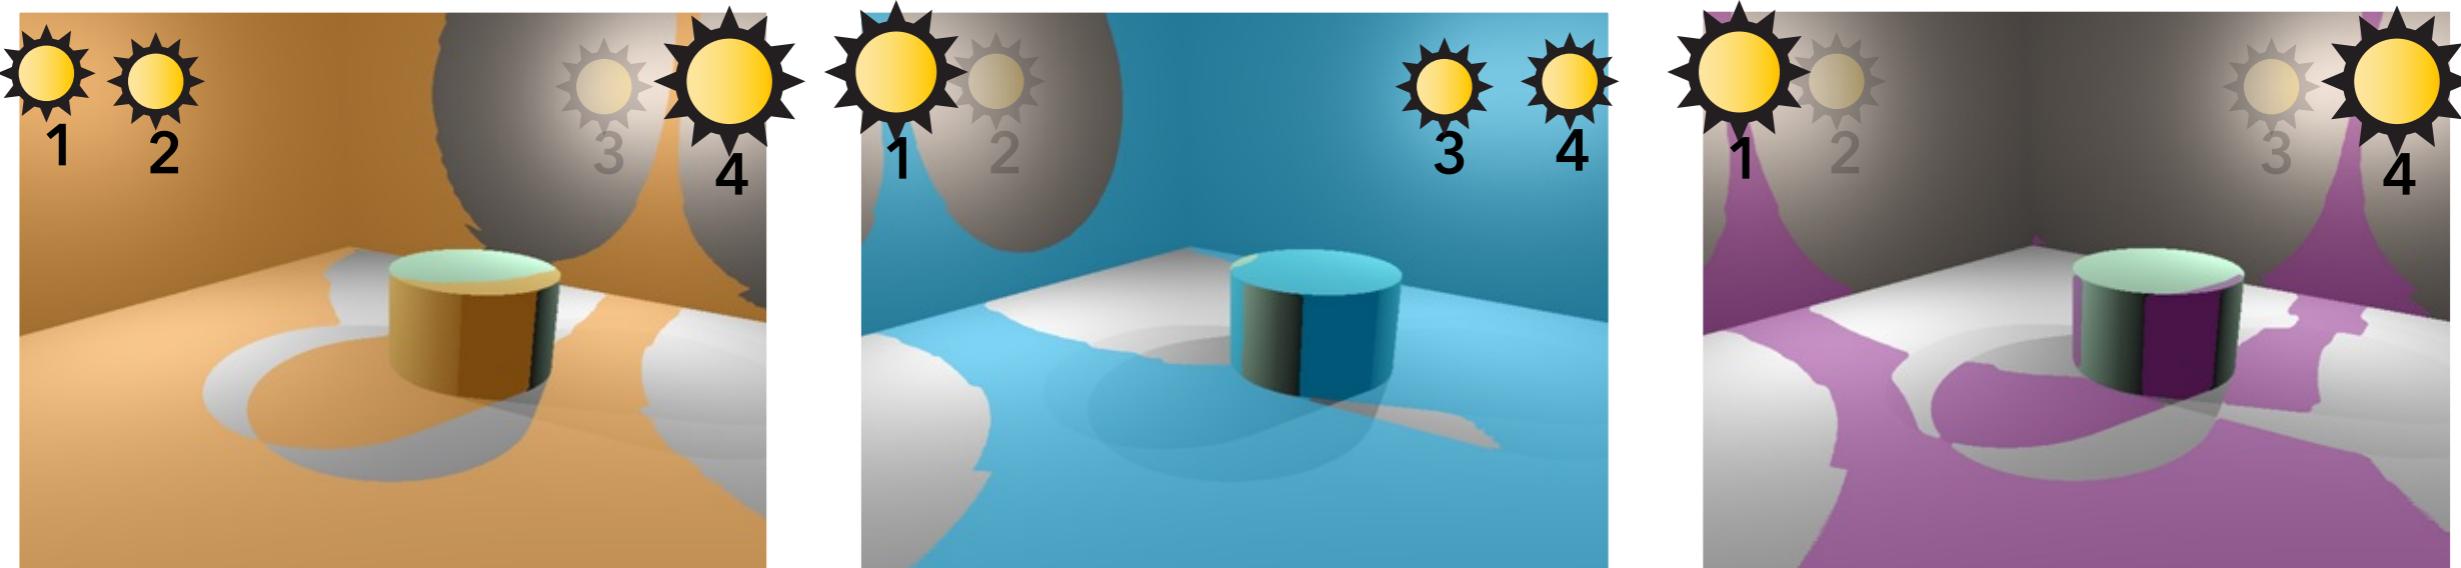
Lightcuts

- Preprocess: build a *tree hierarchy*
- Rendering: adaptive clustering of VPLs



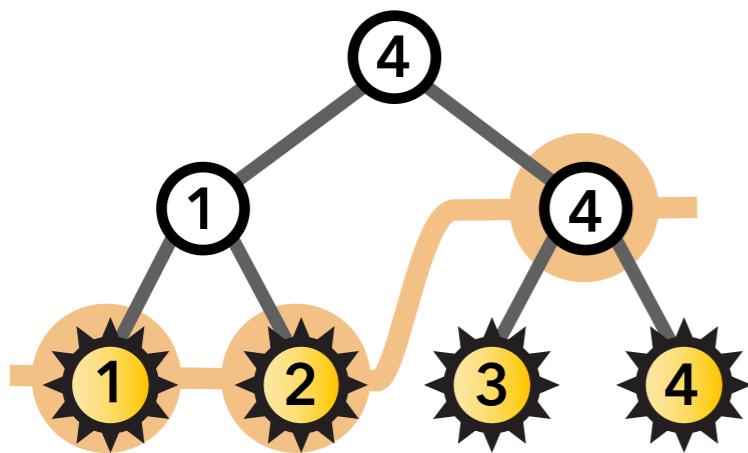
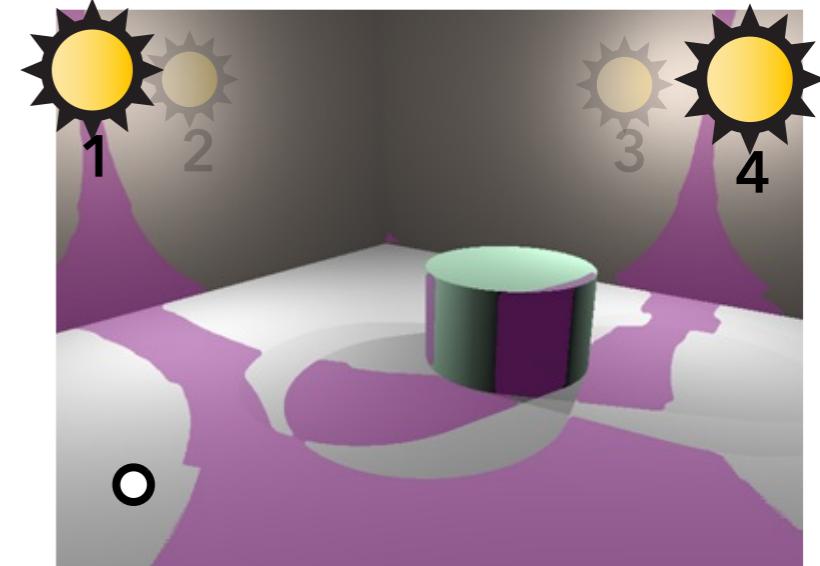
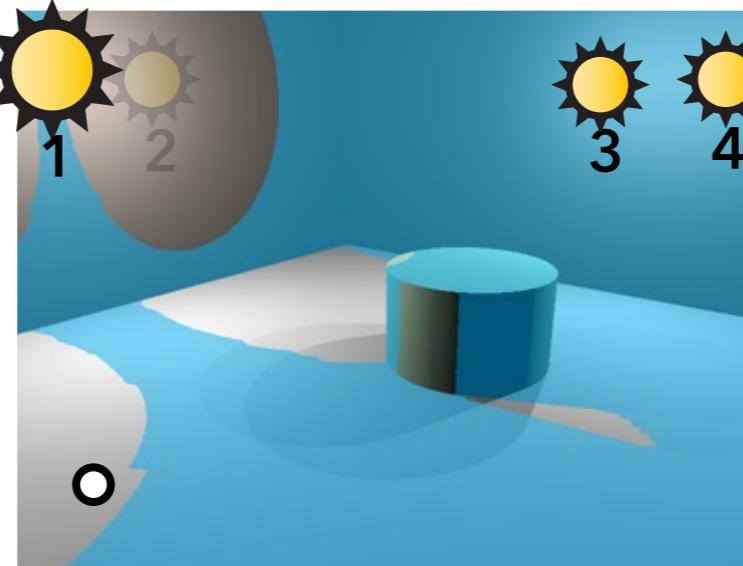
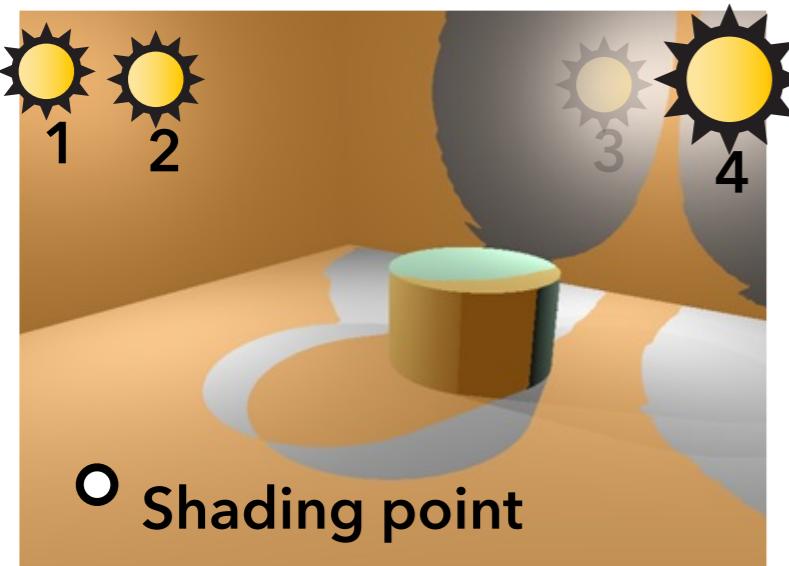
Lightcuts

Three example cuts

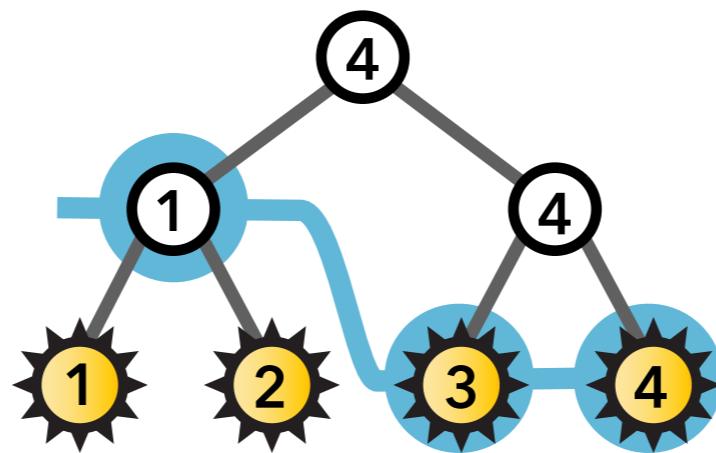


Lightcuts

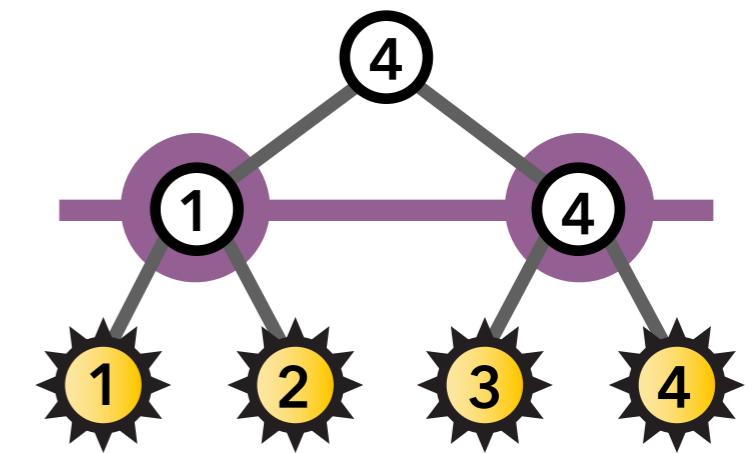
Three example cuts



GOOD



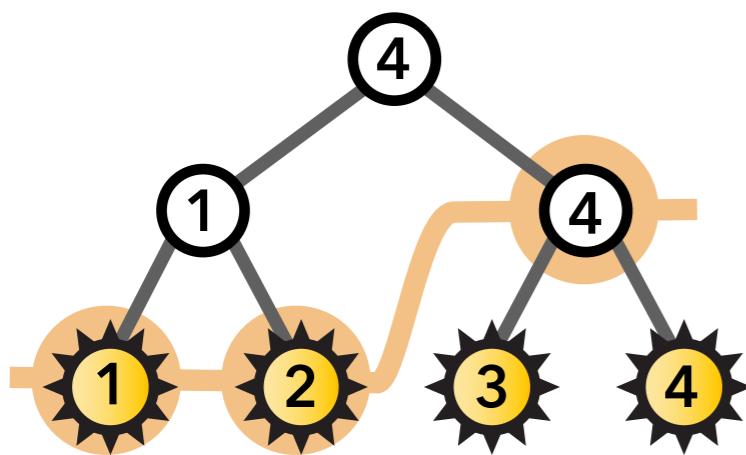
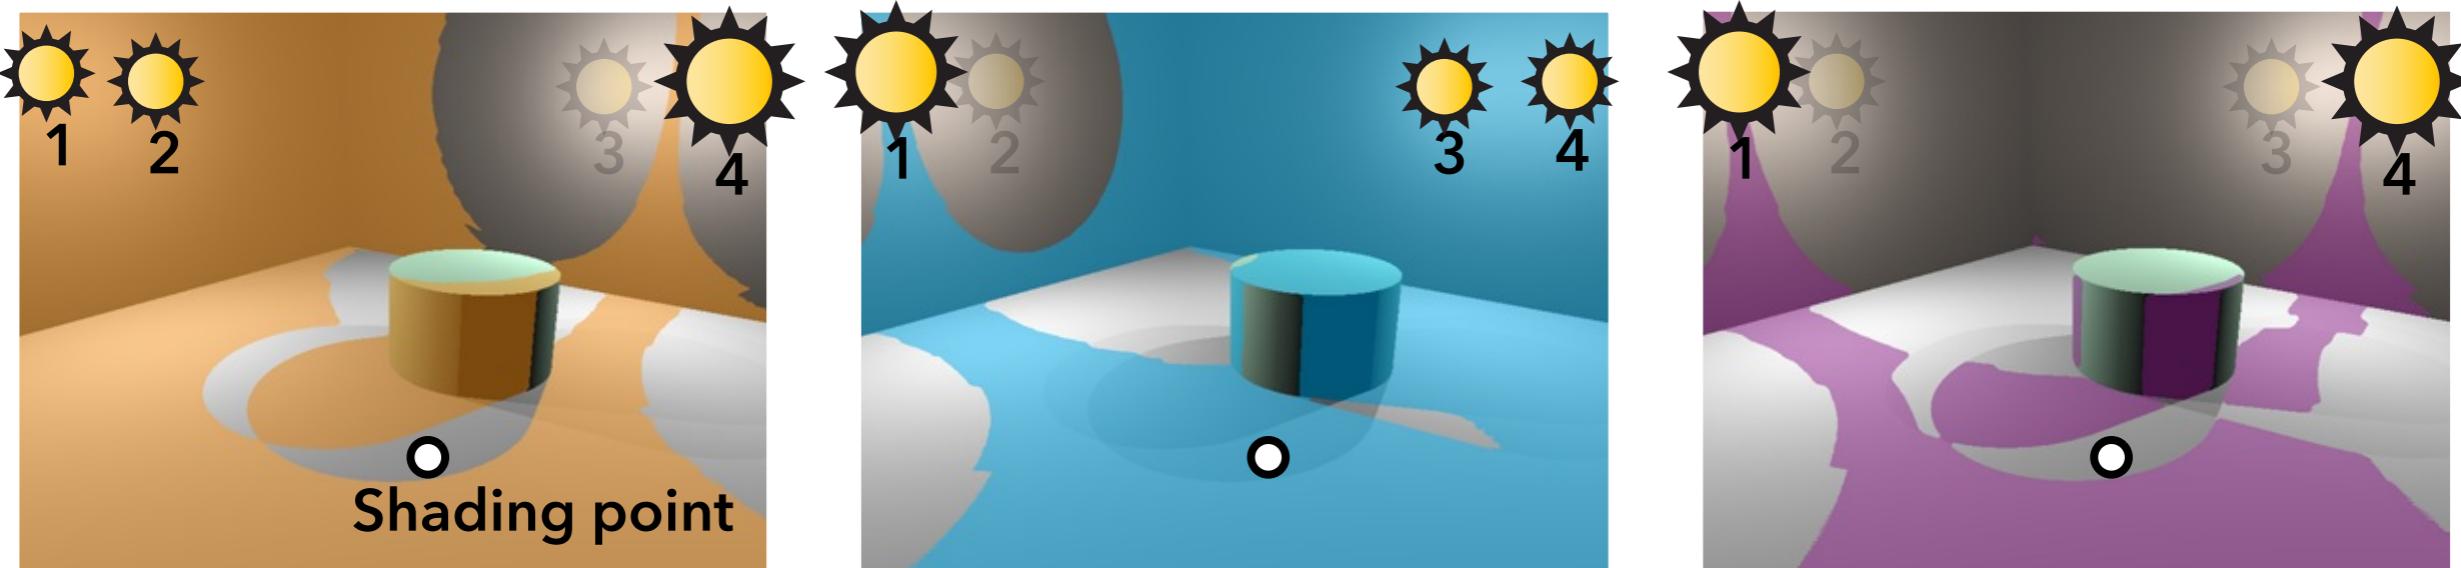
BAD



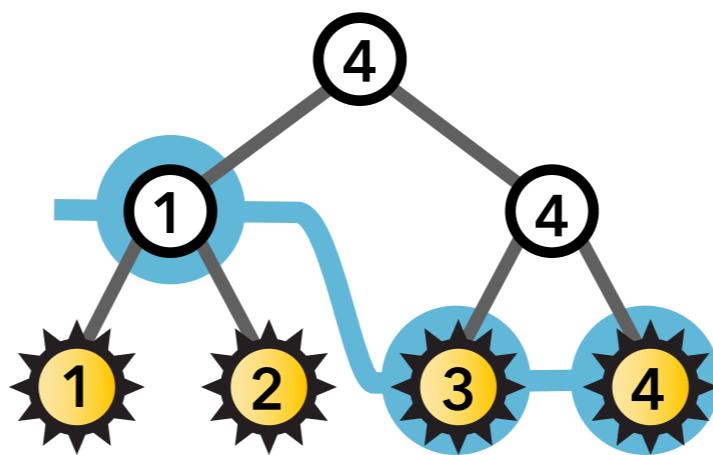
BAD

Lightcuts

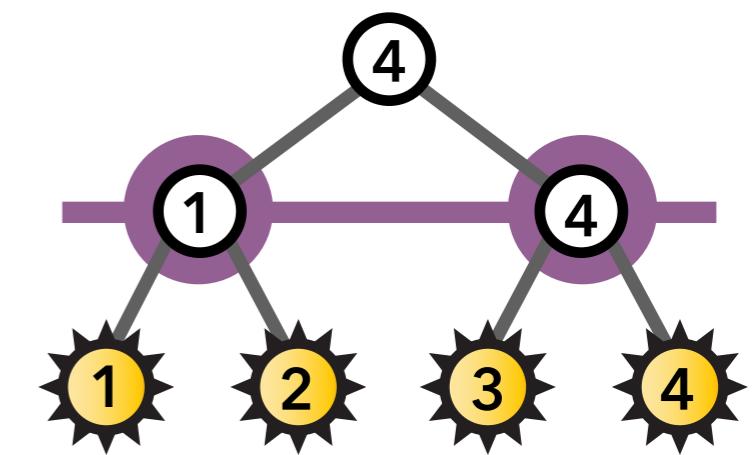
Three example cuts



BAD



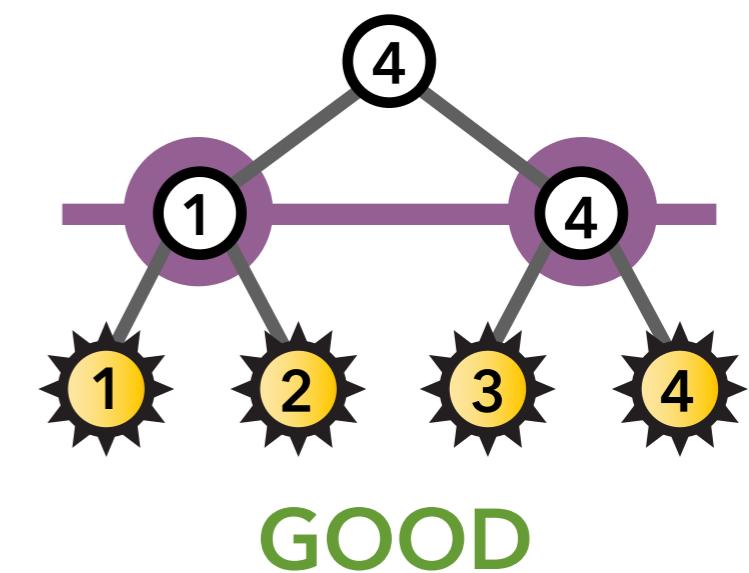
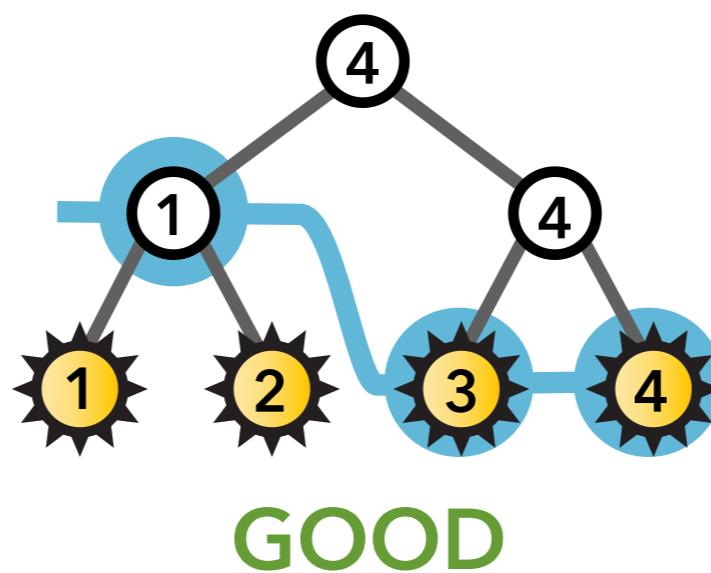
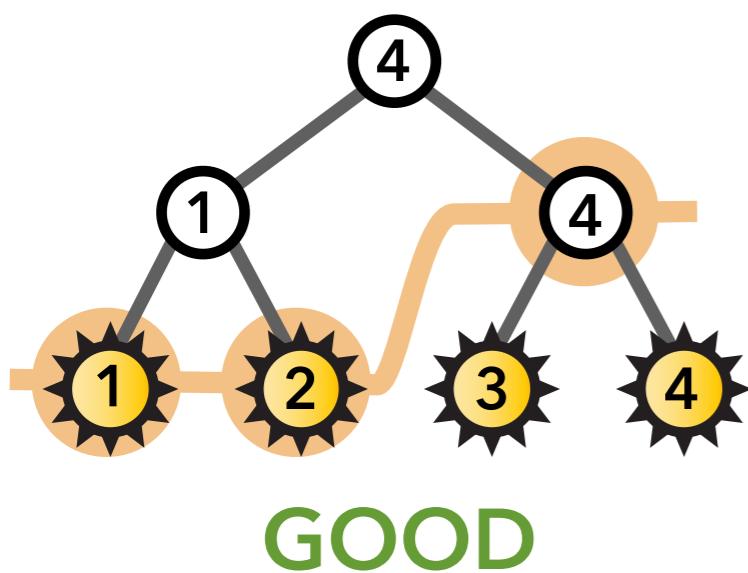
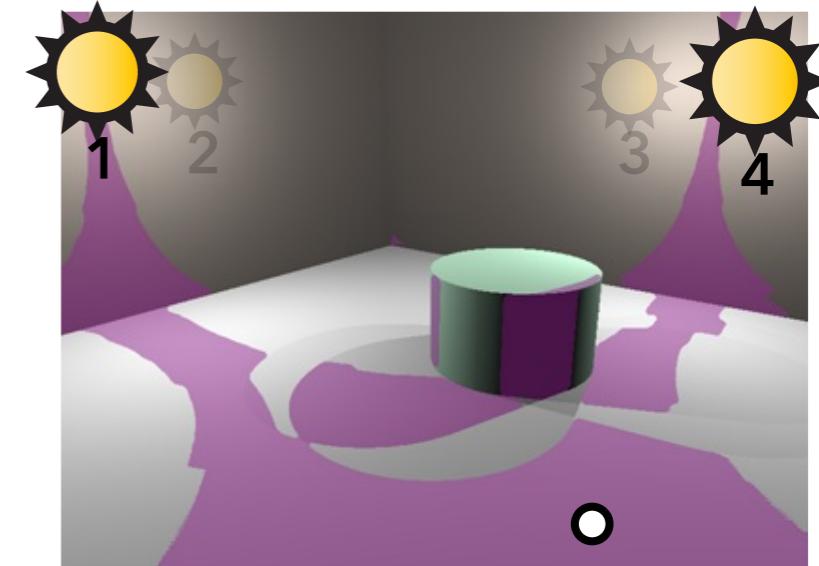
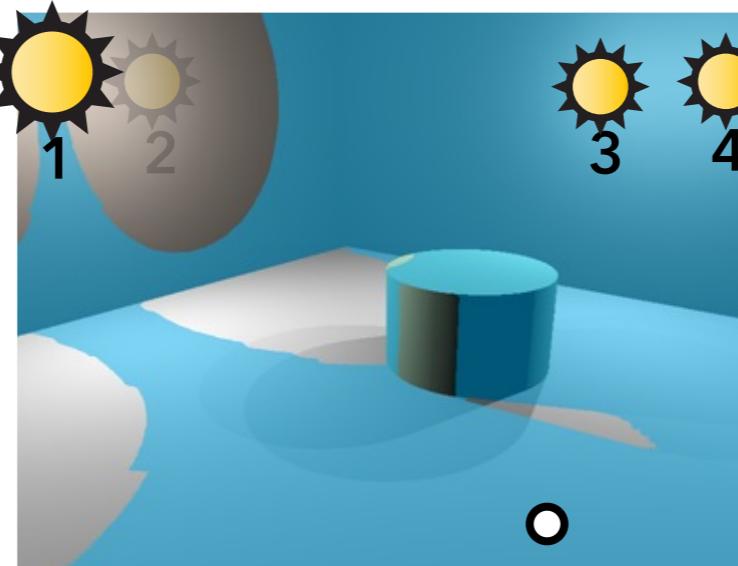
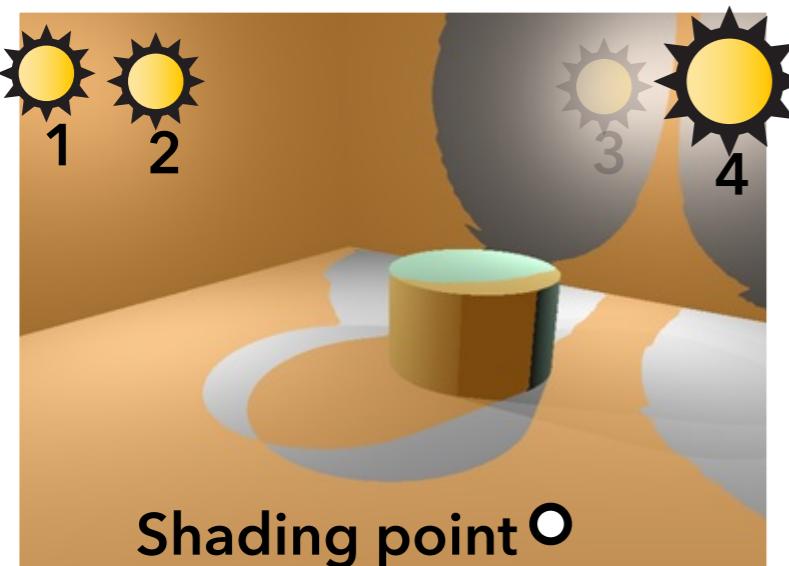
GOOD



BAD

Lightcuts

Three example cuts

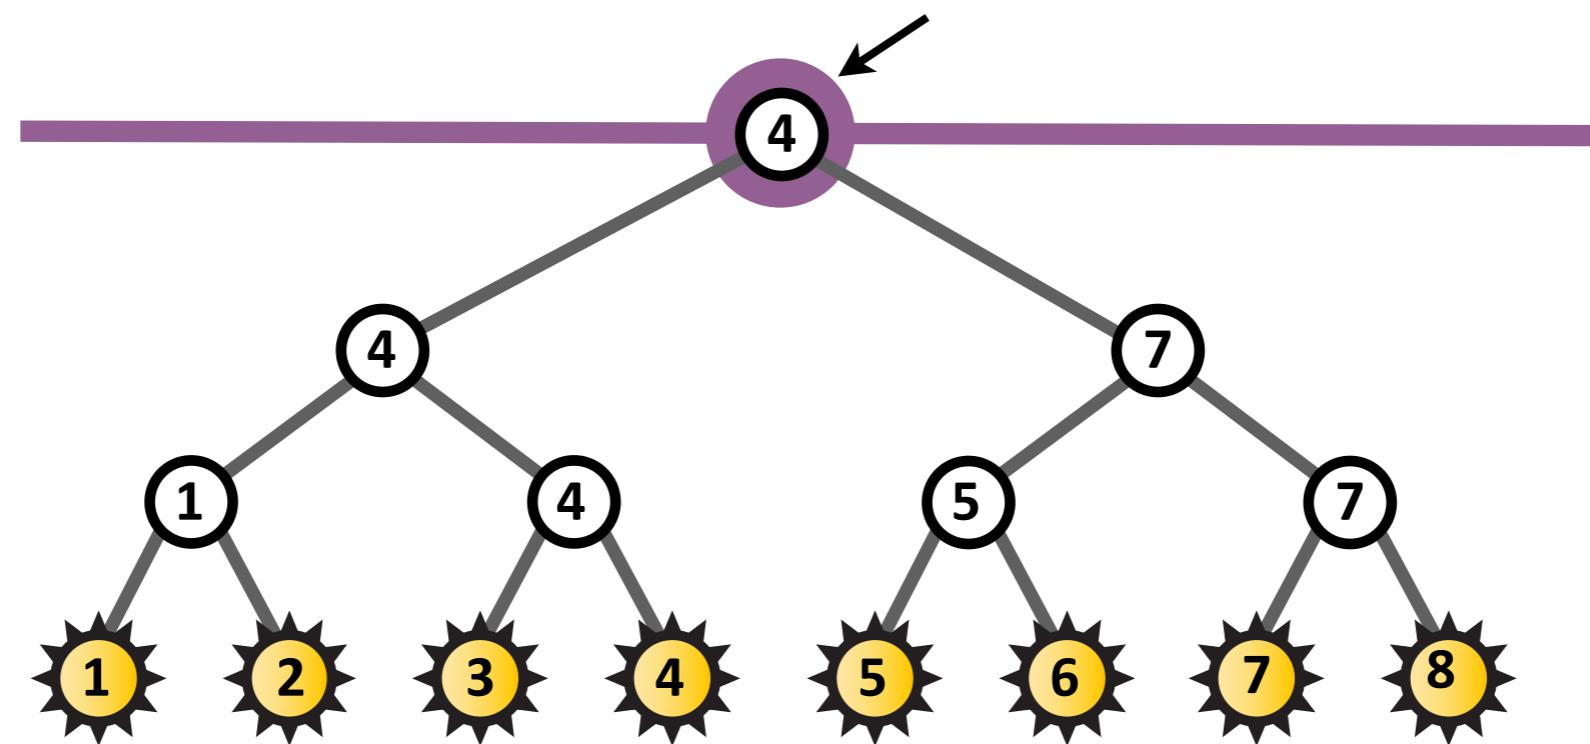


Lightcuts

- Preprocess
 - convert illumination to points
 - build a light tree
- Rendering:
 - choose a cut for each shading point
 - avoid visible artifacts (bound maximum predicted error)
 - refine the cut if necessary

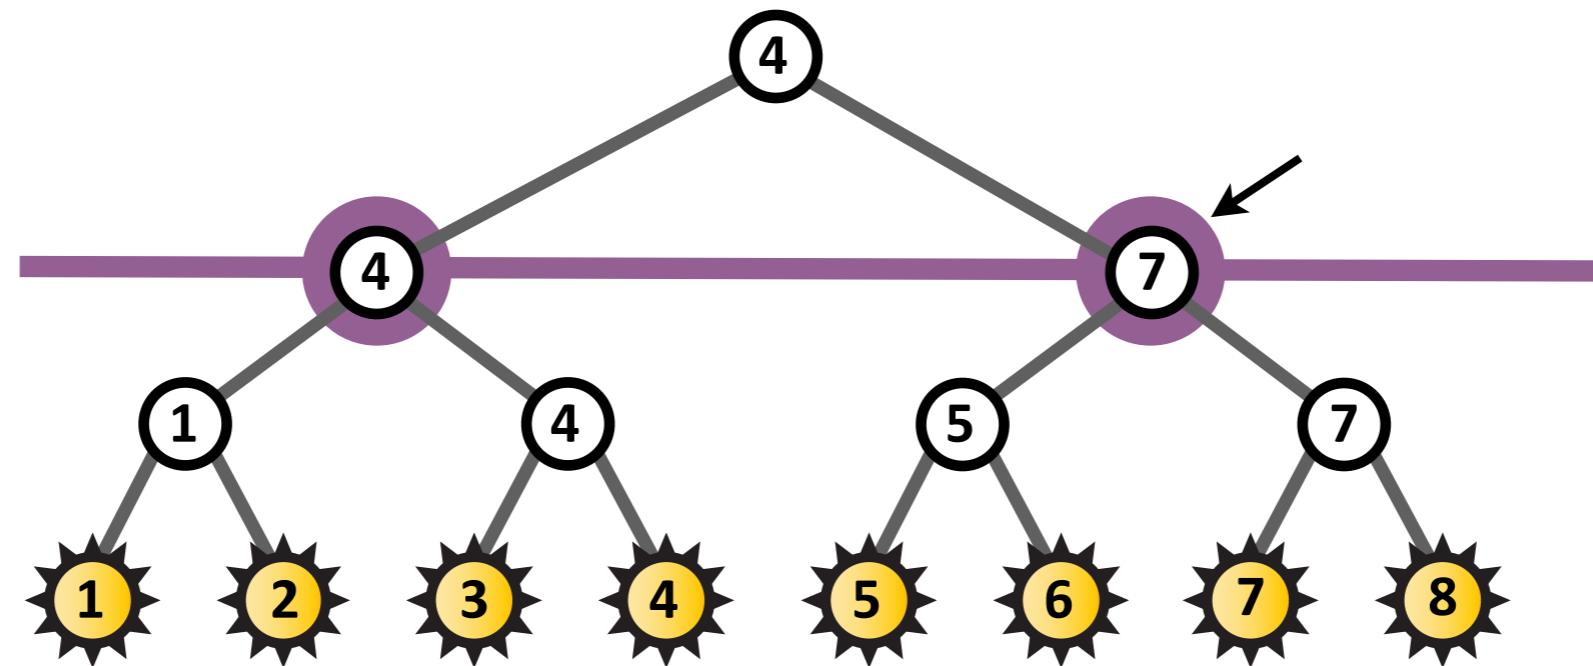
Lightcuts

Building a cut



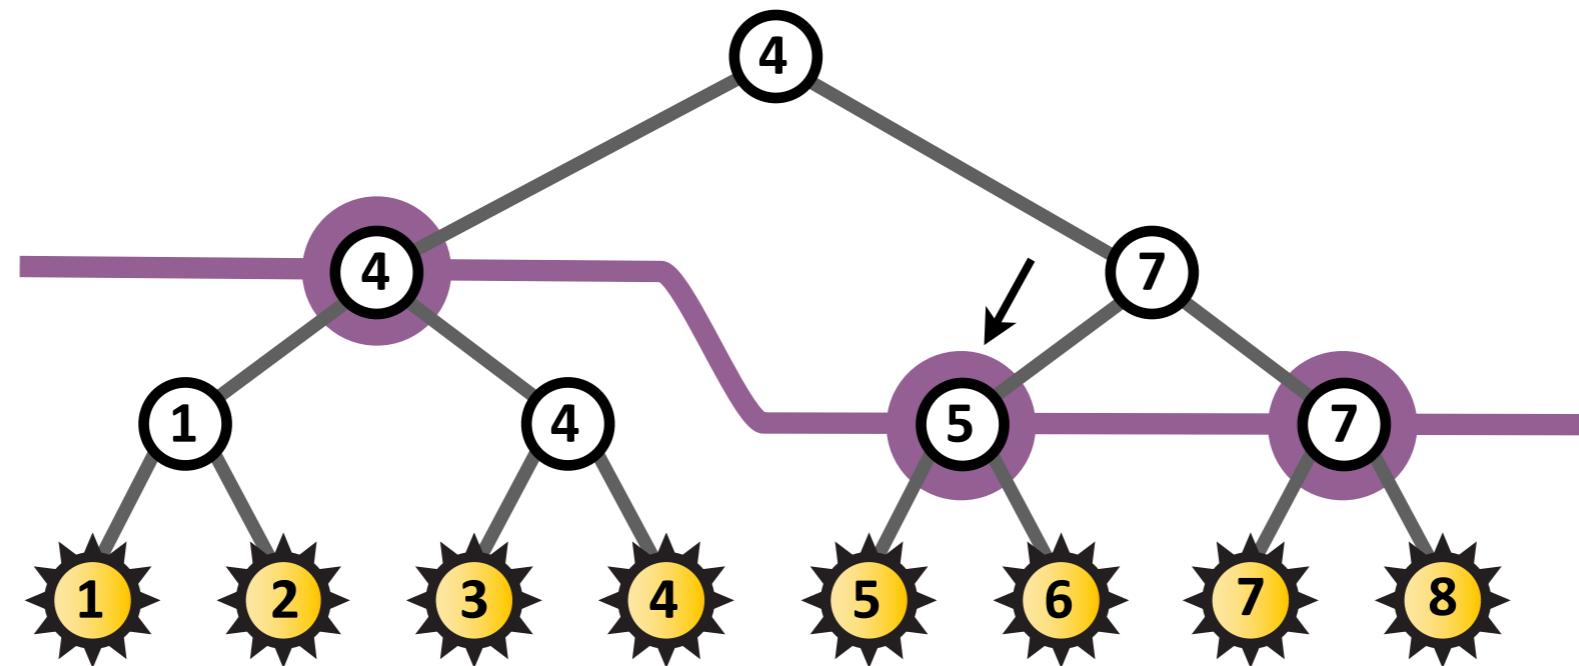
Lightcuts

Building a cut



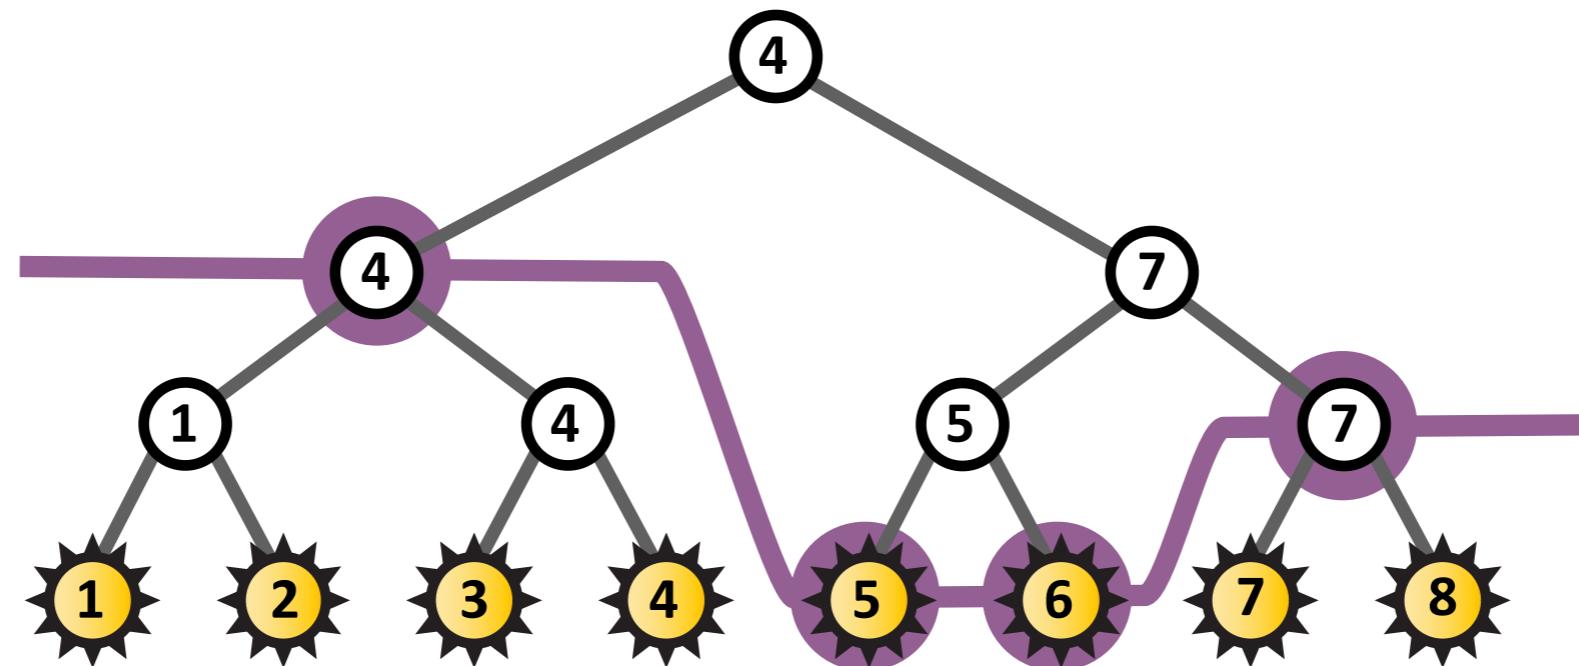
Lightcuts

Building a cut

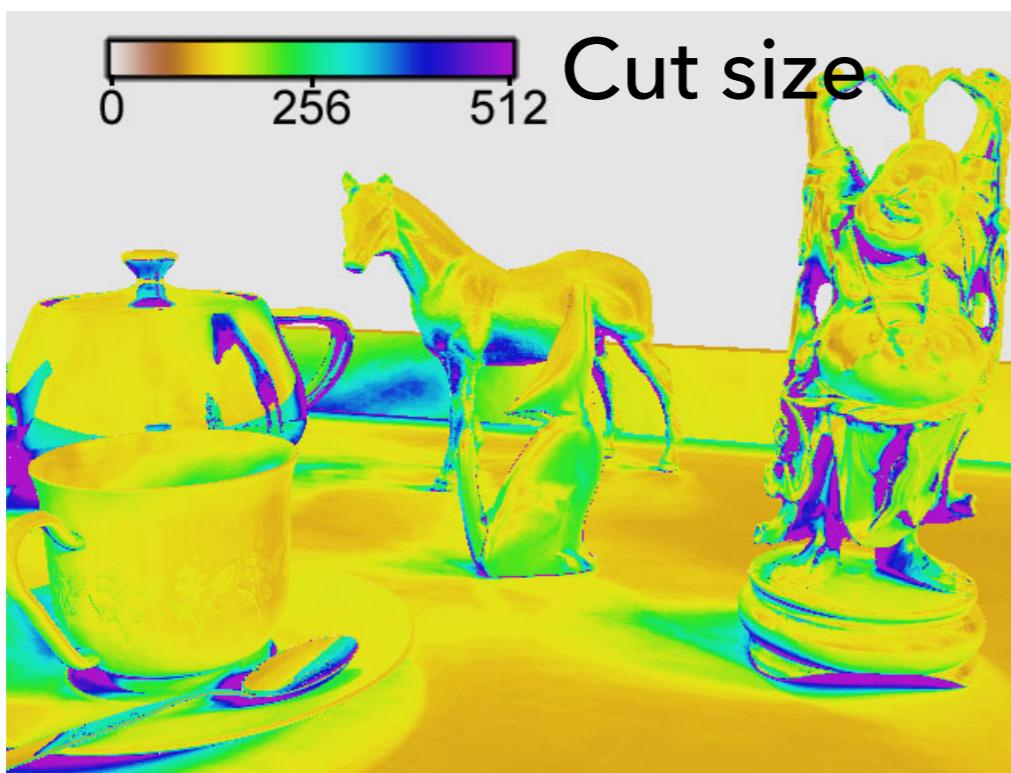


Lightcuts

Building a cut

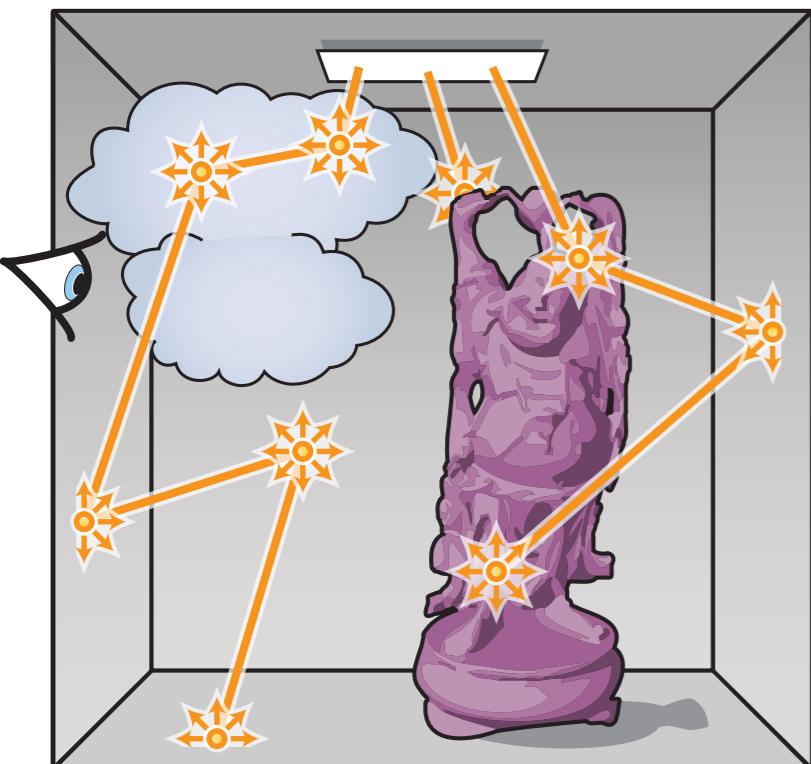


Lightcuts

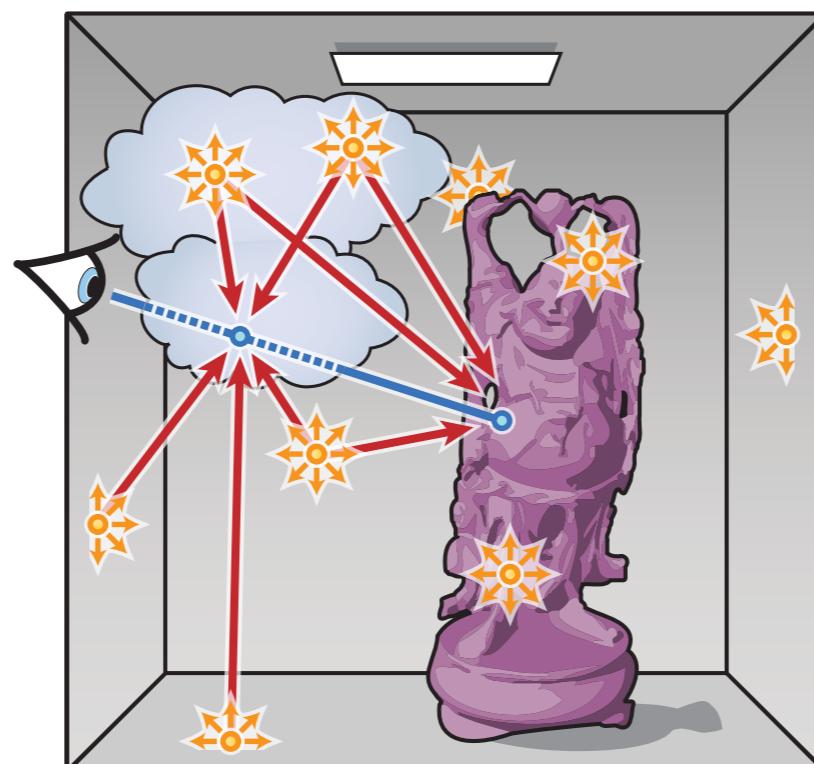


Many-Light Rendering

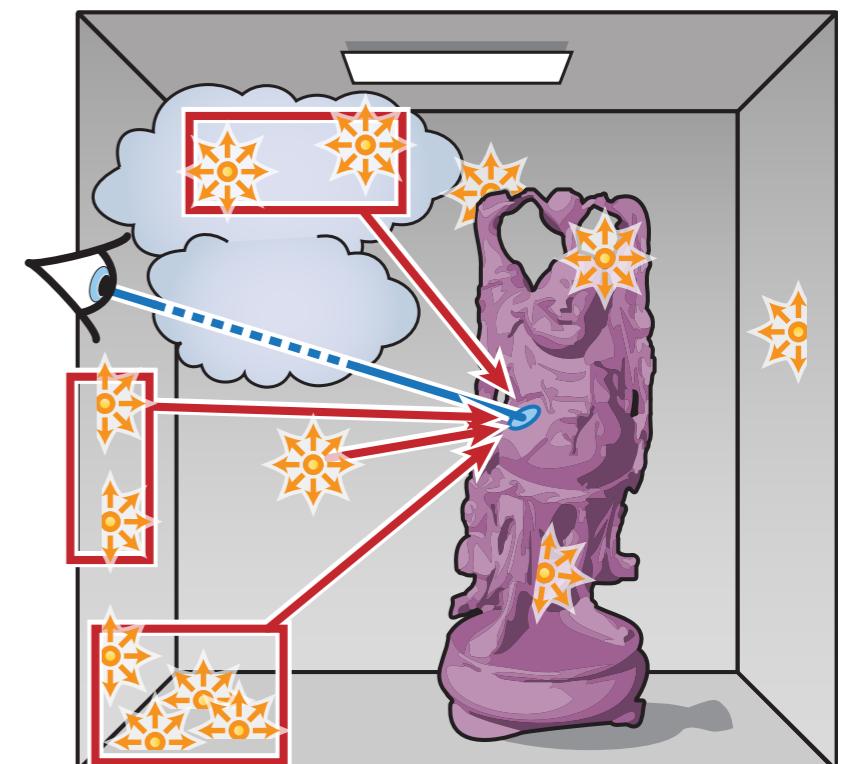
Generation of VPLs



Lighting with



Clustering of



Many-Light Rendering: Summary

- Advantages
 - basic method is easy to implement
 - great scalability
 - low amount of noise
 - GPU friendly, can be accelerated using shadow maps
 - popular in real-time, interactive, and fast offline rendering
- Disadvantages
 - horrible temporal stability shows up as flickering
 - hard to capture high-frequency transport (e.g. glossy reflections)

Reference	Abbreviated Paper Title	Category / Goal	Speed	Materials	Clustering	Remark
[Kel97]	Instant Radiosity	diffuse GI	real-time	D	no	basic method
[DS05]	Reflective Shadow Maps	single-bounce GI	real-time	D	no	no VPL visibility, importance sampling of VPLs
[DS06]	Splatting Indirect Illumination	single-bounce GI	real-time	D,(G)	no	no VPL visibility, importance sampling of VPLs
[LLK07]	Incremental Instant Radiosity	single-bounce GI	real-time	D	no	reuse VPLs over frames, for static scenes only
[NW09,10]	Multiresolution Splatting	single-bounce GI	real-time	D,(G)	yes	hierarchical shading, no VPL visibility
[DGRS09]	Clustered Visibility	reduce banding	real-time	D,(G)	k-means	reduce banding with virtual area lights
[PKD12]	Reflective Shadow Map Clustering	single-bounce, reduce banding	real-time	D	k-means	virtual area lights with no visibility computation
[KK04]	Illumination Presence of Weak Singularities	bias compensation	offline	D,G	no	bias compensation using final gathering
[NED11]	Screen-Space Bias Compensation	bias compensation	real-time	D,G	no	image-space technique for approximate bias compensation
[RSK08]	Unbiased GI with Participating Media	bias compensation	offline	PM	no	basic bias compensation technique for media
[ENSD12]	Approximate Bias Compensation	bias compensation	offline/interactive	PM	no	approximate bias compensation for participating media
[WKB12]	Bidirectional lightcuts	bias compensation	offline	D,G,SSS,PM	yes	short-range indirect illumination, uses Multidimensional Lightcuts
[HKWB09]	Virtual Spherical Lights	avoid singularities	offline	D,G	yes	inflate VPLs to spherical lights, uses MRCS for clustering
[DKH*10]	Combining Global and Local Virtual Lights	reduce singularities	offline	D,G	yes	local lights for short-range indirect illumination
[NNDJ12b]	Virtual Ray Lights	reduce singularities	offline	D,G,PM	no	use light path segments as virtual ray lights
[NNDJ12a]	Progressive Virtual Beam Lights	avoid singularities	offline	D,G,PM	no	inflate virtual ray lights
[HPB07]	Matrix Row-Column Sampling	scalability	precompute	D,G	yes	compute a global clustering of VPLs
[HVAPB08]	Tensor Clustering for Many-Light Animations	scalability	precompute	D,G	yes	compute a global clustering of VPLs, support for animations
[OP11]	Lightslice: Matrix Slice Sampling	scalability	precompute	D,G	yes	localized/clustered cuts
[WFA*05]	Lightcuts	scalability	offline	D,G	yes	per-shading point local cut of VPL hierarchy
[WABG06]	Multidimensional Lightcuts	scalability	offline	D,G,PM	yes	cut of VPLs and sensor points for depth of field, motion blur
[DGS12]	Progressive Lightcuts for GPU	scalability	offline/interactive	D,G	yes	progressive, GPU-friendly variant of [WFA*05]
[SIMP06a]	Bidirectional Instant Radiosity	VPL generation	real-time	D,G	no	bidirectional VPL generation
[SIP07]	Metropolis Instant Radiosity	VPL generation	real-time	D,G	no	bidirectional VPL generation with Metropolis sampling
[GS10]	Iterative Importance Sampling of VPLs	VPL generation	real-time	D,G	no	rejection sampling of VPL paths
[GKPS12]	Importance Caching for Complex Illumination	VPL selection	offline	D,G	no	importance caching of VPL contributions and improved selection
[RGK*08]	Imperfect Shadow Maps	visibility	real-time	D,G	no	fast approximate shadow maps for VPLs
[REH*11]	Making ISMs View-Adaptive	visibility	real-time	D,G	no	extension of [RGK*08], view-adaptive ISMs and VPL placement
[REG*09]	Micro-Rendering	visibility	interactive	D,G	dna	compute shadow maps for VPLs or final-gather from VPLs
[HREB11]	ManyLODs	visibility	real-time	D,G	dna	fast many-view rasterization
[PGSD13]	Adaptive Quantization Visibility Caching	visibility	offline/interactive	D,G	dna	general visibility cache

Autodesk 360

