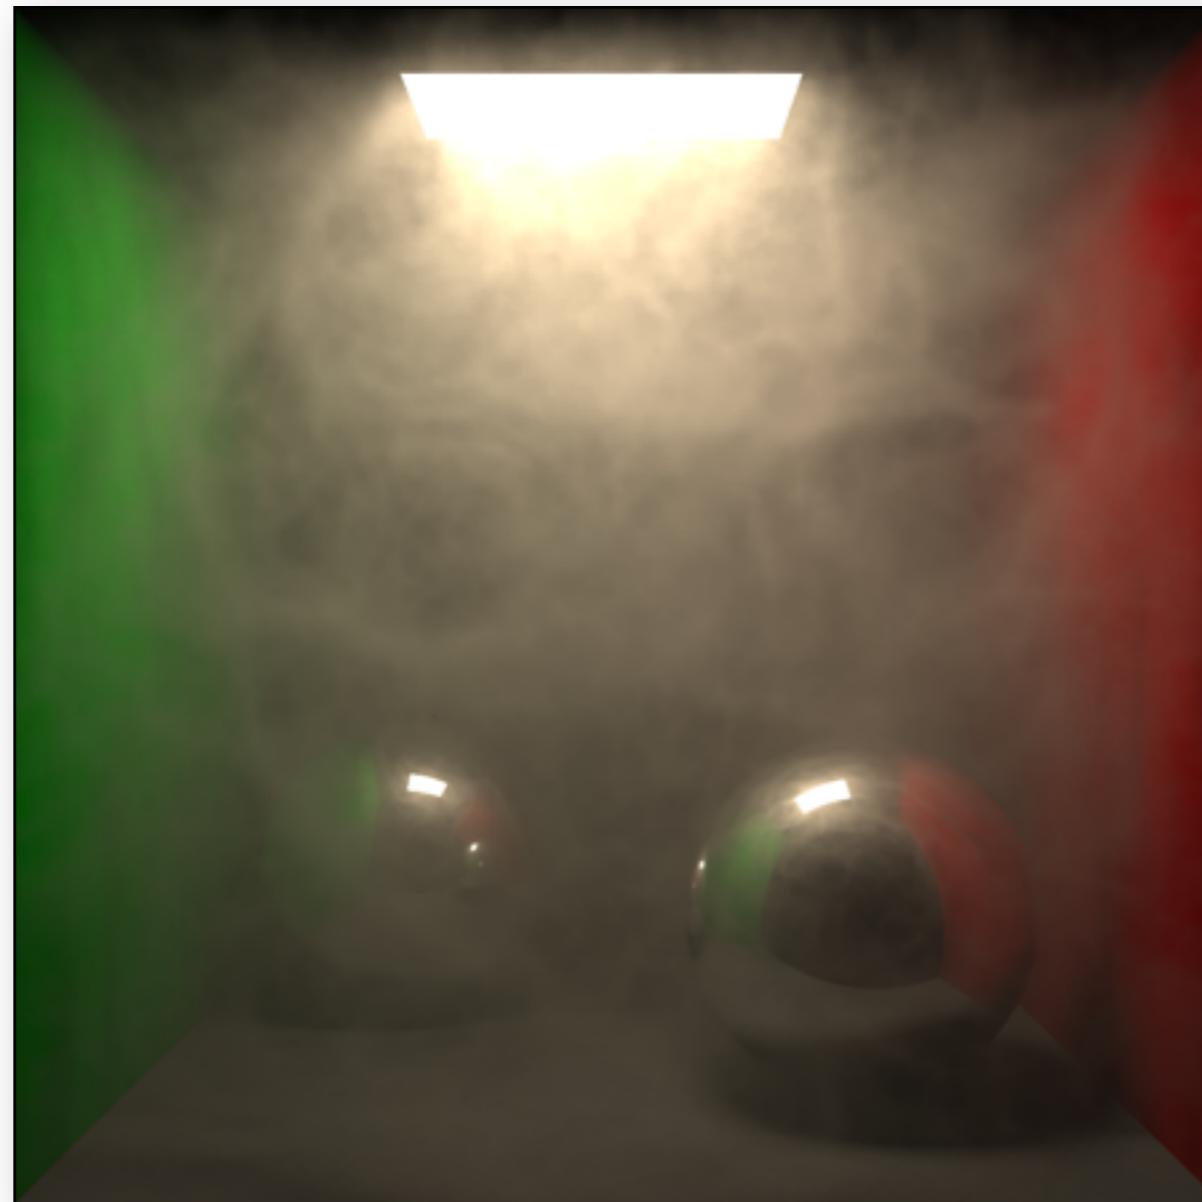


CS 87/187, Spring 2016

RENDERING ALGORITHMS

Caching Illumination

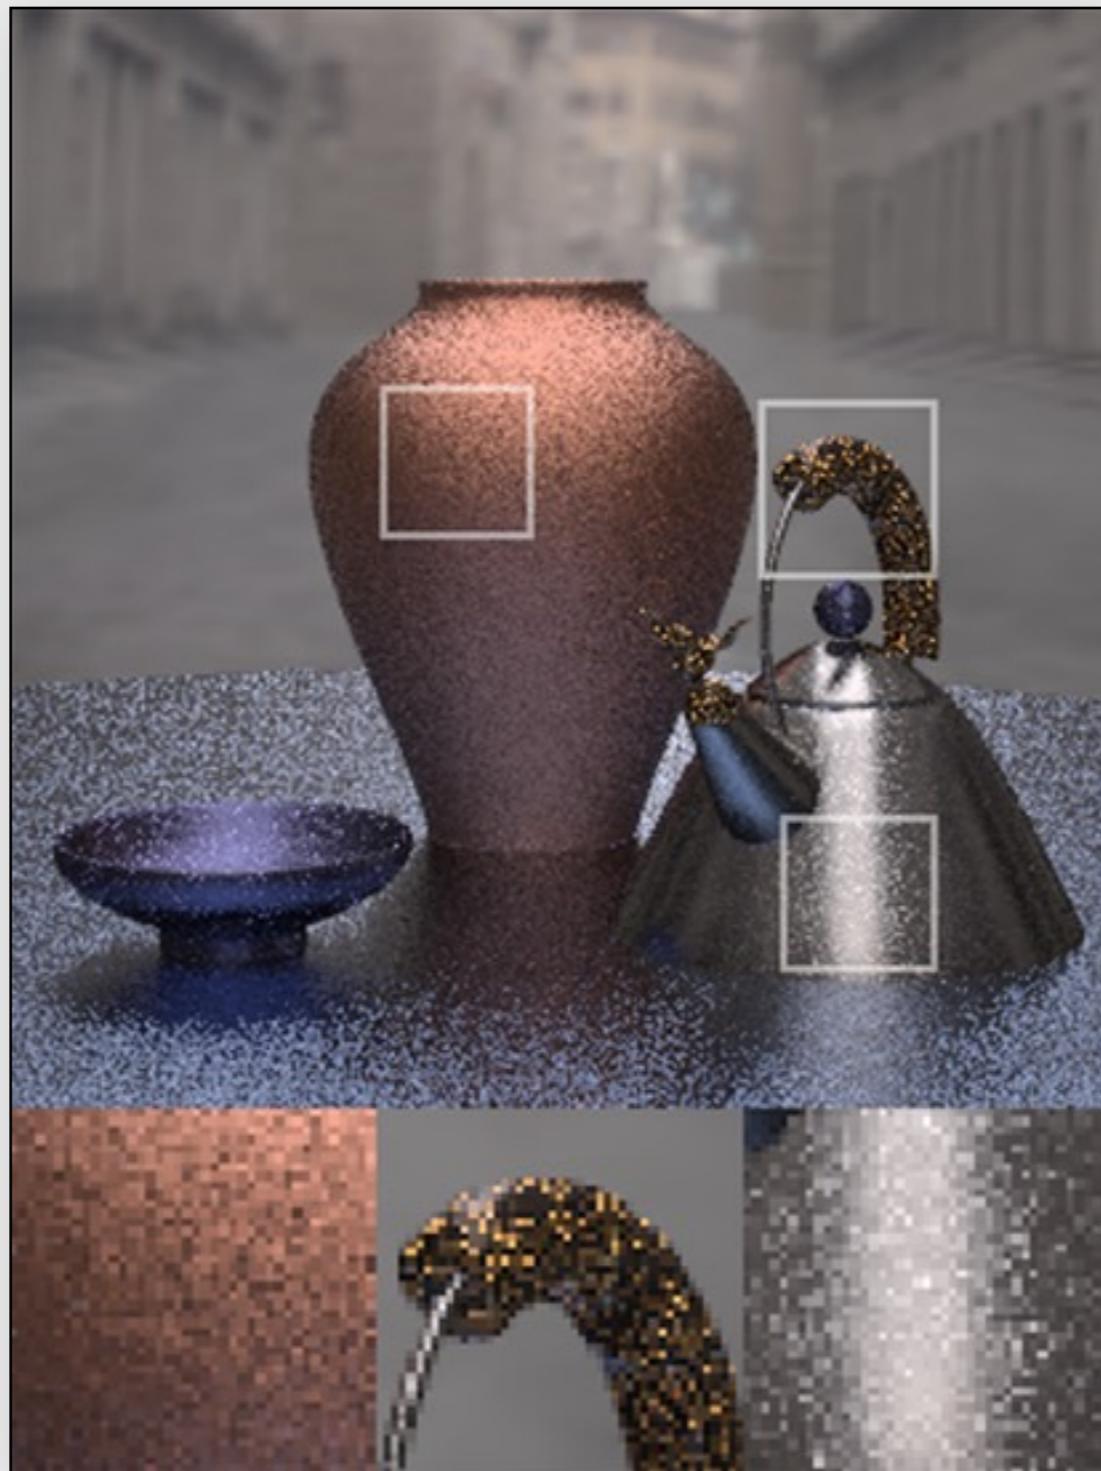


Prof. Wojciech Jarosz
wojciech.k.jarosz@dartmouth.edu

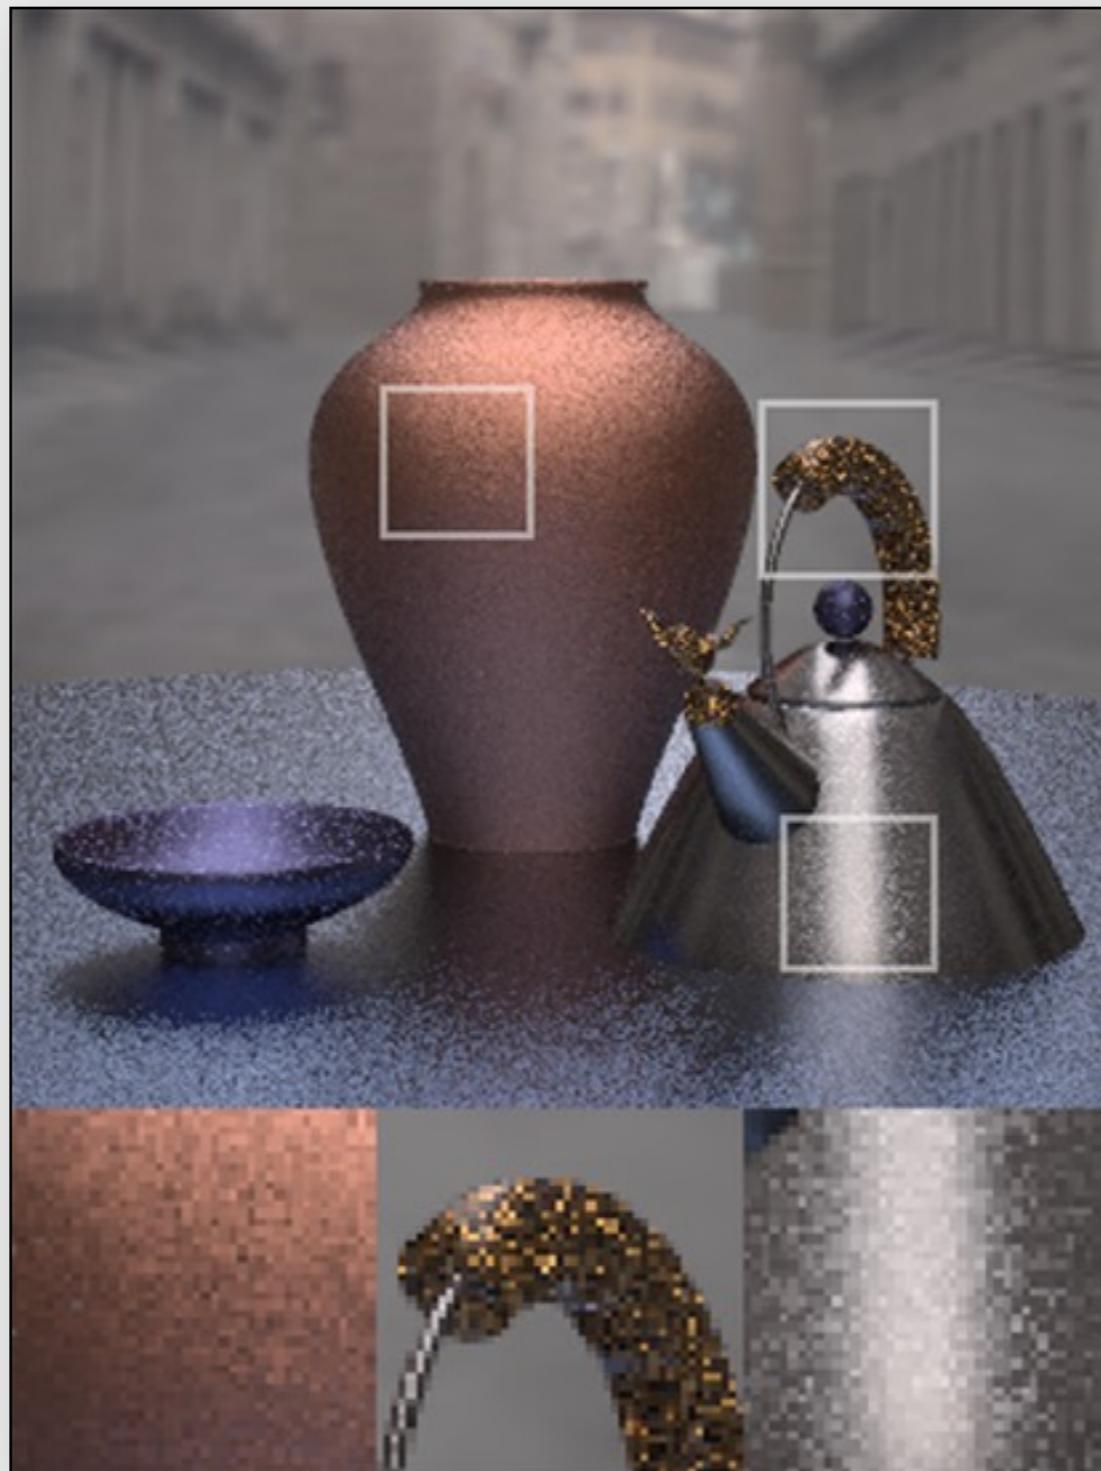
Today's menu

- Irradiance caching and extensions

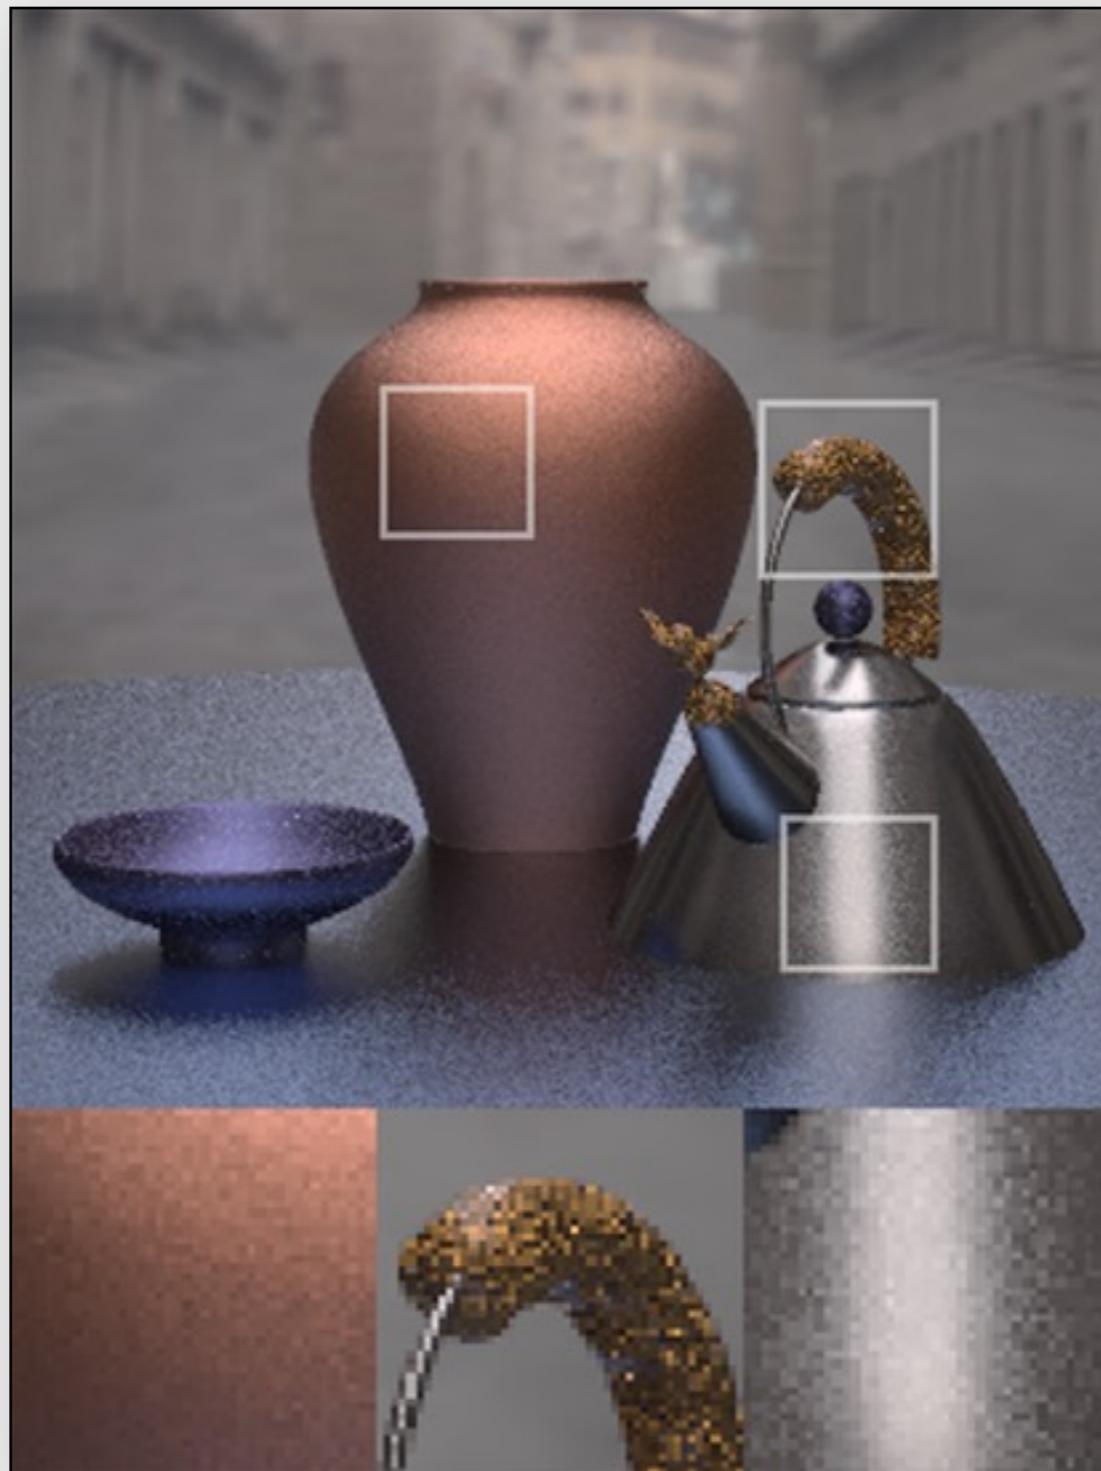
Path tracing - 40 samples per pixel



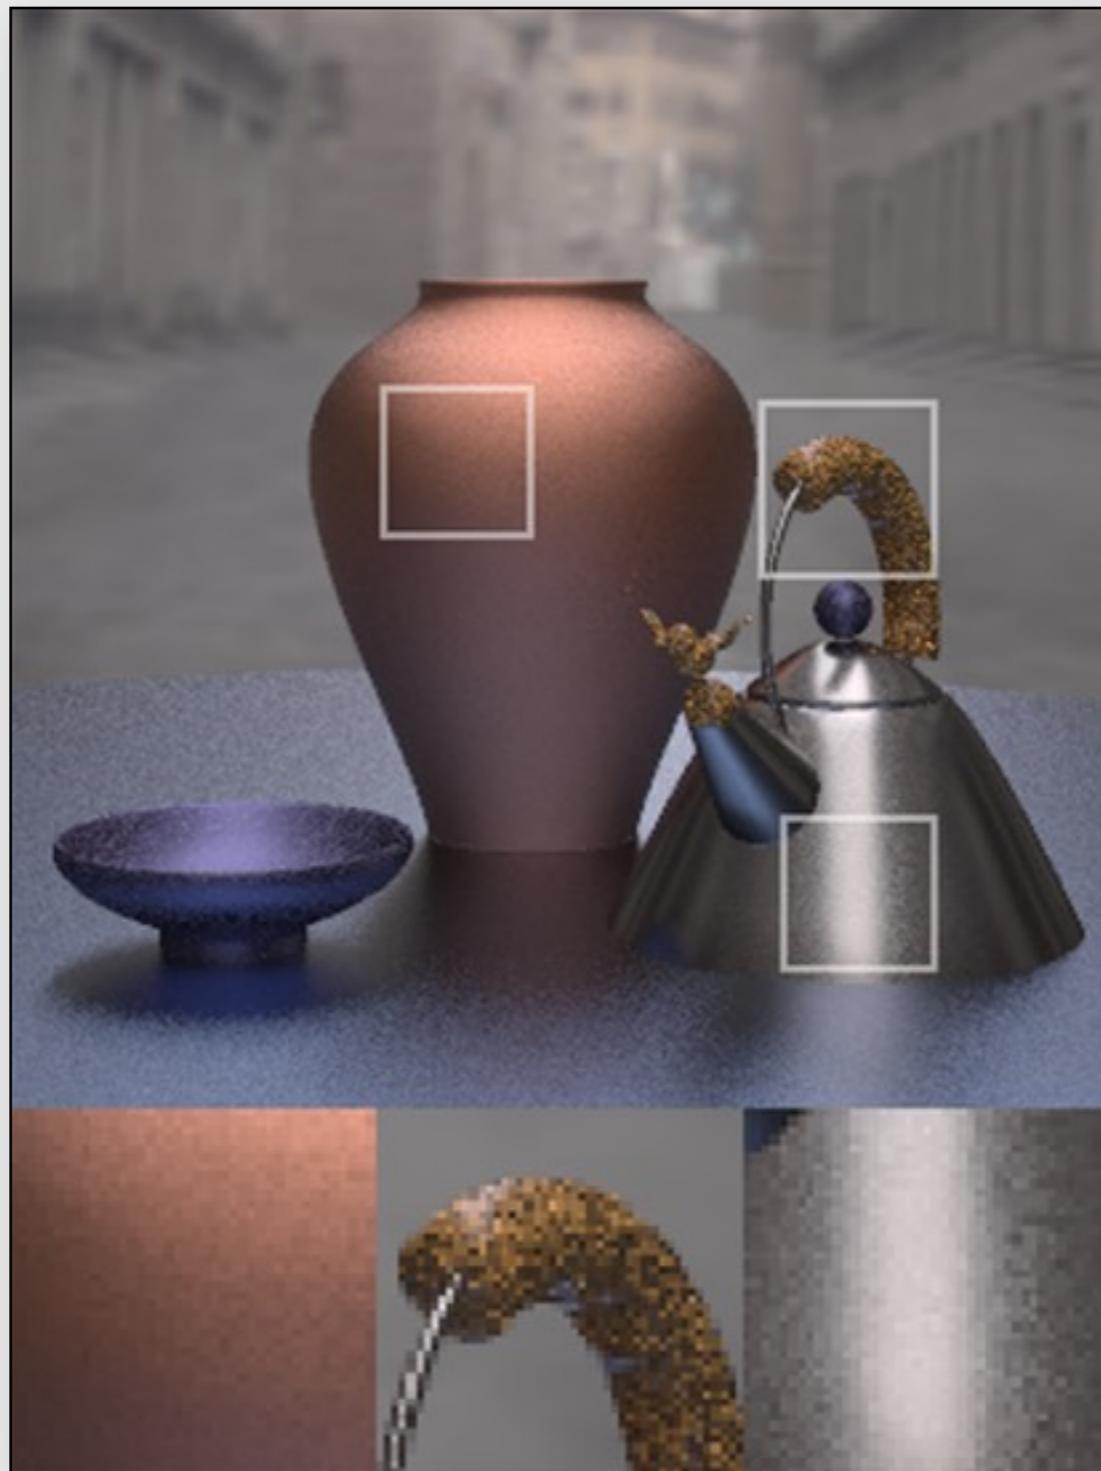
Path tracing - 100 samples per pixel



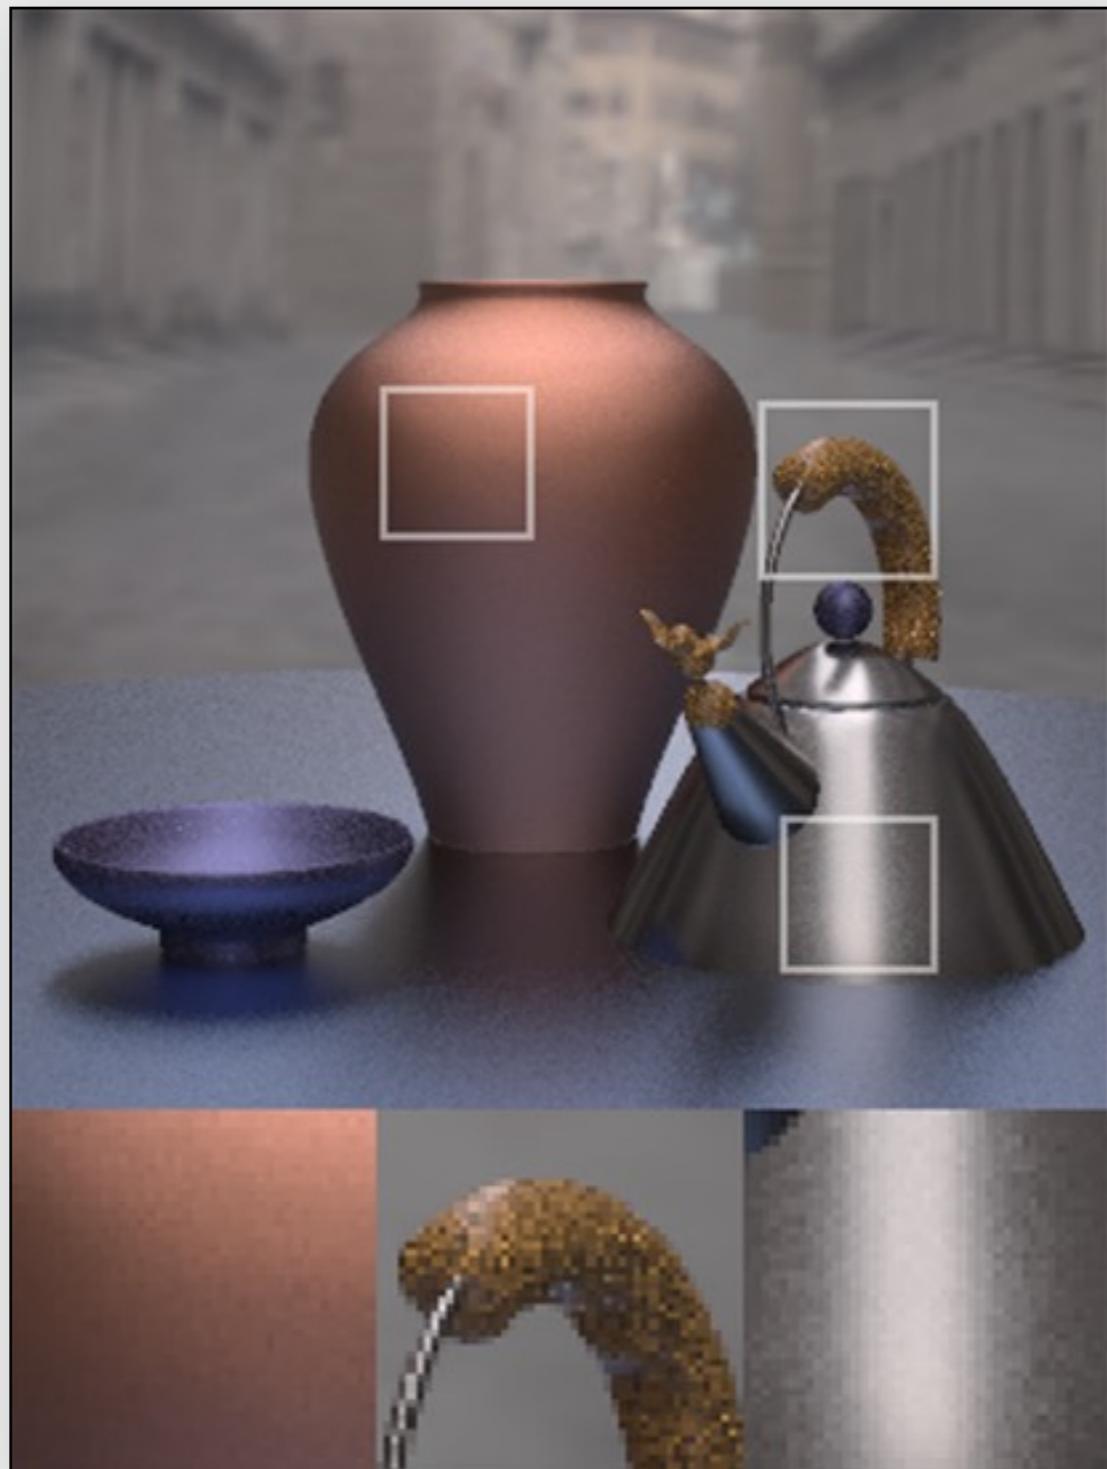
Path tracing - 300 samples per pixel



Path tracing - 600 samples per pixel

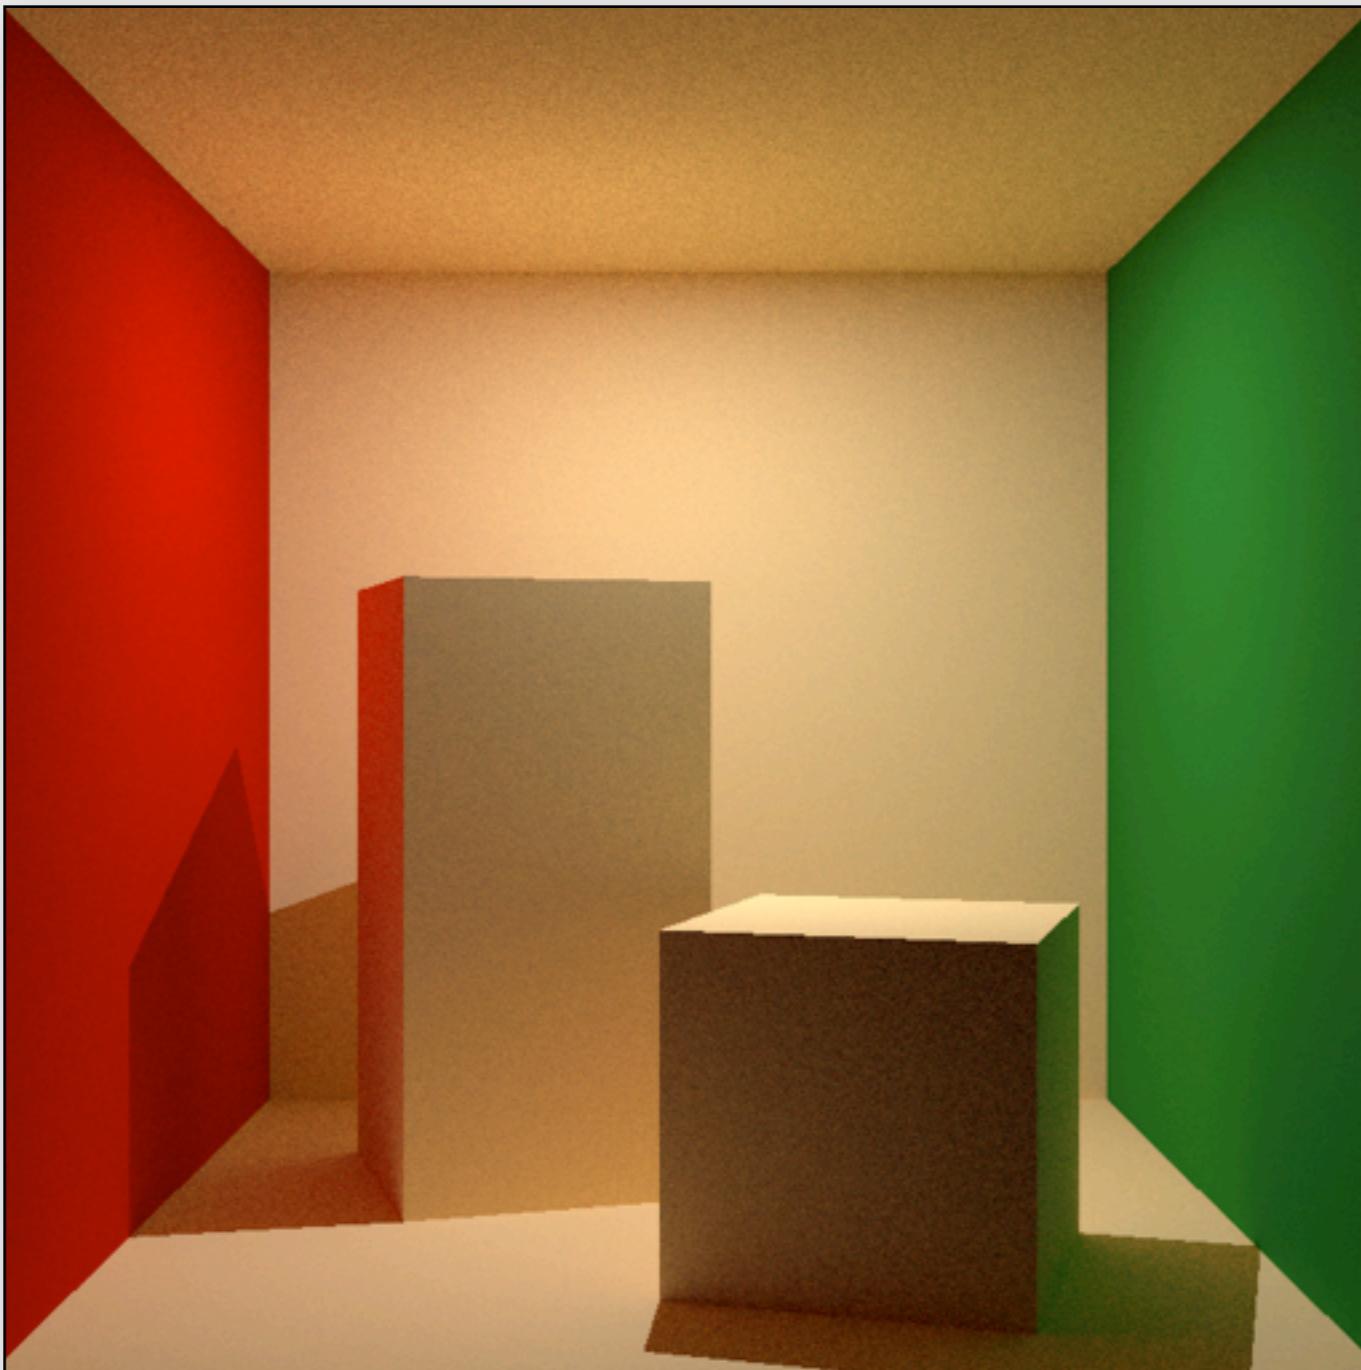


Path tracing - 1200 samples per pixel



Still noisy!
Slow convergence

Path tracing - diffuse scene



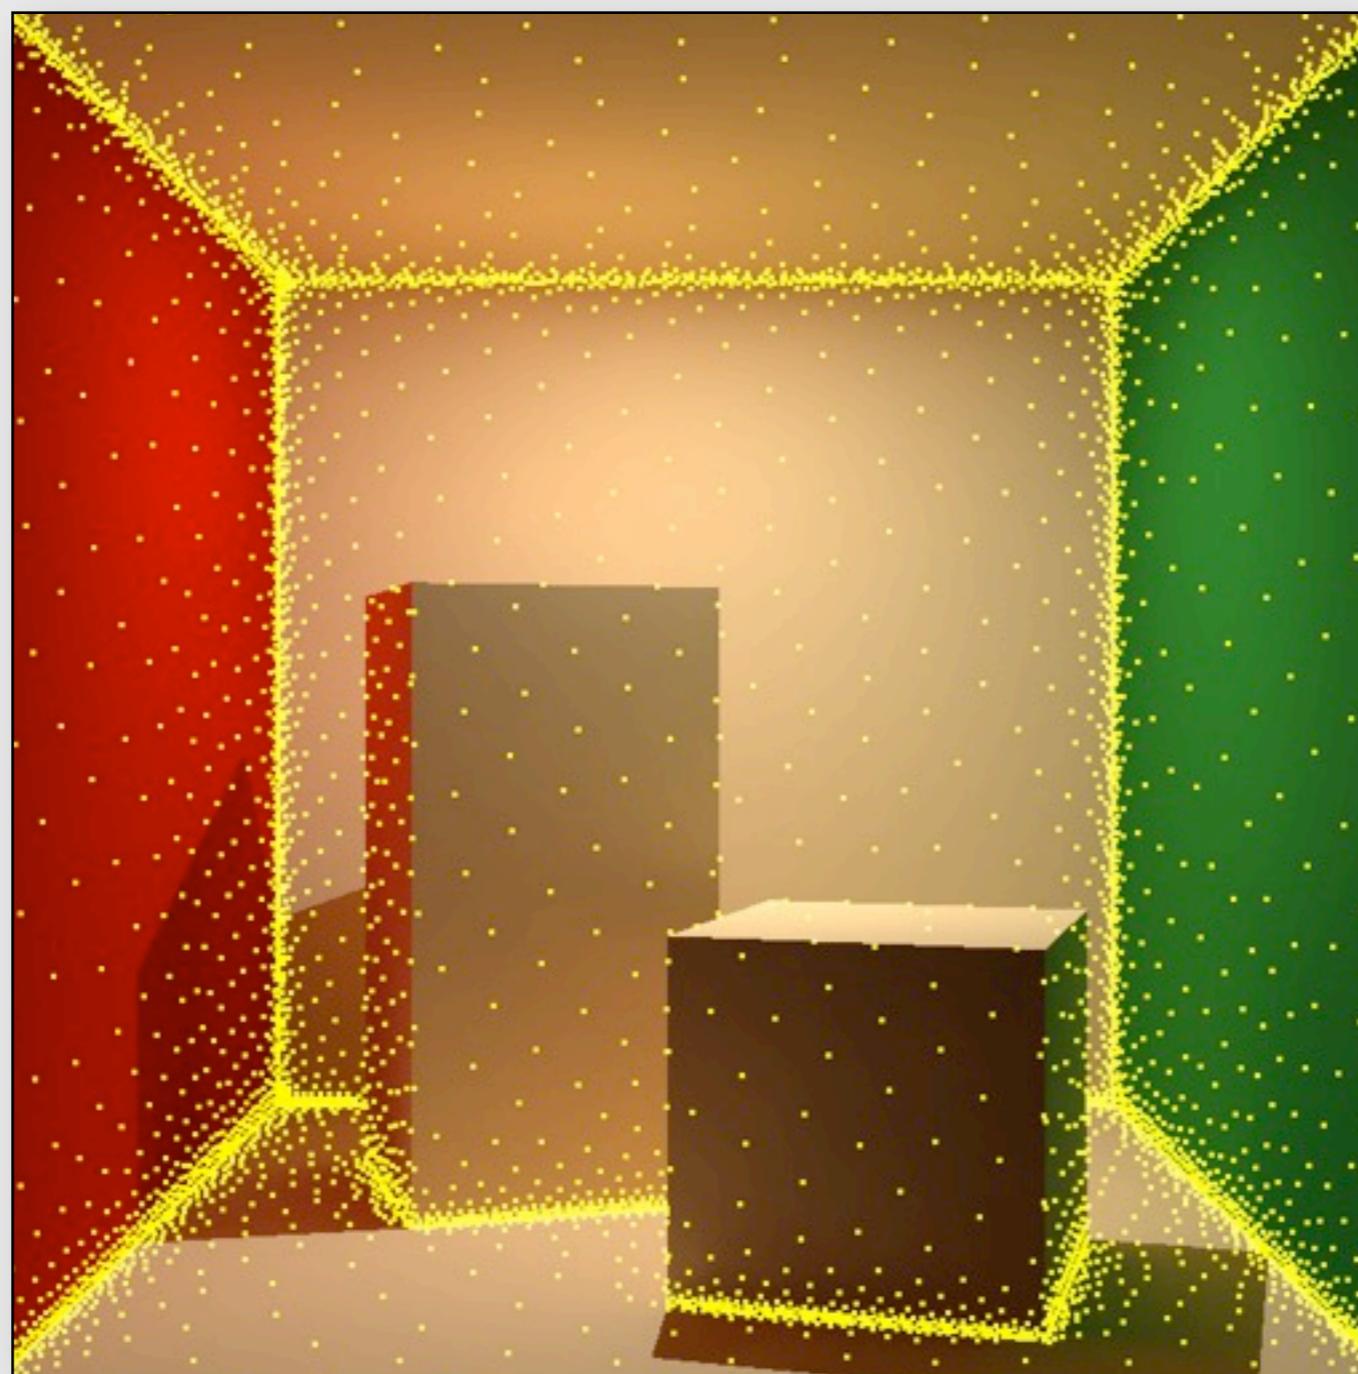
128 paths/pixel

Diffuse indirect illumination is smooth



Perfect candidate for sparse sampling and interpolation

Interpolated indirect illumination

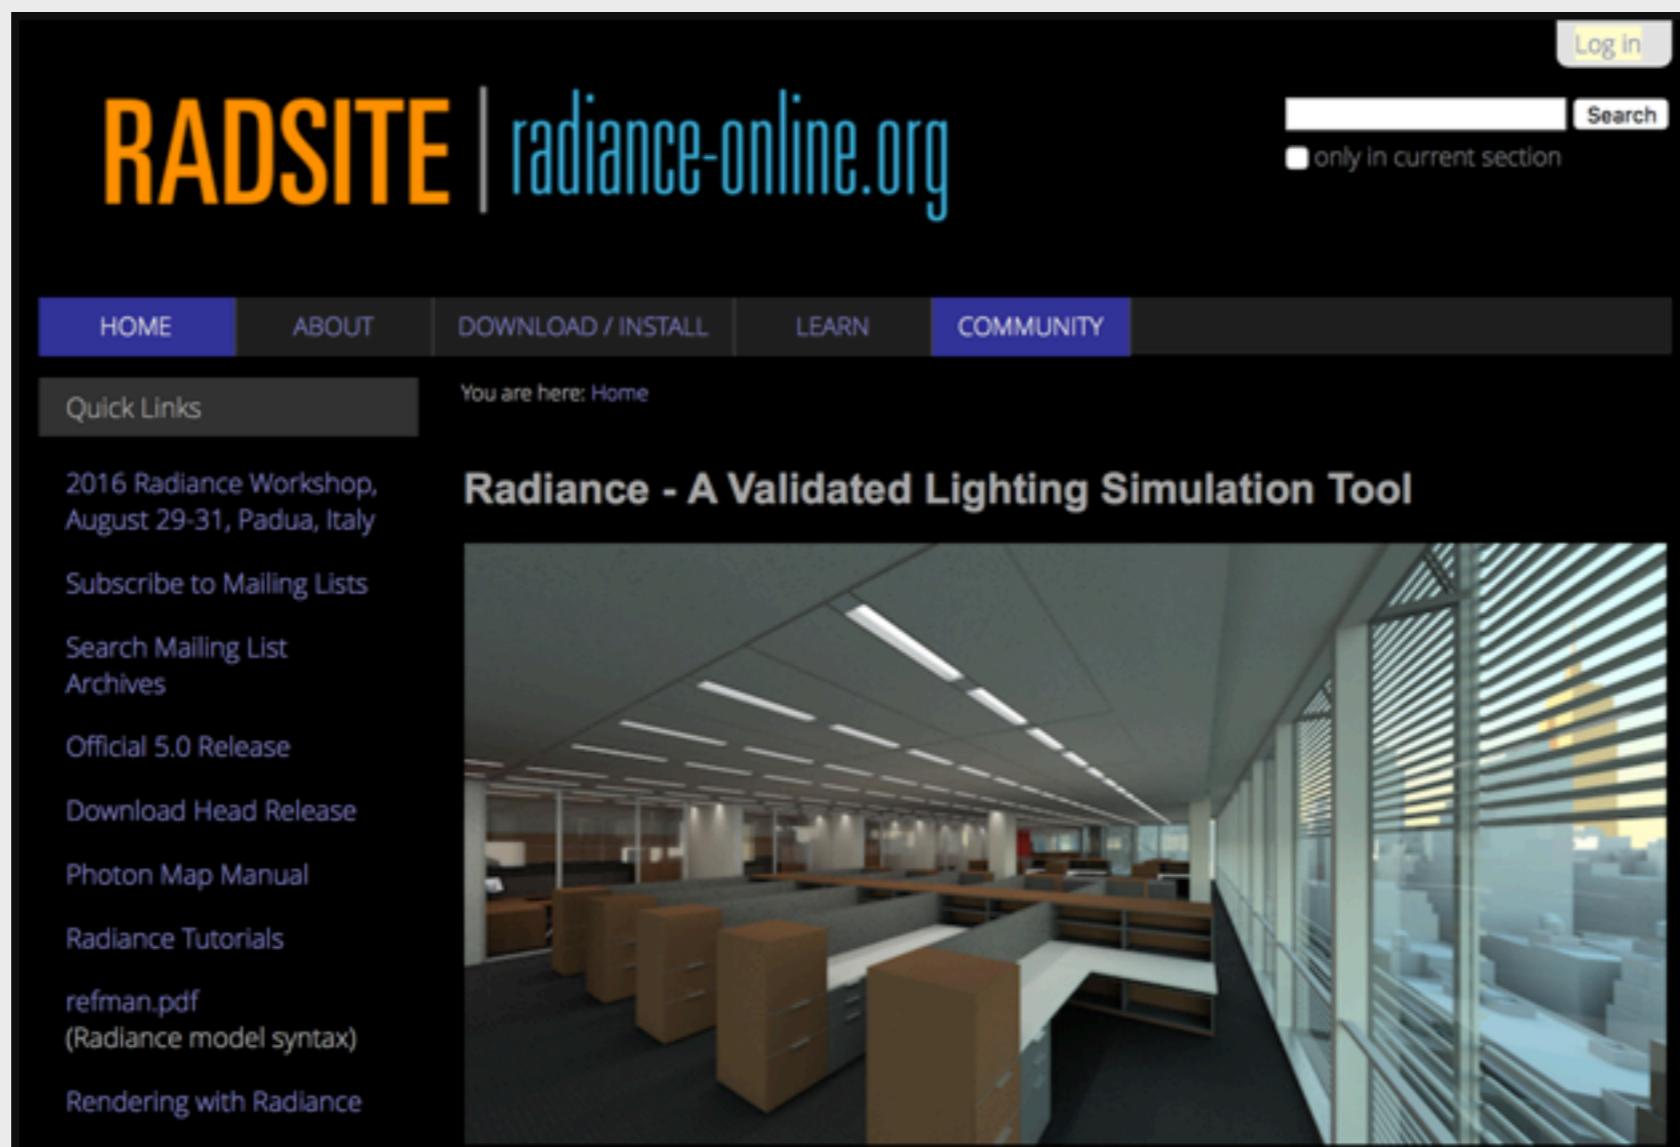


Irradiance Caching
[Ward et al. 1988]

1M pixels - 4K cache points

Irradiance Caching in RADIANCE

- Originally developed by Greg Ward for RADIANCE architectural lighting simulation package



Irradiance Caching Algorithm

```
if at least one cached illumination value found near x then
    Interpolate illumination from the cached value(s).
else
    Compute and cache a new illumination value at x.
```

- Some questions that remain:
 - What do we cache?
 - What makes a cache point “nearby”?
 - How do we interpolate the nearby cached values?

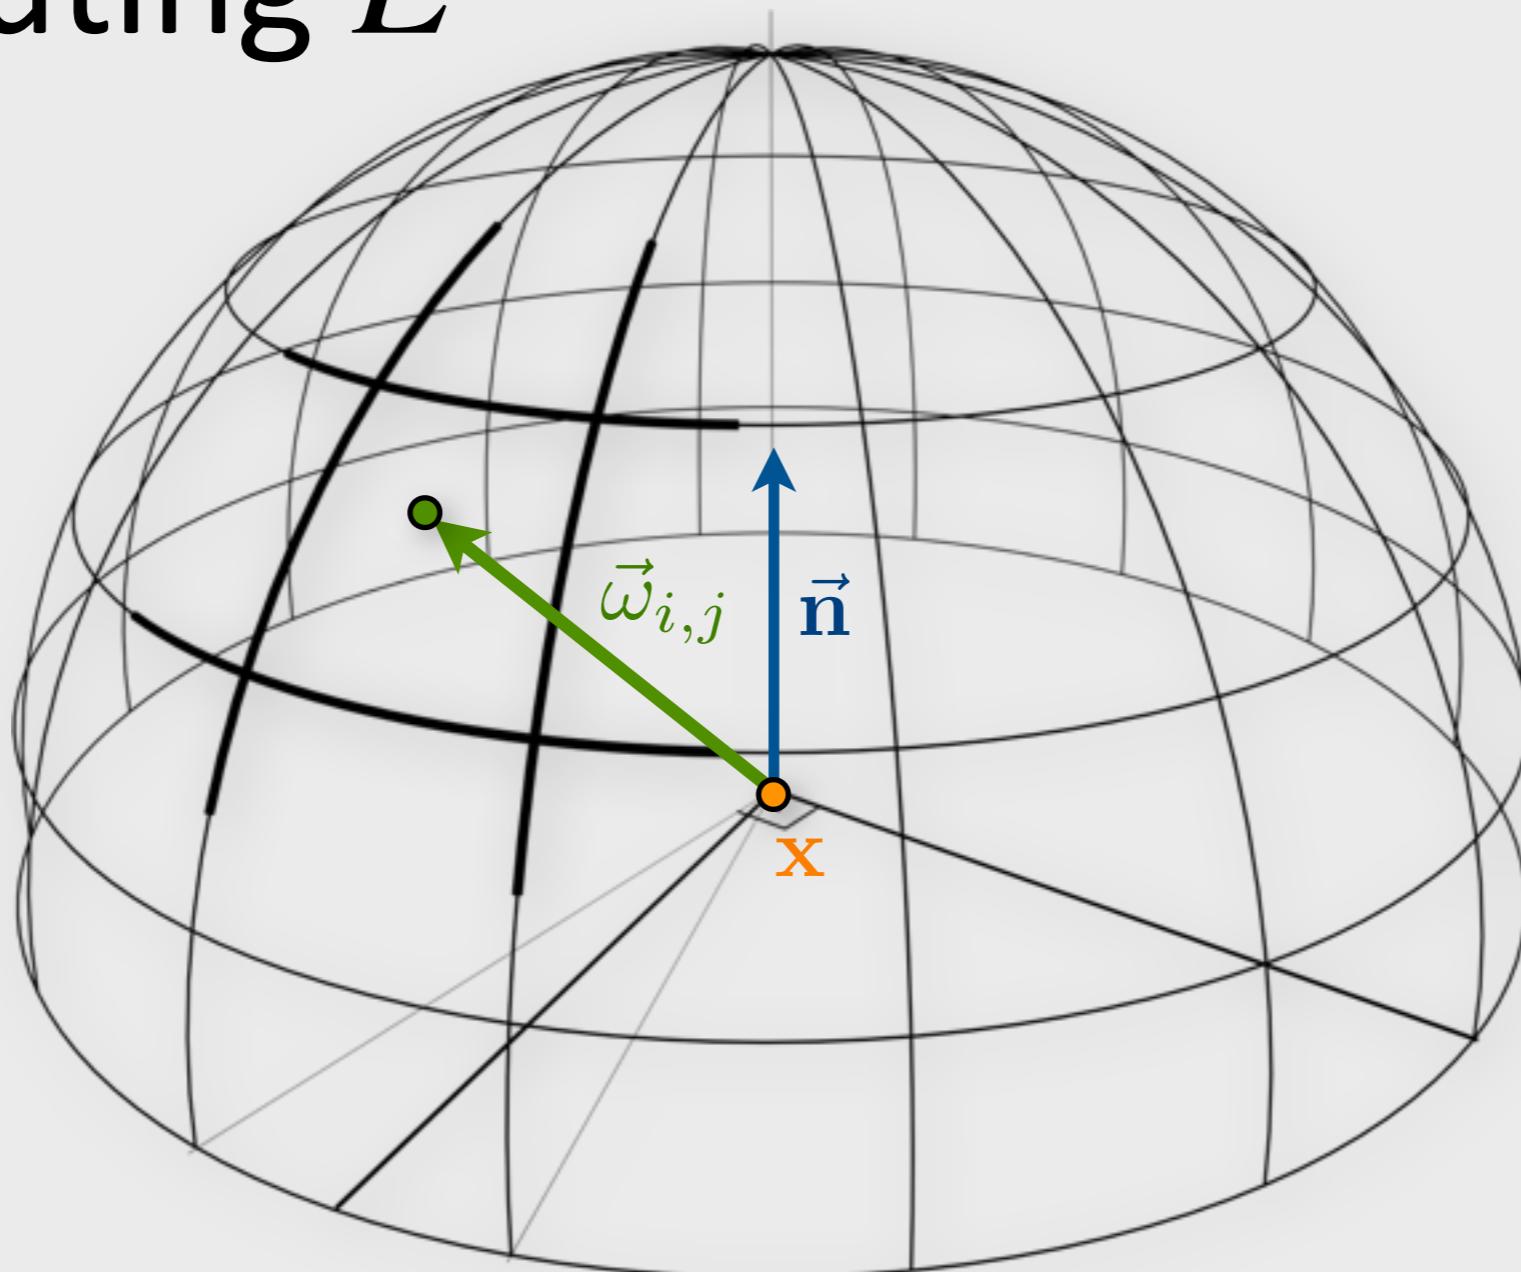
Lambertian assumption

- Indirect illumination on a Lambertian surface:

$$\begin{aligned} L_r(\mathbf{x}, \vec{\omega}_r) &= \int_{H^2} f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_r) L_i(\mathbf{x}, \vec{\omega}_i) (\vec{n} \cdot \vec{\omega}_i) d\vec{\omega}_i \\ &= \underbrace{\frac{\rho}{\pi} \int_{H^2} L_i(\mathbf{x}, \vec{\omega}_i) (\vec{n} \cdot \vec{\omega}_i) d\vec{\omega}_i}_{E(\mathbf{x}, \vec{n}) \rightarrow \text{Irradiance}} \end{aligned}$$

$$E(\mathbf{x}, \vec{n}) \approx \frac{\pi}{N} \sum_{j=1}^N L_i(\mathbf{x}, \vec{\omega}_{i,j})$$

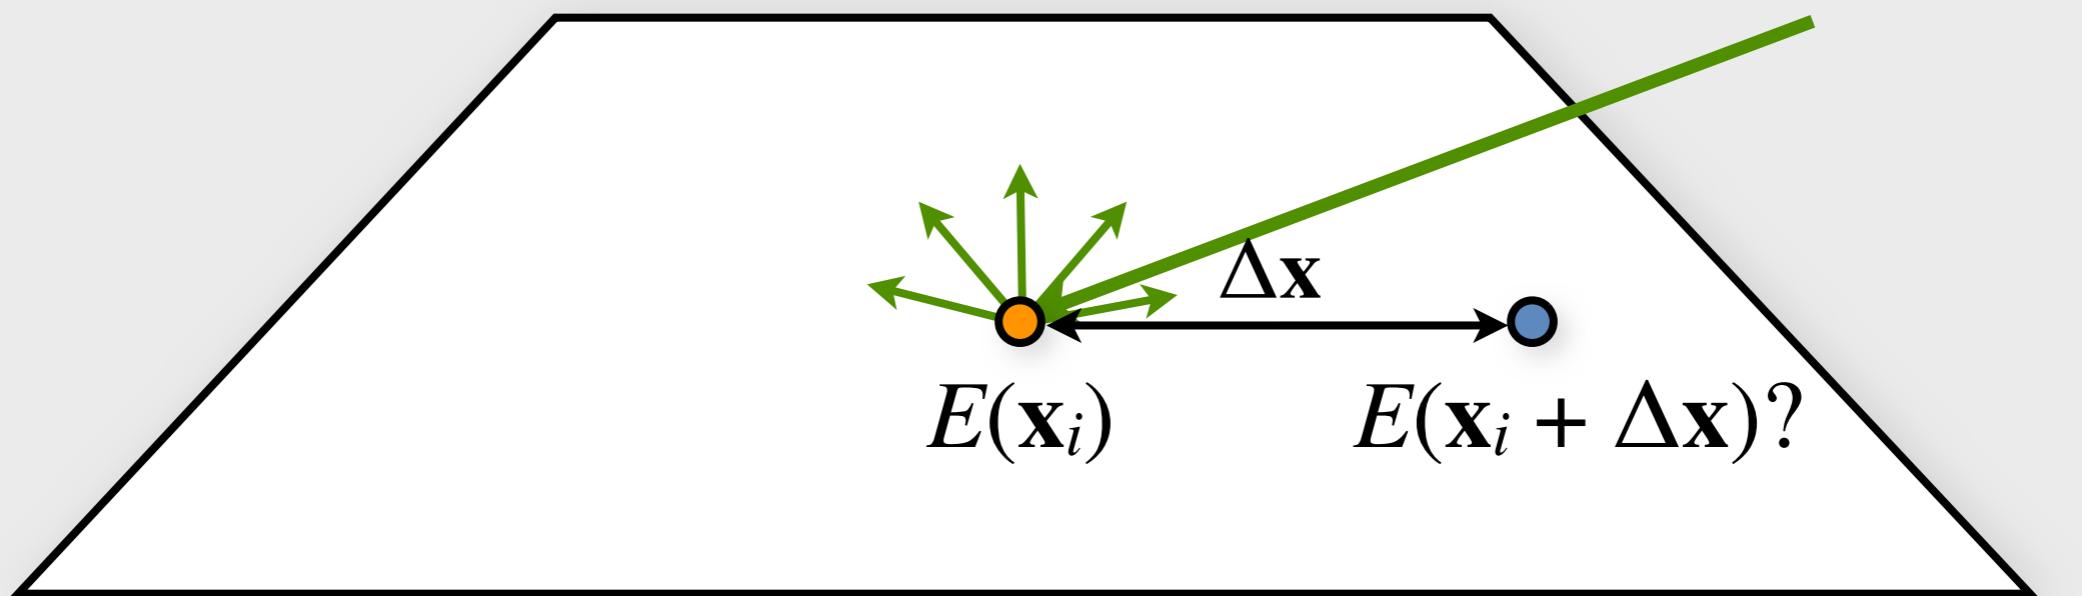
Computing E



$$E(\mathbf{x}, \vec{n}) \approx \frac{\pi}{N_1 N_2} \sum_{i=1}^{N_1} \sum_{j=1}^{N_2} L_i(\mathbf{x}, \vec{\omega}_{i,j})$$

Interpolating Irradiance

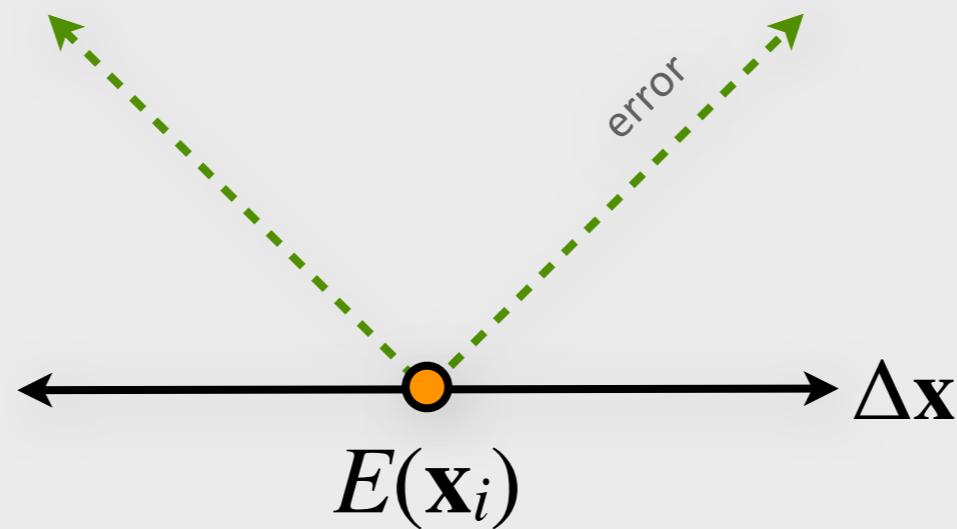
- Irradiance computation costly, reuse whenever possible
- How far away can we reuse a cached value?



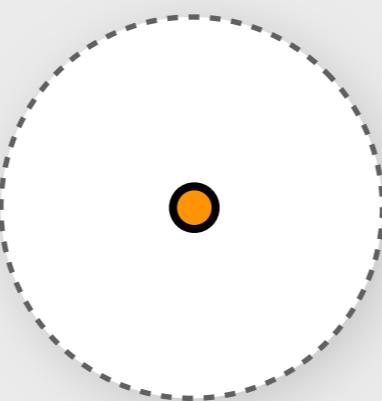
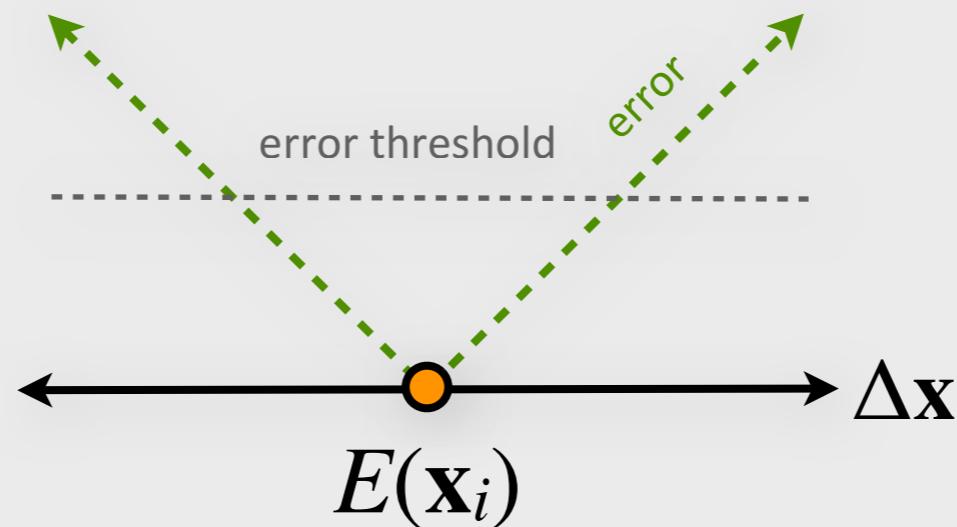
$$E(\mathbf{x}) \approx E(\mathbf{x}_i) + \left(\frac{\partial E}{\partial \mathbf{x}} \Delta\mathbf{x} + \frac{\partial E}{\partial \vec{\mathbf{n}}} \Delta\vec{\mathbf{n}} \right)$$

error

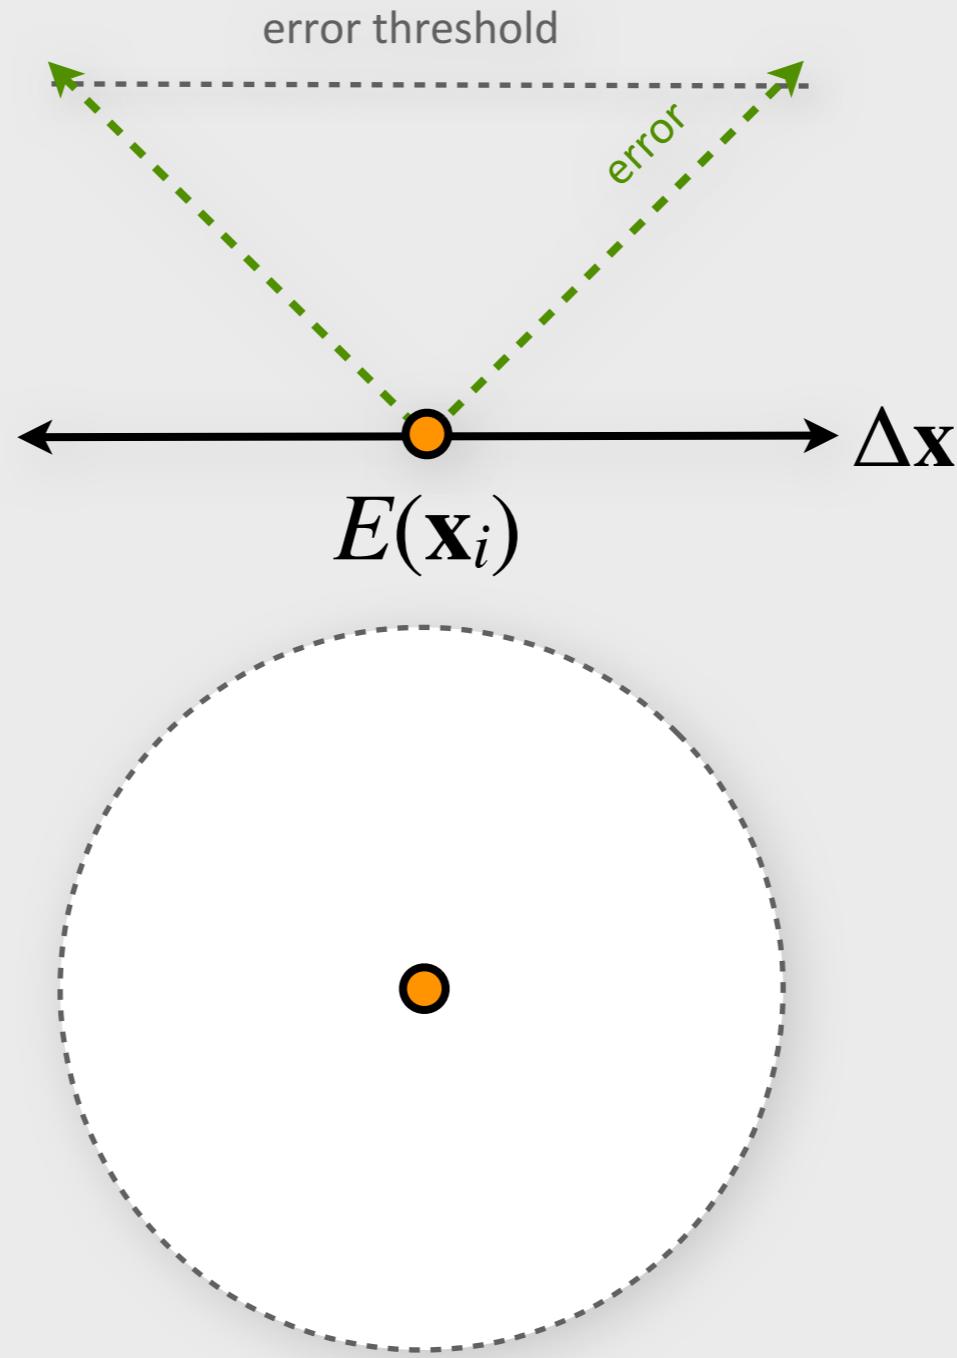
Interpolating Irradiance



Interpolating Irradiance



Interpolating Irradiance

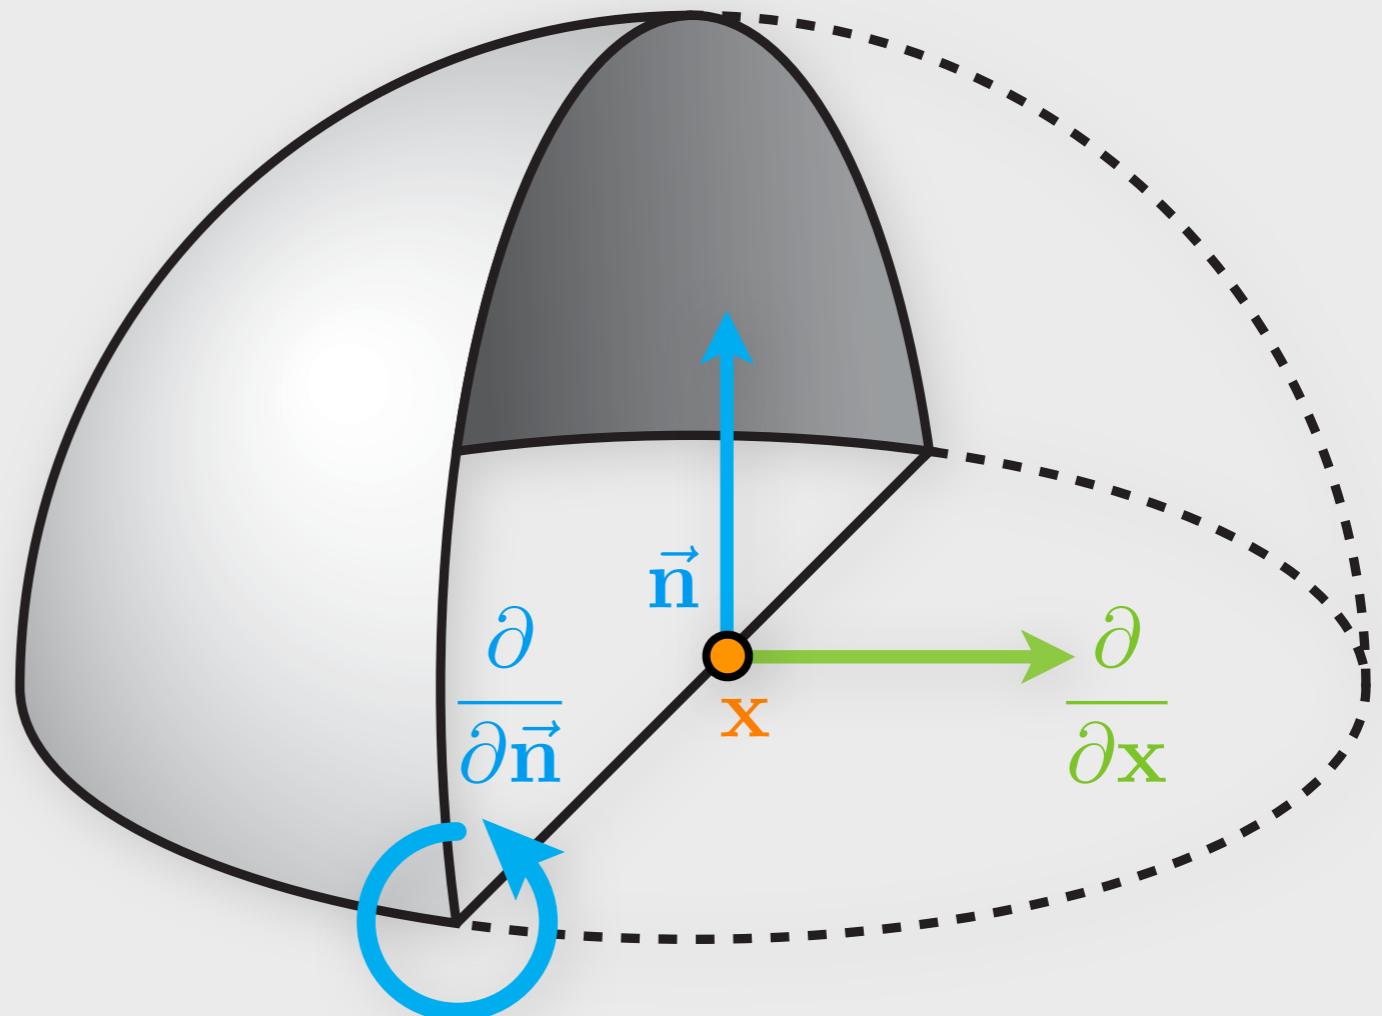


The “Split-Sphere” Heuristic

- To compute valid region, need to estimate change in irradiance

$$\varepsilon_i \lesssim \left| \frac{\partial E}{\partial \mathbf{x}} \Delta \mathbf{x} + \frac{\partial E}{\partial \vec{n}} \Delta \vec{n} \right|$$

- Consider hypothetical, worst-case scene:
the “Split-Sphere”

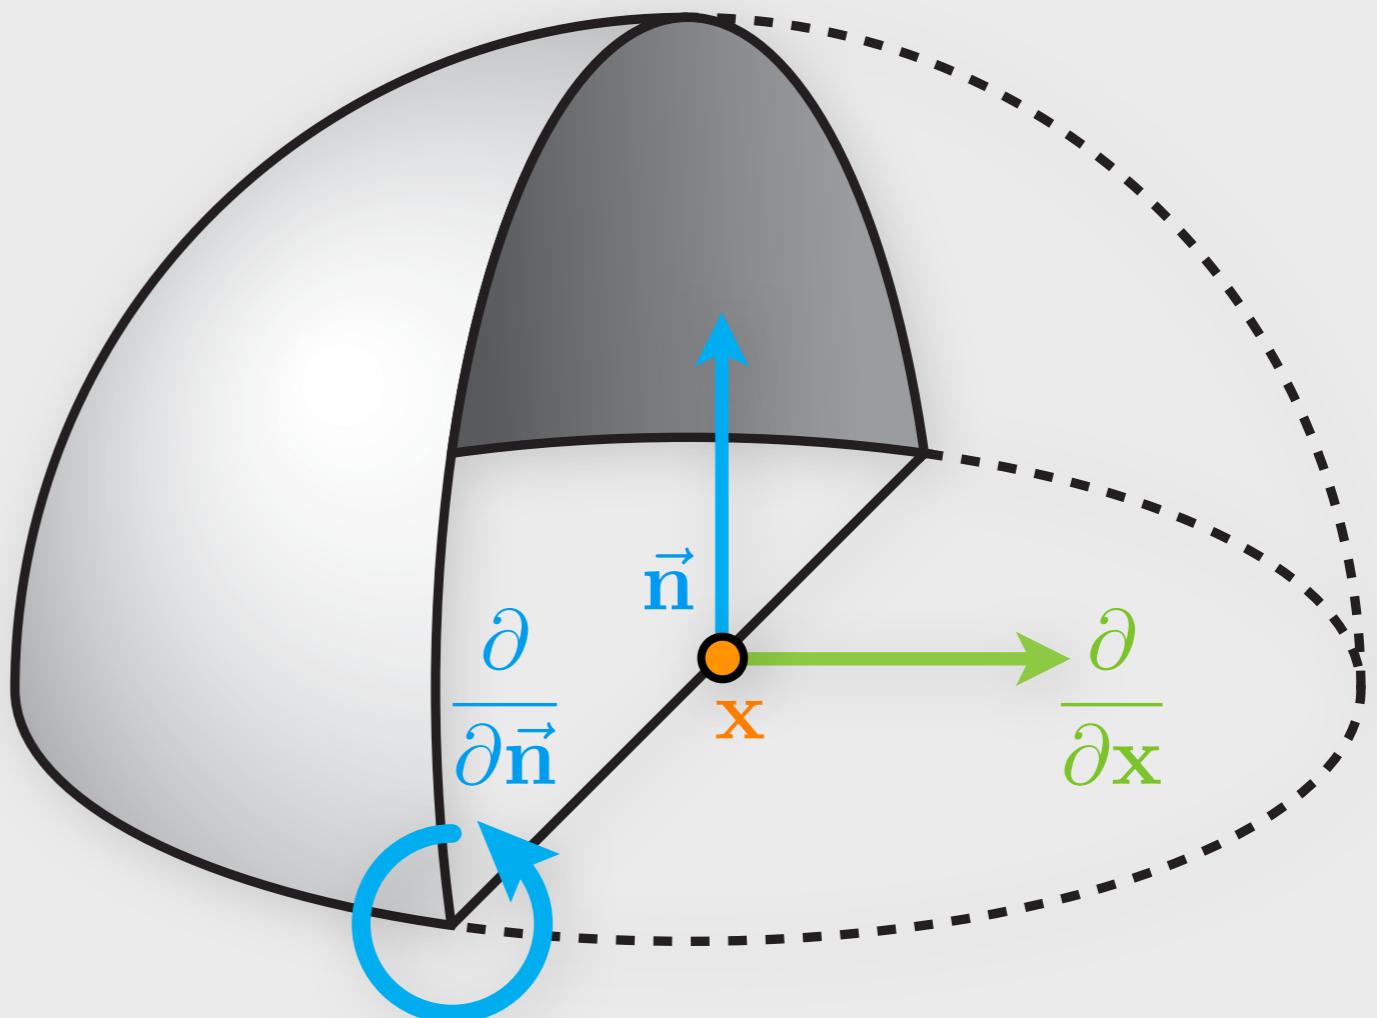


The “Split-Sphere” Heuristic

- To compute valid region, need to estimate change in irradiance

$$\varepsilon_i \lesssim \left| \frac{\partial E}{\partial \mathbf{x}} \Delta \mathbf{x} + \frac{\partial E}{\partial \vec{n}} \Delta \vec{n} \right|$$

- Consider hypothetical, worst-case scene:
the “Split-Sphere”

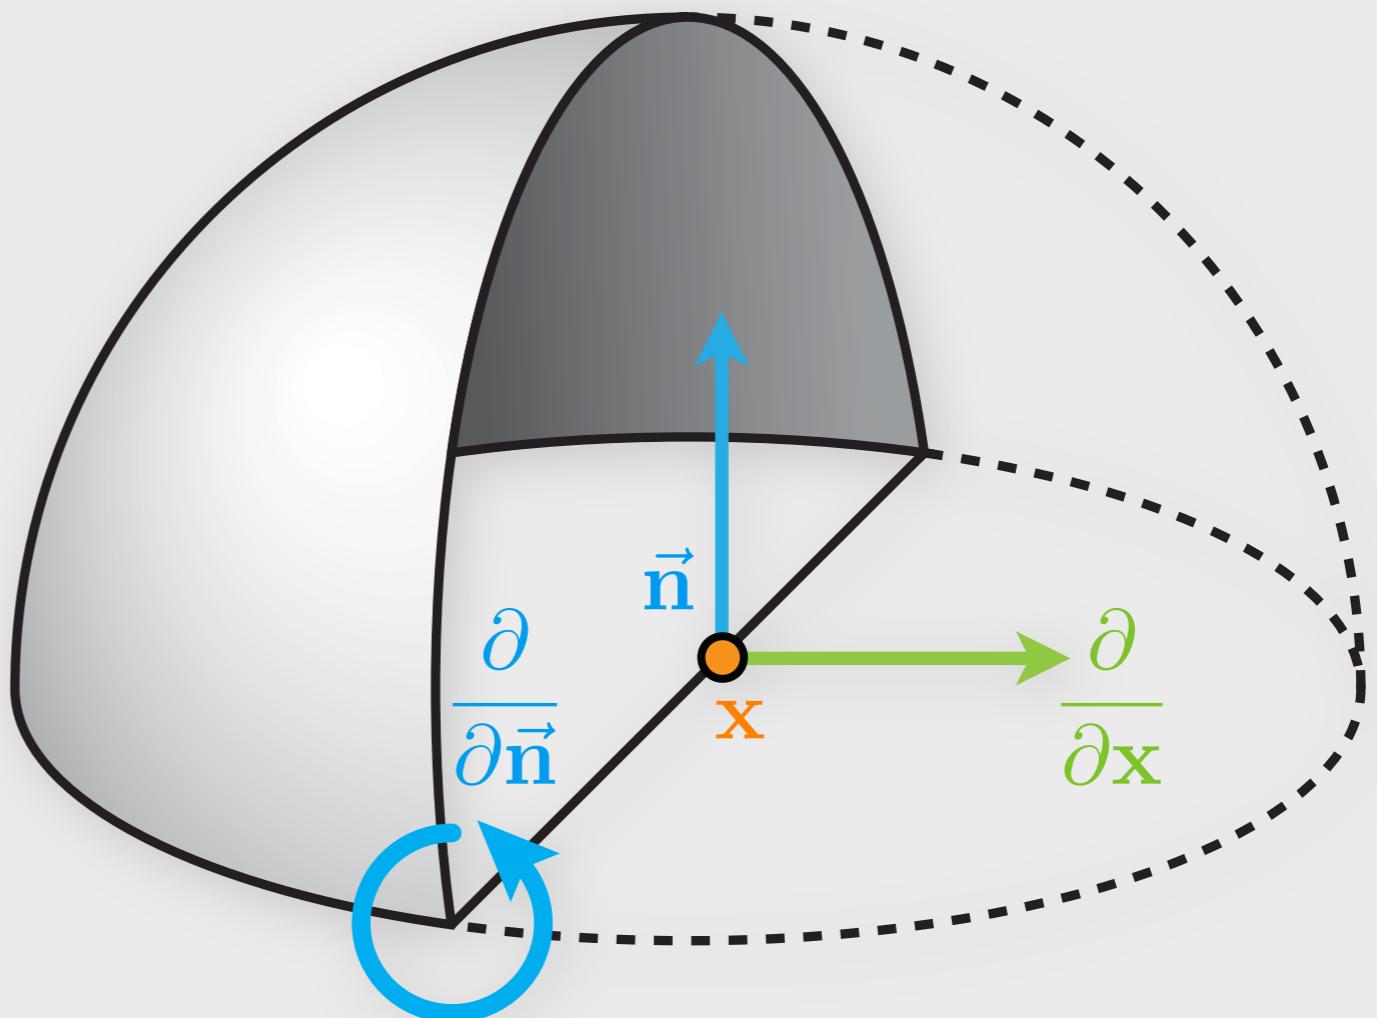


The “Split-Sphere” Heuristic

- To compute valid region, need to estimate change in irradiance

$$\varepsilon_i \lesssim \left| \frac{\partial E}{\partial \mathbf{x}} \Delta \mathbf{x} + \frac{\partial E}{\partial \vec{\mathbf{n}}} \Delta \vec{\mathbf{n}} \right|$$

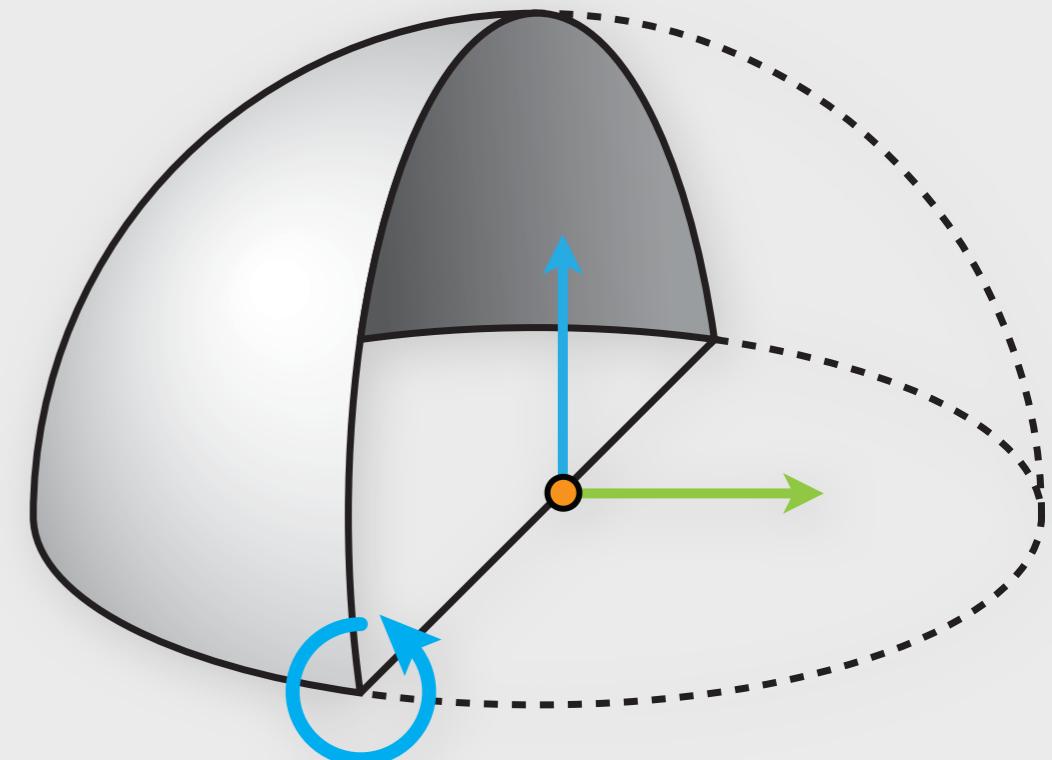
- Consider hypothetical, worst-case scene:
the “Split-Sphere”



The “Split-Sphere” Heuristic

- In the “Split-Sphere” environment the error becomes:

$$\varepsilon_i \lesssim \left| \frac{\partial E}{\partial \mathbf{x}} \Delta \mathbf{x} + \frac{\partial E}{\partial \vec{\mathbf{n}}} \Delta \vec{\mathbf{n}} \right|$$



The “Split-Sphere” Heuristic

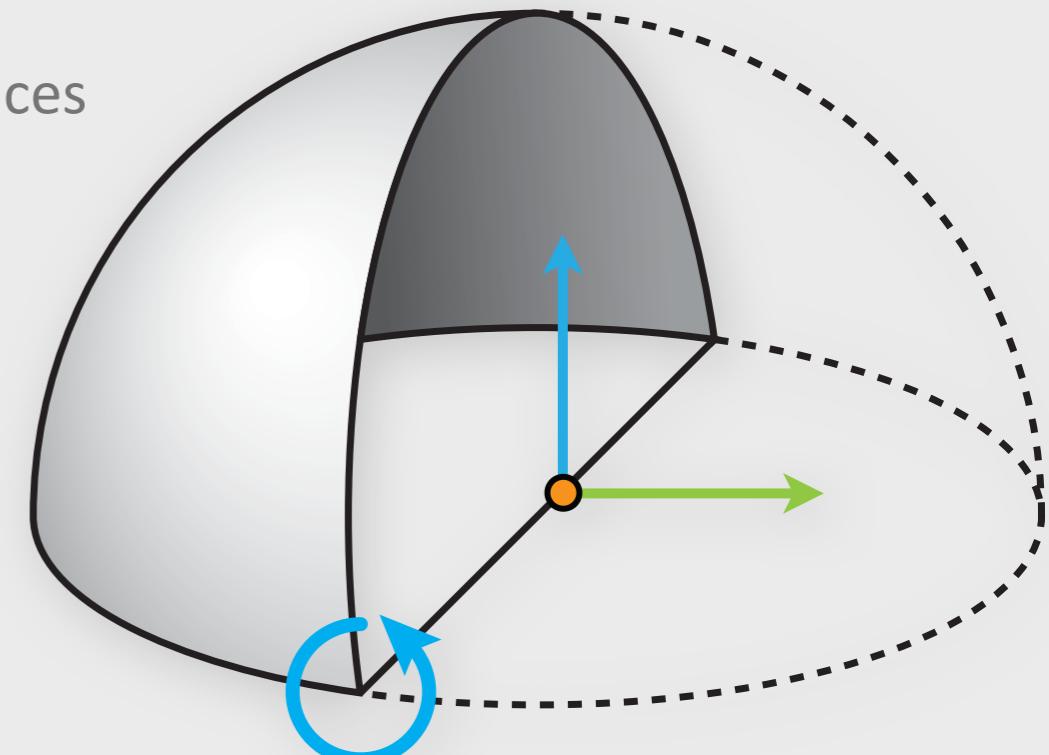
- In the “Split-Sphere” environment the error becomes:

$$\varepsilon_i \lesssim E_i \left(\frac{4}{\pi} \frac{\|\mathbf{x} - \mathbf{x}_i\|}{R_i} + \sqrt{1 - (\vec{\mathbf{n}} \cdot \vec{\mathbf{n}}_i)} \right)$$

position difference, relative to
radius of sphere

orientation difference

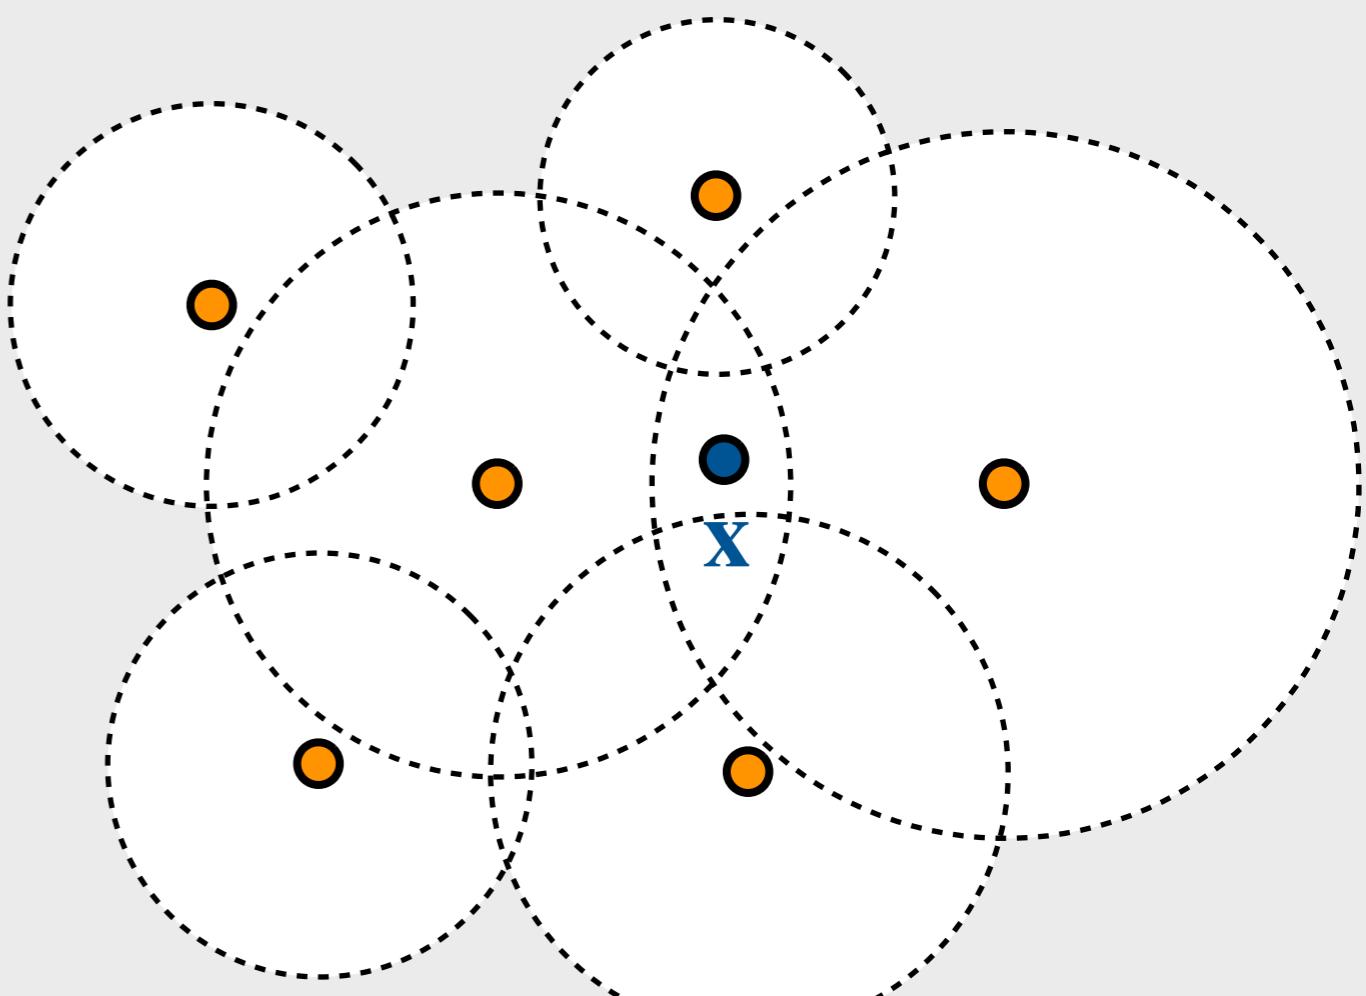
“average” distance to visible surfaces



Interpolating Irradiance

- At each **shading location**, perform a weighted average of all **cached values** which have an error below some threshold.

$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) E_i}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$



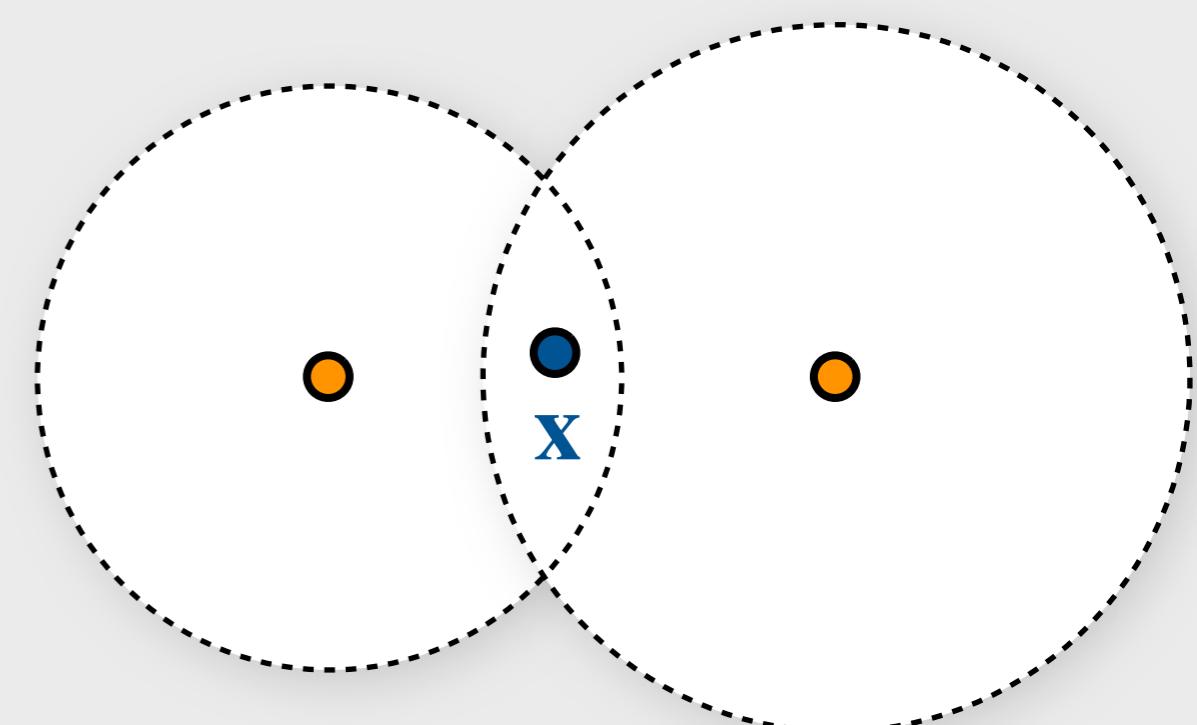
Interpolating Irradiance

- At each **shading location**, perform a weighted average of all **cached values** which have an error below some threshold.

$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) E_i}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$

where:

$$S = \{i : \epsilon_i(\mathbf{x}, \vec{\mathbf{n}}) < a\}$$



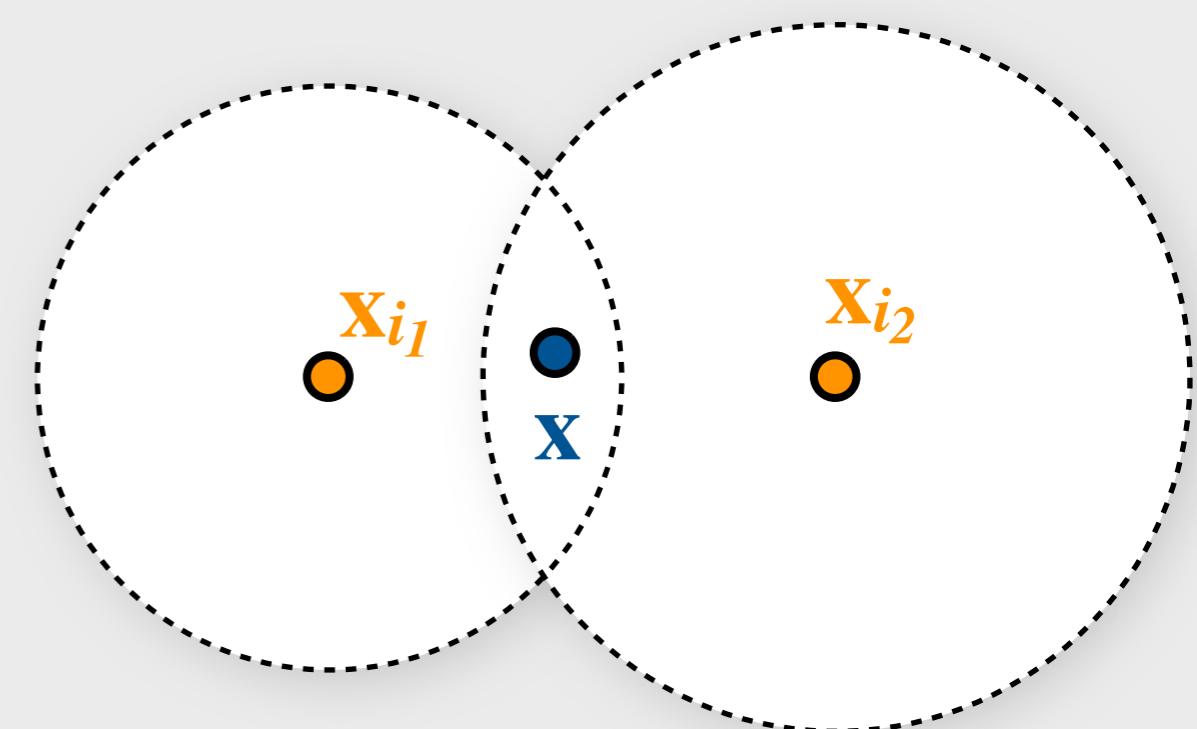
Interpolating Irradiance

- At each **shading location**, perform a weighted average of all **cached values** which have an error below some threshold.

$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) E_i}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$

where:

$$S = \{i : \epsilon_i(\mathbf{x}, \vec{\mathbf{n}}) < a\}$$



Interpolating Irradiance

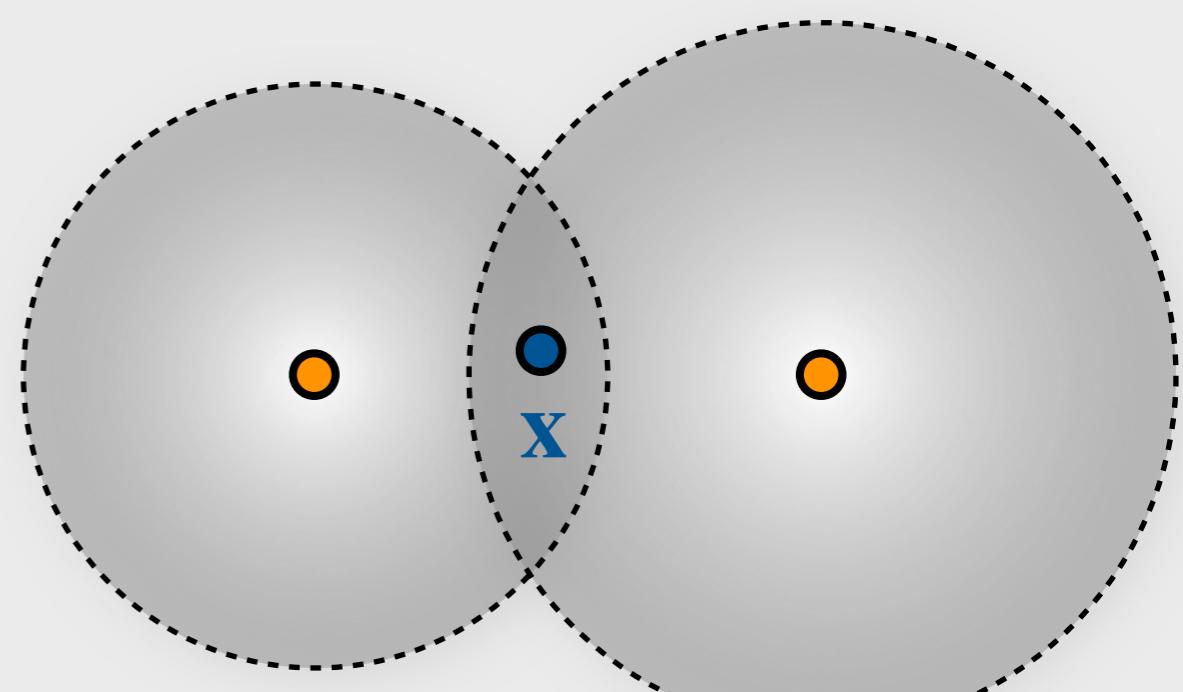
- At each **shading location**, perform a weighted average of all **cached values** which have an error below some threshold.
- Reciprocal of the error is used as the weight

$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) E_i}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$

where:

$$S = \{i : \epsilon_i(\mathbf{x}, \vec{\mathbf{n}}) < a\}$$

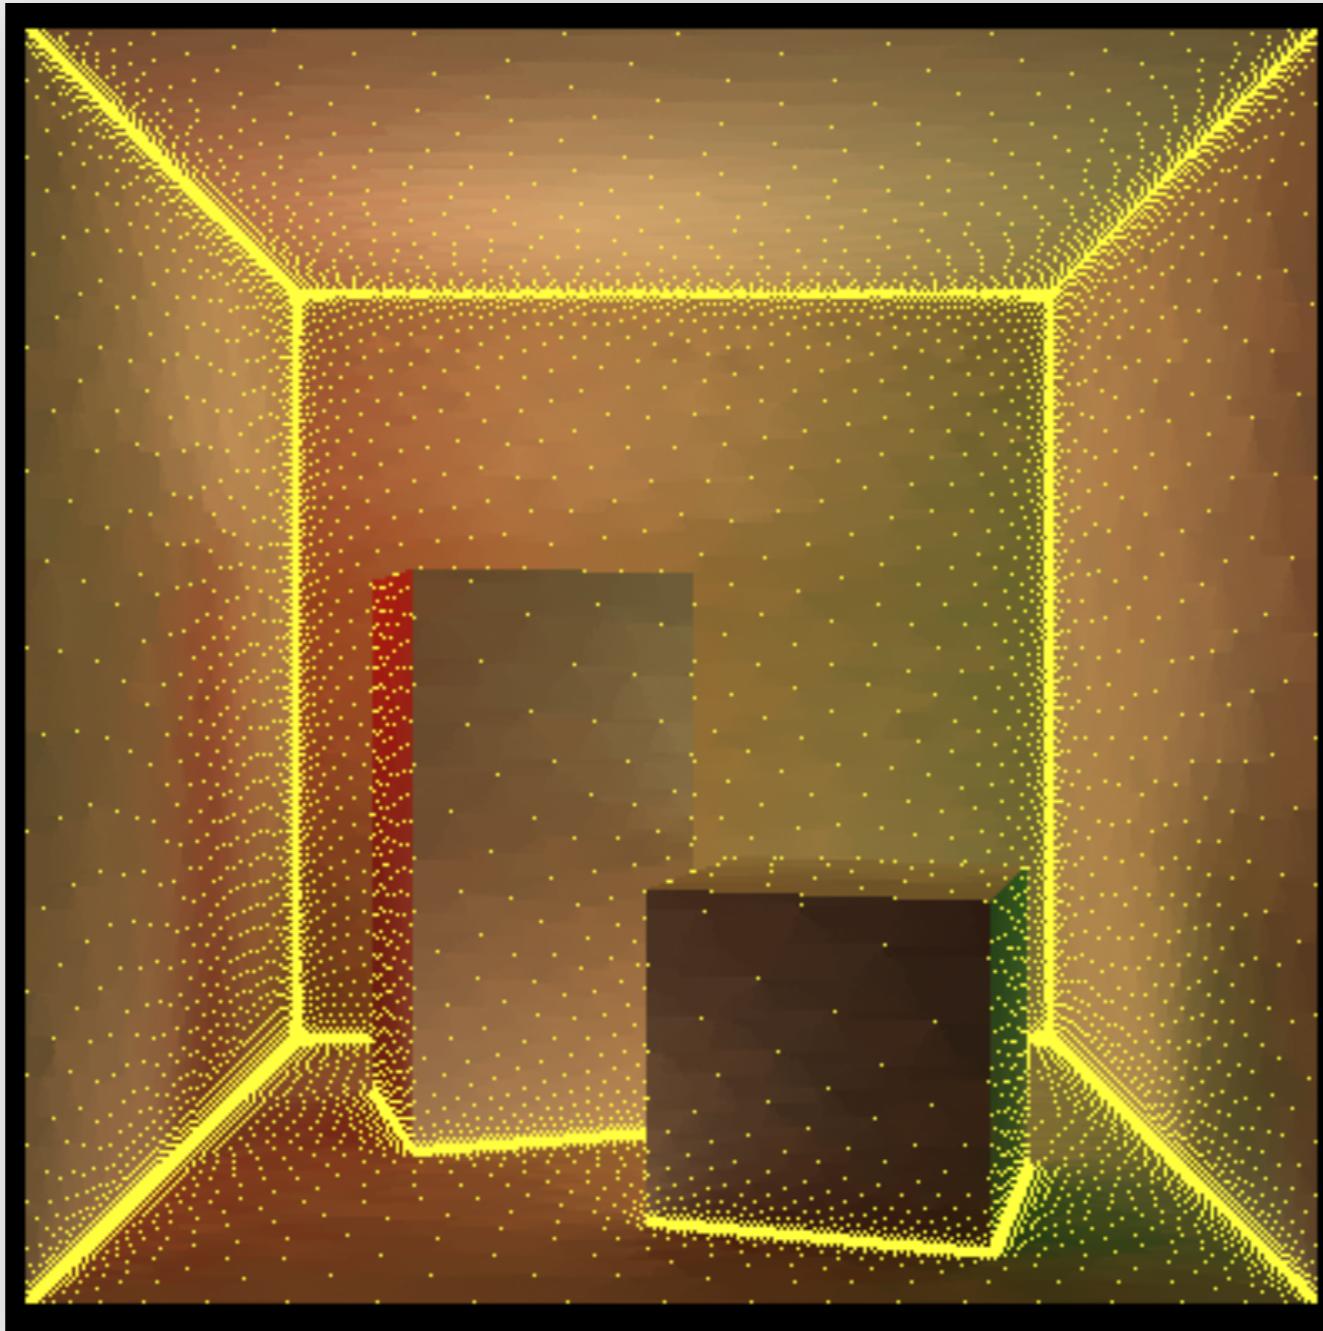
$$w_i(\mathbf{x}, \vec{\mathbf{n}}) = \frac{1}{\epsilon_i(\mathbf{x}, \vec{\mathbf{n}})} - \frac{1}{a}$$



Practical Considerations

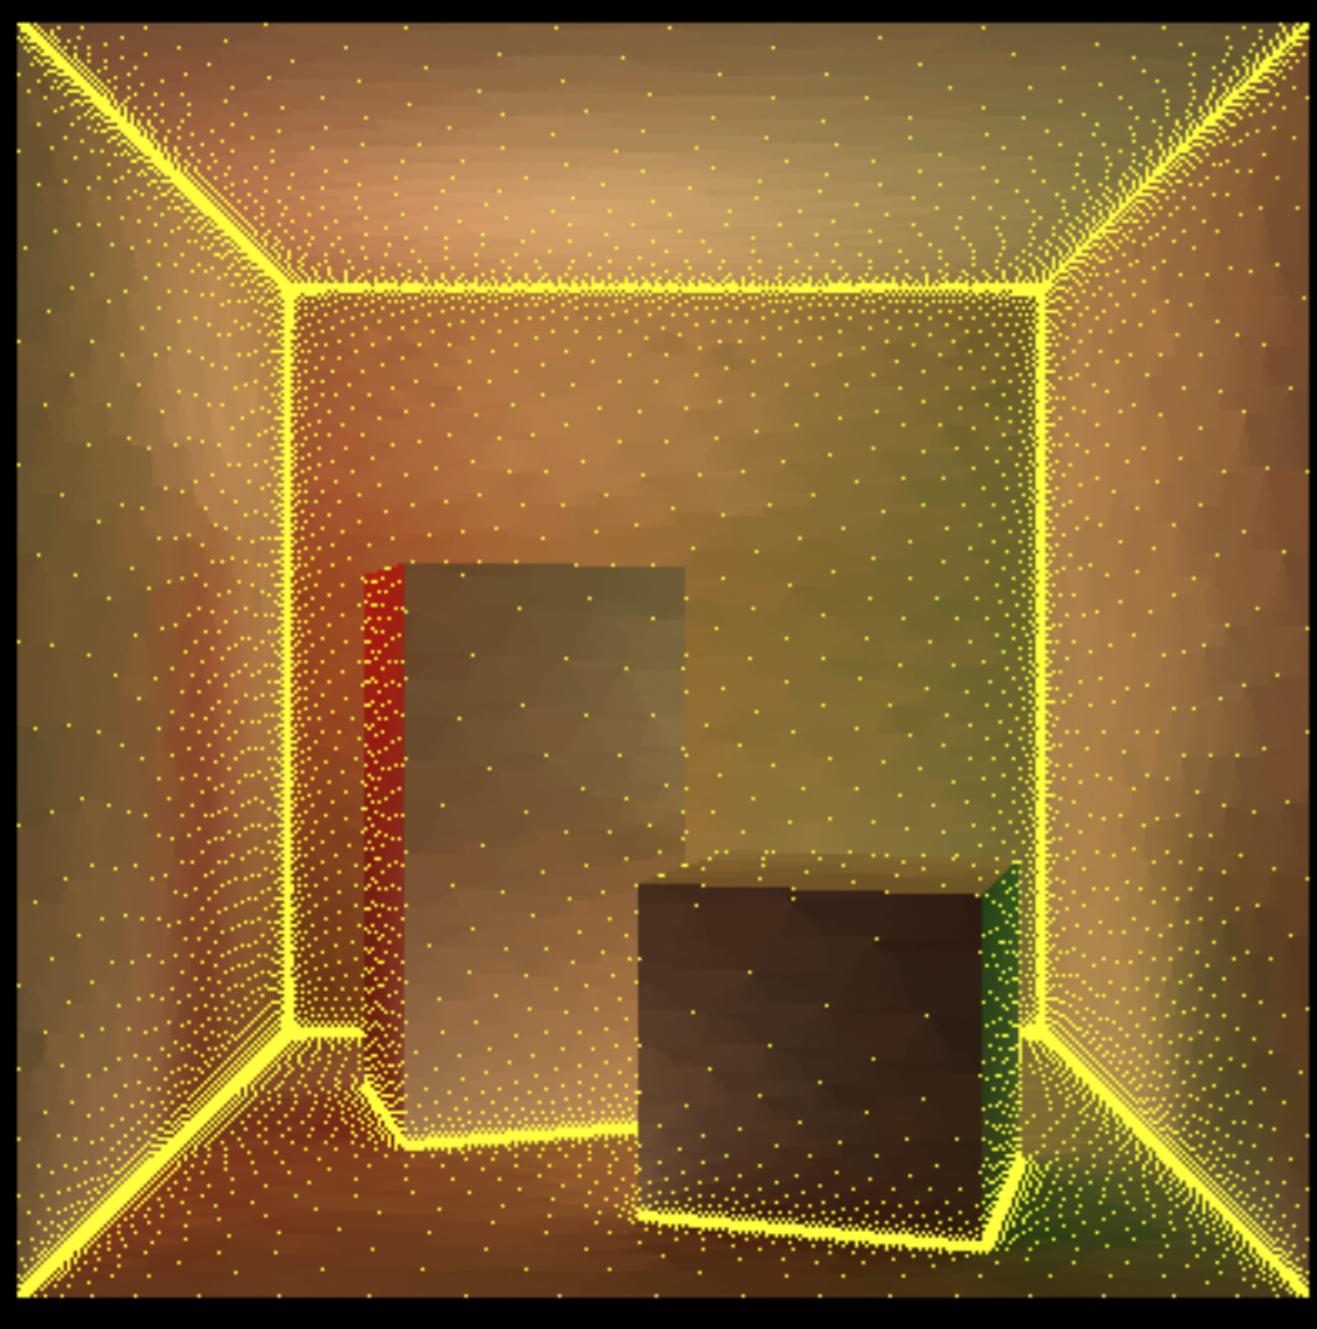


Over-Sampling/Clumping

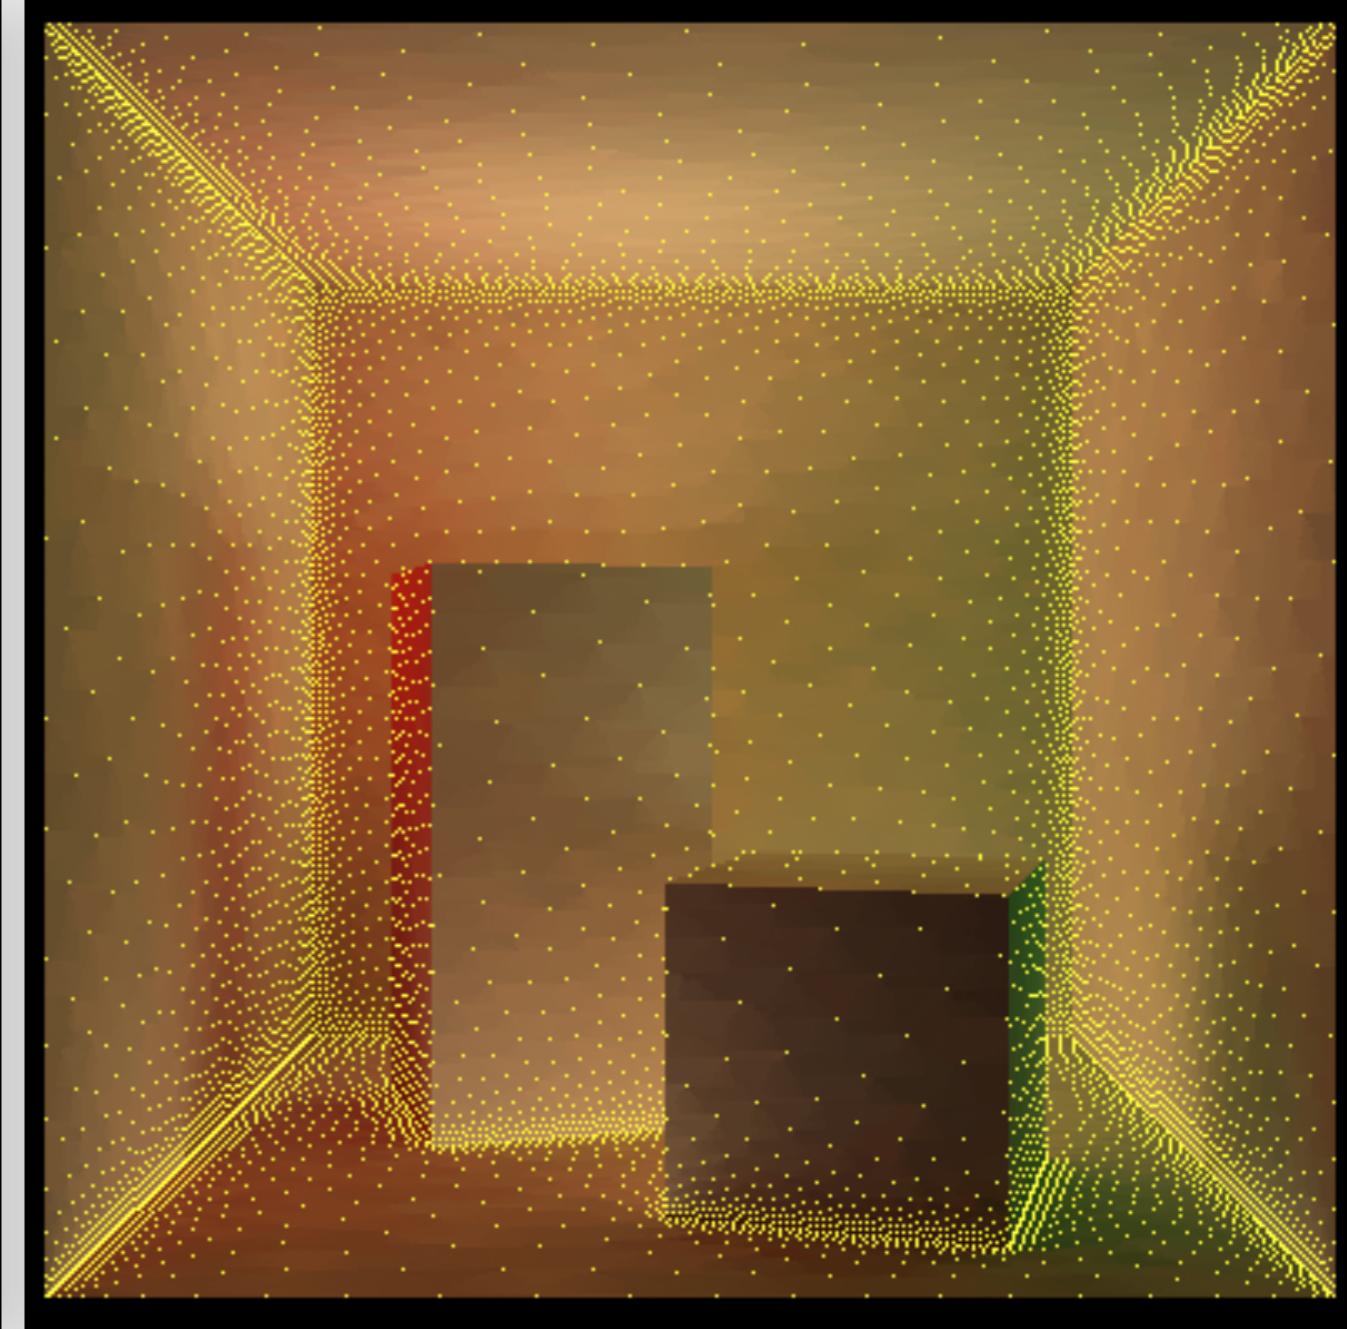


split-sphere metric

Over-Sampling/Clumping

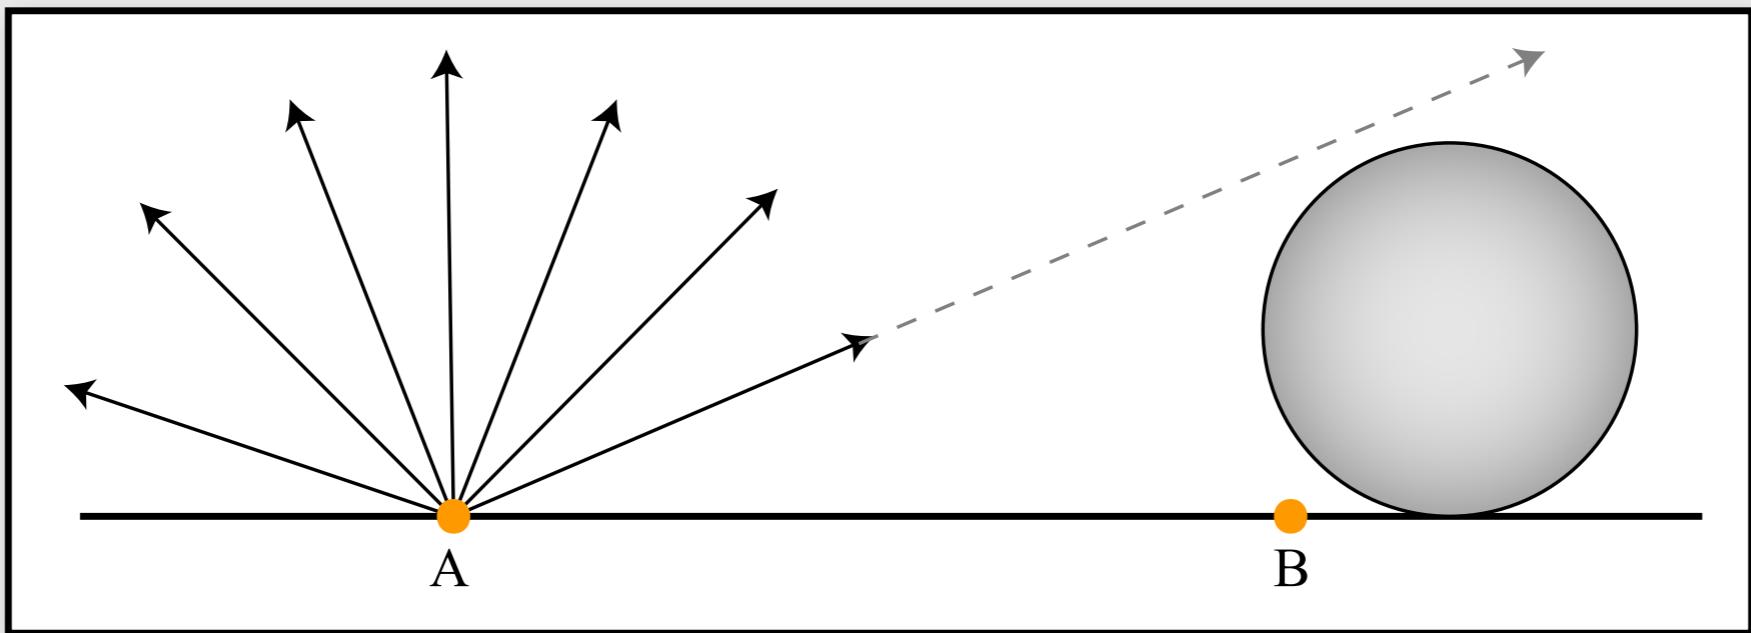


split-sphere metric

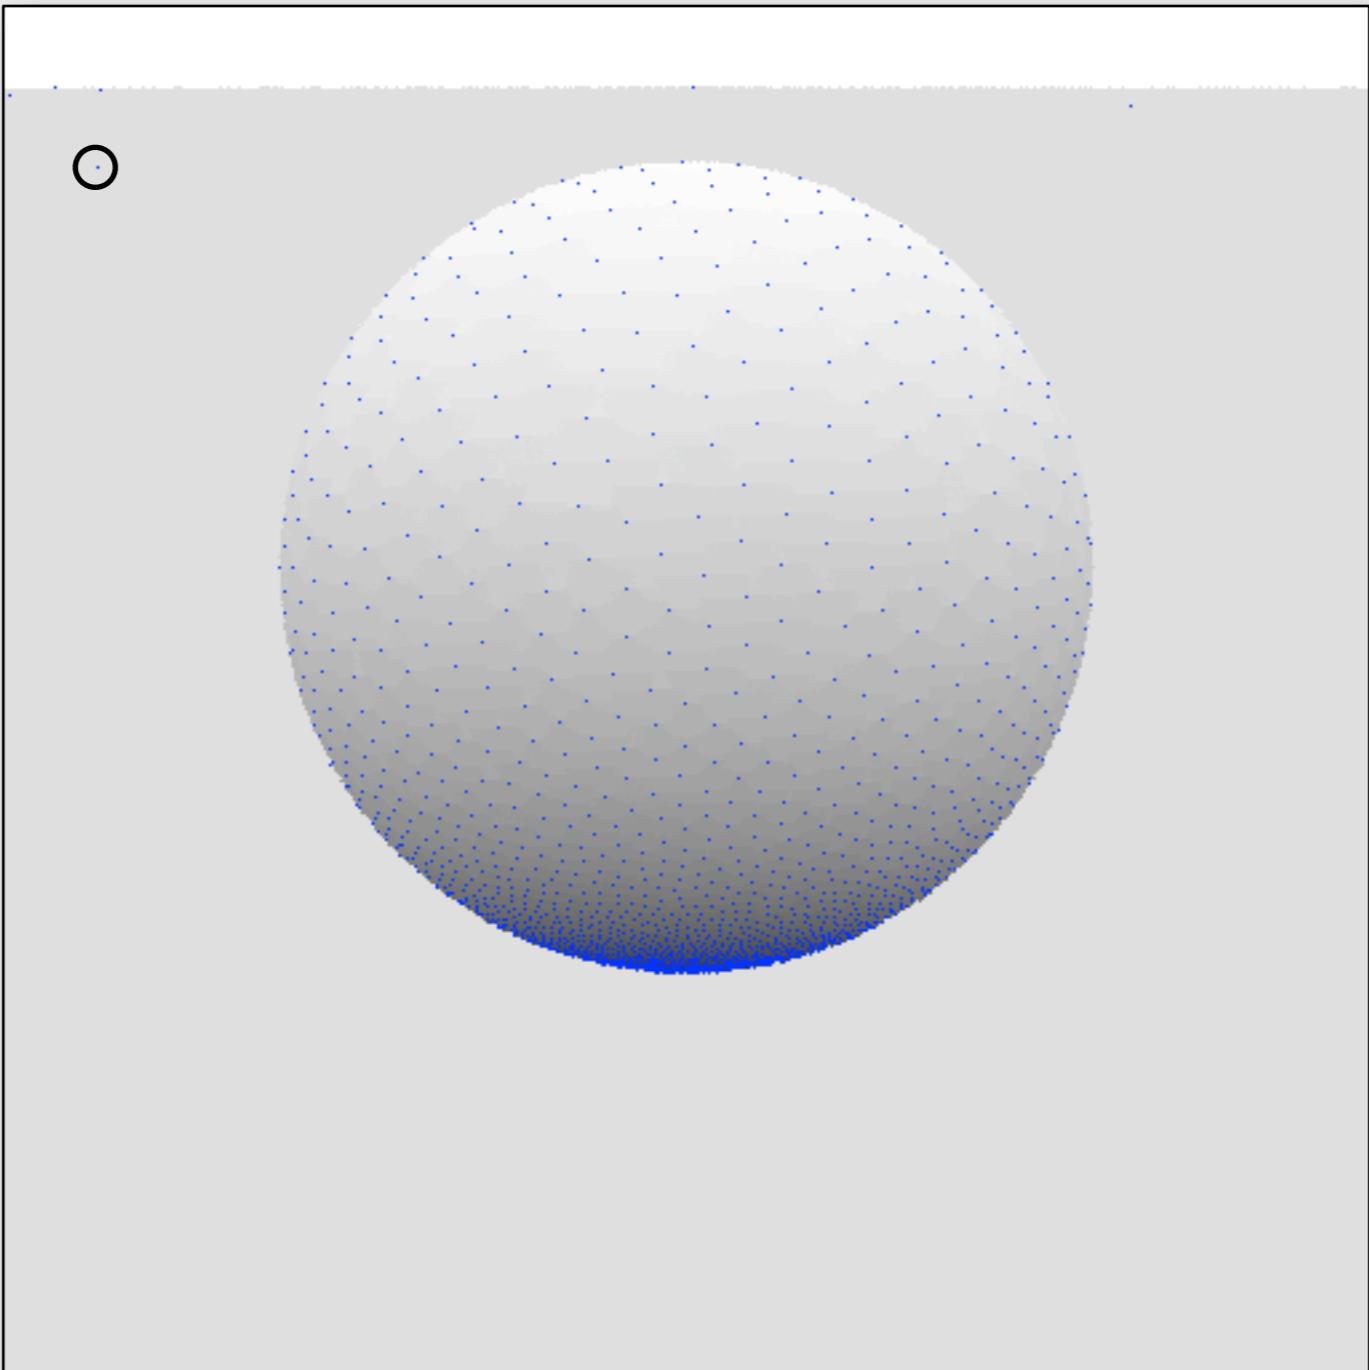


with minimum radius clamping

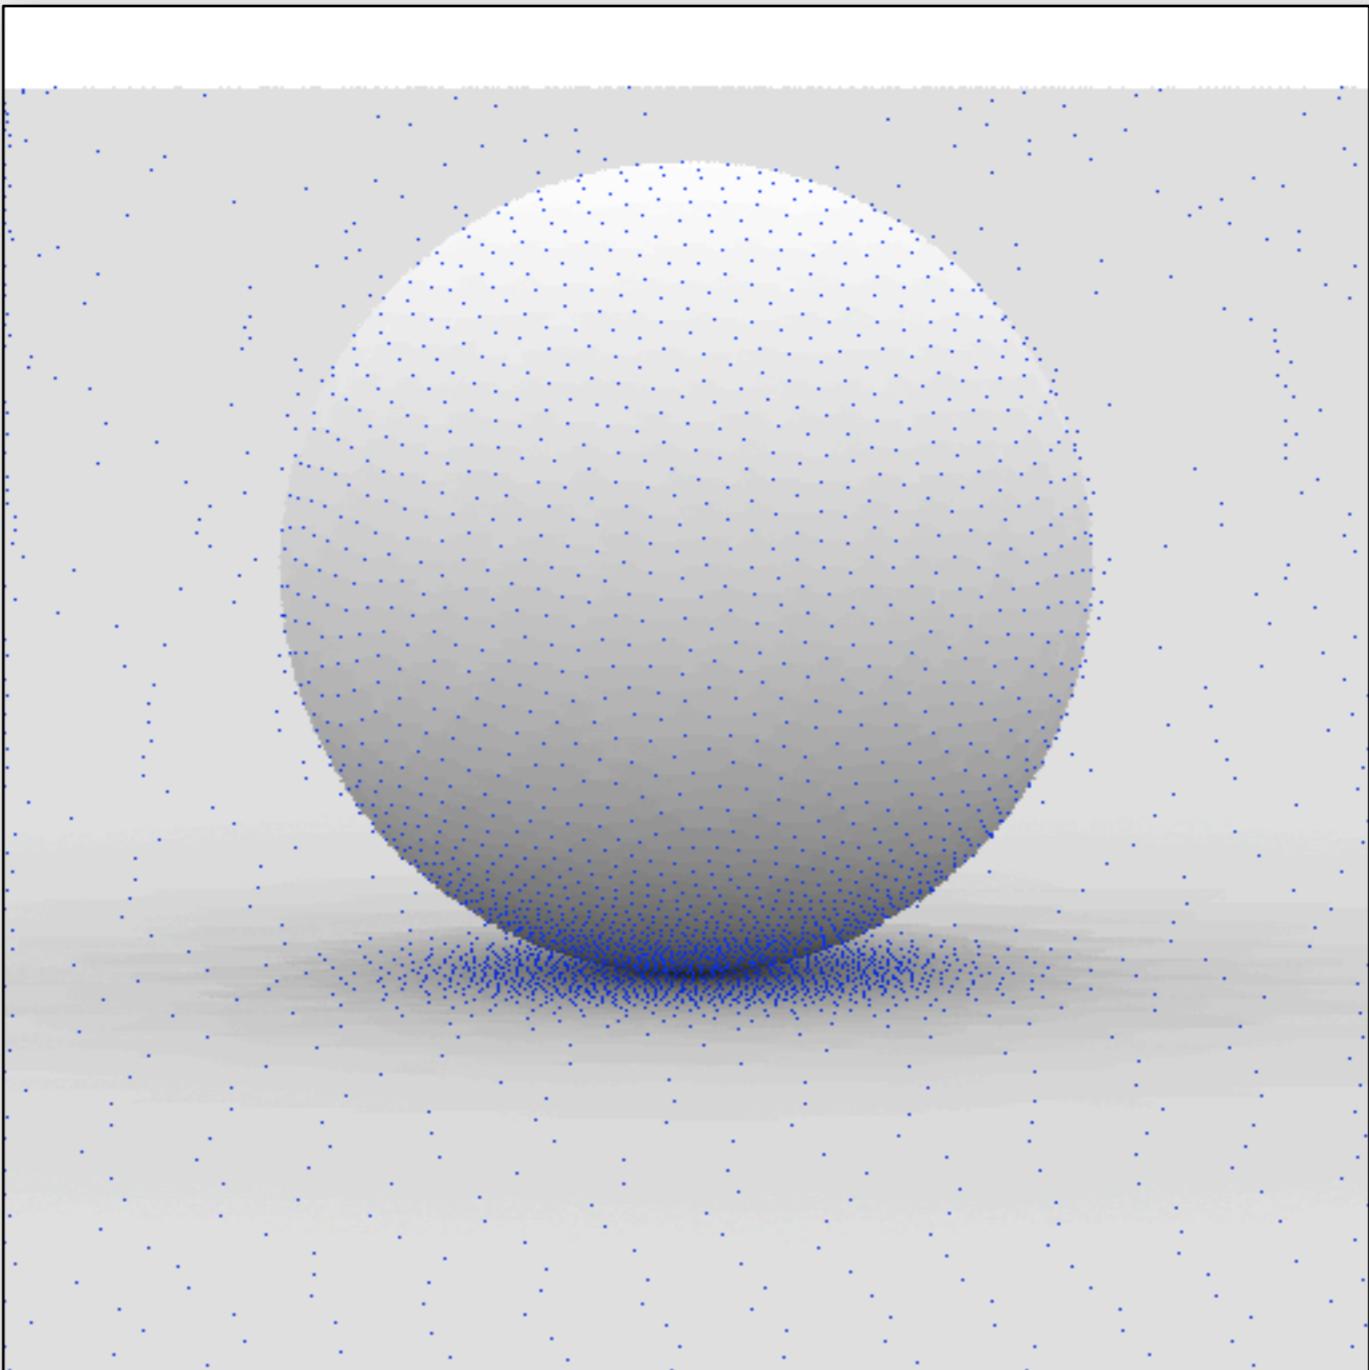
Under-Sampling



Under-Sampling



Under-Sampling



Split-Sphere Heuristic

- In practice, must also specify:
 - Min/max screen-space & world-space radii

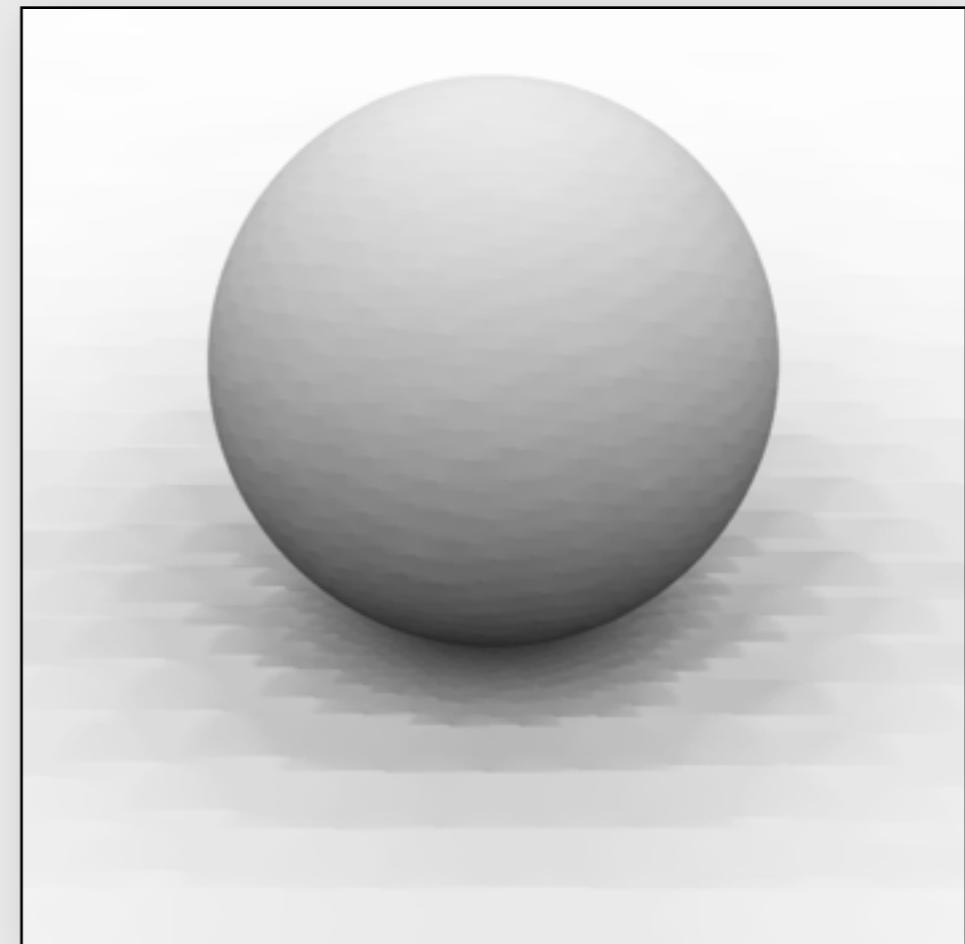
Irradiance Caching

- Pros:
 - Independent of resolution.
 - Computation amortized across many pixels
 - Concentrates computation in visible regions where illumination changes rapidly

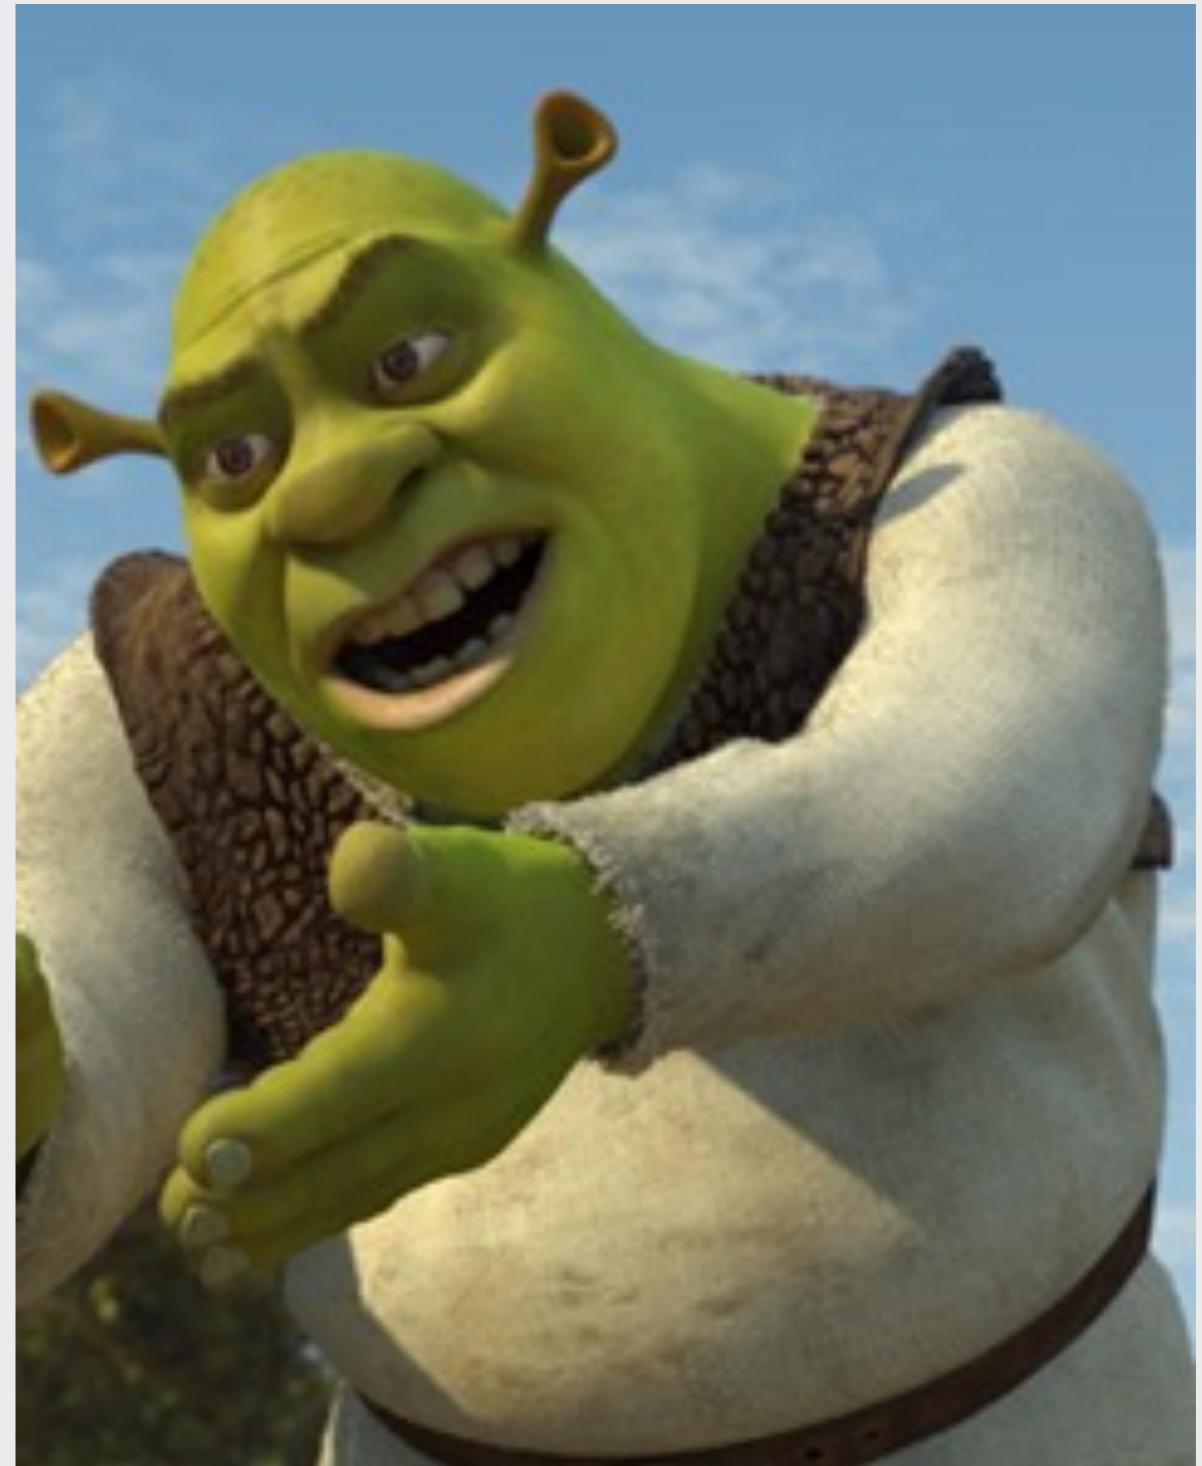


Irradiance Caching

- Cons:
 - Interpolation/extrapolation can introduce visible artifacts
 - Valid radius metric not always robust
 - Limited to Lambertian (matte) surfaces



Use in the film industry



Improvements/Extensions

- Many extensions:
 - Ward and Heckbert '92 - better interpolation
 - Křivánek et al. '05a, '05b - glossy surfaces
 - Jarosz et al. '08 - participating media
 - Jarosz et al. '12 - irradiance Hessians
 - Schwarzhaupt et al. '12 - better error control
 - ...

Irradiance Caching Artifacts

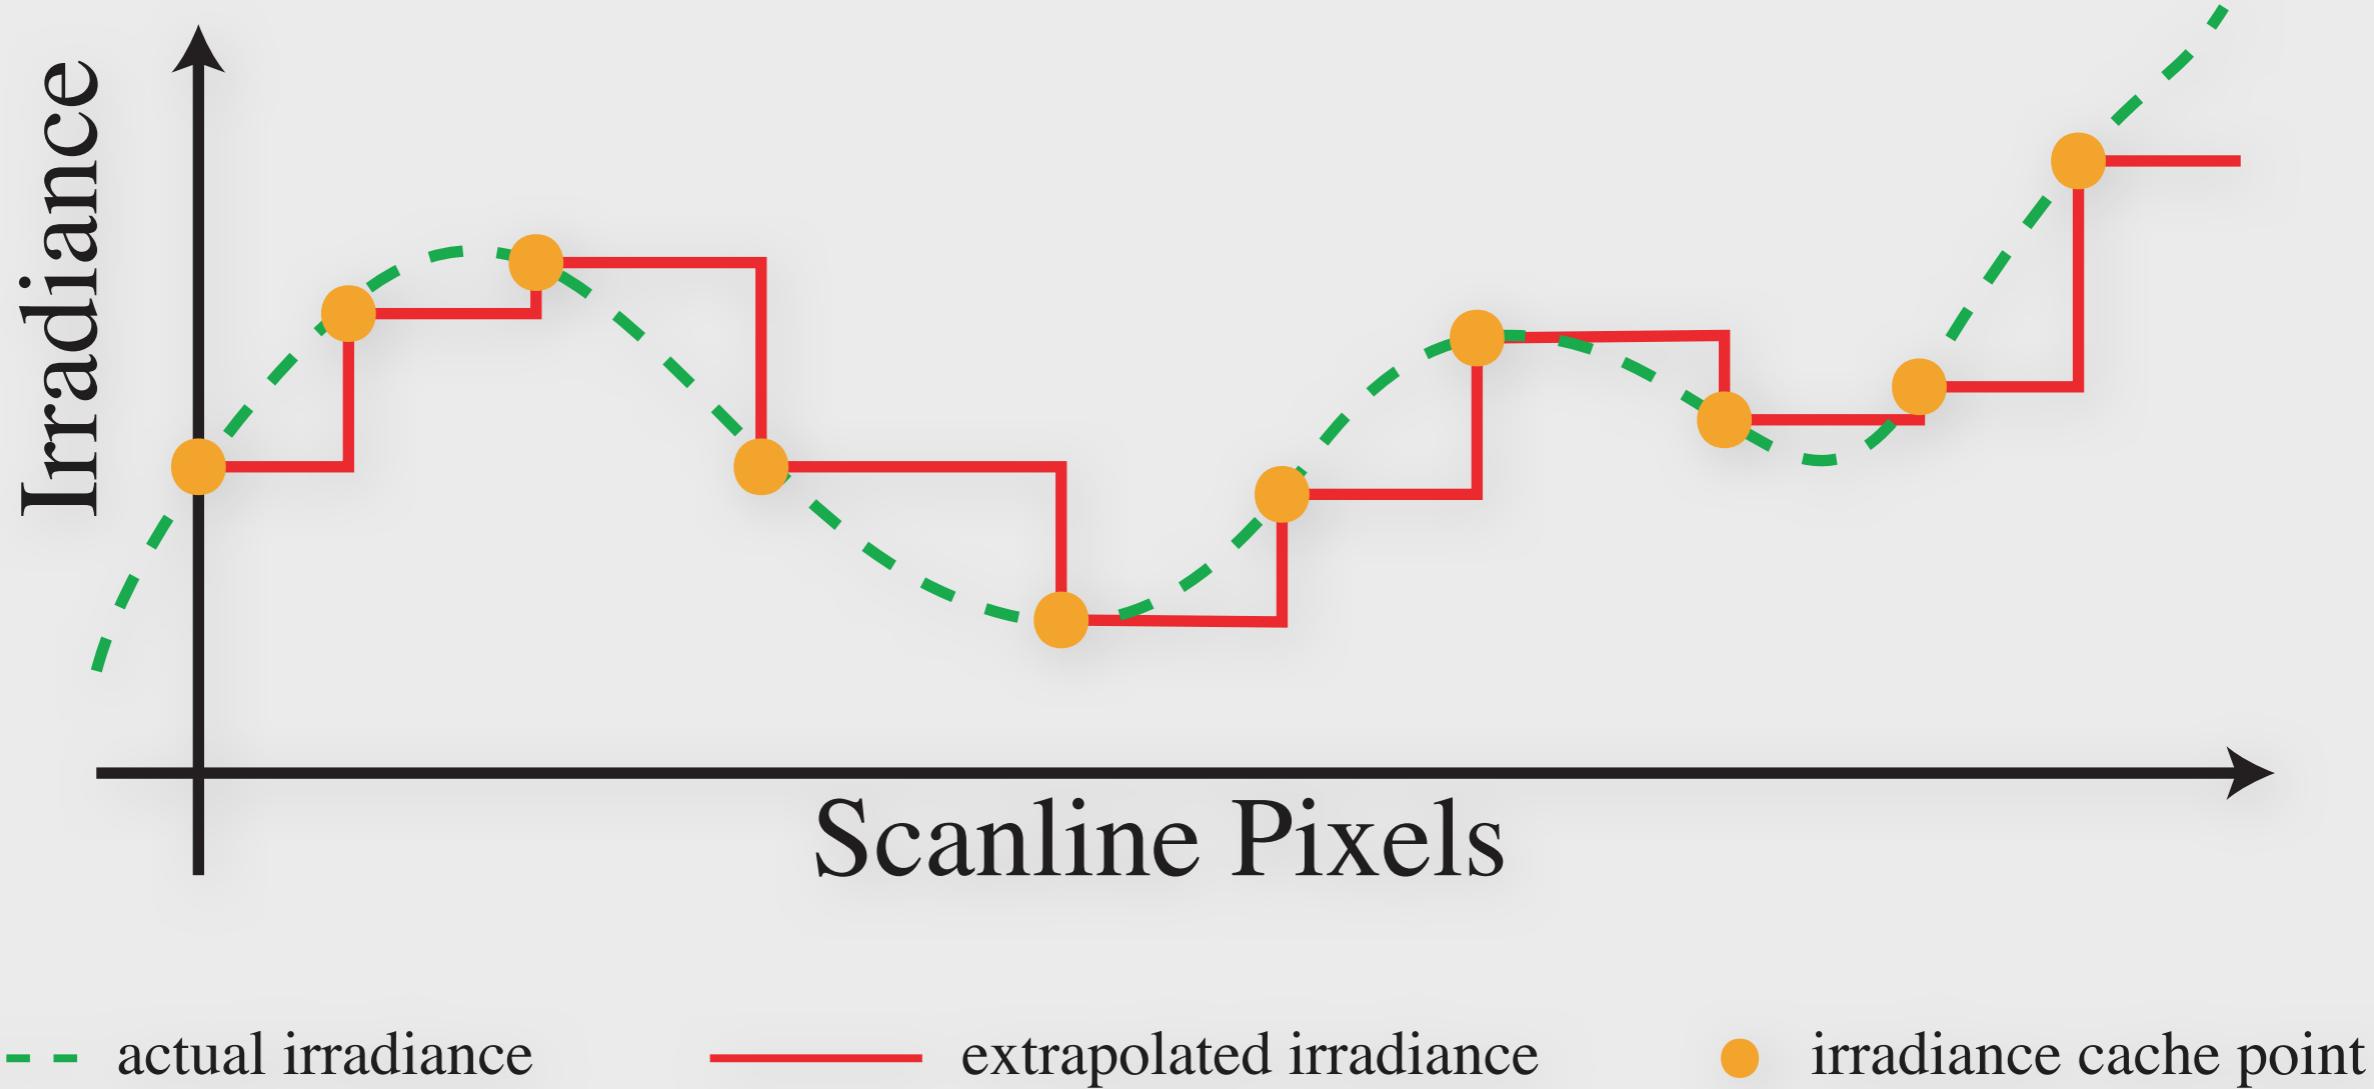


Irradiance gradients

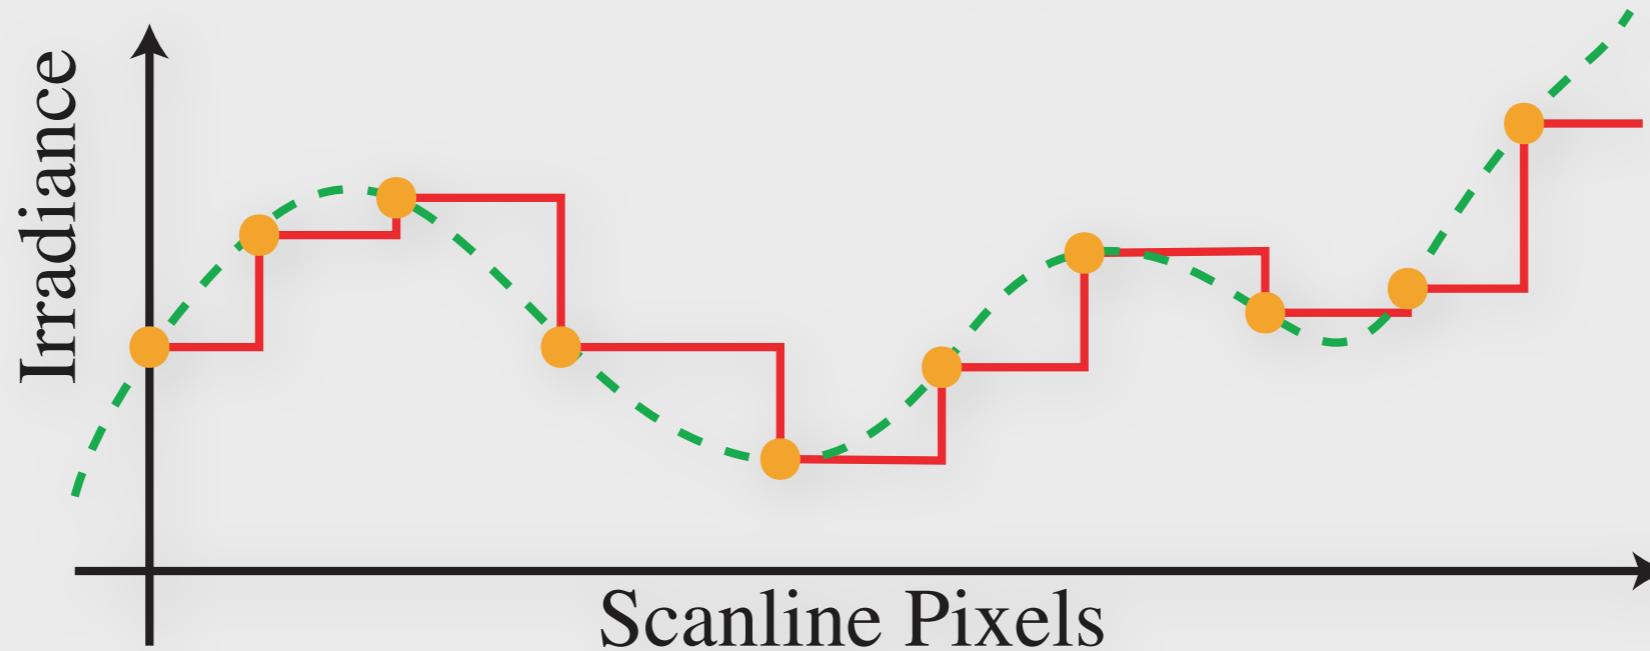
- Improve interpolation/extrapolation quality using gradients
- Irradiance Gradients [Ward and Heckbert 1992]
 - Estimate an actual derivative to the irradiance
 - Apply this derivative to the weighted average



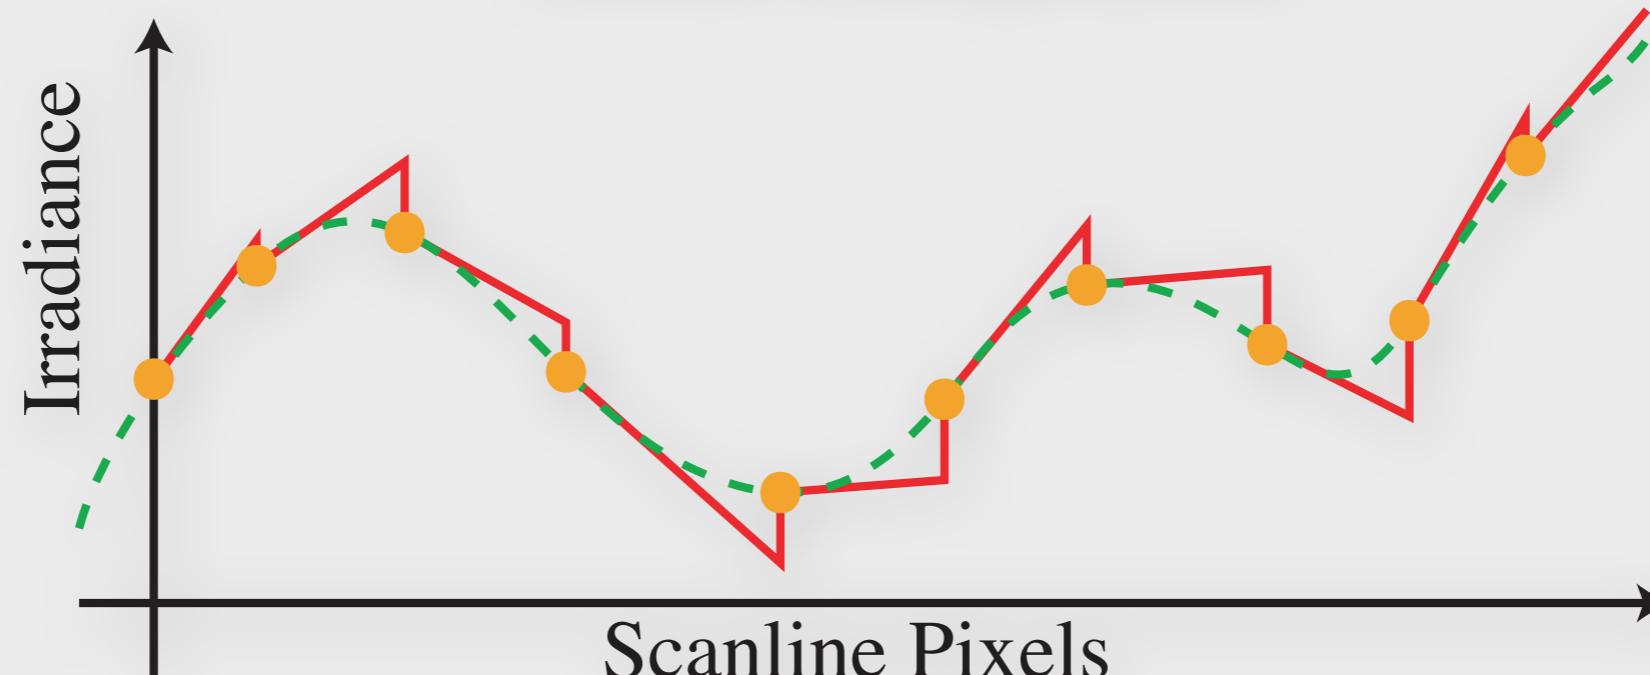
Constant Extrapolation



Constant Extrapolation



Linear Extrapolation

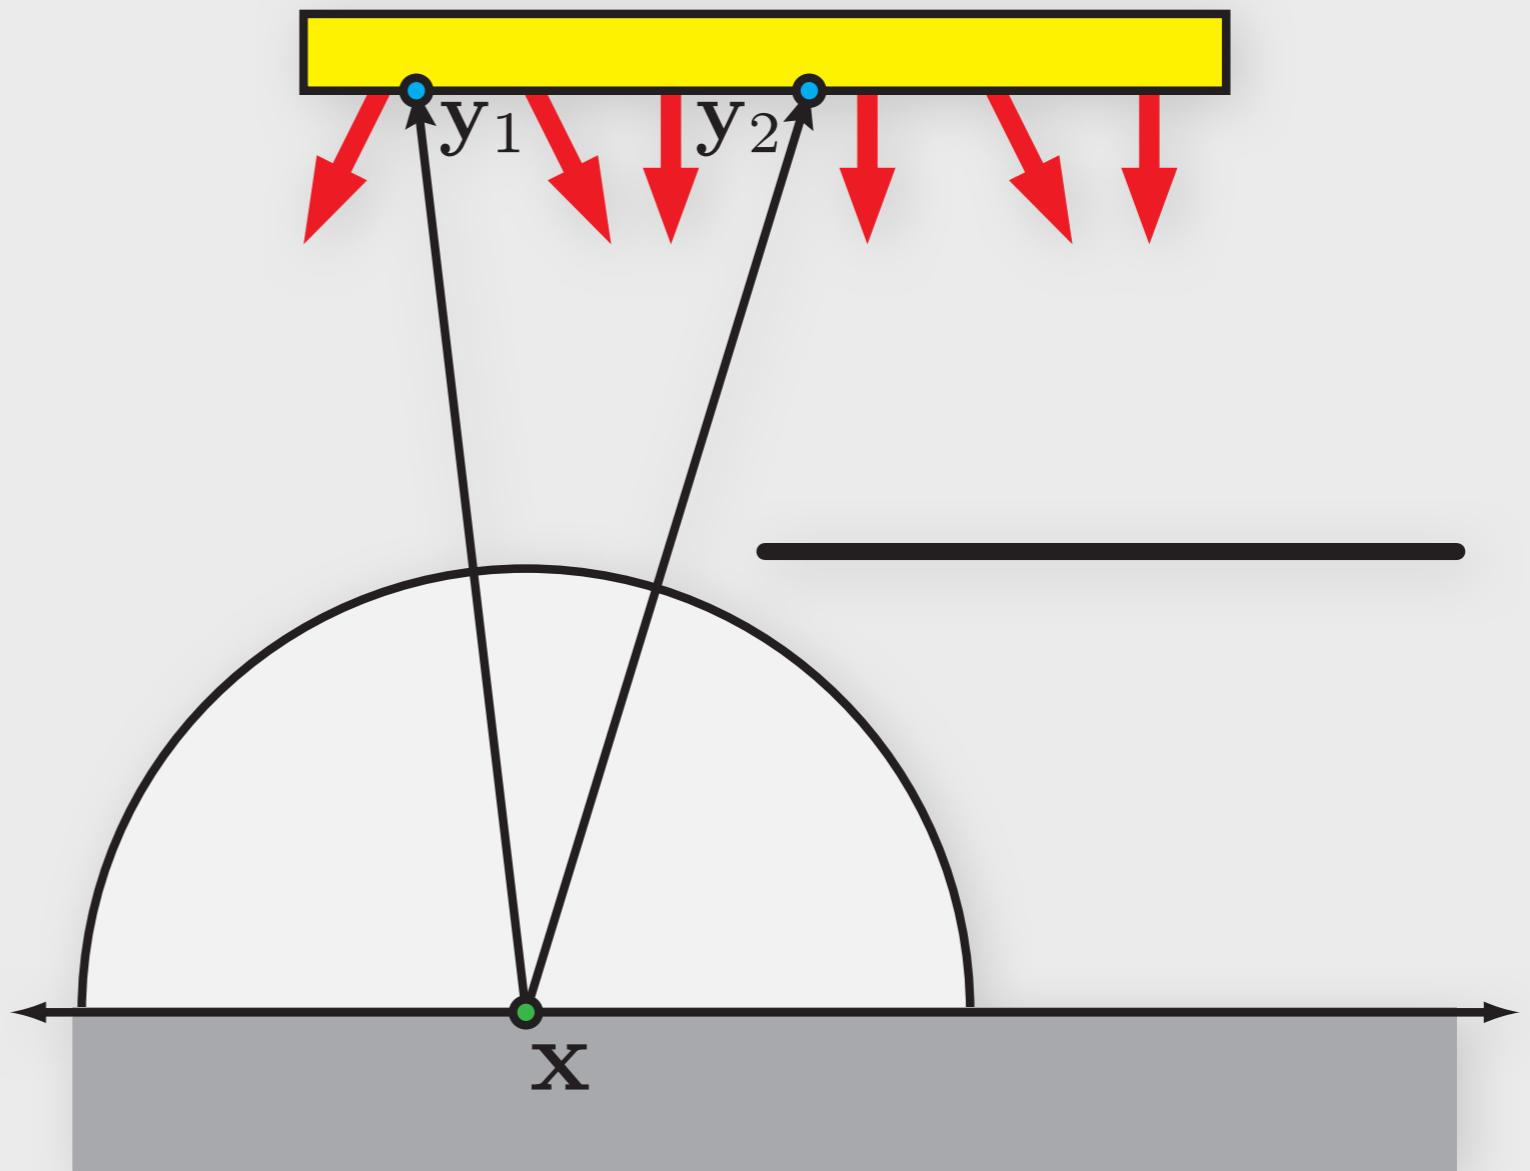


--- actual irradiance

— extrapolated irradiance

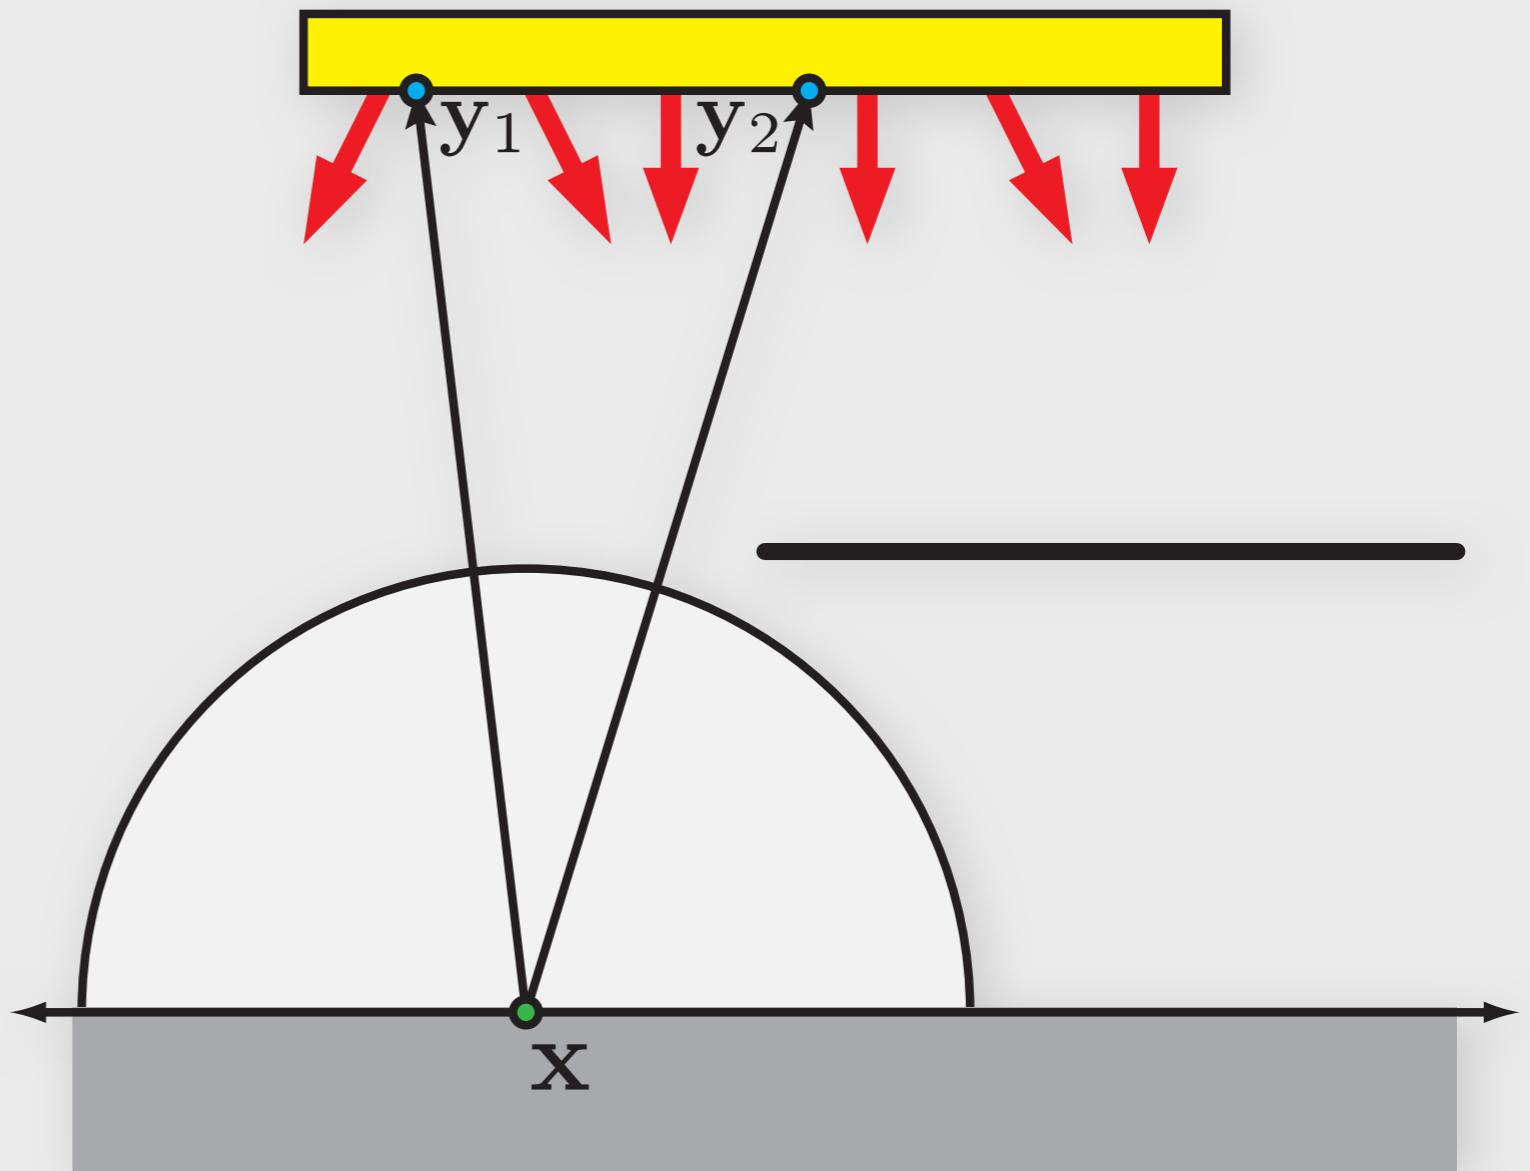
● irradiance cache point

Gradients (surface-area formulation)



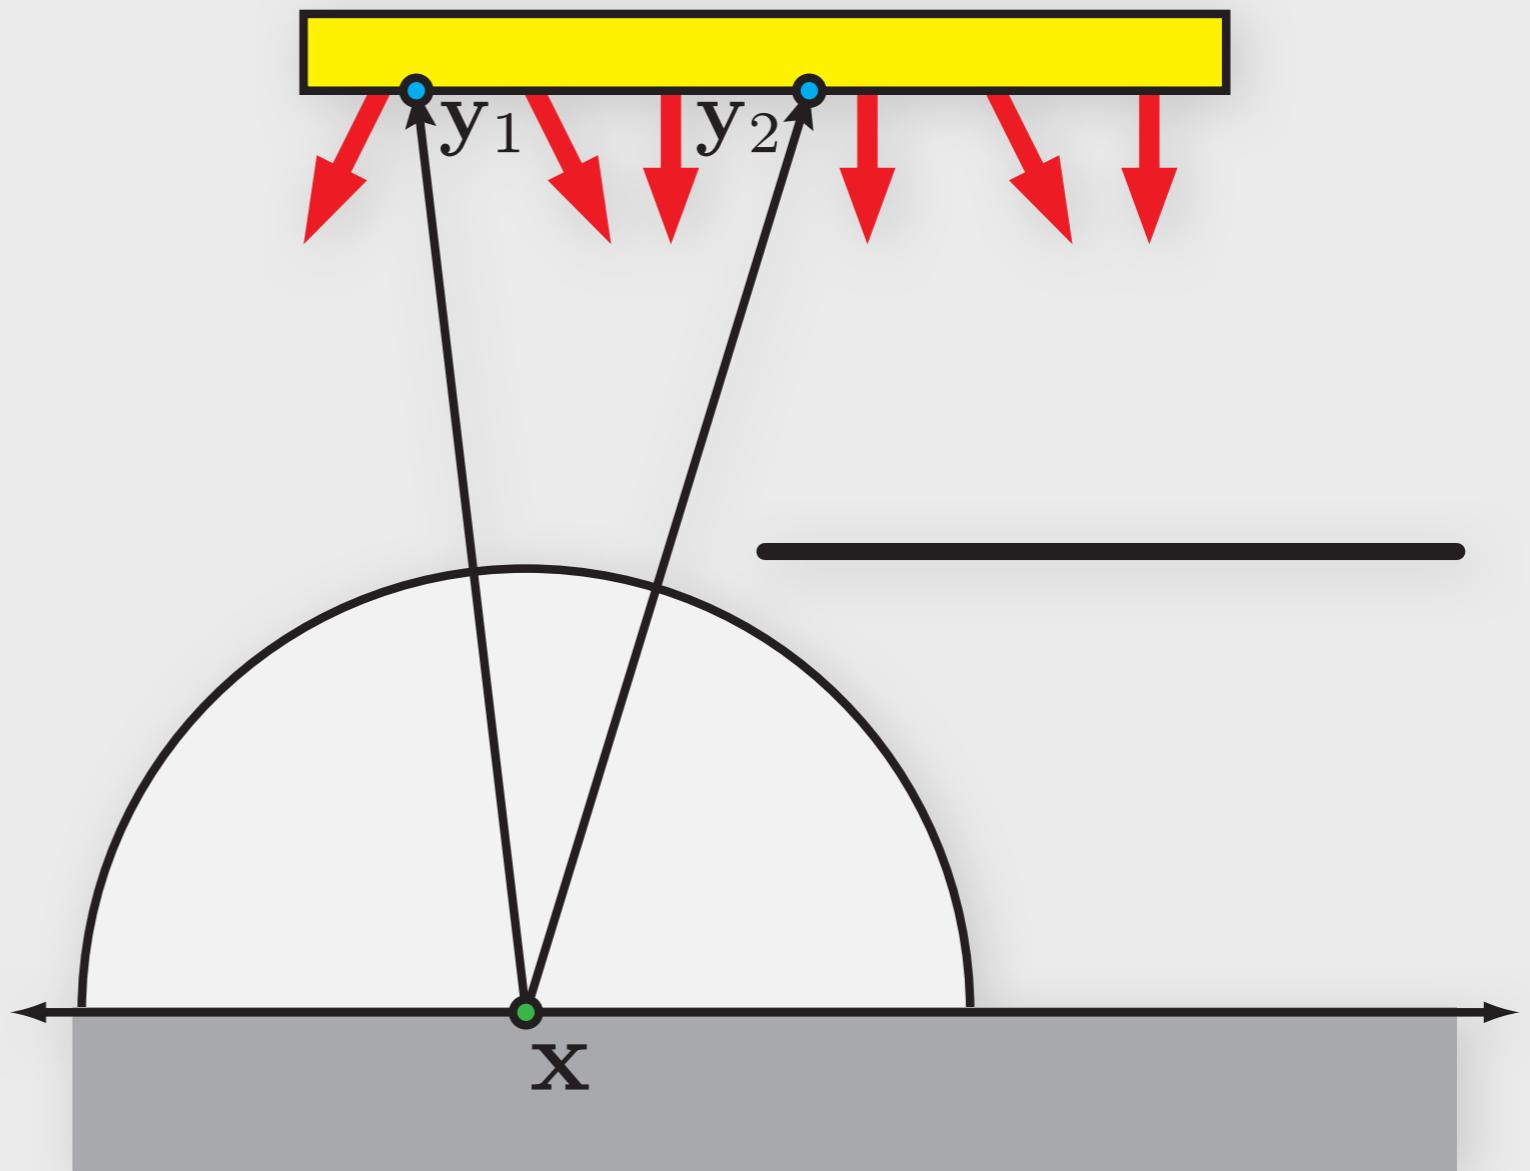
$$\nabla E(\mathbf{x}) = \nabla \int_A L(\mathbf{x} \leftarrow \mathbf{y}) V(\mathbf{x}, \mathbf{y}) G(\mathbf{x}, \mathbf{y}) d\mathbf{y}$$

Gradients (surface-area formulation)



$$\nabla E(\mathbf{x}) = \int_A \cancel{\nabla L V G} + \cancel{L \nabla V G} + L V \nabla G \, d\mathbf{y}$$

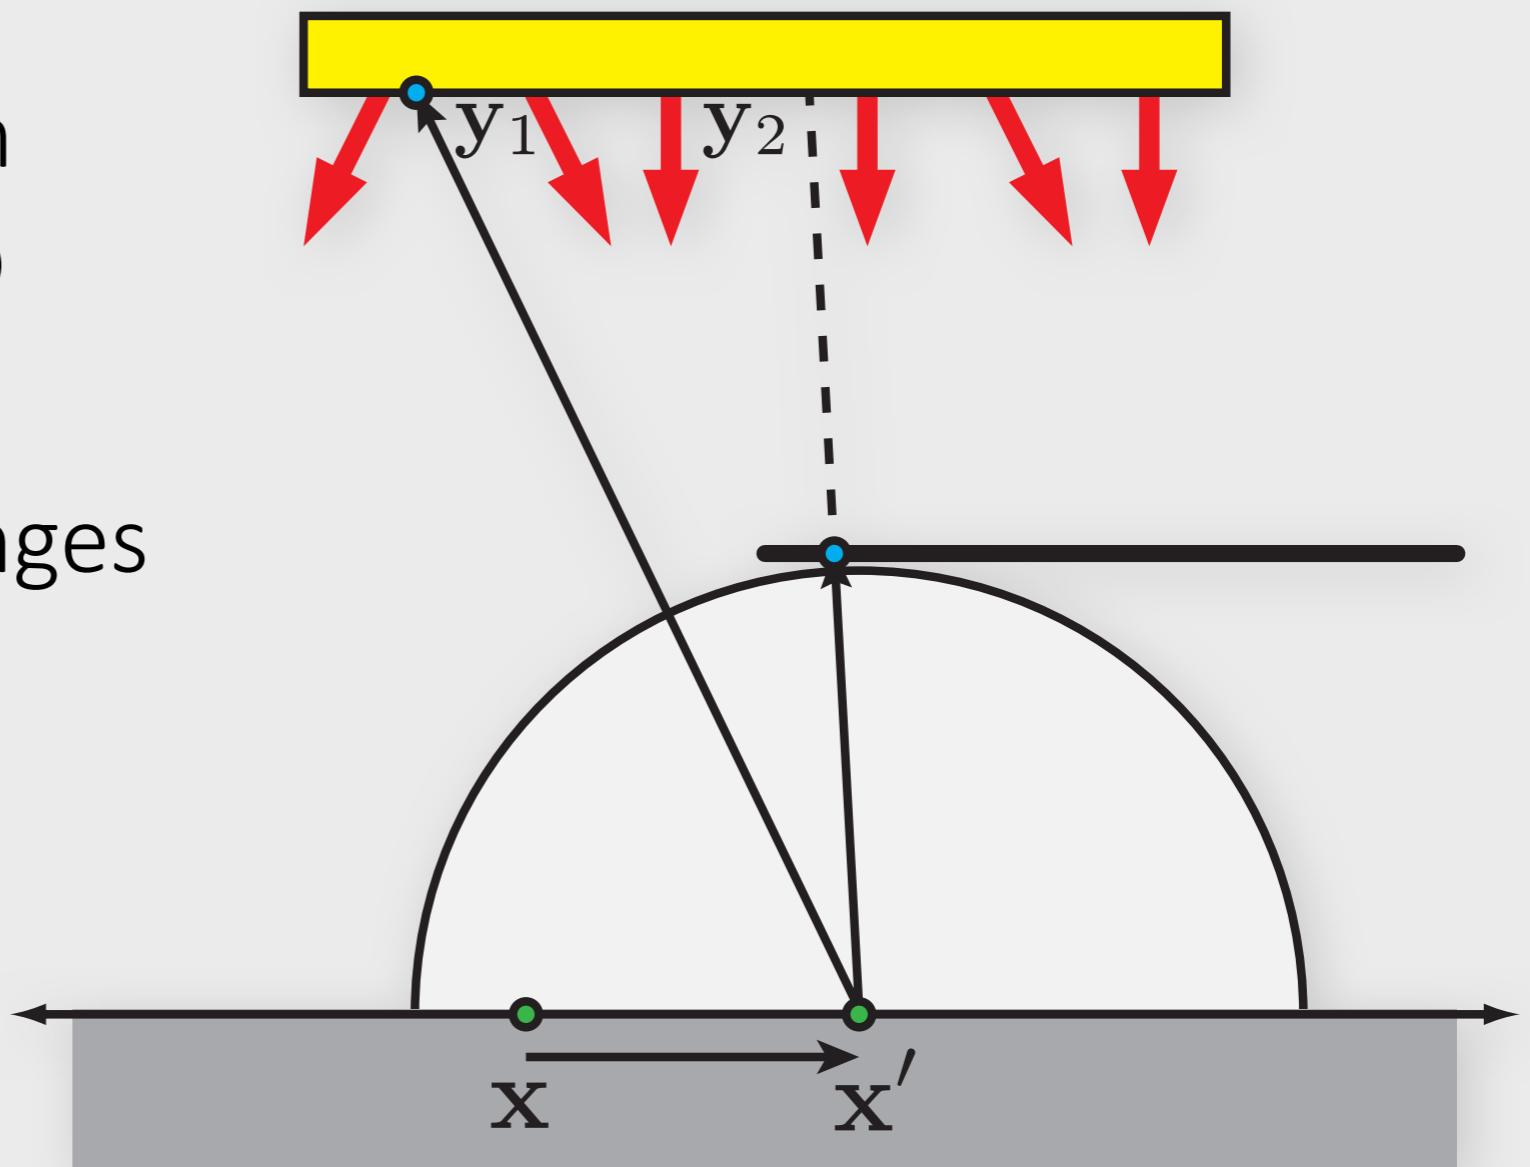
Gradients (surface-area formulation)



$$\nabla E(\mathbf{x}) \approx \int_A L(\mathbf{x} \leftarrow \mathbf{y}) V(\mathbf{x}, \mathbf{y}) \nabla G(\mathbf{x}, \mathbf{y}) d\mathbf{y}$$

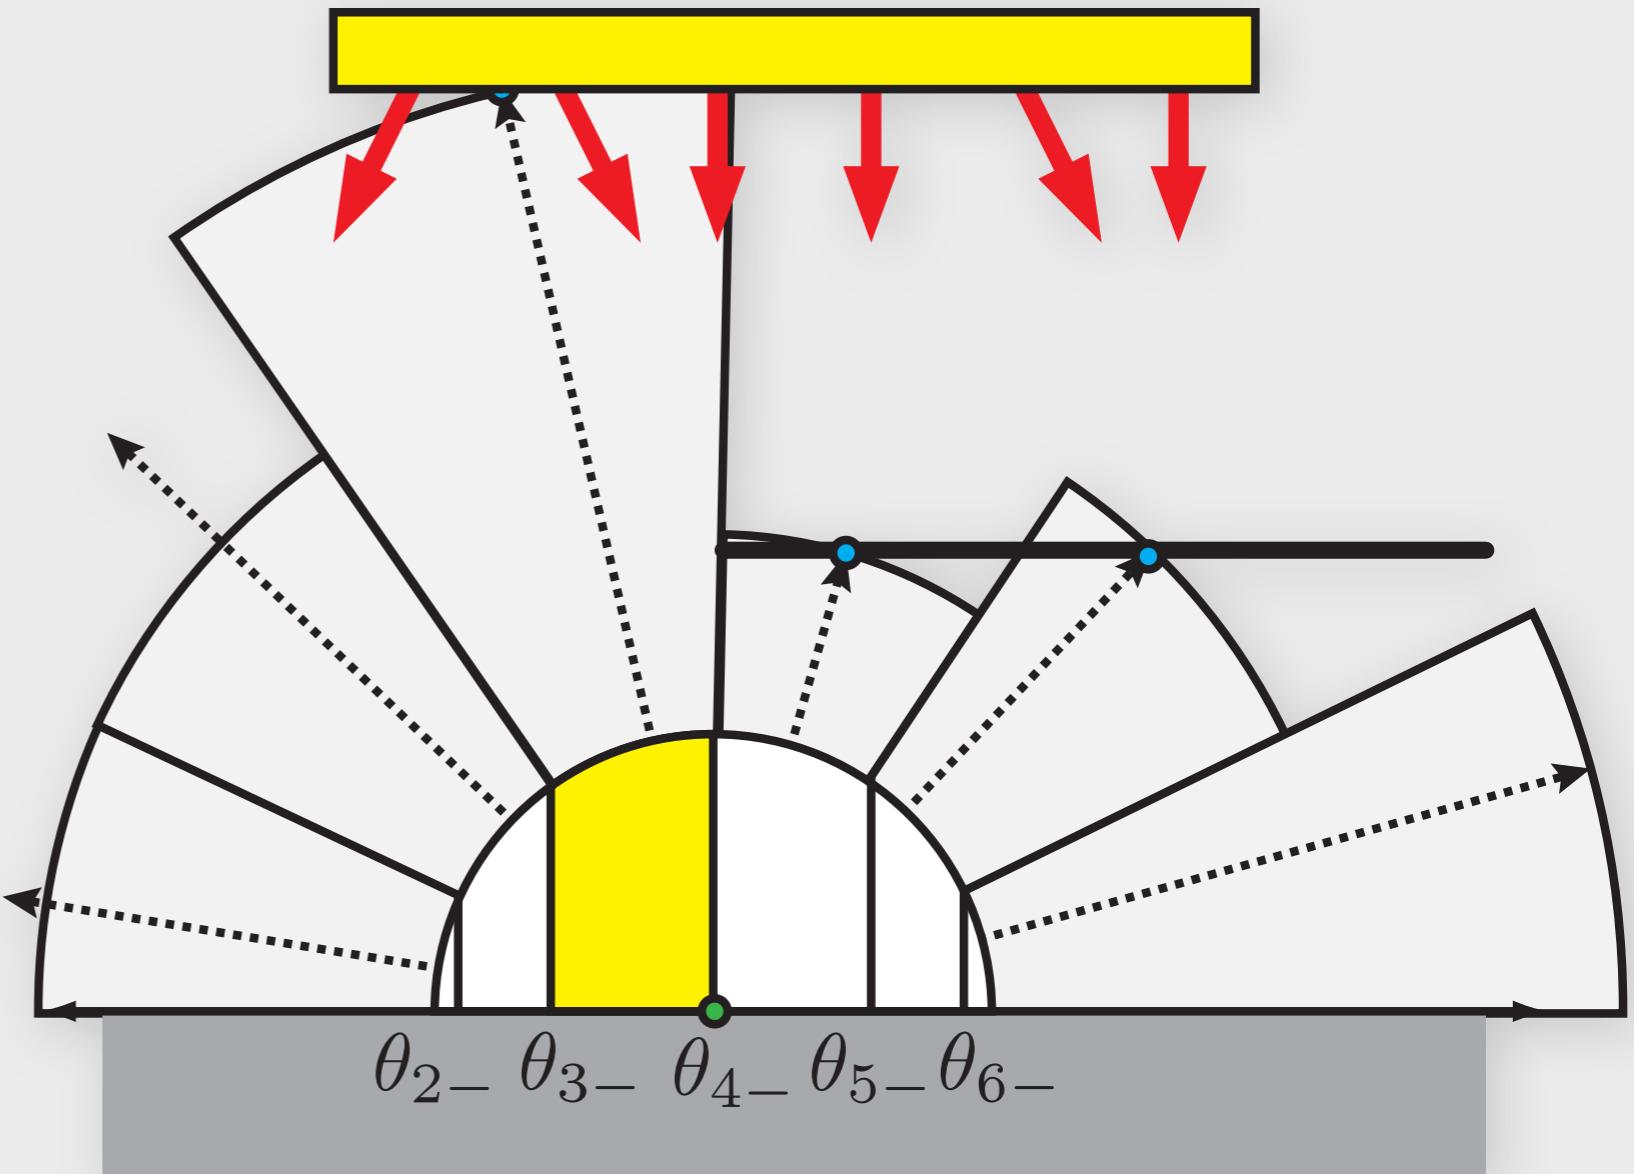
Gradients (surface-area formulation)

- Accounts for change in geometric relationship between x & y
- Ignores occlusion changes



$$\nabla E(\mathbf{x}) \approx \int_A L(\mathbf{x} \leftarrow \mathbf{y}) V(\mathbf{x}, \mathbf{y}) \nabla G(\mathbf{x}, \mathbf{y}) d\mathbf{y}$$

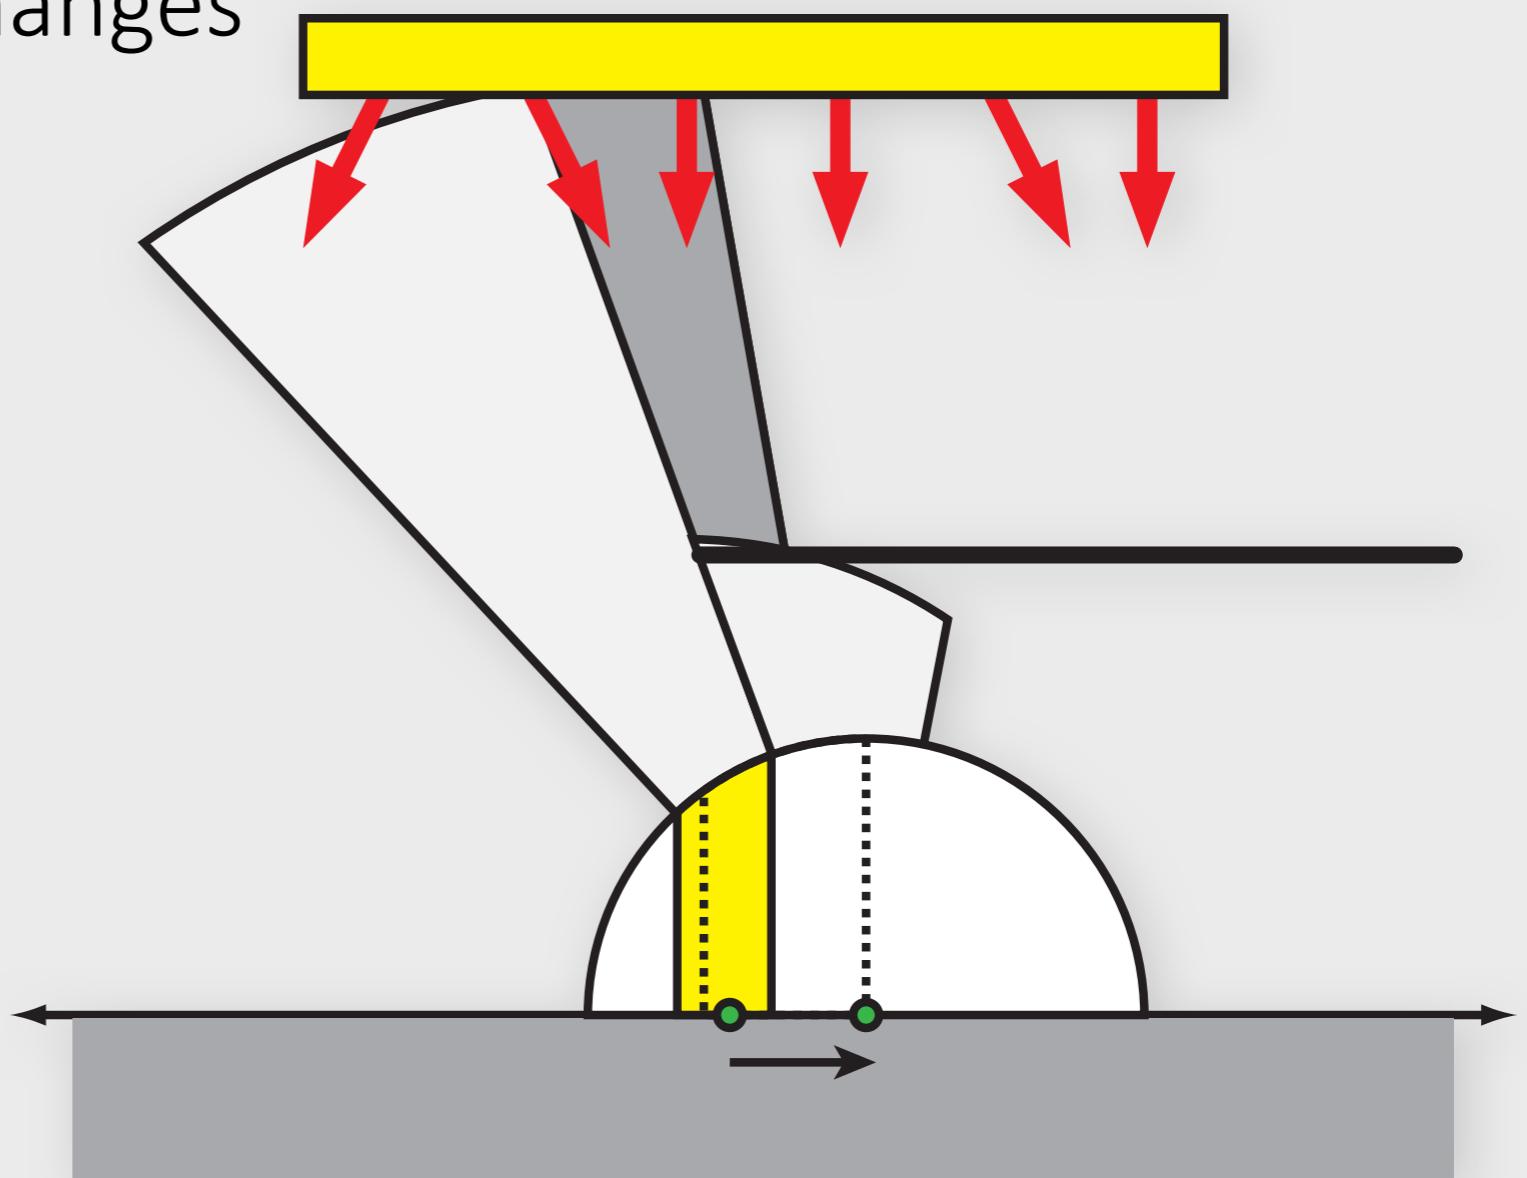
Gradients (stratified formulation)



Gradients (stratified formulation)

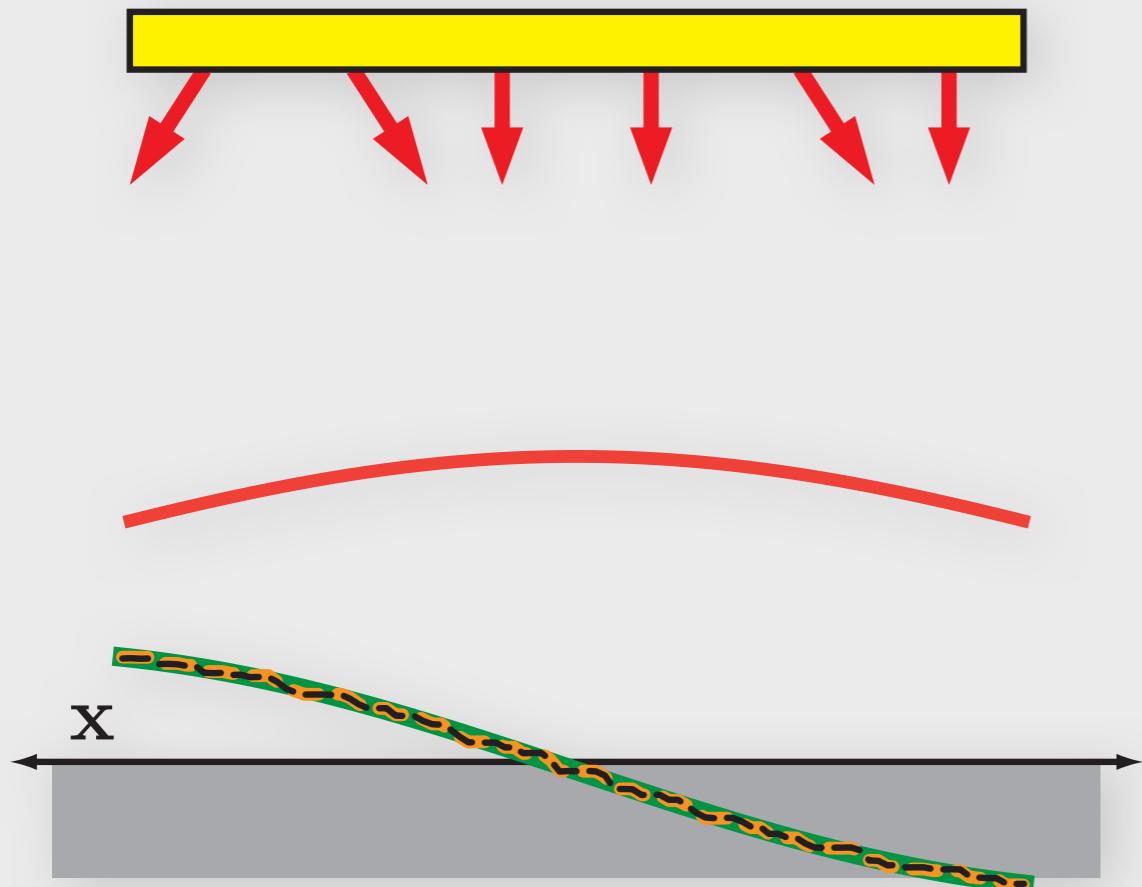
- Considers occlusion changes

Very
Important!

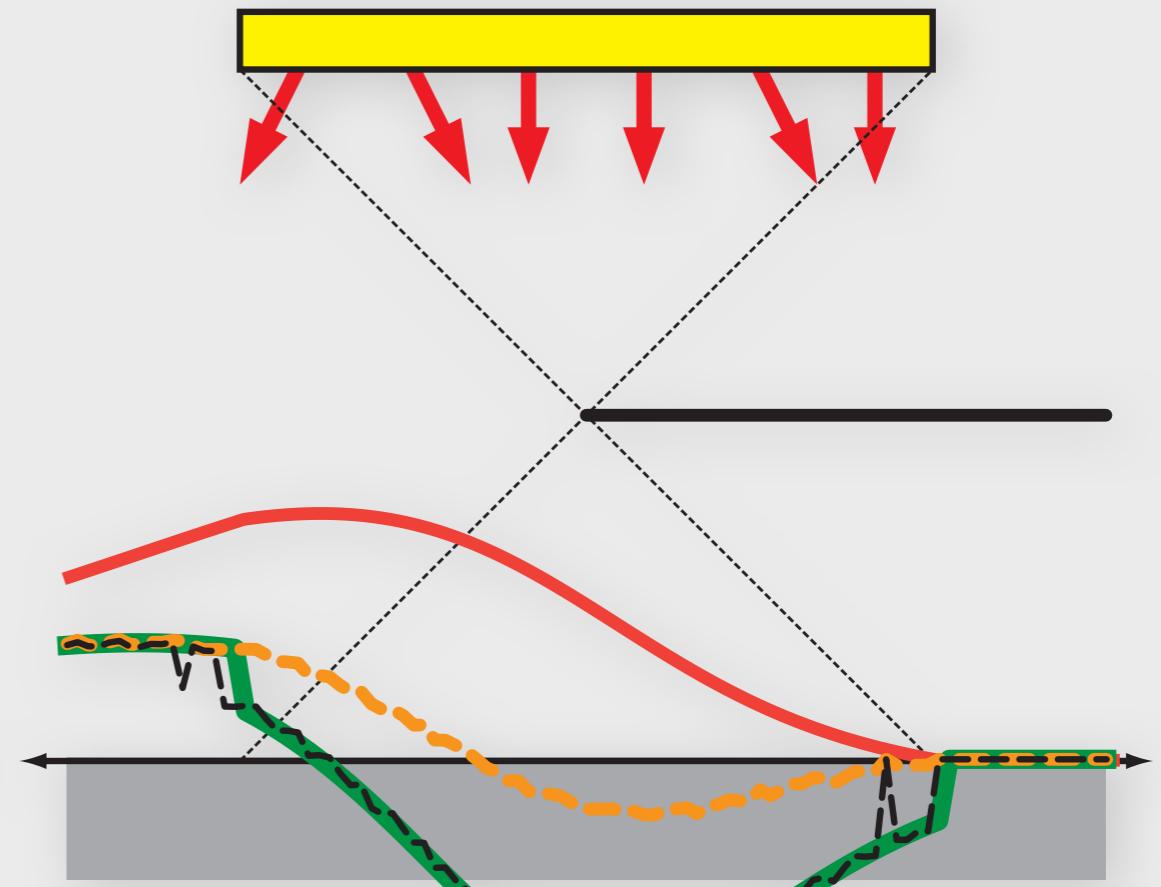


Irradiance gradients comparison

Without Occluder



With Occluder



Irradiance



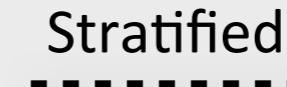
True Gradient



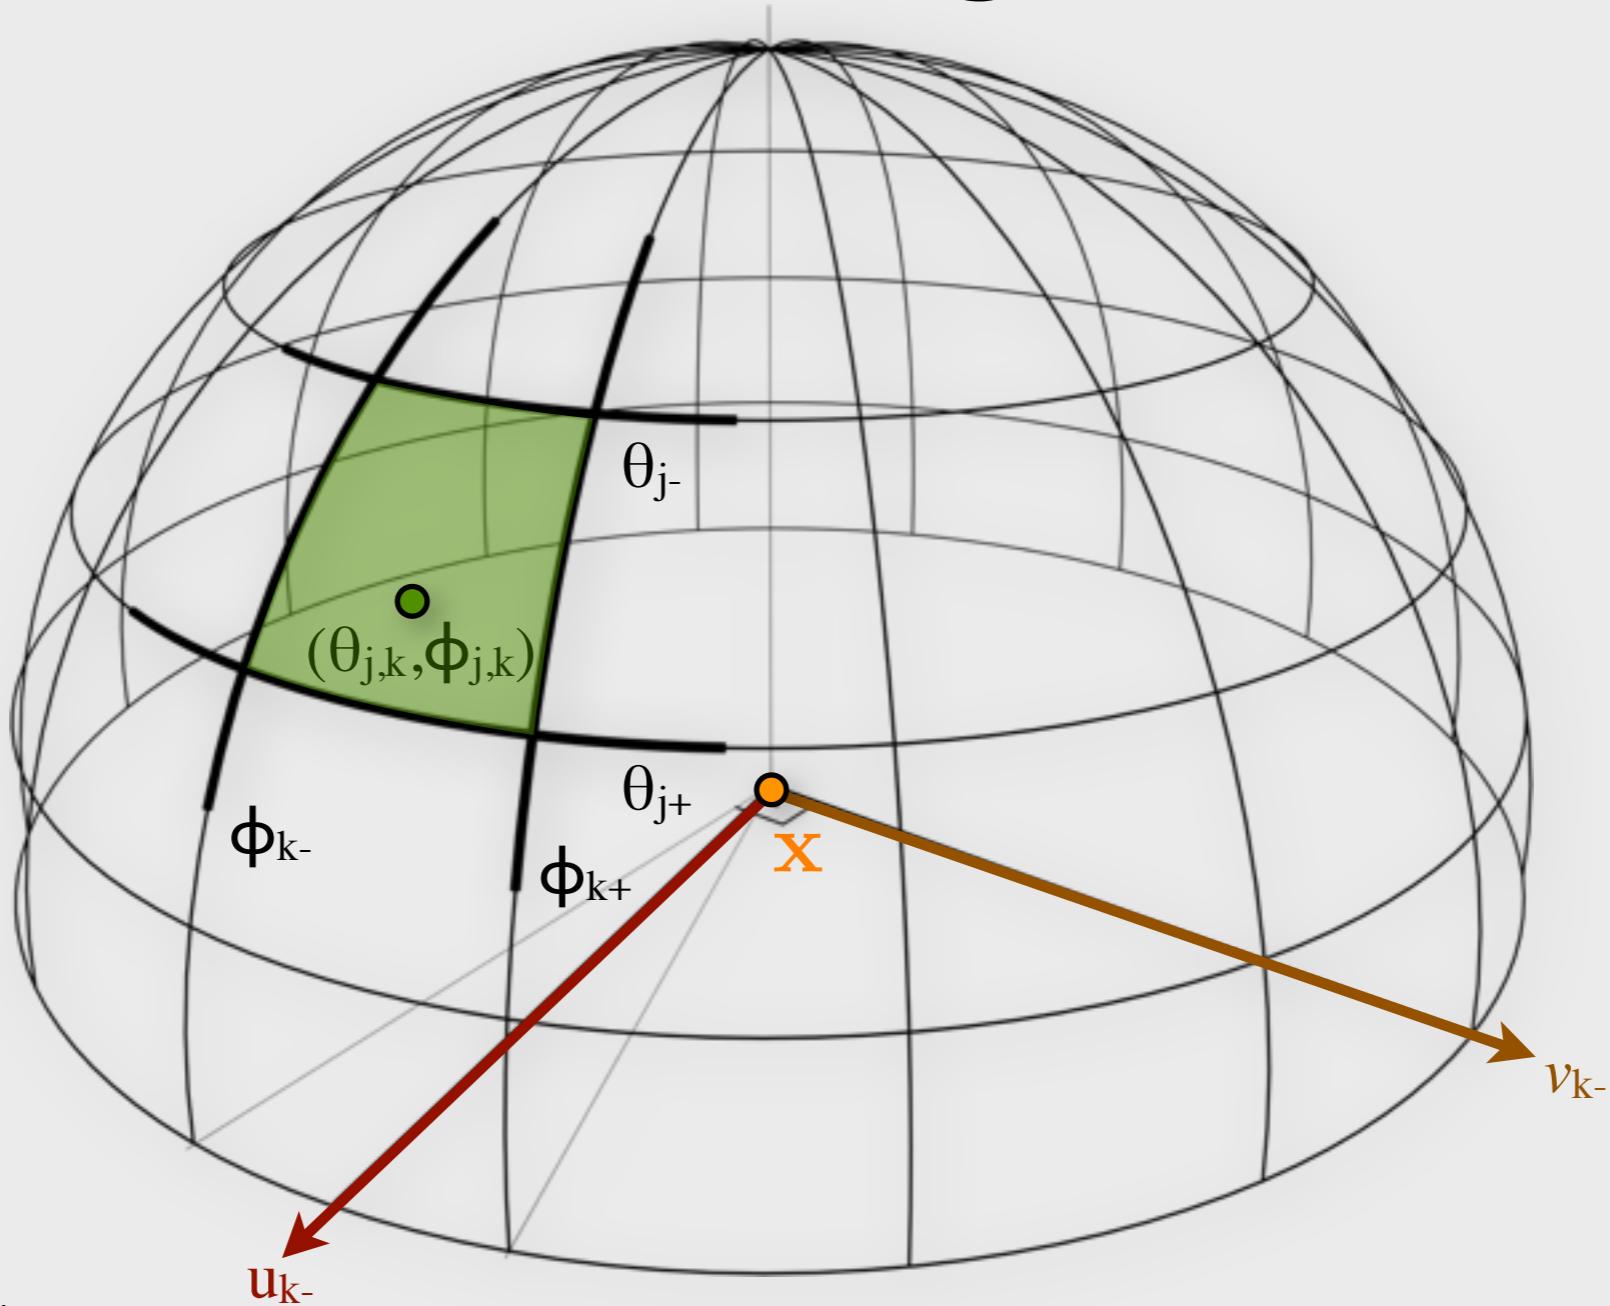
Surface-area



Stratified



Stratified irradiance gradient



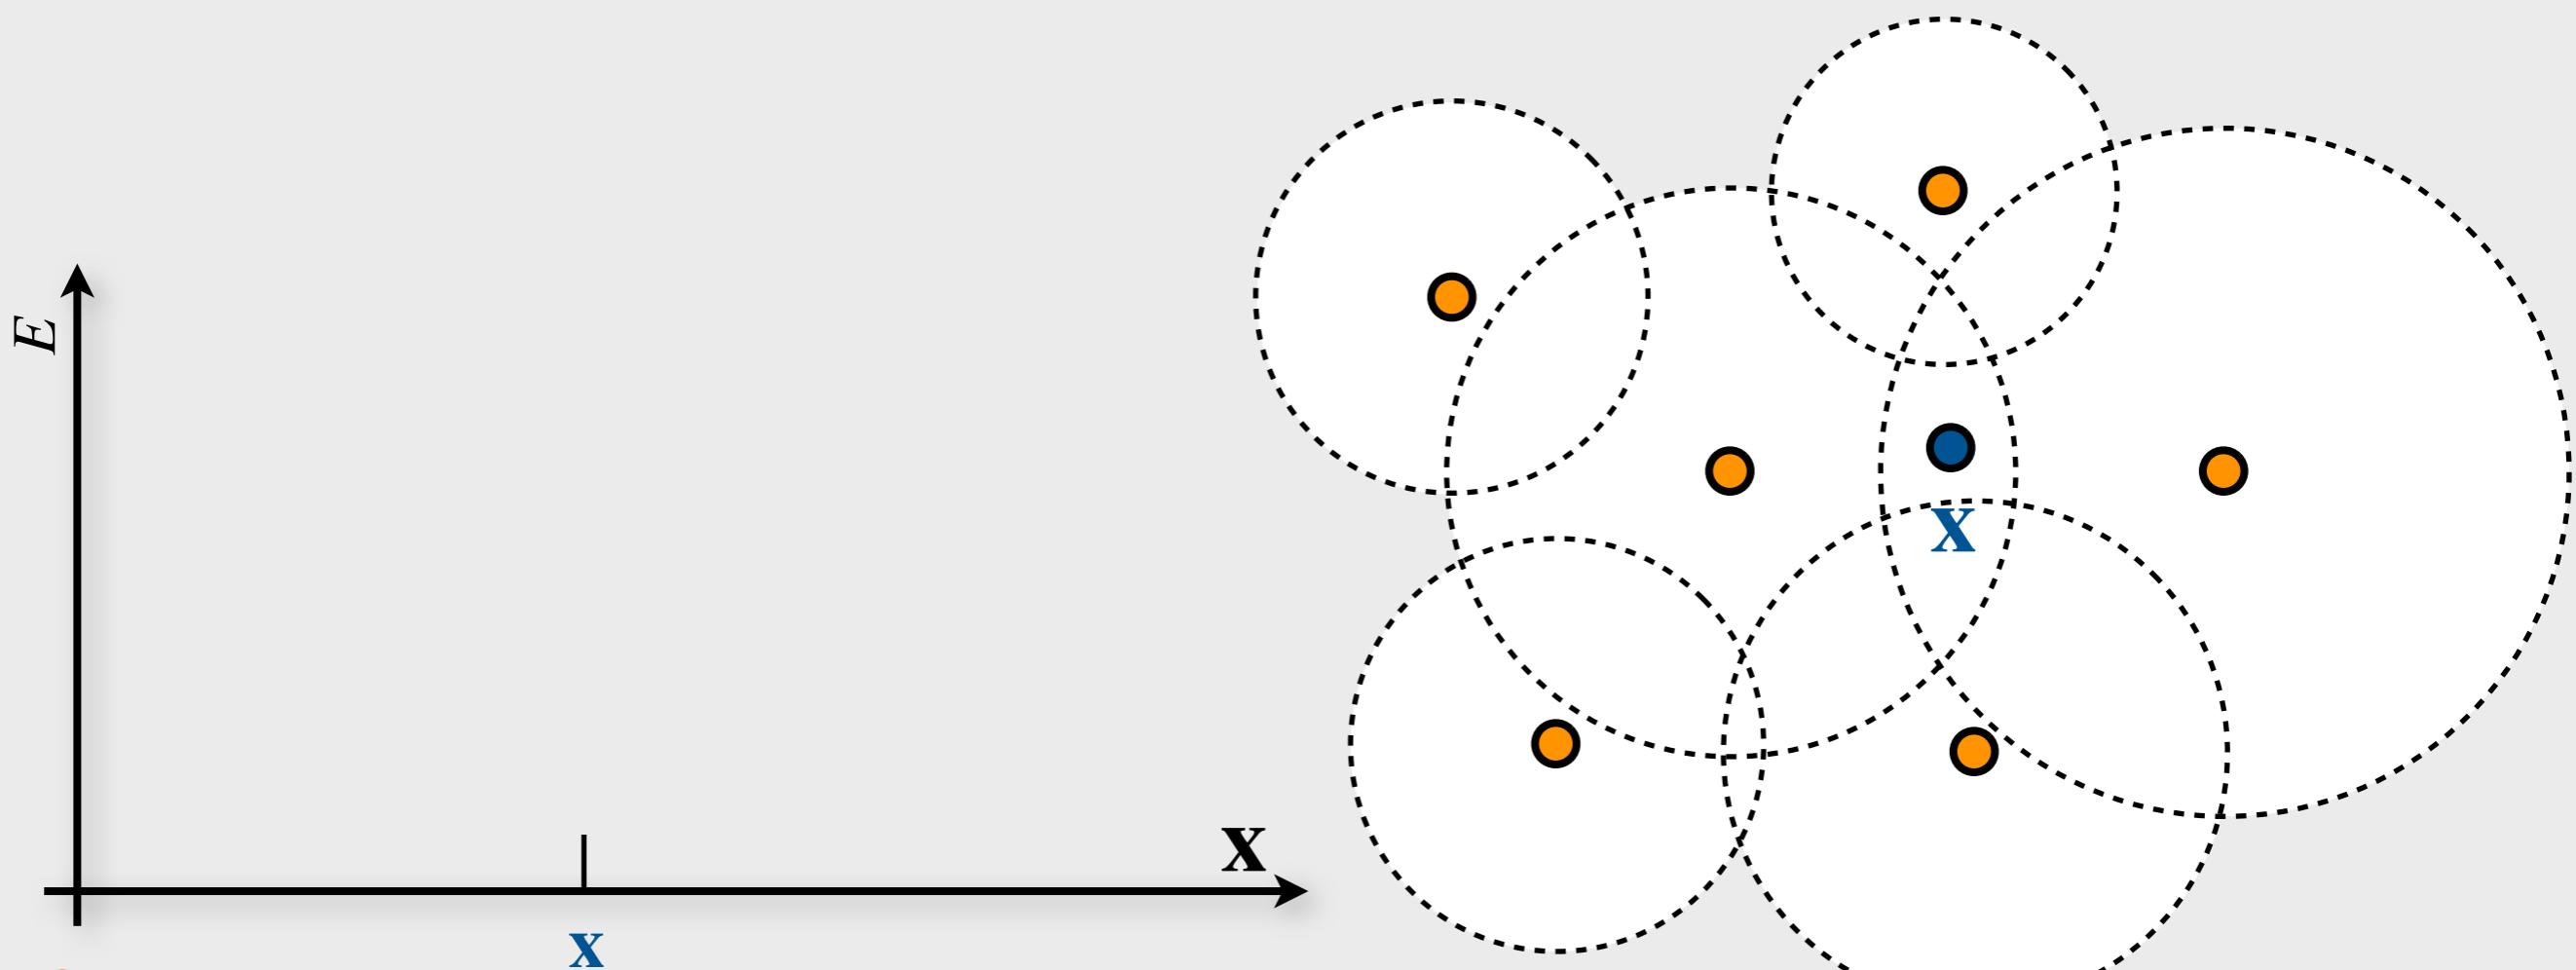
$$\nabla_t E(\mathbf{x}) = \sum_{k=1}^{N_1} \left(\hat{u}_k \sum_{j=2}^{N_2} \nabla_{\hat{u}_k} A_{j-,k} (L_{j,k} - L_{j-1,k}) \cos \theta_{j-} + \hat{v}_{k-} \sum_{j=1}^{N_2} \nabla_{\hat{v}_{k-}} A_{j,k-} (L_{j,k} - L_{j,k-1}) \cos \theta_j \right)$$

Interpolating with gradients

$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) \left(E_i + (\vec{n}_i \times \vec{\mathbf{n}}) \cdot (\vec{\nabla}_r E_i) + (\mathbf{x} - \mathbf{x}_i) \cdot (\vec{\nabla}_t E_i) \right)}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$

Interpolating with gradients

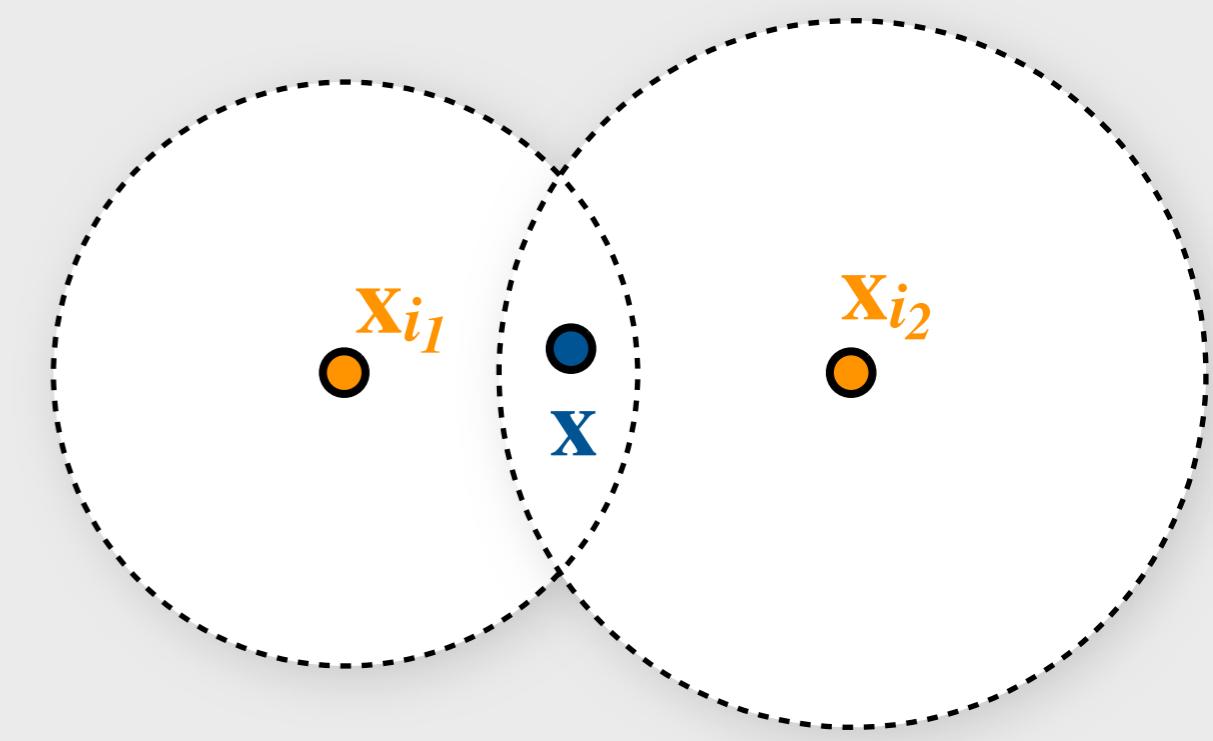
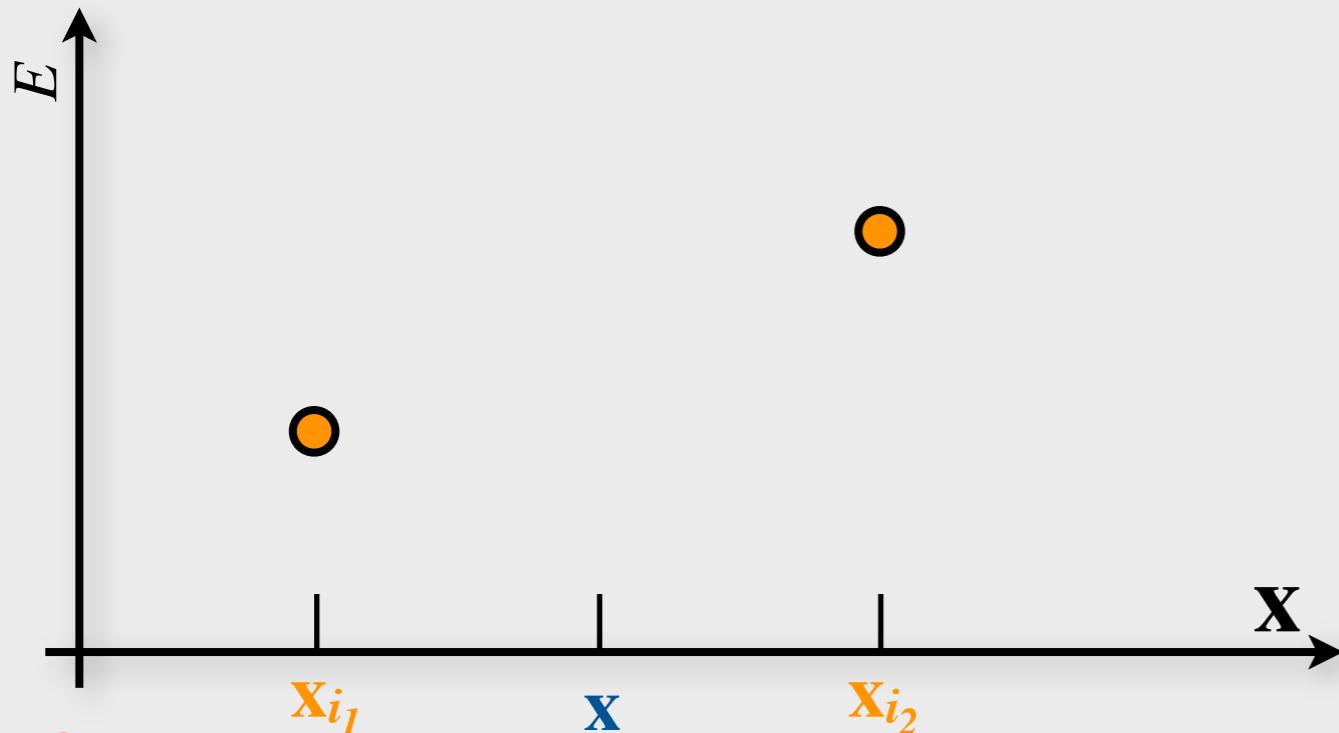
$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) \left(E_i + (\vec{n}_i \times \vec{\mathbf{n}}) \cdot (\vec{\nabla}_r E_i) + (\mathbf{x} - \mathbf{x}_i) \cdot (\vec{\nabla}_t E_i) \right)}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$



Interpolating with gradients

$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) \left(E_i + (\vec{n}_i \times \vec{\mathbf{n}}) \cdot (\vec{\nabla}_r E_i) + (\mathbf{x} - \mathbf{x}_i) \cdot (\vec{\nabla}_t E_i) \right)}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$

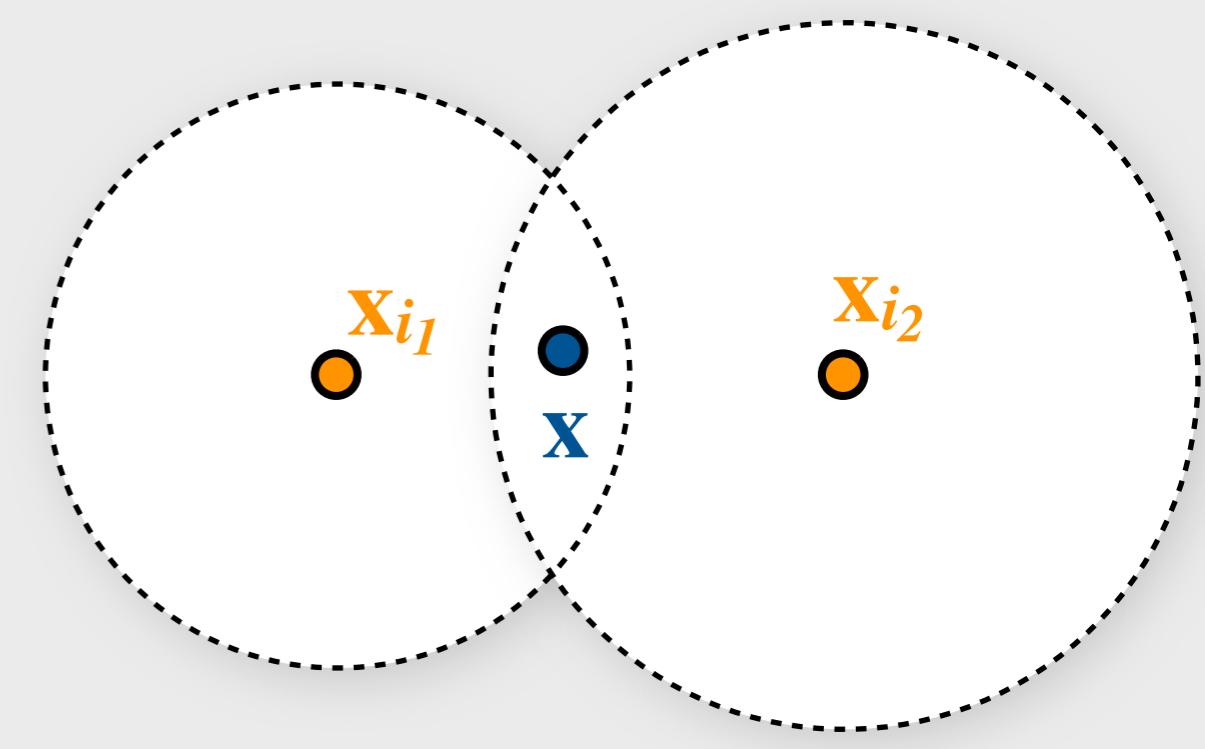
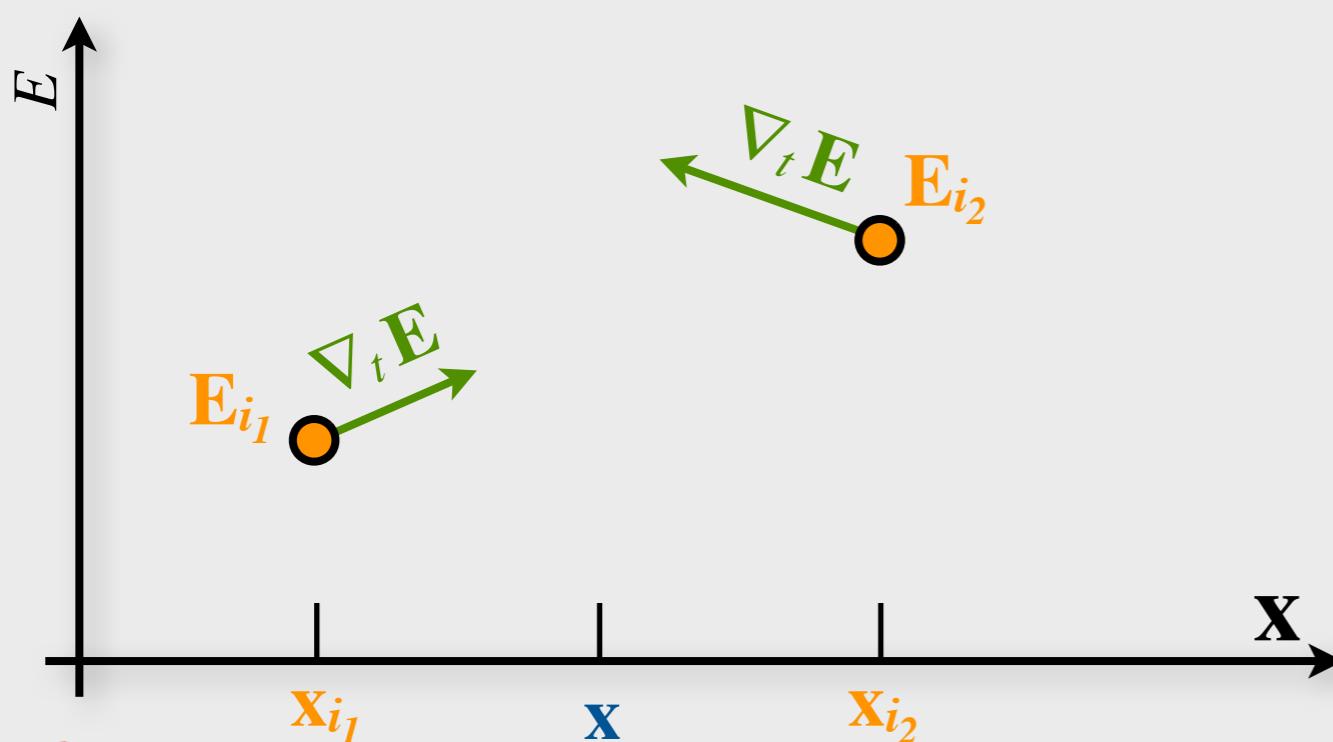
Find overlapping cache records



Interpolating with gradients

$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) \left(E_i + (\vec{n}_i \times \vec{\mathbf{n}}) \cdot (\vec{\nabla}_r E_i) + (\mathbf{x} - \mathbf{x}_i) \cdot (\vec{\nabla}_t E_i) \right)}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$

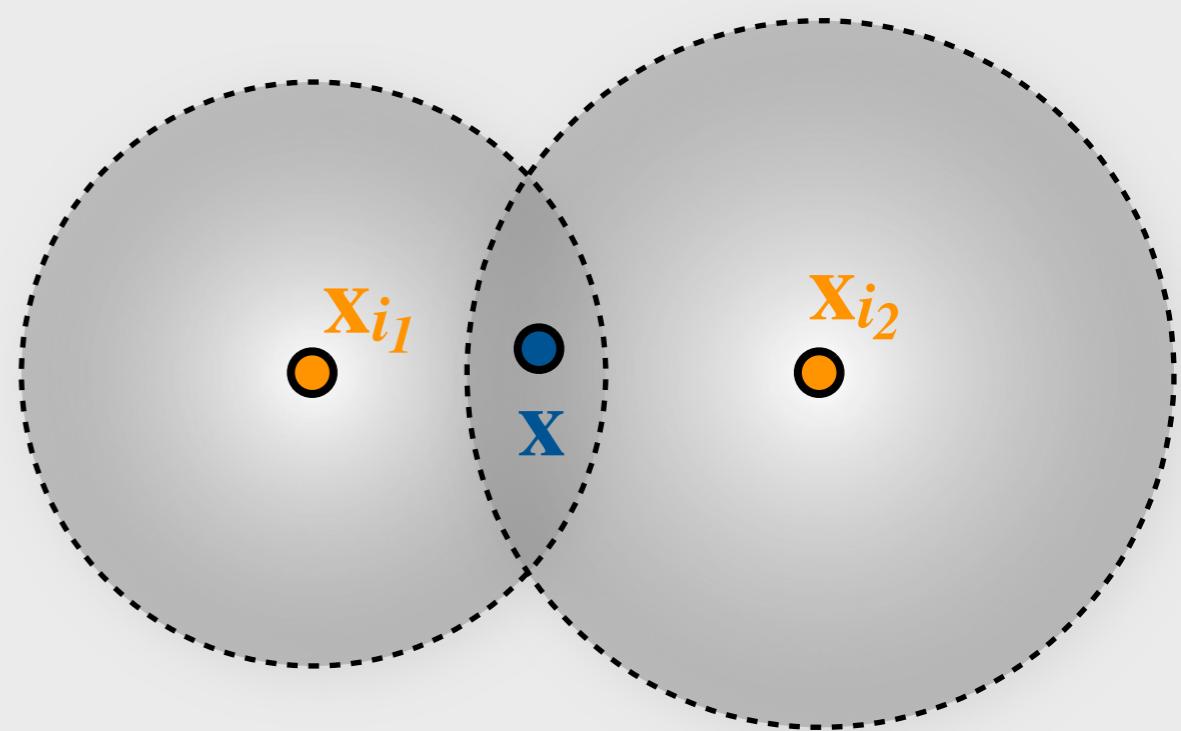
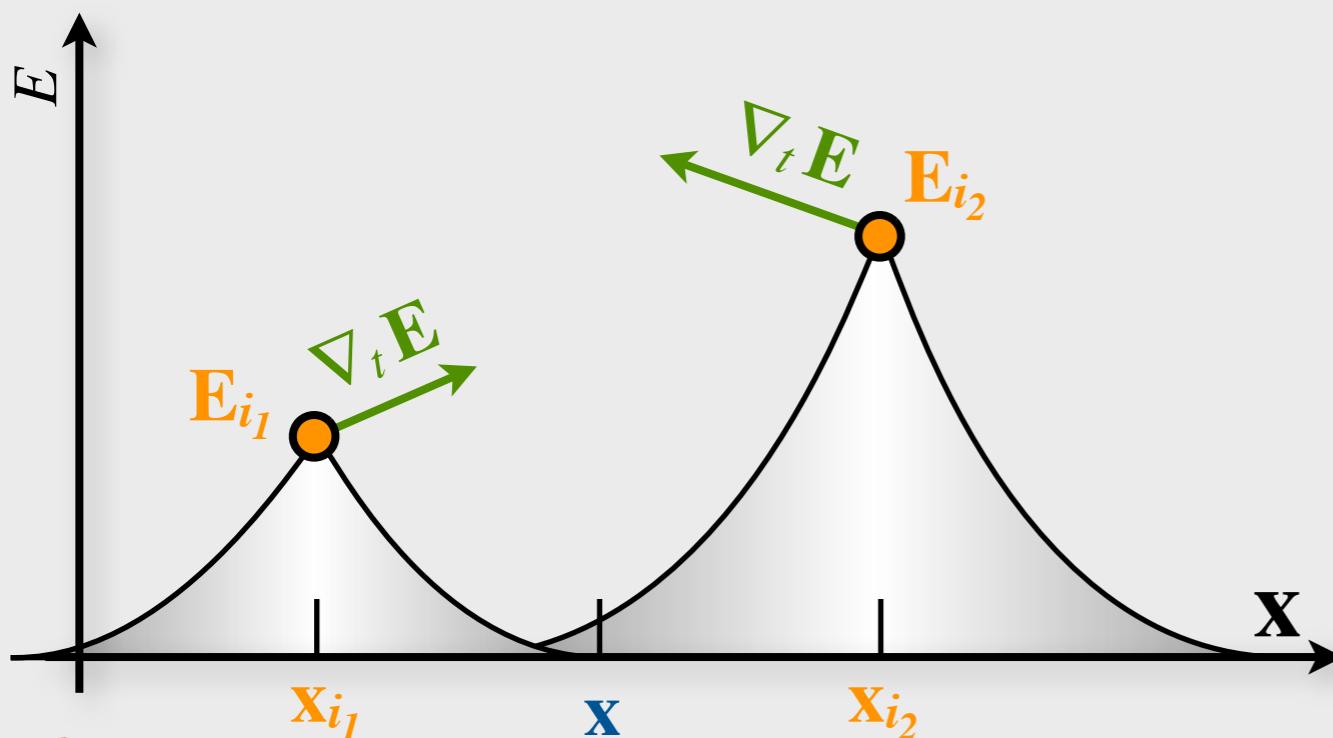
Extrapolate along gradients



Interpolating with gradients

$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) \left(E_i + (\vec{n}_i \times \vec{\mathbf{n}}) \cdot (\vec{\nabla}_r E_i) + (\mathbf{x} - \mathbf{x}_i) \cdot (\vec{\nabla}_t E_i) \right)}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$

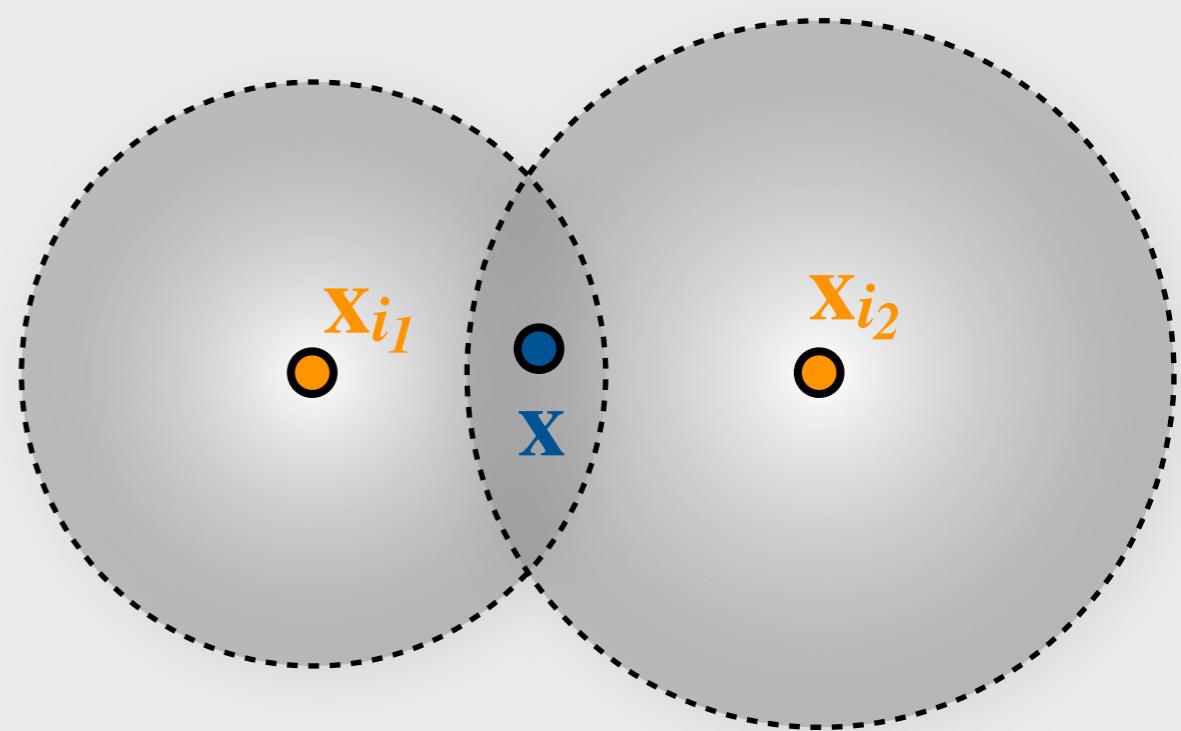
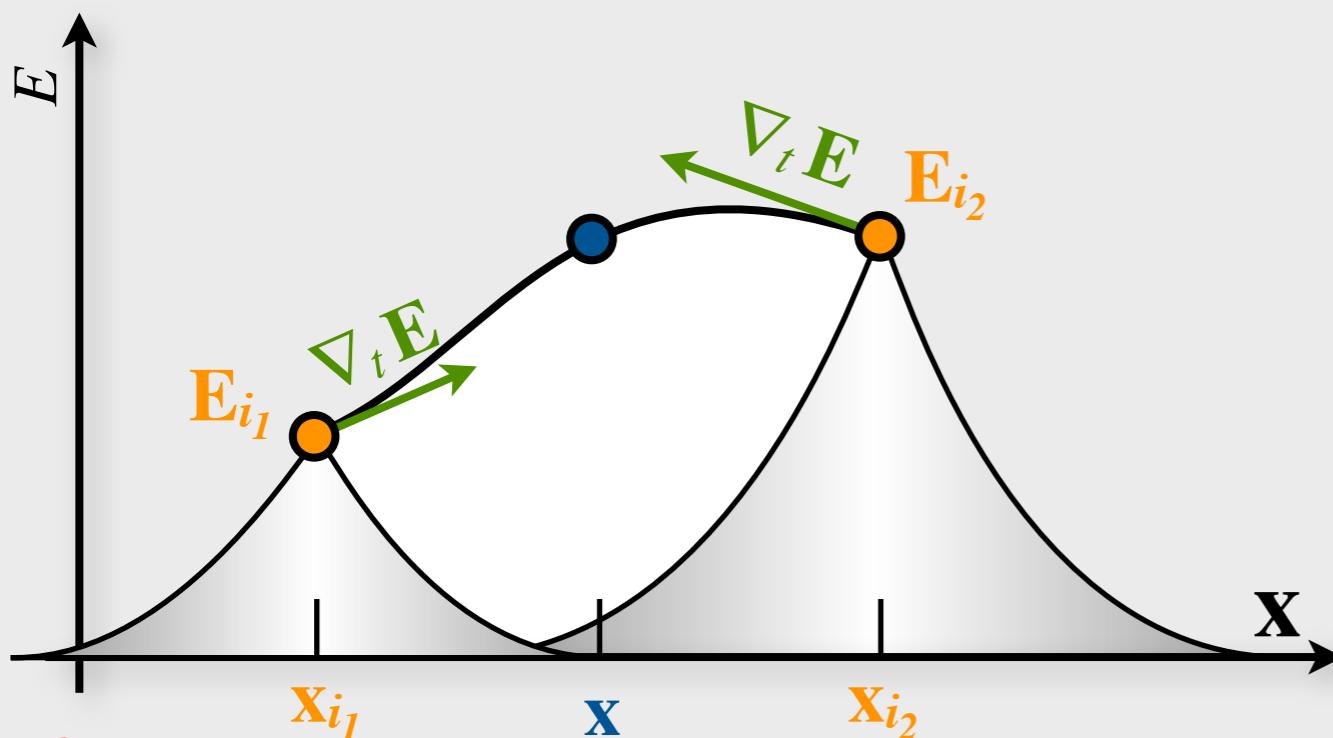
Weight contributions



Interpolating with gradients

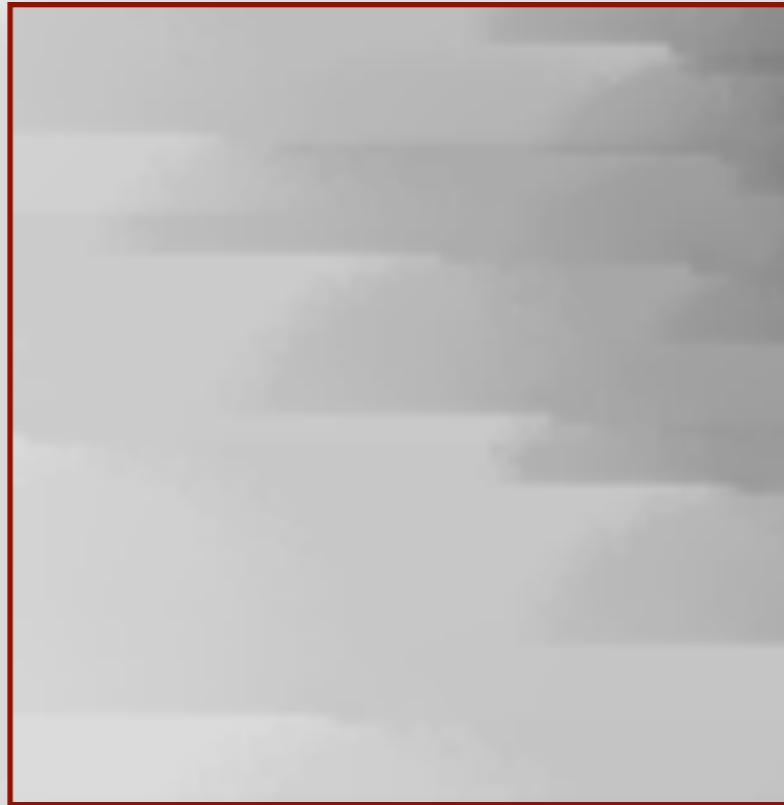
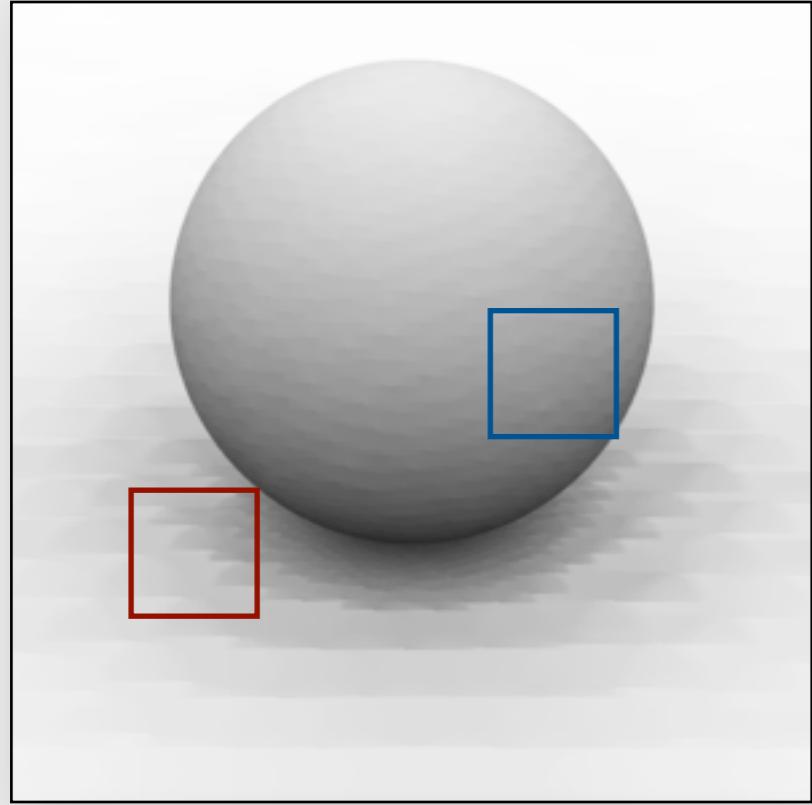
$$E(\mathbf{x}, \vec{\mathbf{n}}) \approx \frac{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}}) \left(E_i + (\vec{n}_i \times \vec{\mathbf{n}}) \cdot (\vec{\nabla}_r E_i) + (\mathbf{x} - \mathbf{x}_i) \cdot (\vec{\nabla}_t E_i) \right)}{\sum_{i \in S} w_i(\mathbf{x}, \vec{\mathbf{n}})}$$

Sum extrapolated values

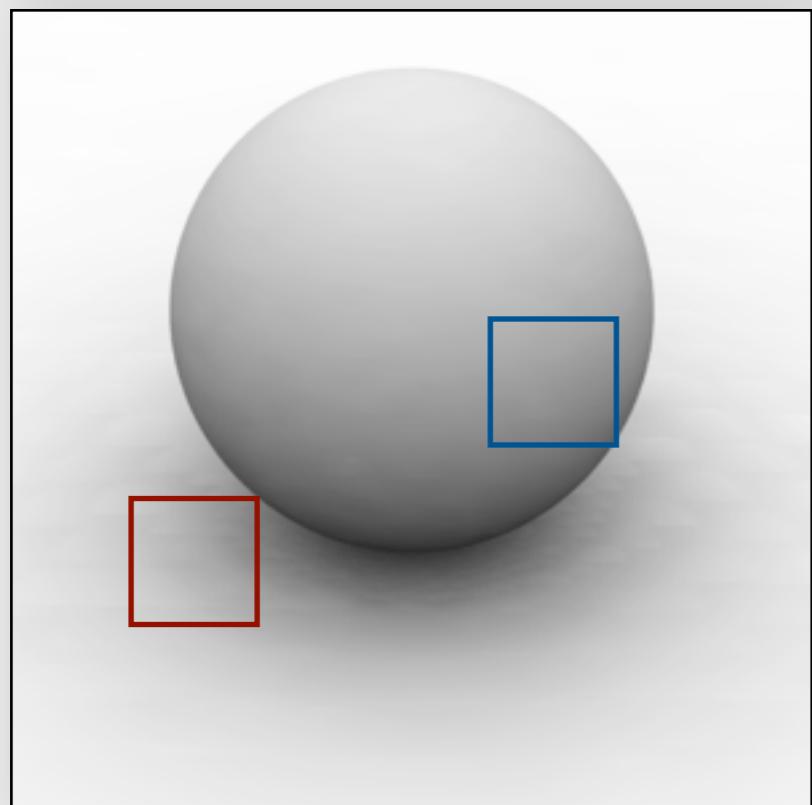


Irradiance Gradients

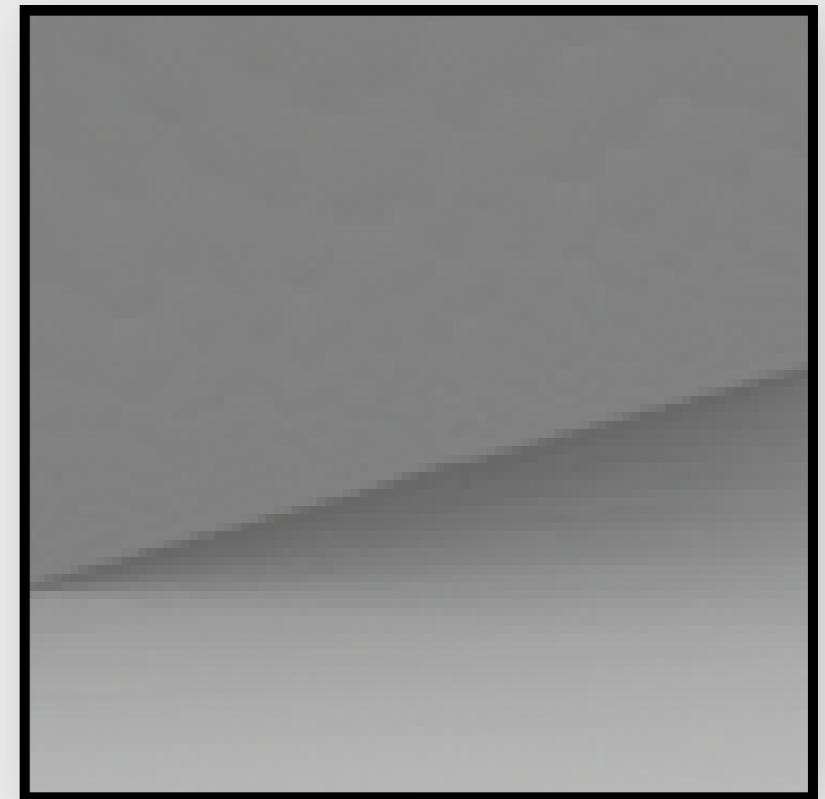
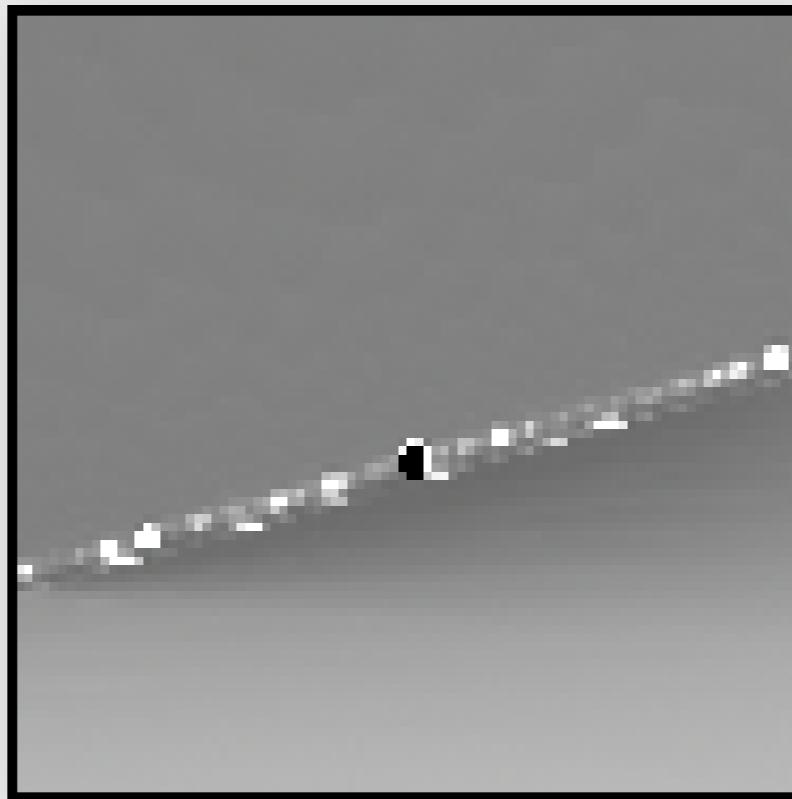
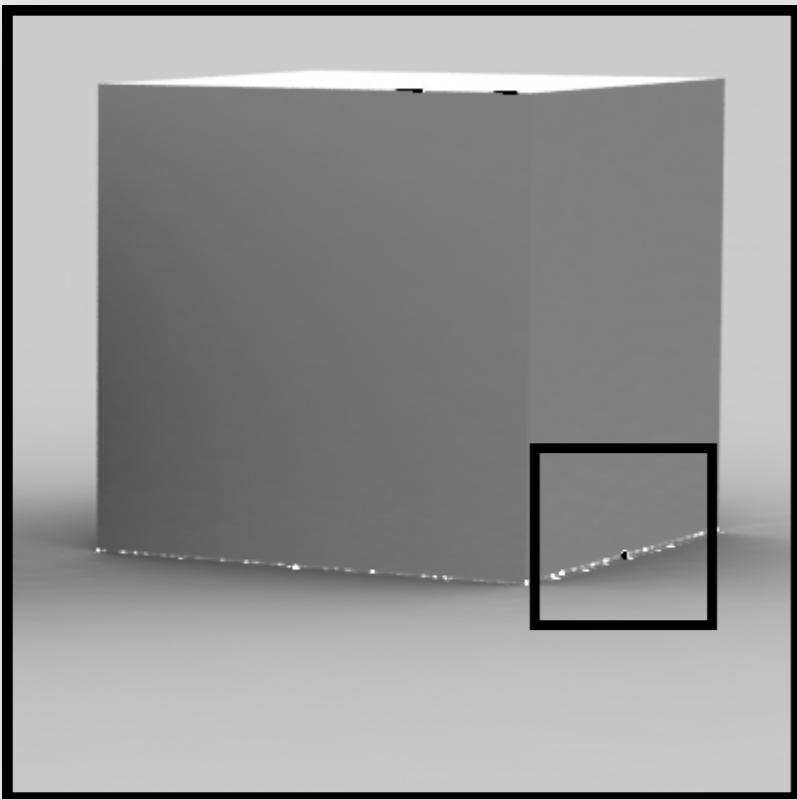
w/o gradients



w/ gradients



Practical issue: over-shooting



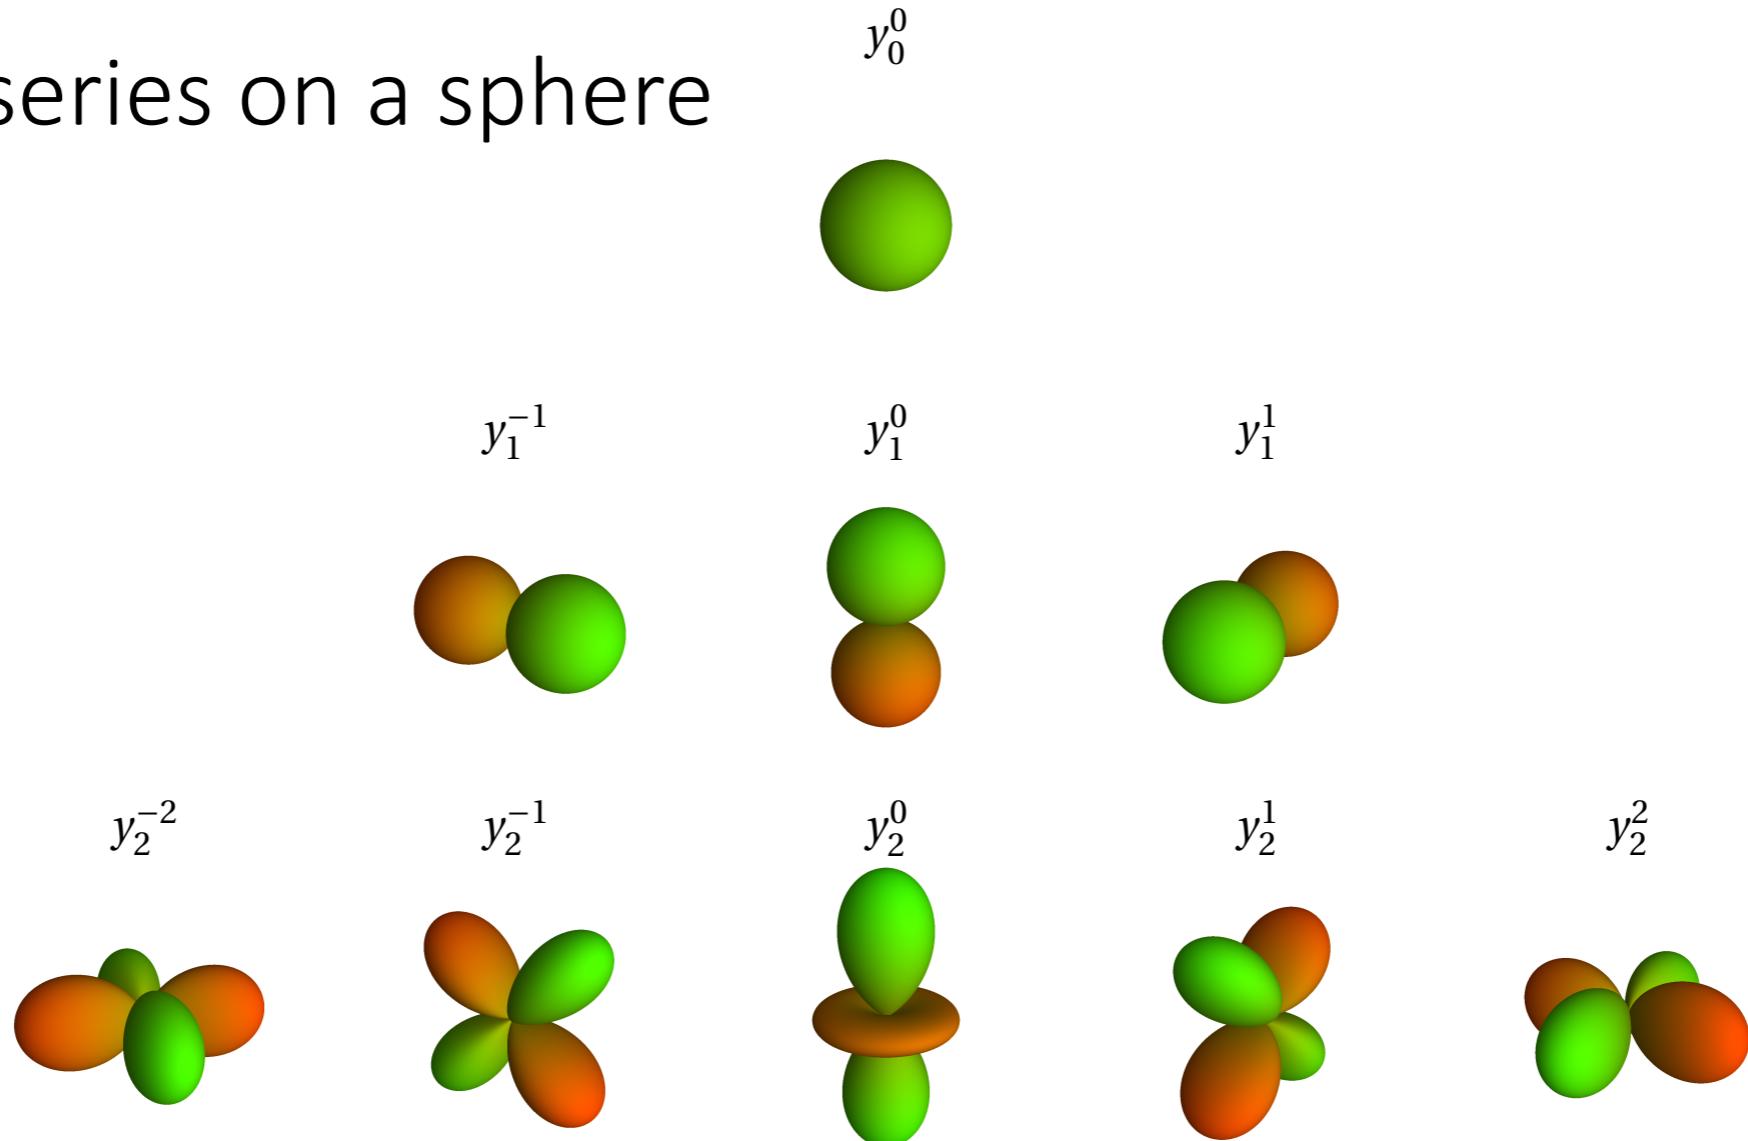
Beyond Lambertian surfaces

- Generalization to glossy surfaces
- Radiance Caching [Křivánek et al. 2005a,2005b]
 - Can no longer cache just the irradiance value
 - Cache full hemispherical *radiance* field at sparse locations

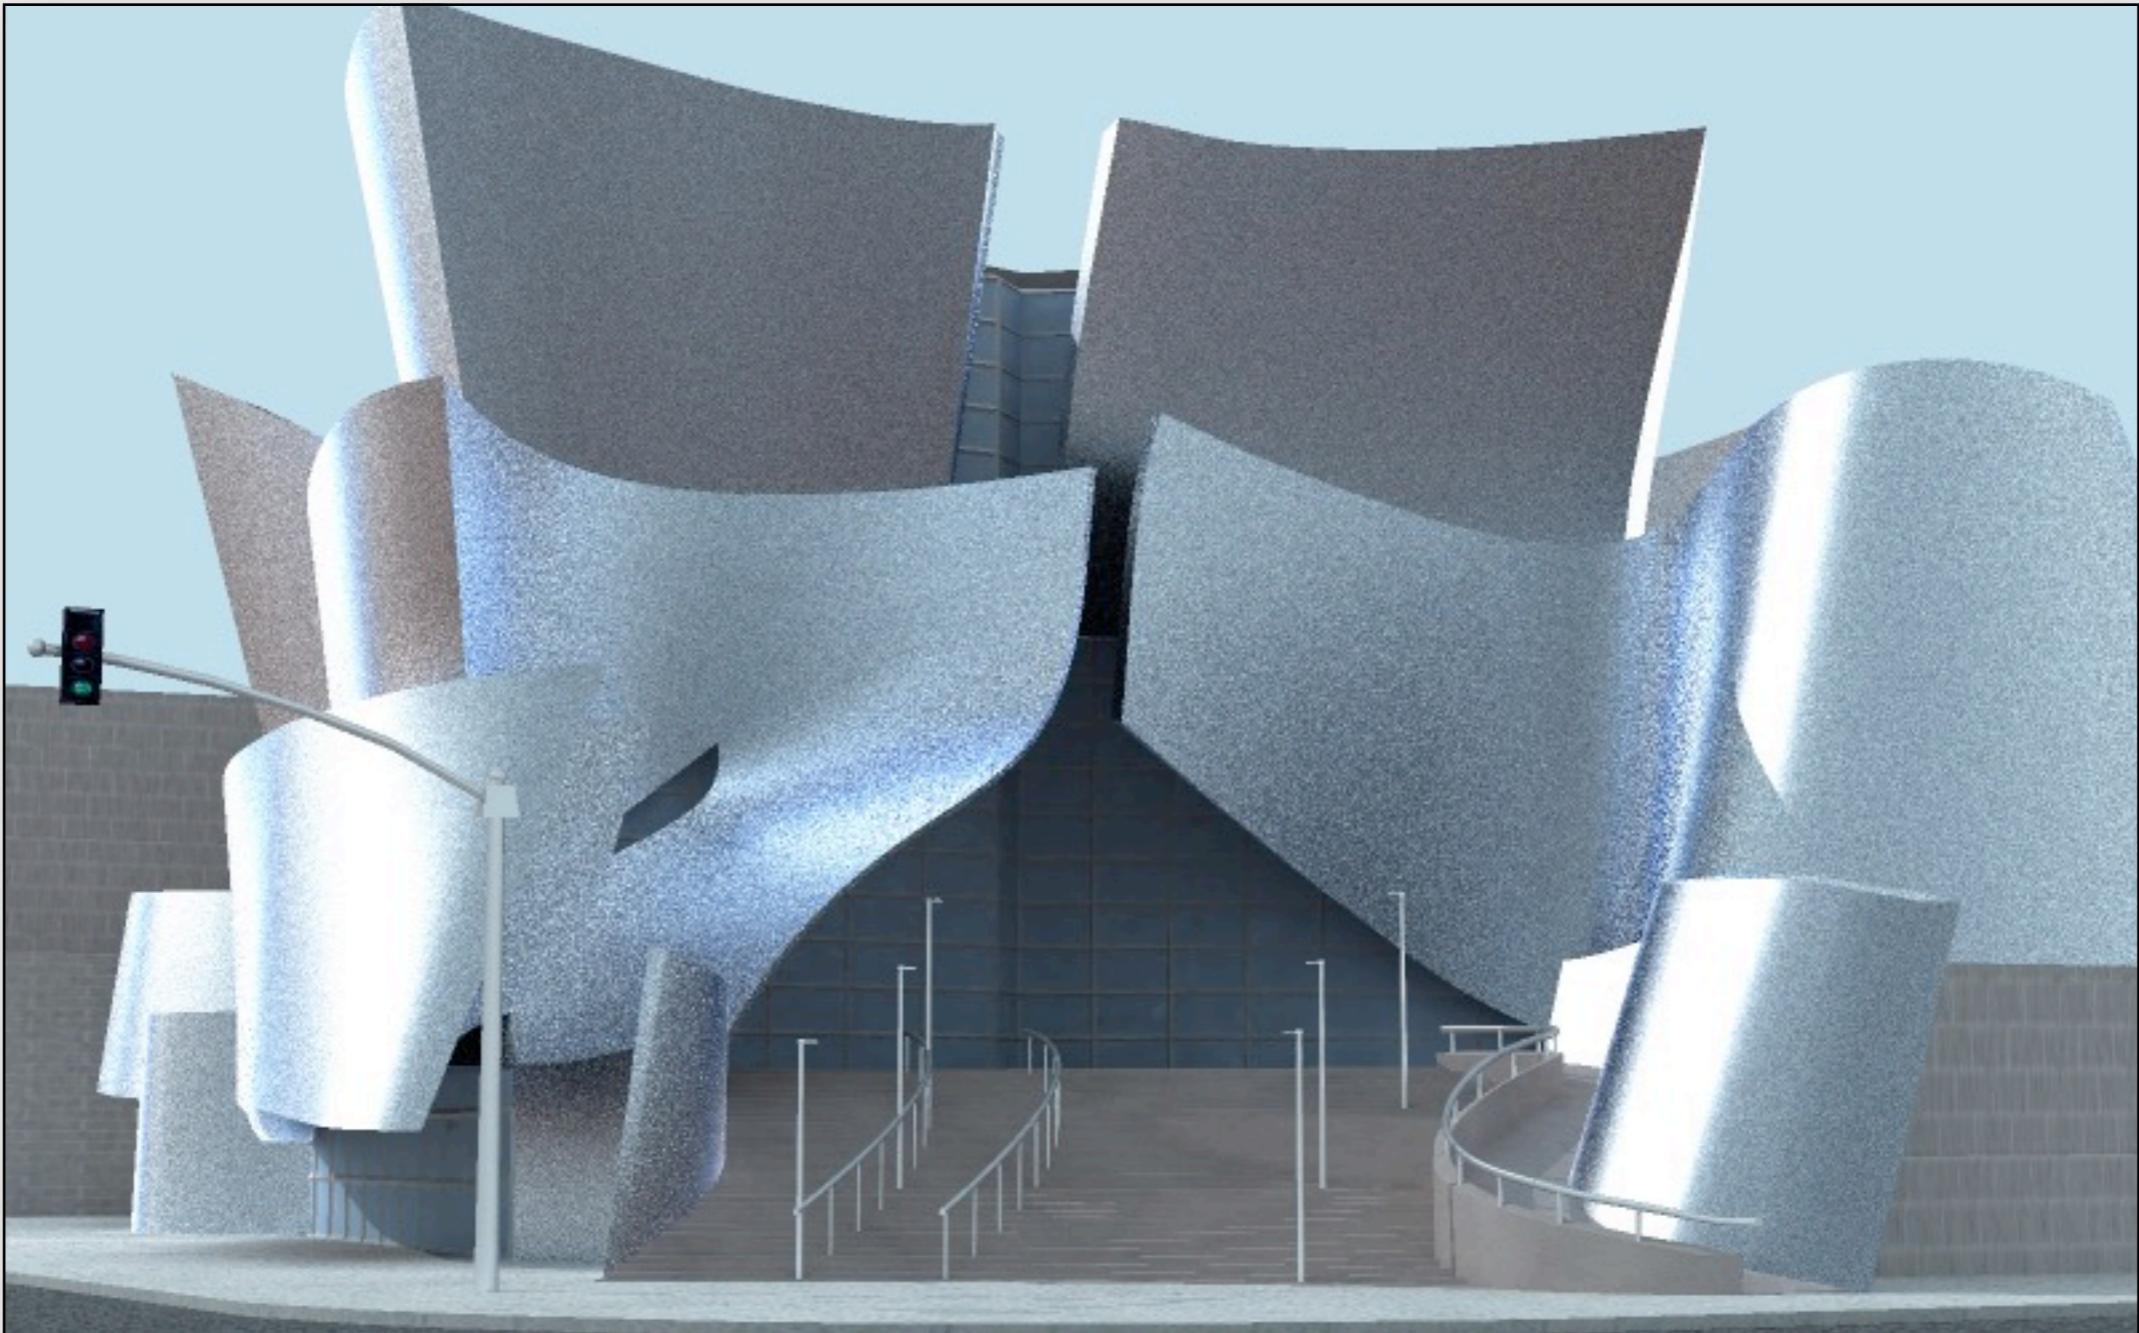


Radiance Storage

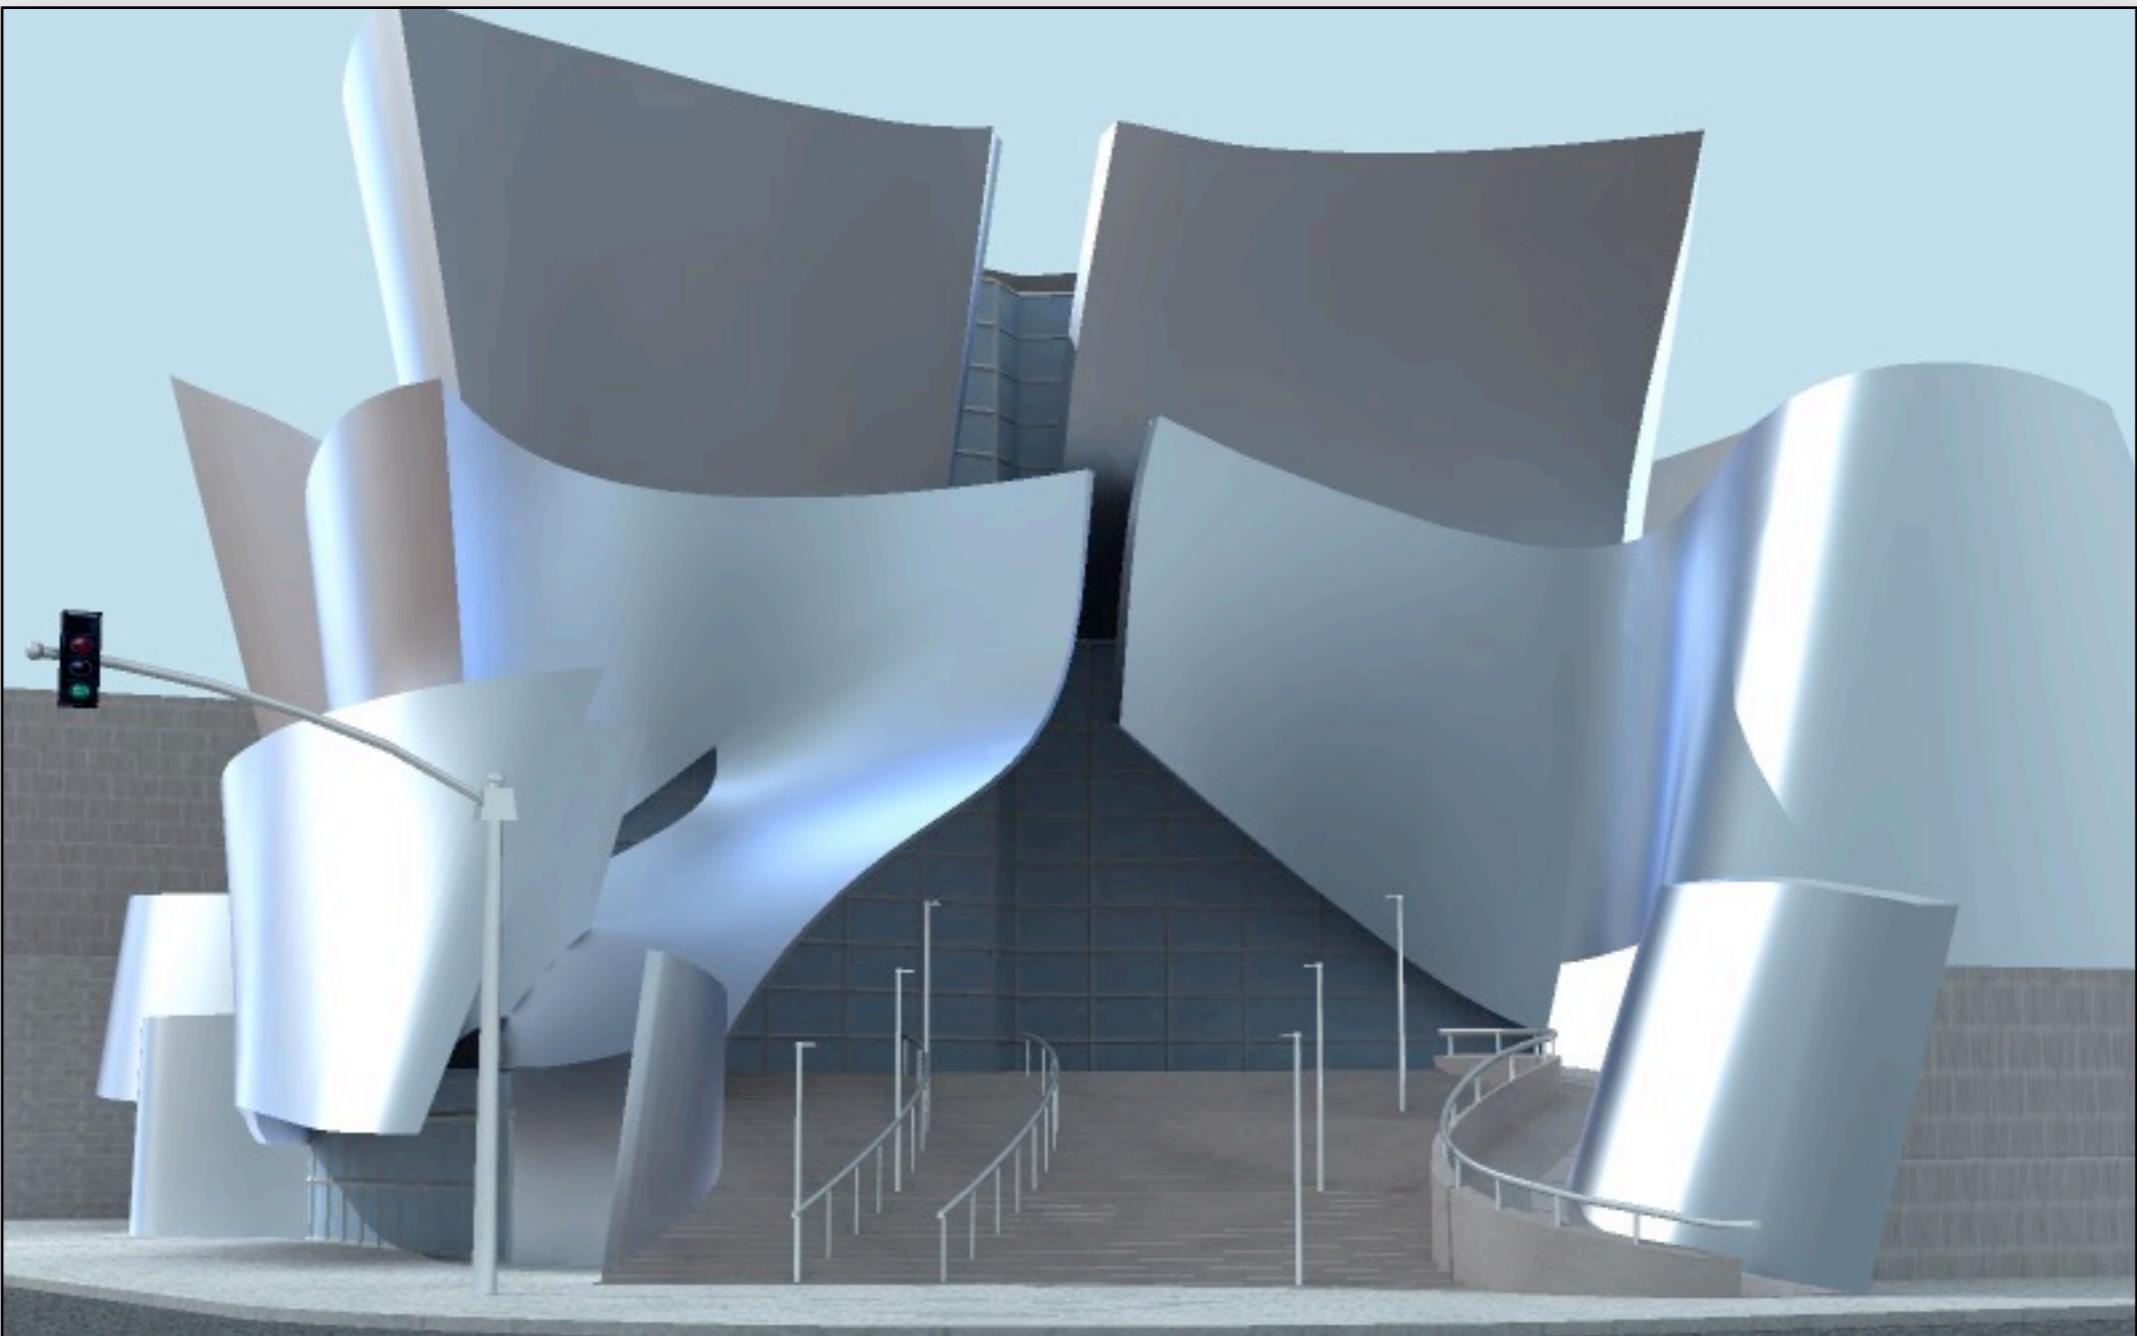
- Use spherical or hemispherical harmonics
- Approximates smooth functions with a few coefficients
- Fourier series on a sphere



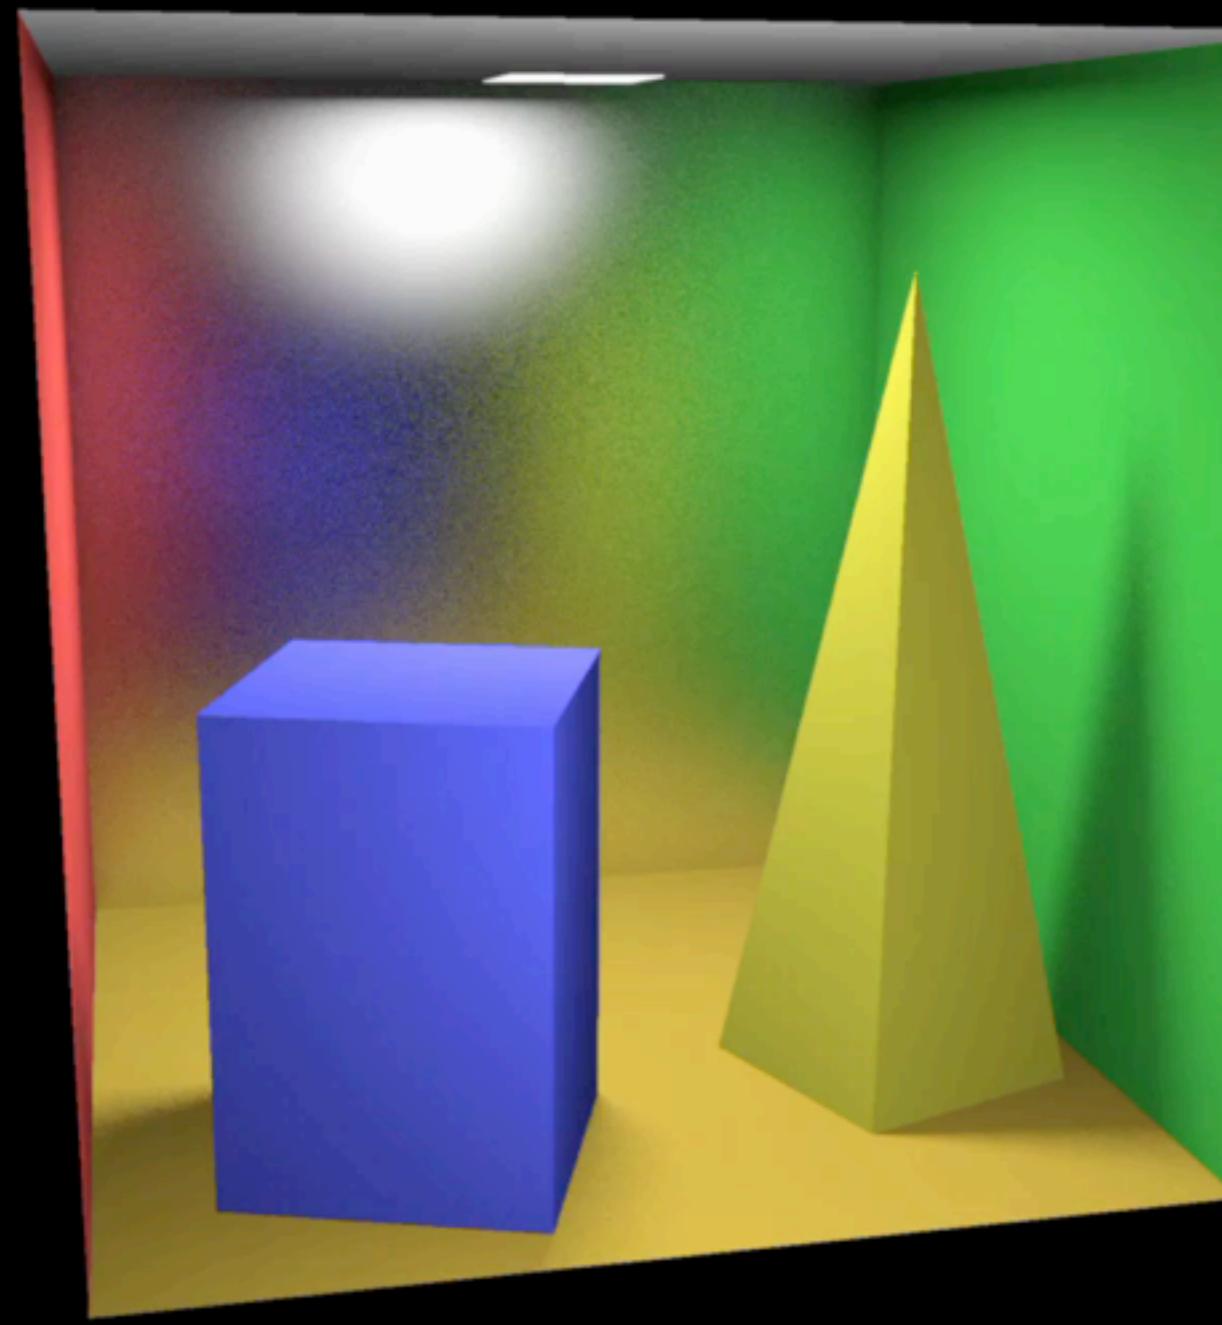
Monte Carlo



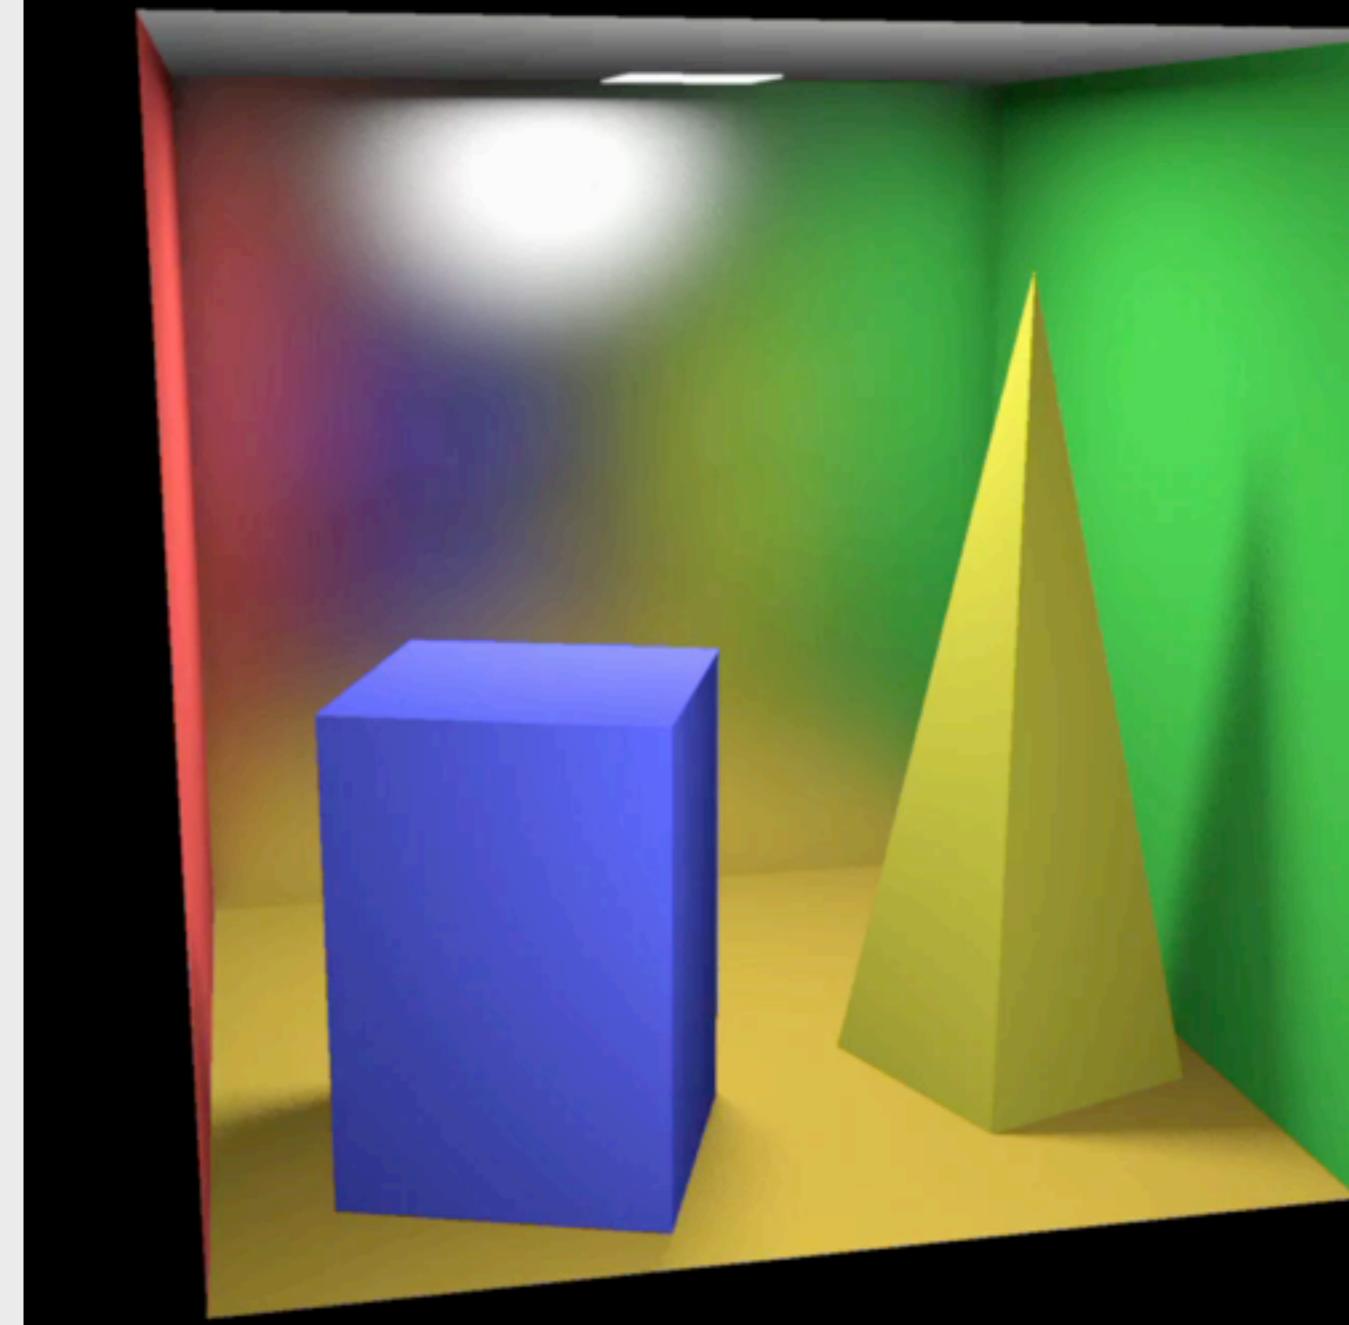
Radiance Caching



Glossy Surfaces



Path Tracing



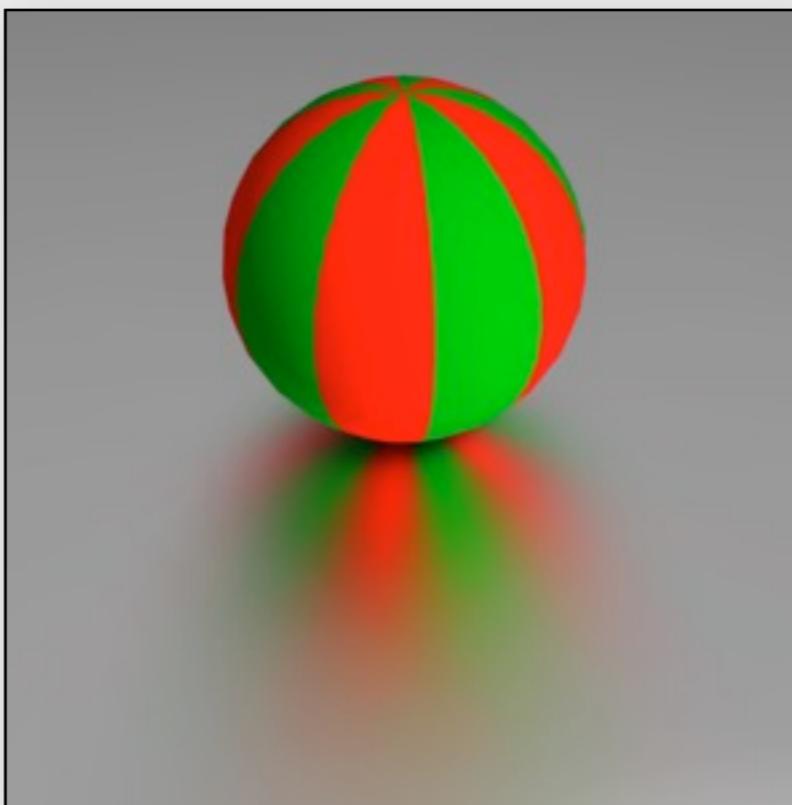
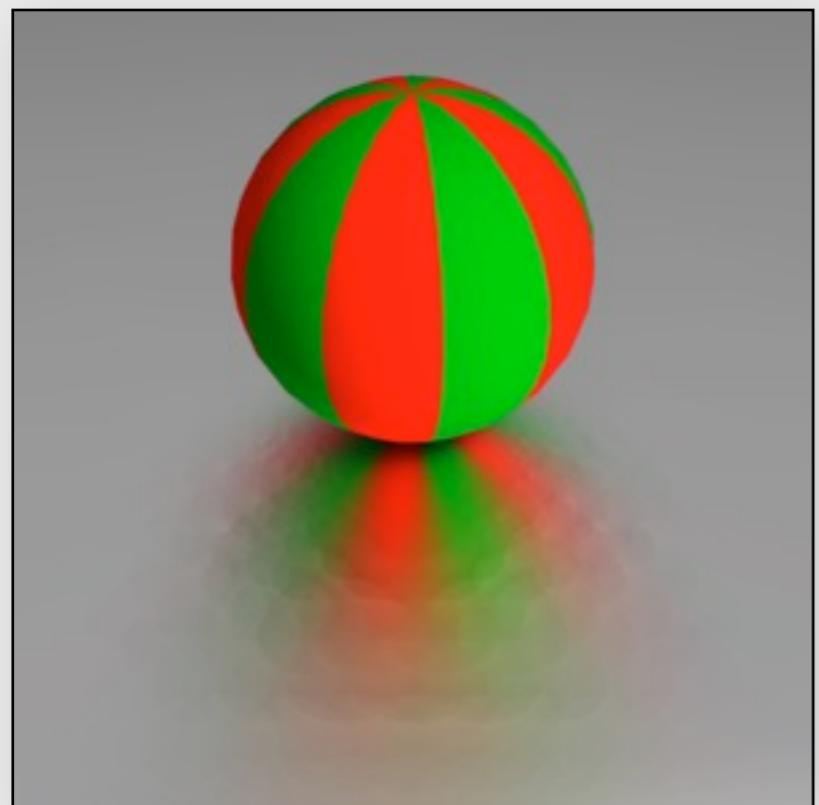
Radiance Caching



Radiance Gradients

- Improve interpolation quality by storing gradient of incoming radiance field





occlusion-unaware

occlusion-aware



[Krivanek et al. 2005a]

[Krivanek et al. 2005b]

occlusion-unaware

occlusion-aware



Eurographics 2015

The 36th Annual Conference of the
European Association for Computer Graphics

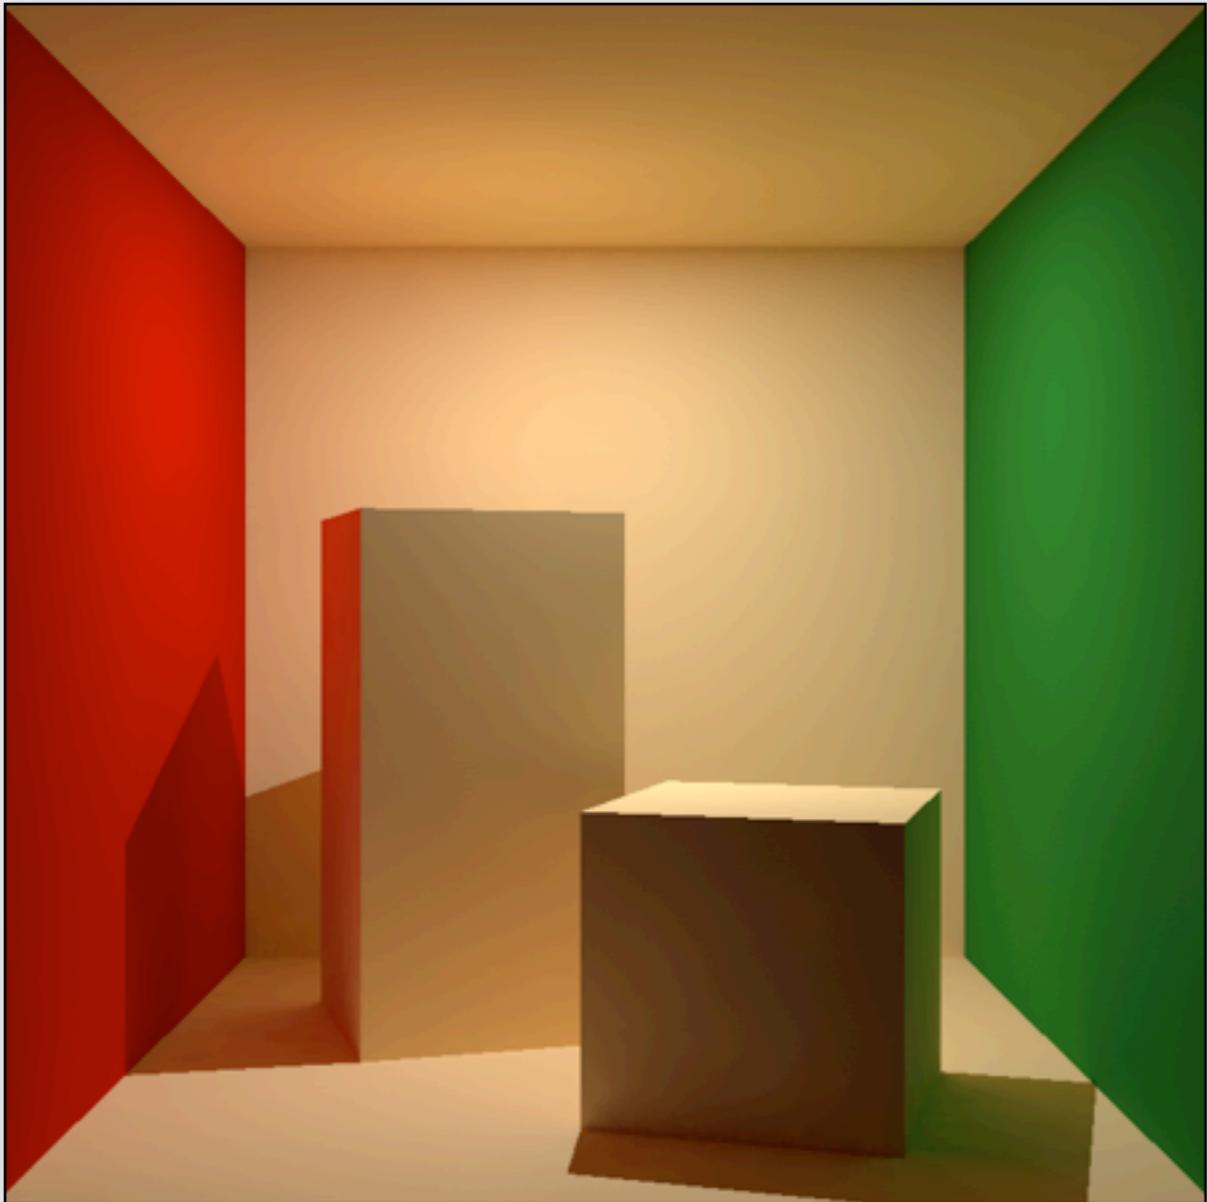
[Křivánek et al. 2005a,2005b] ⁶⁶

Beyond surfaces

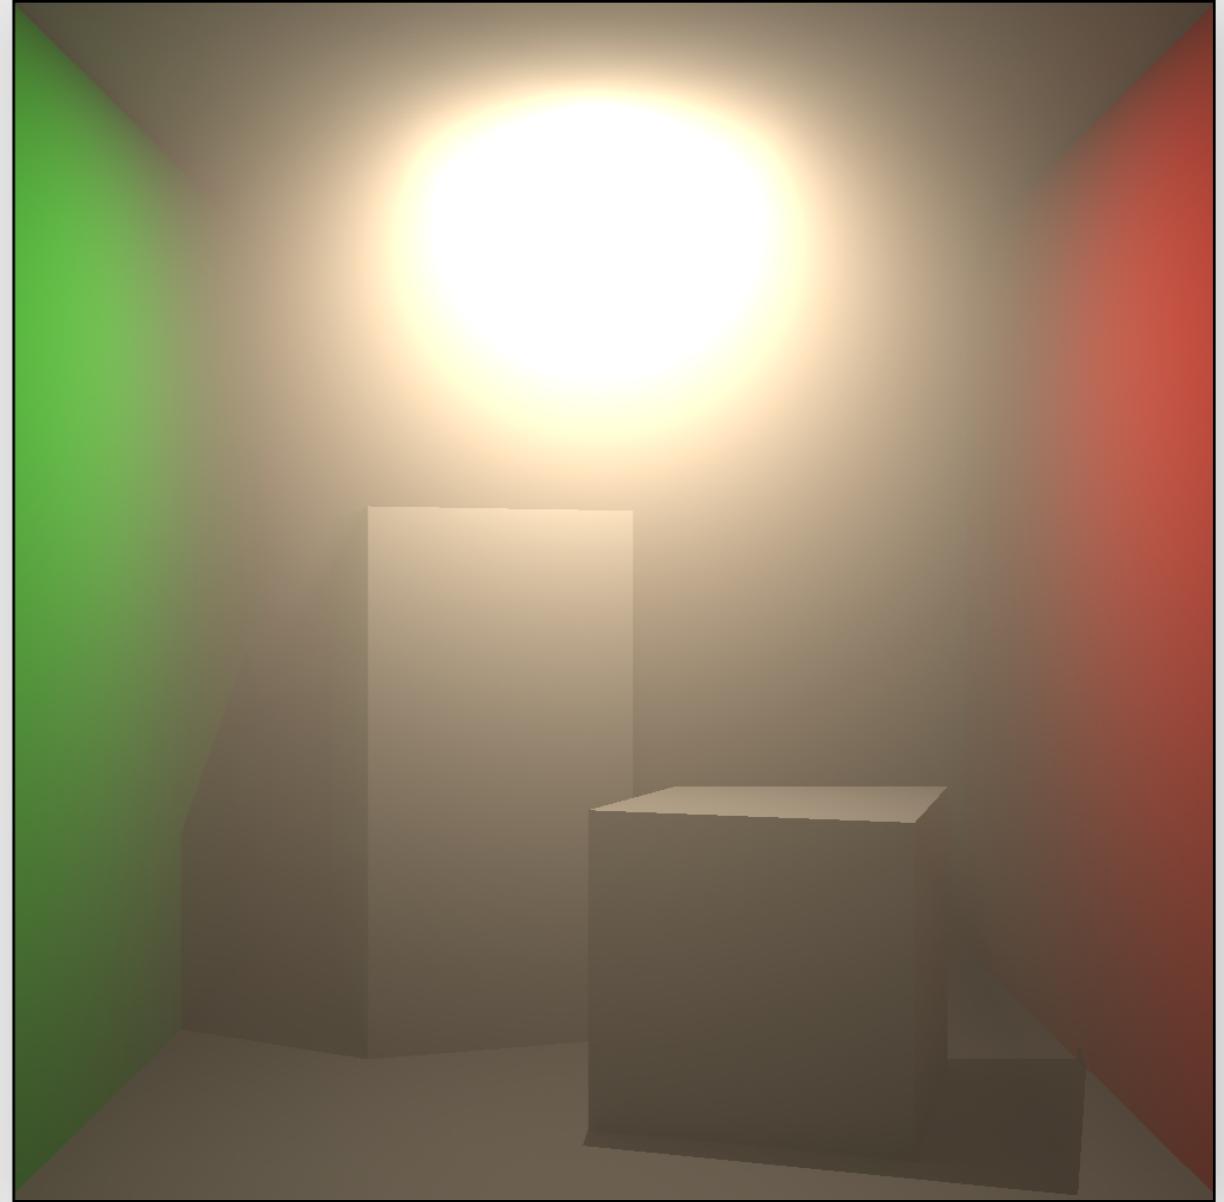
- Generalizations to participating media
- Volumetric Radiance Caching [Jarosz et al. 2008a, 2008b]
 - Cache radiance and gradients within volume



Participating Media

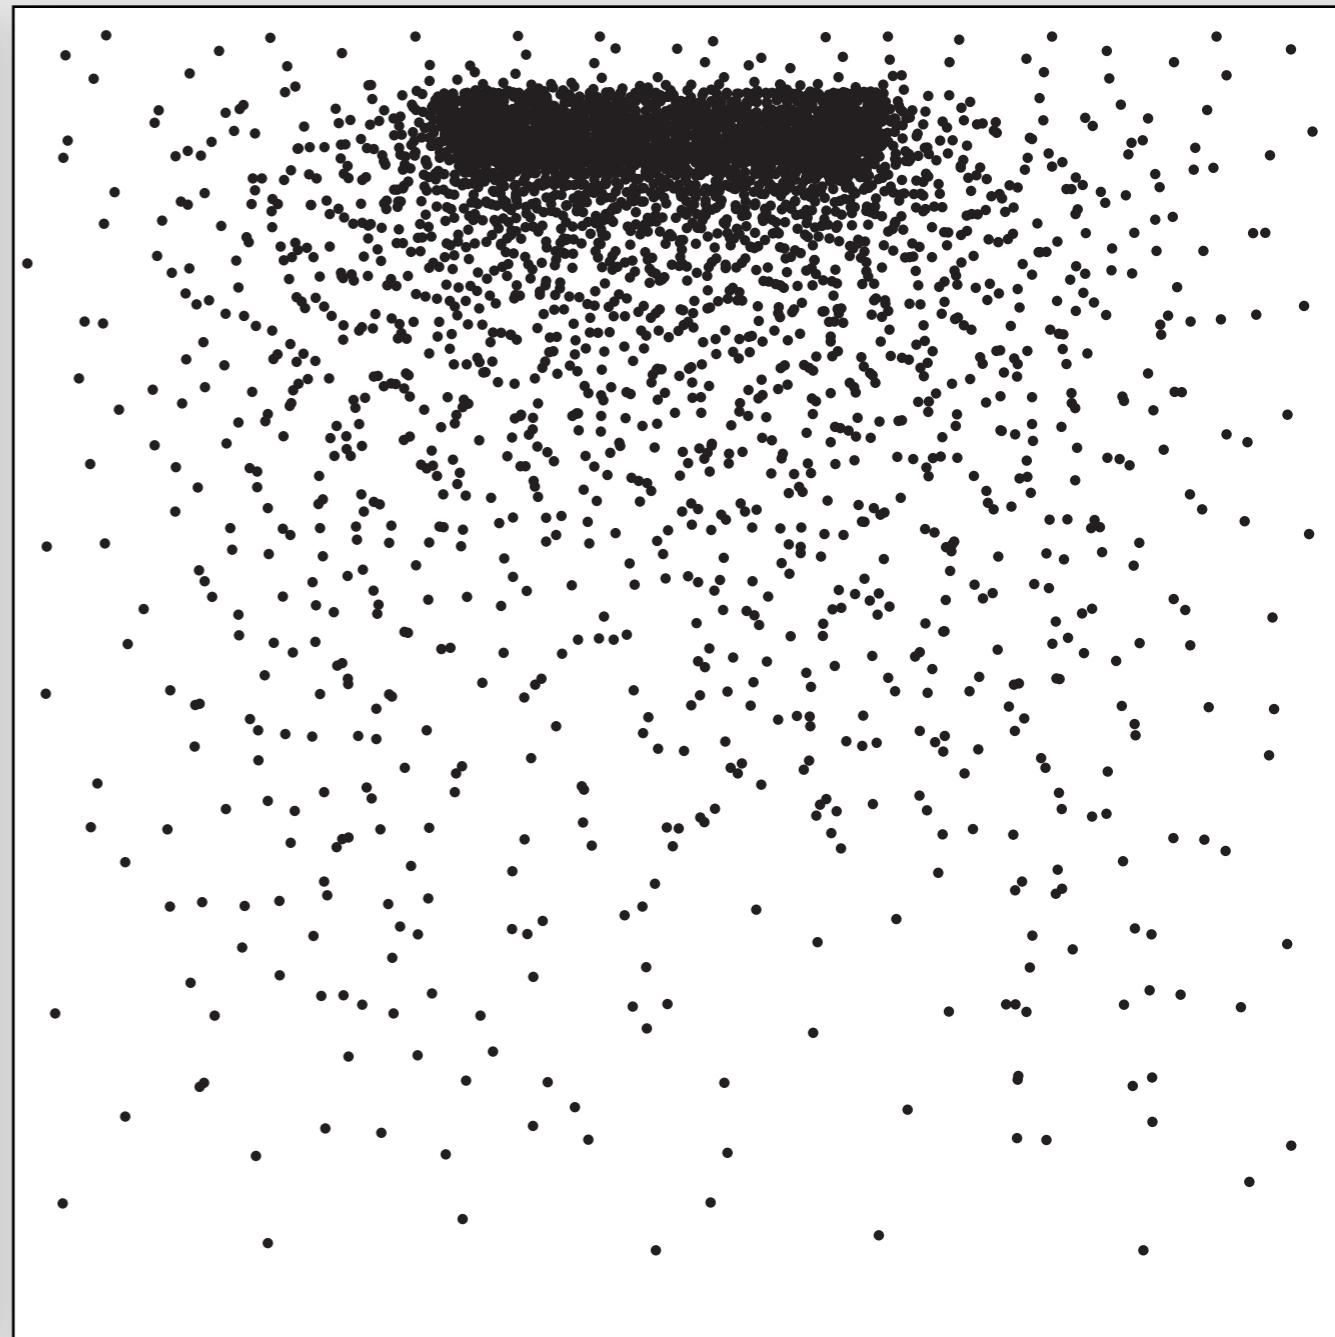
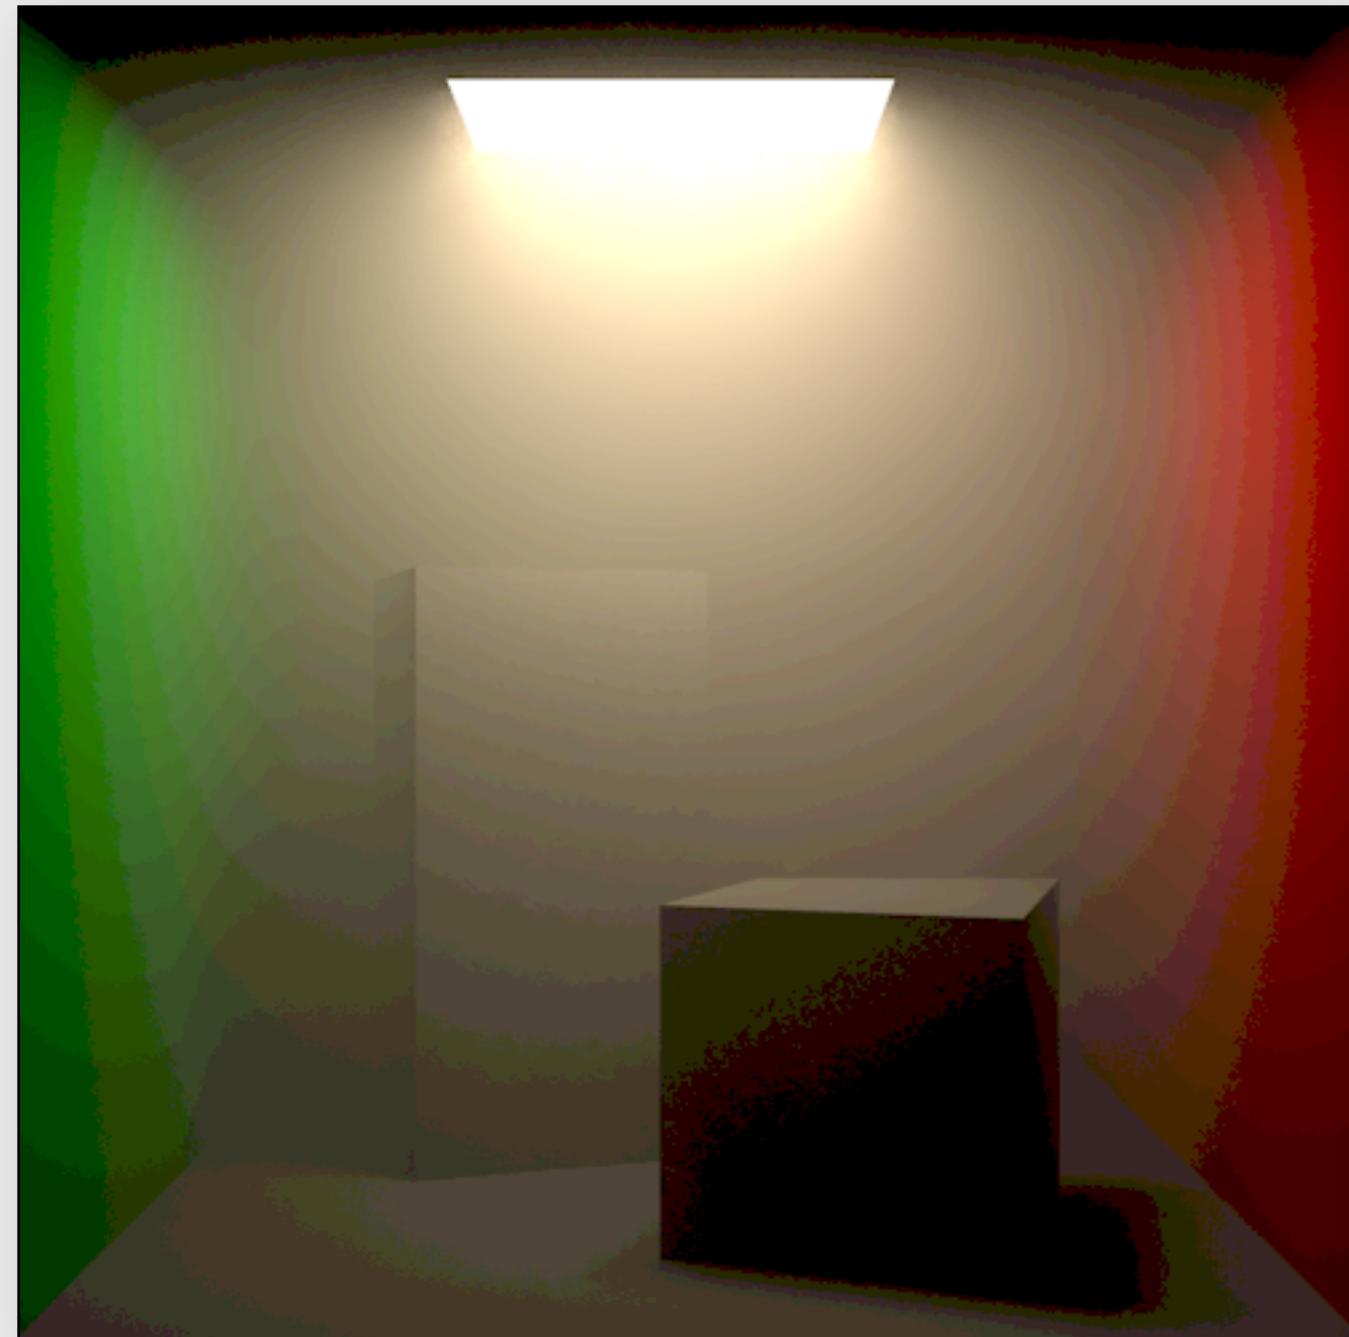


no media

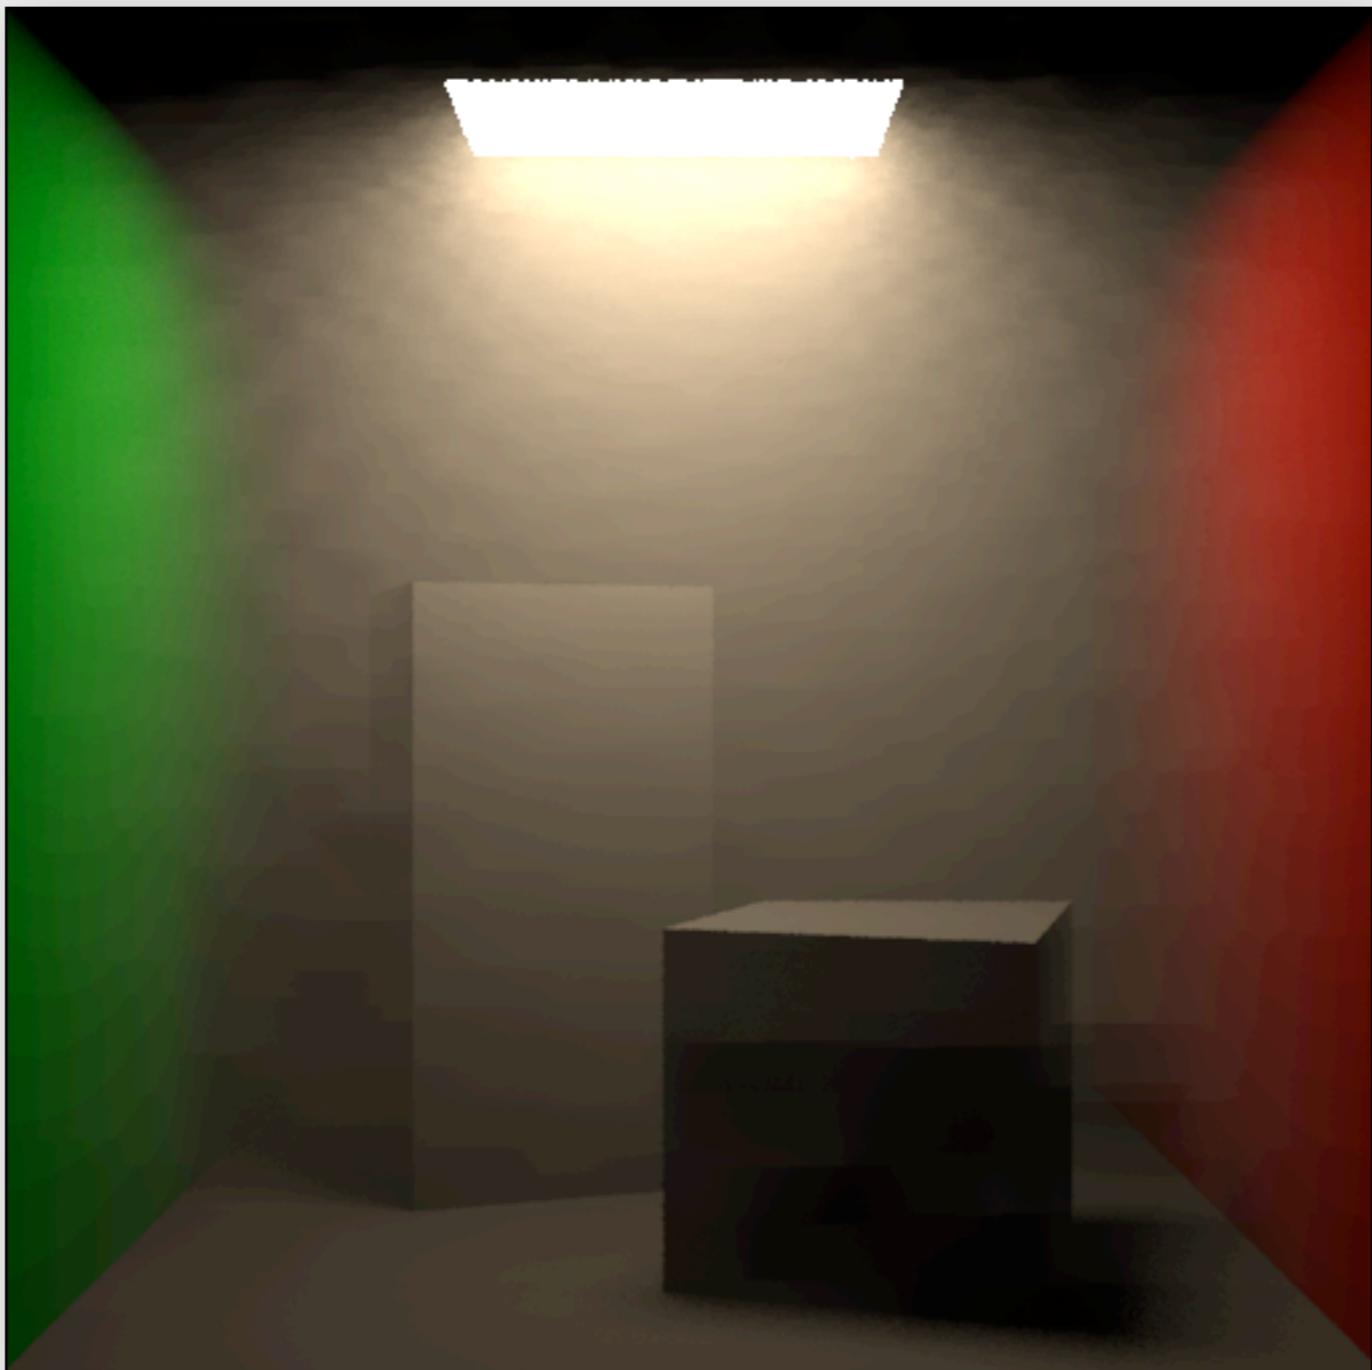


with media

Valid Radius



Gradients



no gradients

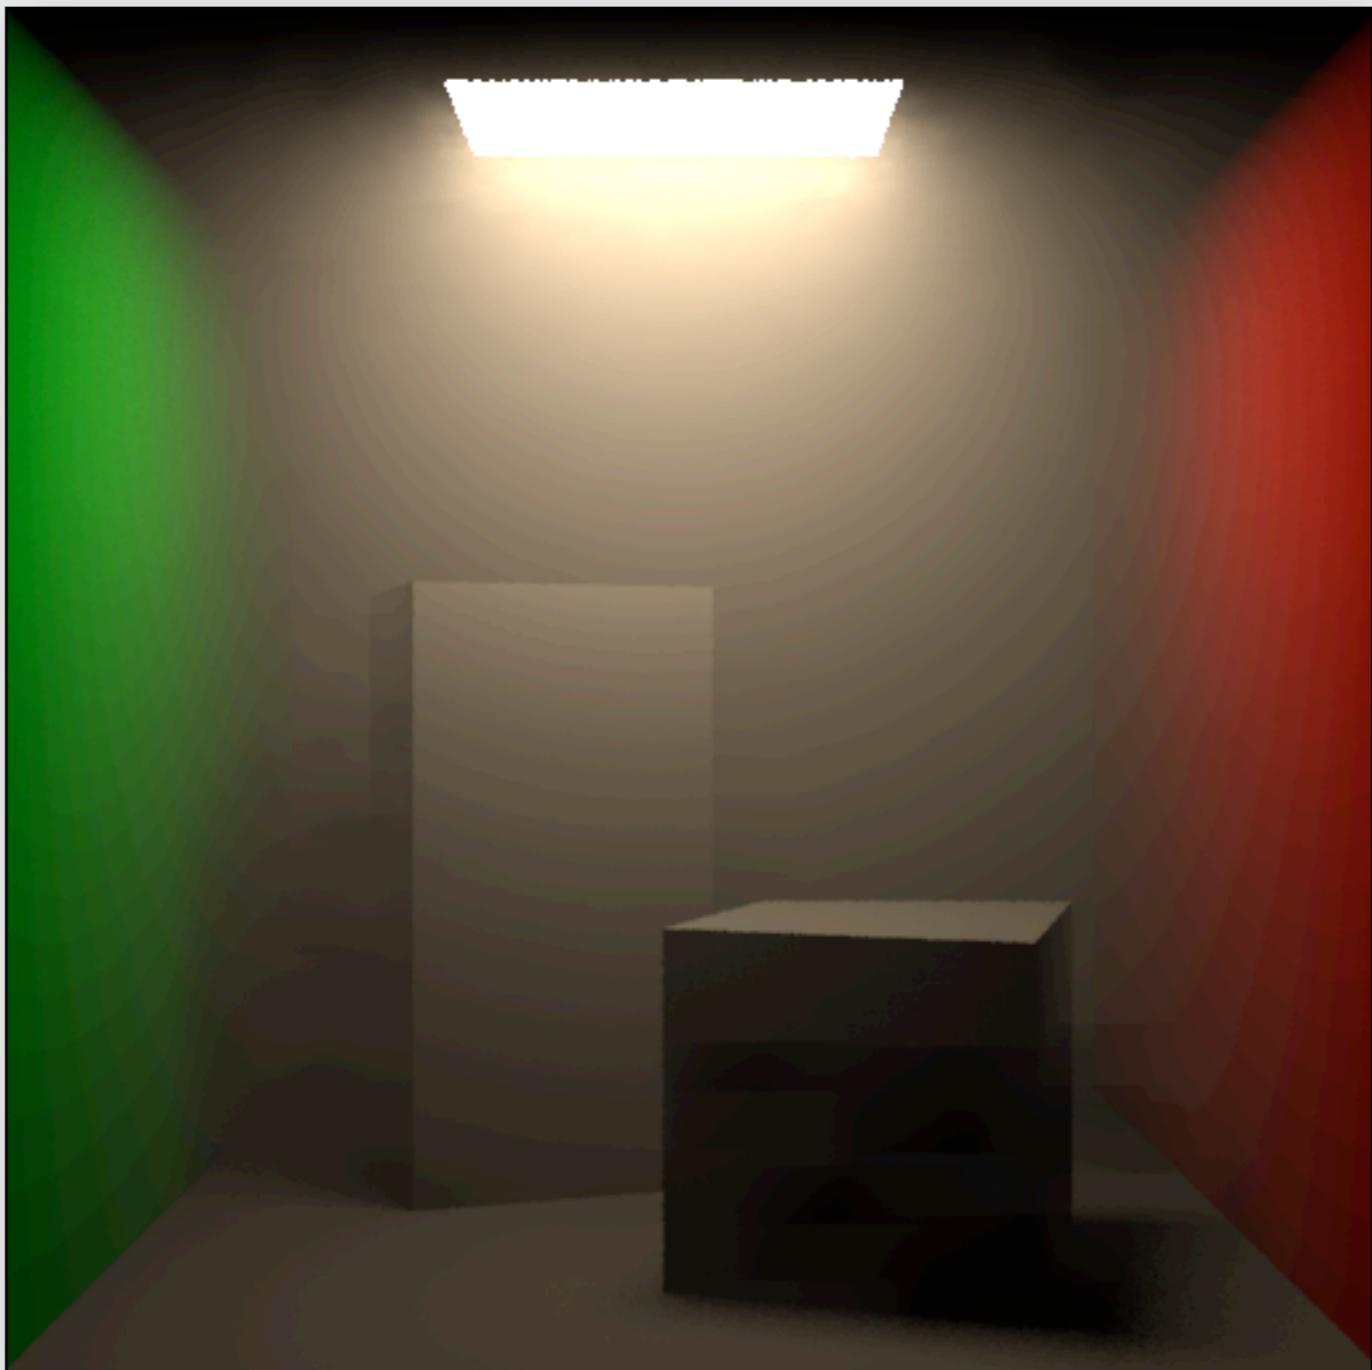


Eurographics 2015

The 36th Annual Conference of the
European Association for Computer Graphics

[Jarosz et al. 2008a] ⁷⁰

Gradients



with gradients

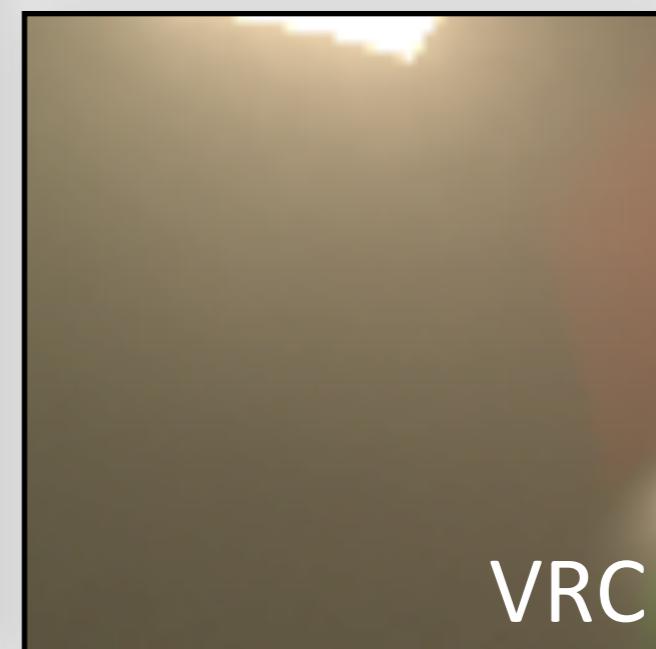
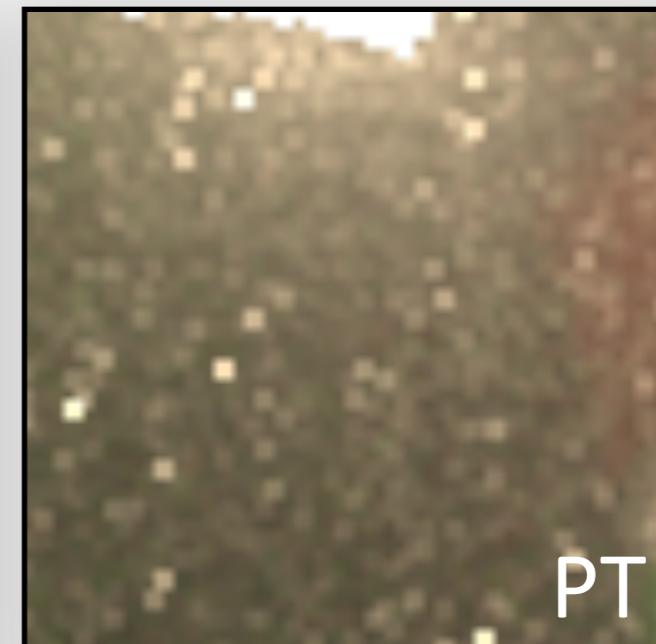
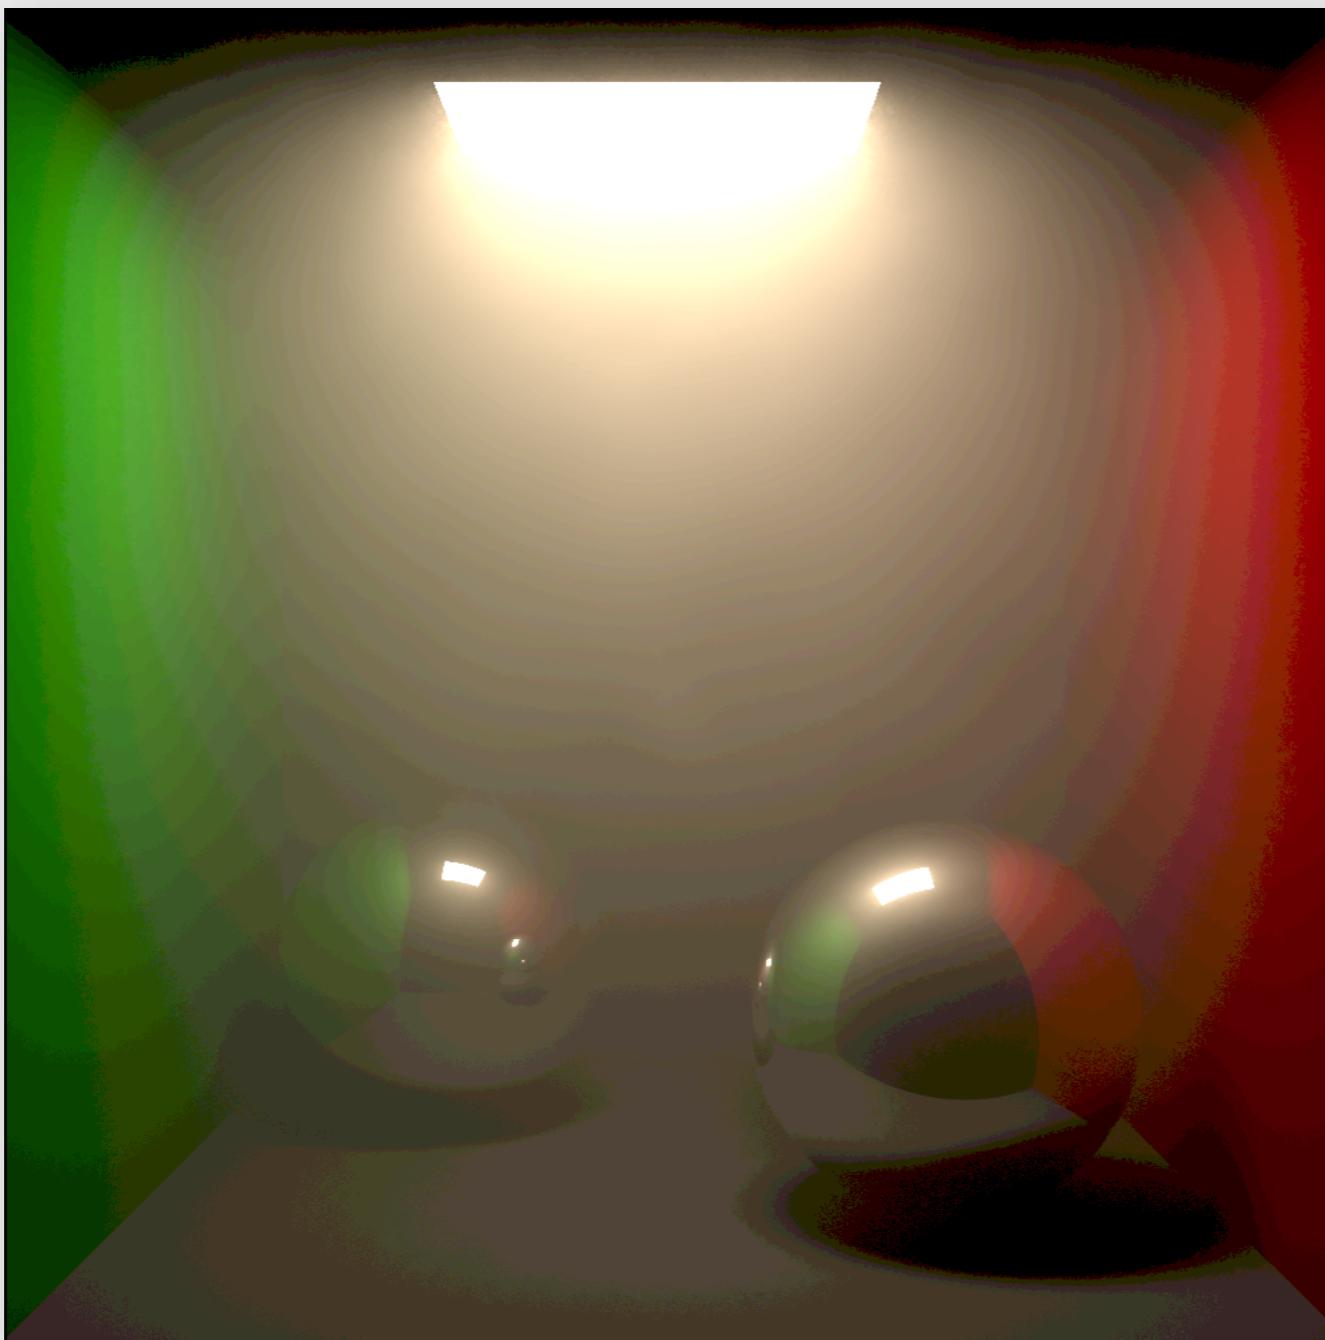


Eurographics 2015

The 36th Annual Conference of the
European Association for Computer Graphics

[Jarosz et al. 2008a] ⁷¹

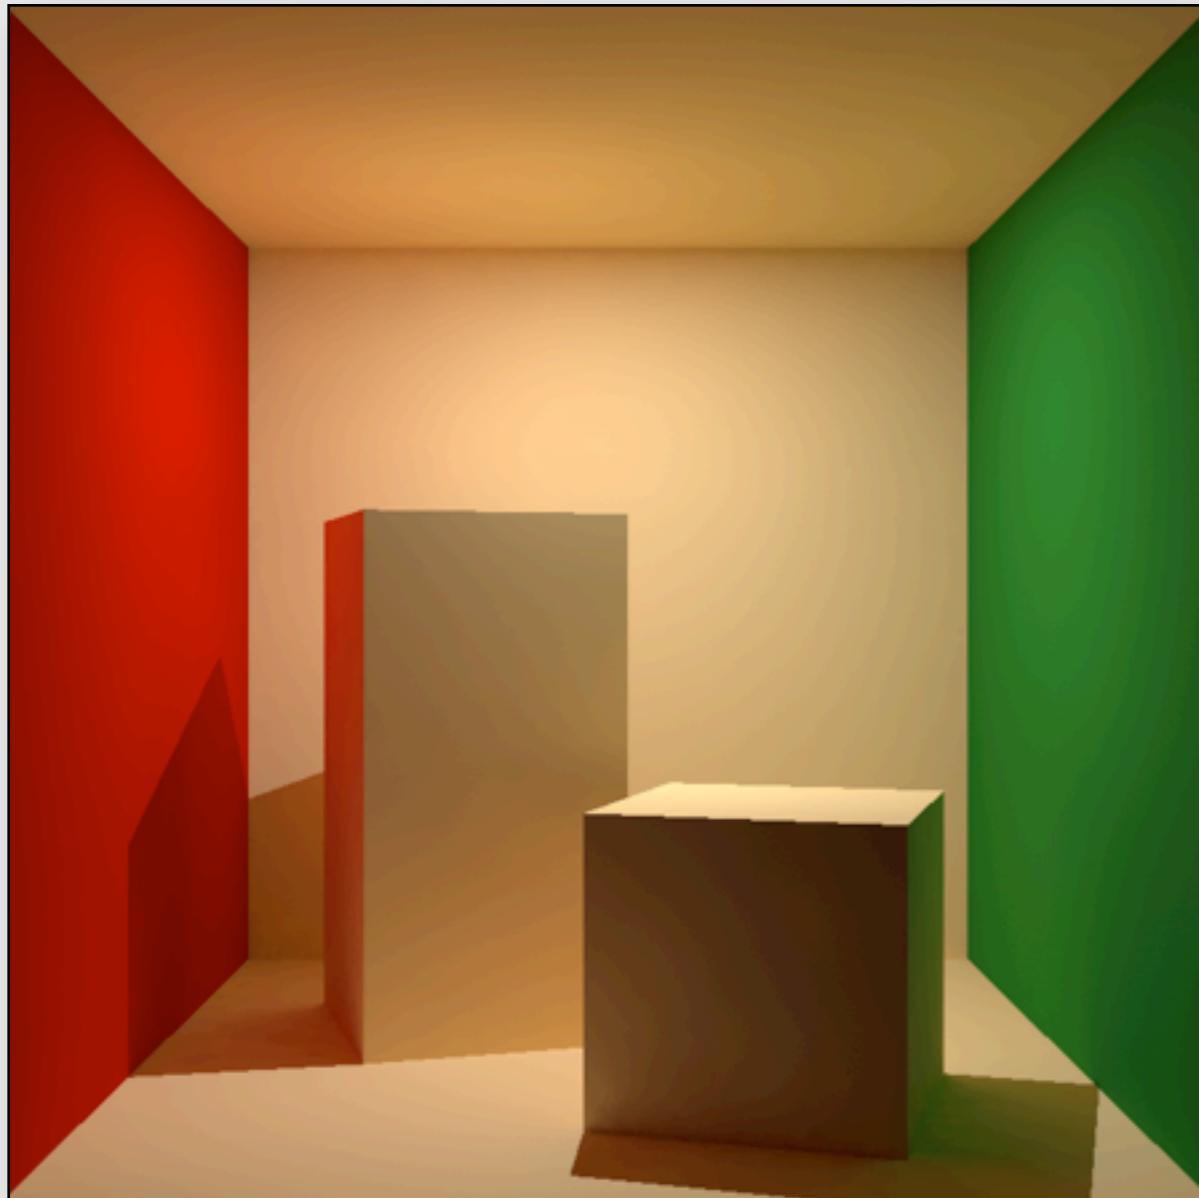
Results



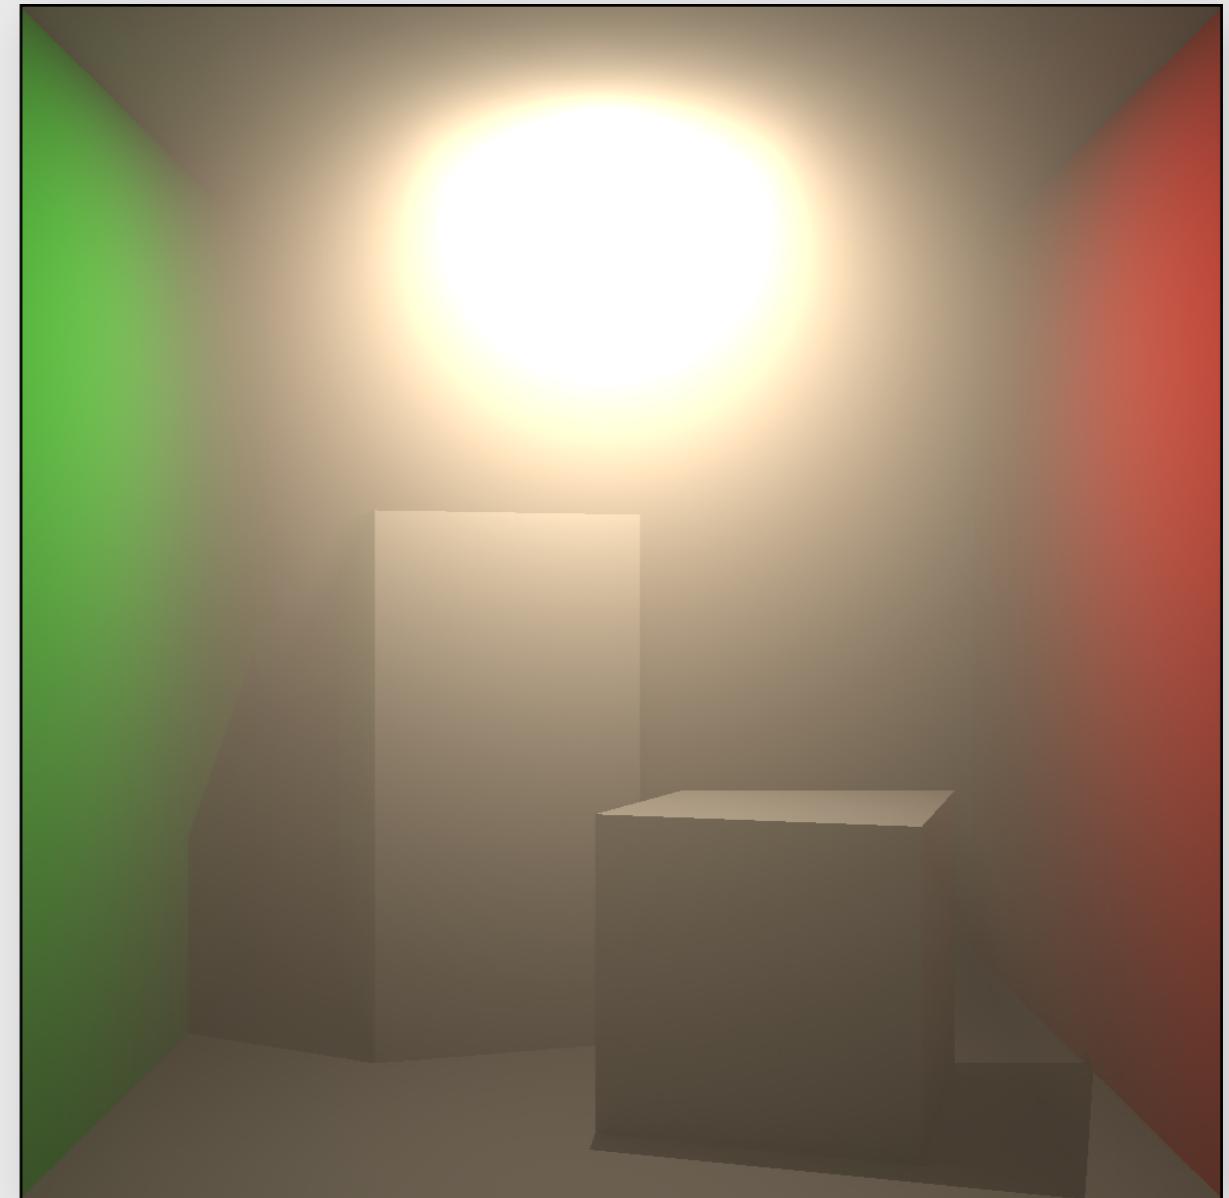
Results



Participating media



no media

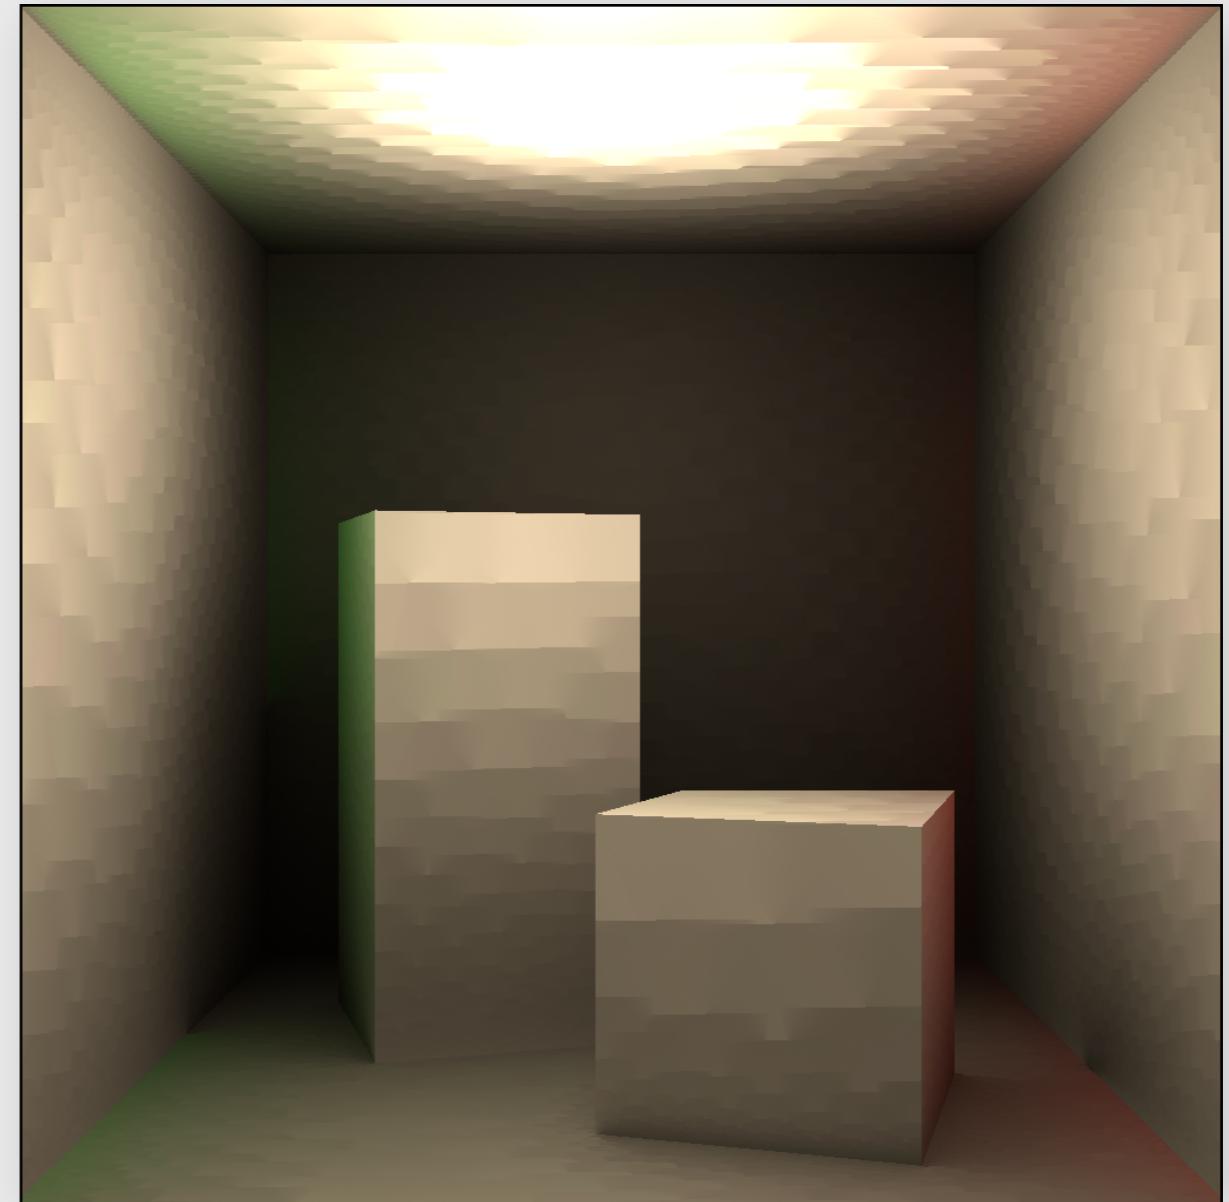


with media

Surfaces in participating media

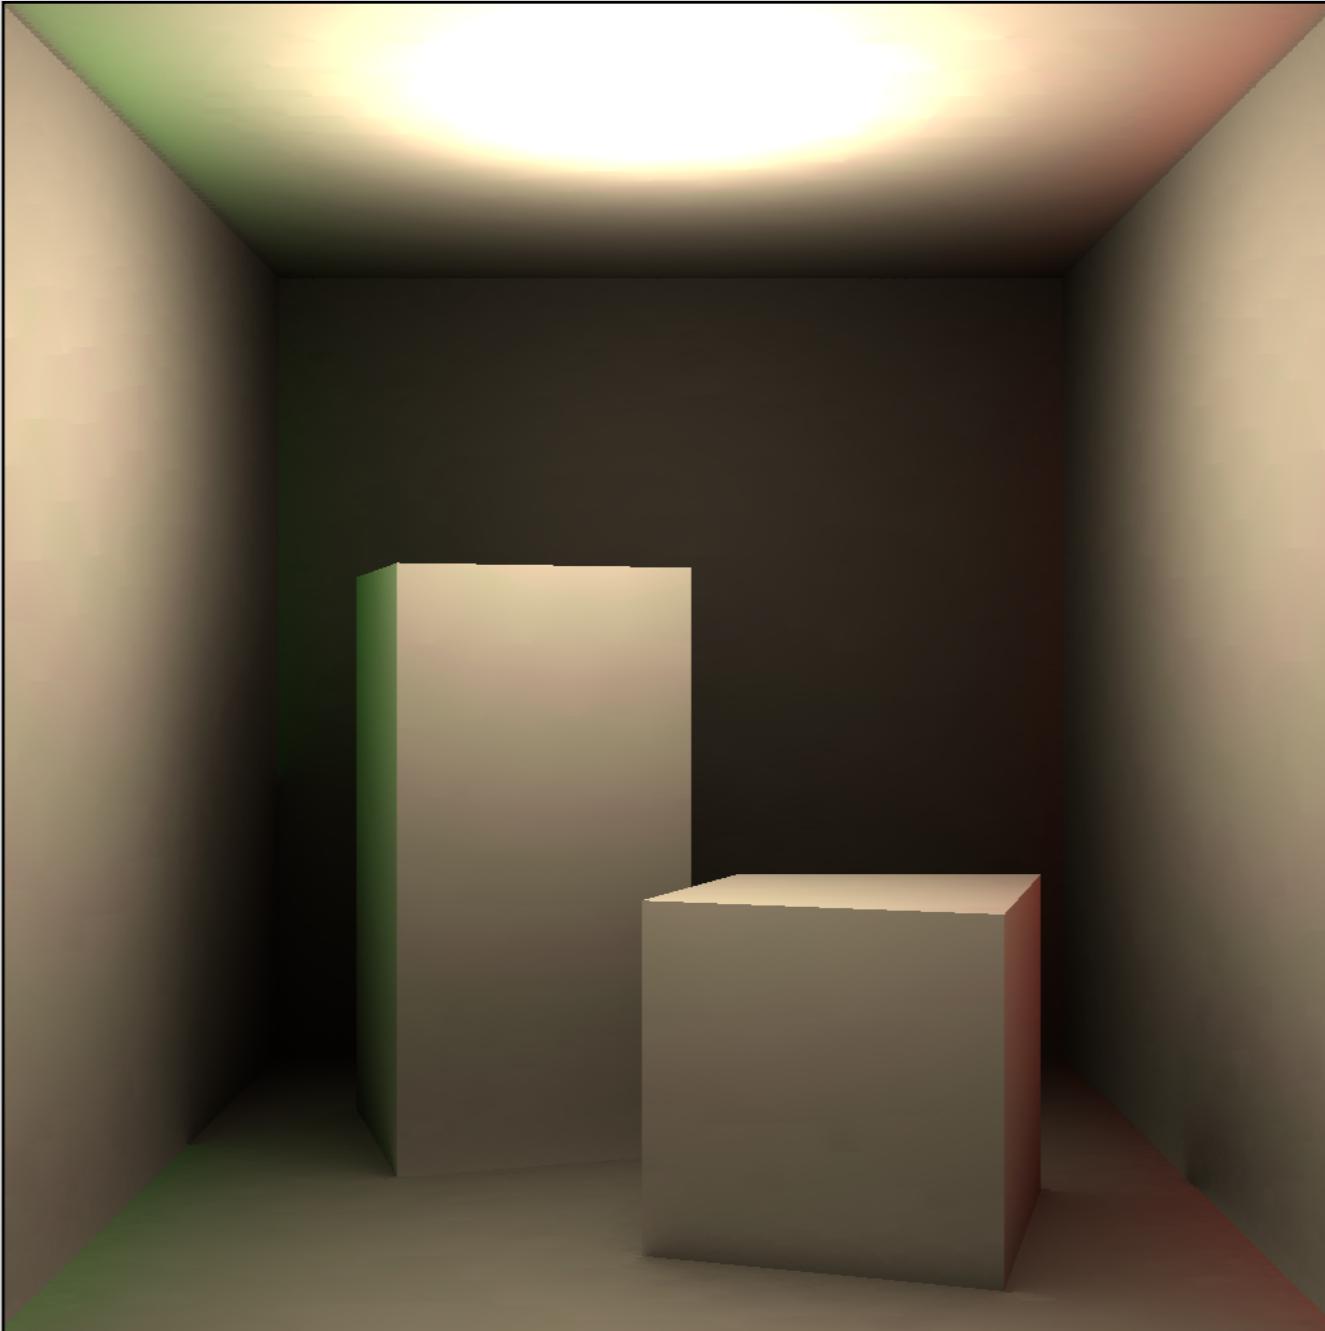


no media
(indirect irradiance)



with media
(indirect irradiance)

Surfaces in participating media



Occlusion, transmission, and the media with no gradients
[Wojciech et al. 2008b]

Sun beam through window



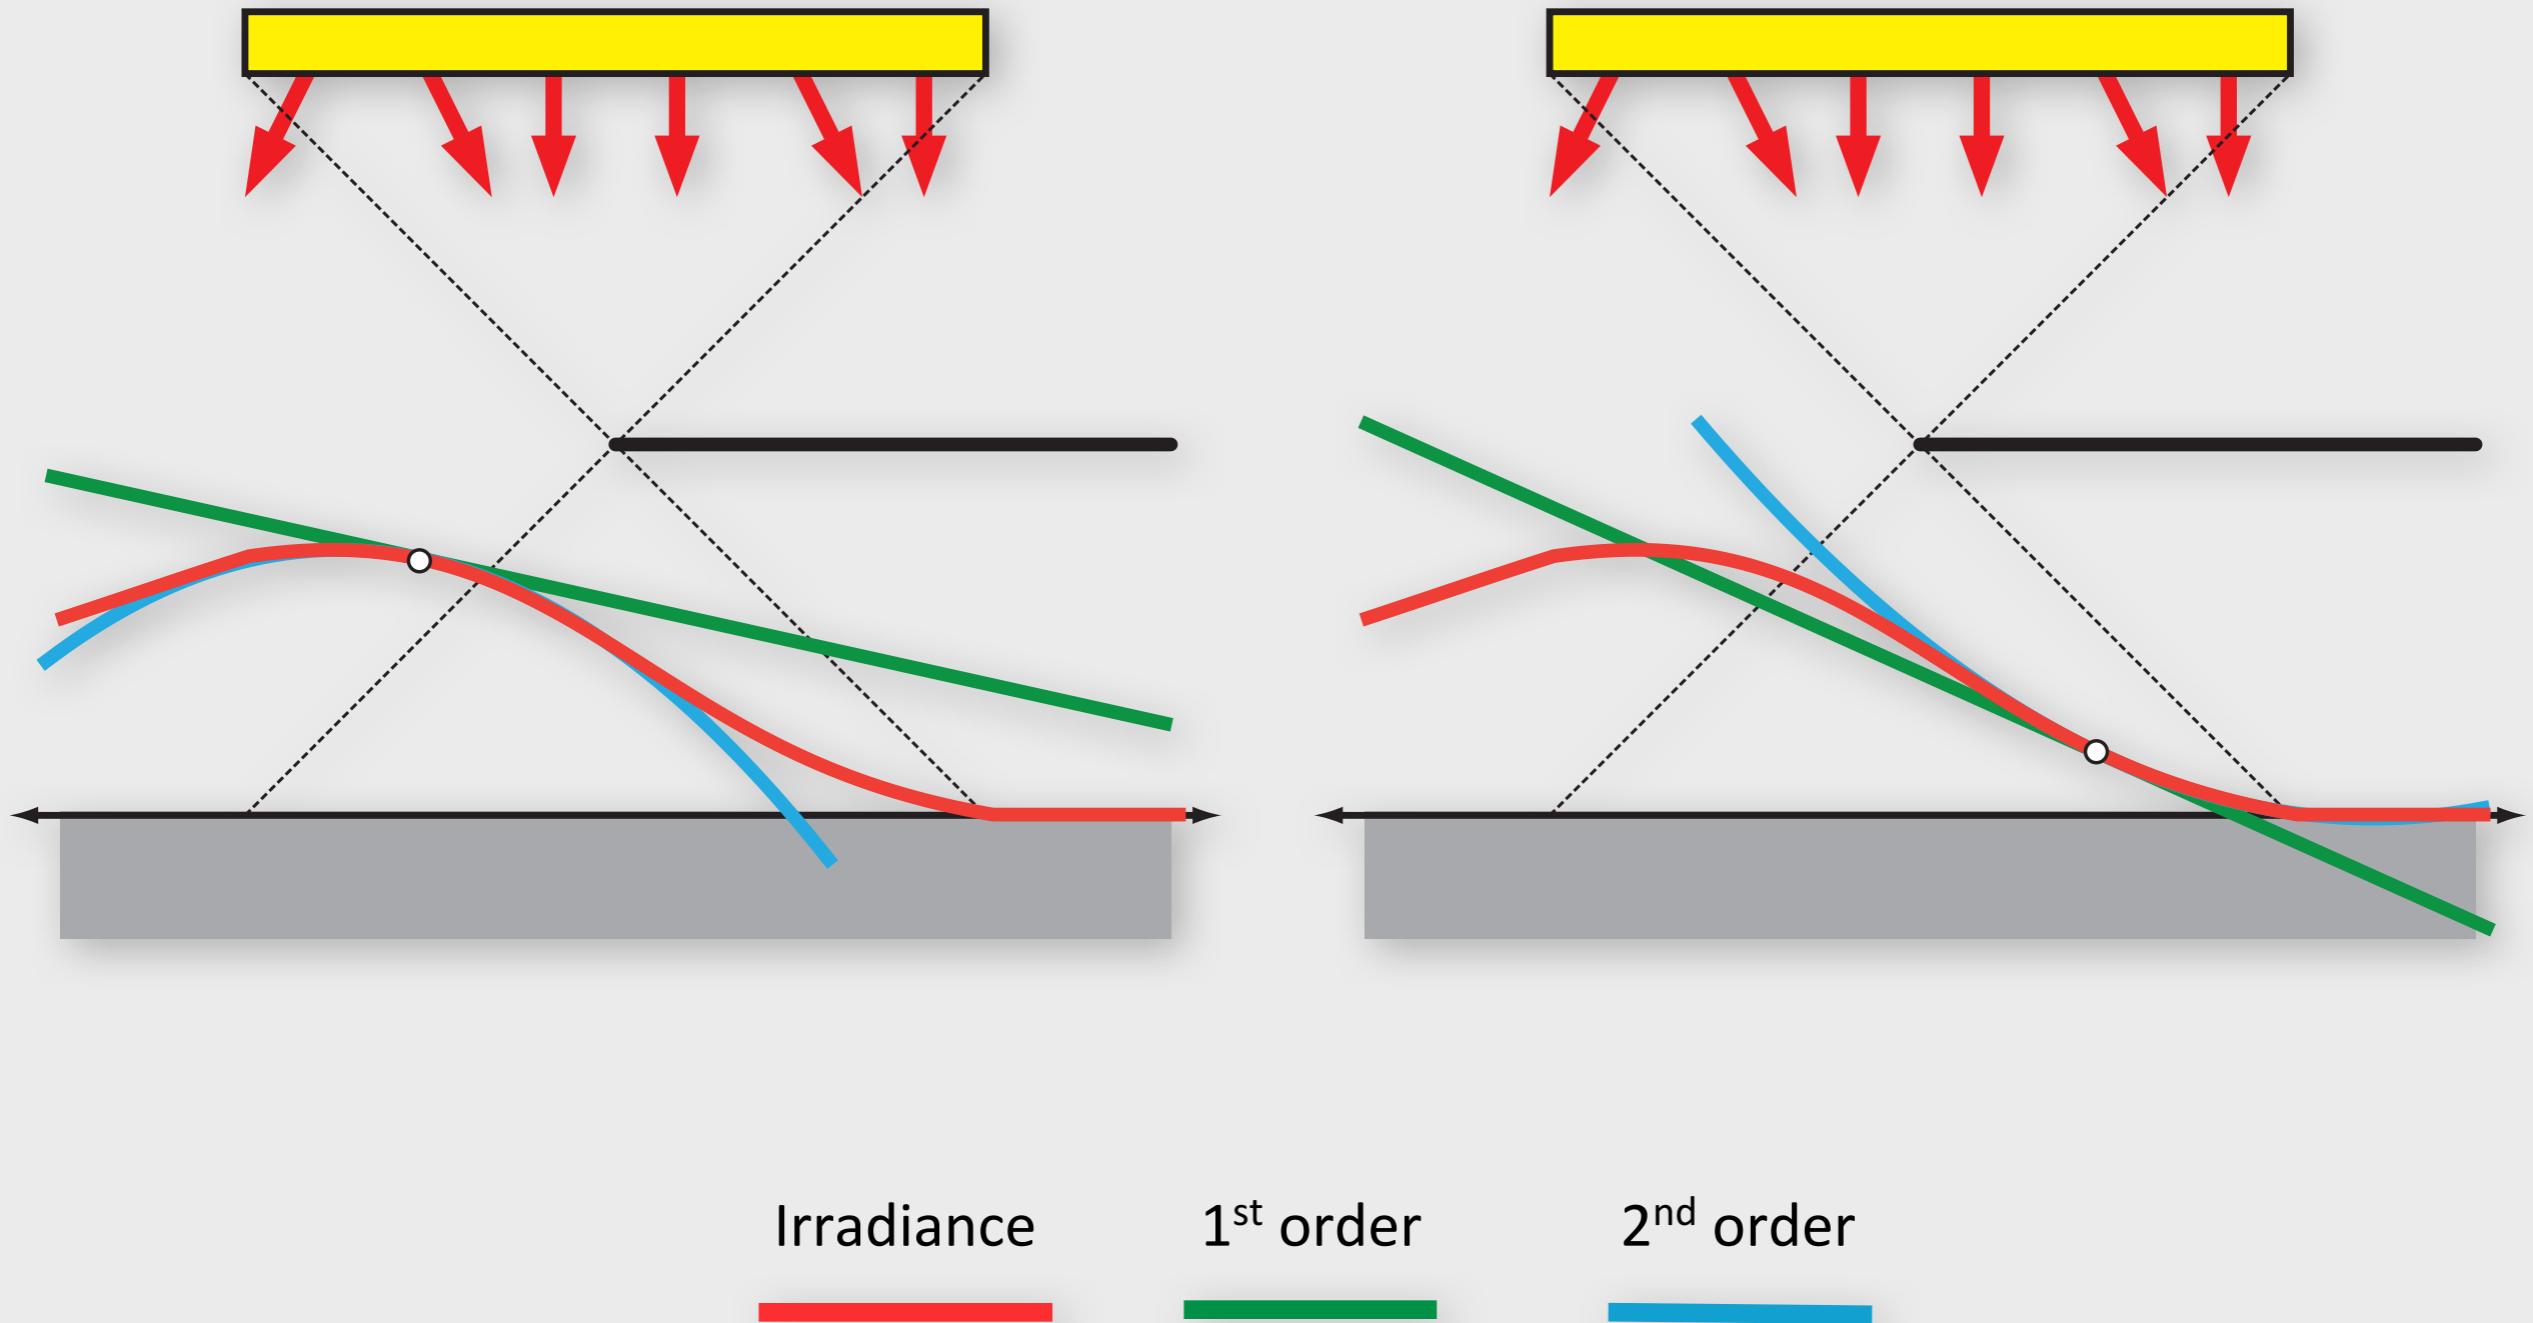
Gradien~~t~~ts by [Ward and Heckbert 1992]



Higher-order derivatives

- Exploit higher-order derivatives for better error control
 - [Jarosz et al. 2012] - Hessians (occlusion-unaware)
 - [Schwarzhaupt et al. 2012] - occlusion-aware Hessians & practical details

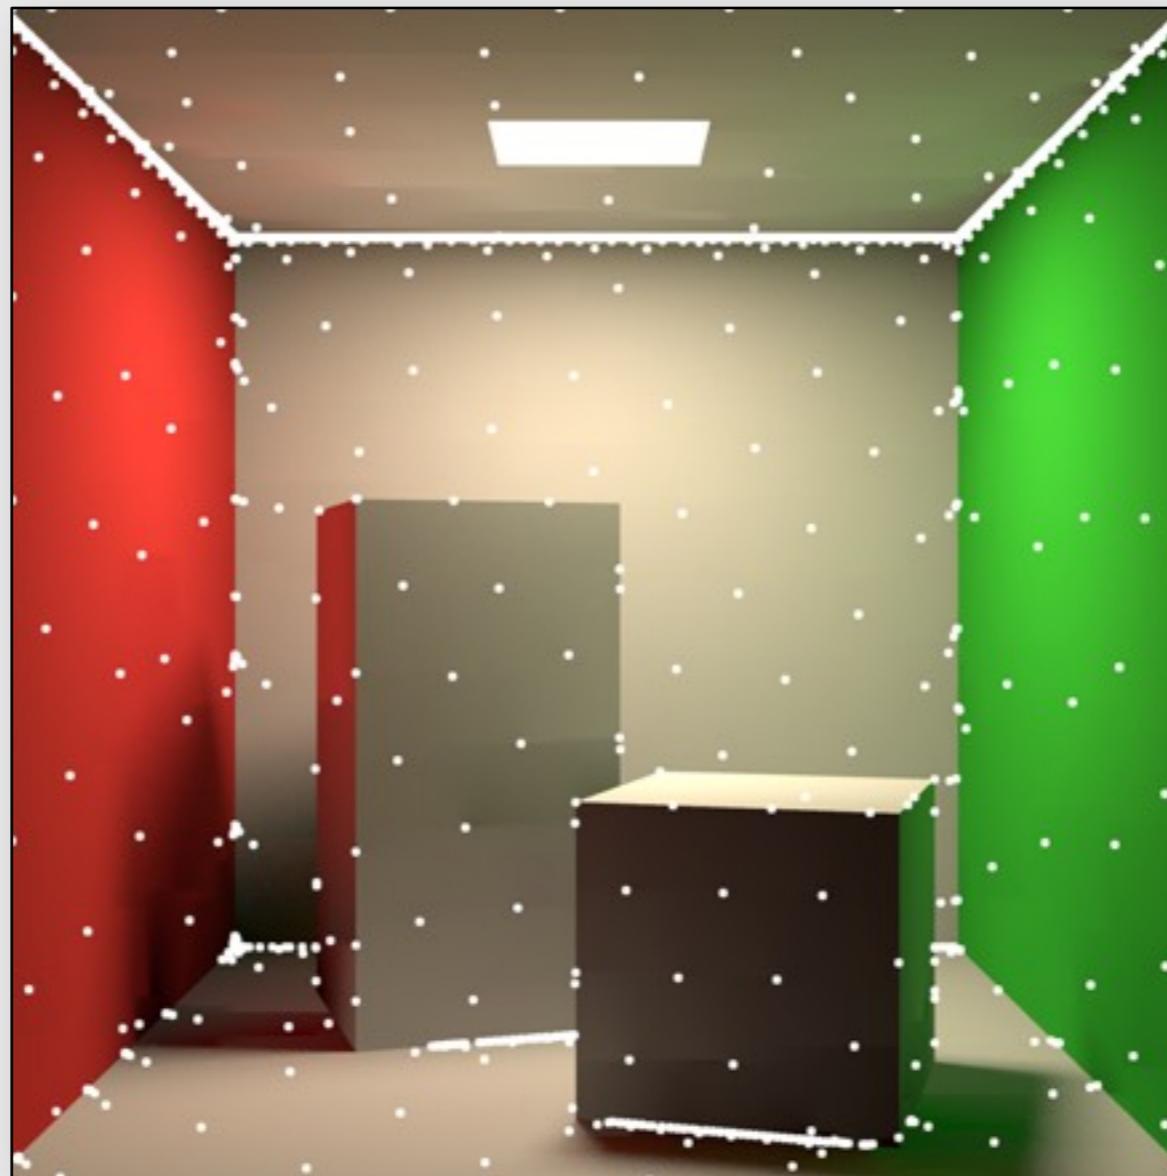
1st & 2nd Order Extrapolation



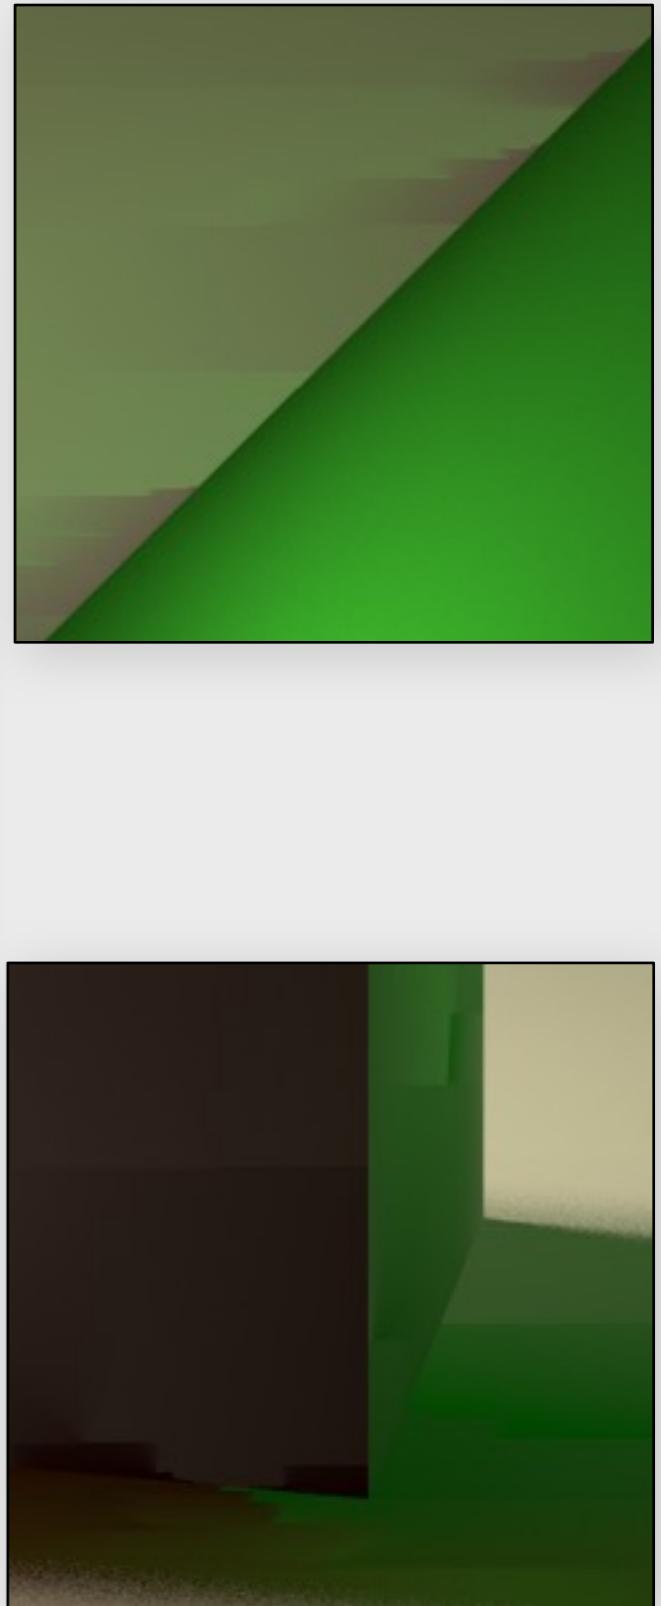
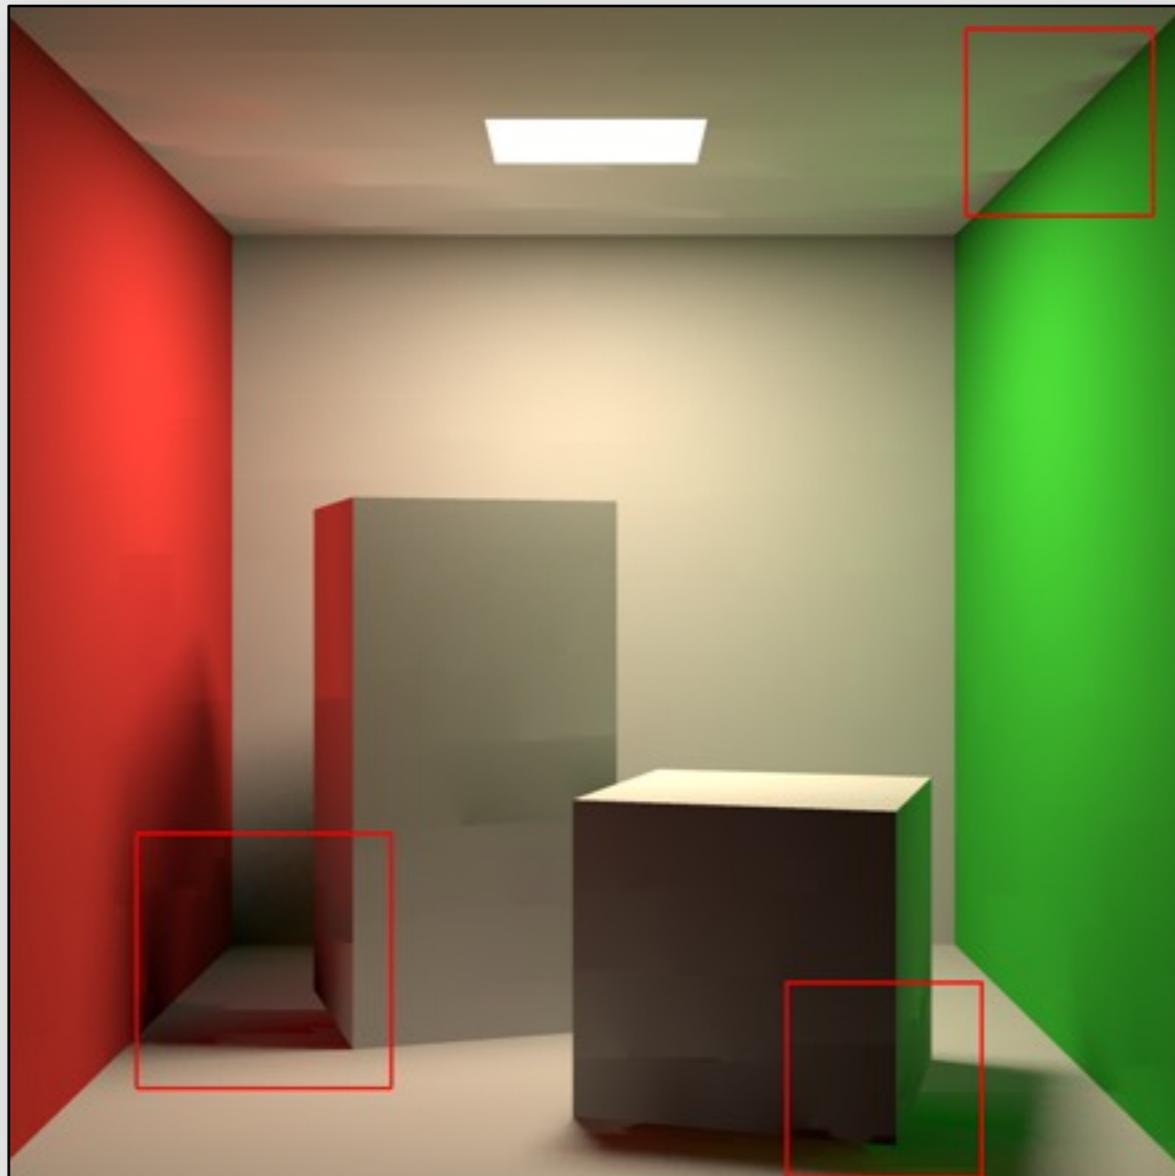
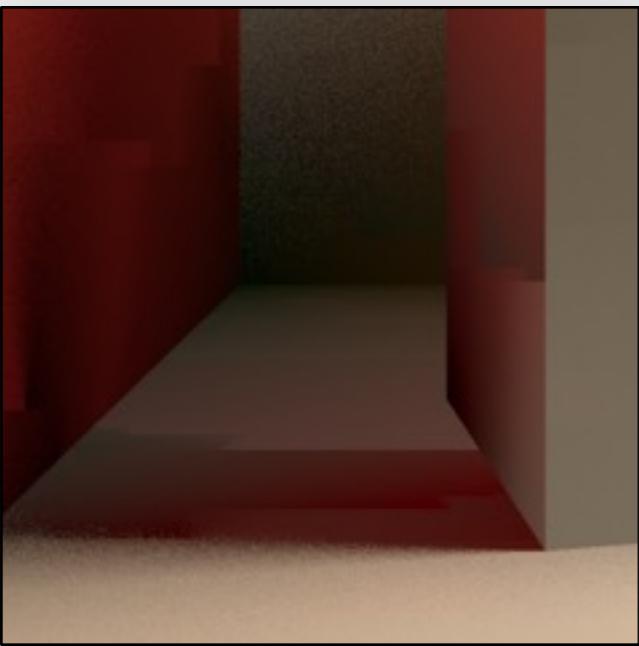
Split-Sphere Heuristic

- Basis for most irradiance caching algorithms for 20+ years
- Fix-ups to original metric lead to many parameters
 - error threshold
 - min/max screen-space radii
 - min/max world-space radii
 - gradient clamping
 - ...
- Hard to control!

The Split Sphere

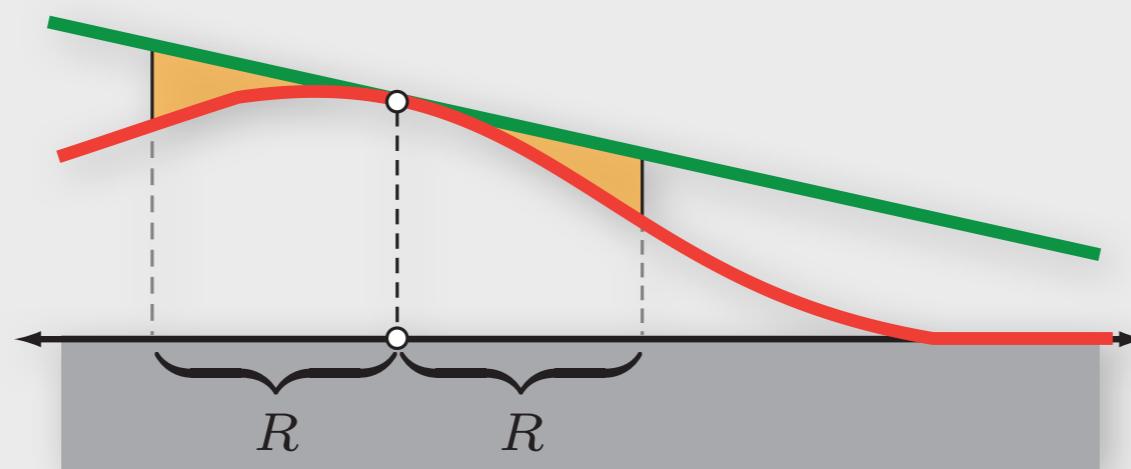


The Split Sphere



Better Error Control

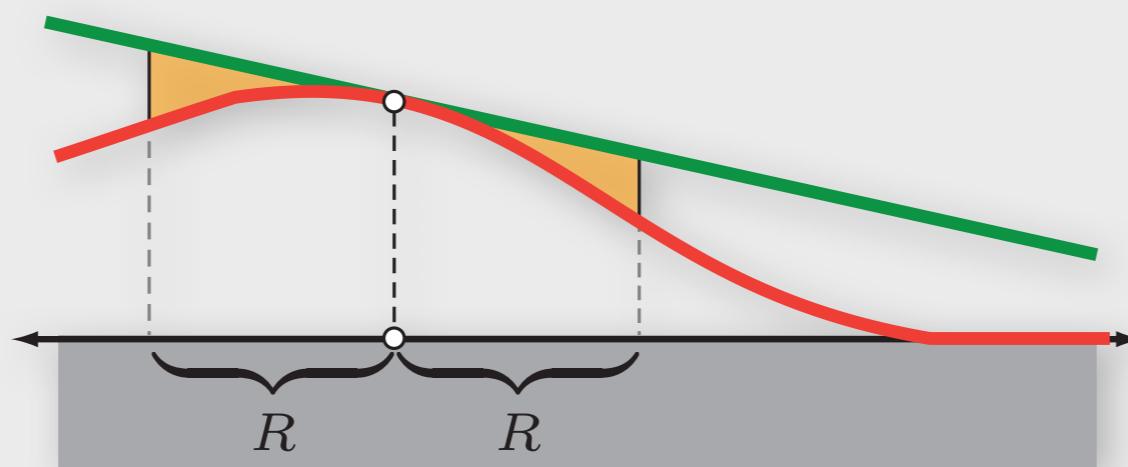
- **total error** ϵ^t = integrated **difference** between **extrapolated** and **correct** irradiance



$$\epsilon^t = \int_{-R_i}^{R_i} |E(\mathbf{x}_i + x) - E'(\mathbf{x}_i + x)| dx$$

Better Error Control

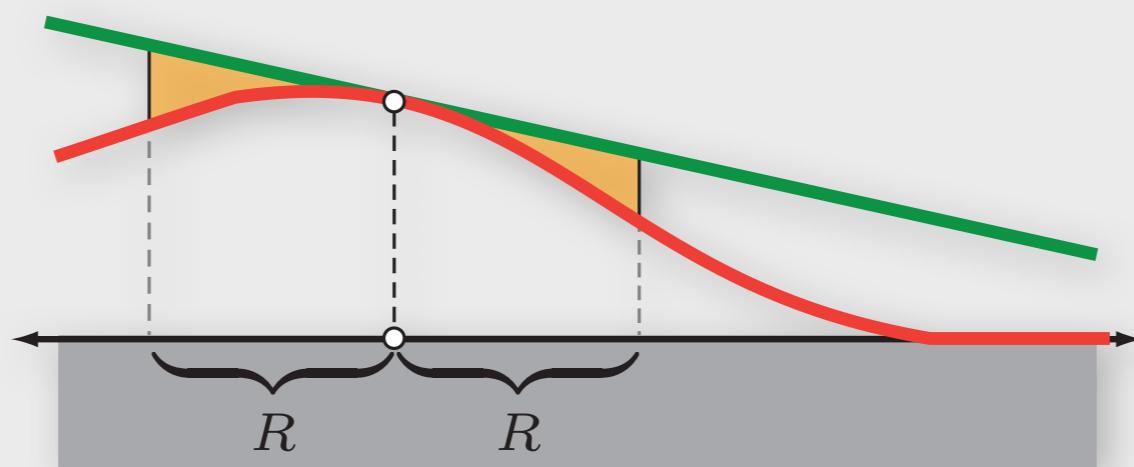
- E' is 1st-order Taylor extrapolation
- E is unknown!



$$\epsilon^t = \int_{-R_i}^{R_i} |E(\mathbf{x}_i + x) - E'(\mathbf{x}_i + x)| dx$$

Better Error Control

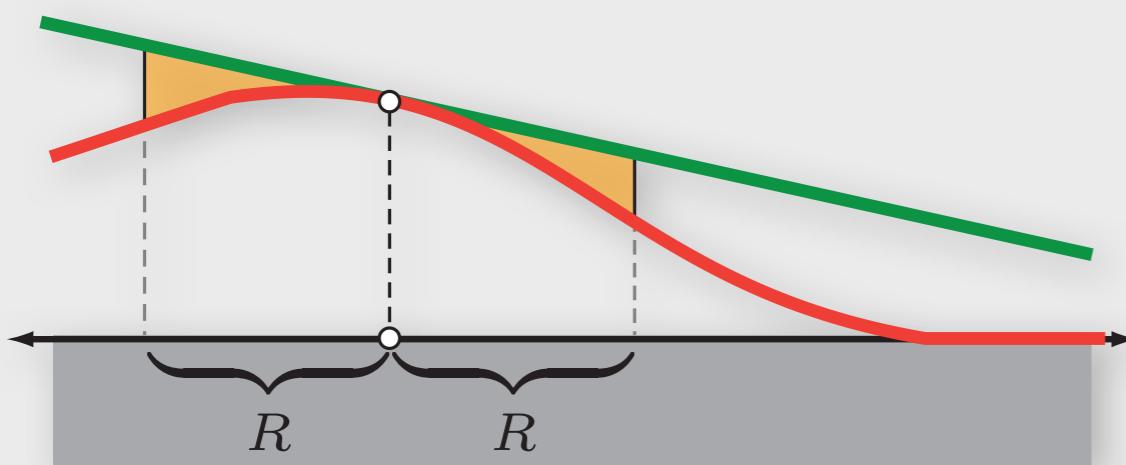
- E' is 1st-order Taylor extrapolation
- E is unknown!



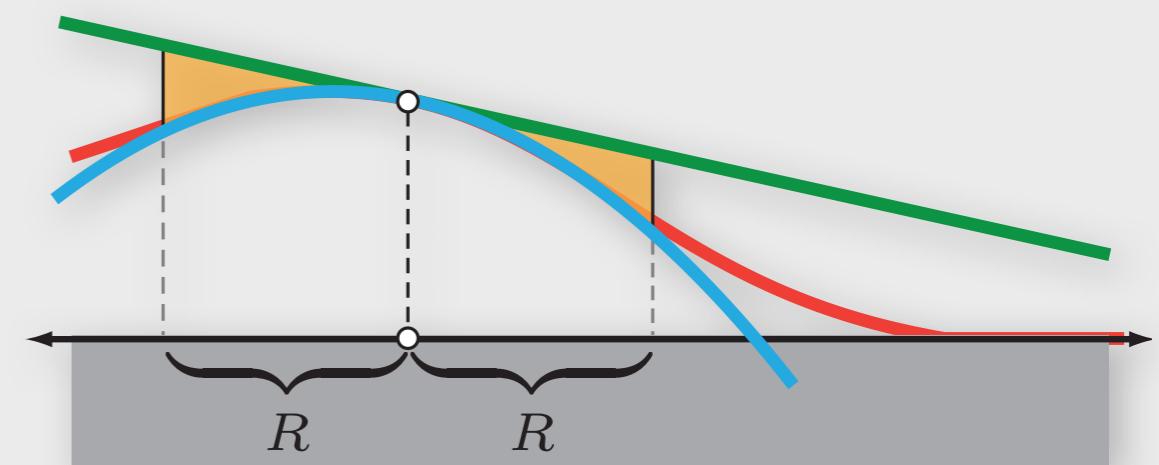
$$\epsilon^t = \int_{-R_i}^{R_i} |E(\mathbf{x}_i + x) - E'(\mathbf{x}_i + x)| dx$$

Better Error Control

- E' is 1st-order Taylor extrapolation
- E is unknown!



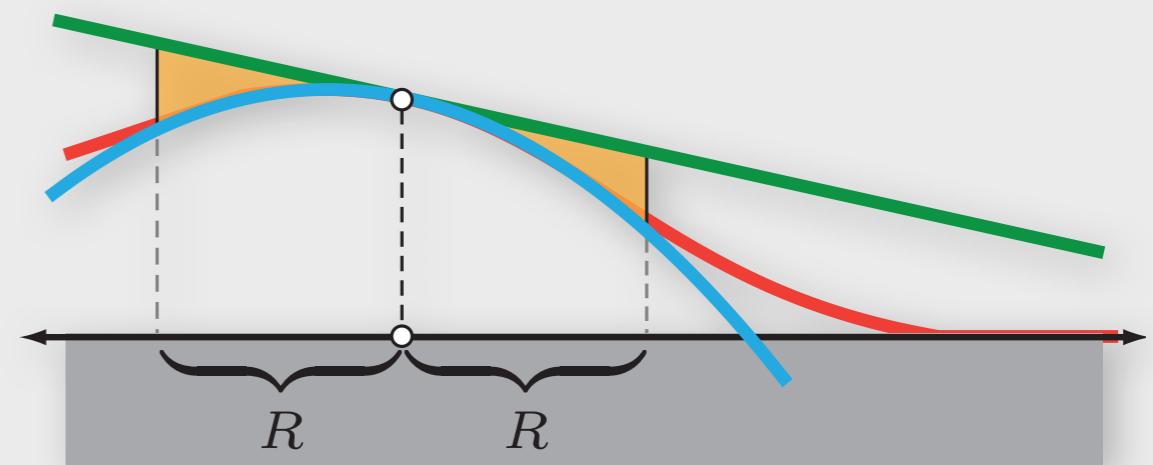
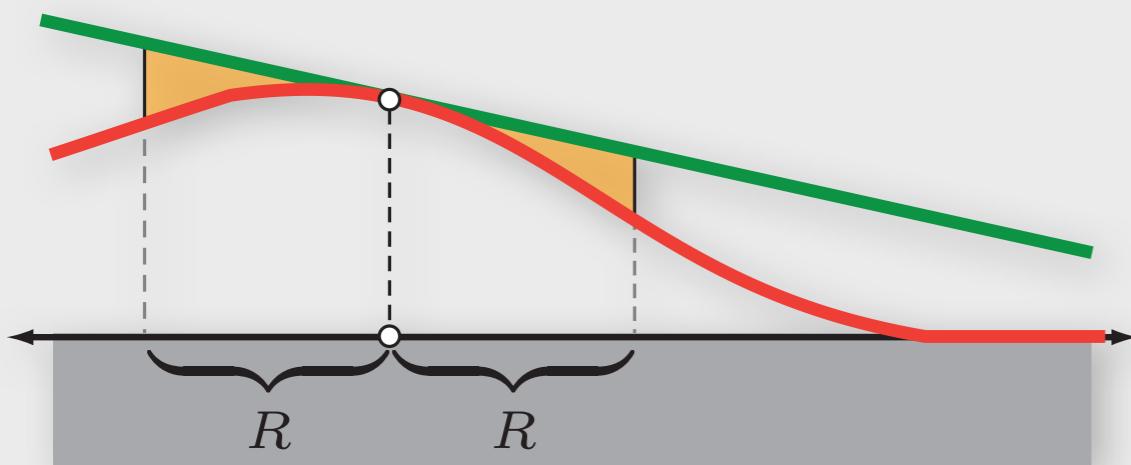
2nd-order Taylor extrapolation



$$\epsilon^t = \int_{-R_i}^{R_i} |E(\mathbf{x}_i + x) - E'(\mathbf{x}_i + x)| dx$$

Hessian-based Error Control

- E' is 1st-order Taylor extrapolation
- 2nd-order Taylor extrapolation approximates E



$$\epsilon^t = \int_{-R_i}^{R_i} |E(\mathbf{x}_i + x) - E'(\mathbf{x}_i + x)| dx \approx \hat{\epsilon}^t = \frac{1}{2} \int_{-R_i}^{R_i} |x \mathbf{H}_{\mathbf{x}}(E_i) x| dx$$

Hessian-based Error Control

$$\hat{\epsilon}^t = \frac{1}{2} \int_{-R_i}^{R_i} |x \mathbf{H}_{\mathbf{x}}(E_i) x| dx = \frac{1}{3} |h_{\mathbf{x}}(E_i)| R_i^3$$

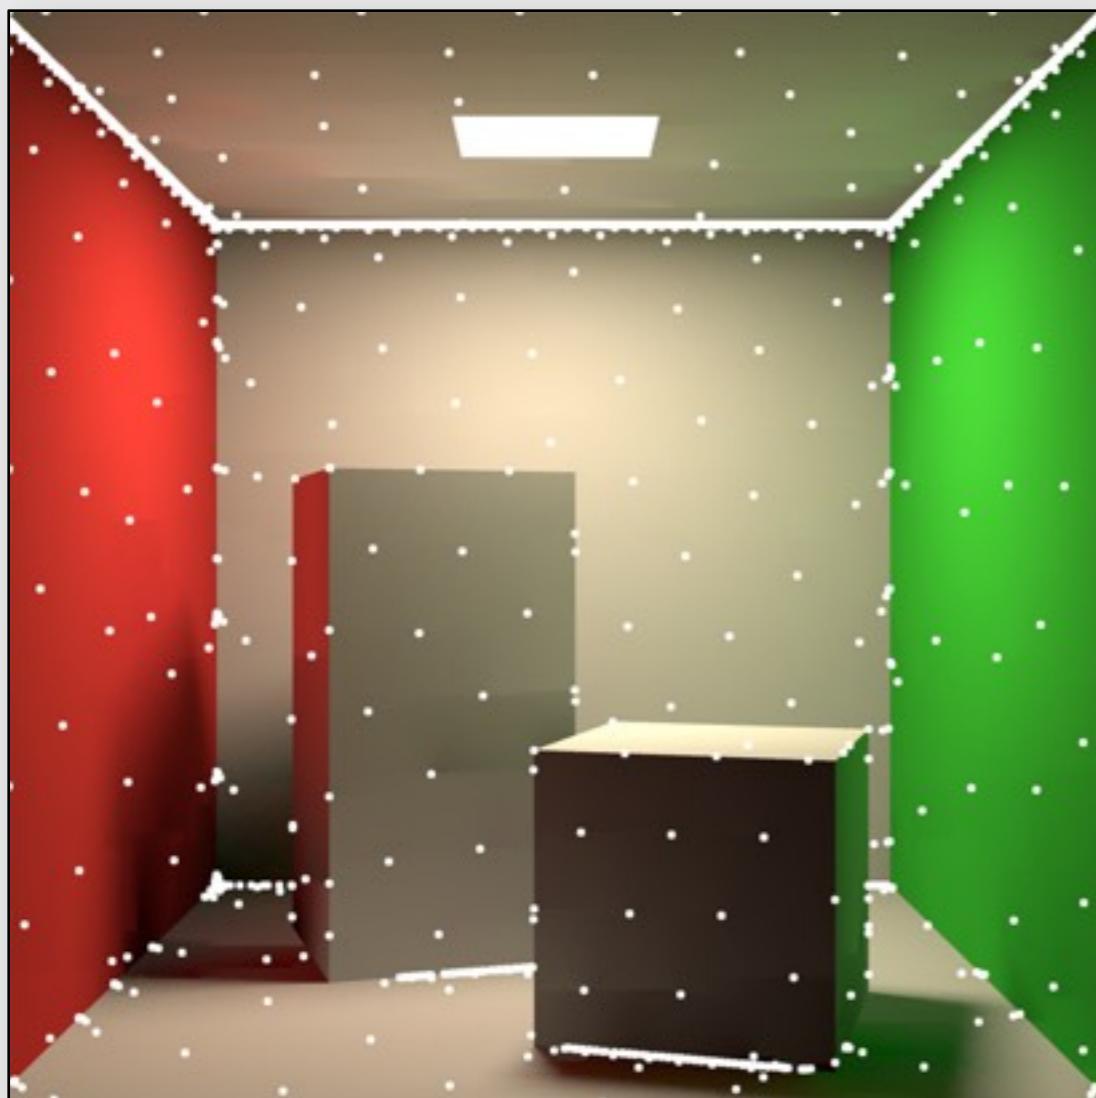
- fix $\hat{\epsilon}^t$, solve for R_i

second derivative

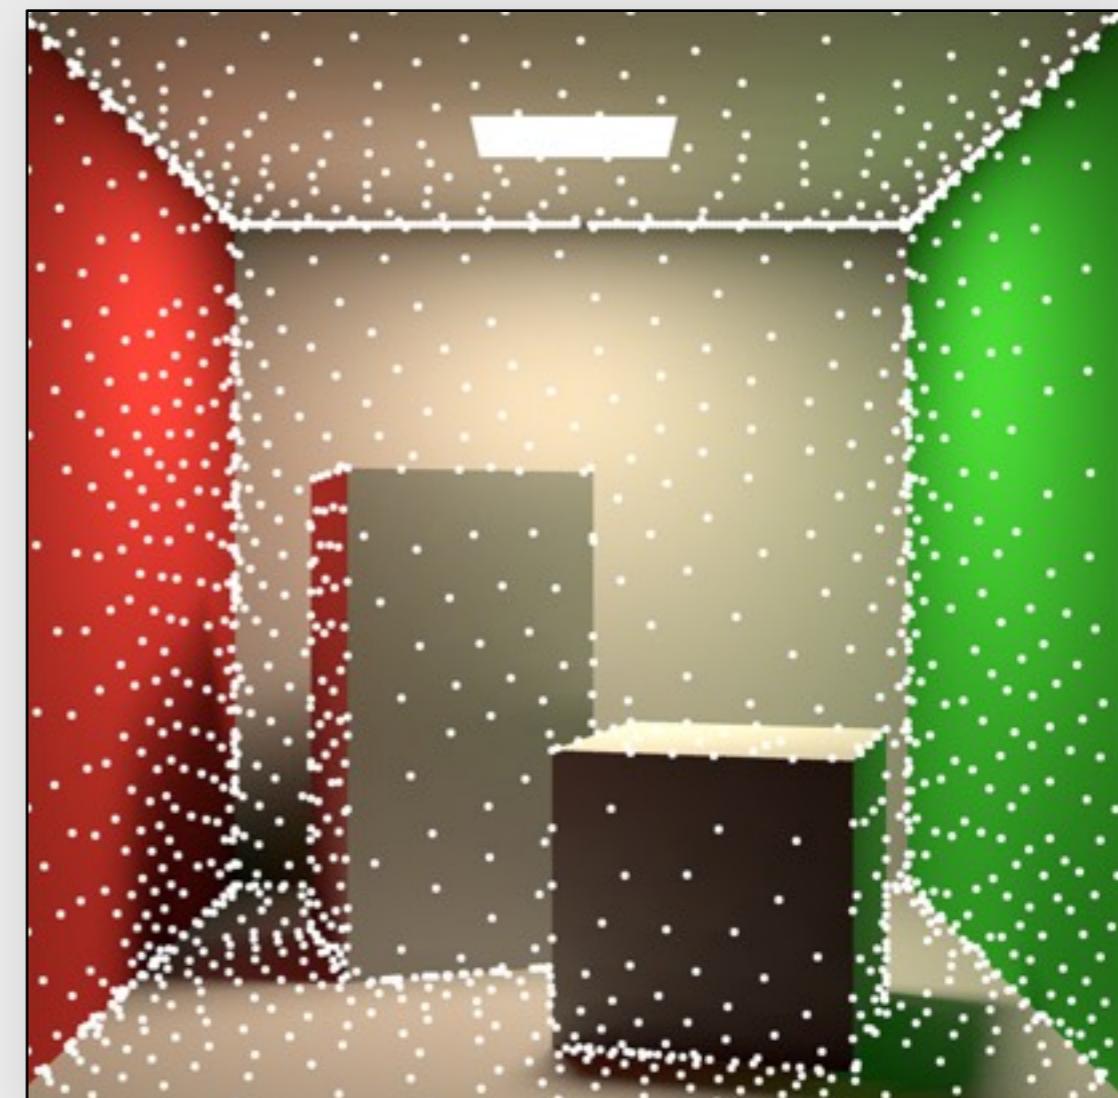
$$R_i = \sqrt[3]{\frac{3\hat{\epsilon}^t}{|h_{\mathbf{x}}(E_i)|}}$$

Beyond the Split-Sphere

~1,700 Cache Points



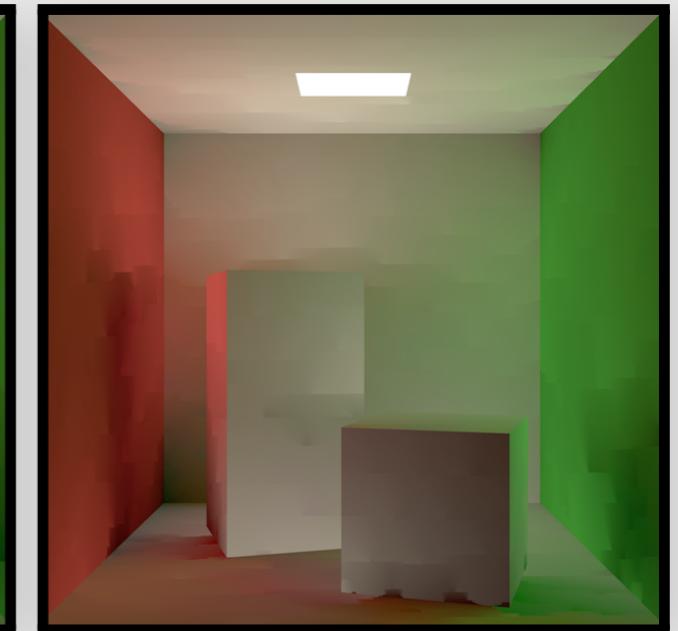
Split-Sphere



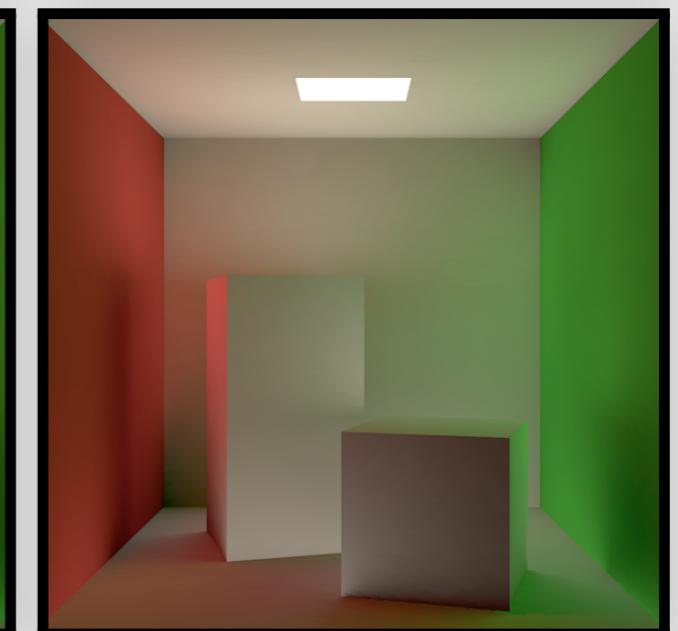
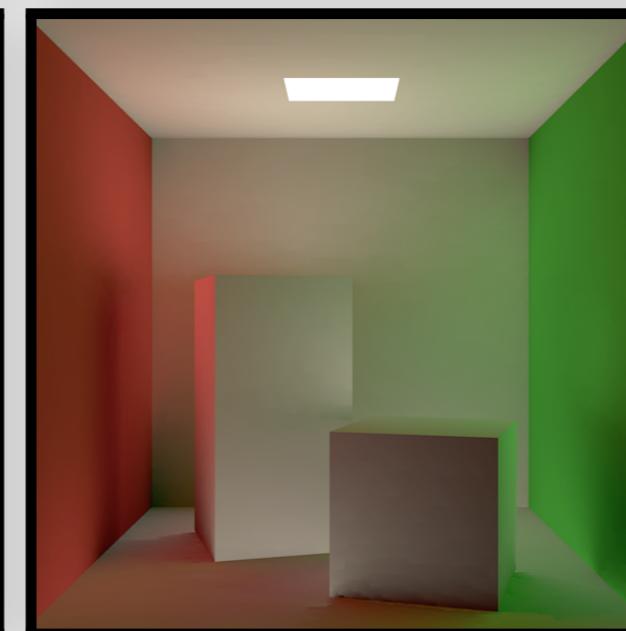
Hessian-based

Split-Sphere vs Hessian-based

split-sphere



Hessian-based



500 Records

1K Records

2K Records

4K Records

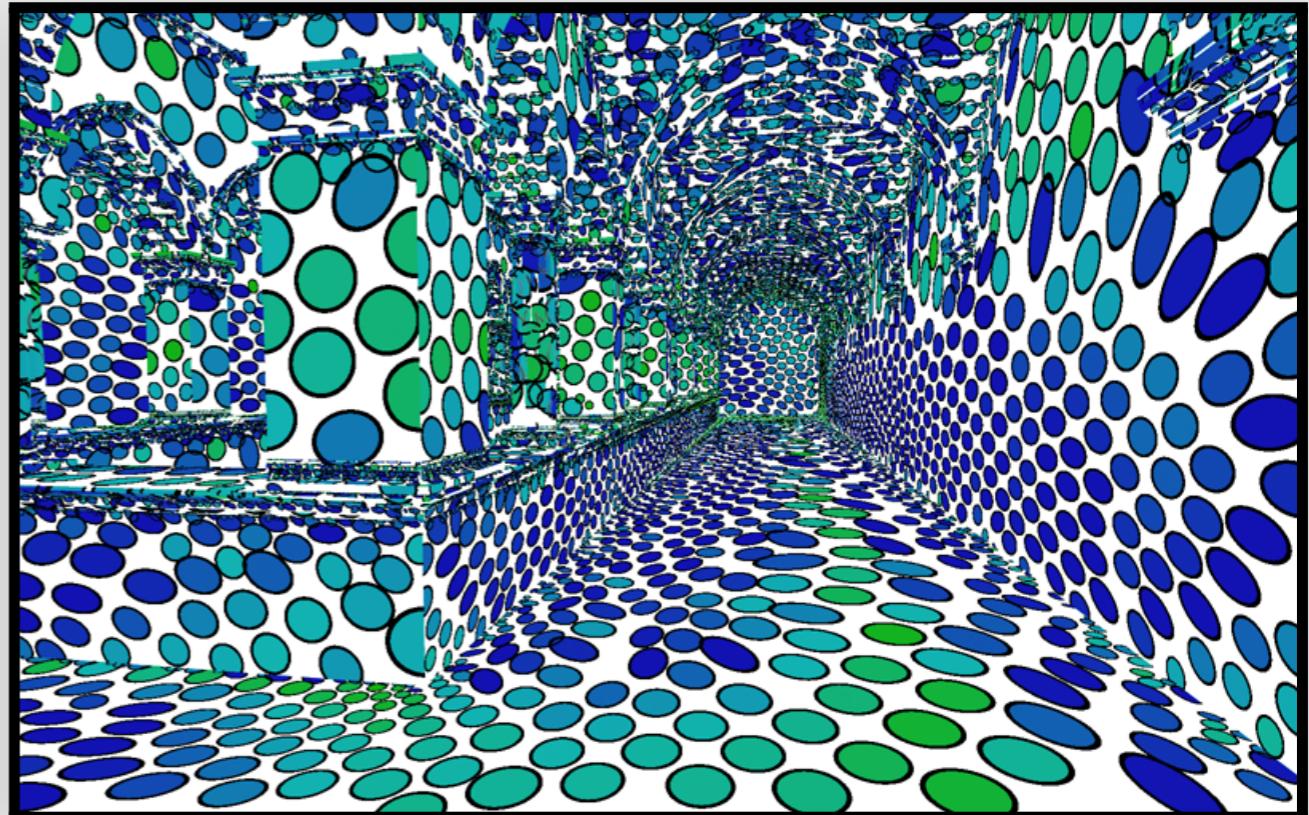
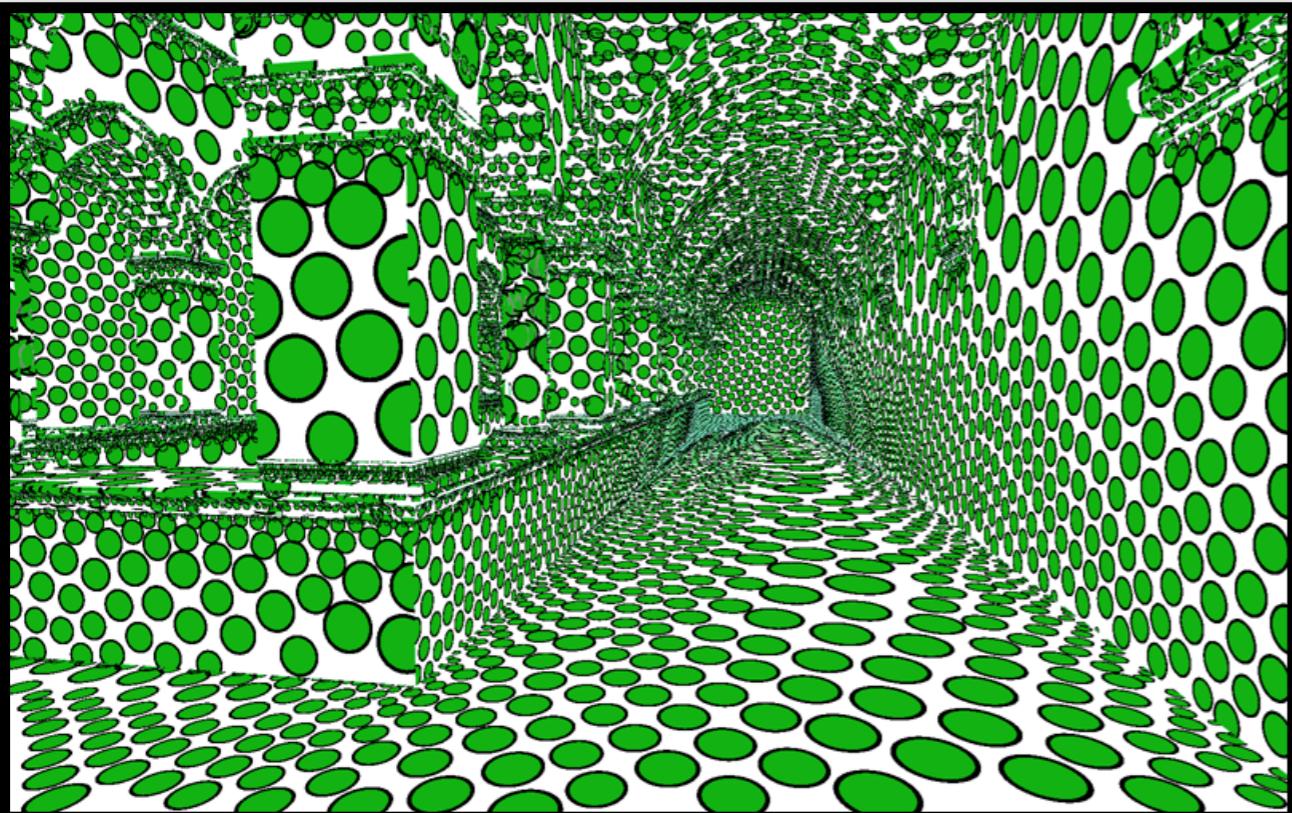


Eurographics 2015

The 36th Annual Conference of the
European Association for Computer Graphics

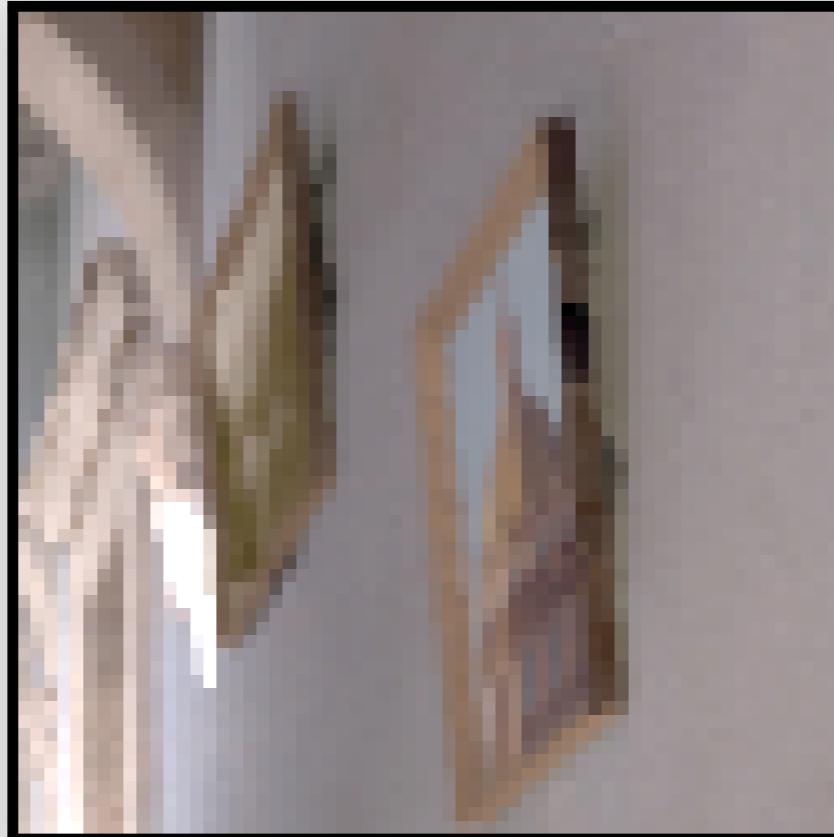
[Schwarzhaft et al. 2012] 90

Anisotropic Cache Records

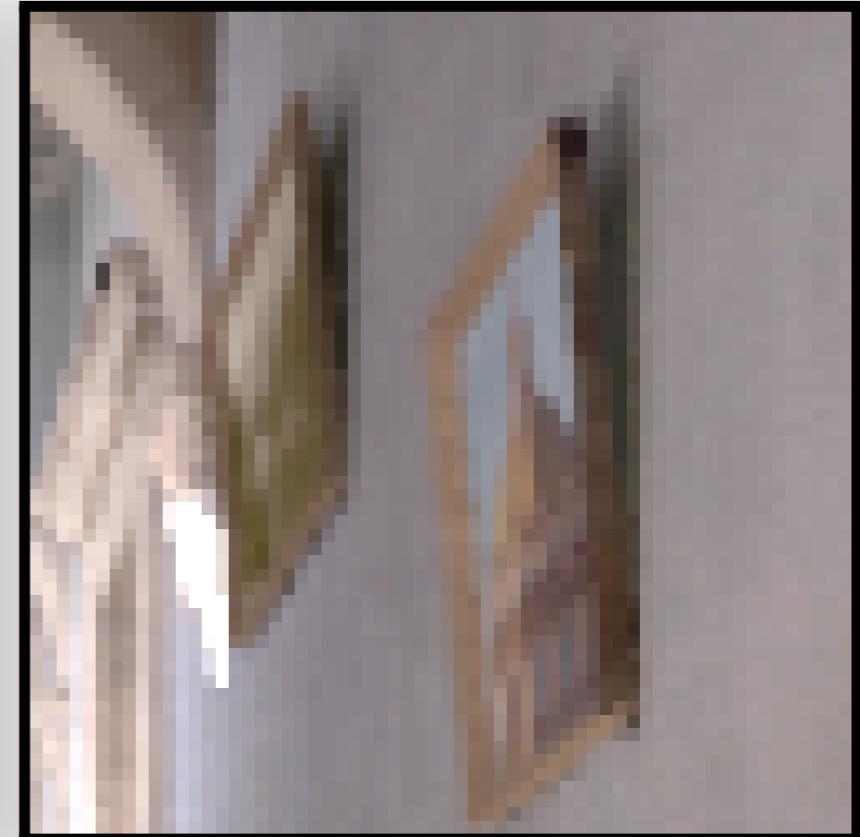




Reference



Bounded Split-Sphere



Occlusion Hessian

Hessians in RADIANCE

- Now the default approach used in RADIANCE

The screenshot shows the official website for Radiance, titled "RAD SITE | radiance-online.org". The top navigation bar includes links for "Log in", "Search", and "only in current section". Below the title, there is a main menu with tabs for "HOME" (highlighted in blue), "ABOUT", "DOWNLOAD / INSTALL", "LEARN", and "COMMUNITY". A "Quick Links" sidebar on the left provides links to various resources: "2016 Radiance Workshop, August 29-31, Padua, Italy", "Subscribe to Mailing Lists", "Search Mailing List Archives", "Official 5.0 Release", "Download Head Release", "Photon Map Manual", "Radiance Tutorials", "refman.pdf (Radiance model syntax)", and "Rendering with Radiance". The main content area features a large image of a modern office interior with rows of cubicles and large windows overlooking a city skyline, with the text "Radiance - A Validated Lighting Simulation Tool" overlaid.

Summary

- Derivatives can estimate local function smoothness
- Amortize illumination computation across many pixels
- Accounting for occlusions is challenging but critical
- Specialized techniques only for diffuse or moderately glossy